



Iter 131 -- G=4 vs G=32 vs the Broader Sweep
(Wu et al. 2025 97.6% retention is BUDGET-CONDITIONAL; falsified at T>=4M on Qwen3-8B GSM8K)


A1. Reward vs G (Qwen2.5-0.5B / arithmetic, n=3)



A2. ZVF vs G (same task)



B. Broader-scale G=4 ~ G=32 (Qwen3-8B / GSM8K)



C. Retention loss vs log-T (mechanism)

