


Extended Pillar 3 (iter7): reward-vs-G is flat on easy tasks; G=4 retention of G=32 falls below Wu et al.'s 97.6% as T grows

Reward vs G on measured Qwen2.5-0.5B / arithmetic
(40 steps × 3 seeds, near-ceiling task)



G=4 vs G=32 retention vs token budget T
(illustrative Qwen3-8B / GSM8K reanalysis)

