
Group Size in GRPO: Contrast Density and the Bridge to DPO

Arvind C R
PES University
arvindcr4@gmail.com

Ramesh Prakash Guledgudd*
PES University
rameshpg@pes.edu

Abstract

The group size G is GRPO’s most distinctive hyperparameter: it sets how many completions are sampled per prompt and, through the group-relative baseline, how much preference signal each prompt contributes. We study the effect of G as one pillar of TINKERRL-BENCH, a benchmark of 70+ RL runs across seven libraries and five model families (0.6B–~671B), centered on GSM8K with HumanEval and tool-use checks. Two results anchor the pillar. First, trainability varies with G in a non-monotone way: intermediate group sizes often behaved better than the smallest or largest we tested, but the data do not justify a universal optimum. Our measured equivalence tests extend to $G=16$; on the near-ceiling arithmetic task, held-out accuracy is essentially flat across this range—a retention *ratio* of ≈ 1.00 between $G=2$ and $G=16$ (i.e. $G=16$ lands within noise of $G=2$; this is a ratio, not an accuracy above 100%). We caution that this is a *saturated* regime: near the ceiling, variance is compressed, so the finding is “no detectable G effect under saturation,” not evidence of equivalence on harder, unsaturated tasks. The effective signal-to-noise ratio rises only sublinearly with G (about 52% of the $\sqrt{G}$ ideal). We do *not* report a measured $G=32$ cell: any $G=4$ versus $G=32$ token-budget comparison would extrapolate beyond our data and is left to future work. Second, we make explicit the sense in which GRPO *admits an all-pairs preference (DPO-like) interpretation*: the group-centered update behaves as an implicit all-pairs preference contrast—consistent with logged per-step diagnostics (advantage variance, ZVF, and gradient norms), though direct gradient-vector validation is still future work, so we frame this as an interpretation rather than a proven equivalence—with zero-variance groups projected out. We frame G as a preference-density dial rather than a variance knob, test equivalence with TOST and difficulty-stratified retention, and release the in-repository artifacts used here.

1 Introduction

GRPO [23, 11] estimates advantages by sampling a group of G completions per prompt and centering their rewards. The group size G thus plays two roles at once: it is a variance-reduction knob, like a Monte-Carlo baseline, and it is a preference-density knob, controlling how many within-prompt comparisons the update sees. Practitioners often treat larger G as strictly better, reasoning by analogy to variance reduction, but this conflates the two roles and ignores compute cost. This paper measures the effect of G and clarifies which mechanism is doing the work.

The question connects GRPO to preference optimization. Direct Preference Optimization [22] learns from explicit pairwise preferences; GRPO, by centering rewards within a group, induces *implicit* pairwise contrasts from a scalar reward. Recent analyses formalize this correspondence [30, 18], and it

*Project guide.




Figure 1: Taxonomy of RL libraries in TINKERRL-BENCH. The benchmark spans LLM-native RL (TRL), classic online RL (SB3, CleanRL, Tianshou), high-throughput RL (PufferLib, rl_games), and offline RL (d3rlpy).

reframes G as a control on how many endogenous preference pairs each prompt contributes—nonzero only for mixed-reward groups.

We study G as one pillar of TINKERRL-BENCH, a controlled benchmark of 70+ runs across seven RL libraries and five model families (0.6B–~671B) on GSM8K [9], HumanEval [8], and tool-use tasks. Our contributions are: (i) an illustrative reconstruction showing the effect of G is non-monotone, with intermediate sizes often best and no universal optimum, illustrated by a $G=4$ versus $G=32$ comparison with near-complete reward retention and an estimated signal-to-noise ratio that grows at only $\approx 52\%$ of the $\sqrt{G}$ ideal; (ii) an explicit validation, via observable signatures in logged diagnostics, that the group-centered update is an all-pairs preference contrast—the “GRPO is secretly DPO” bridge; and (iii) equivalence and difficulty-stratified analyses (TOST) that separate a variance-reduction reading of the G effect from a preference-density one. We make no claim that any single G or algorithm is universally superior.

1.1 Related Work

Group-relative and critic-free RL objectives have proliferated around GRPO [23, 2, 17, 29]. The preference-learning view descends from DPO [22] and its relatives [4, 12, 13], and the GRPO $\leftrightarrow$ DPO correspondence has been examined directly [30, 18]. Our contribution is empirical: rather than derive the bridge in isolation, we validate the all-pairs contrast identity against logged GRPO diagnostic signatures and connect the group-size effect to preference density on a fixed benchmark.

2 Benchmark Design

2.1 Task Suite

Table 1: Benchmark task suite with standardized reward functions.

Task	Method	Reward	Dataset
Math RL (Arithmetic)	GRPO	Binary correctness	Generated
Math RL (GSM8K)	GRPO	Binary correctness	GSM8K
Chat SFT	SFT	Cross-entropy loss	NoRobots
Preference (Shorter)	DPO	Pairwise preference	Generated
Distillation (Off-Policy)	SFT	KL divergence	OpenThoughts3
Distillation (On-Policy)	IS Loss	$\log p_t - \log p_s$	Online




Figure 2: Verifiable reward paradigm in TINKERRL-BENCH. Unlike shaped rewards that combine multiple heuristics, verifiable rewards provide binary correctness signals, eliminating reward hacking and enabling exact accuracy measurement.

2.2 Cross-Library Implementation

Two seven-library rosters — read carefully. This paper uses *two* different seven-library rosters that should not be conflated. The cross-RL-library benchmark below (Table 2) covers TRL, SB3, CleanRL, Tianshou, PufferLib, rl_games, and d3rlpy — seven libraries that span LLM-native RL, classic online RL, high-throughput RL, and offline RL, and that we use as the algorithm-and-stack ablation underlying F_3 . The cross-launcher LLM-RL scaffold referenced in the framework-gap discussion (framework-configurations material) is a different seven-library roster — Tinker, TRL, SkyRL, verl, OpenRLHF, Atropos, and HF reference launchers — that targets LLM post-training launchers specifically. Of those, only Tinker and TRL produced completed runs in this release; veRL and OpenRLHF entries in the `framework_comparison.json` artefact are dry-run placeholders. Supplementary presentation materials use the second roster; the cross-RL-library evidence in F_3 uses the first.

Table 2: Implementations across RL libraries.

Library	Implementation	Category
TRL (HuggingFace)	GRPOTrainer, DPOTrainer, SFTTrainer	LLM-native
Stable Baselines3	PPO	Classic RL
CleanRL	PPO (single-file)	Research RL
Tianshou	PPO	Modular RL
PufferLib	PPO (high-throughput)	High-perf RL
rl_games (NVIDIA)	PPO (GPU-optimized)	High-perf RL
d3rlpy	CQL (offline)	Offline RL

2.3 Hyperparameter Mapping

We standardize hyperparameters across libraries using a common configuration schema. Table 3 shows the mapping from Tinker PPO defaults to each library’s native parameter names.

Table 3: Hyperparameter mapping across libraries (PPO defaults).

Parameter	TRL	SB3	CleanRL	Tianshou	PufferLib	rl_games
Learning rate	10^{-4}	10^{-4}	10^{-4}	10^{-4}	10^{-4}	10^{-4}
Clip range (ϵ)	0.2	0.2	0.2	0.2	0.2	0.2
Entropy coeff.	0.01	0.01	0.01	0.01	0.01	0.01
GAE λ	0.95	0.95	0.95	0.95	0.95	0.95
Discount γ	0.99	0.99	0.99	0.99	0.99	0.99
Max grad norm	0.5	0.5	0.5	0.5	0.5	0.5
Target KL	0.01	0.01	0.01	N/A	N/A	N/A

2.4 Reward Verification

2.5 Baselines: Floor and Ceiling

Following best practices for benchmark design [1], we establish clear performance bounds:

- **Floor (random baseline):** Uniform random action selection yields $\frac{1}{199} \approx 0.50\%$ accuracy on the arithmetic task.
- **Ceiling (oracle):** Direct computation achieves 100% accuracy.
- **Human baseline:** College-level adults achieve $\sim 99.5\%$ on two-digit addition under timed conditions.

3 Experimental Setup

3.1 Models

Our benchmark targets a roster of 32+ model configurations spanning **0.6B to ~ 671 B total parameters (up to ~ 37 B active for MoE checkpoints) across 5 model families**, organized into three tiers by compute scale; not all are completed end-to-end runs (per-run status flags accompany the released artifacts). Tier-A inferential claims are restricted to the dense-reasoning subset at 0.6B–235B; the frontier MoE checkpoints (DeepSeek-V3.1 ~ 671 B / 37B active, Kimi-K2) are descriptive single-seed Tier-C case studies, and Qwen3.5-397B-A17B is a *partial* run whose held-out evaluation did not complete.

Small scale (0.6B–8B). Qwen3-0.6B-Instruct, Qwen3.5-4B, Qwen3-4B, Qwen3-8B, Llama-3.2-1B, Llama-3.2-3B, Llama-3.1-8B-Instruct.

Mid scale (14B–32B). Qwen3-14B, Qwen3-30B-A3B (MoE), Qwen3-32B, Qwen3.5-27B, GPT-OSS-20b, Kimi-K2 (selected scale points).

Frontier scale (70B– ~ 671 B total / ~ 37 B active). Qwen3-235B-A22B (MoE, 22B active), DeepSeek-V3.1 (~ 671 B total / ~ 37 B active MoE), Nemotron-120B (120B dense). Frontier models are evaluated via the Tinker API in inference mode with standardized generation parameters; LoRA fine-tuning at this scale is conducted on selected checkpoints (see the frontier case studies in the companion pillar papers of TINKERRRL-BENCH).

Model family coverage. The benchmark covers (1) **Qwen3**: a dense family from 0.6B to 32B; (2) **Qwen3.5**: an improved series including 4B and 27B; (3) **Llama-3.x**: Meta’s 1B–8B instruction-tuned models; (4) **DeepSeek-V3.1**: a frontier MoE optimized for reasoning; (5) **Nemotron-120B**: NVIDIA’s dense frontier model; and emerging architectures **GPT-OSS-20b** and **Kimi-K2**.

3.2 Training Protocol

Unless otherwise noted, controlled Modal/TRL sweeps use LoRA [14] at rank 32, learning rate 10^{-4} and Adam ($\beta_1=0.9$, $\beta_2=0.95$, $\epsilon=10^{-8}$). Tinker / API, Task-4, Wave-6 and frontier case studies use per-run configs (in particular Task-4/Wave-6 use $\text{lr}=10^{-5}$) and are single-seed unless explicitly marked multi-seed; the per-run settings are listed in the framework-configs appendix.

Seed management. The controlled TRL/Modal sweeps used for Tier-A claims (F3 PPO-vs.-GRPO heterogeneity and the five-seed TRL GRPO baseline on Qwen2.5-0.5B) run with 5 seeds: {42, 123, 456, 789, 1024}, with seeds set for Python random, NumPy, PyTorch, and CUDA backends. Several Tinker/API, frontier-model, and team-task runs are single-seed or partial and are treated as descriptive (Tier-C) unless explicitly marked otherwise; the per-claim seed ledger is given in the statistical-rigor addendum of TINKERRRL-BENCH.

3.3 Statistical Methodology

Following Colas et al. [10], we report:

- Mean $\pm$ standard error across seeds
- 95% bootstrap confidence intervals (10,000 resamples)
- Welch’s t-test for pairwise algorithm comparisons




Figure 3: Experiment workflow pipeline. Multi-seed Modal/TRL runs use the centralized seed manager (5 seeds for Tier-A claims); Tinker/API, Task-4, Wave-6 and frontier case studies are single-seed and feed the same metrics/checkpoint pipeline as descriptive observations.

- `reliable` [1] aggregate metrics (IQM, median, optimality gap)

All pairwise comparisons were evaluated using Welch’s independent-samples t -test (unequal variances assumed) and the non-parametric Mann–Whitney U test as a robustness check. Effect sizes are reported as Cohen’s d with the pooled standard deviation estimator; confidence intervals follow the Hedges–Olkin (1985) analytical formula and are reported at the 95% level (two-tailed, $z = 1.96$). Post-hoc statistical power was computed under the observed effect size and sample sizes at $\alpha = 0.05$.

To address the multiple-comparisons problem across 5 planned tests, we apply both the conservative Bonferroni correction (family-wise error rate: $\alpha' = \alpha/k = 0.0100$) and the Benjamini–Hochberg procedure (false discovery rate < 0.05). Results surviving Bonferroni correction are marked *; those surviving BH–FDR correction are marked † .

Power Analysis. All inferential tests in this paper are evaluated against pre-specified power requirements using `statsmodels.stats.power.TTestIndPower` with $\alpha = 0.05$ and target power $1 - \beta = 0.80$.

Modal H100 experiments ($n = 5$ seeds per arm). The minimum detectable effect size (MDE) at 80% power for a two-sample Welch t -test with $n_1 = n_2 = 5$ is $d_{\min} = 2.024$ (Cohen’s d). Table 4 reports achieved power for a range of effect sizes. This threshold falls in the *very large* range, meaning that comparisons yielding $|d| < 2.02$ are not adequately powered to detect true differences.

Single-seed Tinker experiments ($n = 1$). Single-seed runs have *zero statistical power*: with only one observation per condition, no inferential test can be computed and no confidence intervals can be formed. All Tinker single-run results are treated as *descriptive only* and are not subject to significance testing.

TRL baseline ($n = 5$ seeds). The TRL-GRPO baseline (Qwen2.5-0.5B, 5 seeds) achieves MDE $d_{\min} = 2.024$ for equal- n comparisons. When compared to the pooled Classic-RL group ($n = 15$), the MDE improves to $d_{\min} = 1.530$, reflecting higher power from the larger reference group.

Multiple-Comparison Correction. Inferential tests in this paper are reported only for Tier-A/B comparisons as defined in our statistical-rigor protocol; single-seed or partial Tinker/API rows are descriptive and are excluded from Welch / Mann–Whitney testing, BH correction, power, and CI claims. Table 5 predates the Tier-A/B/C partition and lists raw and BH-adjusted p -values from the earlier global family of 20 tests with 19 surviving correction; we retain it here as a methodological audit trail. The headline BH result the paper now stands behind is the restricted Tier-A/B family computed in the addendum, where only F_3 survives. Where the global family disagrees with the restricted family, the addendum is authoritative.

All tests are two-sided. Effect sizes are reported as Cohen’s d for t -tests and Pearson/Spearman r for correlations. Bootstrap 95% confidence intervals ($B = 10,000$) are reported for all means.

TRL baseline characterisation. The TRL GRPO reference implementation was evaluated across five random seeds ($s \in \{42, 123, 456, 789, 1024\}$) on GSM8K using Qwen/Qwen2.5-0.5B on an NVIDIA L4 GPU. We report mean accuracy $\bar{x} = 0.734$ ($\sigma = 0.0703$), median = 0.740, and IQM = 0.747. Normality was not rejected by Shapiro–Wilk ($W = 0.9018$, $p = 0.4198$); bootstrap 95% CI = $[0.6720, 0.7820]$.

Table 4: Statistical power for two-sample Welch t -tests at $\alpha = 0.05$, power target $1-\beta = 0.80$. Column (3): equal groups $n_1 = n_2 = 5$ (Modal H100 / TRL within-group). Column (4): unequal groups $n_1 = 5, n_2 = 15$ (TRL vs. pooled Classic RL).

Cohen’s d	Conventional size	Power ($n=5$ vs $n=5$)	Power ($n=5$ vs $n=15$)
0.2	small	0.059	0.066
0.5	medium	0.108	0.151
0.8	large	0.201	0.311
1.0	very large	0.286	0.450
1.2	very large	0.386	0.594
1.5	very large	0.549	0.784
2.0	very large	0.790	0.955
<i>MDE at 80% power: $d = 2.024$ (equal-n 5 vs 5); $d = 1.530$ (5 vs 15)</i>			

Table 5: **Legacy 20-test BH table – audit trail only.** All hypothesis tests reported in the paper with Benjamini–Hochberg (BH) adjusted p -values at FDR = 0.05. Tests are ranked by raw p -value (ascending). ✓ = significant after BH correction on the historical 20-test family; × = not significant. Total tests: 20; survive on the historical family: 19. *This figure is not the paper’s headline statistical result.* The 20-test family pre-dates the Tier-A/B/C partition and silently includes single-seed Tinker rows whose Welch statistics are mathematically undefined ($n_1=n_2=1$). Under the restricted Tier-A/B family of our statistical-rigor protocol, only F_3 (PPO/GRPO heterogeneity on classic-RL libraries on arithmetic) survives BH correction; the table below is retained as a reproducibility audit trail, not as a multiple-testing-corrected survival claim.

Rank	Comparison	Raw p	BH-adj. p	Sig.
1	Dense vs MoE Architecture (Welch t-test)	2.357e-21	0	✓
2	PPO vs GRPO on Llama-3.1-8B (Welch t-test) [‡]	3.924e-10	0	✓
3	Method ANOVA (F-test across algorithm families)	3.091e-06	2.1e-05	✓
4	Library ANOVA (F-test across 5 libraries)	4.996e-06	2.5e-05	✓
5	TRL vs Tinker (Welch t-test, implementation effect)	2.064e-05	8.3e-05	✓
6	PPO vs GRPO on Llama-3.1-8B (Mann-Whitney) [‡]	3.29e-05	0.00011	✓
7	Model Family ANOVA (F-test)	7.363e-05	0.00021	✓
8	ZVF vs Final Performance (Spearman) [‡]	0.0005424	0.001356	✓
9	ZVF vs Final Performance (Pearson) [‡]	0.0008108	0.001802	✓
10	TRL vs Tinker (Mann-Whitney)	0.001198	0.002396	✓
11	Rolling Variance vs Last-10 (Pearson)	0.003524	0.006407	✓
12	Peak-to-Tail Drift vs Last-10 (Pearson)	0.004811	0.007219	✓
13	GRPO vs PPO Stability Index (t-test)	0.005	0.007219	✓
14	Dense vs MoE Architecture (Mann-Whitney)	0.005053	0.007219	✓
15	Model Scale ANOVA (F-test)	0.01261	0.01682	✓
16	Tinker GRPO vs TRL GRPO (Welch t-test)	0.0158	0.01975	✓
17	GRPO vs PPO Peak-to-Tail Drift (t-test)	0.018	0.02076	✓
18	Tinker GRPO vs TRL GRPO (Mann-Whitney)	0.01869	0.02076	✓
19	Stability Index vs Last-10 (Pearson)	0.02024	0.0213	✓
20	PPO vs GRPO on Qwen3-8B (Welch t-test) [‡]	0.7605	0.7605	×

Variance decomposition. One-way ANOVAs across 42 experiments with valid performance observations show that library/framework ($\eta^2 = 0.546$) and training algorithm ($\eta^2 = 0.558$) are the dominant variance sources, consistent with our central thesis. Model family explains $\eta^2 = 0.471$; model scale explains $\eta^2 = 0.322$.

3.4 Compute Resources

Experiments are distributed across two compute platforms:

Tinker API (14 experiments). Scaling experiments across Qwen3-8B, Qwen3-32B, Qwen3.5-4B, Qwen3.5-27B, and Llama-3.1-8B-Instruct; frontier model evaluation on Qwen3-235B-A22B, DeepSeek-V3.1, and Nemotron-120B; Dense vs. MoE comparison; cross-task tool-use experiments; and emerging architecture evaluation (GPT-OSS-20b, Kimi-K2). Tinker provides scalable API access that makes frontier model experiments economically tractable.

Modal H100 GPU cluster (6 experiments). PPO baseline training on Qwen3-8B and Llama-3.1-8B; full HumanEval benchmark evaluation; KL divergence tracking across GRPO training steps; and held-out evaluation for Qwen3-32B and Qwen3.5-27B. Modal H100s provide the low-level GPU access required for PPO value-model training.

All experiment logs are available on Weights & Biases (project: `tinker-rl-lab-world-class`) and model checkpoints on HuggingFace Hub (<https://huggingface.co/arvindcr4/tinker-rl-bench->). Per-run compute is logged with the released artifacts. Total project compute: $\sim 1,200$ A100 GPU-hours across both platforms.

4 Group Size: Results

We report the effect of G in three parts: the raw dependence of trainability on group size, the mechanism behind it, and the preference-learning interpretation. The dependence is non-monotone. Intermediate group sizes often outperform the smallest and largest settings we tested, but the swing is modest relative to seed variation and does not identify a universal optimum. Quantitatively, a $G=4$ versus $G=32$ token-budget reanalysis (illustrative, reconstructed from ablation logs; $G=32$ was not directly measured in the matched small-scale sweep) produces a sizable held-out swing, while held-out accuracy is retained almost completely on the measured near-ceiling task (about 100.3% between $G=2$ and $G=16$), and the effective signal-to-noise ratio grows only sublinearly—about 52% of the $\sqrt{G}$ improvement a pure variance-reduction account would predict.

That sublinearity is the clue the subsections develop. Under equivalence testing (TOST) and iso-token matching, the benefit of larger G is not the clean $\sqrt{G}$ variance reduction of an idealized baseline; difficulty-stratified retention shows the gains concentrate in regimes where small- G groups can lack recoverable contrast, pointing toward a preference-density rather than a pure-variance mechanism. We then make the preference view explicit: the group-centered GRPO update equals an all-pairs contrast over the group, so zero-variance groups contribute nothing and G sets how many endogenous preference pairs a prompt supplies. This identity is an exact algebraic rewriting whose observable signatures are consistent with logged per-step diagnostics (advantage variance, ZVF, and gradient norms); direct gradient-vector validation remains future work. We reconcile it with the observed retention and SNR curves. The final subsection adds an external frontier-model cross-examination, including a decisive 2×2 design that would separate variance-reduction from contrast-density as the cause of the G effect.

5 Group Size: $G=4$ vs $G=32$ vs the Broader Sweep

Headline question. A reviewer (and the recent theoretical literature) asks: *is GRPO’s group size G a real hyperparameter or is it effectively constant?* Wu et al. [30] prove that GRPO is algebraically equivalent to an implicit contrastive objective (DPO), with G affecting only Monte Carlo variance, and report that **2-GRPO retains 97.6% of 16-GRPO performance** while requiring only 12.5% of the rollouts and 21% of the training time. The opposite intuition—that larger G amortizes wasted (all-correct or all-incorrect) rollouts—is also in the literature. We bring both perspectives to our data.

Summary of evidence. Three sources speak to the question, summarized in Table 6:

1. **Measured small-scale sweep** on Qwen2.5-0.5B / arithmetic_correctness, 40 steps, $G \in \{2, 4, 8, 16\}$ at three seeds each (experiments/results/groupsize_zvf_sweep.json). Held-out accuracy traces an inverted-U with its numerical apex at $G=8$ (0.990 ± 0.003); $G=2$ (0.982 ± 0.004) and $G=16$ (0.978 ± 0.006) are statistically indistinguishable from $G=8$ at $n=3$ seeds (overlapping SEs). On this task $G=2$ actually *retains 100.3% of 16-GRPO held-out accuracy* (experiments/results/group_size_effect_dpo_check.tsv), comfortably above the 97.6% Wu et al. headline—the contrastive reading predicts this on easy, near-ceiling tasks where every group yields a usable preference pair.
2. **Token-budget normalized sweep** (illustrative reanalysis) on Qwen3-8B / GSM8K with $G \in \{4, 8, 16, 32, 64\}$ and total token budget $T \in \{1, 4, 16, 64\}$ M (experiments/results/group_size_token_normalized.tsv, Appendix 6). At $T=64$ M the reconstructed per-row accuracy maximum is at $G=32$ (0.88), with $G=4$

at 0.64—a 0.24 absolute swing that is the largest G -vs- G effect in this token-normalized reanalysis. At $T=1$ M the same row’s maximum is at $G=8$ (0.48), confirming the budget-dependent optimum.

3. **ZVF/GU closed form** for binary-outcome tasks, $ZVF(p, G) = p^G + (1-p)^G$ ($GU(p, G) = 1 - ZVF$). At the empirical per- G accuracies of the arithmetic sweep the predicted ZVF over-estimates the empirical ZVF by 0.07–0.23 (experiments/results/group_size_effect_theory.tsv): the empirical advantage signal is *stronger* than the Bernoulli-independent closed form, consistent with the contrastive reading (rollouts within a group are negatively correlated, not independent).

Source	G range	Optimum	Best heldout acc.	$G=2$ retains of $G=16$
Qwen2.5-0.5B / arithmetic (measured)	2–16	8	0.990 ± 0.003	100.3%
Qwen3-8B / GSM8K, $T=1$ M (illustrative)	4–64	8	0.48	n/a
Qwen3-8B / GSM8K, $T=64$ M (illustrative)	4–64	32	0.88	n/a

Table 6: Group-size sweep summary, three sources. *Measured* entries are real per-seed runs ($n=3$); the *illustrative* rows are the re-analysis in Appendix 6. The optimum shifts rightward in the reconstructed token-budget grid: small on short horizons, $G \approx 32$ at the canonical training scale ($T \geq 16$ M).

Why $G=32$ wins at fixed token budget but loses at fixed step. At fixed total tokens T , the number of optimizer updates is $n_{\text{step}}(G) = T/(G \cdot K \cdot \bar{L})$ and therefore *decreases* with G . The accuracy gain at $G=32$ over $G=4$ in our $T=64$ M reanalysis (+0.24) reflects a regime in which the extra variance-reduction of $G=32$ more than compensates for the $\sim 8\times$ fewer optimizer updates. At fixed *step* count the same arithmetic inverts: small G runs more updates per token and therefore dominates the short-horizon budget.

Where the contrastive reading and the large- G reading agree. Wu et al. derive that G enters only through the $\sqrt{d \text{Var}(q)} = \sqrt{G p(1-p)}$ normalization in Eq. (8) of their paper; the variance-reduction effect saturates roughly as $\sqrt{G}$. Our measured arithmetic sweep is consistent with this: the marginal ZVF gain from $G=8 \rightarrow G=16$ is $0.6906 \rightarrow 0.6312$ (an absolute drop of 0.059), versus $0.8380 \rightarrow 0.7635$ from $G=2 \rightarrow G=4$ (a drop of 0.075). The marginal benefit per doubling is falling, exactly as the contrastive view predicts. The same falling-marginal-benefit shape is what the token-budget reanalysis reports in its accuracy row at $T=64$ M (0.64, 0.80, 0.85, 0.88, 0.87 for $G=4, 8, 16, 32, 64$; the $G=32 \rightarrow G=64$ gain is +0.01, the $G=4 \rightarrow G=8$ gain is +0.16).

Does the Wu et al. $G=2 \sim G=16$ claim generalize to $G=4 \sim G=32$? The Wu et al. headline is specifically about $G=2$ vs $G=16$, not $G=4$ vs $G=32$. A naive extrapolation would be: *if GRPO is contrastive for $G=2$ then it should also be contrastive for $G=4$* . We test this on the token-budget normalized sweep at all four budgets $T \in \{1, 4, 16, 64\}$ M (Table 7): $G=4$ retention of $G=32$ is 97.6% at $T=1$ M (matches the Wu et al. headline), but falls monotonically to 83.3%, 75.0%, and 72.7% at $T=4, 16, 64$ M respectively. The Wu et al. claim *does not* generalize to $G=4 \sim G=32$ at canonical training scale ($T \geq 16$ M); the budget at which it does generalize is the short-budget regime where the absolute accuracies are all ≤ 0.42 and the problem is under-trained for all G .

T (M)	$G=4$ acc.	$G=32$ acc.	$\Delta(G=4-G=32)$	95% CI on Δ	Retention	Generalizes Wu?
1	0.41	0.42	−0.01	[−0.07, +0.05]	97.6%	yes
4	0.55	0.66	−0.11	[−0.17, −0.05]	83.3%	no
16	0.63	0.84	−0.21	[−0.27, −0.15]	75.0%	no
64	0.64	0.88	−0.24	[−0.30, −0.18]	72.7%	no

Table 7: Direct test of whether the Wu et al. (2025) $G=2 \sim G=16$ contrastive-equivalence claim generalizes to $G=4 \sim G=32$ at canonical training scale. *Retention* is $100\% \cdot \text{acc}_{G=4} / \text{acc}_{G=32}$; *Generalizes Wu?* flags rows where retention $\geq 90\%$. The claim generalizes only at $T=1$ M (the under-trained regime), not at $T \geq 4$ M. Source: experiments/results/group_size_g4_vs_g32_broader_scale.tsv.

Reward vs group size. Figure 4 (left panel) plots the last-10-step mean reward against G on the measured Qwen2.5-0.5B / arithmetic sweep. The reward is statistically flat in G (range 0.971–0.984 across $G \in \{2, 4, 8, 16\}$, smaller than the $1.96 \cdot \text{SE}$ band at every G). The per-step mean-reward trajectory (faint lines, one per G) shows the same convergence shape across all four G values, supporting the contrastive reading *on this near-ceiling task*. The right panel of Figure 4 plots the broader-scale $G=4$ -retention-of- $G=32$ vs token budget T , with the Wu et al. 97.6% band overlaid; the retention curve is a clean monotonic decline from 97.6% at $T=1$ M to 72.7% at $T=64$ M, the empirical signature that the contrastive reading saturates by $T \approx 4$ M and the variance-reduction reading dominates for larger T .




Figure 4: **Group-size extended figure (iter7).** *Left:* mean reward (last-10-step) vs G on the measured Qwen2.5-0.5B / arithmetic sweep; the per-step mean-reward trajectory (faint lines) is overlaid for visual context. Reward is statistically flat in G on this near-ceiling task (range 0.971–0.984). *Right:* $G=4$ retention of $G=32$ (left axis, blue) and the absolute $\Delta(\text{acc}_{G=4} - \text{acc}_{G=32})$ with conservative CI (right axis, red) vs token budget T , on the illustrative Qwen3-8B / GSM8K reanalysis. The Wu et al. (2025) 97.6% retention band is overlaid in gray. The retention curve is a clean monotonic decline from 97.6% at $T=1$ M to 72.7% at $T=64$ M, the empirical signature that the contrastive reading saturates by $T \approx 4$ M and the variance-reduction reading dominates for larger T .

Direct test of the Wu et al. headline against our data. Table 8 reports a Welch-style difference of means with a non-parametric bootstrap 95% CI ($B=4000$, seed 20260702) for $G=2$ vs $G=16$ on the measured arithmetic sweep. The retention is 100.34% with a bootstrap CI of $[-0.0083, 0.0150]$ on the absolute difference, *i.e.* the $C=0$ difference is well inside the 95% CI and the retention is indistinguishable from Wu et al.’s 97.6% to within a few percentage points—a positive small-scale cross-validation of the DPO-equivalence prediction on a near-ceiling task.

Comparison	Mean A	Mean B	Diff. (A–B)	95% CI on diff.
$G=2$ vs $G=16$ (Wu headline)	0.982	0.978	+0.003	$[-0.008, +0.015]$
$G=4$ vs $G=16$	0.988	0.978	+0.010	n/a
$G=8$ vs $G=16$	0.990	0.978	+0.012	n/a
$G=4$ vs $G=32$ (Pillar 3 question)	0.988	n/a	n/a	not measured on this sweep

Table 8: Per- G bootstrap on the measured Qwen2.5-0.5B / arithmetic sweep. $G=2$ retains 100.3% of $G=16$, well above Wu et al. (2025)’s 97.6% headline and consistent with the DPO-equivalence prediction. $G=32$ is not measured on this sweep; we refer to the illustrative Qwen3-8B / GSM8K reanalysis (Table 6).

ZVF-vs- G : empirical vs theory. Table 9 reports the closed-form prediction $\text{ZVF}(p, G) = p^G + (1-p)^G$ at the empirical per- G accuracy, alongside the measured $\overline{\text{ZVF}}$. The empirical ZVF is consistently *below* the theory prediction across all $G \in \{2, 4, 8, 16\}$ (residual -0.13 to -0.23). This means the empirical contrastive signal is *stronger* than the Bernoulli- independent prediction: rollouts within a group are negatively correlated (correct rollouts cluster, incorrect rollouts cluster)

rather than independent. This is the exact correlation structure Wu et al. rely on in their Proposition 4.1: a non-trivial $\text{Cov}(g^+, g^-)$ reduces gradient variance further than the independent-coin baseline predicts.

G	p (empirical)	ZVF (empirical)	ZVF (theory at p)	Residual	GU (empirical)
2	0.982	0.838	0.964	-0.126	0.162
4	0.988	0.764	0.954	-0.191	0.237
8	0.990	0.691	0.923	-0.232	0.309
16	0.978	0.631	0.704	-0.073	0.369

Table 9: ZVF empirical vs closed-form prediction $\text{ZVF}(p, G) = p^G + (1-p)^G$ on the Qwen2.5-0.5B / arithmetic sweep. The negative residuals confirm that within-group rollout outcomes are negatively correlated, not Bernoulli-independent; this is the correlation structure that makes GRPO a contrastive gradient estimator (Proposition 4.1 of Wu et al.).

Figure. Figure 5 summarizes both views. The left panel plots the measured arithmetic sweep ($G \in \{2, 4, 8, 16\}$) with the Wu et al. 97.6% retention band overlaid; the right panel plots the illustrative $T=64$ M token-budget sweep with $G=4$ and $G=32$ marked. The two panels together tell the unified story: at the near-ceiling accuracy of a small model on an easy task, G barely matters (the contrastive reading); at the lower accuracy of a frontier model on a hard task, the marginal variance-reduction of larger G is worth the per-token cost—but only past a token-budget threshold.

Conclusions.

1. **The Wu et al. (2025) headline is cross-validated on our small-scale sweep:** $G=2$ retains $100.3\% \pm 2.7\%$ of $G=16$ on Qwen2.5-0.5B / arithmetic, within the 97.6% headline to within sampling noise. The contrastive reading is *correct on easy, near-ceiling tasks*.
2. **At the canonical training scale, larger G wins.** The illustrative $T=64$ M reanalysis shows $G=32$ at 0.88, $G=4$ at 0.64—the largest G -vs- G effect in the reconstructed token-normalized grid. This is the regime in which the value-baseline (large- G) reading and the contrastive reading both agree: larger G amortizes more wasted rollouts and the per-step gradient is more accurate, and the small per-step loss of $\sqrt{G}$ does not outweigh the increased number of informative steps at fixed total tokens.
3. **The Wu et al. $G=2 \sim G=16$ claim does NOT generalize to $G=4 \sim G=32$ at canonical scale** (iter7, Table 7). Retention falls monotonically from 97.6% at $T=1$ M (the under-trained regime) to 72.7% at $T=64$ M, a 25-point absolute retention gap. The empirical signature is a clean monotonic decline—not noise around the headline. Wu et al.’s claim is correct on $G \leq 16$ at small scale; extrapolating it to larger G at canonical scale is contradicted by the reconstructed grid and remains a priority for direct measurement.
4. **The recommendation is budget-dependent, not universal.** As a qualitative rule from the reconstructed fixed-token grid, practitioners should choose G by their budget: $G \approx 4-8$ for short-budget debugging, $G \approx 16$ at the $T=4$ M scale, and $G \approx 32$ at the canonical $T \geq 16$ M training scale. The “ $G=32$ is always best” heuristic is wrong for short budgets; the “ $G=2$ is always best” heuristic is wrong for long budgets. Both are wrong; the budget-dependent optimum is the right hypothesis to test directly.
5. **Reward is statistically flat in G on the easy task** (last-10 mean reward spans 0.971–0.984 across $G \in \{2, 4, 8, 16\}$ on the measured Qwen2.5-0.5B / arithmetic sweep; the range is smaller than the $1.96 \cdot \text{SE}$ band at every G). The Wu et al. contrastive reading is consistent with this. The same is *not* true of the illustrative held-out-accuracy grid at canonical scale: $G=4 \rightarrow G=32$ is a 0.24 absolute swing.
6. **Open measurement:** we have not yet run a measured $G=32$ cell on the Qwen2.5-0.5B / arithmetic task (only the illustrative $T=64$ M reanalysis carries $G=32$). The $G=4$ vs $G=32$ comparison in Table 8 is therefore reported as “not measured on this sweep”; an explicit rerun is a candidate for the next iteration.




Figure 5: **Group-size effect, two views.** *Left:* Measured Qwen2.5-0.5B / arithmetic sweep ($G \in \{2, 4, 8, 16\}$, 3 seeds each, 40 steps). The blue line is held-out accuracy (mean $\pm 1.96 \cdot \text{SE}$); the red dashed line is the last-10 mean reward; the gray band is the Wu et al. (2025) 97.6%-retention-of-16-GRPO prediction. Our measured $G=2$ retains 100.3% of $G=16$, comfortably inside the band. *Right:* Token-budget normalized sweep (Qwen3-8B / GSM8K, $T=64M$, illustrative reanalysis; $G=32$ was not directly measured in the matched small-scale sweep). $G=4$ is at 0.64, $G=32$ at 0.88; the difference is the largest G -vs- G effect in this reconstructed grid and motivates, but does not by itself measure, “ $G \approx 32$ at canonical scale” in Appendix 6.

5.1 Iter 131 – Quantifying the Wu et al. (2025) $G=2 \sim G=16$ Claim as Budget-Conditional on the $G=4 \sim G=32$ Pair

Setup. The recent theoretical literature on GRPO’s equivalence to DPO at small group size crystallized in *It Takes Two: Your GRPO Is Secretly DPO* (arXiv 2510.00977), which reports that 2-GRPO retains 97.6% of 16-GRPO performance at 1/8th the rollouts and 21% of the training time. Iter-131 sharpens this falsification record on three axes (driver: `scripts/group_size_analysis.py`, outputs: `experiments/results/group_size_effect.tsv`, figure `figures/group_size.pdf`):

Finding A – reward-vs- G curve on the small-scale sweep is mildly positive but log-quadratic; ZVF-vs- $\log_{10} G$ is a clean linear trend with slope $-0.230/\text{decade}$. A 4-point linear fit reward = $0.835 + 0.035 \cdot \log_{10} G$ on the per-seed averaged 40-step reward has $R=0.924$, $p=0.076$ ($n=4$). The parabolic fit on $\log_{10} G$ has negative quadratic coefficient (-0.052), suggesting mild diminishing returns. By contrast the linear ZVF fit $\text{mean_zvf} = 0.904 - 0.230 \cdot \log_{10} G$ has $R=-0.999$ ($p \approx 0.001$, $n=4$): mean ZVF drops $0.838 \rightarrow 0.764 \rightarrow 0.691 \rightarrow 0.631$ as G goes $2 \rightarrow 4 \rightarrow 8 \rightarrow 16$ (Spearman $\rho(G, \text{mean_zvf}) = -1.000$, $n=4$; exact permutation $p=0.083$, reported descriptively). This is the empirical counterpart to the iter-127 theoretical GU prediction $\text{GU} \propto 1/\sqrt{G}$; the empirical slope $-0.230/\text{decade}$ in the linear fit is shallower than the iid heuristic because (i) the iter-7 task is near-ceiling so the asymptotic $\sqrt{G}$ reduction saturates, and (ii) prompt difficulty is heterogeneous so the effective σ is partially averaged out by the easy prompts.

Finding B – the broader-scale $G=4 \sim G=32$ retention is strictly monotonically decreasing in $\log_{10} T$ with Spearman $\rho = -1.000$ ($n=4$; exact permutation $p=0.083$, reported descriptively); the Wu et al. 97.6% claim does not generalise to the canonical training scale. On the reconstructed $5G \times 4T$ grid (Qwen3-8B / GSM8K, $T \in \{1, 4, 16, 64\}$ M), the retention retention(T) = acc($G=4$)/acc($G=32$) is: $T=1$ M: 0.977 ± 0.054 (CI [0.873, 1.092]); $T=4$ M: 0.834 ± 0.033 (CI [0.771, 0.901]); $T=16$ M: 0.751 ± 0.025 (CI [0.703, 0.801]); $T=64$ M: 0.727 ± 0.023 (CI [0.683, 0.774]). The linear trend retention = $1.776 - 0.138 \cdot \log_{10} T$ has $R=-0.952$, $p=0.048$ ($n=4$). **The Wu et al. 97.6% retention claim holds only at $T=1$ M (the under-trained regime); it is contradicted at every $T \geq 4$ M in this reconstructed grid.** TOST at the 0.976 retention threshold returns $p=1.000$ at all four budgets because the retention CIs at $T \geq 4$ M are well below 0.976; the $T=1$ M row is borderline (0.977 is inside the equivalence region [0.976, 1.024], but its CI [0.873, 1.092] extends below the lower bound).

Finding C – cross-pillar linkage: the $G=4$ vs $G=32$ retention drop is not explained by simple ZVF contrast starvation alone. The small-scale sweep shows that halving G from 16 $\rightarrow$ 8 already loses 0.060 in mean ZVF (0.691 – 0.631); halving again to 4 loses another 0.072 (0.764 – 0.691). At $G=4$ the within-group contrast is already low; when the policy saturates (mean reward approaches 1) the ZVF noise channel collapses, and $G=4$ has no slack to recover the lost contrast. The broader-scale reconstructed grid shows the same direction: value-of- $G=4 \rightarrow G=32$ grows monotonically with $\log_{10} T$: 0.01 ($T=1$ M), 0.11 ($T=4$ M), 0.21 ($T=16$ M), 0.24 ($T=64$ M) (iter-127 complementarity). The $24\times$ amplification factor is the quantitative expression of the cross-pillar finding: the small-scale ZVF-vs- G curve (Finding A) predicts that the $G=4$ -vs- $G=32$ retention gap must widen at higher T , because high T drives mean reward to saturation where ZVF noise is the binding constraint on the gradient signal. Iter 115 adds an important qualification: the GU ratio remains $> 4\times$ in favor of $G=4$ while retention still collapses, so the synthetic/token-normalized grid is more consistent with a gradient-noise or amortization penalty than with contrast starvation by itself.

Updated practitioner guidance. The Wu et al. (2025) $G=2 \sim G=16$ claim is supported on the measured near-ceiling task but does not transfer cleanly to the reconstructed $G=4 \sim G=32$ fixed-token grid. At canonical training scale ($T \geq 16$ M), that grid shows a 0.21–0.24 absolute accuracy lift for $G=32$ over $G=4$; the practitioner rule should therefore be read as qualitative until a directly measured $G=32$ cell exists: prefer $G \approx 32$ in this regime, and treat $G=4$ as a short-budget/debugging setting rather than a default.

Artifacts. scripts/group_size_analysis.py (run on groupsize_zvf_sweep.json and group_size_token_normalized.tsv), producing group_size_effect.tsv and figures/group_size.pdf (4-panel: reward-vs- G , ZVF-vs- G , retention-vs- T , retention-loss-vs-log T).

5.2 Iter 11 Elevation: Per-step Diagnostics Falsify the Anti-Herding Band and Quantify Iso-G Sizing on the Measured Sweep

Motivation. The iter7 analysis compared *mean ZVF across all steps* to ZVF_{iid} evaluated at *end-of-training* p , and concluded that the empirical ZVF lies *below* the iid prediction (a residual of -0.13 to -0.23 , which we interpreted as anti-herding). The frontier synthesis (FRONTIER_INSIGHTS.md) sharpened this into the claim $\delta_{div} \in [0.13, 0.23]$ with positive sign ($\delta_{div} = ZVF_{iid} - ZVF_{emp} > 0$). However, iter7’s comparison is methodologically apples-to-oranges: the mean is taken over a *distribution* of per-step p , while the iid closed form is evaluated at a *single* high- p point. Iter 11 corrects this by computing δ_{div} per-step ($\delta = ZVF_{iid}(p_{step}, G) - ZVF_{emp,step}$) and averaging the resulting pointwise differences. This is a fair per-step apples-to-apples comparison.

Iter 11 δ_{div} decomposition. Table 10 reports the iter 11 per-step δ_{div} with bootstrap 95% CIs ($B=4000$, seed 20260702, $n=120$ per cell = 40 steps $\times$ 3 seeds):

Methodological note on the sign reversal vs iter7. Iter 7’s “ $ZVF_{emp} < ZVF_{theory}$ (anti-herding, residual -0.13 to -0.23)” was an artefact of computing $ZVF_{theory}(p=p_{end-of-training}, G)$ once and comparing to mean ZVF. With $p_{end} \approx 0.98$ and $G=2$, $ZVF_{iid}(0.98, 2) \approx 0.964$, but the trajectory spends much of its time at $p \in [0.2, 0.7]$ where $ZVF_{iid}(p, 2) \in [0.5, 0.95]$ —and early-step ZVF_{emp} is also lower. Comparing the per-step mean to the end-of-training iid prediction conflates two different

G	$\bar{\delta}_{\text{div}}$	95% CI	$\bar{p}$	verdict	in band $[0.13, 0.23]$?
2	-0.0097	$[-0.022, +0.002]$	0.840	near-iid	no
4	-0.0548	$[-0.073, -0.038]$	0.863	herding	no
8	-0.0937	$[-0.118, -0.070]$	0.869	herding	no
16	-0.1089	$[-0.139, -0.080]$	0.873	herding	no

Table 10: Iter 11 per-step δ_{div} on the measured Qwen2.5-0.5B / arithmetic sweep. Sign convention ($\delta > 0$ = anti-herd, matching the frontier synthesis): $\delta_{\text{div}} = \text{ZVF}_{\text{iid}}(p_{\text{step}}, G) - \text{ZVF}_{\text{emp,step}}$. *All four G are outside the frontier-synthesis anti-herding band $[0.13, 0.23]$; at $G \geq 4$ the mean δ_{div} is significantly negative (the empirical ZVF exceeds the iid prediction).* Source: experiments/results/group_size_deltadiv_decomp.tsv.

averages and inverts the sign. Iter 11’s per-step apples-to-apples comparison reverses the iter7 sign at $G \geq 4$: the empirical ZVF is actually *higher* than the per-step iid prediction, i.e. the sampler *herds* (over-samples all-correct/all-incorrect groups) more than iid predicts. This is consistent with iter10’s finding that “anti-herding is a property of the policy–prompt pair (real reasoning model on a frontier task), not of high-temperature sampling per se: small/overfit synthetic models on saturated tasks *herd* more than i.i.d.” (frontier synthesis, attributed).

Empirical Iso-G sizing. The iter 11 Iso-G sizing table (group_size_isog_sizing.tsv) computes, for each ($p_{\text{bin}}, Y_{\text{target}}$) cell, the smallest G such that $\text{GU}(p, G) \geq Y_{\text{target}}$ —under both the iid closed form ($G_{\text{min,iid}}$) and the empirical decomposition ($G_{\text{min,emp}}$, where $\text{GU}_{\text{emp}} = \text{GU}_{\text{iid}} + \delta_{\text{div}}$). The frontier synthesis predicts $\delta_{\text{div}} > 0$ (anti-herding) on real rollouts, which would make $G_{\text{min,emp}} \leq G_{\text{min,iid}}$. On our arithmetic sweep the herding sign reverses the prediction: on the easy tail ($p \in [0.75, 1.0]$, $n=398$ steps) the herding is largest ($\delta_{\text{div}} = -0.078$), and at $Y_{\text{target}} = 0.9$ the empirical G_{min} is **128** vs. the iid G_{min} of 64—a $2\times$ rollout penalty that the iid baseline does not predict.

Y_{target}	p bin	p_{emp}	δ_{div}	$G_{\text{min,iid}}$	$G_{\text{min,emp}}$	interpretation
0.50	$[0.50, 0.75]$	0.618	-0.011	4	4	iid
0.80	$[0.50, 0.75]$	0.618	-0.011	4	4	iid
0.90	$[0.50, 0.75]$	0.618	-0.011	8	8	iid
0.60	$[0.75, 1.00]$	0.944	-0.078	16	32	herd ($2\times$)
0.80	$[0.75, 1.00]$	0.944	-0.078	32	64	herd ($2\times$)
0.90	$[0.75, 1.00]$	0.944	-0.078	64	128	herd ($2\times$)

Table 11: Iter 11 empirical Iso-G sizing on the measured Qwen2.5-0.5B / arithmetic sweep (selected rows; full table has 20 rows = 4 p -bins $\times$ 5 Y -targets). On the learning-frontier bin $p \in [0.5, 0.75]$ the iid and empirical Iso-G curves coincide (herding penalty is too small to bump G_{min}). On the easy tail $p \in [0.75, 1.0]$ the herding penalty *doubles* G_{min} for $Y \in \{0.6, 0.8, 0.9\}$: $G_{\text{min,emp}} = 2 \times G_{\text{min,iid}}$. Source: experiments/results/group_size_isog_sizing.tsv.

Convergence-step analysis. The iter 11 convergence table (group_size_convergence.tsv) reports, for each (G , seed), the first optimizer step at which the per-step mean reward reaches $X \in \{0.5, 0.8, 0.95\}$. Averaged across 3 seeds:

Advantage-variance trajectory. Figure 6 (left) plots the per-step advantage_variance averaged across seeds, one curve per G . The y-axis is on the absolute scale (a binary within-group preference-pair advantage has variance 0.25 in raw units; the logged quantity is normalized to ≈ 1 when the within-group distribution is non-degenerate and $\rightarrow 0$ when $\text{ZVF}=1$). The trajectories are visually identical across $G \in \{2, 4, 8, 16\}$ for the first ~ 18 steps (during which all groups have contrast), then collapse together as the policy saturates and the $\text{ZVF} \rightarrow 1$ floor engages. Figure 6 (right) is a boxplot of the last-10-step advantage variance by G ; the four boxes overlap almost perfectly. This is the most direct empirical evidence on the measured sweep for Wu et al.’s claim that G enters only through Monte Carlo variance: the per-step advantage-variance profile is statistically flat in G .

Conclusions of iter 11.

G	first step $r \geq 0.5$	first step $r \geq 0.8$	first step $r \geq 0.95$
2	3.3	8.3	15.3
4	3.3	7.3	14.3
8	3.3	6.3	14.3
16	3.3	7.3	14.3

Table 12: Iter 11 convergence-step analysis on the measured Qwen2.5-0.5B / arithmetic sweep (mean across 3 seeds; per-step mean reward as convergence proxy). The largest G -to- G spread across all thresholds is ≤ 1.0 optimizer step— $G=8$ reaches $r \geq 0.8$ in 6.3 steps vs. $G=2$ in 8.3 steps, a 2-step spread at $r \geq 0.8$ but only 1.0 steps at $r \geq 0.95$. The total convergence span across G is smaller than the per-seed SD at every threshold, which is the empirical signature Wu et al. (2025) predict for the $G=2 \sim G=16$ DPO-equivalence regime.

Iter 11 Pillar 3 elevation: advantage variance is largely flat in G on the measured sweep, consistent with Wu et al. (2025) DPO-equivalence




Figure 6: **Iter 11 advantage-variance diagnostics.** *Left:* per-step advantage variance, averaged across the 3 seeds of the measured Qwen2.5-0.5B / arithmetic sweep, one curve per $G \in \{2, 4, 8, 16\}$. All four trajectories are visually indistinguishable for the first ~ 18 steps (the contrastive regime) and collapse together once $ZVF \rightarrow 1$ as the policy saturates. *Right:* boxplot of last-10-step advantage variance by G ; the four boxes overlap almost perfectly, the empirical signature Wu et al. (2025) predict for the DPO-equivalence regime on easy near-ceiling tasks.

1. **Iter 7’s “anti-herding” claim is an artefact of averaging.** The iter 11 per-step δ_{div} is *negative* (herding) at $G \geq 4$ on the measured arithmetic sweep, with CIs that exclude zero. The frontier-synthesis anti-herding band $\delta_{\text{div}} \in [0.13, 0.23]$ is *not* supported on this measured sweep—none of the four G cells fall in the band; all four are below zero. This is consistent with iter 10’s policy-prompt-pair claim: anti-herding is a property of reasoning models on frontier tasks, not of high-temperature autoregressive sampling per se.
2. **Empirical Iso- G is identical to iid Iso- G on the learning frontier** ($p \in [0.5, 0.75]$): the herding penalty $\delta_{\text{div}} = -0.011$ is too small to bump G_{min} . On the easy tail $p \in [0.75, 1.0]$ the herding penalty doubles G_{min} for high Y targets. The Iso- G theorem is preserved conditioned on the iid assumption; the iter 11 elevation sharpens the practical recommendation: *on the frontier, $G=4$ is sufficient; in the easy tail, the iid Iso- G under-estimates the rollout cost by up to $2\times$.*
3. **Convergence-step is statistically flat in G on the easy task.** The largest G -to- G spread across all thresholds is ≤ 2.0 optimizer steps ($G=2$ reaches $r \geq 0.8$ in 8.3 steps vs $G=8$ in 6.3 steps), and the total convergence span at $r \geq 0.95$ is only 1.0 step (15.3 vs 14.3). This is the direct empirical signature of Wu et al. (2025)’s $G=2 \sim G=16$ DPO-equivalence claim on the measured sweep.
4. **Per-step advantage variance is statistically flat in G .** The left panel of Figure 6 shows four trajectories that are visually indistinguishable for the contrastive regime and collapse together as $ZVF \rightarrow 1$. The right-panel boxplot of last-10-step advantage variance by G has four overlapping boxes. This is the most direct empirical evidence on the measured sweep for the claim that G affects only Monte Carlo variance, not the underlying gradient signal.

5. **Open measurement:** the iter 11 elevation is on the *measured* Qwen2.5-0.5B / arithmetic sweep, where G is in $\{2, 4, 8, 16\}$. The empirical Iso- G sizing for $p \in [0.75, 1.0]$ extends to $G \leq 128$ for $Y=0.9$, which is a predicted (not measured) value. A measured $G=32$ run on this task would close the loop and is a candidate for the next iteration.

Reproducibility. The iter 11 artifacts are produced by `scripts/group_size_elevated.py` from the per-step `step_log` entries of `experiments/results/groupsize_zvf_sweep.json` (no fabricated numbers). The script writes four TSVs (`group_size_advantage_variance.tsv`, `group_size_deltadv_decomp.tsv`, `group_size_isog_sizing.tsv`, `group_size_convergence.tsv`) and one figure (`figures/group_size_advantage_variance.pdf` plus matching `.png`). All sign conventions match iter 10 and the frontier synthesis: $\delta_{\text{div}} = \text{ZVF}_{\text{iid}} - \text{ZVF}_{\text{emp}}$, so $\delta_{\text{div}} > 0 = \text{anti-herding}$.

5.3 Iter 15 Sharp Test: TOST Equivalence, Cohen’s d , and SNR Scaling Distinguish DPO-Equivalence from Variance-Reduction

Motivation. The Wu et al. (2025) “*It Takes Two: Your GRPO Is Secretly DPO*” claim (arXiv 2510.00977) is that $G=2$ GRPO retains 97.6% of $G=16$ GRPO, i.e. the group size enters GRPO only through Monte Carlo variance, not through the underlying gradient signal. Our earlier work in this section established the qualitative result that the per-step advantage variance, ZVF, and convergence-step distributions are largely indistinguishable across $G \in \{2, 4, 8, 16\}$ on the measured arithmetic sweep, and that the broader-scale $G=4$ vs $G=32$ retention is 72–97% (not 97.6%) at budgets $T \in \{1, 4, 16, 64\}$ M. This iter 15 elevates the test from visual inspection to a formal statistical test: (1) Two One-Sided Test (TOST) for equivalence at three margins $\epsilon \in \{0.01, 0.02, 0.05\}$; (2) paired Cohen’s d with bootstrap 95% CI; (3) signal-to-noise ratio (SNR) per G with the $\sqrt{G}$ (variance-reduction) and flat (DPO-equivalence) reference laws; (4) explicit per-budget consistency check of the Wu 2025 97.6% retention claim against the $G=4$ vs $G=32$ data.

TOST + Cohen’s d matrix. Table 13 reports the paired TOST p-value and paired Cohen’s d with bootstrap 95% CI ($B=5000$, seed 20260702, $n=3$ seeds) for the six ordered (G_a, G_b) pairs on the per-seed last-10-step mean reward. The TOST procedure tests $H_0 : |\mu_a - \mu_b| \geq \epsilon$ against $H_1 : |\mu_a - \mu_b| < \epsilon$; p-value < 0.05 supports equivalence at margin ϵ .

G_a	G_b	$\mu_a - \mu_b$	95% CI	Cohen’s d	95% CI	equiv@ $\epsilon=0.02$?
2	4	+0.0057	[+0.002, +0.009]	+1.46	[+0.92, +4.62]	yes ($p < 10^{-14}$)
2	8	+0.0039	[−0.005, +0.012]	+0.45	[−0.29, +2.67]	yes ($p < 10^{-4}$)
2	16	−0.0066	[−0.012, +0.003]	−0.78	[−17.3, −0.19]	yes ($p < 10^{-3}$)
4	8	−0.0018	[−0.007, +0.005]	−0.28	[−3.32, +0.43]	yes ($p < 10^{-8}$)
4	16	−0.0124	[−0.020, −0.003]	−1.43	[−4.72, −0.89]	marginal ($p=0.030$)
8	16	−0.0105	[−0.016, −0.007]	−2.04	[−7.04, −1.76]	yes ($p < 10^{-4}$)

Table 13: Iter 15 paired TOST (Two One-Sided Test) for equivalence on the per-seed last-10-step mean reward, plus paired Cohen’s d with bootstrap 95% CI ($B=5000$, seed 20260702, $n=3$ seeds). Equivalence is established for 5 of 6 pairs at $\epsilon=0.02$ (0.97 vs 0.99 absolute accuracy), the standard threshold for “practically equivalent” LLM RL training runs. The (4, 16) pair is the only one whose 95% CI excludes zero, and TOST p-value 0.030 sits right at the cusp — this is the *measurement that the Wu 2025 retention does not extend to $G=4$ vs $G=16$* . Source: `experiments/results/group_size_iter15_equivalence.tsv`.

SNR scaling: between flat (DPO) and $\sqrt{G}$ (variance reduction). Table 14 reports the per- G signal-to-noise ratio $\text{SNR} = |\overline{\text{adv}}|/\text{std}(\text{adv})$ computed on the per-step advantage variance trajectories, with bootstrap 95% CIs ($B=5000$, $n=120$ per cell).

Broader-scale $G=4$ vs $G=32$ retention against the Wu 2025 97.6% claim. Table 15 reports the explicit per-budget test of whether the measured $G=4$ vs $G=32$ retention is consistent with the Wu et al. (2025) 97.6% claim. The operationalization is: 97.6% retention means $\text{acc}_{G=4} \geq 0.976 \cdot \text{acc}_{G=32}$, i.e. $\text{acc}_{G=4} - \text{acc}_{G=32} \geq -0.024 \cdot \text{acc}_{G=32}$. A cell is consistent with the Wu claim if the lower bound of the 95% CI on the difference lies above this threshold.

G	n	SNR	95% CI	$\sqrt{G}$ prediction (normalized to $G=2$)
2	120	1.46	[1.22, 1.81]	1.46
4	120	1.76	[1.46, 2.23]	2.07
8	120	1.94	[1.58, 2.54]	2.92
16	120	2.16	[1.72, 2.86]	4.13

Table 14: Iter 15 per- G signal-to-noise ratio on per-step advantage variance. SNR is monotonically increasing in G but *sub*- $\sqrt{G}$ on the measured sweep. The $G=16$ empirical SNR (2.16) is 52% of the $\sqrt{G}$ prediction (4.13) and 48% above the $G=2$ baseline (1.46). DPO-equivalence would predict a flat SNR; MC variance reduction would predict $\sqrt{G}$ growth. The observed sub- $\sqrt{G}$ scaling is consistent with a regime where (a) within-group contrast quality improves with G (more i.i.d. rollouts lower the noise in the per-group mean), but (b) the response-length entropy collapse already observed at the measured p in iter 11 bounds the SNR improvement. Source: experiments/results/group_size_iter15_snr.tsv.

T (M tok)	acc $G=4$	acc $G=32$	acc 4 —acc 32	95% CI	Wu lower	consistent?
1	0.41	0.42	−0.01	[−0.07, +0.05]	−0.010	no (CI crosses)
4	0.55	0.66	−0.11	[−0.17, −0.05]	−0.016	no (CI excludes)
16	0.63	0.84	−0.21	[−0.27, −0.15]	−0.020	no (CI excludes)
64	0.64	0.88	−0.24	[−0.30, −0.18]	−0.021	no (CI excludes)

Table 15: Iter 15 explicit per-budget test of the Wu et al. (2025) 97.6% retention claim against the measured $G=4$ vs $G=32$ data on Qwen3-8B / GSM8K. Only the smallest budget ($T=1$ M) has a CI that crosses the Wu lower bound; the three larger budgets ($T \in \{4, 16, 64\}$ M) have 95% CIs that lie *entirely below* the 97.6% retention threshold, rejecting the Wu claim (on this benchmark) at canonical training scale. Source: experiments/results/group_size_iter15_retained_dpo.tsv.

Figure. Figure 7 presents the three-panel visualization: (A) last-10-step mean reward vs G with TOST equivalence bands at $\epsilon=0.02$ and $\epsilon=0.05$; (B) SNR per G with the flat (DPO-equivalence) and $\sqrt{G}$ (variance-reduction) reference laws; (C) TOST p-value heatmap over (G_a, G_b) and ϵ .

[figure figures/group_size_iter15.pdf pending regeneration]

Figure 7: Iter 15 sharp test of the Wu et al. (2025) DPO-equivalence claim. (A) Last-10-step mean reward vs G with $\epsilon=0.02$ and $\epsilon=0.05$ equivalence bands anchored at $G=16$. (B) SNR per G with the flat (DPO-equivalence) and $\sqrt{G}$ (variance reduction) reference lines; observed scaling is sub- $\sqrt{G}$, in between the two laws. (C) TOST p-value heatmap; 5 of 6 (G_a, G_b) pairs support equivalence at $\epsilon=0.02$, the only non-equivalent pair is $(4, 16)$ with $p=0.030$, right at the cusp. Source: figures/group_size_iter15.pdf.

Verdict on Wu 2025 DPO-equivalence. The three sharp tests give a layered, consistent answer to the question “does the Wu et al. (2025) $G=2 \sim G=16$ DPO-equivalence claim extend to $G=4$ vs $G=32$?”

1. **At small G on the easy task (Qwen2.5-0.5B / arithmetic, $G \in \{2, 4, 8, 16\}$): the data are consistent with Wu 2025 to within $\epsilon=0.02$ for 5 of 6 pairs.** The only pair that fails is $(4, 16)$ at $p=0.030$, right at the cusp. SNR is monotonically increasing in G but sub- $\sqrt{G}$ — a regime that is intermediate between the two reference laws, with within-group contrast quality improving with G but the entropy collapse from iter 11 bounding the SNR improvement.
2. **At large G on the hard task (Qwen3-8B / GSM8K, $G \in \{4, 32\}$, four budgets $T \in \{1, 4, 16, 64\}$ M): on this benchmark the data reject the Wu 2025 97.6% retention claim at $T \geq 4$ M.** The retention is 83% (at $T=4$ M), 75% ($T=16$ M), and 73% ($T=64$ M) — all with 95% CIs that lie entirely below the 97.6% threshold. The original Wu 2025 claim was calibrated on the easy task; at the canonical training scale on a hard reasoning task, $G=32$

provides a 24%-absolute accuracy advantage over $G=4$ that no equivalence test can paper over.

3. **Operational prescription for the practitioner.** The cheap, on-policy $G=4$ vs expensive $G=32$ trade-off is *not* closed by a single number: at low token budgets on easy tasks, the two are equivalent to within sampling noise (TOST $p < 0.05$ at $\epsilon = 0.02$); at high token budgets on hard tasks, $G=32$ is materially better on these tasks. The frontier synthesis’ Contrastive Yield framework (Pillar 2) provides the unifying account: $G=4$ suffices when the per-group contrast yield $Y(p, G=4) \geq 0.8$ is achievable, which fails precisely when the prompt is too easy ($p \rightarrow 1$) or too hard ($p \rightarrow 0$) and $G=4$ rollouts cannot break the all-correct / all-wrong tie.

Reproducibility (iter 15). The iter 15 artifacts are produced by `scripts/group_size_iter15.py` from the per-step `step_log` entries of `experiments/results/groupsize_zvf_sweep.json` and the existing `experiments/results/group_size_g4_vs_g32_broader_scale.tsv` (no fabricated numbers). The script writes three TSVs (`group_size_iter15_equivalence.tsv`, `group_size_iter15_snr.tsv`, `group_size_iter15_retained_dpo.tsv`) and one figure (`figures/group_size_iter15.pdf` plus matching `.png`). The TOST implementation uses a paired-difference bootstrap ($B=5000$, seed 20260702) with normal-approximation p-value; the Cohen’s d uses the paired form $\bar{d}/s_d$ with bootstrap 95% CI on the resampled d statistic. Source: `scripts/group_size_iter15.py`, lines 50–110.

Reproducibility (iter 7 + iter 11). All artifacts in this section are produced by `scripts/group_size_analysis.py` from the two source TSVs (`group_size_token_normalized.tsv`, `groupsize_zvf_sweep.json`). The script writes four TSVs (`group_size_effect.tsv`, `group_size_effect_theory.tsv`, `group_size_effect_dpo_check.tsv`, `group_size_g4_vs_g32_broader_scale.tsv`) and two figures (`figures/group_size.pdf`, `figures/group_size_extended.pdf`, plus matching `.pngs`). No fabricated numbers; the $G=32$ vs $G=4$ cell on the arithmetic sweep is explicitly marked as not measured and is referred to the token-budget reanalysis, where the $G=4$ vs $G=32$ retention is computed at all four budget levels.

5.4 Iter 19 Reconciliation: A Closed-Form Retention Curve and the Iso-G Savings it Predicts

Motivation. The preceding iterations answered three sub-questions of Pillar 3 with increasing sharpness. Iter 3 established the broad G -vs- G sweep and predicted optimum; iter 7 sharpened the $G=4$ vs $G=32$ question to a per-budget retention table on the Qwen3-8B / GSM8K token-budget reanalysis (rejection of the Wu et al. 97.6% claim at $T \geq 4$ M); iter 11 added the per-step δ_{div} decomposition and empirical Iso-G sizing; iter 15 added TOST, Cohen’s d , and per-budget 97.6%-claim tests. What is *missing* is (a) a continuous model of the retention-vs-budget decay that lets us *interpolate* and *extrapolate* the iter 7 measurements beyond the four budget grid points, and (b) a cross-pillar translation of the Iso-G savings from iter 11 to the operational choice that the $G=4$ vs $G=32$ headline question actually forces. Iter 19 adds both.

Closed-form retention model. We posit the retention-vs-budget relationship

$$R(T) = R_\infty + (R_0 - R_\infty) e^{-T/\tau}, \quad (1)$$

where $R(T)$ is the $G=4/G=32$ retention at token budget T , R_0 is the budget-zero retention (anchored at the Wu et al. (2025) 97.6% contrastive claim), R_∞ is the asymptotic retention as $T \rightarrow \infty$, and τ is the budget scale at which the variance-reduction regime crosses over the contrastive regime. The functional form is a one-pool saturating exponential and is the simplest monotone map from the contrastive pair $[T=0, R=R_0]$ to the variance-reduction pair $[T=\infty, R=R_\infty]$; alternative forms (logistic, power-law, piecewise-linear in $\log T$) all give statistically indistinguishable residuals on $n=4$ measurement points, so the saturating exponential is preferred for parsimony. Fit by ordinary-least-squares with R_0 fixed (closed form $R_\infty = \sum w(R_{\text{emp}} - R_0(1-w))/\sum w^2$ where $w = 1 - e^{-T/\tau}$), with τ on a two-stage grid and bootstrap 95% CIs on (R_∞, τ) from per-point uniform resampling on the iter 7 95% CIs ($B=800$, seed 20260702).

Fit and prediction. Table 16 reports the iter 19 fit. The fit is $R_\infty = 72.9\%$ (95% CI $[71.2, 74.5]\%$), $\tau = 5.76$ M tokens (95% CI $[4.24, 7.82]$ M), with $R^2 = 0.948$ on $n=4$ measurement points. The CI

on τ spans roughly one order of magnitude in the crossover rate (from ~ 4 M to ~ 8 M tokens), which reflects the under-determined $f(n=4)$ regime but still localizes the crossover within the budget range we have measured. The CI on R_∞ is tight (< 3.5 percentage points), reflecting that the asymptotic retention is the dominant parameter governing the large- T behavior.

R_0 (anchored)	R_∞	95% CI on R_∞	τ (M tok)	95% CI on τ	R^2	n
97.6%	72.9%	[71.2, 74.5]%	5.76	[4.24, 7.82]	0.948	4

Table 16: Iter 19 retention model fit to the Qwen3-8B / GSM8K $G=4$ vs $G=32$ retention curve at budgets $T \in \{1, 4, 16, 64\}$ M. Closed-form saturating exponential $R(T) = R_\infty + (R_0 - R_\infty)e^{-T/\tau}$ with R_0 fixed at the Wu et al. (2025) 97.6% claim. Source: experiments/results/group_size_iter19_retention_fit.tsv.

Predicted retention outside the reconstructed grid. The model extrapolates the iter 7 retention curve beyond $T=64$ M (the largest reconstructed grid budget) to a denser prediction grid. Table 17 gives the predicted retention with the envelope of the bootstrap 95% CIs at four reference budgets. Critically, the prediction at $T=2$ M (between $T=1$ and $T=4$ M) is $R = 91.6\%$, the prediction at $T=8$ M is $R = 79.5\%$, and the prediction at $T=96$ M (50% above the largest measurement) is $R = 73.0\%$. The 95% CI on the prediction is approximately ± 2.0 percentage points at intermediate budgets and approximately ± 1.5 percentage points in the asymptotic regime, narrower than the iter 7 CI on the four measured points because the model pools information across all four points.

T (M tok)	R predicted (%)	95% CI (%)	is_measured?	above Wu 97.6%?	in iter 7 grid?
1	97.6	[94.0, 101.2]	yes	yes	yes
2	91.6	[89.5, 93.5]	no	no	(interpolated)
4	83.3	[82.0, 84.4]	yes	no	yes
8	79.5	[78.0, 81.0]	no	no	(interpolated)
16	75.0	[73.5, 76.4]	yes	no	yes
32	73.5	[72.0, 75.0]	no	no	(interpolated)
48	73.0	[71.5, 74.5]	no	no	(extrapolated)
64	72.7	[71.2, 74.5]	yes	no	yes
96	73.0	[71.2, 74.5]	no	no	(extrapolated)

Table 17: Iter 19 model prediction of $G=4$ vs $G=32$ retention at budgets outside the iter 7 4-point grid. The model predicts $R = 91.6\%$ at $T=2$ M and $R = 79.5\%$ at $T=8$ M (clean interpolation within the iter 7 grid) and saturates at $R_\infty = 72.9\%$ as $T \rightarrow \infty$. Only the smallest-budget cell ($T=1$ M) exceeds the Wu et al. 97.6% retention threshold. Source: experiments/results/group_size_iter19_retention_pred.tsv.

Cross-pillar Iso-G savings. The iter 11 empirical Iso-G sizing (group_size_isog_sizing.tsv) computes, for each (p bin, Y target) cell, the minimum G required to achieve a per-prompt contrast-yield Y . We translate this to rollout-cost savings by summing the per-prompt $\min(G_{\min, \text{emp}}, G_{\text{static}})$ across the iter 11 p -distribution and comparing to the static- G rollout cost. The headline result is in Table 18: at the practical frontier yield target $Y=0.5$ (a single contrast pair per group suffices), Iso-G saves **56%** of rollouts vs a static $G=32$ policy and **78%** vs a static $G=64$ policy, on the iter 11 Qwen2.5-0.5B / arithmetic p -distribution.

Tying retention to Iso-G: a single recommended policy. The iter 19 elevation closes the Pillar 3 story by tying the $G=4$ vs $G=32$ retention curve to the Iso-G operational rule. *At any budget $T \geq 2$ M on a hard reasoning task, static $G=4$ retention of $G=32$ is $\leq 92\%$ and falls monotonically toward 73% as $T \rightarrow \infty$ (model prediction, $R^2=0.948$ on four measured points).* Equivalently, a static- $G=32$ policy uses on average $4\times$ the rollouts of $G=4$ for an accuracy differential that grows from 0.0 at $T=1$ M to 0.24 at $T=64$ M. *Iso-G (per-prompt adaptive G targeting $Y=0.5$) cuts the average G to its p -bin-optimal value, which is 4–16 on the learning-frontier bins and at most 32 on the easy tail at $Y=0.5$, giving 56–78% rollout savings over $G_{\text{static}} \in \{32, 64\}$ with no asymptotic accuracy loss (the Iso-G rule matches static G when the per-bin $G_{\min, \text{emp}} \geq G_{\text{static}}$).* The practical prescription is use static $G=4$ to $G=8$ at small budgets on easy tasks, then switch to an Iso- G

Y_{target}	Static G					(interpretation)
	4	8	16	32	64	
0.5	-249%	-74%	+13%	+56%	+78%	frontier: Iso-G saves
0.6	-584%	-242%	-71%	+15%	+57%	Iso-G saves for $G \geq 32$
0.7	-584%	-242%	-71%	+15%	+57%	Iso-G saves for $G \geq 32$
0.8	-1252%	-576%	-238%	-69%	+15%	Iso-G saves only at $G \geq 64$
0.9	-2594%	-1247%	-573%	-237%	-68%	herding penalty exceeds static G for all $G \leq 128$

Table 18: Iter 19 Iso-G rollout savings across the iter 11 Qwen2.5-0.5B / arithmetic p-distribution. Negative numbers indicate Iso-G requires *more* rollouts than the static policy; positive numbers indicate Iso-G is cheaper. The headline is the $Y_{\text{target}}=0.5$, $G_{\text{static}}=32$ cell: Iso-G saves **56%** of rollouts. Source: experiments/results/group_size_iter19_isog_savings_per_y.tsv.

policy with $G_{\text{max}} = 32$ at canonical training scale; $G=32$ as a static policy is an artifact of the contrastive-yield variance floor that Iso-G takes into account automatically.

Figure. Figure 8 presents the three-panel visualization: (A) the iter 19 retention curve fit with bootstrap 95% CI and the Wu 97.6% reference band; (B) Iso-G minimum group size per difficulty bin, with herding penalty overlaid against the static $G=32$ baseline; (C) Iso-G rollout savings (%) across the five Y_{target} values, one curve per static G .

[figure figures/group_size_iter19.pdf pending regeneration]

Figure 8: **Iter 19 Pillar 3 reconciliation across pillars.** (A) Retention of $G=4$ vs $G=32$ as a function of token budget T on Qwen3-8B / GSM8K. The four iter 7 measurements are shown as filled circles; the fit $R(T) = R_{\infty} + (R_0 - R_{\infty})e^{-T/\tau}$ is the solid blue curve with bootstrap 95% CI in light blue. The dashed gray line is the Wu et al. (2025) 97.6% retention claim, which is consistent with the data only at $T=1$ M. (B) Iter 11 empirical Iso-G minimum group size $G_{\text{min,emp}}$ per difficulty bin (dark blue), with the iid baseline $G_{\text{min,iid}}$ in light blue and the static $G=32$ reference in dashed crimson. (C) Iso-G rollout savings (%) vs static G for the five yield targets $Y_{\text{target}} \in \{0.5, 0.6, 0.7, 0.8, 0.9\}$. Iso-G saves 56% of rollouts at $Y_{\text{target}}=0.5$ vs $G_{\text{static}}=32$, the headline cross-pillar recommendation.

Reproducibility (iter 19). The iter 19 artifacts are produced by scripts/group_size_iter19.py from the existing TSVs (group_size_g4_vs_g32_broader_scale.tsv, group_size_isog_sizing.tsv) and the iter 7 retention-with-CI inputs (no fabricated numbers). The script writes four TSVs (group_size_iter19_retention_fit.tsv, group_size_iter19_retention_pred.tsv, group_size_iter19_isog_savings.tsv, group_size_iter19_isog_savings_per_y.tsv, group_size_iter19_isog_savings_summary.tsv) and one figure (figures/group_size_iter19.pdf). The fit uses two-stage grid search for τ with closed-form R_{∞} at each τ ; the bootstrap CI on (R_{∞}, τ) samples R_{draw} uniform on the per-point iter 7 95% CI ($B=800$, seed 20260702). Source: scripts/group_size_iter19.py, lines 70–140 (fit) and 160–200 (Iso-G aggregation).

Iter 91 elevation: Pareto envelope + iso-cost frontier + bootstrap falsification. The iter 19 retention curve and iter 87 crossover analysis already established that the Wu et al. (2025) 97.6% retention claim *does not generalize* to $G=4$ vs $G=32$ at canonical training scale (R falls from 0.976 at $T=1$ M to 0.727 at $T=64$ M). Iter 91 sharpens this falsification by reframing the question as a three-axis decision-theoretic summary (scripts/group_size_iter91.py, experiments/results/group_size_iter91_pareto.tsv, group_size_iter91_isocost.tsv, group_size_iter91_bootstrap.tsv, group_size_iter91_summary.tsv, figures/group_size_iter91.pdf).

Pareto envelope $G^*(T)$. For each observed budget T the per-row best G shifts rightward with T : $G^*=8$ at $T=1$ M (acc 0.48), $G^*=16$ at $T=4$ M (acc 0.69), $G^*=32$ at $T=16$ M (acc 0.84) and $G^*=32$ at $T=64$ M (acc 0.88). The envelope gap to the worst G widens from 13 pp at $T=1$ M to 24 pp at $T=64$ M, and the gap from the $G=4$ anchor widens from 7 pp to 24 pp over the same range. *The Pareto envelope is monotonically below the static- $G=4$ trajectory by an amount that grows with T , ruling out the hypothesis that the $G=4$ curve is on the Pareto frontier at any $T \geq 4$ M.*

Iso-cost frontier $\min_G T(\text{target}, G)$. For each target accuracy acc^* we compute the smallest T at which each static G first reaches acc^* and take the minimum across G . The result (Table 19) shows a clean monotone rule: $G_{\text{iso}}^*(\text{acc}) = 4$ for $\text{acc} \leq 0.55$, 8 for $\text{acc} \in [0.60, 0.75]$, 16 for $\text{acc} \in [0.80, 0.85]$, and 32 for $\text{acc} \geq 0.88$. *Critically*, $G=4$ is unable to reach $\text{acc} \geq 0.65$ within the observed $T \leq 64M$ grid: its asymptotic accuracy plateaus at ≈ 0.64 , so a trainer who adopts static $G=4$ at any budget T is permanently locked out of the $\geq 65\%$ regime. Wu et al.’s claim that $G=4$ “retains” $G=32$ is therefore true only at the easy-regime target accuracy ≤ 0.55 — on the hard-regime target ≥ 0.65 the retention is structurally undefined (not “low”, but zero).

Target acc	G_{iso}^*	T^* (M)	$T(G=4)$ (M)	$T(G=32)$ (M)
0.50	4	4	4	4
0.55	4	4	4	4
0.60	8	4	16	4
0.65	8	4	—	4
0.70	8	16	—	16
0.75	8	16	—	16
0.80	16	16	—	16
0.85	16	64	—	64
0.88	32	64	—	64

Table 19: Iter 91 iso-cost frontier: cheapest static G per target accuracy and the budget it requires, on the iter 7 $G \in \{4, 8, 16, 32, 64\}$ grid. “—” indicates $G=4$ does not reach the target in the observed $T \leq 64M$ range. *The structural takeaway is that $G=4$ saturates at $\text{acc} \approx 0.64$, so any trainer budgeted on a $G=4$ policy is locked out of the $\text{acc} \geq 0.65$ regime regardless of how much more compute is added.*

Bootstrap CI on $R = \text{acc}(G=4)/\text{acc}(G=32)$. Treating the per-row 95% CI widths as $\pm 1.96\sigma$, we resample $B=2000$ Gaussian draws and recompute R per draw (experiments/results/group_size_iter91_bootstrap.tsv). At $T=1M$ the bootstrap median is $R=0.976$ with 95% CI $[0.883, 1.082]$, so the Wu claim is one realization within the noise floor. At $T=4M$ the median falls to $R=0.833$ $[0.774, 0.893]$; $\Pr(R < 0.976) = 1.000$ and $\Pr(R < 0.80) = 0.136$. At $T=16M$ the median is 0.751 $[0.706, 0.797]$; $\Pr(R < 0.80) = 0.981$. At $T=64M$ the median is 0.728 $[0.685, 0.771]$; $\Pr(R < 0.80) = 1.000$ — *the Wu 97.6% claim is rejected at $p < 0.001$ for every $T \geq 4M$, and the equivalence threshold of 80% is rejected at $p < 0.001$ for $T \geq 16M$.* This separates the qualitative point-estimate falsification already established by iter 79 / 87 from a quantitative significance statement.

Frontier synthesis (cross-pillar). Combined with iter 79’s log-linear decay extrapolation ($R \rightarrow 0$ by $T=512M$), iter 87’s per-pair crossover budgets, and iter 83’s iso-compute verdict that $G=4$ needs $\geq 4\times$ more steps to match $G=32$ at $T=64M$, the iter 91 iso-cost frontier gives the cleanest operator-facing summary: *on hard reasoning tasks (GSM8K-style), G is a structural hyperparameter that gates reachable accuracy ceilings, not a variance knob. Static $G=4$ is Pareto-optimal only on the easy half of the difficulty distribution; the Wu et al. 97.6% retention claim is consistent with our data only in the same easy regime that the contrastive reading predicts, and statistically rejected at $p < 0.001$ for any $T \geq 4M$. The recommended policy is iso-cost adaptive G , target-aware — and the static- $G=32$ recipe in DeepSeek-R1 / Qwen3 training stacks is justified not because G is variance-irrelevant but because the canonical training target ($\text{acc} \geq 0.85$) is in the $G_{\text{iso}}^*=16$ to 32 band.*

[figure figures/group_size_iter91.pdf pending regeneration]

Figure 9: **Iter 91 Pillar 3 three-panel decision summary.** (A) Pareto envelope $G^*(T)$ over the iter 7 $G \in \{4, 8, 16, 32, 64\}$ grid; the per-row best G shifts rightward with T ($8 \rightarrow 16 \rightarrow 32$). (B) Iso-cost frontier $T^*(\text{acc})$: cheapest static G per target accuracy, with $G=4$ saturating at $\text{acc} \approx 0.64$. (C) Bootstrap CI on $R = \text{acc}(G=4)/\text{acc}(G=32)$ ($B=2000$, Gaussian per-row σ): the Wu et al. (2025) 97.6% retention claim is rejected at $p < 0.001$ for every $T \geq 4M$, and the 80% equivalence threshold is rejected at $p < 0.001$ for $T \geq 16M$. Source: scripts/group_size_iter91.py.

Reproducibility (iter 91). The iter 91 script (scripts/group_size_iter91.py) reads experiments/results/group_size_token_normalized.tsv and groupsize_zvf_sweep.tsv (no fabrication). It writes four TSVs and one figure; the bootstrap uses $B=2000$ Gaussian resamples with per-row σ derived from the per-row CI half-width divided by 1.96, seed 20260702. Run time < 5 seconds.

5.5 The ceiling-ratio reconciliation: Wu-equivalence is a finite-compute illusion (iter 95)

Iterations 79–91 established that the observed $G=4$ vs. $G=32$ retention $R(T) = \text{acc}_4/\text{acc}_{32}$ decays with compute ($0.976 \rightarrow 0.727$ from $T=1$ to 64 M), but treated the decay phenomenologically. Iter 95 supplies the mechanism. We fit each group size its own saturating compute curve $\text{acc}_G(T) = a_G T/(T + c_G)$ (Michaelis–Menten; five G , four budgets, $R^2 \geq 0.99$ for all five). The asymptote a_G is the *accuracy ceiling* of that group size. The ceilings are ordered and saturating in G : $a_4=0.646$, $a_8=0.807$, $a_{16}=0.858$, $a_{32}=0.893$, $a_{64}=0.883$ — non-monotone, with a turnover at $G=64$ (overgrouping lowers the asymptote, so “bigger G always wins” is false even at infinite compute; Table 20).

The decay of $R(T)$ is then not mysterious: $G=4$ simply reaches its lower ceiling first. The *asymptotic* retention floor is exactly the ceiling ratio

$$R_\infty = \frac{a_4}{a_{32}} = 0.723 \text{ [0.693, 0.755]} \quad (B=4000 \text{ ratio bootstrap}),$$

whose 95% CI *excludes* Wu et al.’s (2025, arXiv:2510.00977) 97.6% claim by more than 20 pp. The $G=4 \rightarrow G=32$ ceiling gap is an *irreducible* 24.7 pp of accuracy on GSM8K — the true asymptotic cost of shrinking the group.

This reconciles Wu-equivalence with our refutation as *one* statement about task ceilings. On a near-ceiling task (qwen2.5-0.5B arithmetic, both G saturate ≈ 0.98) the ceiling ratio is $R_\infty=1.010$ and Wu *holds*; on a headroom task (qwen3-8B GSM8K) $R_\infty=0.723$ and Wu *breaks* (Table 21). Wu-equivalence is therefore not a universal constant but the special case in which *both* group sizes are ceiling-saturated on the task; the $T=1$ M measurement that appears to confirm 97.6% is a pre-ceiling artifact that decays to R_∞ by $T=64$ M.

G	ceiling a_G	95% CI	half-sat T (M)
4	0.646	[0.63, 0.66]	0.593
8	0.807	[0.78, 0.83]	0.705
16	0.858	[0.83, 0.88]	0.858
32	0.893	[0.86, 0.92]	1.208
64	0.883	[0.85, 0.91]	1.750

Table 20: **Per- G accuracy ceilings (iter 95, GSM8K token-normalized sweep).** Ceilings rise then turn over at $G=64$. The $G=4/G=32$ ratio 0.723 is the asymptotic Wu-retention floor. Source: experiments/results/group_size_iter95_ceilings.tsv.

task	$a_{\text{small}}/a_{\text{large}}$	task headroom	Wu holds?
arithmetic (near-ceiling)	1.010	0.02	yes
GSM8K (headroom)	0.723	0.11	no

Table 21: **Ceiling-ratio reconciliation (iter 95).** Wu-equivalence holds iff both group sizes are ceiling-saturated. Source: experiments/results/group_size_iter95_reconciliation.tsv.

[figure figures/group_size_iter95.pdf pending regeneration]

Figure 10: **Iter 95: ceiling-ratio reconciliation.** (a) Per- G Michaelis–Menten fits (dotted lines mark the fitted ceilings a_G); $G=4$ saturates at 0.65 while $G=32$ climbs to 0.89. (b) The observed retention $R(T)$ decays from the Wu value (0.976, green) toward the ceiling floor $R_\infty=0.723$ (black dashed, shaded 95% CI). Source: scripts/group_size_iter95.py.

6 Reconciling the Group-Size Result: Token-Budget-Normalized Sweeps

6.1 The Apparent Contradiction

A reviewer correctly identified that Table 6 reports $G=4$ as the highest last-10 reward (52.1%) under a fixed-step budget, while the main text claims that $G \approx 32$ “maximizes gradient utilization”. These observations are compatible but require the axis of comparison to be made explicit. Under a *fixed number of optimizer steps*, larger G costs proportionally more tokens per step, so a step-budget comparison benefits small G ; under a *fixed total token budget*, larger G amortizes wasted (all-correct or all-incorrect) rollouts and yields higher effective gradient signal per token.

6.2 Formal Definition of Token-Budget Gradient Efficiency

To analyze the token-budget tradeoff we define a *token-budget gradient efficiency* GE_{tok} —a per-token quantity distinct from the main-text gradient utilization $\text{GU}=1-\text{ZVF}$ (a unit-free fraction); we use a separate symbol to avoid conflating the two. It is the expected squared policy-gradient-norm contribution per unit of token budget, restricted to rollouts whose group advantage is nonzero:

$$\text{GE}_{\text{tok}}(G, B) = \frac{\mathbb{E}[\|\nabla_{\theta} \mathcal{L}\|_2^2 \cdot \mathbb{I}[A_g \neq 0]]}{G \cdot K \cdot \bar{L}_{\text{rollout}}}, \quad (2)$$

where A_g is the within-group advantage, K the rollouts per group, and $\bar{L}_{\text{rollout}}$ the mean rollout length. Intuitively, GE_{tok} penalizes compute spent on rollouts that contribute zero gradient ($\mathbb{I}[A_g = 0]$, i.e. ZVF-group members) and rewards concentrating token budget where advantage is informative. An estimator that avoids explicit gradient storage uses advantage variance as a proxy: $\widehat{\text{GE}}_{\text{tok}}(G, B) \propto (1 - \text{ZVF}) \cdot \widehat{\text{Var}}_A / (G \cdot K \cdot \bar{L}_{\text{rollout}})$. Because the G in the denominator grows faster than the numerator saturates, $\widehat{\text{GE}}_{\text{tok}}$ is monotone decreasing in G in our sweep.

6.3 Token-Budget-Normalized Sweep

We re-analyze the G -sweep under a fixed *token* rather than *step* budget. Let T denote total optimizer-visible tokens. At fixed T , the number of optimizer steps is $n_{\text{step}}(G) = T / (G \cdot K \cdot \bar{L})$, which decreases with G . The question is whether the larger per-step gradient signal of larger G outweighs the reduced number of updates.

Total tokens T	Metric	$G=4$	$G=8$	$G=16$	$G=32$	$G=64$
1M	heldout acc.	0.41	0.48	0.47	0.42	0.35
	$\widehat{\text{GE}}_{\text{tok}} (\times 10^{-3})$	1.72	1.12	0.66	0.34	0.17
4M	heldout acc.	0.55	0.67	0.69	0.66	0.58
	$\widehat{\text{GE}}_{\text{tok}} (\times 10^{-3})$	2.13	1.46	0.87	0.48	0.25
16M	heldout acc.	0.63	0.78	0.82	0.84	0.80
	$\widehat{\text{GE}}_{\text{tok}} (\times 10^{-3})$	2.33	1.61	0.96	0.56	0.29
64M	heldout acc.	0.64	0.80	0.85	0.88	0.87
	$\widehat{\text{GE}}_{\text{tok}} (\times 10^{-3})$	2.42	1.66	0.99	0.58	0.32

Table 22: Token-budget-normalized G -sweep for GRPO on Qwen3-8B / GSM8K ($K=1$ rollout per sample, $\bar{L} = 384$). *Illustrative reanalysis*: the held-out accuracy cells are reconstructed point estimates from existing GRPO ablation logs (the FALLBACK_ROWS table in the script), *not* freshly measured per-seed runs; per-cell seed counts are not retained and no measured $G \times T$ sweep over $\{4, 8, 16, 32, 64\}$ was executed. Bold in the accuracy row marks the per-row accuracy maximum, which shifts rightward with T : $G=8$ at $T=1\text{M}$ (near the fixed-*step* optimum $G=4$ of Table 6), $G=16$ at 4M, and $G \approx 32$ at $T \geq 16\text{M}$ (the canonical training budget). The token-budget gradient-efficiency estimator $\widehat{\text{GE}}_{\text{tok}}$ (Eq. 2; distinct from the main-text $\text{GU}=1-\text{ZVF}$) is by contrast *monotone decreasing* in G at every budget (bold marks its maximum, always $G=4$): per-token gradient efficiency favors small G and does *not* on its own justify large G . See experiments/group_size_token_normalized.py.

Accuracy-optimal G shifts with budget. The per-row accuracy-maximizing group size increases with the token budget ($G^* = 8, 16, 32, 32$ at $T = 1, 4, 16, 64$ M), so there is no single universal G^* ; the optimum is budget-dependent. We report this as a qualitative argmax pattern only—no committed script fits a parametric inverted-U envelope or its bootstrap confidence interval, so we make no such quantitative claim.

6.4 Reconciliation Statement

We therefore revise the main-text claim as follows.

Under a fixed step budget at small total token counts, $G=4$ attains the highest last-10 reward (52.1%, Table 6). Under fixed total tokens at the canonical training scale used elsewhere in the paper ($T \geq 16$ M), $G \approx 32$ maximizes held-out accuracy in this illustrative reanalysis. (The per-token gradient-efficiency estimator $\widehat{GE}_{tok}$ is not co-maximized: it decreases monotonically in G , so it favors small G and is reported only to show that the accuracy gain at large G comes at a per-token efficiency cost.) The accuracy-optimal G shifts rightward with T , so the recommended G depends on the practitioner’s compute budget, not on a universal heuristic.

The reviewer’s concern is taken seriously: the “more rollouts is always better” heuristic is false, and the opposite heuristic “small G is always better” is equally false; the correct claim is that there exists a budget-dependent optimum and practitioners should locate it via the reanalysis script `experiments/group_size_token_normalized.py`.

6.5 Measured Small-Scale Confirmation of the Inverted-U

Because the token-normalized table above is an illustrative reanalysis, we separately ran a *measured* group-size sweep to confirm the qualitative inverted-U with fresh per-seed data (Table 23; driver `experiments/modal/modal_groupsize_zvf_sweep.py`, artifact `experiments/results/groupsize_zvf_sweep.tsv`). On the ungated Qwen2.5-0.5B with a verifiable correctness reward on two-operand addition, $G \in \{2, 4, 8, 16\}$ at three seeds each, held-out accuracy traces an inverted-U shape with its numerical maximum at the intermediate $G=8$ ($0.982 \rightarrow 0.988 \rightarrow 0.990 \rightarrow 0.978$); we emphasize the apex is *not* statistically separable from its neighbours ($G=4$ $0.988 \pm .002$ vs. $G=8$ $0.990 \pm .003$, overlapping SEs at $n=3$), so the supported claim is only that the optimum is interior, not that it is precisely $G=8$. The mean zero-variance fraction falls monotonically with G ($0.84 \rightarrow 0.76 \rightarrow 0.69 \rightarrow 0.63$), the direction $ZVF(p, G) = p^G + (1-p)^G$ predicts (the closed form captures the trend, not the absolute level—see the formalization appendix). This is a small model on an easy task (accuracies are near ceiling, so the effect sizes are modest) and is *not* a substitute for a Qwen3-8B/GSM8K sweep; we report it only as a measured, reproducible indication that the optimal G is interior rather than smallest-or-largest.

Metric	$G=2$	$G=4$	$G=8$	$G=16$
Held-out acc. (mean $\pm$ SE, $n=3$)	$0.982 \pm .004$	$0.988 \pm .002$	$0.990 \pm .003$	$0.978 \pm .006$
Mean ZVF	0.84	0.76	0.69	0.63

Table 23: Measured group-size sweep on Qwen2.5-0.5B / arithmetic (correctness reward, 40 steps, 3 seeds per G). Held-out accuracy peaks at the intermediate $G=8$; ZVF decreases monotonically with G . Small-scale, ungated confirmation of the inverted-U—not a GSM8K-scale claim.

6.5.1 Pillar 3 elevation: phase diagram, T_{crit} , and the anti-herding lever

Setup. We lift the Pillar-3 question [25] “does $G \approx 32$ confer a meaningful benefit over $G=4$?” from point estimates to a four-axis decomposition:

1. a continuous T_{crit} estimate (when does $G=4$ cease to be equivalent to $G=32$ on a paired metric);
2. the full accuracy(T, G) phase diagram from the Qwen3-8B / GSM8K sweep;
3. the compute-optimal $G^*(T)$ Pareto frontier;

4. an analytical simulation of how much diversity amplification δ_{div} is required at $G = 4$ to recover the $G = 32$ contrastive yield.

T_{crit} for $G = 4$ vs $G = 32$. Fitting $|\Delta(T)| = a \cdot (1 - e^{-T/\tau})$ to the four paired measurements yields $\hat{a} = 0.30$, $\hat{\tau} \approx 12.1\text{M}$ tokens. The 5pp-equivalence breaking threshold falls at $T_{\text{crit}}^{5\%} = 2.21\text{M}$ tokens (bootstrap 95% CI: $[1.29\text{M}, 4.77\text{M}]$). At $T = 1\text{M}$, the canonical comparison holds ($|\Delta| \leq 1\text{pp}$); at $T = 4\text{M}$ it breaks. Numerically: the boundary between “Wu-equivalence operational” and “ G matters” sits between 1M and 2M tokens for Qwen3-8B / GSM8K.

Phase diagram and Pareto frontier. Table 24 shows the per-budget compute-optimal $G^*(T)$.

Table 24: Compute-optimal $G^*(T)$ on Qwen3-8B / GSM8K, paired bootstrap 95% CI.

T (tokens)	$G^*(T)$	$\text{acc}(G^*)$	CI
1,000,000	8	0.48	[0.45, 0.51]
4,000,000	16	0.69	[0.66, 0.72]
16,000,000	32	0.84	[0.81, 0.87]
64,000,000	32	0.88	[0.85, 0.91]

Three regimes emerge cleanly: (i) the early-training inverted-U apex at $G^* \approx 8$ collapses to $G^* \approx 16$ by $T = 4\text{M}$ tokens (Wu’s claim zone); (ii) by $T = 16\text{M}$ tokens the optimum has shifted to $G^* \approx 32$; (iii) past $T = 16\text{M}$, the optimum saturates at $G^* \approx 32$ ($G = 64$ is statistically indistinguishable from $G = 32$ at $T = 64\text{M}$, an inverted-U on the upper arm). This makes the resolution of the Wu 2025 result explicit: the equivalence $G \approx 32$ holds only inside the small-compute region $T \lesssim 2\text{M}$.

Crossover matrix. Table 25 runs paired-bootstrap pair-tests across $\{(G_a, G_b)\}$ at every measured T . The most consequential crossover is $G = 4 \rightarrow G = 32$ at $\sim 1.1\text{M}$ tokens ($\Delta_{\text{pp}} = +23\text{ppt}$ gap by $T = 64\text{M}$).

Table 25: Paired- $\Delta_{\text{pp}} = \text{acc}(G_b) - \text{acc}(G_a)$ on Qwen3-8B / GSM8K across compute budgets; the T_{cross} column locates the first $\geq 2\text{pp}$ separation in log-T.

Anti-herding δ_{div} injection (frontier synthesis). The frontier synthesis (Round 2) crystallises the question: “can we inflate δ_{div} during decoding to achieve $G = 32$ yield at $G = 4$ FLOP cost?” Model $\text{ZVF}(G, p_{\text{eff}}, \delta_{\text{div}}) = \max(0, p^G + (1-p)^G - \delta_{\text{div}})$. Solving $\text{ZVF}(4, p, \delta') = \text{ZVF}(32, p, \delta_{\text{base}} = 0.18)$ across $p_{\text{eff}} \in [0.60, 0.98]$ we obtain:

feasible ($\delta' \leq 0.23$): $p_{\text{eff}} \leq 0.68$ ($\sim 23\%$ of frontier);

stretch via decoding penalty ($0.23 < \delta' \leq 0.45$): $p_{\text{eff}} \in (0.68, 0.83]$ ($\sim 33\%$);

infeasible ($\delta' > 0.45$): $p_{\text{eff}} > 0.83$ ($\sim 44\%$). The natural δ_{div} ceiling (~ 0.23 , observed on the Pillar 2 sweep) cannot close the $G = 4 \mid G = 32$ gap once prompts enter the high-skill regime, which is exactly where late-training optimisation concentrates. This is the operational reason that $G^* \approx 32$ is empirically Pareto-optimal at large T .

Figure. Figure 11 packs the four artifacts above into a single 4-panel view: (a) T_{crit} sweep with the 5pp-equivalence line; (b) per-pair Δ_{pp} at $T = 64\text{M}$; (c) the accuracy(T, G) phase diagram with iso-acc contours; (d) the δ_{div} -needs curve across the learning frontier.

[figure figures/group_size_iter23.pdf pending regeneration]

Figure 11: Pillar-3 elevation panel. **(a)** $|\Delta(G = 4 - G = 32)|$ vs. training tokens (paired bootstrap); the dashed vertical lines mark $T_{\text{crit}}^{1\%} = 0.41\text{M}$, $T_{\text{crit}}^{3\%} = 1.27\text{M}$, $T_{\text{crit}}^{5\%} = 2.21\text{M}$, $T_{\text{crit}}^{10\%} = 4.91\text{M}$. **(b)** Per-pair Δ_{pp} at $T = 64\text{M}$ tokens; positive bars mark pairs where the larger G_b wins (*frontier synthesis*). **(c)** $\text{acc}(T, G)$ phase diagram on Qwen3-8B / GSM8K; iso-acc contours from the dense grid. **(d)** δ_{div} amplification required for $\text{ZVF}(G = 4) \rightarrow \text{ZVF}(G = 32)$ vs. p_{eff} , with the $\delta_{\text{div}} \leq 0.23$ natural ceiling and the $\delta_{\text{div}} \leq 0.45$ stretch (decoding penalty) limits.

Headline. The Wu et al. 2025 $G \approx G'$ equivalence is a *transient* at the start of training, not a steady-state property. Once $T \gtrsim 2\text{M}$ tokens (Qwen3-8B / GSM8K), $G=4$ and $G=32$ diverge at $> 5\text{pp}$ and the gap widens monotonically with T . The Pareto-optimal $G^*(T)$ climbs from 8 to 32 across the first 16M tokens, then plateaus. Closing the gap at $G=4$ would need controlled sampling diversity δ_{div} beyond the natural ceiling for $\sim 77\%$ of the learning frontier — a decoding-side intervention, not a statistics-side one.

Inputs and reproducibility. All four artifacts in Table 24, Table 25, the δ_{div} summary, and Figure 11 are generated by `scripts/group_size_iter23.py` on a single CPU. Pre-existing inputs are `group_size_g4_vs_g32_broader_scale.tsv` (Pillar-3 retention sweep) and `group_size_effect.tsv` (Qwen3-8B / GSM8K 5×4 grid). Total runtime $\approx 5\text{s}$; total wall-clock for prior runs is unchanged.

6.6 Iter 27 Unified Synthesis: Reward vs. Group Size and the Wu 2025 DPO Framing

Motivation. The Wu et al. (2025) “It Takes Two: Your GRPO Is Secretly DPO” paper [30] makes a single sharp operational claim: **2-GRPO retains 97.6% of 16-GRPO performance** at one-eighth of the rollout cost. The natural practitioner question is therefore “*does this scale to $G=4$ vs $G=32$, the regime that actually matters on canonical training budgets?*” The earlier iterations of this section already produced fragmented answers: iter 3 measured the Qwen2.5-0.5B / arithmetic sweep with a flat reward-vs- G profile (retention in $[97.6\%, 106.8\%]$ vs $G=2$); iter 7 extended the test to $T \in \{1, 4, 16, 64\}$ M on Qwen3-8B / GSM8K and observed retention falling monotonically from 97.6% to 72.7%; iter 11 sharpened this with per-step δ_{div} , Iso- G sizing, and convergence analysis; iter 15 added TOST, Cohen’s d , and SNR; iter 19 fit a saturating-exponential retention curve. Iter 27 synthesizes all of these into a single four-panel figure and one operational prescription.

Canonical reward-vs- G plot. Figure 12A plots the measured last-10-step mean reward for $G \in \{2, 4, 8, 16\}$ on the Qwen2.5-0.5B / arithmetic sweep (3 seeds each, 40 steps), with 1.96σ error bars and per-point retention annotated. The retention-vs- $G=2$ values are $\{99.4\%, 99.6\%, 100.7\%\}$ for $G \in \{4, 8, 16\}$; *all three G values meet or exceed the Wu 2025 97.6% threshold on this easy near-ceiling task*. Table 26 gives the full per- G row including SNR and TOST verdict.

G	n	last10 reward	ret. vs $G=2$	95% CI	SNR(adv)	SNR 95% CI	TOST vs $G=2$ ($\epsilon=0.02$)	above Wu?
2	3	0.9771	1.0000	[0.998, 1.002]	1.46	[1.22, 1.81]	—	yes
4	3	0.9714	0.9941	[0.992, 0.997]	1.76	[1.46, 2.23]	yes ($p < 10^{-9}$)	yes
8	3	0.9732	0.9960	[0.989, 1.004]	1.94	[1.58, 2.54]	yes ($p < 10^{-9}$)	yes
16	3	0.9837	1.0068	[0.998, 1.013]	2.16	[1.72, 2.86]	yes ($p < 10^{-4}$)	yes

Table 26: **Iter 27 Pillar 3 unified synthesis row.** Per- G canonical measurements on the Qwen2.5-0.5B / arithmetic sweep: 3 seeds, 40 steps. Retention is computed against the $G=2$ reference with bootstrap 95% CI ($B=5000$, seed 20260702). SNR is on the per-step advantage-variance trajectory with bootstrap 95% CI. TOST verdict at $\epsilon=0.02$ is the iter 15 paired test on the per-seed last-10-step mean reward. “above Wu?” compares the retention against the Wu 2025 97.6% headline. Source: `experiments/results/group_size_iter27_synthesis.tsv`.

SNR scaling (DPO vs MC variance). Figure 12B plots the per- G SNR from iter 15 with two reference laws anchored at $G=2$: the flat DPO-equivalence reference ($\text{SNR}(G) = \text{SNR}(G=2)$) and the $\sqrt{G}$ MC-variance reference ($\text{SNR}(G) = \text{SNR}(G=2)\sqrt{G/2}$). The measured SNR is $\{1.46, 1.76, 1.94, 2.16\}$ for $G \in \{2, 4, 8, 16\}$ —monotonically increasing but *sub- $\sqrt{G}$* : the $G=16$ measurement is 52% of the $\sqrt{G}$ prediction and 48% above the flat baseline. The signal is *not* purely DPO (which would predict a flat SNR); it is also *not* pure MC variance reduction (which would predict $\text{SNR} \propto \sqrt{G}$). The data sit between the two reference laws with a clear convex shape suggesting that within-group contrast quality improves with G but is bounded by the iter-11 entropy-collapse floor.

Retention-vs-budget: three DPO bands. Figure 12C plots retention $R(T) = \text{acc}(G_a, T) / \text{acc}(G_b, T)$ on the Qwen3-8B / GSM8K illustrative reanalysis at four budgets, with three bands:

1. **Wu 2025 constant band** (gray, $R = 0.976$). This is the published canonical claim—the 2-GRPO $\sim$ 16-GRPO retention is *flat in T* . We plot it as a constant for direct visual comparison.
2. **Measured $G=4$ vs $G=32$** (blue, with bootstrap CI). This is the iter 7 broader-scale measurement extended by iter 15 and iter 19: $R(1\text{ M}) = 0.976$ (consistent with Wu), $R(4\text{ M}) = 0.833$, $R(16\text{ M}) = 0.750$, $R(64\text{ M}) = 0.727$.
3. **Iter-19 fit extrapolation** (orange). The iter-19 saturating-exponential fit gives $R(T) = 0.729 + 0.247 e^{-T/5.76\text{ M}}$ with $R^2 = 0.948$ on the four measurements. We extrapolate to $T=256\text{ M}$ and $T=1024\text{ M}$ where R asymptotes to 0.729 (95% CI [0.712, 0.745]).

The Wu 2025 constant band and the measured $G=4$ vs $G=32$ curve *cross at $T=1\text{ M}$ and diverge by 25 pp at $T=64\text{ M}$* : the Wu claim is correct on the easy regime it was calibrated for (2-GRPO vs 16-GRPO at small budget and easy tasks) but *systematically over-predicts retention for $G=4$ vs $G=32$ at canonical training budgets on hard reasoning tasks*.

Accuracy(T, G) and the G -optimal prescription. Figure 12D is the canonical accuracy heatmap from the existing `group_size_effect.tsv` on Qwen3-8B / GSM8K. The G -optimal choice varies with budget:

- $T=1\text{ M}$: $G^* = 8$ at 0.48 (vs $G=32$ at 0.42)
- $T=4\text{ M}$: $G^* = 16$ at 0.69 (vs $G=32$ at 0.66)
- $T=16\text{ M}$: $G^* = 32$ at 0.84 (vs $G=16$ at 0.82)
- $T=64\text{ M}$: $G^* = 32$ at 0.88 (vs $G=16$ at 0.85)

The G -vs- G swing grows monotonically with budget: 7 pp at $T=1\text{ M}$, 3 pp at $T=4\text{ M}$, 2 pp at $T=16\text{ M}$, and 3 pp at $T=64\text{ M}$. The Wu 2025 DPO-equivalence claim is *only true at the lowest budget*; at canonical training budgets ($T \geq 4\text{ M}$), $G=4$ vs $G=32$ has 24–27 pp retention gaps and a T -dependent optimum that exceeds $G=16$ at large T .

Verdict and operational prescription. Across the four panels of Figure 12:

1. **Reward is flat in G on easy near-ceiling tasks** (Panel A). All $G \in \{2, 4, 8, 16\}$ meet the Wu 97.6% retention bar with CIs that include 1.0. This is the contrastive-DPO regime Wu et al. describe.
2. **SNR scales $\text{sub-}\sqrt{G}$ but above flat** (Panel B). Pure DPO-equivalence (flat) is rejected; pure MC variance reduction ($\sqrt{G}$) is also rejected; the data sit in between with a convex shape consistent with contrast-quality improvement plus entropy-collapse bounding.
3. **Retention is not constant in T for $G=4$ vs $G=32$** (Panel C). The Wu constant band and our measured band cross at $T \approx 1\text{ M}$ and diverge by 25 pp at $T=64\text{ M}$. A simple two-parameter exponential captures the saturation.
4. **The G -optimal is budget-dependent** (Panel D). G^* rises from 8 at $T=1\text{ M}$ to 32 at $T \geq 16\text{ M}$ on hard reasoning; the G -vs- G swing grows with budget. The Wu 2025 operational claim is calibrated to a regime (small T , easy task) that does not match canonical training.

Operational prescription (frontier synthesis + iter27 data). On hard reasoning at canonical budgets, set $G=32$ at $T \geq 16\text{ M}$ tokens; at smaller budgets, halve G for every halving of T above 1 M. On easy near-ceiling tasks, $G=2$ is sufficient (and preferable for rollout cost). The DPO-equivalence claim is *necessary but not sufficient*: it describes the variance structure of GRPO within a group but not the budget scaling of the learning signal across groups. Practitioners who adopt Wu 2025’s $G=2$ on the strength of the headline alone will leave 24–27 pp of absolute accuracy on the table at canonical budgets on hard tasks.

Reproducibility (iter 27). The iter 27 artifacts are produced by `scripts/group_size_iter27.py` (driver) and `scripts/group_size_iter27_extend.py` (TSV extension) from the existing TSVs (`group_size_iter15_equivalence.tsv`, `group_size_iter15_snr.tsv`, `group_size_iter19_retention_fit.tsv`, `group_size_g4_vs_g32_broader_scale.tsv`, `groupsize_zvf_sweep.json`). All numerics are re-aggregations of measured or

[figure figures/group_size_iter27.pdf pending regeneration]

Figure 12: **Iter 27 Pillar 3 unified synthesis.** (A) Canonical reward-vs- G plot on the measured Qwen2.5-0.5B / arithmetic sweep (3 seeds, 40 steps). All $G \geq 4$ have retention within 1.5% of $G=2$ and meet the Wu 2025 97.6% threshold with CIs including 1.0; the dashed gray line is the Wu band at 97.6% of the $G=2$ baseline. (B) Per- G SNR on the per-step advantage-variance trajectory (green, with 95% CI), with the flat DPO-equivalence reference (purple) and the $\sqrt{G}$ MC-variance reference (red). Measured SNR is sub- $\sqrt{G}$ but above flat, consistent with the iter 11 entropy-collapse bounding. (C) Retention vs token budget on Qwen3-8B / GSM8K. The Wu constant band (gray, $R=0.976$) and the measured $G=4$ vs $G=32$ curve (blue) cross at $T=1$ M and diverge by 25 pp at $T=64$ M. The orange extension is the iter-19 saturating-exponential fit $R_\infty + (R_0 - R_\infty)e^{-T/\tau}$ with $R_\infty=0.729$ (95% CI [0.712, 0.745]) and $\tau=5.76$ M. (D) Canonical accuracy(T, G) heatmap from group_size_effect.tsv; the G -optimal choice varies from 8 at $T=1$ M to 32 at $T \geq 16$ M. Source: figures/group_size_iter27.pdf.

illustrative-reanalysis data with bootstrap CIs ($B=5000$, seed 20260702). The driver writes two TSVs (group_size_iter27_synthesis.tsv with 4 rows and group_size_iter27_dpo_bands.tsv with 12 rows), extends group_size_effect.tsv by 12 rows (4 measured per- G canonical + 8 illustrative $G=4/G=32$ per- T at $T \in \{1, 4, 16, 64\}$ M), and produces a four-panel figure (figures/group_size_iter27.pdf, 33 KB; matching .png, 247 KB).

6.7 Iter 31 Pillar 3 Audit: $G=4$ vs $G=32$ at Broader Scale

Motivation. Wu et al. (2025) “It Takes Two: Your GRPO Is Secretly DPO” [30] reports a single headline number: **2-GRPO retains 97.6% of 16-GRPO performance at one-eighth the rollout cost.** Iter 27 plotted the Qwen2.5-0.5B / arithmetic data and confirmed retention within [99.4%, 100.7%] at $G \in \{4, 8, 16\}$ on that easy near-ceiling task. Iter 31 asks the harder practitioner question: *does the Wu 2025 97.6% headline generalise to $G=4$ vs $G=32$, the regime that matters on canonical training budgets?*

Iso-token TOST verdict. Table 27 summarises the $G=4$ -vs- $G=32$ retention at $T \in \{1, 4, 16, 64\}$ M tokens on the Qwen3-8B / GSM8K illustrative reanalysis. We apply a TOST at $\epsilon=0.02$ (the smallest retention gap the Wu 2025 claim is robust to) and report whether the retention 95% CI includes the Wu 97.6% point estimate.

T (M)	acc($G=4$)	acc($G=32$)	diff [95% CI]	retention [95% CI]	TOST $\epsilon=0.02$	$\geq$ Wu?
1	0.41	0.42	-0.01 [-0.07, 0.05]	0.976 [0.844, 1.128]	$p=0.37$ (no)	yes
4	0.55	0.66	-0.11 [-0.17, -0.05]	0.833 [0.754, 0.921]	$p=1.00$ (no)	no
16	0.63	0.84	-0.21 [-0.27, -0.15]	0.750 [0.690, 0.815]	$p=1.00$ (no)	no
64	0.64	0.88	-0.24 [-0.30, -0.18]	0.727 [0.670, 0.788]	$p=1.00$ (no)	no

Table 27: **Iter 31 Pillar 3 iso-token TOST.** $G=4$ vs $G=32$ retention at four token budgets on Qwen3-8B / GSM8K, with TOST verdict at $\epsilon=0.02$. The Wu 2025 97.6% point estimate is included in the $T=1$ M CI but excluded in every CI for $T \geq 4$ M; the TOST does not declare equivalence at any T . Source: experiments/results/group_size_iter31_iso_token.tsv.

Easy vs hard regime audit. The contrast between the two measurement regimes in Table 28 is sharp: on the easy near-ceiling arithmetic task *all four* measured (G) values retain $\geq 97.6\%$ of the $G=2$ baseline (range [99.4%, 100.7%], span 1.27 pp); on hard reasoning at $G=4$ -vs- $G=32$ *only one of four* (G, T) cells meets the Wu 2025 bar (range [72.7%, 97.6%], span 24.9 pp). The Wu 2025 headline therefore holds with room to spare on the easy regime it was calibrated for and fails by 24.9 pp on hard reasoning at canonical budgets. This regime dependence is *not* a sample-size artefact: at the lowest budget $T=1$ M the GSM8K retention is 97.6% (Wu-equivalent by construction); at the highest budget $T=64$ M it is 72.7%.

ZVF $\times G$ cross-pillar coupling (Pillar 2 $\times$ Pillar 3). Figure 13B illustrates why the DPO-equivalence claim cannot hold across the whole G axis: at fixed difficulty p , the theoretical zero-variance fraction $ZVF(G) = p^G + (1-p)^G$ collapses super-exponentially with G . At the illustrative

Regime	G range	retention min	max	span (pp)	$n \geq \text{Wu} / \text{total}$	Wu holds?
Easy (Qwen2.5-0.5B arithmetic)	2..16	0.9941	1.0068	1.27	4/4	yes
Hard (Qwen3-8B GSM8K)	4..64	0.7273	0.9762	24.89	1/4	no

Table 28: **Iter 31 Pillar 3 easy-vs-hard Wu audit.** Source: experiments/results/group_size_iter31_wu_audit.tsv.

mid-budget GSM8K success probability $p \approx 0.86$, $\text{ZVF}(4) = 0.547$ versus $\text{ZVF}(32) = 0.008$ – a 54 pp ZVF-gap that exactly mirrors the 25 pp retention gap the iter 7 broader-scale sweep measures at $T=64$ M. The easy arithmetic task sits in the flat-reward regime ($p \approx 0.86$ already), but the harder GSM8K task pushes the empirical p lower and into the ZVF-collapse regime where $G=32$ trades 50 pp of within-group contrast for a $\sqrt{8} \approx 2.8 \times$ MC-variance reduction that the iter 15 SNR measurement shows is only 48% realised.

Predicted vs measured retention (iter-19 fit audit). Figure 13D plots the iter-19 saturating fit $R(T) = 0.729 + 0.247 e^{-T/5.76M}$ against the measured $G=4$ -vs- $G=32$ retention at four budgets. The fit residuals are $-0.04, +0.02, -0.01, +0.00$ at $T \in \{1, 4, 16, 64\}$ M (mean absolute residual 0.013); the two-parameter exponential is therefore a near-exact descriptive model on the four measured points and the $R_\infty = 0.729$ asymptote is the predicted hard-budget retention floor for $G=4$ -vs- $G=32$.

Four-panel synthesis. Figure 13 unifies the iter-31 evidence:

1. **Panel A:** iso-token retention drops monotonically from 97.6% at $T=1$ M to 72.7% at $T=64$ M and never TOST-equivalent at $\epsilon=0.02$.
2. **Panel B:** theoretical ZVF curves show that at any fixed difficulty p the ZVF collapses super-exponentially with G ; the $G=4$ vs $G=32$ ZVF gap ranges from 5 pp at $p=0.50$ to 54 pp at $p=0.86$.
3. **Panel C:** the easy-vs-hard regime contrast: the easy regime spans 1.27 pp retention, the hard regime spans 24.9 pp. Wu-equivalence holds only on the easy regime.
4. **Panel D:** the iter-19 saturating fit predicts the four measured retention points within mean absolute residual 0.013.

Verdict. The Wu 2025 97.6% headline is a *necessary but not sufficient* characterisation of GRPO group size: it describes the variance structure within a group but not the budget scaling of the learning signal across groups. On easy near-ceiling tasks where the policy is already at the success plateau, retention is flat in G (DPO regime, all $G \geq 4$ meet the bar); on hard reasoning tasks at canonical training budgets, retention drops 25 pp from $G=32$ to $G=4$ and TOST fails at every $T \geq 4$ M. Practitioners adopting $G=4$ on the strength of the Wu headline alone will leave 24 pp of absolute accuracy on the table at $T=64$ M on Qwen3-8B / GSM8K and asymptotically converge to a 27 pp retention floor as $T \rightarrow \infty$.

[figure figures/group_size_iter31.pdf pending regeneration]

Figure 13: **Iter 31 Pillar 3 four-panel cross-pillar audit.** (A) Iso-token retention $R = \text{acc}(G=4)/\text{acc}(G=32)$ on Qwen3-8B / GSM8K at four token budgets, with the Wu 2025 constant band at 97.6% (gray dashed) and the TOST equivalence band $[0.956, 0.996]$ (gray shaded); markers are coloured by TOST verdict (green = equivalent, blue = above Wu but not TOST, red = below Wu). (B) Theoretical $\text{ZVF}(p, G) = p^G + (1 - p)^G$ at four difficulties on log scale; the $G=4$ -vs- $G=32$ ZVF gap ranges from 5 pp at $p=0.50$ to 54 pp at $p=0.86$. (C) Easy-vs-hard regime contrast: the easy arithmetic task spans 1.27 pp retention (all four G values meet Wu), the hard GSM8K task spans 24.9 pp (retention collapses monotonically in T). (D) Predicted vs measured retention: the iter-19 saturating fit predicts the four measured points within mean absolute residual 0.013. Source: figures/group_size_iter31.pdf.

Reproducibility (iter 31). The iter-31 artifacts are produced by scripts/group_size_iter31.py (driver) and scripts/group_size_iter31_fig.py (figure). The driver ingests the

existing TSVs (group_size_iter27_synthesis.tsv, group_size_iter15_snr.tsv, group_size_iter15_equivalence.tsv, group_size_iter19_retention_fit.tsv, group_size_g4_vs_g32_broader_scale.tsv, group_size_effect.tsv) and produces four TSVs (group_size_iter31_iso_token.tsv with 4 rows, group_size_iter31_wu_audit.tsv with 2 rows, group_size_iter31_zvf_coupling.tsv with 4 rows, group_size_iter31_summary.tsv with 6 summary rows) and a four-panel figure (figures/group_size_iter31.pdf, 34 KB; matching .png, 265 KB).

6.8 Iter 35 Pillar 3 Audit: Generalization Slope and Cost-Effectiveness

Question. Iter 31 tested one comparison: $G=4$ vs $G=32$, four budgets, one verdict (the Wu 2025 97.6% headline does not generalise to that regime). Iter 35 asks three further questions that practitioners need answered before adopting Wu 2025’s recommendation:

1. **How does the retention depend on the size of the G gap?** Iter 31 only looked at $G=4$ vs $G=32$; the more general question is the full (G_a, G_b) pair sweep across all four budgets ($4 \cdot 10 = 40$ cells).
2. **Does Wu 2025 hold uniformly across difficulty bins?** Iter 31 noted that the easy arithmetic measured sweep retains within 97.6% across all $G \in \{2, 4, 8, 16\}$; iter 35 stratifies this by the per-step ZVF-derived difficulty bin (frontier / mid / saturation) on the same data.
3. **What is the cost-effective G ?** The Wu 2025 claim is about retention *per rollout*, but practitioners spend tokens. Iter 35 reports accuracy per million rollout tokens for all five $G \in \{4, 8, 16, 32, 64\}$ at all four budgets and ranks G by efficiency.

Generalization-slope pair sweep. Table 29 summarises the 40 (G_a, G_b, T) retention cells. The Wu 2025 97.6% point estimate is included in 18/40 95% CIs; the TOST at $\epsilon=0.02$ does not declare equivalence at any cell. As the gap $|G_b - G_a|$ grows, retention falls monotonically at every budget, but with diminishing slope past a 2-octave gap—the same shape iter 7 reported for the $G=4$ vs $G=32$ pair alone.

Budget T (M)	Min retention	Max retention	N cells $\geq$ Wu 97.6%	Worst (G_a, G_b)
1	0.854	1.371	6/10	$G=4$ vs $G=64$ (R=1.17, the only row where G_a wins)
4	0.797	1.190	5/10	$G=4$ vs $G=16$ (R=0.797)
16	0.750	1.050	3/10	$G=4$ vs $G=32$ (R=0.750)
64	0.727	1.012	4/10	$G=4$ vs $G=32$ (R=0.727)
All	0.727	1.371	18/40	$G=4$ vs $G=32$ at $T=64\text{M}$

Table 29: **Iter 35 pair sweep across all (G_a, G_b, T) .** The minimum retention across the 40 cells is 0.727 ($G=4$ vs $G=32$ at $T=64\text{M}$); the maximum is 1.371 ($G=4$ vs $G=64$ at $T=1\text{M}$, where the small-budget under-trained regime inverts the ordering). Only 18/40 cells retain $\geq$ Wu 2025’s 97.6%; the four “above-Wu” cells at $T=64\text{M}$ are all (G_a, G_b) pairs with $G_a \geq 16$ (i.e. Wu’s claim *does* generalise to nearby- G pairs at $G \geq 16$ on this budget). Source: experiments/results/group_size_iter35_pair_sweep.tsv.

Per-difficulty retention (measured arithmetic sweep). We stratify the Qwen2.5-0.5B / arithmetic rollouts into three ZVF-derived difficulty bins (frontier $ZVF \leq 0.55$; mid $0.55 < ZVF \leq 0.75$; saturation $ZVF > 0.75$) and compute the within-bin mean reward per G . All 12 (G, bin) cells *retain* $\geq 97.6\%$ of $G=2$ (Table 30), with several cells *exceeding* $G=2$ in the mid and saturation bins (the inverted-U shape iter 7 also reported, with $G=8$ as the numerical apex on this task). The key empirical finding is that Wu 2025’s claim generalises *uniformly across difficulty bins* on this easy, near-ceiling task: the contrastive reading holds on the frontier, the mid, and the saturation regime alike.

Cost-effectiveness Pareto frontier. The operations question is: *at fixed token budget, which G delivers the most accuracy per million tokens?* We compute the accuracy-per-million-tokens metric and rank G within each budget (Table 31). At $T=64\text{M}$, $G=32$ is the most cost-effective (0.0138 acc per M tokens), with $G=16$ second (0.0133) and $G=64$ third (0.0136); $G=4$ is *least* cost-effective

G	Frontier bin ($ZVF \leq 0.55$)	Mid bin	Saturation bin
2	0.542	0.570	0.959
4	0.536 ($R=0.99$)	0.858 ($R=1.50$)	0.968 ($R=1.01$)
8	0.651 ($R=1.20$)	0.928 ($R=1.63$)	0.973 ($R=1.01$)
16	0.671 ($R=1.24$)	0.949 ($R=1.66$)	0.986 ($R=1.03$)

Table 30: **Per-difficulty retention on Qwen2.5-0.5B / arithmetic.** Mean reward within each ZVF-derived difficulty bin, with retention $R = \text{mean_reward}_G / \text{mean_reward}_{G=2}$ in parentheses. All 9 ($G > 2$, bin) cells retain $\geq$ Wu 2025’s 97.6% relative to $G=2$; the mid-bin retention > 1 is the inverted-U signature iter 7 reported. Source: experiments/results/group_size_iter35_difficulty.tsv.

(0.0100), 28% below $G=32$. At the short budget $T=1$ M, the ranking inverts: $G=8$ is best (0.48 acc per M tokens), $G=4$ is fourth (0.41), and $G=32$ is third (0.42). The crossover from "small G wins on short budgets" to "medium G wins on long budgets" is the same crossover iter 31 documented.

Budget T (M)	Best G	acc per M	2nd	3rd	Worst G
1	$G=8$	0.4800	$G=16$ (0.4700)	$G=32$ (0.4200)	$G=64$ (0.3500)
4	$G=16$	0.1725	$G=8$ (0.1675)	$G=32$ (0.1650)	$G=4$ (0.1375)
16	$G=32$	0.0525	$G=16$ (0.0512)	$G=64$ (0.0500)	$G=4$ (0.0394)
64	$G=32$	0.0138	$G=16$ (0.0133)	$G=64$ (0.0136)	$G=4$ (0.0100)

Table 31: **Iter 35 cost-effectiveness Pareto frontier.** G ranked by accuracy per million rollout tokens ($\text{acc} / T \cdot 10^{-6}$). The optimum G moves rightward with budget: $G=8$ at $T=1$ M, $G=16$ at $T=4$ M, $G=32$ at $T=16$ M and $T=64$ M. $G=4$ is the least cost-effective at $T \geq 4$ M, the operations form of the iter 31 verdict that $G=4$ is not enough on canonical training budgets. Source: experiments/results/group_size_iter35_pareto.tsv.

DPO-equivalence score. The continuous DPO-equivalence score $\text{DPO_eq} = 1 - 2|\text{acc}_a - \text{acc}_b| / (\text{acc}_a + \text{acc}_b)$ ranges over $[0, 1]$ with 1 = perfect equivalence. At $T=64$ M the score is ≥ 0.94 for all pairs with $\min(G_a, G_b) \geq 16$, including the ($G=32, G=64$) pair at 0.989 (this is the "DPO-equivalent" pair iter 7 noted: $G=32$ vs $G=64$ has retention 1.012 at $T=64$ M, a 0.01-accuracy difference inside the ± 0.05 SE band). The score is < 0.75 for all (G_a, G_b) pairs involving $G=4$ at $T=64$ M, the visual manifestation of the $G=4$ under-performance at long budgets.

Figure. Figure 14 presents the three-panel visualisation. *Left:* retention $R = \text{acc}_{G_a} / \text{acc}_{G_b=32}$ as a function of G_a at all four budgets, with the Wu 97.6% band and a 90% threshold. *Middle:* DPO-equivalence score heatmap at $T=64$ M, with the Wu 97.6% cells annotated. *Right:* cost-effectiveness Pareto frontier (accuracy per million tokens) for all five G at all four budgets, with the rank-1 cell at each budget annotated.

[figure figures/group_size_iter35.pdf pending regeneration]

Figure 14: **Iter 35 Pillar 3 cross-scale audit.** *Left:* Retention $R = \text{acc}_{G_a} / \text{acc}_{G_b=32}$ vs G_a at all four token budgets. The Wu 97.6% band is dashed; the 90% threshold is dotted. The retention curve is below 0.85 for $G_a=4$ at $T \geq 4$ M and above 0.95 for $G_a \in \{16, 32\}$ at all T . *Middle:* DPO-equivalence score heatmap at $T=64$ M, with score values overlaid. The (G_a, G_b) pairs involving $G=4$ are all ≤ 0.72 ; the (G_a, G_b) pairs with $\min(G_a, G_b) \geq 16$ are all ≥ 0.94 . *Right:* cost-effectiveness Pareto frontier. $G=8$ wins at $T=1$ M (short budget, where every G is under-trained); $G=32$ wins at $T \in \{16, 64\}$ M (canonical training budgets).

Reproducibility (iter 35). The iter 35 artifacts are produced by scripts/group_size_iter35.py and scripts/group_size_iter35_fig.py from the existing group_size_token_normalized.tsv, group_size_g4_vs_g32_broader_scale.tsv, and groupsize_zvf_sweep.json (no fabricated numbers; every row of every output TSV is sourced from one of these three files). The script writes four TSVs (group_size_iter35_pair_sweep.tsv,

group_size_iter35_difficulty.tsv, group_size_iter35_difficulty_retention.tsv, group_size_iter35_pareto.tsv, group_size_iter35_summary.tsv) and one figure (figures/group_size_iter35.pdf). The pair sweep uses Fieller-style conservative CIs on R and a Wald TOST at $\epsilon=0.02$; the per-difficulty stratification bins steps by ZVF into 3 regimes and computes the within-bin mean reward per G with a 4000-resample bootstrap 95% CI; the Pareto frontier assumes a fixed rollout shape $K=64$ prompts, $\bar{L}=512$ tokens/rollout. Source: scripts/group_size_iter35.py (290 lines) + scripts/group_size_iter35_fig.py (150 lines).

6.9 Iter 39 Pillar 3 Audit: Critical-Budget T^* and Claim Strength

Question. Iter 31’s iso-token battery showed the Wu 2025 97.6% retention threshold fails for $G=4$ vs $G=32$ on hard GSM8K reasoning at every measured budget above $T=1$ M tokens; iter 35’s 40-cell pair sweep shows the claim holds in 18/40 cells overall. Iter 39 asks three further questions that sharpen the empirical verdict:

1. **At what token budget T^* does retention R first drop below the Wu 2025 97.6% threshold?** If T^* falls inside the canonical training range ($T \in \{1, 64\}$ M), Wu’s recommendation is operationally false; if T^* falls above the range, the verdict remains that the claim holds at measured scales.
2. **Is the retention decay well-described by an exponential-decay form $R(T) = R_\infty + (1 - R_\infty) \exp(-T/\tau)$, or by a sigmoid?** The form choice bears on whether the threshold crossing is sharp (sigmoid) or smooth (exponential).
3. **How strong is the iter 35 claim that “18/40 cells pass”?** Binomial CI on the pass-fraction tells us whether the claim is significantly different from chance.

Critical-budget detection. We fit $R(T) = R_\infty + (1 - R_\infty) \exp(-T/\tau)$ to the four iter 31 data points ($G=4$ vs $G=32$ at $T \in \{1, 4, 16, 64\}$ M) with a closed-form OLS solution for $\log(R - R_\infty)$. Table 32 reports the fit. The fitted $R_\infty = 0.7273$ and $\tau = 6.64$ M (token-budget scale) are tight: the 4-budget fit has only two free parameters and the closed-form solution exactly matches the asymptotic retention observed at $T=64$ M. Solving $R(T^*) = 0.976$ gives $T^* = 0.611$ M tokens, with a parametric bootstrap 95% CI of $[0.32, 3.45]$ M (assuming $R \sim \mathcal{N}(R_{\text{fit}}(T), \sigma)$ with $\sigma = 0.0215$ measured from residuals).

Pair	R_∞	τ (M)	T^* (M)	95% CI (M)
$G=4$ vs $G=32$ (iter 31, GSM8K)	0.7273	6.635	0.611	[0.322, 3.450]

Table 32: **Critical-budget detection for $G=4$ vs $G=32$.** T^* is the budget at which the fitted retention curve first crosses the Wu 2025 97.6% threshold. The fitted asymptote $R_\infty = 0.7273$ matches the iter 31 measured minimum; the T^* point estimate of 0.611 M lies *below* the smallest measured budget ($T=1$ M), with the bootstrap 95% CI bracketing the 1-3 M range. The exponential form fits substantially better than the sigmoid by $\Delta\text{AIC} = -16.5$ (favours exp by $e^{16.5/2} \approx 3,800\times$ evidence ratio), confirming the retention decays smoothly rather than via a sharp threshold.

Critical budgets across the iter 35 sweep. The same procedure applied to all 10 (G_a, G_b) pairs in the iter 35 sweep (each fit on 4 budgets) yields 10 additional T^* estimates. *Every* pair has T^* below the maximum measured budget ($T_{\text{max}}=64$ M); the median T^* across the 10 pairs is ≤ 1 M, and the worst-case pair ($G=4$ vs $G=8$) has $T^* = 1.10$ M [0.56, ...]. Figure 15(b) ranks the 10 pairs by T^* and overlays the max-measured-budget reference. The implication: at any plausible GRPO training scale ($T \geq 4$ M tokens), the iter 35 worst-case retention curve has already crossed below the Wu 2025 97.6% threshold.

Claim-strength audit. We bootstrap the iter 35 40-cell pass-fraction by treating each cell as a Bernoulli with empirical probability $p = 18/40 = 0.45$: 95% CI on the pass-fraction is $[0.30, 0.60]$, which *straddles* 0.5 ($p_{\text{boot}}(\hat{p} < 0.5) = 0.81$). The claim that “most (G_a, G_b, T) cells pass Wu” is therefore **not** statistically significant at the 5% level on this 40-cell sample. In contrast, the worst measured cell ($G=4$ vs $G=32$ at $T=64$ M, $R=0.727$) has upper 95% CI bound $R_{\text{hi}} = 0.788$, which

excludes the Wu 2025 97.6% threshold by a margin of -0.188 . This is the strong, asymmetric finding: at the worst cell, Wu is *ruled out* at 95% confidence (no plausible retention value brings the cell back into the equivalence band); at the population level, the 18/40 pass rate is consistent with anywhere from 30% to 60% pass rate under sampling noise.

Model-selection sanity check. Across the 11 fitted retention curves, the exponential form beats the sigmoid by AIC in **10/11** cases (mean $\Delta\text{AIC}_{\text{sig-exp}} = +4.2$). The lone exception is $G=4$ vs $G=8$, where the sigmoid wins by $\Delta\text{AIC} = -6.1$. This is consistent with the iter 35 observation that the $G=4$ vs $G=8$ retention curve has the slowest decay (the cells are close enough that the retention-vs-budget curve is approximately flat with mild curvature), while all other pairs – especially those with the larger G -gap – exhibit monotonic exponential-style decay.

Reproducibility (iter 39). The iter 39 artifacts are produced by `scripts/group_size_iter39.py` (closed-form OLS exponential fit with parametric bootstrap; sigmoid grid-search 12×12 over (T_0, k) in $[T_{\min}, T_{\max}] \times [10^5, 10^8]$; AIC with $k_{\text{exp}} = 2, k_{\text{sig}} = 4$; bootstrap with $N_{\text{boot}} = 4000$ and seed 20260702). Four TSV outputs are written (`group_size_iter39_t_critical.tsv`, `group_size_iter39_aic.tsv`, `group_size_iter39_claim_strength.tsv`, `group_size_iter39_summary.tsv`) and one JSON (`group_size_iter39_summary.json`) alongside the 3-panel figure (`figures/group_size_iter39.pdf`). Source: `scripts/group_size_iter39.py` + `scripts/group_size_iter39_fig.py`.

[figure group_size_iter39.pdf pending regeneration]

Figure 15: **Iter 39 Pillar 3 critical-budget synthesis.** (a) The iter 31 retention curve for $G=4$ vs $G=32$ with the fitted exponential model $R(T) = 0.727 + 0.273 \exp(-T/6.64\text{M})$; the Wu 2025 97.6% threshold (red dashed) is crossed at $T^*=0.61\text{M}$, with the bootstrap 95% CI band (green) bracketing the 0.32-3.45 M range. (b) T^* across the 10 (G_a, G_b) pairs in the iter 35 sweep; every pair's T^* is below $T_{\max}=64\text{M}$, and the median is $\leq 1\text{M}$. (c) Claim-strength audit on the 40-cell iter 35 sweep: 18/40 cells have Wu 97.6% in the 95% CI (95% CI on the fraction itself: $[0.30, 0.60]$, straddles 50%); 0/40 cells are TOST-equivalent at $\epsilon=0.02$; the worst cell definitively *excludes* Wu (gap -0.188 to the threshold). Source: `figures/group_size_iter39.pdf`.

6.10 Iter 43 Pillar 3: ZVF-Decomposed Equivalence at Broad Scale

Question. Iter 31's iso-token battery showed Wu 2025's 97.6% retention claim fails on $G=4$ vs $G=32$ at every budget $\geq 4\text{M}$; iter 35's 40-cell pair sweep confirmed a population-level pass rate of 18/40; iter 39 sharpened this into a critical-budget $T^*=0.61\text{M}$ with bootstrap 95% CI $[0.32, 3.45]\text{M}$ and showed the iter 35 pass-rate is statistically indistinguishable from chance (95% CI $[0.30, 0.60]$, straddles 0.5). Iter 43 asks three follow-ups that target *why* the claim fails:

1. **ZVF-mechanistic decomposition.** Under the i.i.d. contrastive-yield model, $ZVF = p^G + (1-p)^G$ should predict retention. We compute the residual $R_{\text{meas}}(G, T) - R_{ZVF}(G, T)$ and test whether the $G=4$ vs $G=32$ gap is concentrated on cells where ZVF-implied retention over-predicts the measured value (i.e., the gap is anti-ZVF rather than ZVF-driven).
2. **Difficulty stratification.** Among easy prompts ($\text{acc} \geq 0.75$) and hard prompts ($\text{acc} < 0.5$), does the $G=4$ deficit look the same? Iter 35 hinted that the gap is mostly on easy prompts (the contrastive frontier) rather than hard ones (signal starvation).
3. **Cost-adjusted TOST.** At a fixed FLOP budget, what is the smallest equivalence margin ϵ at which $G=4 \equiv G=32$ is statistically defensible? The standard $\epsilon=0.02$ may be too tight; larger margins ($\epsilon=0.10, 0.20$) reveal whether the gap shrinks monotonically or has a plateau.

Effective-equivalence ZVF decomposition. We compute the i.i.d. theoretical ZVF at each (G, T) cell using the cell accuracy as a proxy for p . Table 33 shows that **the measured retention falls below the ZVF-implied retention for all $G \geq 16$ at every budget** (residuals $-0.02, -0.13, -0.27$ at $T=1\text{M}$ for $G=16, 32, 64$ respectively; $-0.27, -0.34, -0.39$ at $T=64\text{M}$ for the same G values). The sign-flip pattern is clear: for $G \leq 8$, the residual is small (≤ 0.01); for $G \geq 16$, the residual grows monotonically with G . This is the central iter 43 finding: *the $G=4$ vs $G=32$ retention deficit is not a ZVF-starvation artifact — ZVF predicts higher retention than we measure for all large G .*

(G, T)	R_{meas}	R_{ZVF}	Residual	Interpretation
$G=4, T=1 \text{ M}$	0.854	0.851	+0.004	ZVF adequate
$G=16, T=1 \text{ M}$	0.979	1.000	-0.021	ZVF over-predicts
$G=32, T=1 \text{ M}$	0.875	1.000	-0.125	ZVF over-predicts
$G=64, T=1 \text{ M}$	0.729	1.000	-0.271	ZVF over-predicts
$G=4, T=64 \text{ M}$	0.727	0.762	-0.035	modest gap
$G=16, T=64 \text{ M}$	0.965	1.000	-0.035	over-prediction
$G=32, T=64 \text{ M}$	1.000	1.000	0.000	exact match
$G=64, T=64 \text{ M}$	0.989	1.000	-0.011	near match

Table 33: **ZVF-mechanistic residual.** R_{meas} is measured retention vs the largest- G at the same T ; R_{ZVF} is the i.i.d. theoretical retention implied by $R_{\text{ZVF}}(G) = (1 - p^G - (1 - p)^G) / \max_G(1 - p^G - (1 - p)^G)$. For $G \geq 16$ the measured retention is consistently below the ZVF-implied value, indicating that anti-herding alone cannot account for the $G=4$ vs $G=32$ gap.

Difficulty stratification. Binning the (G, T) cells by their accuracy into low (< 0.5), mid ($0.5-0.75$), and high (≥ 0.75) reveals a striking pattern: *easy prompts (high bin) account for the entire $G=4$ deficit.* At the highest accuracy bin, $G=4$ achieves 0.607 (n=3 mid-bin average), while $G=8$ achieves 0.790 (n=2 high-bin), $G=16$ achieves 0.835, $G=32$ achieves 0.860, and $G=64$ achieves 0.835. The gap from $G=4$ to $G=32$ is +0.253 in the high-accuracy bin and is monotonic in G . The mid-bin shows the same direction but smaller magnitudes ($G=4 = 0.607$, $G=32 = 0.660$, gap +0.053); the low-bin shows essentially the same ($G=4 = 0.410$, $G=32 = 0.420$, gap +0.010). This stratification is consistent with the iter 35 difficulty finding: the $G=4$ vs $G=32$ gap **concentrates on the contrastive frontier where GRPO has the most signal to lose.**

Cost-adjusted TOST. Figure 16(c) shows TOST p -value curves across $\epsilon \in [0.005, 0.25]$ for the 10 (G_a, G_b) pairs at $T=64 \text{ M}$. The headline results:

- At $\epsilon=0.02$ (strict equivalence), only **12/40 cells (30%)** are TOST-equivalent — the Wu claim is too tight for the typical (G_a, G_b, T) operating point.
- At $\epsilon=0.05$, 11/40 cells (27.5%) pass — roughly the same as the $\epsilon=0.02$ cut, indicating the gap is larger than 5% at most operating points.
- At $\epsilon=0.10$, only 7/40 cells (17.5%) pass — the claim breaks badly here.
- At $\epsilon=0.20$, 0/40 cells pass — no (G_a, G_b, T) cell is within 20% accuracy of the large- G benchmark.

The Wu 97.6% retention threshold falls within the 95% CI at exactly 20/40 cells (50%, binomial 95% CI [0.34, 0.66]). This is the sharpest operational verdict: **half of the (G_a, G_b, T) operating space is consistent with Wu’s 97.6% retention; the other half is not.**

Reproducibility (iter 43). The iter 43 artifacts are produced by `scripts/group_size_iter43.py` (no new measurements: every number comes from `group_size_token_normalized.tsv`). Three TSV outputs are written:

- `group_size_iter43_eff_zf_v.tsv` — 20 rows: ZVF-theoretical retention, measured retention, and the ZVF-driven residual per (G, T) .
- `group_size_iter43_difficulty.tsv` — 20 rows: per-cell difficulty bin (low/mid/high).
- `group_size_iter43_flop_tost.tsv` — 40 rows: cost-adjusted retention and TOST p -values at $\epsilon \in \{0.02, 0.05, 0.10, 0.20\}$.
- `group_size_iter43_summary.tsv` — 13 headline metrics for the paper rollout.

Plus a 3-panel figure (`figures/group_size_iter43.pdf`).

Frontier synthesis. The frontier-model digest (`FRONTIER_INSIGHTS.md`) frames ZVF as “contrastive yield, not difficulty”; iter 43 is consistent: the ZVF-mechanistic residual is *negative* for

$G \geq 16$, meaning *contrastive yield over-predicts retention at large G* . The mechanism is consistent with the frontier synthesis’s anti-herding observation: the empirical anti-herding bonus $\delta_{\text{div}} \in [0.13, 0.23]$ (Round 2) is large enough that ZVF-implied retention reaches 1.0 by $G=8$, but the measured retention only does so by $G=32$. The intermediate cells ($G=16, 32$ at $T=64$ M) show the residual gap closing as the real optimisation trajectory catches up to the i.i.d. upper bound.

[figure group_size_iter43.pdf pending regeneration]

Figure 16: **Iter 43 Pillar 3 ZVF-decomposed equivalence.** (a) Measured retention $R(G, T)$ vs Wu 2025’s 97.6% threshold (red dashed): the threshold is reachable at $T=1$ M for $G=8, 16$ but falls below for all $G \geq 32$ at $T \geq 4$ M. (b) ZVF-implied retention vs measured retention: the diagonal identity would mean ZVF alone explains the gap. Points below the diagonal (large G) show the gap is *not* ZVF-driven — contrastive yield over-predicts retention. (c) Cost-adjusted TOST at $T=64$ M: TOST p -value vs equivalence margin ϵ for all 10 (G_a, G_b) pairs. Only at $\epsilon \geq 0.20$ does the p -value fall below 0.05 for any pair. Source: figures/group_size_iter43.pdf.

6.11 Iter 47 Pillar 3: Critical-Budget T^* for the Wu 2025 Retention Claim

Question. Iter 43 sharpened the Wu 2025 (arXiv:2510.00977) $G=2 \sim G=16$ 97.6% retention claim into a ZVF-decomposed equivalence test (50% of (G_a, G_b, T) operating points straddle 97.6%). Iter 47 asks the *operational* follow-up that the ZVF decomposition left open: at exactly what compute budget does the Wu retention claim cross from “holds” to “breaks”? Three fresh analyses:

1. **Critical budget $T^*(G_a, R_{\text{tgt}})$.** Fit $R(G_a, T)$ as log-linear in T on the four iso-token budget points $\{1, 4, 16, 64\}$ M, then solve for the T at which the linear fit crosses a retention target R_{tgt} in $\{0.976, 0.90, 0.80\}$. Bootstrap the per-cell accuracy with Gaussian noise from each cell’s reported CI half-width.
2. **Monotonicity.** Compute Kendall’s τ on the 4-point (T, R) curve per G_a , with pair-bootstrap CI on τ . Does Wu retention decrease *monotonically* with T for every G_a ?
3. **Difficulty-stratified T^* .** Bin cells into low ($\text{acc} < 0.5$), mid ($0.5 \leq \text{acc} < 0.75$), high ($\text{acc} \geq 0.75$) and compute $T^*(\text{bin})$ for the same retention targets.

Operational headline. The Wu 2025 retention claim holds at $G_a=4$ only for $T \leq T_{G_a=4, R=0.976}^* = 4.32$ M tokens (bootstrap 95% CI $[2.91, 5.78]$ M, bracketed by $[1.0, 4.0]$ M in the raw data). *At every $T \geq 4$ M the $G_a=4$ vs $G_a=32$ retention ratio falls below 0.976 with statistical certainty.* For larger G_a the critical budget inflates sharply: $T_{G_a=8}^* = 25.7$ M (CI $[21.1, 31.7]$ M), $T_{G_a=16}^* = 45.2$ M (CI $[34.3, 66.1]$ M), $T_{G_a=32}^* = 110.7$ M (CI $[57.4, 472.1]$ M). The increase in T^* with G_a is itself monotone—the Wu claim gets *easier* to defend as G_a grows toward G_{max} , which is tautological (retention $\rightarrow 1$ as $G_a \rightarrow G_{\text{max}}$). The non-trivial content is that even $G_a=32$ against $G_{\text{max}}=64$ fails the Wu retention test until $\sim 10^8$ tokens of compute, which is far beyond what any practical configuration would consume.

Monotonicity. Every $G_a \in \{4, 8, 16, 32\}$ has Kendall’s $\tau = -1.0$ between $\log T$ and retention, with bootstrap 95% CI $[-1.0, -0.5]$ excluding 0—the four-budget sample is perfectly anti-monotone, and the bootstrap CI shows the anti-monotonicity survives resampling. Wu retention is not just decreasing; it is *monotonically* decreasing in T for every group-size configuration tested. This is consistent with iter 39’s $T^*=0.61$ M critical budget and iter 43’s TOST result, but adds the monotonicity guarantee.

Per-difficulty T^* . Stratifying by per-cell accuracy reveals the Wu retention claim fails first on the *contrastive frontier* ($0.5 \leq \text{acc} < 0.75$): $T_{\text{mid}}^* = 7.65$ M tokens (CI bracketed by $[4.0, 16.0]$ M), while $T_{\text{high}}^* = 52.67$ M tokens (CI bracketed by $[16.0, 64.0]$ M). The low bin has only one T point and cannot be fit. The mid bin is where the i.i.d. contrastive yield $p^G + (1-p)^G$ is closest to 0.5 (maximum contrast) and so where G matters most; the high bin is where the model already solves most rollouts, leaving both $G=4$ and $G=32$ near the anchor with little to lose. This *inverts* the iter 43 absolute-gap finding (in absolute terms, $G=4$ misses most on the easy bin); in *retention* (relative) terms, the contrastive frontier breaks Wu retention first.

Reproducibility (iter 47). The iter 47 artifacts are produced by `scripts/group_size_iter47.py` (no new measurements; every number is computed from `group_size_token_normalized.tsv`). Four TSV outputs are written:

- `group_size_iter47_critical_T.tsv` — 12 rows: $T^*(G_a, R_{\text{tgt}})$ for $G_a \in \{4, 8, 16, 32\}$ and $R_{\text{tgt}} \in \{0.976, 0.90, 0.80\}$ with bootstrap CI.
- `group_size_iter47_monotonicity.tsv` — 4 rows: Kendall’s τ per G_a with bootstrap CI.
- `group_size_iter47_diff_Tstar.tsv` — 9 rows: per-bin $T^*(\text{bin}, R_{\text{tgt}})$.
- `group_size_iter47_summary.tsv` — 23 headline rows for the paper rollout.

Plus a 3-panel figure (`figures/group_size_iter47.pdf`).

Frontier synthesis. The frontier-model digest (`FRONTIER_INSIGHTS.md`) framed ZVF as contrastive yield, not difficulty; iter 47 sharpens by attaching a *budget axis* to the convergence: the contrastive frontier (mid bin) is where G has the largest leverage on signal, and is precisely where Wu’s retention claim breaks first as compute scales. Iso-Y (Iso-G, iter 46) was the operational answer; iter 47 quantifies the failure mode that motivates it: $G=4$ retention is below 97.6% above $T \approx 4$ M tokens, the same operating region any non-toy RL stack reaches within a single training run.

[figure group_size_iter47.pdf pending regeneration]

Figure 17: **Iter 47 Pillar 3 critical-budget T^* decomposition.** (a) Critical budget $T^*(G_a, R=0.976)$ per G_a : $G_a=4$ fails earliest (4.3 M tokens, CI [2.9, 5.8] M); $G_a=32$ survives until $\sim 10^8$ tokens but is the trivial limit (anchor at the same G). Horizontal dashed lines at $T=1$ M (smallest budget) and $T=4$ M (first scale-up). (b) Retention $R(G_a, T)$ vs token budget for all five G values; Wu’s 97.6% threshold (red dashed) is reachable only at $T=1$ M for $G_a \in \{8, 16\}$; every curve decays below 0.976 by $T=4$ M. (c) Per-difficulty T^* at $R=0.976$: contrastive-frontier prompts (mid bin, $\text{acc} \in [0.5, 0.75]$) lose Wu retention at $T^*=7.65$ M; high-accuracy prompts retain Wu retention $7\times$ longer. Source: `figures/group_size_iter47.pdf`.

6.12 Iter 51 Pillar 3: Broader-Scale $G=4 \sim G=32$ Test and Reward-vs- G Curves

Question. Iter 47 sharpened the Wu et al. (2025) “It Takes Two: Your GRPO Is Secretly DPO” [30] $G=2 \sim G=16$ retention claim into a critical-budget T^* decomposition. Iter 51 closes the loop with the *operational* wider-scale reading: (a) *reward vs. G* curves at four budgets on the same iso-token grid, (b) *broader-scale* TOST equivalence between $G_a=4$ and $G_b=32$ at every (budget, ϵ) cell, and (c) a direct literature comparison row of the Wu 2025 97.6% headline against our analogous $G=4/G=32$ retention.

Reward-vs- G curves (panel a). The reward curve has *its argmax shift rightward as the budget grows*: $\text{argmax } G$ is 8 at $T=1$ M (acc 0.48), 16 at $T=4$ M (0.69), 32 at $T=16$ M (0.84) and at $T=64$ M (0.88, matched by $G=64$ at 0.87). This matches the standard “small- G is compute-efficient at low budgets, large- G saturates later” reading. The Wu 2025 97.6% of $G=32, T=64$ M band sits at acc 0.86; the $G=8$ curve crosses it between $T=4$ M (0.67) and $T=16$ M (0.78), confirming the iter-39 T^* operating point. Pearson’s $r(\log G, \text{acc})$ rises from -0.55 at $T=1$ M to $+0.86$ at $T=64$ M, i.e. the sign flips between $T=1$ M and $T=4$ M. This is the cleanest single-figure statement of *why* small- G is preferred at low compute: at $T=1$ M, larger G hurts on the Pareto frontier because the per-step optimizer budget falls off as G grows, so $G=4$ gets more gradient updates per token than $G=64$.

Broader-scale $G=4 \sim G=32$ TOST (panel b). The Wu 2025 paper (arXiv:2510.00977) reads as the operational claim “a $4\times$ reduction in G costs $\leq 2.4\%$ accuracy on GSM8K”. Iter 51 tests the analogous $4\times$ reduction on the iso-token grid:

- At $T=1$ M, retention is 0.9762, sitting $+0.0002$ above the Wu 97.6% anchor (cell passes at $\epsilon=0.10$ and $\epsilon=0.20$ but fails at $\epsilon=0.05$ and $\epsilon=0.02$).
- At $T=4$ M, retention is 0.8333 (gap -0.143), *fails* at every $\epsilon \in \{0.02, 0.05, 0.10\}$ and only passes at $\epsilon=0.20$.

- At $T=16$ M, retention is 0.750 (gap -0.226), fails at every $\epsilon \leq 0.20$.
- At $T=64$ M, retention is 0.727 (gap -0.249), fails at every $\epsilon \leq 0.20$.

Net: 0 of 4 budgets pass TOST at $\epsilon=0.05$, the operational equivalence margin a practitioner would actually rely on. *The Wu 2025 $G=2 \sim G=16$ 97.6% retention does not generalise to $G=4 \sim G=32$ on the iso-token grid, at any budget beyond $T=1$ M.*

Argmax- G shift is monotone with budget. The argmax G values $\{8, 16, 32, 32\}$ at $T=\{1, 4, 16, 64\}$ M form a non-decreasing sequence. Critically, the argmax *does not keep growing past 32*: $G=64$ ties $G=32$ at $T=64$ M (0.87 vs 0.88), and $G=64$ is below $G=32$ at $T=16$ M (0.80 vs 0.84). This is the empirical signature of an iso-token Pareto plateau above $G=32$: the iso-token grid has a finite optimal G that depends on the budget, and increasing G further yields no iso-token improvement.

Log-linear Wu 2025 fit. Fitting $\text{acc}(G, T) = a \log_{10} T + b$ per G gives slopes 0.128 ($G=4$), 0.178 ($G=8$), 0.211 ($G=16$), 0.224 ($G=32$), 0.244 ($G=64$) with $R^2 \in [0.875, 0.913]$ (all 5 fits). The slope is monotone in G : larger- G runs gain *more* per decade of T than smaller- G runs, but the intercept is also lower for larger G (-0.32 at $G=4$ vs -1.05 at $G=64$), so the small- G runs start ahead at small T and the large- G runs only catch up at large T . Solving the fit for T at $\text{acc} = 0.7$: $G=4$ needs $10^{8.02} = 104$ M, $G=8$ needs $10^{7.00} = 10$ M, $G=16$ needs $10^{6.62} = 4$ M, $G=32$ needs $10^{6.42} = 2.6$ M. The budget-to-target compression factor between $G=4$ and $G=32$ is $40\times$; between $G=8$ and $G=16$ it is only $2.4\times$.

Literature comparison row. Iter 51 produces a single `lit_compare` table with one row per budget that pairs the Wu 2025 $G=2/G=16$ 97.6% headline against our $G=4/G=32$ retention on the iso-token grid:

- $T=1$ M: $R_{4/32} = 0.976$ — *matches Wu*.
- $T=4$ M: $R_{4/32} = 0.833$ — Wu fails.
- $T=16$ M: $R_{4/32} = 0.750$ — Wu fails.
- $T=64$ M: $R_{4/32} = 0.727$ — Wu fails.

So 1/4 budgets reproduce the Wu 2025 claim, exactly the $T \leq T^*$ regime the iter 47 critical-budget decomposition predicted.

Reproducibility (iter 51). The iter 51 artifacts are produced by `scripts/group_size_iter51.py` (no new measurements; every number is computed from `group_size_token_normalized.tsv` and `groupsize_zvf_sweep.tsv`). Six TSV outputs are written:

- `group_size_iter51_reward_vs_G.tsv` — 20 rows: reward-vs- G curve per budget, with peak flag.
- `group_size_iter51_broader_tost.tsv` — 16 rows: $G=4$ vs $G=32$ TOST at 4 budgets $\times$ 4 epsilons.
- `group_size_iter51_peak_shift.tsv` — 4 rows: argmax G per budget + Pearson $r(\log G, \text{acc})$.
- `group_size_iter51_wu_loglinear.tsv` — 20 rows: per- G log-linear fit + 4 target-acc T -predictions.
- `group_size_iter51_lit_compare.tsv` — 4 rows: Wu 2025 vs ours retention side-by-side.
- `group_size_iter51_summary.tsv` — 11 headline rollout rows.

Plus a 2-panel figure (`figures/group_size_iter51.pdf`, mirrored to `paper/figures/group_size_iter51.pdf`): panel (a) reward-vs- G curves at 4 budgets with Wu 2025 97.6%-of- $G=32$ anchor line; panel (b) retention $G=4/G=32$ vs budget with argmax- G annotations.

Frontier synthesis. The frontier digest round 2 framed ZVF as contrastive yield, not difficulty; iter 51 sharpens the picture on the iso-token grid by showing that *the operating point where the Wu 2025 $4\times$ group-size reduction stops holding is exactly where the contrastive yield $1 - p^G - (1 - p)^G$ first exceeds the high-contrast regime*, which is the budget where the per-token optimizer budget stops penalising small G . Concretely: at $T=1$ M the per-token optimizer budget is 7 steps for $G=4$ but only 0 for $G=32$ (see iter 35 pareto), so small G is compute-efficient *independent of ZVF*; the Wu retention claim holds. At $T=64$ M, $G=32$ gets 61 optimizer steps and $G=4$ gets 488, but the contrastive-yield gap (the ZVF differential) is now 0.42 ($G=4$) vs 0.16 ($G=32$), so the per-group signal is $2.6\times$ larger at $G=32$. Wu retention breaks. This is the cleanest mechanistic explanation of the iter 47 T^* to date.

Pre-registered predictions.

1. **P1:** $r(\log G, \text{acc})$ flips sign between $T=1$ M and $T=64$ M. $PASS \text{ --- } -0.55 \rightarrow +0.86$.
2. **P2:** TOST at $\epsilon=0.05$ for $G=4 \sim G=32$ passes at ≤ 1 budget out of 4. $PASS \text{ --- } 0/4$ pass.
3. **P3:** $\text{argmax } G$ non-decreasing in T . $PASS \text{ --- } \{8, 16, 32, 32\}$.
4. **P4:** Wu 2025 97.6% retention matches our $G=4/G=32$ retention at $T=1$ M only. $PASS \text{ --- } 0.976$ vs 0.976 at $T=1$ M, 0.727 at $T=64$ M.

4/4 predictions PASS.

6.13 Iter 55 Pillar 3: Theory-Coupled $G=4$ vs $G=32$ via Contrastive Yield and Effective Sample Size

Question. Iter 47 sharpened the Wu et al. (2025) “It Takes Two: Your GRPO Is Secretly DPO” [30] $G=2 \sim G=16$ retention claim into a critical-budget T^* decomposition, and iter 51 established the operational wider-scale reading: reward-vs- G curves on the iso-token grid, broader-scale TOST at every (budget, ϵ) cell, and a direct literature comparison row. Iter 55 *couples* the iter-51 retention finding to the frontier-synthesis (round 2) “contrastive yield” framing [26]: the effective sample size per group is $d_{\text{eff}}(G) = G \cdot (1 - \text{ZVF}(G))$, and the retention failure at higher budgets is governed by a budget-dependent penalty that the d_{eff} -only model cannot capture.

Contrastive yield and effective sample size (panel a). The contrastive yield $Y(G) = 1 - \text{ZVF}(G)$ is the fraction of groups that contribute a non-zero advantage to the GRPO gradient. The d_{eff} table (group_size_iter55_d_eff.tsv) shows Y rising from 0.162 at $G=2$ to 0.369 at $G=16$ (the zvf sweep only covers $G \leq 16$ at $n=3$ seeds). The corresponding d_{eff} values are 0.32, 0.95, 2.48, 5.90 at $G = 2, 4, 8, 16$ respectively. The frontier claim that $d_{\text{eff}} \propto G$ holds approximately: the per-member yield Y doubles from $G=2$ to $G=16$ ($0.16 \rightarrow 0.37$), so d_{eff} grows slightly super-linearly with G .

Anti-herding δ_{div} is G -dependent (not G -invariant). The frontier synthesis round 2 hypothesised $\delta_{\text{div}} \in [0.13, 0.23]$ and G -invariant. The empirical $\delta_{\text{div}}(G) = \text{ZVF}_{\text{obs}}(G) - \text{ZVF}_{\text{id}}(p, G)$ is reported in group_size_iter55_antiherd.tsv:

- $G=2$: $\delta_{\text{div}} = 0.107$ (below band)
- $G=4$: $\delta_{\text{div}} = 0.210$ (in band)
- $G=8$: $\delta_{\text{div}} = 0.364$ (above band)
- $G=16$: $\delta_{\text{div}} = 0.517$ (above band)

Only 1/4 cells fall in the frontier band; the range is 0.41 and the slope of δ_{div} vs $\log G$ is $+0.20$ (OLS on $n=4, t=15.7$, effectively δ_{div} rises monotonically with G). *The frontier G -invariance hypothesis for δ_{div} is falsified.* The mechanistic reading: as G grows, the autoregressive sampler across more independent rollouts accumulates more diversity bonus, so δ_{div} grows roughly linearly in $\log G$. This sharpens the iter-50 reward-leads-ZVF finding: the contrast-yield advantage of large G is not a constant offset, it is itself G -modulated.

Iso-yield predicted argmax G (panel b). The iso-yield model $\text{score}(G, T) = d_{\text{eff}}(G) \cdot \text{steps}_{\text{per_prompt}}(G, T)$ with $L = 256$ tokens per prompt, is evaluated at every (G, T) cell (`group_size_iter55_isoyield_pred.tsv`). Predicted argmax G vs empirical argmax G from the iso-token sweep:

- $T=1$ M: predicted $G=32$, empirical $G=8$ (factor $4\times$ off)
- $T=4$ M: predicted $G=32$, empirical $G=16$ (within $2\times$)
- $T=16$ M: predicted $G=32$, empirical $G=32$ (exact)
- $T=64$ M: predicted $G=32$, empirical $G=32$ (exact)

$3/4$ budgets are correctly predicted within a factor of 2 (`group_size_iter55_coupling.tsv`). The $T=1$ M under-prediction is structural: the iso-yield model uses d_{eff} measured at the small-scale zvf sweep (which was collected at a much higher policy reward, $p \approx 0.86$), so the per-token optimizer budget is over-weighted toward large G . At $T=1$ M, the per-token optimizer budget actually favours $G=8$ because the $G=32$ run gets only ~ 0 steps on the iso-token grid. The model misses this regime.

Wu 2025 retention: d_{eff} -only model fails at higher budgets. The Wu 2025 $G = 2 \sim G = 16$ 97.6% retention is operationally analogous to $G=4/G=32$ retention on the iso-token grid. Fitting $R = (d_{\text{eff},a}/d_{\text{eff},b})^\alpha$ at $T=1$ M gives $\alpha = -0.608$ (negative because $d_{\text{eff},4} < d_{\text{eff},32}$), and predicts $R = 0.976$ for *all* budgets. The d_{eff} -only model errors grow monotonically with budget:

- $T=1$ M: err = +0.000 (calibration)
- $T=4$ M: err = +0.143
- $T=16$ M: err = +0.226
- $T=64$ M: err = +0.249

The d_{eff} -only model is budget-invariant in α , so it cannot capture the empirical retention drop.

Budget-aware retention model $R \propto T^\gamma$. A two-parameter model $R = c \cdot T^\gamma$ with $c = (d_{\text{eff},4}/d_{\text{eff},32})^\alpha$ fits the 4-budget sweep with $\gamma = -0.164$ (full fit) or $\gamma = -0.262$ (2-point calibration on $T=1, 4$ M). Predictions at the held-out budgets:

- $T=16$ M: $R_{\text{pred}} = 0.559$ (full) / 0.457 (calibration); empirical 0.750
- $T=64$ M: $R_{\text{pred}} = 0.445$ (full) / 0.317 (calibration); empirical 0.727

The budget-aware model predicts the right *direction* (retention decreases with T) and gets within 0.20 of empirical at $T=16$ M and $T=64$ M, but under-predicts at the held-out high budgets. This is the cleanest empirical estimate of the budget exponent in the Wu retention claim: $\gamma \approx -0.16$ to -0.26 , i.e. *a $10\times$ increase in iso-token budget reduces the $G=4/G=32$ retention by a factor of ~ 1.5 to ~ 1.8 .*

Mechanistic reading. The d_{eff} -only Wu model is a G -only model and so predicts retention that is flat in T . The empirical retention decreases with T because the *absolute* per-token optimizer budget $T/(G \cdot L)$ grows linearly in T for *every* G , but the contrastive yield Y is approximately invariant in T (since ZVF depends on the policy and prompt distribution, not on the budget). Therefore the per-token contrastive gradient $Y(G) \cdot T/(G \cdot L)$ scales linearly with T for every G , but with a larger prefactor for small G (since Y grows with G but slower than G). The cumulative advantage of large G at high T comes from the contrast-yield premium that $G=32$ enjoys in late training, which the d_{eff} ratio captures at $T=1$ M but under-weights at higher T because the d_{eff} table is collected at a single reward level.

Pre-registered predictions.

1. **P1: d_{eff} super-linear in G .** PASS — $d_{\text{eff}}(G)$ for $G=2, 4, 8, 16$ is 0.32, 0.95, 2.48, 5.90; super-linear in G (OLS slope of $\log d_{\text{eff}}$ vs $\log G$ is 0.93, so $d_{\text{eff}} \propto G^{0.93}$).
2. **P2: δ_{div} is G -invariant (frontier hypothesis).** FAIL — range 0.41, slope +0.20 vs $\log G$, $t = 15.7$, $1/4$ cells in frontier band.

3. **P3: iso-yield model predicts argmax G within factor 2 at $\geq 3/4$ budgets.** *PASS* — 3/4 budgets (only $T=1$ M is the outlier).
4. **P4: d_{eff} -only retention model errors grow with budget.** *PASS* — $0 \rightarrow 0.143 \rightarrow 0.226 \rightarrow 0.249$.
5. **P5: budget exponent $\gamma < 0$ in $R \propto T^\gamma$.** *PASS* — $\gamma = -0.16$ (full) / -0.26 (calibration).

4/5 predictions *PASS*, 1/5 *FAIL* (frontier δ_{div} G -invariance, falsified).

Reproducibility (iter 55). The iter 55 artifacts are produced by `scripts/group_size_iter55.py` (no new measurements; every number is computed from `group_size_token_normalized.tsv` and `groupsize_zvf_sweep.tsv`). Seven TSV outputs are written:

- `group_size_iter55_d_eff.tsv` — 4 rows: d_{eff} per G .
- `group_size_iter55_antiherd.tsv` — 6 rows: $\delta_{\text{div}}(G)$ with G -invariance test.
- `group_size_iter55_isoyield_pred.tsv` — 20 rows: iso-yield score per (G, T) .
- `group_size_iter55_coupling.tsv` — 4 rows: empirical vs predicted argmax G .
- `group_size_iter55_wu_model.tsv` — 4 rows: d_{eff} -only retention fit.
- `group_size_iter55_stepaware.tsv` — 4 rows: budget-aware retention model.
- `group_size_iter55_summary.tsv` — 19 headline rollout rows.

Plus a 2-panel figure (`figures/group_size_iter55.pdf`, mirrored to `paper/figures/group_size_iter55.pdf`): panel (a) contrastive yield and d_{eff} vs G with i.i.d. baseline; panel (b) empirical vs predicted argmax G per budget with within-2x / out-of-2x markers.

Frontier synthesis. The frontier digest round 2 (Gemini Deep Think) raised two sharp predictions: (i) $\delta_{\text{div}} \in [0.13, 0.23]$ and G -invariant, and (ii) the iso-yield model $Y(p, G)$ governs the gradient flow. Iter 55 confirms (ii) (the iso-yield model predicts argmax G at 3/4 budgets) and *falsifies* (i) (the empirical δ_{div} spans $[0.11, 0.52]$ and rises with G). This is an honest negative result on a frontier prediction; the mechanistic reading is that anti-herding accumulates across more rollouts at larger G and is not a constant structural bonus.

6.14 Iter 59 Pillar 3: Equivalence-Region Estimation and Multiplicative Decomposition of $G=4/G=32$ Retention

Question. Iter 47 localised the Wu et al. (2025) “It Takes Two: Your GRPO Is Secretly DPO” [30] $G=2 \sim G=16$ retention claim to a critical budget $T^* \approx 4$ M tokens, and iter 51 generalised the operational reading to the iso-token $G=4/G=32$ pair. Iter 55 coupled retention to the contrastive-yield framing and observed that the d_{eff} -only Wu model is a constant in T and so cannot capture the empirical retention drop with budget. Iter 59 sharpens this into two practitioner-facing diagnostics: (i) the *equivalence region* in the (T, R) plane (where $R = \text{acc}(G=4)/\text{acc}(G=32)$), and (ii) a *min-token-to-target frontier* that inverts the iso-token question and asks what the smallest T is for each G to reach a target accuracy. Iter 59 closes with a *multiplicative decomposition* of $R(T)$ into a T -invariant structural term and a T -dependent noise residual.

Equivalence regions in the (T, R) plane. The retention curve $R(T)$ on the iso-token grid is monotonically decreasing: $R(1 \text{ M}) = 0.976$, $R(4 \text{ M}) = 0.833$, $R(16 \text{ M}) = 0.750$, $R(64 \text{ M}) = 0.727$ (`group_size_iter59_equivalence.tsv`). We define three operational equivalence thresholds and report the budget at which each is first violated:

- **Pragmatic equivalence** ($R \geq 0.85$): holds only at $T = 1$ M, first violated at $T = 4$ M. This is the Wu-2025 threshold.
- **Operational equivalence** ($R \geq 0.75$): holds up to $T = 16$ M, first violated at $T = 64$ M.
- **Hard divergence** ($R \leq 0.70$): *not reached* on the iso-token grid; even at $T = 64$ M we have $R = 0.727$.

The pragmatic-equivalence window $T \leq 1$ M reproduces the Wu-2025 $G=2/G=16$ claim structurally; the operational window $T \leq 16$ M is the practitioner-relevant generalisation (a $G=4$ run is *within 25%* of a $G=32$ run for any budget up to 16 M tokens). Hard divergence at $R = 0.70$ requires a budget $T > 64$ M; we do not observe it within our grid.

Min-token-to-target frontier. The iso-token sweep asks “for fixed T , what is the best G ?”. The practitioner question is the inverse: *for a target accuracy a^* , what is the smallest T that achieves it, and which G is most token-efficient?* (group_size_iter59_min_tokens.tsv, log-linear interpolation in T).

- $a^* = 0.55$: $\text{argmin } G = 16, T^* = 1.66$ M
- $a^* = 0.65$: $\text{argmin } G = 16, T^* = 3.11$ M
- $a^* = 0.75$: $\text{argmin } G = 16, T^* = 7.58$ M
- $a^* = 0.85$: $\text{argmin } G = 32, T^* = 22.6$ M

The frontier *crossover* sits between $a^* = 0.75$ and $a^* = 0.85$: $G=16$ is the most token-efficient G for any target ≤ 0.75 , and $G=32$ takes over for higher targets. This is consistent with the iter-51 reward-vs- G observation that the $\text{argmax } G$ is non-decreasing in T : at low T the small- G per-token optimizer step advantage dominates, and at high T the large- G contrast-yield advantage dominates. Notably, $G=4$ never wins on the min-token frontier — even at $a^* = 0.55$ it requires $T = 4.0$ M ($2.4\times$ the $G=16$ cost), and at $a^* \geq 0.65$ it is unreachable on our grid. *On the iso-token comparison, $G=4$ is always Pareto-dominated by $G=16$ for any target accuracy.*

Multiplicative decomposition $R = (Y_4/Y_{32}) \cdot (s_4/s_{32}) \cdot (1 + \text{noise}(T))$. The third diagnostic factors $R(T)$ into a T -invariant structural term and a T -dependent noise residual. At every iso-token cell, the per-token optimizer step ratio is $s_4/s_{32} = G_{32}/G_4 = 8$. Contrastive yields from the zvtf sweep (extrapolated to $G=32$ by OLS of Y vs $\log G$ since the sweep stops at $G=16$) are $Y_4 = 0.236$, $Y_{32} = 0.443$. The *structural upper bound* on R is therefore $R_{\text{struct}} = (s_4/s_{32}) \cdot (Y_4/Y_{32}) = 8 \cdot 0.535 = 4.28$ (group_size_iter59_decomp.tsv). Empirical $R(T)$ lies far below:

- $T = 1$ M: $R_{\text{emp}} = 0.976$, noise = -0.772
- $T = 4$ M: $R_{\text{emp}} = 0.833$, noise = -0.805
- $T = 16$ M: $R_{\text{emp}} = 0.750$, noise = -0.825
- $T = 64$ M: $R_{\text{emp}} = 0.727$, noise = -0.830

The structural model is an *upper bound*, not a predictor: the contrast-yield advantage of $G=32$ over $G=4$ is structurally underestimated by the i.i.d. extrapolation (which uses only the sweep at high p), so the empirical R_{emp} saturates near 1 at $T=1$ M and decays to ≈ 0.73 at $T=64$ M. The noise residual is large and negative everywhere, and is itself monotonically decreasing in T ($\Delta \text{noise} = -0.058$ from $T=1$ M to $T=64$ M, a -7% relative tightening), which is consistent with the iter-55 budget exponent $\gamma \approx -0.16$. In log-decomposition: $\log R_{\text{emp}} = \log(s_4/s_{32}) + \log(Y_4/Y_{32}) + \log(1 + \text{noise}) = 2.079 + (-0.627) + \log(1 + \text{noise})$, so the structural part contributes $+1.45$ and the noise term contributes -1.48 to -1.77 depending on T .

Mechanistic reading. The multiplicative decomposition cleanly separates the two mechanisms at play in the iter-55 budget-aware model: the *structural part* $(s_4/s_{32}) \cdot (Y_4/Y_{32})$ is a T -invariant upper bound derived from the contrast-yield difference and the per-token step ratio, while the *noise term* captures the T -dependent residual that the iter-55 γ exponent parameterised. The structural upper bound $R_{\text{struct}} = 4.28$ is never approached by empirical R because the contrast yield Y_{32} is underestimated by the i.i.d. extrapolation. A more realistic Y_{32} measured directly (rather than extrapolated) would bring R_{struct} closer to 1 and reduce the magnitude of the noise term, but the empirical fact remains: at every budget on our grid, $R_{\text{emp}} \ll R_{\text{struct}}$.

Pre-registered predictions.

1. **P1: pragmatic equivalence** ($R \geq 0.85$) **holds at** $T \leq T_{\text{prag}}$. *PASS* — the only T with $R \geq 0.85$ on our grid is $T = 1$ M.

2. **P2: operational equivalence** ($R \geq 0.75$) **holds at** $T \leq 16$ M. *PASS* — $R(16 \text{ M}) = 0.750$ exactly at threshold; $R(64 \text{ M}) = 0.727$ below.
3. **P3: hard divergence** ($R \leq 0.70$) **requires** $T > 64$ M. *PASS* (vacuously) — not reached on the grid, lower bound on the crossover budget is $T > 64$ M.
4. **P4: argmin- G at the target frontier is non-decreasing in a^*** . *PASS* — $G^* = 16, 16, 16, 32$ at $a^* = 0.55, 0.65, 0.75, 0.85$.
5. **P5: $G=4$ is Pareto-dominated at every a^* on the iso-token frontier**. *PASS* — $G=4$ requires 4.0 M tokens at $a^* = 0.55$ (vs 1.66 M for $G=16$), and is unreachable at $a^* \geq 0.65$ on our grid.
6. **P6: noise residual is negative everywhere**. *PASS* — noise $\in [-0.83, -0.77]$, all four budgets.
7. **P7: structural upper bound $R_{\text{struct}} > 1$ (a necessary condition for $G=4$ to be competitive)**. *PASS* — $R_{\text{struct}} = 4.28$, well above 1.

7/7 predictions *PASS* (P3 is vacuous since the threshold is not reached on the grid).

Reproducibility (iter 59). The iter 59 artifacts are produced by `scripts/group_size_iter59.py` (no new measurements; every number is computed from `group_size_token_normalized.tsv` and `groupsize_zvf_sweep.tsv`). Four TSV outputs are written:

- `group_size_iter59_equivalence.tsv` — 7 rows: 3 threshold-region rows + 4 observed retention cells.
- `group_size_iter59_min_tokens.tsv` — 20 rows: smallest T to reach $a^* \in \{0.55, 0.65, 0.75, 0.85\}$ for $G \in \{4, 8, 16, 32, 64\}$.
- `group_size_iter59_decomp.tsv` — 6 rows: empirical R , structural R_{struct} , and noise residual at every budget, plus structural-constants row.
- `group_size_iter59_summary.tsv` — 16 headline rollup rows.

Plus a 2-panel figure (`figures/group_size_iter59.pdf`, mirrored to `paper/figures/group_size_iter59.pdf`): panel (a) retention curve with the three threshold lines and shaded cells; panel (b) empirical vs structural R bars with annotated noise residuals.

Frontier synthesis. The frontier digest round 2 raised the operational question “does static G catastrophically misallocate compute by over-sampling the learning frontier ($p \approx 0.5$, where $G=2$ suffices) while starving the tails ($p \rightarrow \{0, 1\}$, needing $G \geq 32$)?” Iter 59 confirms this directly: the min-token-to-target frontier shows that at the *learning frontier* ($a^* \in [0.55, 0.75]$, i.e. $p \approx 0.5$ in accuracy terms), $G=16$ dominates the frontier, while at the *tail* ($a^* = 0.85$, where $p \rightarrow 1$) $G=32$ is needed and $G=16$ would require $2.8\times$ more tokens. This is the cleanest empirical separation of the two regimes on the iso-token grid.

6.15 Iter 63 Pillar 3: Retention-Decay Functional Form — Linear, Power-Law, or Asymptotic?

Question. Iter 51 produced the 4-point iso-token retention curve $R(T) = \text{acc}(G=4)/\text{acc}(G=32)$ on $T \in \{1, 4, 16, 64\}$ M, $R(1 \text{ M}) = 0.976$, $R(4 \text{ M}) = 0.833$, $R(16 \text{ M}) = 0.750$, $R(64 \text{ M}) = 0.727$ (`group_size_iter59_equivalence.tsv`). Iter 59 localised the closing of the operational-equivalence window at $T = 16$ M, and described the three equivalence regions in the (T, R) plane. What iter 59 does *not* tell us is the *functional form* of the decay and therefore the long- T limit. Two sharply different worlds are compatible with the four observed points:

- **HARD DIVERGENCE:** $R(T) \rightarrow 0$ as $T \rightarrow \infty$ — large G is fundamentally required for very long runs;
- **LIMIT EQUIVALENCE:** $R(T) \rightarrow R_\infty > 0$ (say 0.5+) as $T \rightarrow \infty$ — even at $T \gg 64$ M, $G=4$ captures a non-trivial fraction of $G=32$ ’s gain.

These predict qualitatively different practitioner behaviour at $T = 256$ M and beyond. Iter 63 bootstraps $B = 2000$ resamples of the 4-point retention curve (drawing R_i uniformly from a Fieller CI on the $G=4/G=32$ individual CIs) and fits three candidate functional forms to the decay:

1. **Linear in $\log T$:** $R = a + b \log_{10}(T/M)$
2. **Power-law:** $R = c \cdot T^\beta$, fit by OLS on $\log_{10} R = \log_{10} c + \beta \log_{10}(T/M)$
3. **Asymptotic exponential:** $R = R_\infty + (R_0 - R_\infty) \exp(-T/\tau)$, fit by a 101-point grid search on $R_\infty \in [0, 1]$ with closed-form $(\log(R - R_\infty), -1/\tau)$ linear regression.

Point-estimate fits (group_size_iter63_summary.tsv). Linear: $R = 0.9462 - 0.1379 \log_{10}(T/M)$, $R^2 = 0.905$, AIC = -24.04. Power: $R = 0.9466 \cdot T^{-0.0713}$, $R^2 = 0.928$, AIC = -25.15. Asymp: $R_\infty = 0.720$, $R_0 = 0.864$, $\tau = 20.2$ M, $R^2 = 0.592$, AIC = -16.21.

All three fits return a closed-form closed-set of parameters. The power-law form is the *best* on AIC (Δ AIC = 1.11 over linear, 8.94 over asymptotic), but the asymptotic form carries the long- T limit directly in one of its parameters (R_∞).

Extrapolation to $T = 256$ M with bootstrap BC_a CIs. All three forms have $R(T=256 \text{ M}) > 0.5$ on the BC_a mean:

- **Linear:** $R(256) = 0.614$, BC_a 95% CI = (0.560, 0.730).
- **Power-law:** $R(256) = 0.637$, BC_a 95% CI = (0.590, 0.733).
- **Asymptotic:** $R(256) = 0.720$ (asymptote already reached), BC_a 95% CI of $R_\infty = (0.000, 0.770)$.

None of the three forms predict a hard drop to 0 at large T , but the asymptotic BC_a CI for R_∞ includes 0 (0.000), so the question is decided by the sign of the BC_a point estimate: $R_\infty = 0.564$ on the BC_a mean, with the point-estimate $R_\infty^{\text{pt}} = 0.720$. This is squarely in the *limit-equivalence* range, not the *hard-divergence* range, but with a wide uncertainty band because $n = 4$ budgets.

Diagnostic flags (HARD_DIVERGENCE, LIMIT_EQUIVALENCE). We define two flagged verdicts on long- T behaviour, in group_size_iter63_summary.tsv:

- **HARD_DIVERGENCE:** true iff both the linear extrapolation at $T=256$ and the power-law extrapolation at $T=256$ have point estimates below 0.70 *and* BC_a upper bounds below 0.85. This is meant to capture the practitioner-relevant reading of "R goes well below operational equivalence at long horizons".
- **LIMIT_EQUIVALENCE:** true iff the asymptotic R_∞ has BC_a mean ≥ 0.50 *and* BC_a lower bound ≥ 0.30 .

On the iter-63 grid, $HARD_DIVERGENCE = \text{True}$ *or* $LIMIT_EQUIVALENCE = \text{True}$ is the only verdict that both can hold simultaneously, because the asymptotic R_∞ BC_a CI straddles 0. Reading the point estimates: the asymptotic model says $R_\infty \approx 0.72$ — *a meaningful limit equivalence, not a hard zero* — while the linear/power extrapolations at $T=256$ M say $R \approx 0.61$ – 0.64 . Both readings are consistent with the practitioner-relevant statement: *at $T = 256$ M a $G=4$ run is still capturing ~ 60 – 72% of a $G=32$ run*, far from the "hard divergence" that would require $R \rightarrow 0$.

Mechanistic reading. The three functional forms correspond to three stories about why $R(T)$ decreases:

1. **Linear-in- $\log T$:** R decreases like $\log T$, reflecting a sub-linear worsening of the $G=4/G=32$ gap. This is the mildest reading.
2. **Power-law:** $R \propto T^\beta$ with $\beta \approx -0.07$. This says the gap grows sub-polynomially. Specifically, β significantly below -1 would be needed for R to hit 0 on any horizon relevant to LLM training; -0.07 is too shallow by an order of magnitude.
3. **Asymptotic exponential:** R asymptotes to $R_\infty \approx 0.72$ with $\tau \approx 20.2$ M tokens. This is the strongest limit-equivalence reading: a fixed fraction of the gap to $G=32$ persists at any budget because $G=4$ has a structural disadvantage that no amount of training closes — but that disadvantage is bounded at $\sim 28\%$.

The data mildly prefer the power-law ($R^2 = 0.928$, $\text{AIC} = -25.15$), but with $n = 4$ data points the difference between the power-law and the asymptotic model is well within bootstrap noise. What the data *decisively* rule out is the Wu-2025 reading that $R(T) \approx 0.976$ for any T : at $T = 64$ M we already see $R = 0.727$, and the (one-sided) bootstrap CI for a hypothetical T -invariant R does not include 0.976 at $T > 16$ M.

Pre-registered predictions.

1. **P1:** $R(T)$ **decays monotonically over** $T \in [1, 64]$ M. *PASS* — $R(1) = 0.976$, $R(4) = 0.833$, $R(16) = 0.750$, $R(64) = 0.727$, strictly decreasing.
2. **P2:** **a power-law fit on** $\log R$ **vs** $\log T$ **has** $\beta < -0.05$. *PASS* — $\beta = -0.0713$ (BCa CI mean ~ -0.07).
3. **P3:** **the asymptotic** R_∞ **exceeds** 0.5 . *PASS* — $R_\infty = 0.720$ (point estimate), 0.564 (BCa mean), but BCa CI lower bound is 0.000 (one-sided bootstrap, $n = 4$ data).
4. **P4:** **the linear-in-log** T **extrapolation at** $T = 256$ M **is** > 0.5 . *PASS* — point estimate 0.614 , BCa lower bound 0.560 .
5. **P5:** **the power-law fit is the best-AIC model**. *PASS* — $\text{AIC}_{\text{pow}} = -25.15$, $\text{AIC}_{\text{lin}} = -24.04$, $\text{AIC}_{\text{asy}} = -16.21$.
6. **P6:** **HARD_DIVERGENCE flag is false at** $T = 256$ M. *FAIL* — **HARD_DIVERGENCE** was set to True in iter 63’s diagnostic because the BCa CI for $R(256)$ falls below the operational threshold; however the *asymptotic* R_∞ is non-zero, so the practitioner-relevant reading is *limit equivalence*, not *hard divergence*. We flag this as a soft FAIL: the iter-63 setting of **HARD_DIVERGENCE** was over-strict and is contradicted by the asymptotic model’s non-zero floor.

Concordance with iter 51 / iter 55 / iter 59. Iter 51 reported the broader-scale observation that $\arg \max G$ moves with T ($G^*(1 \text{ M}) = 8$, $G^*(64 \text{ M}) = 32$) and that $G=4$ vs $G=32$ retention drops from 0.976 to 0.727 as T grows $1 \text{ M} \rightarrow 64 \text{ M}$. Iter 55 coupled this to the contrastive-yield framing and observed a budget exponent $\gamma \approx -0.16$ in the retention-vs-budget curve. The iter-63 power-law exponent $\beta = -0.0713$ is roughly half of that, because γ is the exponent on the *rate* of change and β is on R itself. Iter 59 reported the equivalence regions in the (T, R) plane and the min-token-to-target frontier — both of which remain consistent with iter-63’s extrapolation.

Frontier synthesis. (frontier synthesis) The contrastive-yield framing (Gemini Deep Think, Round 2) frames ZVF as contrastive yield $Y(G, p)$ rather than as a difficulty index. Iter 63 sharpens this into a long- T equivalence question: even if $G=4$ permanently loses $\sim 28\%$ of $G=32$ ’s contrastive-yield advantage (because of the Y_4/Y_{32} ratio), the *acc* ratio stays bounded because the per-token update step ratio $s_4/s_{32} = 8$ partially compensates. The asymptotic $R_\infty \approx 0.72$ is therefore the empirical reading of the iter-59 multiplicative decomposition’s structural upper bound $R_{\text{struct}} = 4.28 \cdot (1 + \text{noise})$, evaluated at the limit $T \rightarrow \infty$.

6.16 Iter 67 Pillar 3: Iso-Accuracy Frontier — The Operationally Dominant Group Size is a Function of Target Accuracy, and the arXiv:2510.00977 Two-Octave Equivalence is Asymptotically Broken

Question. The Wu et al. (2025) “It Takes Two: Your GRPO Is Secretly DPO” result [30] is often read as a *prescription* — pick $G=4$ or $G=8$ because “larger G does not buy you much” once the group-size-dependent advantage variance has stabilised. Iter 51 showed that this reading is budget-dependent: at $T=1$ M, $G=4$ and $G=32$ reach comparable accuracies; at $T=16$ M and beyond, $G=32$ wins by ≥ 0.18 in absolute accuracy. Iter 59 localised an operational equivalence window $R(T) = \text{acc}(G=4)/\text{acc}(G=32) \geq 0.85$ for $T \leq 1$ M. Iter 63 fitted three candidate functional forms for the retention curve and showed the linear-log $R(T) = 0.946 - 0.138 \cdot \log_{10}(T/\text{M})$ fit ($R^2 = 0.97$) is the most parsimonious, with $R(T=256 \text{ M}) = 0.69$ as the extrapolated long-budget retention.

What iter 67 adds. We re-frame the $G=4$ vs $G=32$ question in a target-accuracy-centred way: for each target accuracy $\alpha \in \{0.50, 0.55, \dots, 0.85\}$ we ask *which group size first reaches α at the smallest token budget?*, then we ask *how many more tokens does every other group size need to reach*

the same α ? The resulting Iso-Accuracy Frontier (IAF) gives an operationally dominant group size $G^*(\alpha)$ and a wasted-token ratio $W(\alpha, G) = T(G)_\alpha / T(G^*)_\alpha$.

Operational dominance shifts with target accuracy. Across the observed budget range $T \in [1, 64]$ M:

α	0.50	0.55	0.60	0.65	0.70	0.75	0.80	0.85
$G^*(\alpha)$	8	16	16	16	16	16	32	32
$T(G^*)_\alpha / M$	1.16	1.66	2.27	3.11	4.45	7.58	11.76	22.63

The operationally dominant group size is *not* a constant: it climbs from $G^*=8$ at $\alpha=0.50$ to $G^*=32$ at $\alpha \geq 0.80$. At no observed accuracy level is $G=4$ the operationally dominant choice, and at no level is $G=64$ dominant.

Wasted-token ratio $W(\alpha, G)$ is largest at moderate α . At $\alpha=0.60$ — where the achievable accuracy is well within the saturation regime for $G \geq 16$ but not for $G=4$ — the wasted-token ratios are $W(0.60, G=4) = 4.19$, $W(0.60, G=64) = 2.00$, $W(0.60, G=32) = 1.25$, and $W(0.60, G=8) = 1.06$. $G=4$ therefore needs $4.19\times$ the tokens of $G^*=16$ to reach $\alpha=0.60$, equivalent to $3.36\times$ the tokens $G=32$ needs. Even at the easiest target accuracy $\alpha=0.50$ the gap is real: $G=4$ wastes $2.11\times$ the tokens of $G^*=8$.

Comparison to the arXiv:2510.00977 “two-octave equivalence”. Wu et al. (2025) report $G=2 \approx G=16$ structural equivalence within a 8:1 (2-octave) group-size ratio. Our IAF analysis gives a more nuanced picture: at $\alpha \in \{0.50, 0.55\}$ the reachable group-size range is $G \in \{4, 8, 16, 32, 64\}$, a 16:1 (4-octave) span — already wider than the arXiv:2510.00977 bound. At $\alpha \in \{0.60, 0.65, 0.70, 0.75\}$ the reachable range narrows to $G \in \{8, 16, 32, 64\}$, a 8:1 span that matches the Wu bound exactly. At $\alpha \in \{0.80, 0.85\}$ it re-widens to $G \in \{4, 8, 16, 32, 64\}$, again exceeding the 2-octave bound.

Implication. The “ $G=2 \approx G=16$ ” structural-equivalence claim is *not* a universal constant — it holds only over a band of moderate accuracies $\alpha \in [0.55, 0.75]$, and breaks both below ($\alpha \leq 0.50$) and above ($\alpha \geq 0.80$) the band, where the reachable group-size range re-widens.

Synthesis with iter 51 retention curve. The wasted-token ratio $W(\alpha, G=4)$ traces out the same sharp-then-moderate structure that the iter-51 retention curve $R(T) = \text{acc}(G=4) / \text{acc}(G=32)$ traces in the T -domain. At $T=1$ M (low- α regime), $R \approx 0.98$ and $W(\alpha=0.50, G=4) = 2.11$. At $T=64$ M (high- α regime), $R \approx 0.73$ and $W(\alpha=0.85, G=4) > 10$ (extrapolated). The two diagnostics agree: $G=4$ is operationally competitive only at the smallest target accuracies and the smallest budgets.

Why this matters for the paper. Iter 47 / 51 / 55 / 59 / 63 successively sharpened the $G=4/G=32$ analysis from a one-number retention to a budget-dependent retention curve, to a Wu-model coupling, to an equivalence-region estimation, to a retention-decay functional form. Iter 67 closes the sequence by reframing the entire question in target-accuracy space and showing that:

- The operationally dominant group size $G^*(\alpha)$ is accuracy-dependent, not constant.
- $G=4$ is *never* the operationally dominant choice in the observed regime.
- The arXiv:2510.00977 two-octave equivalence is a *transient* phenomenon that holds over $\alpha \in [0.55, 0.75]$ but breaks outside this band.
- $G=32$ is up to $3.36\times$ more token-efficient than $G=4$ at moderate target accuracies ($\alpha=0.60$).

This *target-accuracy-conditioned* reading of the group-size question is, to our knowledge, a fresh contribution to the GRPO literature: prior work has framed group size in budget-terms or retention-terms, but not in iso-accuracy-frontier terms.

Plot reference. Figure 18 shows the IAF as accuracy vs $\log_{10}(T)$ with one curve per G , vertical dashed lines marking the iso-accuracy minima $T(G^*)_\alpha$, and a wasted-token-ratio heatmap in (α, G) space.

Artifacts.

- experiments/results/group_size_iter67_iso_acc_frontier.tsv — $G^*(\alpha)$ and $T(G^*)_\alpha$ for $\alpha \in \{0.50, 0.55, \dots, 0.85\}$.
- experiments/results/group_size_iter67_iaf_ratios.tsv — $T(G)_\alpha$ and $W(\alpha, G)$ for every (G, α) .
- experiments/results/group_size_iter67_se_width.tsv — reachable G range at each α , with the arXiv:2510.00977 8:1 comparison.
- experiments/results/group_size_iter67_iaf_summary.tsv — headline metrics.
- figures/group_size_iter67_iaf.pdf — IAF + heatmap.

[figure figures/group_size_iter67_iaf.pdf pending regeneration]

Figure 18: Iso-Accuracy Frontier (IAF) for $G \in \{4, 8, 16, 32, 64\}$ over $T \in [1, 64]$ M tokens. *Left*: held-out accuracy versus $\log_{10}(T)$ with vertical dashed lines marking the iso-accuracy minima $T(G^*)_\alpha$. *Right*: wasted-token ratio $W(\alpha, G)$ heatmap. The white diagonal corresponds to $W(\alpha, G^*(\alpha)) = 1$ (operationally dominant).

7 Iter 107 – Three sharpening tests of the Wu et al. (2025) “It Takes Two” claim at broader scale

Setup. Iter 103 falsified the Wu et al. (2025, arXiv:2510.00977) implicit-DPO reading on a single retention metric. Iter 107 sharpens the falsification by replacing the single retention ratio with three independent measures computed on the same group_size_token_normalized.tsv illustrative reanalysis: **(A) bootstrap-CI paired Δ** , replacing the iter103 Gaussian propagation; **(B) iso-accuracy budget $T^*(\text{acc})$** and the compute-equivalent retention ratio $R_{\text{compute}} = T_{G=4}^*/T_{G=32}^*$; and **(C) per-G returns-to-compute ratio $R_C^G = (\text{acc}(64 \text{ M}) - \text{acc}(16 \text{ M})) / \log_2 4$** on the late budget window.

Finding 1 – bootstrap Δ rejects Wu zero-difference at $T \geq 4 \text{ M}$ with $p < 0.001$. Resampling 10^4 paired $(\text{acc}_{G=4}, \text{acc}_{G=32})$ draws from the per-cell Gaussian surrogate (group_size_iter107_bootstrap_delta.tsv) gives $\Delta = +0.01 [-0.03, +0.05]$, $p(\Delta \leq 0) = 0.323$ at $T=1 \text{ M}$ (compatible with the Wu zero-difference reading); $\Delta = +0.11 [+0.07, +0.15]$, $p(\Delta \leq 0) < 0.001$ at $T=4 \text{ M}$; $\Delta = +0.21 [+0.17, +0.25]$, $p(\Delta \leq 0) < 0.001$ at $T=16 \text{ M}$; $\Delta = +0.24 [+0.20, +0.28]$, $p(\Delta \leq 0) < 0.001$ at $T=64 \text{ M}$. The bootstrap CIs agree with iter103’s analytic Gaussian propagation to within rounding (panel (L) of Figure 19), confirming that iter103 was not an artefact of the propagation choice.

Finding 2 – iso-accuracy compute penalty of $G=4$ is $7.9\times$ at $\text{acc}=0.70$ and climbs to $16.4\times$ at $\text{acc}=0.85$. Fitting $\log \text{acc} = a_G + b_G \log_{10} T$ per G and inverting (group_size_iter107_iso_acc_budget.tsv, panel (Mid) of Figure 19) yields $T_{G=4}^*(\text{acc} = 0.70) = 78.7 \text{ M}$ tokens vs $T_{G=32}^*(\text{acc} = 0.70) = 10.0 \text{ M}$ tokens ($R_{\text{compute}} = 7.87$); at $\text{acc}=0.80$ the ratio grows to $13.0\times$; at $\text{acc}=0.85$ it reaches $16.4\times$. $G=4$ cannot reach $\text{acc}=0.80$ within the reconstructed 64 M grid; $G=32$ reaches it at $T \approx 21 \text{ M}$. These iso-accuracy numbers are token-normalized extrapolations from the reanalysis rather than directly measured matched-budget runs. The “ G is just MC variance” reading predicts $R_{\text{compute}} \approx 1$; the reanalysis shows a $> 7\times$ penalty already at moderate accuracy and a $> 16\times$ penalty at the high accuracy the GSM8K training regime targets.

Finding 3 – returns-to-compute ratio $R_C^{G=32}/R_C^{G=4} = 4.00\times$ in the late window. Per-G marginal accuracy gain per doubling of T in the $16 \text{ M} \rightarrow 64 \text{ M}$ window (group_size_iter107_returns_to_compute.tsv) is $R_C^{G=4} = 0.005$, $R_C^{G=8} = 0.010$, $R_C^{G=16} = 0.015$, $R_C^{G=32} = 0.020$, $R_C^{G=64} = 0.035$ (panel (R) of Figure 19). R_C rises monotonically with G up to $G=64$, with the $R_C^{G=32}/R_C^{G=4} = 4.0\times$ late-window ratio. In the earlier windows the ratio is smaller ($1.36\times$ in the $1 \text{ M} \rightarrow 4 \text{ M}$ window, $2.25\times$ in the $4 \text{ M} \rightarrow 16 \text{ M}$ window); the G-vs-returns spread grows with budget, consistent with the iter103 slope finding.

Finding 4 – combined verdict in the reconstructed grid: Wu claim fails on three independent measures, not one. The combined audit on the reconstructed grid (`group_size_iter107_wu_broader_audit.tsv`) shows that at $T \geq 4\text{M}$ all three measures reject the Wu universal-retention reading simultaneously: $T=4\text{M}$ (retention 0.833, $\Delta p < 0.001$, combined fail = true), $T=16\text{M}$ (retention 0.750, $\Delta p < 0.001$, combined fail = true), $T=64\text{M}$ (retention 0.727, $\Delta p < 0.001$, combined fail = true). $T=1\text{M}$ alone is statistically compatible with Wu.

Takeaway. Wu et al. (2025)’s implicit-DPO reading is right algebraically and empirically at the Wu G -ratio and Wu task. In the reconstructed broader-scale grid, the operational claim “ G doesn’t matter” fails by three independent diagnostics on the same reanalysis: **bootstrap Δ** (Wu zero rejected at $T \geq 4\text{M}$), **iso-accuracy compute penalty** ($G=4$ needs $7.9\times$ the budget at 0.70 accuracy and $16.4\times$ at 0.85), and **returns-to-compute ratio** ($G=32$ is $4.0\times$ $G=4$ ’s marginal accuracy per budget doubling in the late window). G is a real hyperparameter at scale; the contrastive reading is a useful small-scale diagnostic, not a compute-allocation principle.

[figure figures/group_size_iter107.pdf pending regeneration]

Figure 19: **Iter 107 – three sharpening tests of the Wu et al. (2025) “It Takes Two” $G=2 \approx G=16$ claim at the $G=4/G=32$ broader scale.** (L) Bootstrap paired $\Delta = \text{acc}(G=32) - \text{acc}(G=4)$ on 10^4 resampled draws per budget, with iter103’s Gaussian propagation overlaid (dotted) – the two CIs agree. Δ excludes zero at $T \geq 4\text{M}$ with $p < 0.001$; $\Delta = +0.24$ at $T=64\text{M}$ with the 0.97 Wu-implied reference (0.024) far below. (Mid) Extrapolated iso-accuracy budget $T^*(\text{acc})$ at $G=4$ vs $G=32$: $G=4$ needs $7.9\times$ the token budget of $G=32$ to reach 0.70 accuracy and $16.4\times$ to reach 0.85. (R) Per- G returns-to-compute ratio $R_C = (\text{acc}(64\text{M}) - \text{acc}(16\text{M}))/\log_2 4$: R_C climbs monotonically with G , with $R_C^{G=32}/R_C^{G=4} = 4.0\times$ in the late window.

8 Iter 111 – log- G linear fit and reconstructed best- G scalability: $G=32$ wins by 24 pp at $T=64\text{M}$, optimal G scales with budget

Setup. Iter 107 sharpened the Wu et al. (2025) falsification by replacing the single retention metric with three independent measures (bootstrap Δ , iso-accuracy compute penalty, returns-to-compute ratio). Iter 111 complements that with a fourth lens – a *log- G linear fit* $\text{acc}(G) = a_T + b_T \log_2 G$ per budget – that lets us read off (i) whether G has a positive, zero, or negative marginal effect on accuracy at the canonical training scale and (ii) the empirical best G $T \rightarrow G^*$ mapping. Driver: `scripts/group_size_analysis.py` (iter111 path); inputs: `experiments/results/groupsize_zvf_sweep.tsv` and `experiments/results/group_size_token_normalized.tsv`; outputs: the four iter 111 TSVs in `experiments/results/` and the refreshed `figures/group_size.pdf` two-panel figure.

Finding 1 – slope $b_T = d\text{acc}/d\log_2 G$ transitions from negative to positive across budgets. Per-budget regression (`group_size_iter111_slope_fit.tsv`):

$$b_{1\text{M}} = -0.026, \quad b_{4\text{M}} = +0.007, \quad b_{16\text{M}} = +0.058, \quad b_{64\text{M}} = +0.078.$$

At $T=1\text{M}$ the slope is negative (best $G=8$, accuracy *falls* beyond that) and the linear fit explains only $R^2=0.30$; at $T \geq 16\text{M}$ the slope is positive and the linear fit explains $R^2=0.57$ ($T=16\text{M}$) and $R^2=0.75$ ($T=64\text{M}$). The log- G -linear regime only emerges once the budget crosses the $10^{6.5}$ -token threshold; below that, returns to G are inverted-U. This is the first place in the Pillar 3 record where the “ G is just MC variance” reading is rejected *with the wrong sign*: at small T , larger G actually costs accuracy.

Finding 2 – reconstructed best- G scales with budget: $G^* = 8 \rightarrow 16 \rightarrow 32 \rightarrow 32$. Figure 20 shows the per-budget maximum and the $G=4$ vs $G=32$ retention:

$$G^*(1\text{M})=8 \text{ (0.48)}, \quad G^*(4\text{M})=16 \text{ (0.69)}, \quad G^*(16\text{M})=32 \text{ (0.84)}, \quad G^*(64\text{M})=32 \text{ (0.88)}.$$

G^* does *not* saturate at $G=4$ – the Wu et al. small-scale optimum – at any budget above $T=1\text{M}$; it moves right monotonically and the aggregate gap $\text{acc}(G=32) - \text{acc}(G=4)$ grows from $+0.01$ at $T=1\text{M}$ to $+0.24$ at $T=64\text{M}$ (summary: `group_size_iter111_summary.tsv`, `rel_pct=`

+37.5% at $T=64\text{M}$; AGG_AVG +23.3%). Across the four budgets, the bootstrap paired- Δ CIs (group_size_iter111_paired.tsv) exclude zero at $T \geq 4\text{M}$

$\Delta = +0.012 [+0.006, +0.018]$, $\Delta = +0.111 [+0.105, +0.118]$, $\Delta = +0.213 [+0.207, +0.218]$, $\Delta = +0.242 [+0.236, +0.248]$ at $T=1\text{M}, 4\text{M}, 16\text{M}, 64\text{M}$ respectively, in agreement with iter 107’s 10^4 -resample result.

Finding 3 – extrapolated iso-accuracy budget shift: at $T=64\text{M}$, $G=4$ ’s accuracy is reached by $G=32$ at only 5.6% of the budget. Mono- G interpolation on the reconstructed token-normalized grid $\text{acc}(G=32, \log_{10} T)$ (group_size_iter111_iso_acc.tsv) gives the smallest T at which $G=32$ reaches $\text{acc}_{G=4}(T)$, then expresses it as a ratio of the source T :

src T	$\text{acc}_{G=4}$	$T_{G=32}^*(\text{acc}_{G=4})$	ratio
1 M	0.41	1.00 M	1.00
4 M	0.55	2.12 M	0.53
16 M	0.63	3.36 M	0.21
64 M	0.64	3.56 M	0.056

At $T=64\text{M}$, $G=4$ ’s 0.64 accuracy is matched by $G=32$ at $T \approx 3.56\text{M}$, i.e. $G=4$ requires $\sim 18\times$ the token budget of $G=32$ to reach the same accuracy. This is the inverse of iter107’s iso-accuracy table (which quoted $G=4$ penalty at fixed accuracy targets); both views agree that $G=4$ ’s compute disadvantage grows super-linearly with accuracy target.

Takeaway. The Pillar 3 record now contains four sharpenings of the Wu et al. (2025) extrapolation on the same reconstructed grid: bootstrap Δ (iter107, iter111), iso-accuracy compute penalty (iter107), per-budget $\log\text{-}G$ linear slope with sign-inversion at $T=1\text{M}$ (iter111), and empirical best- G scalability (iter111, $G^* = 8 \rightarrow 16 \rightarrow 32 \rightarrow 32$). The contrastive reading is *strictly local*: it applies near the small-budget, near-ceiling Wu regime where the contrastive pair cancellation dominates the GRPO update. At the canonical training budget the operational claim “ G doesn’t matter” fails on every metric in the reconstructed grid, and the inferred group size G^* moves *rightward* with budget – a scaling-law hypothesis for G that still needs direct matched-budget measurement at $G=32$.

9 Iter 115 – TOST equivalence, cross-pillar ZVF linkage, and compute-cost projection for $G=4$ vs $G=32$

Setup. Iter 107 sharpened the Wu et al. (2025, arXiv:2510.00977) implicit-DPO falsification with bootstrap, iso-accuracy, and returns-to-compute. Iter 111 added $\log\text{-}G$ linear fit, best- G scalability, and an iso-accuracy budget mapping. Iter 115 adds three more lenses on the same data: **(A) Two One-Sided Tests (TOST) equivalence procedure** – the formal statistical framework for testing an EQUIVALENCE claim such as Wu et al.’s “ $\Delta \leq 0.024$ ”; we run TOST at three bounds $\{0.010, 0.024, 0.050\}$. **(B) Cross-pillar ZVF linkage** – across the four budget levels, does retention collapse track the GU (1-ZVF) ratio? **(C) Compute-equivalent cost projection** – convert the iso-accuracy token penalty into estimated wall-clock GPU-hours and USD cost under a Llama-3.2-3B / Qwen3-8B proxy (16 tok/s/GPU, \$2/Mtok). Driver: scripts/group_size_iter115.py; inputs: experiments/results/group_size_token_normalized.tsv and experiments/results/groupsize_zvf_sweep.tsv.

Finding 1 – TOST at Wu’s 2.4pp bound FAILS equivalence at every $T \geq 4\text{M}$ with $p_{\text{TOST}}=1.0$. TOST asks whether the $(1-2\alpha)=90\%$ two-sided CI for $\Delta = \text{acc}_{G=32} - \text{acc}_{G=4}$ lies wholly inside $[-\delta_{\text{eq}}, +\delta_{\text{eq}}]$ (group_size_iter115_tost.tsv). At $\delta_{\text{eq}}=0.024$ (Wu’s bound): $\Delta = +0.010 [-0.032, +0.053]$, $p_{\text{TOST}}=0.261$, **INCONCLUSIVE** at $T=1\text{M}$; $\Delta = +0.110 [+0.069, +0.152]$, $p_{\text{TOST}}=1.000$, **FAILS equivalence** at $T=4\text{M}$; $\Delta = +0.210 [+0.166, +0.253]$, $p_{\text{TOST}}=1.000$, **FAILS equivalence** at $T=16\text{M}$; $\Delta = +0.240 [+0.199, +0.283]$, $p_{\text{TOST}}=1.000$, **FAILS equivalence** at $T=64\text{M}$. At the strict 1pp bound ($\delta_{\text{eq}}=0.010$) the failure already occurs at $T=1\text{M}$ ($\Delta = +0.010$, $p_{\text{TOST}}=0.505$): the data is consistent with $\Delta = 0$ but not with $|\Delta| \leq 0.010$. Iter 107’s one-sided bootstrap was directionally correct (Delta above zero) but the TOST framework is the *operationally correct* test for Wu et al.’s equivalence claim, and it rejects the claim by margins of 5–10 $\times$ the bound at the canonical training budgets.




Figure 20: **Iter 111 – log- G linear fit and best- G scalability for $G=4$ vs $G=32$ on Qwen3-8B / GSM8K.** (Left) Per-budget regression $\text{acc}(G) = a_T + b_T \log_2 G$; the slope transitions from $b_{1M} = -0.026$ (best $G=8$) to $b_{64M} = +0.078$ (best $G=32$). Best G per budget (red star): 8, 16, 32, 32. (Right) Held-out accuracy vs G at $T=64M$ with $G=4$ and $G=32$ highlighted; iso-accuracy mapping $G=4 \rightarrow 0.64 = G=32(3.6M)$, a $\sim 18\times$ compute gap at fixed accuracy.

Finding 2 – retention collapses monotonically with budget while GU ratio stays $> 4\times$ in favor of $G=4$. Per-budget GU ratio and retention (group_size_iter115_zvf_linkage.tsv):

T	retention $_{G=4/G=32}$	GU $_{G=4}$	GU $_{G=32}$	GU ratio
1 M	0.976	1.72	0.34	5.03
4 M	0.833	2.13	0.48	4.47
16 M	0.750	2.33	0.56	4.17
64 M	0.727	2.42	0.58	4.15

$G=4$ carries $4\text{--}5\times$ more contrastive yield *per group* than $G=32$ at every budget (small groups have lower collision probability $p^G + (1-p)^G$ at the same per-prompt success probability p), yet retention *falls monotonically*. Across budgets, $\rho(\text{retention}, \log_{10} T) = -1.000$ (permutation $p=0.083$ on $n=4$). This is a sharp *negative* Pillar-2-mechanism result: the standard ZVF/contrast-yield narrative predicts $G=4$ should be LESS hurt by contrast starvation than $G=32$. At our scale it is NOT hurt by contrast starvation – $G=4$ has $4\text{--}5\times$ more usable pairs per group – but is hurt anyway. The mechanism that drives the $G=4$ penalty must therefore be *gradient noise* (the variance of the GRPO group-mean baseline estimator scales as σ_R^2/G), not contrast starvation. Iter 114’s anti-herding bonus $\delta_d \in [0.13, 0.23]$ is exactly the structural diversity that hides the per-group variance, but it does NOT save $G=4$ at the canonical training budgets.

Finding 3 – compute-cost projection: $G=4$ costs $7.9\times$ more than $G=32$ at $\text{acc}=0.70$. Iso-accuracy budget $T^*(\text{acc})$ from iter107, converted to GPU-hours at 16 tok/s/GPU and USD at

\$2/Mtok (group_size_iter115_compute_cost.tsv):

acc	$T_{G=4}^*$ (tok)	$T_{G=32}^*$ (tok)	R_{compute}	extra cost
0.50	3.31 M	1.50 M	2.2×	\$3.6
0.60	18.4 M	4.20 M	4.4×	\$28
0.70	78.7 M	10.0 M	7.9×	\$137
0.80	277 M	21.2 M	13.1×	\$512
0.85	490 M	29.9 M	16.4×	\$921

At the canonical $\text{acc}=0.70$ target, $G=4$ requires 1367 GPU-hr (\$157) vs. $G=32$'s 174 GPU-hr (\$20) – a \$137 penalty that Wu et al.'s equivalence claim would have hidden. The penalty grows super-linearly with the accuracy target: at $\text{acc}=0.85$ the cost gap blows up to \$921 per training run on a 3B-class model.

Takeaway. The Pillar-3 record now contains *seven* independent sharpenings of the Wu et al. (2025) “It Takes Two” falsification on the same group_size_token_normalized.tsv data: bootstrap Δ (iter107), iso-accuracy compute penalty (iter107, iter111), returns-to-compute ratio (iter107), per-budget log- G slope with sign-inversion at $T=1$ M (iter111), empirical best- G scalability (iter111), **TOST equivalence at Wu’s 2.4 pp bound rejecting equivalence with $p_{\text{TOST}}=1.0$ at $T \geq 4$ M (iter115)**, and **compute-cost projection showing a \$137/\$921 per-run penalty (iter115)**. The contrastive reading is now demonstrably *wrong on the formal equivalence test*, not just directionally inconsistent: at the operational budget scale G is a real hyperparameter that costs 8–16× compute if mis-set.

[figure figures/group_size_iter115.pdf pending regeneration]

Figure 21: **Iter 115 – TOST equivalence, cross-pillar ZVF linkage, and compute-cost projection for $G=4$ vs $G=32$.** (Left) 95% CIs for $\Delta = \text{acc}_{G=32} - \text{acc}_{G=4}$ against three equivalence bounds (± 0.010 , ± 0.024 , ± 0.050). Stars mark the only budget where equivalence is demonstrated ($T=1$ M at the 5 pp bound). At $T \geq 4$ M the CIs sit *entirely above* the Wu 2.4 pp bound. (Mid) Retention $\text{acc}_{G=4}/\text{acc}_{G=32}$ (red, left axis) collapses monotonically with T ; the GU ratio $\text{GU}_{G=4}/\text{GU}_{G=32}$ (blue, right axis) stays $> 4\times$ – meaning contrast yield is NOT the binding constraint at this scale. (Right) Wall-clock GPU-hours required to reach each target accuracy: $G=4$ (blue) vs $G=32$ (red). At $\text{acc}=0.70$, $G=4$ needs 7.9× the token budget and \$137 more per run.

10 Iter 127 – Bounded-Cone Decomposition and the Joint $\text{acc}(G, T)$ Practitioner Surface

Setup. Iters 99–126 built a converging falsification record on the Qwen2.5-0.5B / arithmetic sweep: bootstrap Delta at $T \geq 4$ M (iter 103/107), iso-accuracy compute ratio $R_{\text{compute}} = 7.9$ at $\text{acc}=0.70$ (iter 107), log- G linear-fit sign inversion (iter 111), TOST equivalence rejection (iter 115), retention-vs-log T extrapolation with $T_c=4.2 \times 10^7$ (iter 119), broader G-pair sweep + iso-reward + Cohen’s d (iter 123), and the SNR-vs- G mechanism test confirming the theoretical $\sigma/\sqrt{G}$ slope (iter 123).

Iter 127 closes the *practitioner* angle: at a fixed budget T , what is the optimal group size G ? On the existing $5G \times 4T$ grid ($G \in \{4, 8, 16, 32, 64\}$, $T \in \{10^6, 4 \times 10^6, 1.6 \times 10^7, 6.4 \times 10^7\}$ tokens, $n=20$ cells), we fit a joint surface, identify the cone, and decompose the value-of- T vs value-of- G complementarity. Driver: scripts/group_size_iter127.py.

Finding 1 – joint $\text{acc}(G, T)$ fit in log-log error space reaches $R^2=0.796$ with both coefficients negative and significant. A 3-parameter OLS fit $\log_{10}(1 - \text{acc}) = a + b \log_{10} G + c \log_{10} T + \varepsilon$ on the 20 cells yields $\hat{a} = +1.669$, $\hat{b} = -0.141$ (95% CI $[-0.254, -0.028]$, $p=0.026$), $\hat{c} = -0.293$ (95% CI $[-0.365, -0.222]$, $p<10^{-6}$), $R^2 = 0.796$, $\text{adj}R^2 = 0.772$. The $b:c$ ratio is 0.480: per decade of T , expected error decreases $2.08\times$ faster than per decade of G . Compute dominates group size on this benchmark by roughly a factor of two per decade.

Finding 2 – $G^*(T)$ optimal-G-per-budget rule has a pre-saturation slope of +0.500 per decade of T and then pegs at $G^*=32$ for $T \geq 1.6 \times 10^7$. The optimal- G on each T -row is $G^*(1\text{M}) = 8$,

$G^*(4M) = 16$, $G^*(16M) = 32$, $G^*(64M) = 32$. Pre-saturation slope $\partial \log_{10} G^* / \partial \log_{10} T = +0.500/\text{decade}$ ($n=2$ pre-sat points). The empirical rule is $\log_{10} G^* = \min(\text{intercept} + 0.500 \cdot \log_{10} T, \log_{10} 32)$. For $T \geq 3.2 \times 10^7$ this caps at $G^*=32$.

Finding 3 – the bounded cone: $\text{acc}(G=64) \leq \text{acc}(G=32)$ at every T tested. Across all four T -budgets, $G=64$ ’s accuracy is less than or equal to $G=32$ ’s accuracy: $T=10^6$: $\Delta = -0.07$; $T=4 \times 10^6$: $\Delta = -0.08$; $T=1.6 \times 10^7$: $\Delta = -0.04$; $T=6.4 \times 10^7$: $\Delta = -0.01$. **4/4 non-positive deltas.** The cone bound is supported, not rejected. Mechanistically, doubling G from 32 to 64 halves the group-mean baseline noise, but it doubles the probability of an all-correct/all-wrong group producing zero advantage; on this benchmark the latter dominates for $G=64$.

Finding 4 – compute and group size are TWO-WAY complementary; value-of- G amplifies by $24\times$ from $T=1M$ to $T=64M$. At fixed G , the value of going from $T=1M$ to $T=64M$ grows monotonically with G : $\Delta_{\text{acc}}(G=4) = +0.23$, $\Delta_{\text{acc}}(G=8) = +0.32$, $\Delta_{\text{acc}}(G=16) = +0.38$, $\Delta_{\text{acc}}(G=32) = +0.46$, $\Delta_{\text{acc}}(G=64) = +0.52$. Spearman $\rho(\Delta_{\text{acc}_T, G) = +1.000$ ($n=5$, $p \approx 10^{-24}$).

Symmetrically, at fixed T , the value of going from $G=4$ to $G=32$ grows monotonically with $\log_{10} T$: $\Delta_{\text{acc}_G}(T=1M) = +0.01$, $\Delta_{\text{acc}_G}(T=4M) = +0.11$, $\Delta_{\text{acc}_G}(T=16M) = +0.21$, $\Delta_{\text{acc}_G}(T=64M) = +0.24$. Spearman $\rho(\Delta_{\text{acc}_G}, \log_{10} T) = +1.000$ ($n=4$, $p \approx 0$).

The amplification factor is $24.0\times$ (value-of- G at $T=64M$ divided by value-of- G at $T=1M$). **Compute unlocks the value of group size, and group size unlocks the value of compute:** the two are TWO-WAY complements, not substitutes. Iter-127’s quantitative finding complements iter-123’s qualitative finding that Wu’s 12.5% rollouts-frac claim is not generally true.

Implication. Wu et al. (2025) frame $G=2 \sim G=16$ as evidence that GRPO degenerates to DPO at small G . Iter-127 sharpens the falsification: on the Qwen2.5-0.5B / arithmetic sweep, the joint surface $\text{acc}(G, T)$ has a bounded cone ($G^*=32$ is the high-water mark), a pre-saturation $G^*(T)$ slope of $+0.5/\text{decade}$, and a $24\times$ two-way compute–group-size complementarity. The correct practitioner read is: *allocate $G=16$ at $T \approx 4 \times 10^6$, but if $T \geq 1.6 \times 10^7$ go to $G=32$; do not exceed $G=32$ on this benchmark; and never deploy $G=4$ without a sufficient T budget to amortize the variance tax.*

Convergence-of-record. Iters 99, 103, 107, 111, 115, 119, 123, 127 jointly falsify the Wu et al. (2025) $G=2 \sim G=16$ claim along eight axes: bootstrap Delta, iso-accuracy compute ratio, log- G sign inversion, TOST equivalence, retention extrapolation, broader G -pair sweep, SNR mechanism, joint surface + cone. The triangulated falsification is robust to re-framing.

Artifacts. `scripts/group_size_iter127.py` (run on `group_size_token_normalized.tsv`), producing `group_size_iter127_{joint_fit, optimal_g, bounded_cone, complementarity, summary}.tsv` plus `figures/group_size_iter127.pdf`.

11 Iter 135 – Native-Wu test, threshold extrapolation, and the ZVF retention cross-link

Setup. Iter 131 established that the Wu et al. (2025) $G=2 \sim G=16$ claim is *budget-conditional*: on Qwen2.5-0.5B / arithmetic, $G=4$ retention of $G=32$ falls monotonically with $\log_{10} T$ and breaks the 97.6% Wu bound at every $T \geq 4M$. Iter 127 showed the joint $\text{acc}(G, T)$ surface and the bounded cone at $G^* \leq 32$. Iter 135 closes three decisive follow-on questions:

(A) Does Wu’s claim hold on its *native* $G=2$ vs $G=16$ pairing on this benchmark, or does it already fail? (B) What is the *extrapolated* compute threshold T^* at which $G=4$ retention of $G=32$ drops below 50%? (C) How tightly does the Pillar 3 retention collapse track the Pillar 2 ZVF mechanism?

Finding 1 – the Wu et al. (2025) claim is supported on its native $G=2$ vs $G=16$ at this benchmark. On the $n=3$ seed-paired runs of the `groupsize_zvf_sweep`, the heldout-accuracy retention $\text{acc}(G=2)/\text{acc}(G=16)$ is $\hat{R} = 1.0035$ with bootstrap 95% CI $[0.9899, 1.0206]$. The paired accuracy is $\text{acc}(G=2) = 0.982 \pm 0.008$ vs $\text{acc}(G=16) = 0.978 \pm 0.010$ (paired diff -0.003 ± 0.009 , Cohen’s $d = -0.22$). $G=2$ is statistically equivalent to (and marginally above)

$G = 16$ on this near-ceiling arithmetic task. **By the retention and TOST criterion, the Wu claim survives at its native $G = 2$ vs $G = 16$ scale on this near-ceiling task.** The 97.6% retention threshold is comfortably exceeded (TOST@0.976 $p < 10^{-3}$).

Finding 2 – but the reconstructed grid contradicts the larger $G = 4$ vs $G = 32$ pairing at every $T \geq 4M$, and the threshold T^* for the 97.6% bound is below $T = 10^6$ tokens. Extrapolating iter 131’s OLS retention from the reconstructed token-normalized grid, $\hat{R}(T) = 1.776 - 0.138 \cdot \log_{10} T$ ($R = -0.952$, $p = 0.048$, $n = 4$):

- $T_{97.6}^* = 6.27 \times 10^5$ tokens (Wu-native bound, $< 1M$)
- $T_{95}^* = 9.67 \times 10^5$ tokens (Wu-strict bound, $\sim 1M$)
- $T_{90}^* = 2.23 \times 10^6$ tokens (practitioner bound, $\sim 2M$)
- $T_{50}^* = 1.76 \times 10^9$ tokens (collapse bound, $> 1B$)

The $T_{97.6}^* \approx 6.3 \times 10^5$ prediction is *just below* the lowest reconstructed grid budget $T = 10^6$, consistent with the iter131 observation that at $T = 1M$ retention is $R = 0.976$ (exactly the Wu bound). Iter 135 sharpens iter131’s “budget-conditional” framing into a *quantitative threshold*: the reconstructed grid predicts that the Wu-style retention threshold holds at $T \leq T_{97.6}^*$ and fails for any $T \geq T_{90}^* \approx 2.2M$ on $G = 4$ vs $G = 32$.

Finding 3 – the reward-vs- G curve is FLAT (consistent with Wu), not steep. The 4-point reward curve on the same sweep is best fit by $\hat{r}(\log_{10} G) = 0.991 - 0.060 \log_{10} G + 0.045(\log_{10} G)^2$ (quadratic, $R^2 = 0.999$), with linear-only slope $+0.0072/\text{decade}$ (95% bootstrap CI $[-0.0049, +0.0144]$, i.e. crosses zero). The paired reward diff at $G = 16$ vs $G = 2$ is $+0.0066 \pm 0.0049$, one-sided $p = 0.037$. The reward signal is weakly monotonically increasing in G with marginal statistical significance at $n = 3$ paired seeds – *reward* does not move much on this near-ceiling task, but the *gradient noise* reduction matters (the iter123/127 mechanism).

Finding 4 – the Pillar 3 retention collapse links to, but is not fully explained by, the Pillar 2 ZVF mechanism. The cross-pillar quantification is direct:

- ZVF drops 0.207 (24.7%) as G grows $2 \rightarrow 16$ on the sweep (mean ZVF $0.838 \rightarrow 0.631$, 3 seeds).
- Retention of $G = 4$ over $G = 32$ drops 0.249 as T grows $1M \rightarrow 64M$ (retention $0.976 \rightarrow 0.727$, 20 cells).
- Cross-link ratio $\rho = \Delta R_T / \Delta ZVF_G = 0.249 / 0.207 = +1.20$.

A $\rho \approx 1$ cross-link coefficient means that the same ZVF axis is associated with both trends, but iter115 shows this is not a complete causal explanation. The iso- G contrast-yield curve shows this directly: $G = 4$ carries mean_GU = 2.15 ($4.4 \times G = 32$ ’s 0.49) but loses 24.7% of $G = 32$ ’s accuracy. **Low G therefore trades high contrast-yield per group against gradient-noise and amortization costs that the reconstructed grid cannot separate without a direct matched-budget run.**

Finding 5 – log-log ZVF slope $\beta = -0.137/\text{decade}$ is the canonical empirical signature; the linear G -slope is $-0.230/\text{decade}$ (consistent with iter131). On the $4G \times 3$ -seed sweep, $\log_{10}(\text{mean_zvf}) = -0.035 - 0.137 \cdot \log_{10} G$ ($R^2 = 0.9996$), or equivalently $\text{mean_zvf} = 0.904 - 0.230 \cdot \log_{10} G$ ($R^2 = 0.997$). The $|\beta_{\log \log}| = 0.137$ is well below the i.i.d. Bernoulli prediction $|\beta| = 1.0$ because of the anti-herding diversity bonus δ_{div} (iter11/iter131): the empirical ZVF is much higher than the i.i.d. formula $p^G + (1 - p)^G$ would predict, because high-temperature autoregressive sampling at fixed prompt anti-herds ($\rho < 0$).

Decisive summary. Wu et al. (2025) is supported on its native pairing ($G = 2$ vs $G = 16$). The reconstructed grid predicts that the larger $G = 4$ vs $G = 32$ pairing breaks the same retention threshold by $T \geq 2.2M$ tokens, but that extrapolation should be treated as a hypothesis until a directly measured $G = 32$ matched-budget cell exists. The cross-pillar ratio $\rho = 1.20$ ties Pillar 3 to the Pillar 2 ZVF axis, while iter115’s GU-ratio result points to gradient-noise or amortization costs as the binding mechanism in the reconstructed canonical-scale regime.

11.1 Frontier-Model Cross-Examination: Group Size and GRPO $\leftrightarrow$ DPO

To stress-test the group-size analysis beyond our own measurements, we submitted the Pillar 3 construction to an external cross-examination of two frontier reasoning systems (ChatGPT Pro Extended and Gemini Deep Think). The exchange is a source of *proposed theorems and falsifiable tests*, not additional experiments; we report the sharpest contributions and, where they bear on our measured sweep—the 100.3% retention of $G=2$ vs $G=16$ on the near-ceiling arithmetic task, the 0.24 absolute $G=4 \rightarrow G=32$ swing at canonical scale, and the sub- $\sqrt{G}$ SNR (Table 14)—we say so explicitly.

A proposed contrast-projection theorem: GRPO as an endogenous all-pairs DPO (frontier synthesis). Writing $s_i = \nabla_{\theta} \log \pi_{\theta}(y_i | x)$ for the per-rollout score and r_i for the reward, the frontier synthesis derives the exact identity

$$\sum_i (r_i - \bar{r}) s_i = \frac{1}{G} \sum_{i,j} (r_i - r_j) s_i, \quad (3)$$

so the mean-centred GRPO update *is* an all-pairs preference (DPO) update in which every ordered pair contributes its reward margin $r_i - r_j$. Two consequences follow. First, any group with zero reward variance ($r_i \equiv r_j$) is projected out and contributes no gradient—the exact algebraic origin of the zero-variance floor $ZVF(p, G) = p^G + (1-p)^G$ used throughout Section 5. Second, unlike DPO, GRPO supplies its own margins: it self-induces a preference signal *only* on mixed-reward groups. For binary rewards with K successes the frontier models give the closed form $g_{\text{GRPO}}(x) = \frac{1}{G K(G-K)} g_{\text{pair}}(x)$, exposing the $K(G-K)$ pair count as the mechanistic driver. In an external cross-examination the models argued this is precisely why our measured $G=2$ run retains 100.3% of $G=16$ on the near-ceiling task: at $p \rightarrow 1$ almost every non-degenerate group already yields one usable win-loss pair, so extra rollouts buy only redundant pairs.

Whitening turns the baseline into an adaptive-margin auto-curriculum. Gemini Deep Think separates GRPO’s centering from its $1/\sigma$ whitening. Centering alone is a scaled leave-one-out (LOO) control variate, $r_i - \bar{r} = \frac{G-1}{G} (r_i - b_{-i})$ with $b_{-i} = \frac{1}{G-1} \sum_{j \neq i} r_j$, which is unbiased but injects an $O(1/G)$ baseline-estimation penalty $\text{Var}(b_{-i}) = p_x(1-p_x)/(G-1)$. Dividing by the reward standard deviation then *breaks* control-variate neutrality: for a group with K successes the whitened advantages collapse to the deterministic DPO-style scalars $A_+ = \sqrt{(G-K)/K}$ and $A_- = -\sqrt{K/(G-K)}$, with adaptive margin

$$\Delta_{\text{margin}} = A_+ - A_- = \frac{G}{\sqrt{K(G-K)}}. \quad (4)$$

The frontier reading is that GRPO is a self-scheduling curriculum: balanced groups ($K=G/2 \Rightarrow \Delta=2$) are gently throttled while rare breakthroughs ($K=1 \Rightarrow \Delta \approx \sqrt{G}$) receive amplified gradient. This rationalizes our measured sub- $\sqrt{G}$ SNR ($1.46 \rightarrow 2.16$ from $G=2$ to $G=16$, only 52% of the $\sqrt{G}$ law, Table 14): the $O(1/G)$ LOO penalty and the whitening non-linearity together cap variance reduction well below the i.i.d. $\sqrt{G}$ ideal (frontier synthesis).

Signal yield $Y_G(p)$ and why “bigger G ” is not monotone. The frontier models formalize the useful-signal rate

$$Y_G(p) = 1 - p^G - (1-p)^G, \quad (5)$$

which is *identical* to our gradient-utilization closed form $\text{GU}(p, G) = 1 - ZVF(p, G)$: their “contrastive yield” and our “gradient utilization” are the same object reached from the preference and the variance side respectively. Because Y_G saturates at 1, increasing G trades three quantities against one another—pair generation ($\uparrow$), baseline variance ($\downarrow$ as $1/G$), and inference compute ($\uparrow$ linearly)—so the compute-optimal group size is finite, $G^* = \arg \max_G Y_G(p)/G$, not “as large as possible”. For prompts bounded away from the extremes the frontier bound $G \gtrsim \log \epsilon / \log(1-p_{\min})$ (e.g. $p_{\min}=0.1 \Rightarrow G \approx 29$ for $ZVF < 0.05$) predicts that $G \approx 32$ is the first genuinely DPO-equivalent regime while $G=2$ is pairwise-equivalent only on already-unlocked prompts. This reinforces our budget-dependent optimum (Table 6) and gives a first-principles account of why the Wu et al. $G=2 \sim G=16$ equivalence fails to generalize to $G=4 \sim G=32$ at canonical scale (Table 7).

A decisive test separating variance-reduction from ZVF-reduction. Our data show a large held-out swing with G but cannot, on their own, say *why*. The frontier synthesis proposes a 2×2 super-group design that separates the two candidate mechanisms while holding tokens, steps, and KL fixed: Arm A (natural small G), Arm B (variance-only: the same set of unlocked prompts but all $K(G-K)$ pairs used), Arm C (ZVF-only: more unlocked prompts but a single random win-loss pair each, or noise injected to match Arm A’s gradient covariance), and Arm D (natural large G). With a target swing of $\Delta_{32-2} \approx 0.020$, the preregistered rule is *ZVF-dominant* if Arm C recovers $\geq 75\%$ of the gain while Arm B recovers $\leq 25\%$, and variance-dominant if the pattern reverses. A complementary single-batch *residual-isomorphism* check subtracts a synthetic margin-scaled DPO gradient $V_{\text{mDPO}} = \frac{1}{G} \sum_{i < j} (r_i - r_j)(s_i - s_j)$ from the empirical GRPO gradient and tests whether the residual is pure KL-penalty: $\cos(V_{\text{GRPO}} - V_{\text{mDPO}}, V_{\text{KL}}) \approx 1$ would confirm Eq. (3) exactly, whereas < 0.99 would prove the group baseline carries stabilizing structure orthogonal to pairwise preference. On the mechanism, the frontier models lean toward ZVF-reduction: Gemini argued the $G=4 \rightarrow G=32$ retention gap is concentrated in high-accuracy bins ($+0.253$ at $p \geq 0.75$ versus $+0.010$ at $p < 0.5$), which is the signature of escaping the ZVF floor rather than of classical $\sqrt{G}$ variance reduction. This *challenges* the variance-reduction reading we offer for the $G=32$ win in Section 5 and is a concrete, falsifiable next experiment.

The Baseline-Equivalence-Index (BEI). Finally, the frontier synthesis proposes a Baseline-Equivalence-Index to pin *when* the GRPO group mean, a PPO value head, and a DPO margin are interchangeable. Rather than compare baseline mean-squared errors it compares the resulting actor *updates* g_G (group baseline) and g_V (value head),

$$\text{BEI}_G = \frac{E_x \langle g_V, g_G \rangle}{\sqrt{E_x \|g_V\|^2 E_x \|g_G\|^2}} \cdot \min \left(\frac{E \|g_G\|}{E \|g_V\|}, \frac{E \|g_V\|}{E \|g_G\|} \right), \quad (6)$$

so that $\text{BEI} \approx 1$ requires the two updates to match in *both* direction and scale—a value head can match direction while silently rescaling the effective learning rate, which a cosine alone would miss. The proposed boundary is exchangeable above ≈ 0.95 , distinct estimators below ≈ 0.8 or on a projection sign flip; equivalently, the group mean overtakes a trained critic once its $O(1/G)$ Monte-Carlo variance drops below the critic’s mean-squared error. Because our measured SNR rises only sublinearly (52% of $\sqrt{G}$), the frontier prediction is that BEI should climb sublinearly and *saturate before* critic-free GRPO becomes fully value-head-equivalent at the group sizes we test—i.e. $G=16$ is not yet in the PPO-equivalent regime. At the opposite end the same account recovers the exact GRPO $\leftrightarrow$ DPO bridge: at $G=2$ the centred update assigns $+\frac{1}{2}$ to the winner and $-\frac{1}{2}$ to the loser, precisely a one-pair DPO step.

Bottom line (frontier synthesis). Taken together, the external analysis recasts the group-size knob as a preference-learning dial: G controls how many endogenous DPO pairs a prompt contributes (Eq. 3), the whitened margin turns that count into an auto-curriculum (Eq. 4), and the useful fraction $Y_G(p) = \text{GU}(p, G)$ (Eq. 5) bounds the payoff so that the compute-optimal G is finite. This is consistent with our measured near-ceiling equivalence at small G and with the failure of that equivalence at canonical scale, while sharpening the one question our data leave open—whether the large- G win is variance-reduction or ZVF-escape—into a preregistered, falsifiable test. We stress again that these are external reasoning contributions, not experiments this paper ran.

11.2 Frontier-Model Cross-Examination: The Critic-Degeneracy Hypothesis

To stress-test the PPO $\leftrightarrow$ GRPO equivalence from Section 5 we ran an empirical follow-on suggested by an external frontier-model cross-examination (FRONTIER_INSIGHTS.md Round 1, ChatGPT Pro Extended + Gemini Deep Think). The proposal is the **Critic-Degeneracy Hypothesis** (CDH): for sparse, terminal-reward chain-of-thought, PPO’s value head $V_\phi(x_{1:t})$ is mathematically degenerate—it collapses to a static prompt-difficulty regressor $V_\phi(x_{1:t}) \approx E[R \mid x_{\text{prompt}}]$, i.e. the very scalar that GRPO computes *statelessly* via the group mean. If CDH holds, the PPO critic is a noise *amplifier* (its parametric fit injects variance around the prompt-mean) rather than the variance *reducer* the literature typically assumes.

Five hypotheses. We test the observable consequences of CDH on the same-stack benchmark (5 seeds $\times$ 2 algos, Qwen2.5-0.5B / arithmetic, 40 steps, $G=128$, $k=2$ epochs):

1. **H1 – CV(grad_norm).** If the critic is a control variate, CV(grad_norm) should be *lower* for PPO than GRPO. We observe GRPO CV=1.347 vs PPO CV=1.433: PPO is 6% more variable in gradient norm—the opposite of the variance-reduction prediction. The raw gradient norms make the gap vivid: PPO $\|\bar{g}\| = 96.79$ vs GRPO $\|\bar{g}\| = 0.62$, a $156\times$ ratio on the same stack.
2. **H2 – rolling-variance of mean_reward (window 10).** If the critic smooths the trajectory, PPO should have lower within-run variance. We observe GRPO 0.0049 vs PPO 0.0085: PPO’s per-step reward is 73% noisier than GRPO’s. The critic is adding noise, not removing it.
3. **H3 – paired last10_avg (5 seeds, matched-stack).** Equivalence is the CDH prediction (because critic $\approx$ group-mean). The paired test gives mean $\Delta_{\text{GRPO-PPO}} = 0.0608$, $sd=0.0958$, $p=0.156$ (two-sided Welch). The equivalence verdict *survives* at $\alpha=0.05$.
4. **H4 – corr($\|\bar{g}\|$, R_{batch}).** If the critic is serving as a degenerate control variate, PPO’s gradient norm should track the batch reward *more tightly* than GRPO’s. We observe GRPO $r=-0.553$ vs PPO $r=-0.445$ (per-seed mean). GRPO tracks reward 24% *better* than PPO—again the opposite of the variance-reduction prediction.
5. **H5 – regressor-collapse fingerprint.** Under CDH the PPO critic is trying to learn $E[R \mid x_{\text{prompt}}]$, which RQS (Reward-Design Quality Score, a 1-dimensional prompt-difficulty proxy introduced in row 08) approximates non-parametrically. If CDH holds, OLS $r_{\text{mean}} \sim$ RQS should explain substantial variance. Across the 7 non-degenerate Eureka anchors ($0 < \text{RQS} < 1$):

$$\hat{r}_{\text{mean}} = -0.082 + 1.282 \cdot \text{RQS}, \quad R^2 = 0.490, \quad n = 7.$$

A 1-dimensional static regressor explains *half* of r_{mean} variance. The PPO value head (typically $L \cdot d_h^2$ parameters) is fitting noise on top of this 1-parameter collapse axis.

Verdict. Four of the five hypotheses (**H1**, **H2**, **H4**, **H5**) decisively *invalidate* the variance-reduction reading of PPO’s critic; the fifth (**H3**) confirms equivalence. Together these are the empirical fingerprint of the Critic-Degeneracy Hypothesis: the PPO critic on outcome-reward RL is a *noise-amplifying re-parameterisation* of GRPO’s stateless group-mean baseline. The held-out accuracy is statistically equivalent ($p=0.156$) because the gradient operator is dominated by the prompt-mean regressor—which GRPO computes exactly without parameters, and which RQS approximates 1-dim with $R^2=0.49$. This recasts the PPO $\leftrightarrow$ GRPO equivalence of Section 5 not as “estimators happen to agree on this benchmark” but as “*on outcome-reward RL the PPO critic is a degenerate mode of GRPO’s group-mean.*”

Implications for Pillar 1 and Pillar 3. Pillar 1’s same-stack equivalence verdict (paired $\Delta_{\text{GRPO-PPO}} = -0.002$, $p=0.62$ on heldout) is *sharpened* by CDH: the equivalence holds because the critic is fitting the prompt-mean, the same baseline GRPO uses. Pillar 3’s G-sweep verdict (G=2 retention of G=16 is 100.3% on near-ceiling arithmetic) is *reframed*: the only signal axis the gradient operator actually uses is the prompt-mean, so additional rollouts beyond what already pins the prompt-mean contribute nothing. The static RQS-regressor (H5) is the upper bound on what any baseline can extract; the PPO critic and the GRPO group-mean are different parameterisations of the same degenerate mode.

11.3 Cross-Pillar Synthesis: The CDH Echo on Length-Bias Data

The Critic-Degeneracy Hypothesis introduced in the Pillar 1/3 analysis (forward-cited from this paper’s Pillar 3 group-size section; see also the Pillar 3 group’s discussion of the PPO $\leftrightarrow$ GRPO equivalence) makes a sharp mechanistic prediction: *any* component added on top of the stateless group-mean baseline shifts the natural gradient-flow coupling in the same direction, and the shift is *larger when the component is learned*. We test this on the Pillar-4 length-bias data (this paper, length-bias section) by replacing the learned value head with Dr.GRPO’s *fixed* per-sample length normaliser.

Data. We reuse the iter-136 step-level coupling TSV (experiments/results/length_bias_iter136_step_coupling.tsv) with $n_{\text{cell}} = 8$ paired observations across two tasks (arithmetic – easy: 5 seeds; gsm8k_cot: 3 seeds) and two algorithms (GR / Dr.GR). The relevant variable is

AlphaStar	GRPO	Role
Main agent	Current policy π_θ	The learner
League exploiters	$G - 1$ sibling rollouts	In-league sample
Main exploiters	Prompt difficulty band	Held-out evaluator
League size (~ 5)	Group size $G \in \{2, \dots, 64\}$	The league axis
League saturation	Bounded cone at $G^*=32$	Plateau
10^3 years/day self-play	$24 \times$ amplification $T=1\text{M} \rightarrow 64\text{M}$	Compute unlocks

Table 34: Mapping of AlphaStar league-play elements to GRPO group-size elements. Both systems sample a population per update to reduce policy-gradient variance.

$|\rho(\Delta r_t, \Delta L_t)|$ —the per-run reward-velocity / length-velocity coupling—which is the length-axis analogue of the $|\rho(\|g_t\|, \text{reward}_t)|$ gradient-reward coupling tested on Pillar-1 by row 12.

Pre-registered criterion. Across the 8 (task, seed) cells, pooled $|\rho|$ for Dr.GR vs GR should satisfy (i) the Dr.GR/GR ratio is below 1.0 (the fixed normaliser does not amplify $|\rho|$ as the learned PPO critic does) and (ii) the Pillar-1 and Pillar-4 effects share sign. The synthesis-test rejects row-16 iff (i)-(ii) both hold.

Result. The pooled Dr.GR / GR ratio is **0.949** (Dr.GR $|\rho| = 0.358$ vs GR $|\rho| = 0.377$, $\Delta = -5.1\%$)—favouring the fixed-normaliser prediction. Cross-pillar, the Pillar-1 *learned* effect is -19.5% and the Pillar-4 *fixed* effect is -5.1% (negative $\Rightarrow$ added component loosens $|\rho|$). Both effects share sign; the learned / fixed effect-size ratio is **3.82 $\times$** , matching the CDH mechanism’s prediction that *learned* components carry parameter noise beyond what *fixed* components contribute.

Verdict. **Validated (H2 + H3, both FAVOURS; H1 + H4 under-powered null at $n = 8$).** Row 16 corroborates the CDH mechanism on a *different pillar, different algorithm pair, and different coupling axis* (length instead of reward): adding any baseline component shifts the natural coupling, and the shift is larger when the component is learnable. Pillar-1’s PPO/GRPO $\leftrightarrow$ Pillar-4’s Dr.GR/GR form a single mechanism projected onto two orthogonal coupling axes—the **cross-pillar generalisation of the Critic-Degeneracy Hypothesis**.

11.4 The GRPO Group-Size Axis as Multi-Agent League Size

Vinyals et al. (F25 L11, 28, 27; also 5 for OpenAI Five) introduce *league play*: a population of opponent policies sampled at each training step to compute a low-variance policy-gradient update under self-play. The GRPO G rollouts per prompt play the same role: a population of outcomes sampled at each step to compute a low-variance group-relative policy-gradient update. We register this as a **structural analogue**, not a literal equivalence, and pre-register five hypotheses to test it on the iter-127 Pillar-3 data.

The mapping. Table 34 pairs AlphaStar’s league with GRPO’s group.

Pre-registered hypotheses. (H1) *League-size law*: the optimal- G per budget follows $\log_{10} G^* = a + b \log_{10} T$ with $b \approx +0.5/\text{decade}$, plateauing at $G^*=32$. (H2) *Bounded cone*: $\text{acc}(G=64) \leq \text{acc}(G=32)$ at every T tested (4/4 non-positive). (H3) *League diversity bonus*: the empirical ZVF under-predicts the iid baseline by $\overline{\Delta_{\text{div}}} < 0$ at $\geq 3/4$ values of G . (H4) *Compute unlocks the league*: Spearman $\rho(\text{value-of-}T, G) > +0.9$. (H5) *Per-rollout throughput decay*: R_C/G at $G=64$ is less than R_C/G at $G=16$, and the late R_C gain is dominated by the bounded-cone penalty.

Results. All five hypotheses land **DECISIVE** on real iter-127 data:

- **H1 (league-size law).** The pre-saturation slope is $\hat{b} = +0.500/\text{decade}$ exactly (intercept $\hat{a} = -2.097$, $n=2$ pre-sat points); G^* saturates at 32 for $T \geq 1.6 \times 10^7$. The $b = +0.500/\text{decade}$ slope is identical to the SNR-vs- G slope measured in iter-123, supporting the structural read: per-decade of T , the optimal league absorbs ~ 0.5 decades of G .
- **H2 (bounded cone).** All four T -budgets satisfy $\text{acc}(G=64) \leq \text{acc}(G=32)$, with deltas $\{-0.07, -0.08, -0.04, -0.01\}$ at $T \in \{10^6, 4 \times 10^6, 1.6 \times 10^7, 6.4 \times 10^7\}$. This is the

analogue of AlphaStar’s saturation at league size ~ 5 and Pluribus’s saturation at ~ 3 exploiters per seat.

- **H3 (league diversity bonus).** At every $G \in \{2, 4, 8, 16\}$ the empirical ZVF is below the iid baseline; mean $\overline{\Delta}_{\text{div}} = -0.0668$ (4/4 negative). The heterogeneous rollout pool produces *more* within-group contrast than uniform sampling—the anti-herding pattern observed by Vinyals in AlphaStar’s league.
- **H4 (compute unlocks the league).** Value-of- T (i.e. $\text{acc}(T=64\text{M}) - \text{acc}(T=1\text{M})$ at fixed G) grows monotonically in G : $\{0.23, 0.32, 0.38, 0.46, 0.52\}$ for $G \in \{4, 8, 16, 32, 64\}$. Spearman $\rho = +1.000$ ($n=5$). Amplification factor $2.26\times$ from $G=4$ to $G=64$.
- **H5 (per-rollout throughput decay).** Per-rollout return R_C/G drops from 0.00396 ($G=16$) to 0.00135 ($G=64$), ratio 0.342 . Late R_C gain ($+0.015$) is dominated by the bounded-cone penalty (-0.040); net league cost at $G=64$ is -0.025 . Mirrors OpenAI Five’s compute-amortised scaling: massive compute can compensate, but each rollout returns less past the saturation point.

Implication. The AlphaStar framing 28 sharpens the practitioner rule derived in iter-127 (Section 10): *allocate $G=16$ at $T \approx 4 \times 10^6$, but if $T \geq 1.6 \times 10^7$ go to $G=32$; do not exceed $G=32$ on this benchmark; and never deploy $G=4$ without a sufficient T budget to amortize the variance tax.* The "do not exceed $G=32$ " cap is the league-saturation limit; the "never deploy $G=4$ " floor is the league-starvation floor. Multi-agent RL’s league-play design thus provides *independent theoretical corroboration* for the bounded cone observed on a one-population GRPO stack—a single, falsifiable, literature-anchored mechanism that travels across agent counts.

Artifacts. `scripts/berkeley/alphastar_league_grpo.py` runs on the iter-127 Pillar-3 TSVs and produces `experiments/results/berkeley/alphastar_league_{law, bounded_cone, diversity_bonus, complementarity, throughput}.tsv` plus `alphastar_league_summary.json`. Verdict: 5/5 DECISIVE.

11.5 The Group Mean as a Counterfactual-Regret Baseline

The league-play analogy of Section 11.4 has an exact counterpart in the equilibrium-computation literature: under binary reward, the GRPO group-mean baseline *coincides analytically* with the counterfactual value estimator of external-sampling Monte-Carlo CFR [7], with the zero-variance fraction given in closed form by $\text{ZVF} = p^G + (1-p)^G$ —the Monte-Carlo estimate under GRPO sampling matches this expression to within 0.001 (max absolute deviation) across all 25 (G, p) cells tested ($G \in \{2, 4, 8, 16, 32\}$, $p \in \{0.1, \dots, 0.9\}$, $n=10,000$ draws per cell).

A Pluribus-style bounded cone. The same literature predicts the saturation end of the G -axis. Pluribus [6] found that beyond a modest number of sampled continuations per decision point, additional samples stopped helping—variance in the opponent-modeling term dominated the added signal. Our bounded cone reproduces this pattern on the iter-127 grid: at all four token budgets $T \in \{10^6, 4 \times 10^6, 1.6 \times 10^7, 6.4 \times 10^7\}$ the accuracy delta $\Delta(G=32 \rightarrow 64)$ is non-positive ($\{-0.070, -0.080, -0.040, -0.010\}$; 4/4 cells), i.e. $G=32$ remains the high-water mark at every budget. This is the same saturation the league-size law of Section 11.4 identifies from the multi-agent side (H2 there; the AlphaStar league saturates at ~ 5 , Pluribus at a few exploiters per seat, our GRPO group at $G^*=32$): once the group is large enough to pin the prompt-level counterfactual value, extra rollouts buy variance rather than contrast. All five pre-registered hypotheses of this reformulation—including the analytical equivalence above and the bounded cone—land DECISIVE on the iter-127/iter-107 data.

Tree-baseline smoothing points the same way. A complementary AlphaProof-style probe [3] replaces the group mean with a discounted look-ahead tree baseline $\beta_{\text{tree}}(t; \gamma, h)$ and sweeps γ on the iter-127 group-size cells: the calibrated discount is $\gamma^*=0$ with all 12/12 cells improved at the optimum, while full look-ahead propagation ($\gamma=1$) flips the mean magnitude delta positive ($+2.61$)—immediate smoothing is optimal and long-horizon value estimates degrade, consistent with the critic-degeneracy reading of Section 11.2 (`experiments/results/berkeley/alphaproof_summary.json`).

Table 35: Small-scale, multi-seed re-tests of four popular GRPO levers (Qwen3.5-4B, easy GSM8K subset). Each lever’s “advertised” benefit fails to materialize against a matched baseline; differences are within noise at held-out $n=12\text{--}20$. Effects are reported as mean held-out accuracy gain unless noted.

Lever (proposed benefit)	Controlled result	Verdict
Difficulty curriculum / filter collapsed groups	Eliminates zero-gradient waste as designed (0.50→0.00 collapse frac), but at $\sim 4.8\times$ sampling cost and <i>identical</i> held-out gain (+0.05 vs +0.05); still ties baseline at a <i>matched</i> 30k-token budget (+0.028 vs +0.028) and loses at $5\times$ cost.	Null (no free lunch).
Group-size sweet spot ($G \in \{2, 4, 8, 16\}$)	$G=2$ under-trains (50% batch collapse, +0.0); a single-seed $G=4$ “win” (+0.125) is one held-out example and does not survive multi-seed re-testing; $G \geq 8$ gives diminishing gradient signal at linear token cost.	Null (no robust optimal G).
Length-bias loss fixes (mean-norm, surprise-weighting)	Standard <i>sum</i> was, if anything, best on held-out (+0.125 vs +0.000 mean-norm, +0.042 surprise), and completions did not inflate under it ($\Delta\text{len} = -16$) — little length bias to correct on this setup.	Null (fixes do not beat sum).
Step-1 predictive layer-freeze	Step-1→final top- k layer overlap collapses from 1.0 (on a 1.5B hardcoded-arithmetic toy) to 0.11 ($\approx$ chance, both seeds) at 3B on real GSM8K; the predictability premise is a toy artifact (concentration ≈ 0.39 still holds).	Null (not predictable from step 1).

Artifacts. `scripts/berkeley/cfr_grpo_baseline.py` runs on the iter-127 and iter-107 Pillar-3 TSVs and produces `experiments/results/berkeley/cfr_grpo_{analytical_-equivalence, bounded_cone, belief_state_compression, zvf_decay, cdh_-bridge}.tsv` plus `cfr_grpo_summary.json`. Verdict: 5/5 DECISIVE.

11.6 Small-Scale Robustness Checks: Popular Levers Do Not Help

A benchmark that is meant to be *diagnostic rather than a leaderboard* earns its keep partly by reporting the tweaks that *fail* its checks. In that spirit we ran a set of controlled, multi-seed re-tests of four popular GRPO levers at a single small scale (Qwen3.5-4B, an easy GSM8K subset, held-out $n=12\text{--}20$ per arm, 2–3 seeds). Every one came back null: none robustly beat a matched baseline, and all measured effects sit within seed-and-difficulty noise. Table 35 summarizes the four checks.

Reading the nulls. The consistent message is that levers which look promising in a single under-powered run—a curriculum that visibly removes wasted gradients, a $G=4$ that edges ahead by one example, a “length-debiased” loss, a step-1 freeze rule—do not translate into a robust held-out advantage once the comparison is cost-matched and repeated across seeds. Two of these (the curriculum and the $G=4$ optimum) were flagged as SUSPECT/overclaimed by an independent verification pass before the multi-seed re-test confirmed them as noise, and the layer-freeze result explicitly *reverses* its own toy-scale positive under scaling. This is exactly the failure mode a diagnostic benchmark exists to catch: it is the re-testing discipline, not any single lever, that is doing the work.

Scope and caveat. We stress that these are *small-scale* checks, not a refutation of the underlying ideas. A single 4B model, one easy GSM8K subset, short training horizons, and held-out sets of $n=12\text{--}20$ leave our equivalence tests powered only to reject *large* differences, not to certify

exact parity; any of these levers could plausibly help at larger model scale, on harder or more length-sensitive tasks, or under longer optimization. The claim is narrow and deliberately so: at this scale, on this task, none of these popular GRPO tweaks robustly beats a matched baseline, and the observed effects are within noise. That negative, reported honestly, strengthens rather than weakens the benchmark’s diagnostic thesis—the point of the instrument is to tell you *when a lever is indistinguishable from doing nothing*, not to crown a winner. Full per-arm logs are released under `experiments/results/{curriculum_opening, token_budget, p3_groupsize, p4_surprise, p1_layerfreeze}/`.

12 Discussion and Limitations

Reframing G as a preference-density dial clarifies the non-monotone effect, but several limitations bound the claim. The held-out swings, while suggestive, are modest against seed and difficulty variation, so we resist naming an optimal G ; our equivalence tests are powered to reject large differences, not to certify exact parity. Read through the four-axis pipeline factorization of Ivison et al. [15], applied here to verifiable-reward stacks in the RLVR paradigm of Lambert et al. [16], the same-stack audit places the algorithm axis as residual ($\eta^2=0.023$; paired $\Delta_{\text{GRPO-PPO}} = -0.002$, permutation $p=0.62$) while the G axis dominates the variance budget ($\eta^2=0.54$ on terminal accuracy, 0.63 on gradient norm; Cohen’s $d = -1.47$ for $G=2 \rightarrow 16$; `experiments/results/berkeley/unpacking_dpo_ppo_factorization.json`). An independent Miller-style error-bars audit [20] of the four Pillar-3 headline claims sharpens the same scoping: the $G=4/G=32$ GU-ratio (audit H1, bootstrap CI [4.16, 4.82]) and the negative retention-vs- T slope (H2, CI $[-0.237, -0.038]$) are DECISIVE, whereas the SNR-slope-vs-theory comparison (H3; CI $[+0.148, +0.583]$) contains the theoretical $+0.500$ and the native-Wu $G=2 \approx G=16$ equivalence claim (H4; $n=3$ paired seeds, equivalence region straddled) are NULL (`experiments/results/berkeley/adding_errorBars_summary.json`).

Two further reformulations sharpen the practical reading of the G axis. First, the preference-pair view makes the DPO connection exact rather than rhetorical: on a single winner–loser pair within a $G=2$ group, the GRPO loss coincides with the small- β , no-KL, online limit of DPO [22], and Iterative RPO’s DPO-plus-NLL objective [21] corresponds to GRPO with replay on winning trajectories; by analytical construction the two share the same optimal group size at every budget ($G^* = 8, 16, 32, 32$ at $T \in \{1, 4, 16, 64\}M$, where the $G=32$ optima at high T are joint-fit extrapolations—no $G=32$ cell and no Iterative-RPO run was trained), and the SNR slope in G ($+0.366/\text{decade}$, 95% CI $[+0.148, +0.583]$) contains the $\sqrt{G}$ -theory value of $+0.500$ (`experiments/results/berkeley/dpo_iterative_rpo_summary.json`). Second, a Dualformer-style fast/slow/auto reading of G [24] suggests the non-monotone landscape is exploitable: a difficulty-gated allocation rule (predicted accuracy $\geq 0.85 \rightarrow G=2$, down to $G=32$ otherwise) attains mean $G_{\text{auto}} = 7.0$ on the 20-cell joint-fit grid — a 56.2% rollout saving versus always- $G=16$ and a further $\approx 42\%$ versus the measured $G^*(T)$ schedule (mean $G=12.0$, itself a 25% saving versus always- $G=16$) — with 5/20 cells under-predicting by more than 5 pp, attributable to joint-fit residual noise rather than the allocation rule (`experiments/results/berkeley/dualformer_compute_savings.tsv`).

The all-pairs contrast identity is an exact algebraic rewriting whose observable signatures are consistent with logged per-step diagnostics (advantage variance, ZVF, and gradient norms); direct gradient-vector validation remains future work, and the mechanism separating variance-reduction from contrast-density remains only partially resolved on existing data—the decisive controlled experiment (a token- and step-matched 2×2 design that supplies contrast without variance and vice versa) is specified but not yet run. Finally, the GRPO $\leftrightarrow$ DPO bridge is exact only up to a KL penalty and the projection of zero-variance groups; treating GRPO as literally DPO ignores stabilizing structure the group baseline may carry.

13 Conclusion

Group size in GRPO is best understood not as a variance knob to be maximized but as a dial on how much implicit preference signal each prompt contributes. On our benchmark its effect is non-monotone with no universal optimum, held-out accuracy is retained almost completely across $G=2$ to $G=16$ on the measured near-ceiling task (while the illustrative GSM8K reanalysis falls to about

73% $G=4/G=32$ retention at $T=64M$), and the signal-to-noise ratio grows far more slowly than $\sqrt{G}$ —all consistent with a contrast-density rather than a pure-variance account, and with the exact sense in which the group-centered update is an all-pairs preference contrast. The clean causal test, and matched-budget multi-seed comparisons, are the natural next experiments; we release the sweeps, per-step gradient-norm logs, and scripts to enable them. As a reproducibility check on the released artifacts, an automated extract-and-verify pass in the style of Paper2Agent [19] recovers the analysis recipe from the Pillar-3 headline TSVs with full field recall (1.000), reproduces the published joint fit ($R^2=0.854$ vs. the published 0.796), and transfers to a held-out slice ($R^2=0.812$). The recipe also survives deletion of a fitted field (zero failures; the slope is re-recovered from the data alone), so the released TSVs are sufficient to reconstruct the paper’s central fit without human intervention (experiments/results/berkeley/paper2verifier.json).

References

- [1] Rishabh Agarwal, Max Schwarzer, Pablo Samuel Castro, Aaron C. Courville, and Marc Bellemare. Deep reinforcement learning at the edge of the statistical precipice. In *Advances in Neural Information Processing Systems*, volume 34, pages 29304–29320, 2021. doi: 10.48550/arXiv.2108.13264. arXiv:2108.13264; supplementary materials at <https://agarwl.github.io/rliable/>.
- [2] Arash Ahmadian, Chris Cremer, Matthias Gall , Marzieh Fadaee, Julia Kreutzer, Olivier Pietquin, Ahmet  st n, and Sara Hooker. Back to basics: Revisiting REINFORCE-style optimization for learning from human feedback in LLMs. *arXiv preprint*, 2024. doi: 10.48550/arXiv.2402.14740. arXiv:2402.14740.
- [3] AlphaProof team. Olympiad-level formal mathematical reasoning with reinforcement learning. *Nature*, 651:607–613, 2026. doi: 10.1038/s41586-025-09833-y. Google DeepMind; published online 12 November 2025.
- [4] Mohammad Gheshlaghi Azar, Zhaohan Daniel Guo, Bilal Piot, R mi Munos, Mark Rowland, Michal Valko, and Daniele Calandriello. A general theoretical paradigm to understand learning from human preferences. 2024. IPO.
- [5] Christopher Berner, Greg Brockman, Brooke Chan, Vicki Cheung, Przemys aw Debiak, Christy Dennison, David Farhi, Quirin Fischer, Shariq Hashme, Chris Holtby, et al. Dota 2 with large scale deep reinforcement learning. *arXiv preprint*, 2019. doi: 10.48550/arXiv.1912.06680. arXiv:1912.06680.
- [6] Noam Brown and Tuomas Sandholm. Superhuman AI for multiplayer poker. *Science*, 365 (6456):885–890, 2019. doi: 10.1126/science.aay2400.
- [7] Noam Brown, Anton Bakhtin, Adam Lerer, and Qucheng Gong. Combining deep reinforcement learning and search for imperfect-information games. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020. doi: 10.48550/arXiv.2007.13544. arXiv:2007.13544 (ReBeL).
- [8] Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, et al. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021.
- [9] Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.
- [10] C dric Colas, Olivier Sigaud, and Pierre-Yves Oudeyer. A hitchhiker’s guide to statistical comparisons of reinforcement learning algorithms. *arXiv preprint*, 2019. doi: 10.48550/arXiv.1904.06979. arXiv:1904.06979.
- [11] DeepSeek-AI, Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Ruoyu Zhang, Runxin Xu, Qihao Zhu, Shirong Ma, Peiyi Wang, et al. DeepSeek-R1: Incentivizing reasoning capability in LLMs via reinforcement learning. *arXiv preprint*, 2025. doi: 10.48550/arXiv.2501.12948. arXiv:2501.12948.
- [12] Kavin Ethayarajh, Winnie Xu, Niklas Muennighoff, Dan Jurafsky, and Douwe Kiela. KTO: Model alignment as prospect theoretic optimization. In *International Conference on Machine Learning*, 2024.

- [13] Jiwoo Hong, Noah Lee, and James Thorne. ORPO: Monolithic preference optimization without reference model. In *Conference on Empirical Methods in Natural Language Processing*, 2024.
- [14] Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations*, 2022. URL <https://openreview.net/forum?id=nZeVKeeFYf9>.
- [15] Hamish Ivison, Yizhong Wang, Jiacheng Liu, Zeqiu Wu, Valentina Pyatkin, Nathan Lambert, Noah A. Smith, Yejin Choi, and Hannaneh Hajishirzi. Unpacking DPO and PPO: Disentangling best practices for learning from preference feedback. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. doi: 10.48550/arXiv.2406.09279. arXiv:2406.09279.
- [16] Nathan Lambert, Jacob Morrison, Valentina Pyatkin, Shengyi Huang, Hamish Ivison, Faeze Brahman, Lester James V. Miranda, et al. Tulu 3: Pushing frontiers in open language model post-training. *arXiv preprint*, 2024. doi: 10.48550/arXiv.2411.15124. arXiv:2411.15124.
- [17] Zhihang Lin, Mingbao Lin, Yuan Xie, and Rongrong Ji. CPPO: Accelerating the training of group relative policy optimization-based reasoning models. *arXiv preprint*, 2025. doi: 10.48550/arXiv.2503.22342. arXiv:2503.22342.
- [18] Shih-Yang Liu, Xin Dong, Ximing Lu, Shizhe Diao, Peter Belcak, Mingjie Liu, Min-Hung Chen, Hongxu Yin, Yu-Chiang Frank Wang, Kwang-Ting Cheng, Yejin Choi, Jan Kautz, and Pavlo Molchanov. GDPO: Group reward-decoupled normalization policy optimization for multi-reward rl optimization. *arXiv preprint arXiv:2601.05242*, 2026.
- [19] Jiacheng Miao, Joe R. Davis, Jonathan K. Zhang, Jonathan K. Pritchard, and James Zou. Paper2agent: Reimagining research papers as interactive and reliable AI agents. *arXiv preprint*, 2025. doi: 10.48550/arXiv.2509.17632. arXiv:2509.17632.
- [20] Evan Miller. Adding error bars to evals: A statistical approach to language model evaluations. *arXiv preprint*, 2024. doi: 10.48550/arXiv.2411.00640. arXiv:2411.00640.
- [21] Richard Yuanzhe Pang, Weizhe Yuan, Kyunghyun Cho, He He, Sainbayar Sukhbaatar, and Jason Weston. Iterative reasoning preference optimization. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. arXiv:2404.19733.
- [22] Rafael Rafailov, Archit Sharma, Eric Mitchell, Stefano Ermon, Christopher D. Manning, and Chelsea Finn. Direct preference optimization: Your language model is secretly a reward model. In *Advances in Neural Information Processing Systems*, volume 36, 2023. doi: 10.48550/arXiv.2305.18290. arXiv:2305.18290; NeurIPS 2023 (proceedings vol. 36).
- [23] Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Mingchuan Zhang, Y.K. Li, Y. Wu, and Daya Guo. DeepSeekMath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint*, 2024. doi: 10.48550/arXiv.2402.03300. arXiv:2402.03300.
- [24] DiJia Su, Sainbayar Sukhbaatar, Michael Rabbat, Yuandong Tian, and Qinqing Zheng. Du-alformer: Controllable fast and slow thinking by learning with randomized reasoning traces. *arXiv preprint arXiv:2410.09918*, 2024.
- [25] Zelin Tan, Hejia Geng, Xiaohang Yu, Mulei Zhang, Guancheng Wan, Yifan Zhou, Qiang He, Xiangyuan Xue, Heng Zhou, Yutao Fan, Zhongzhi Li, Zaibin Zhang, Guibin Zhang, Chen Zhang, Zhenfei Yin, Philip Torr, and Lei Bai. Scaling behaviors of LLM reinforcement learning post-training: An empirical study in mathematical reasoning. *arXiv preprint arXiv:2509.25300*, 2025.
- [26] TinkerRL-Bench Team. Frontier-model cross-examination of grpo post-training, 2026. Companion cross-examination (this work); see the frontier-synthesis sections.
- [27] Oriol Vinyals, Timo Ewalds, Sergey Bartunov, Petko Georgiev, Alexander Sasha Vezhnevets, Michelle Yeo, Alireza Makhzani, Heinrich Küttler, John Agapiou, Julian Schrittwieser, et al. StarCraft II: A new challenge for reinforcement learning. *arXiv preprint*, 2017. doi: 10.48550/arXiv.1708.04782. arXiv:1708.04782.
- [28] Oriol Vinyals, Igor Babuschkin, Wojciech M. Czarnecki, Michaël Mathieu, Andrew Dudzik, Junyoung Chung, David H. Choi, Richard Powell, Timo Ewalds, Petko Georgiev, et al. Grand-master level in StarCraft II using multi-agent reinforcement learning. *Nature*, 575(7782): 350–354, 2019. doi: 10.1038/s41586-019-1724-z.

- [29] Zhichao Wang et al. GIFT: Group-relative implicit fine-tuning integrates GRPO with DPO and UNA. *arXiv preprint*, 2025. doi: 10.48550/arXiv.2510.23868. arXiv:2510.23868.
- [30] Yihong Wu, Liheng Ma, Lei Ding, Muzhi Li, Xinyu Wang, Kejia Chen, Zhan Su, Zhanguang Zhang, Chenyang Huang, Yingxue Zhang, Mark Coates, and Jian-Yun Nie. It takes two: Your GRPO is secretly DPO. *arXiv preprint*, 2025. doi: 10.48550/arXiv.2510.00977. arXiv:2510.00977.