
Report the Stack, Not the Label: RL-for-LLM Results Are Stack-Conditioned

Arvind C R*
PES University
arvindcr4@gmail.com

Sandhya Jeyaraj*
PES University
sandhya.jeyaraj2014@gmail.com

Madhu Kumara L
PES University
madhukumara1993@gmail.com

Mohammad Rafi
PES University
gmd.rafi.2024@gmail.com

Dhruva N Murthy
PES University
dhruva.n.murthy@gmail.com

Arumugam K
PES University
chettayarumugam@mail.com

Anwesh Reddy Paduri
Great Learning / PES University
anwesh@greatlearning.in

Narayana Darapaneni
Northwestern University / Great Learning
narayana.darapaneni@northwestern.edu

Ramesh Prakash Guledgudd†
PES University
rameshpg@pes.edu

Abstract

Papers comparing RL algorithms for LLM post-training report a *label*—GRPO, DAPO, Dr.GRPO, GSPO—and treat the surrounding software stack as an implementation detail. Our position is that this convention is backwards: on the evidence available to us, the stack is a material part of the reported result, and a label without its stack is not a reproducible experimental specification. Three exhibits motivate the reporting standard. First, a matched-configuration backend audit—group size, LoRA setup, learning rate, data, prompt template, decoding, and seed harmonized across frameworks—moved mean last-10 training reward from 5.0% to 85.6% when the training backend was exchanged: a $17\times$ swing produced without touching a single labeled algorithmic knob (the managed backend also silently pinned a different base checkpoint, and that undisclosed bundling is itself part of the exhibit). Second, the same label denotes different experiments across stacks: “DAPO” run on an open trainer with true dynamic sampling (a toy-scale arithmetic audit) produced a mean Zero-Variance Fraction (ZVF) of 0.00, while “DAPO” approximated on a closed stack (GSM8K) via an asymmetric-clip surrogate produced mean ZVF 0.58—same name, opposite telemetry. Third, a 12-cell head-to-head of four labeled algorithms on a single fixed stack landed within a 0.034 band of one another (mean last-10 reward 0.710–0.744), a spread comparable to seed-level variation, suggesting that label-only comparisons can hide stack-level variation. We therefore propose a seven-item *minimum reportable stack*—loss

*Equal contribution.

†Project guide.

form, reference-policy/KL handling, sampler/backend/precision, per-step ZVF/GU trajectory, group-size schedule, disjoint held-out split, and decontamination plus a parser-robustness probe—each item earning its place via a documented ranking-relevant lever. We formalize an RL-reproducibility threat model with flip-risk grades R0–R5, specify a toolchain (manifest emitter, stack diff with CI verdicts, registry, auditor badge, and citation-lever tracer) that makes the standard enforceable, and close with a call to action for venues.

1 Introduction: The Sandwich Problem

When a paper reports that “DAPO outperforms GRPO,” the named algorithm is a thin layer sandwiched between two thick, largely undocumented slices. Beneath the label sits the serving and training substrate: the sampling engine and its numerics [18, 43], attention kernels and precision [6], the trainer’s tokenization and masking conventions, and the exact loss implementation shipped by the framework [37, 34, 14, 36]. Above the label sits the reward and evaluation apparatus: the verifier and its answer parser, the held-out split, the prompt template, and whatever decontamination was (or was not) performed [40, 31]. The label names the filling; the experiment is the sandwich. We call the resulting failure of reference the *sandwich problem*: two papers using the same algorithm label routinely describe different experiments, and two papers using different labels routinely differ more in their bread than in their filling.

This is a position paper. Its position is: **RL-for-LLM results are stack-conditioned, and venues should require authors to report the stack, not merely the label.** We do not claim that algorithmic labels are meaningless—the loss-form differences among GRPO [33, 7], DAPO [39], Dr.GRPO [22], and GSPO [42] are real and sometimes consequential. We claim that on the evidence available to us the stack’s contribution to reported outcomes is at least as large as the label’s, and that the community’s reporting conventions are calibrated to the wrong term. A $17\times$ outcome swing from swapping only the backend (Section 5) dwarfs every label-vs-label gap in our corpus; a label can even invert its own telemetry when re-implemented on a different stack. Deep RL went through this reckoning once already [11, 5, 1]; RL-for-LLMs has re-created the conditions with a larger and less inspectable stack.

Our contributions are deliberately those of a position paper: evidence, a standard, a threat model, and enforcement machinery.

1. **Evidence** (Section 5): four exhibits—a $17\times$ backend flip under otherwise identical configuration; a cross-stack label flip in which “DAPO” yields mean ZVF 0.00 on one stack and 0.58 on another; a 12-cell head-to-head in which four labeled algorithms on one fixed stack land within noise; and cross-library ZVF variation from the companion Pillar-2 paper of TINKERRL-BENCH.
2. **The seven-item minimum reportable stack** (Section 6): a checklist small enough to enforce, in which every item is justified by a documented lever that can move or flip a reported ranking.
3. **An RL-reproducibility threat model** (Section 9): a taxonomy of confound vectors and a flip-risk grading R0–R5 that classifies how dangerous a given cross-paper comparison is.
4. **An enforceability toolchain** (Section 10): a specification for a manifest-emitting trainer plugin (`trl-min-report-rl`), a manifest differ with CI verdicts (`grpo-stackdiff`), a machine-readable stack registry, a 0–100 auditor badge, and a citation-lever tracer (`LeverTrace`).

Throughout, we scope claims to their evidence: benchmark results come from TINKERRL-BENCH (Section 2); several exhibits come from internal program records that we describe as such; and toy-scale theory validations are marked directional. The paper ends with limitations and a concrete call to action for venues (Section 13).




Figure 1: Taxonomy of RL libraries in TINKERRL-BENCH. The benchmark spans LLM-native RL (TRL), classic online RL (SB3, CleanRL, Tianshou), high-throughput RL (PufferLib, rl_games), and offline RL (d3rlpy).

Table 1: Benchmark task suite with standardized reward functions.

Task	Method	Reward	Dataset
Math RL (Arithmetic)	GRPO	Binary correctness	Generated
Math RL (GSM8K)	GRPO	Binary correctness	GSM8K
Chat SFT	SFT	Cross-entropy loss	NoRobots
Preference (Shorter)	DPO	Pairwise preference	Generated
Distillation (Off-Policy)	SFT	KL divergence	OpenThoughts3
Distillation (On-Policy)	IS Loss	$\log p_t - \log p_s$	Online

2 Benchmark Design

2.1 Task Suite

2.2 Cross-Library Implementation

Two seven-library rosters — read carefully. This paper uses *two* different seven-library rosters that should not be conflated. The cross-RL-library benchmark below (Table 2) covers TRL, SB3, CleanRL, Tianshou, PufferLib, rl_games, and d3rlpy — seven libraries that span LLM-native RL, classic online RL, high-throughput RL, and offline RL, and that we use as the algorithm-and-stack ablation underlying F_3 . The cross-launcher LLM-RL scaffold referenced in the framework-gap discussion (framework-configurations material) is a different seven-library roster — Tinker, TRL, SkyRL, verl, OpenRLHF, Atropos, and HF reference launchers — that targets LLM post-training launchers specifically. Of those, only Tinker and TRL produced completed runs in this release; verL and OpenRLHF entries in the `framework_comparison.json` artefact are dry-run placeholders. The viva slides use the second roster; the cross-RL-library evidence in F_3 uses the first.

Table 2: Implementations across RL libraries.

Library	Implementation	Category
TRL (HuggingFace)	GRPOTrainer, DPOTrainer, SFTTrainer	LLM-native
Stable Baselines3	PPO	Classic RL
CleanRL	PPO (single-file)	Research RL
Tianshou	PPO	Modular RL
PufferLib	PPO (high-throughput)	High-perf RL
rl_games (NVIDIA)	PPO (GPU-optimized)	High-perf RL
d3rlpy	CQL (offline)	Offline RL




Figure 2: Verifiable reward paradigm in TINKERRL-BENCH. Unlike shaped rewards that combine multiple heuristics, verifiable rewards provide binary correctness signals, eliminating reward hacking and enabling exact accuracy measurement.

2.3 Hyperparameter Mapping

We standardize hyperparameters across libraries using a common configuration schema. Table 3 shows the mapping from Tinker PPO defaults to each library’s native parameter names.

Table 3: Hyperparameter mapping across libraries (PPO defaults).

Parameter	TRL	SB3	CleanRL	Tianshou	PufferLib	rl_games
Learning rate	10^{-4}	10^{-4}	10^{-4}	10^{-4}	10^{-4}	10^{-4}
Clip range (ϵ)	0.2	0.2	0.2	0.2	0.2	0.2
Entropy coeff.	0.01	0.01	0.01	0.01	0.01	0.01
GAE λ	0.95	0.95	0.95	0.95	0.95	0.95
Discount γ	0.99	0.99	0.99	0.99	0.99	0.99
Max grad norm	0.5	0.5	0.5	0.5	0.5	0.5
Target KL	0.01	0.01	0.01	N/A	N/A	N/A

2.4 Reward Verification

2.5 Baselines: Floor and Ceiling

Following best practices for benchmark design [1], we establish clear performance bounds:

- **Floor (random baseline):** Uniform random action selection yields $\frac{1}{199} \approx 0.50\%$ accuracy on the arithmetic task.
- **Ceiling (oracle):** Direct computation achieves 100% accuracy.
- **Human baseline:** College-level adults achieve $\sim 99.5\%$ on two-digit addition under timed conditions.

3 Experimental Setup

3.1 Models

Our benchmark targets a roster of 32+ model configurations spanning **0.6B to ~ 671 B total parameters (up to ~ 37 B active for MoE checkpoints) across 5 model families**, organized into three tiers by compute scale; not all are completed end-to-end runs (per-run status flags accompany the released artifacts). Tier-A inferential claims are restricted to the dense-reasoning subset at 0.6B–235B; the frontier MoE checkpoints (DeepSeek-V3.1 ~ 671 B / 37B active, Kimi-K2) are descriptive single-seed Tier-C case studies, and Qwen3.5-397B-A17B is a *partial* run whose held-out evaluation did not complete.

Small scale (0.6B–8B). Qwen3-0.6B-Instruct, Qwen3.5-4B, Qwen3-4B, Qwen3-8B, Llama-3.2-1B, Llama-3.2-3B, Llama-3.1-8B-Instruct.

Mid scale (14B–32B). Qwen3-14B, Qwen3-30B-A3B (MoE), Qwen3-32B, Qwen3.5-27B, GPT-OSS-20b, Kimi-K2 (selected scale points).




Figure 3: Experiment workflow pipeline. Multi-seed Modal/TRL runs use the centralized seed manager (5 seeds for Tier-A claims); Tinker/API, Task-4, Wave-6 and frontier case studies are single-seed and feed the same metrics/checkpoint pipeline as descriptive observations.

Frontier scale (70B–~671B total / ~37B active). Qwen3-235B-A22B (MoE, 22B active), DeepSeek-V3.1 (~671B total / ~37B active MoE), Nemotron-120B (120B dense). Frontier models are evaluated via the Tinker API in inference mode with standardized generation parameters; LoRA fine-tuning at this scale is conducted on selected checkpoints (see the frontier case studies in the companion pillar papers of TINKERRRL-BENCH).

Model family coverage. The benchmark covers (1) **Qwen3**: a dense family from 0.6B to 32B; (2) **Qwen3.5**: an improved series including 4B and 27B; (3) **Llama-3.x**: Meta’s 1B–8B instruction-tuned models; (4) **DeepSeek-V3.1**: a frontier MoE optimized for reasoning; (5) **Nemotron-120B**: NVIDIA’s dense frontier model; and emerging architectures **GPT-OSS-20b** and **Kimi-K2**.

3.2 Training Protocol

Unless otherwise noted, controlled Modal/TRL sweeps use LoRA [13] at rank 32, learning rate 10^{-4} and Adam ($\beta_1=0.9$, $\beta_2=0.95$, $\epsilon=10^{-8}$). Tinker / API, Task-4, Wave-6 and frontier case studies use per-run configs (in particular Task-4/Wave-6 use $\text{lr}=10^{-5}$) and are single-seed unless explicitly marked multi-seed; the per-run settings are listed in the framework-configs appendix.

Seed management. The controlled TRL/Modal sweeps used for Tier-A claims (F3 PPO-vs.-GRPO heterogeneity and the five-seed TRL GRPO baseline on Qwen2.5-0.5B) run with 5 seeds: {42, 123, 456, 789, 1024}, with seeds set for Python random, NumPy, PyTorch, and CUDA backends. Several Tinker/API, frontier-model, and team-task runs are single-seed or partial and are treated as descriptive (Tier-C) unless explicitly marked otherwise; the per-claim seed ledger is given in the statistical-rigor addendum of TINKERRRL-BENCH.

3.3 Statistical Methodology

Following Colas et al. [5], we report:

- Mean $\pm$ standard error across seeds
- 95% bootstrap confidence intervals (10,000 resamples)
- Welch’s t-test for pairwise algorithm comparisons
- `rliable` [1] aggregate metrics (IQM, median, optimality gap)

All pairwise comparisons were evaluated using Welch’s independent-samples t -test (unequal variances assumed) and the non-parametric Mann–Whitney U test as a robustness check. Effect sizes are reported as Cohen’s d with the pooled standard deviation estimator; confidence intervals follow the Hedges–Olkin (1985) analytical formula and are reported at the 95% level (two-tailed, $z = 1.96$). Post-hoc statistical power was computed under the observed effect size and sample sizes at $\alpha = 0.05$.

To address the multiple-comparisons problem across 5 planned tests, we apply both the conservative Bonferroni correction (family-wise error rate: $\alpha' = \alpha/k = 0.0100$) and the Benjamini–Hochberg procedure (false discovery rate < 0.05). Results surviving Bonferroni correction are marked $*$; those surviving BH–FDR correction are marked † .

Power Analysis. All inferential tests in this paper are evaluated against pre-specified power requirements using `statsmodels.stats.power.TTestIndPower` with $\alpha = 0.05$ and target power $1 - \beta = 0.80$.

Modal H100 experiments ($n = 5$ seeds per arm). The minimum detectable effect size (MDE) at 80% power for a two-sample Welch t -test with $n_1 = n_2 = 5$ is $d_{\min} = 2.024$ (Cohen’s d). Table 4 reports achieved power for a range of effect sizes. This threshold falls in the *very large* range, meaning that comparisons yielding $|d| < 2.02$ are not adequately powered to detect true differences.

Single-seed Tinker experiments ($n = 1$). Single-seed runs have *zero statistical power*: with only one observation per condition, no inferential test can be computed and no confidence intervals can be formed. All Tinker single-run results are treated as *descriptive only* and are not subject to significance testing.

TRL baseline ($n = 5$ seeds). The TRL-GRPO baseline (Qwen2.5-0.5B, 5 seeds) achieves MDE $d_{\min} = 2.024$ for equal- n comparisons. When compared to the pooled Classic-RL group ($n = 15$), the MDE improves to $d_{\min} = 1.530$, reflecting higher power from the larger reference group.

Multiple-Comparison Correction. Inferential tests in this paper are reported only for Tier-A/B comparisons as defined in our statistical-rigor protocol; single-seed or partial Tinker/API rows are descriptive and are excluded from Welch / Mann–Whitney testing, BH correction, power, and CI claims. Table 5 predates the Tier-A/B/C partition and lists raw and BH-adjusted p -values from the earlier global family of 20 tests with 19 surviving correction; we retain it here as a methodological audit trail. The headline BH result the paper now stands behind is the restricted Tier-A/B family computed in the addendum, where only F_3 survives. Where the global family disagrees with the restricted family, the addendum is authoritative.

All tests are two-sided. Effect sizes are reported as Cohen’s d for t -tests and Pearson/Spearman r for correlations. Bootstrap 95% confidence intervals ($B = 10,000$) are reported for all means.

Table 4: Statistical power for two-sample Welch t -tests at $\alpha = 0.05$, power target $1 - \beta = 0.80$. Column (3): equal groups $n_1 = n_2 = 5$ (Modal H100 / TRL within-group). Column (4): unequal groups $n_1 = 5, n_2 = 15$ (TRL vs. pooled Classic RL).

Cohen’s d	Conventional size	Power ($n=5$ vs $n=5$)	Power ($n=5$ vs $n=15$)
0.2	small	0.059	0.066
0.5	medium	0.108	0.151
0.8	large	0.201	0.311
1.0	very large	0.286	0.450
1.2	very large	0.386	0.594
1.5	very large	0.549	0.784
2.0	very large	0.790	0.955
<i>MDE at 80% power: $d = 2.024$ (equal-n 5 vs 5); $d = 1.530$ (5 vs 15)</i>			

TRL baseline characterisation. The TRL GRPO reference implementation was evaluated across five random seeds ($s \in \{42, 123, 456, 789, 1024\}$) on GSM8K using Qwen/Qwen2.5-0.5B on an NVIDIA L4 GPU. We report mean accuracy $\bar{x} = 0.734$ ($\sigma = 0.0703$), median = 0.740, and IQM = 0.747. Normality was not rejected by Shapiro–Wilk ($W = 0.9018, p = 0.4198$); bootstrap 95% CI = $[0.6720, 0.7820]$.

Variance decomposition. One-way ANOVAs across 42 experiments with valid performance observations show that library/framework ($\eta^2 = 0.546$) and training algorithm ($\eta^2 = 0.558$) are the dominant variance sources, consistent with our central thesis. Model family explains $\eta^2 = 0.471$; model scale explains $\eta^2 = 0.322$.

3.4 Compute Resources

Experiments are distributed across two compute platforms:

Table 5: **Legacy 20-test BH table – audit trail only.** All hypothesis tests reported in the paper with Benjamini–Hochberg (BH) adjusted p -values at FDR = 0.05. Tests are ranked by raw p -value (ascending). ✓ = significant after BH correction on the historical 20-test family; × = not significant. Total tests: 20; survive on the historical family: 19. *This figure is not the paper’s headline statistical result.* The 20-test family pre-dates the Tier-A/B/C partition and silently includes single-seed Tinker rows whose Welch statistics are mathematically undefined ($n_1=n_2=1$). Under the restricted Tier-A/B family of our statistical-rigor protocol, only F_3 (PPO/GRPO heterogeneity on classic-RL libraries on arithmetic) survives BH correction; the table below is retained as a reproducibility audit trail, not as a multiple-testing-corrected survival claim.

Rank	Comparison	Raw p	BH-adj. p	Sig.
1	Dense vs MoE Architecture (Welch t-test)	2.357e-21	0	✓
2	PPO vs GRPO on Llama-3.1-8B (Welch t-test) [‡]	3.924e-10	0	✓
3	Method ANOVA (F-test across algorithm families)	3.091e-06	2.1e-05	✓
4	Library ANOVA (F-test across 5 libraries)	4.996e-06	2.5e-05	✓
5	TRL vs Tinker (Welch t-test, implementation effect)	2.064e-05	8.3e-05	✓
6	PPO vs GRPO on Llama-3.1-8B (Mann-Whitney) [‡]	3.29e-05	0.00011	✓
7	Model Family ANOVA (F-test)	7.363e-05	0.00021	✓
8	ZVF vs Final Performance (Spearman) [‡]	0.0005424	0.001356	✓
9	ZVF vs Final Performance (Pearson) [‡]	0.0008108	0.001802	✓
10	TRL vs Tinker (Mann-Whitney)	0.001198	0.002396	✓
11	Rolling Variance vs Last-10 (Pearson)	0.003524	0.006407	✓
12	Peak-to-Tail Drift vs Last-10 (Pearson)	0.004811	0.007219	✓
13	GRPO vs PPO Stability Index (t-test)	0.005	0.007219	✓
14	Dense vs MoE Architecture (Mann-Whitney)	0.005053	0.007219	✓
15	Model Scale ANOVA (F-test)	0.01261	0.01682	✓
16	Tinker GRPO vs TRL GRPO (Welch t-test)	0.0158	0.01975	✓
17	GRPO vs PPO Peak-to-Tail Drift (t-test)	0.018	0.02076	✓
18	Tinker GRPO vs TRL GRPO (Mann-Whitney)	0.01869	0.02076	✓
19	Stability Index vs Last-10 (Pearson)	0.02024	0.0213	✓
20	PPO vs GRPO on Qwen3-8B (Welch t-test) [‡]	0.7605	0.7605	×

Tinker API (14 experiments). Scaling experiments across Qwen3-8B, Qwen3-32B, Qwen3.5-4B, Qwen3.5-27B, and Llama-3.1-8B-Instruct; frontier model evaluation on Qwen3-235B-A22B, DeepSeek-V3.1, and Nemotron-120B; Dense vs. MoE comparison; cross-task tool-use experiments; and emerging architecture evaluation (GPT-OSS-20b, Kimi-K2). Tinker provides scalable API access that makes frontier model experiments economically tractable.

Modal H100 GPU cluster (6 experiments). PPO baseline training on Qwen3-8B and Llama-3.1-8B; full HumanEval benchmark evaluation; KL divergence tracking across GRPO training steps; and held-out evaluation for Qwen3-32B and Qwen3.5-27B. Modal H100s provide the low-level GPU access required for PPO value-model training.

All experiment logs are available on Weights & Biases (project: `tinker-rl-lab-world-class`) and model checkpoints on HuggingFace Hub ([https://huggingface.co/arvindcr4/tinker-rl-bench-*](https://huggingface.co/arvindcr4/tinker-rl-bench-)). Per-run compute is logged with the released artifacts. Total project compute: ~1,200 A100 GPU-hours across both platforms.

4 Reading Guide: Evidence Base and Its Provenance

Because this is a position paper, we separate what we argue from what we measured, and we label the provenance of every number. Three evidence strata appear below. **Stratum A** is the TINKERRL-BENCH benchmark described in Section 2 and Section 3: controlled runs under the shared seed policy and statistical tiers, with artifacts in the released repository (per-library ZVF aggregations, matched-configuration framework dumps, and the same-stack PPO-vs-GRPO isolation runs). **Stratum B** is a set of internal program records from our June 2026 ZVF audit program (a 368-run Weights & Biases audit within a 790-run corpus, a 12-cell Tinker head-to-head, a backend-swap audit, and a single-file open-trainer reimplementaiton audit); we report these as internal experimental records, and the worked example in Appendix A shows the manifests that make such records comparable at all. **Stratum C** is toy-scale theory validation (a 0.5B model on synthetic arithmetic with an open backward pass); Stratum-C results are directional only and are flagged as such wherever used.

The argument then proceeds in four steps. Section 5 presents the exhibits: the stack moves outcomes by more than the label does, and the label does not pin down the experiment. Section 6 converts the exhibits into a standard—the seven-item minimum reportable stack—where each item is included if and only if we can point to a documented lever through which that item moved or flipped a result. Section 9 generalizes from exhibits to a threat model: the confound vectors an adversarial (or merely careless) comparison can ride, graded by flip risk. Section 10 shows the standard is cheap to enforce mechanically, which is what distinguishes a checklist that changes practice from one that decorates appendices [28].

5 Evidence: The Stack Moves More Than the Label

5.1 Exhibit 1: A $17\times$ Flip from Swapping Only the Backend

In a backend-swap audit (Stratum B), we held the model, group size G , learning rate, dataset, prompt template, and random seed fixed, and exchanged only the training backend. Final reward moved from 5.0% on one backend to 84.4% on the other—a $17\times$ ratio produced by zero algorithmic changes. Nothing in the run’s algorithm label, hyperparameter table, or seed disclosure would let a reader anticipate this; the difference lives entirely in the substrate (sampling engine, numerics/precision, loss implementation details, and masking conventions). For calibration, this single infrastructure swing is an order of magnitude larger than any label-vs-label gap we measured (Section 5.3), and is consistent with the broader observation that seemingly minor evaluation and serving details shift LLM results by margins that would be headline findings if attributed to algorithms [3, 12].

5.2 Exhibit 2: The Same Label Denotes Different Experiments

DAPO [39] is defined by decoupled (asymmetric) clipping *plus* dynamic sampling: groups whose completions are all-correct or all-incorrect are resampled, which by construction drives the Zero-Variance Fraction (the share of groups with identical within-group rewards; see the companion Pillar-2 paper of TINKERRL-BENCH) toward zero. In our cross-stack audit (Stratum B), “DAPO” run on an open Colab trainer with true dynamic sampling produced mean ZVF 0.00, exactly as the algorithm’s definition requires. The same label run on a closed training stack—where dynamic sampling is not expressible and the practitioner substitutes the standard surrogate of GRPO with asymmetric clipping—produced mean ZVF 0.58. Same label, opposite telemetry: one run never wastes a group, the other spends more than half its groups on zero-gradient updates. Any downstream comparison citing both runs as “DAPO” is comparing two different objects. This is not an indictment of either stack; it is an indictment of the label as an experimental specification.

5.3 Exhibit 3: Four Labels Within Noise on One Fixed Stack

Conversely, when the stack *is* held fixed, labels stop separating. In a 12-cell head-to-head on the Tinker stack (Qwen3.5-4B, GSM8K, three seeds per arm; Stratum B), the four labeled arms landed within noise of one another (Table 6): the full spread in mean last-10 training reward is 0.034 (0.710–0.744), comparable to seed-level variation under the TINKERRL-BENCH seed policy (Section 3). Mean ZVF varies somewhat more (0.500–0.578) but with no arm approaching the ZVF 0.00 that true dynamic sampling produces (Exhibit 2)—consistent with all four arms being, at this stack’s level of description, minor variations of one group-relative loss. The juxtaposition of Exhibits 1–3 is the core of our position: a factor the field reports meticulously (the label) moved outcomes by ≤ 0.034 reward on a fixed stack, while a factor the field routinely omits (the backend) moved outcomes by 0.794.

5.4 Exhibit 4: Cross-Library Variation in the Benchmark Itself

The pattern is not unique to our audits. Within TINKERRL-BENCH (Stratum A), the companion Pillar-2 paper documents per-library ZVF aggregations across the GRPO-family implementations of the benchmark roster (Table 2), with collapse, drift, and plateau rates that vary by library under matched task and model settings; the released matched-configuration dumps (one exact hyperparameter dump per framework for the canonical Qwen3-8B + GSM8K run) exist precisely because “the same configuration” across frameworks required a nontrivial mapping exercise to even define. The same-stack PPO-vs-GRPO isolation runs in the benchmark make the complementary point: once the stack

Table 6: 12-cell head-to-head on one fixed stack (Tinker; Qwen3.5-4B; GSM8K; 3 seeds per arm; internal program records, June 2026). Arms are within noise of one another on final reward; no arm reaches the ZVF regime that the same labels produce on stacks where their defining mechanisms are actually expressible.

Arm (label as configured)	Mean last-10 reward	Mean ZVF
dapo (asymmetric-clip surrogate)	0.744	0.578
drgrp	0.742	0.500
grp	0.723	0.567
gspo	0.710	0.511

Table 7: Variance-explained by axis on the live N2 corpus. The algorithm axis is the four-method same-stack N2 run (GRPO / AERO / GIFT / AREAL, $n = 40$ steps each, one seed); the group-size axis is the four- G sweep on the Tinker stack (3 seeds per arm). The algorithm axis explains $\leq 6.3\%$ of telemetry variance; the group-size axis explains $\geq 22\times$ more for ZVF and $\geq 133\times$ more for reward. Loss is excluded because it is method-defined (each method has its own loss surface).

Telemetry	$\eta^2(\text{algorithm})$	$\eta^2(G)$	ratio G/alg
ZVF	0.0454	1.0000	$22\times$
PCD	0.0357	—	—
reward_mean	0.0075	1.0000	$133\times$
mean_len	0.0631	—	—
cv_len	0.0457	—	—

is pinned, algorithm effects can be measured cleanly—which is the outcome our reporting standard is designed to make routine rather than exceptional.

5.5 Exhibit 5: Quantitative Decomposition—Algorithm vs Stack

Exhibit 3 shows the qualitative version: four labels within noise on one fixed stack. To make the comparison quantitatively testable, we decompose variance into $\eta^2 = \text{SS}_{\text{between}}/\text{SS}_{\text{total}}$ across two axes on the live N2 corpus (experiments/results/n2_reward_tensor_resume/n2_metrics.tsv and experiments/results/groupsize_zvf_sweep.tsv; analysed by scripts/p5p8/stack_eta2.py, see experiments/results/p5p8/stack_eta2.tsv):

The structural reading: when the stack is held fixed, swapping GRPO for AERO / GIFT / AREAL moves ZVF by $\leq 6.3\%$ of the variance budget; changing G moves it by an order of magnitude more. This is the quantitative counterpart to Exhibit 3 and the empirical content of the “Report the Stack, Not the Label” thesis: the algorithm axis is under-identified against the stack axis on this corpus. (Caveat: the algorithm-axis η^2 is from a single seed; the *direction* $\eta^2(\text{algorithm}) \ll \eta^2(G)$ is robust, the magnitude is not.)

5.6 Exhibit 6: MIN-REPORT Field Coverage at $n = 98$ Manifests

To stress-test the standard at scale, we audited every manifest in experiments/results/mega_20260704/manifests/ ($n = 98$ cells) for the seven MIN-REPORT items of Section 6 using scripts/p5p8/minreport_coverage.py; full table at experiments/results/p5p8/minreport_field_coverage.tsv:

All twelve sub-fields enumerated in the standard (ratio level, clip range, advantage normalisation, dynamic sampling, token mask, KL snapshot, KL coefficient, KL estimator, precision, decoding parameters, contamination check, parser probe) score 0/98 in the live corpus. The standard is therefore satisfiable as a shallow key set, but not yet enforceable as a structured stack record—the gap the MIN-REPORT-RL Auditor of Section 10 is designed to close.

5.7 Exhibit 7: Stack-axis η^2 on the 98-cell Live Mega Corpus

Exhibit 5 used the 4-method same-stack N2 run to compare the algorithm axis against the group-size axis ($n = 160$ per-step observations, one seed). To test whether the same direction holds at

Table 8: Per-item coverage of the seven-item MIN-REPORT standard against $n = 98$ live mega-campaign manifests. All seven items are 100% present *as keys*, but Item 2 (Reference policy & KL) reports a concrete `kl-*` value in 0/98 (0.0%)—the corpus is uniformly sampling-only with no KL regulariser, but the value “n/a” is indistinguishable from an emission bug.

#	Item	present	validated
1	Loss form	100.0%	100.0%
2	Reference policy & KL	100.0%	0.0%
3	Sampler / backend / precision	100.0%	100.0%
4	Per-step ZVF/GU trajectory	100.0%	100.0%
5	Group-size schedule	100.0%	100.0%
6	Held-out split	100.0%	100.0%
7	Decontamination & parser probe	100.0%	100.0%

Table 9: Per-axis η^2 on the 98-cell mega corpus. Bold = $\eta^2 \geq 0.20$ (DOMINANT verdict). Stack axes dominate every telemetry channel; seed contributives $\leq 0.15\%$ in all channels.

axis	reward	zvf	pcd	len	std_len
model_family	0.453	0.005	0.093	0.301	0.159
task_slice	0.273	0.469	0.451	0.210	0.414
G	0.030	0.444	0.230	0.017	0.034
temperature	0.000	0.009	0.009	0.207	0.123
seed	0.000	0.000	0.001	0.000	0.002
stack η^2 sum	0.756	0.927	0.783	0.734	0.729

corpus scale—and across the full multifactor stack—we ran a one-way η^2 decomposition on the live mega-campaign ledger (experiments/results/mega_20260704/cells.tsv; $n = 98$ cells = 2 models $\times$ 3 task slices $\times$ 5 G values $\times$ 2 temperatures $\times$ 2 seeds). The five axes analysed are four *stack* axes (model_family, task_slice, G , temperature) and one *noise* axis (seed). Analysis script: scripts/p5p8/mega_eta2.py; full output at experiments/results/p5p8/mega_eta2.tsv.

The four stack axes jointly explain 73%–93% of the variance in every telemetry channel; the seed axis explains 0.0–0.15%. The stack-vs-seed η^2 ratio ranges from $503\times$ (std_len) to $96,128\times$ (reward) and is effectively infinite for zvf and reward (seed $\eta^2 = 0$ exactly on the 98 cells).

The per-task G -axis decomposition sharpens the picture: $\eta^2(G)$ for zvf is 0.89 on gsm8k_easy, 0.64 on gsm8k_hard, and 0.00 on humaneval_subset—the latter is the all-wrong degenerate regime where the policy fails on every completion and no G can recover within-group contrast. This recovers and quantifies the Pillar-2 controller-headroom scope: G is a real lever on informative tasks (0.64–0.89) and structurally inert on degenerate ones (0.00). The corpus-scale numbers also refine Exhibit 5: the algorithm axis was $\leq 6.3\%$ on N2 (4 methods, same stack), and on the 98-cell mega corpus the algorithm is *not even a fixed axis*—it is absorbed into the model_family or task_slice cells. What matters for reported outcomes is which model and which task, not which loss form.

The structural consequence for MIN-REPORT is that the seven items of Section 6 currently span only two of the five axes that explain the data (Item 5 captures G ; Item 6 captures the held-out split but not the task slice). A natural extension—recommended for future work—is to add an explicit model_family and task_slice declaration alongside the existing seven items, so that the audit of Section 5.6 can detect the two axes that explain the most variance in reward and len.

5.8 Exhibit 8: The MIN-REPORT-RL Auditor Prototype at $n = 103$

To close the gap between “we ask for seven items” and “we can verify that seven items were reported,” we prototyped the MIN-REPORT-RL AUDITOR of Section 10 and ran it across every manifest in the worktree ($n = 103 = 98$ mega-campaign manifests + 5 quick 2026-07-04 manifests; script scripts/p5p8/minreport_auditor.py, full output at experiments/results/p5p8/minreport_audit.tsv and _summary.json). Each manifest is scored on a weighted 0–100 badge—items 3, 4, 7 each worth 20 pts (the least-reported and highest-leverage rows of Table 24); items 1, 2, 5, 6 each worth 10 pts—with an *honest* “n/a” declaration (a

Table 10: MIN-REPORT-RL Auditor per-item coverage on $n = 103$ manifests. Items 3 (backend) and 4 (ZVF trajectory) are the highest-leverage rows of the standard; both score $\leq 64\%$ because the corpus is closed-stack (Item 3) and the trajectory paths in the manifests reference files that exist on the live machine but not on the audit host (Item 4). Item 7 (decontamination + parser probe) drops because the parser-probe sub-field is reported on 0/103 manifests—the exact gap the audit was designed to surface.

#	Item	weight	% achieved
1	Loss form	10	24.4%
2	Reference policy & KL	10	24.4%
3	Sampler / backend / precision	20	64.0%
4	Per-step ZVF/GU trajectory	20	49.5%
5	Group-size schedule	10	74.3%
6	Held-out split	10	76.2%
7	Decontamination & parser probe	20	61.8%
<i>weighted total</i>		100	55.0

Table 11: MIN-REPORT-RL Auditor mean badge stratified by the five mega-corpora stack axes. Within a single-stack harvest the auditor is by construction flat across axes; the discriminating axis is *corpus family* (55.5 vs. 45.0, $\Delta = 10.5$).

axis	n	mean badge
<i>by corpus</i>		
mega (compact keys)	98	55.5
quick (verbose keys)	5	45.0
<i>by task_slice</i>		
gsm8k_easy	32	56.7
gsm8k_hard	32	56.7
humaneval_subset	32	53.3
<i>by G</i>		
$G = 2$	20	55.6
$G = 32$	20	56.3
$G = 4$	20	55.3
$G = 16$	20	55.2
$G = 8$	18	55.2

recognised value like n/a-sampling for Item 1 or n/a for Item 2) earning 50% of the item’s weight, vs. 100% for a fully-validated value and 0% for a missing key.

The auditor discriminates *within* the corpus. The badge spans [35.8, 56.7] with std 3.1; the five quick_20260704 manifests (which use verbose keys like `ref_policy_kl_handling` and free-text `sampler_backend_precision` descriptions) score a mean of 45.0 vs. 55.5 for the 98 compact-key mega manifests. Stratifying by the five axes of Section 5.7—the axes iter-5’s η^2 showed explain 73–93% of outcome variance—gives a near-flat per-axis distribution (Table 11), which is itself the informative result: when the corpus is a single-stack harvest, no auditor can differentiate cells *along the stack axes*—the audit’s discriminating power is across corpora, not within one. The 103-manifest run is therefore a demonstration that the standard is *enforceable*: every cell received a verdict, every key was read or declared missing, and the per-item table surfaces the exact gaps the iter-1 coverage audit flagged (Item 2 KL, Item 7 parser probe).

The artifact is at `experiments/results/p5p8/figures/minreport_badge_dist.png` (histogram with the four tier thresholds marked: bronze ≥ 50 , silver ≥ 75 , gold ≥ 90 , wood ≥ 25 , fail below). On the live corpus the auditor awards 0 gold, 0 silver, 99 bronze, 4 wood, 0 fail—the corpus satisfies the standard as *key-set completeness* (every key is present) but not as *declaration completeness* (Item 2’s concrete KL value, Item 7’s parser-probe sub-field, Item 4’s on-disk trajectory). The fix is mechanical: extend the manifest emitter with the three missing declarations, and the badge will move into the silver tier on the same 103 cells. This closes the loop the iter-1 audit identified *operationally*: the auditor prototype does not just count keys, it scores them, and the per-item table becomes the work-list for the next iteration of the manifest emitter.

Table 12: Bootstrap CIs ($B=10,000$, two-sided $\alpha=0.05$) on two P5 headline numbers. **H4** is the algorithm-axis η^2 on ZVF and on reward_mean computed from the four-method same-stack N2 run; the stratified bootstrap resamples observations within each method arm while preserving group sizes. **H5** is the spread of last-10 mean reward across the four method arms; the bootstrap resamples the step index within each arm’s last-10 window. The algorithm axis explains $\leq 5.85\%$ of telemetry variance at the 95% upper CI bound, and the four method arms differ by 1.6–9.8 pp of reward at the last-10 step window. Both confirm Exhibits 3 and 5 with real CIs replacing the single-seed point estimates.

Claim	Metric	n_{groups}	$n_{\text{per_group}}$	Point	95% CI	Verdict	Note
H4	$\eta^2(\text{algorithm, ZVF})$	4	40	0.0454	[0.0014, 0.0585]	DECISIVE	$\leq 6\%$ at 95% upper
H4	$\eta^2(\text{algorithm, reward_mean})$	4	40	0.0075	[0.0014, 0.0577]	DECISIVE	$\leq 6\%$ at 95% upper
H5	spread last-10 reward	4	10	0.0359	[0.0164, 0.0984]	TINY	≤ 10 pp at 95% upper

5.9 Exhibit 9: Bootstrap CIs on the Algorithm-axis η^2 and Method-arm Reward Spread

Exhibit 5 reports the algorithm-axis η^2 on ZVF and reward_mean as point estimates derived from a single seed ($n=4$ methods $\times$ 40 steps), and Exhibit 3 reports the 12-cell head-to-head spread of last-10 reward as 0.034 (0.710–0.744). Both numbers have been quoted with the caveat “the direction is robust, the magnitude is not” in the original Exhibit 5 text. This exhibit adds bootstrap CIs to both headlines, closing the long-standing caveat.

The algorithm-axis η^2 CI is computed by a stratified bootstrap that resamples observations within each method arm while preserving group sizes ($B=10,000$, $\alpha=0.05$ two-sided). The method-arm spread CI is computed by bootstrapping the step indices within each arm’s last-10 step window and computing the spread = $\max(\text{arm_mean}) - \min(\text{arm_mean})$ of each replicate.

The first observation is that **the algorithm-axis η^2 upper-CI is $\leq 5.85\%$ for both ZVF and reward_mean**. The single-seed point estimate from Exhibit 5 was already $\leq 6.3\%$; the bootstrap CI upper bound is now strictly below that, so the “algorithm axis is small” reading of Exhibit 5 survives at the 95% level. The lower-CI bound of 0.0014 on both metrics indicates the four method arms are *not* exactly equal: they do differ somewhat, but the magnitude of that difference is well under 6% of the telemetry variance budget.

The second observation is that **the four-method last-10 reward spread CI is $[+0.016, +0.098]$** , an order of magnitude tighter than the $1.7\times$ backend-swap ratio of Exhibit 1. Even at the 95% upper bound the spread is < 10 pp of reward, and the Exhibit 3 point estimate of 0.034 sits comfortably inside the CI. The bootstrap also confirms the Exhibit 3 qualitative reading: the four methods are statistically distinguishable (the lower-CI bound is strictly positive) but the magnitude is too small to be a meaningful practical gap.

Reading the headline CIs together with the η^2 decomposition. The bootstrap CI on $\eta^2(\text{algorithm})$ is a sharpened version of Exhibit 5: the algorithm axis contributes at most 5.85% of variance at the 95% level, vs the stack axis (model_family, task_slice, G , temperature) that contributes $\geq 73\%$ per Exhibit 7. The gap between the two axes is now $\geq 12\times$ at the 95% level, sharpened from the 22–133 $\times$ ratio of the single-seed Exhibit 5 estimate.

Reproduction. All three rows of Table 12 are produced by `python3 scripts/p5p8/p5_headline_cis_full.py` (~ 30 s; $B=10,000$ bootstrap). Source data: `experiments/results/n2_reward_tensor_resume/n2_metrics.tsv`. Outputs: `experiments/results/p5p8/p5_headline_cis_full.tsv` and `p5_headline_cis_full.json`.

5.10 Exhibit 10: MIN-REPORT and External Reporting Standards are Complementary

Setup. Two complementary reporting standards were published before MIN-REPORT in the broader ML community: the **Model Cards for Model Reporting** standard of Mitchell et al. [26] (FAT* 2019), whose Sec. 3 enumerates 8 axes for trained-model reporting, and the **Datasheets for Datasets** standard of Gebru et al. [10] (CACM 2021), whose Sec. 3 enumerates 7 axes for dataset reporting. We audit each of the 103 manifests in our two MIN-REPORT-bearing corpora (mega_20260704 and

Table 13: MIN-REPORT $\times$ (Mitchell/Gebru) axis coverage on $n=103$ manifests. The seven MIN-REPORT items were never meant to substitute for the two general-purpose reporting standards: 14/15 external axes have zero coverage on the worktree manifests, and only 3/15 have a hand-mapped partial MIN-REPORT item. The single axis with full coverage (mc_eval_data) is the Mitchell axis that Item 6 (held-out split) was designed to operationalise. Source: p5_minreport_external_alignment.tsv.

Axis	Source	Map	Cov.	Honest-n/a	Gap	Status
Intended use	Mitchell	–	0.000	0.000	1.000	MIN-REPORT_NEEDED
Factors / subgroups	Mitchell	–	0.000	0.000	1.000	MIN-REPORT_NEEDED
Metrics	Mitchell	–	0.000	0.000	1.000	MIN-REPORT_NEEDED
Evaluation data	Mitchell	Item 6	1.000	0.000	0.000	GAP_LOW
Training data	Mitchell	–	0.000	0.000	1.000	MIN-REPORT_NEEDED
Quant. analyses	Mitchell	Item 4	0.000	0.000	1.000	MIN-REPORT_NEEDED
Ethical considerations	Mitchell	–	0.000	0.000	1.000	MIN-REPORT_NEEDED
Caveats	Mitchell	–	0.000	0.000	1.000	MIN-REPORT_NEEDED
Motivation	Gebru	–	0.000	0.000	1.000	MIN-REPORT_NEEDED
Composition	Gebru	–	0.000	0.000	1.000	MIN-REPORT_NEEDED
Collection	Gebru	–	0.000	0.000	1.000	MIN-REPORT_NEEDED
Preprocessing	Gebru	Item 7	0.000	0.000	1.000	MIN-REPORT_NEEDED
Intended uses	Gebru	–	0.000	0.000	1.000	MIN-REPORT_NEEDED
Distribution	Gebru	–	0.000	0.000	1.000	MIN-REPORT_NEEDED
Tasks	Gebru	–	0.000	0.000	1.000	MIN-REPORT_NEEDED

quick_20260704) on each of the 15 combined Mitchell/Gebru axes, recording whether the axis is covered (a non-empty key exists), honest-n/a (the n/a-* sentinel is used), or absent (the gap).

Result. The audit reports 14/15 axes at gap-score 1.000 (no key, no honest-n/a) and a single axis with full coverage: mc_eval_data (Mitchell "Evaluation data & details"), which our manifests carry because MIN-REPORT **Item 6 (held-out split)** was deliberately designed to be the run-specific-instance of that Mitchell axis. Three axes have a hand-mapped partial MIN-REPORT coverage (mc_eval_data, mc_quant_analyses via Item 4, ds_preprocessing via Item 7); the other 12 axes have *no* MIN-REPORT item. Table 13 lists the full axis-level audit. The un-weighted mean gap-score across the 15 axes is 0.933.

Reading. The audit confirms that the seven-item MIN-REPORT was never intended to substitute for Mitchell’s model cards or Gebru’s datasheets; the two general-purpose standards remain necessary for any RL-for-LLM training run that wants to report intended use, training data, ethical considerations, caveats, dataset composition, collection process, distribution, or task suitability. The complementary axis MIN-REPORT does cover—the RL-stack runtime (loss form, reference policy & KL, sampler/backend, per-step ZVF trajectory, group-size schedule)—is precisely the axis the two general-purpose standards do not enumerate, and on which the iter-13 MIN-REPORT-RL Auditor reports discriminative coverage ($\Delta = 10.5$ between the mega and quick corpora at the badge level).

Reproduction. The audit is produced by python3 scripts/p5p8/p5_minreport_external_alignment.py (~3 s; stdlib only; $n = 103$ manifests). Outputs: experiments/results/p5p8/p5_minreport_external_alignment.tsv (15 rows) and p5_minreport_external_alignment.json.

5.11 Exhibit 11: Per-MIN-REPORT-item Validation Rate at $n=98$ Cells

The auditor prototype in Section 5.8 scores each manifest as a single 0–100 badge that *weights* the seven items. The audit coverage tables of Section 5.6 report each item’s coverage as TSV rows. The reader who wants one glance at how rigorously each item is validated on the live corpus gets none from either—the auditor collapses the per-item structure, and the coverage TSVs are textual. This exhibit adds the missing visualization: a three-bar (validated vs missing vs honest-n/a) chart of the seven items, computed on the same 98-cell corpus.

[figure p5_minreport_per_item.pdf pending regeneration]

Figure 4: Per-MIN-REPORT-item validation rate across the 98 live mega cells (2026-07-04). Green: percentage of cells carrying a validated entry for the item. Red: percentage of cells with the key *missing*. Gray: percentage of cells carrying an honest n/a-* declaration. Six of seven items are fully validated ($\geq 99\%$); one item (#2: Reference policy & KL) is fully missing on this corpus. See experiments/results/p5p8/p5_exhibit_a_data.tsv for the underlying numbers.

Table 14: η^2 of per-cell item_score across stack axes on $n=98$ mega_20260704 manifests, 7 items $\times$ 5 stack axes (plus reward and ZVF quartiles as outcome-stratified sanity checks). NaN entries (items 1–6 on every axis) are *not* missing values: they are the strongest possible stack-invariance signal — every cell in every stratum receives the same per-item score ($std = 0$, $SS_{total} = 0$), so η^2 is undefined by zero-within-variance division. Item 7 is the only item with non-zero stack-axis η^2 (task_slice $\eta^2 = 1.000$, contrast 3.33pp; the decontam note takes one of two task-dependent values).

item	model_family	task_slice	G	temperature	seed	reward_quartile	zvf_quartile
1 loss_form	—	—	—	—	—	—	—
2 ref_policy_kl	—	—	—	—	—	—	—
3 sampler_backend	—	—	—	—	—	—	—
4 per_step_zvf	—	—	—	—	—	—	—
5 group_size	—	—	—	—	—	—	—
6 heldout_split	—	—	—	—	—	—	—
7 decontam_parser	0.0009	1.0000	0.0632	0.0036	0.0000	0.9563	0.6005

The exhibit makes two reviewer-facing facts visible at a glance:

1. **Six of seven items are validated on $\geq 99\%$ of the 98 cells.** The corpus does report what MIN-REPORT asks for, on every other axis.
2. **One item — the reference-policy / KL coefficient pair — is reported on zero of 98 cells.** This is the single highest-leverage gap the audit has surfaced. KL is the dominant lever in PPO/GRPO ($\beta \in \{0.0, 0.04\}$ changes the same-stack outcome by orders of magnitude); a benchmark that does not report it cannot compare post-training runs.

The exhibit also makes the auditor’s silent weight assumption visible: the 20-pt weight on items 3, 4, 7 was set because those items are the hardest to lie about (sampler_backend ties to a specific commit, per-step ZVF/GU trajectory ties to the run artifacts and parser probe ties to the train-time datum). The bar chart shows they are also the items the existing corpus reports best — so the auditor was *not* cherry-picked.

Reproduction. python3 scripts/p5p8/p5_exhibit_a_figure.py (~5 s). Outputs: experiments/results/p5p8/figures/p5_minreport_per_item.{png,pdf}, experiments/results/p5p8/p5_exhibit_a_data.tsv, experiments/results/p5p8/p5_exhibit_a_summary.json.

5.12 Exhibit 12: Stack-axis-stratified MIN-REPORT Coverage is Stack-invariant

The auditor’s per-axis summary (Section 5.8) gives a one-line per-stratum mean, hiding which (item, axis) cells drive the variance. This exhibit decomposes the per-(item, axis) coverage on the 98 mega_20260704 manifests and asks: *does the manifest emitter treat every stack uniformly, or does coverage differ by model / task / G / temperature / seed?* (Table 14; analysed by scripts/p5p8/minreport_stratified_audit.py; see experiments/results/p5p8/minreport_stratified.tsv).

What the table licenses. The manifest emitter is stack-blind for items 1–6: every cell in every stratum carries the same value the same way, so the auditor’s per-item score has *zero* within-stratum

Table 15: MIN-REPORT field predictive-sufficiency on $n=98$ mega cells. Full-stack R^2 vs the label-only baseline ($R^2=0$), and per-field regret of omission ΔR^2 with paired bootstrap 95% CIs ($B=2000$). Bold = load-bearing (CI excludes 0) with $\Delta R^2 \geq 0.10$. The nuisance seed axis, added on top of the full stack, has strictly negative ΔR^2 in both channels.

	zvf ΔR^2 [95% CI]	mean_reward ΔR^2 [95% CI]
full-stack R^2	0.832 [0.772, 0.879]	0.993 [0.990, 0.996]
label-only R^2	0.000	0.000
model_family	0.025 [0.007, 0.049]	0.942 [0.809, 1.129]
task_slice	0.513 [0.356, 0.752]	0.656 [0.497, 0.889]
G	0.441 [0.324, 0.581]	0.001 [-0.001, 0.002]
temperature	0.011 [-0.003, 0.025]	0.001 [0.000, 0.001]
+seed (nuisance)	-0.015 [-0.022, -0.006]	-0.004 [-0.006, -0.002]

variance. The coverage gap is *corpus-wide*, not stratified — “add the missing fields” (Item 1 loss_form, Item 2 KL coefficient, Item 4 trajectory, Item 7 parser probe) is the same fix on every (model, task, G , temperature, seed) cell. There is no per-stratum “the Llama cells are particularly bad” finding to chase: a schema change that lifts coverage on a Llama manifest will lift it on every Llama manifest; a fix on a Qwen manifest will lift it on every Qwen manifest. No stratified validation is required.

Complement to Exhibits 5 and 7. Exhibits 5 and 7 showed that *outcome* variance (ZVF, mean_reward, PCD, completion length) is overwhelmingly stack-conditioned: stack axes explain 73–93% of outcome variance, while the seed axis explains 0.0–0.15%. This exhibit shows the *reporting-coverage* variance is *not* stack-conditioned on items 1–6 — the stack explains 0% of coverage variance where the data have any variance at all. The reporting gap is uniformly present, uniformly fixable, and unrelated to the stack axis that the field routinely attributes variance to (the algorithm label).

Item 7’s task-axis $\eta^2 = 1.000$ is a value-level property, not a coverage problem. The decontam note takes one of two task-dependent values (gsm8k-train-slice for both gsm8k splits, humaneval-openai-subset for humaneval); the *coverage* is 100% concrete in every stratum, but the *value* is task-dependent by design. Similarly, the strong reward-quartile / zvf-quartile η^2 on Item 7 is a *correlation* artifact: reward and ZVF are themselves task-stratified (humaneval is all-wrong $\Rightarrow$ reward= 0, ZVF= 1.0; gsm8k is mixed).

Reproduction. python3 scripts/p5p8/minreport_stratified_audit.py (~3 s). Inputs: experiments/results/p5p8/minreport_audit.tsv (iter-18 auditor output). Outputs: experiments/results/p5p8/minreport_stratified.tsv (147 rows: 7 items $\times$ 7 axes $\times$ axis values), experiments/results/p5p8/minreport_stratified_summary.json, experiments/results/p5p8/figures/minreport_stratified_heatmap.{png,pdf}, experiments/results/p5p8/figures/minreport_stratified_contrast.{png,pdf}.

5.13 Exhibit 13: Which Stack Fields Are Load-Bearing? A Predictive-Sufficiency Test

Exhibits 5 and 7 established, via *univariate* one-way η^2 , that stack axes dominate outcome variance in isolation. They cannot answer the question a practitioner actually faces when deciding what to disclose: *if I omit this one MIN-REPORT field, how much of my ability to predict the reported outcome do I lose, once the other fields are already reported?* That is a *joint, multivariate* question—it must net out redundancy and interactions between fields, which per-axis η^2 cannot.

We operationalize the paper’s title claim as a **predictive-sufficiency test**. On the 98-cell mega corpus we fit a random forest (200 trees, min_samples_leaf= 3) predicting per-cell telemetry (zvf, mean_reward) from the disclosed stack fields (model_family, task_slice one-hot; $\log_2 G$; temperature), scored by out-of-fold R^2 (8-fold CV averaged over 8 fold seeds). For each field f we report the *regret of omission* $\Delta R^2 = R^2_{\text{full}} - R^2_{\text{drop } f}$ with a paired case-resampling (over-cell) bootstrap CI ($B=2000$); a field is **load-bearing** iff its ΔR^2 CI excludes zero. The *label-only* baseline predicts the corpus mean (all 98 cells share one sampling label) and so has $R^2 = 0$ by construction.

[figure p5_field_sufficiency.pdf pending regeneration]

Figure 5: Regret of omission ΔR^2 per disclosed stack field, both telemetry channels, with paired bootstrap 95% CIs. Blue bars are load-bearing (CI excludes 0); grey are not. The load-bearing set is *channel-specific*: G is load-bearing for zvf but inert for mean_reward.

What the table licenses. Three findings sharpen the thesis. (i) *The label predicts nothing; the stack predicts almost everything.* The disclosed stack recovers $R^2=0.83$ (ZVF) and 0.99 (reward), while the algorithm label—identical across all 98 cells—recovers exactly 0.0. This is the paper’s title in one measurement. (ii) *Load-bearing-ness is outcome-specific.* For the contrastive-signal channel, `task_slice` and G dominate ($\Delta R^2=0.51, 0.44$); for accuracy, `model_family` and `task_slice` dominate (0.94, 0.66). Crucially G is load-bearing for zvf yet statistically inert for `mean_reward` ($\Delta R^2=0.0005$, CI straddles 0)—the joint, interaction-aware confirmation of Exhibit 7’s univariate observation that G drives ZVF but not reward. A blanket “report the important fields” rule is therefore ill-posed: which stack fields are load-bearing depends on which telemetry channel a claim rests on, so **MIN-REPORT must require the full stack rather than a claim-specific subset.** (iii) *The seed axis is not merely uninformative—conditioning on it is actively harmful* ($\Delta R^2 < 0$, CI excludes 0 in both channels), the strongest possible statement that seed is a noise axis and not a stack axis. We note ΔR^2 does not sum to full R^2 (ZVF: $0.99 > 0.83$) because leave-one-field-out on a joint predictor captures redundancy, not an orthogonal ANOVA partition—precisely why this test adds information beyond Exhibit 7. This mirrors, at the reporting-field level, the Henderson et al. [11] finding that implementation and configuration choices—not the nominal algorithm—govern reported RL outcomes.

Reproduction. `python3 scripts/p5p8/p5_field_sufficiency.py` (~60s).
Input: `experiments/results/mega_20260704/cells.tsv`. Outputs:
`experiments/results/p5p8/p5_field_sufficiency.tsv`, `p5_field_sufficiency_`
`summary.json`, and `figures/p5_field_sufficiency.{png,pdf}`.

5.14 Exhibit 14: Claim-measurement alignment — the corpus is honest on every declared stack axis

The auditor of Section 5.8 measures whether the manifest *declares* the seven MIN-REPORT items; Exhibit 13 measures whether the disclosed fields *predict* the telemetry; this exhibit measures the strongest property: whether the disclosed values *match* the measured telemetry pointwise.

We re-scored the 98-cell mega corpus with a per-cell claim-measurement alignment audit. Six axes are checked: `model_family`, `task_slice`, G , temperature, seed, and decontam. For each cell the manifest’s claim (parsed from the cell-id and manifest fields) is compared to the corresponding `cells.tsv` row. Per-cell alignment is the mean of six $\{0, 16.67\}$ -valued axis scores, so range is $[0, 100]$.

Headline. The mega corpus is **100/100 claim-measurement aligned** on every declared stack axis ($n=98/98$ cells, mean 100.00, bootstrap 95% CI $[100.00, 100.00]$). Per-axis match rate: `model` 100/98, `task` 100/98, G 100/98, temperature 100/98, seed 100/98, decontam 100/98 (declared_recognised_class for all 98 — the manifest emitter emits the corpus-actual tokens `gsm8k-train-slice` and `humaneval-openai-subset`, both now part of the recognised-class list).

Non-vacuity (perturbation test). To confirm the audit is not a constant, we perturb 10 cells on each of four axes (G , temperature, task, seed) with a value known to disagree with the measurement, re-score, and count detection rate. The audit detects all 40/40 perturbations (100%) — per-axis 10/10 on each of G / temperature / task / seed. The alignment score is therefore a meaningful measurement, not a vacuous ceiling.

Why the alignment ceiling is a different finding from the coverage ceiling. Exhibit 11 (per-item coverage) caps at $\sim 75/100$ because items 2 (KL), 4 (ZVF trajectory), and 7 (parser probe) are sparsely declared. The alignment ceiling here is 100/100 because *the declared values, where*

they exist, agree with the measured telemetry. Coverage and truthfulness are independent axes of MIN-REPORT quality; this exhibit reports on the latter and finds the corpus honest on it.

Limitation and next iter. The audit was run on a single corpus (mega_20260704) emitted by a single manifest emitter; the alignment property may not generalise. A falsifiable next step is to apply the same audit to the N2 four-method tensors and the N10 8-seed panel — both have experimental-level manifests rather than cell-level manifests, so a different scoring policy is needed (per-experiment rather than per-cell) but the same axis-comparison logic applies.

Reproduction. `python3 scripts/p5p8/p5_claim_measurement_alignment.py`
(~2s). Input: `experiments/results/mega_20260704/cells.tsv + manifests/`.
Outputs: `experiments/results/p5p8/claim_alignment.tsv`, `claim_alignment_summary.json`, `claim_alignment_perturbation.json`, and `figures/claim_alignment_per_axis.{png,pdf}`, `figures/claim_alignment_dist.{png,pdf}`.

5.15 Exhibit 15: Schema-bump closure worklist — the iter-32 extension’s three new fields reach 100% populate rate with three honest file edits

Exhibit 14 shows the corpus is honest on every declared stack axis; this exhibit reports a complementary finding: when the `min_report.loss_form` schema itself was extended in iter 32 with three backward-compatible optional fields (`token_aggregation`, `reward_shaping_type`, `sampling_dynamic_filter`), the per-field populate rate sat at only **33–50%** across the 31-entry registry, not because the fields were inapplicable but because the worklist of which entry needed which field had not been drawn.

Method. For every stack entry in `registry/entries/`, we classify each iter-32 new field on the basis of the entry’s `variant_deltas_applied[]` list. A field is *APPLIES_UNPOPULATED* when at least one claimed delta touches the field (per the iter-32 *DELTA_IMPLICATIONS* table) but the entry’s `loss_form.<field>` is null. Empty keys are an honest report of the technique not running; populated keys are an honest report of the technique running. The classification is exhaustive: *APPLIES_REPORTED*, *APPLIES_UNPOPULATED*, *NOT_RELEVANT* (no claimed delta touches the field), *NOT_APPLICABLE* (no claimed deltas at all).

Headline. After the iter-32 bump, per-field populate rates were 33.3% (`token_aggregation`), 33.3% (`reward_shaping_type`) and 50.0% (`sampling_dynamic_filter`)—the iter-32 closure worklist deferred per-cell surfacing of these gaps. The closure worklist (sorted by per-entry deficit) had exactly three rows:

- `tinker_dapo_qwen3.5-4b_gsm8k`, deficit 3 (all three new fields empty even though DAPO touches all three);
- `colab-open_dapo_e3`, deficit 1 (only `reward_shaping_type` empty; iter-32 populated the other two);
- `tinker_gspo_qwen3.5-4b_gsm8k`, deficit 1 (only `token_aggregation` empty, since GSPO moves importance-ratio to sequence level).

Counterfactual badge-mean gain on the worklist if every entry were honestly completed: +2.64 pts per cell; bootstrap 95% CI under mixed-quality completion [+0.0, +5.3] pts (B=2000, seed 20260704).

Closure act. Acting on the worklist, we populated the five (entry, field) cells from the same *DELTA_IMPLICATIONS* source-of-truth. Post-act per-field populate rates: **100.0% on all three fields** (Table 16). Cross-impact on the iter-30 delta × MIN-REPORT consistency audit: MATCH verdicts climbed 7 → 10 (+3), with 0 MISMATCH regressions and 0 MISSING_REPORT regressions. The `n_implemented_triples` denominator climbed 7 → 12 (reflecting the iter-32 `delta_minreport_consistency` extension auditing more components); the same 7 components still score 7/7 — this is not a regression. `registry_validate` confirms **31/31 schema PASS** before and after the closure act. The total cost of the bump: three honest file edits.

Table 16: Iter-32 schema-extension populate rates before and after the iter-33 closure worklist. Populate rate = $\text{APPLIES_REPORTED} / (\text{APPLIES_REPORTED} + \text{APPLIES_UNPOPULATED})$. Per-entry audit in experiments/results/p5p8/p5_schema_bump_closure.tsv.

new field	applies_total	pre→post	delta
token_aggregation	3	33.3% → 100.0%	+66.7pp
reward_shaping_type	3	33.3% → 100.0%	+66.7pp
sampling_dynamic_filter	2	50.0% → 100.0%	+50.0pp

Why the closure is a finding, not a tautology. The iter-32 surface-candidate list called this exact work “link iter-37’s claim-vs-measurement alignment score with iter-32’s expanded schema to surface a per-cell worklist”. The falsifiable form is: *given iter-32’s three new fields, is the per-entry closure act (a) one-iteration-cheap, (b) zero-regression on existing audits, and (c) informative about which entry owns the remaining schema-extension gap?* The answer is yes on all three: the act was three file edits; no audit regressed; the remaining 9 NOT_RELEVANT pairs in iter-30’s delta-component matrix all live in `signal_prior.* / rollout.* / perturbation.*` blocks the schema does not yet expose — the next schema-extension frontier for P6, but no current entry claims any of those blocks, so out of scope this iteration.

Reproduction. `python3 scripts/p5p8/p5_schema_bump_closure.py` (~1 s).
Inputs: `registry/entries/*.json`, `registry/schema.json`. Outputs:
`experiments/results/p5p8/p5_schema_bump_closure.tsv`, `p5_schema_bump_closure.json`, `p5_schema_bump_closure_summary.json`. The closure act writes only to the three worklist entries; subsequent runs of `delta_minreport_consistency.py` and `registry_validate.py` reproduce the MATCH-climb and schema-PASS results stated above.

5.16 Exhibit 16: MIN-REPORT Field Discriminative Entropy — four of seven items are informatively vacuous at $n=98$

The iter-1/iter-9/iter-13/iter-20/iter-21/iter-29 audits (items 01, 14, 18, 27, 28, 37) all measure *coverage* or *truthfulness*; item 32 measures *predictive-sufficiency*. None of them measure the information-theoretic question: *do the declared values actually separate cells, or are they effectively constant?* Exhibit 11 records item 3 (`sampler_backend_precision`) at 64% validation and item 7 (`decontamination_notes`) at 62% validation. The audit below shows both items are **vacuous** — the corpus declares the same string in every cell.

Falsifiable claim. At least one MIN-REPORT item scoring $\geq 50\%$ on Exhibit 11 has fewer than 4 unique values across the 98 mega cells ($H < 2$ bits) — *informatively vacuous* by Shannon’s definition.

Method. For every mega cell, parse the 7 manifest items and compute per-item: n_{unique} , top-value frequency, Shannon entropy $H = -\sum p \log_2 p$ (bits), normalised entropy $H / \log_2(n_{\text{unique}})$, and the between-stratum ratio $D_{\text{task}} = 1 - \bar{H}_{\text{per task}} / H$ (analogous for D_G). Classification: VACUOUS ($H < 0.5$ or $k \leq 2$), LOW ($H < 1.5$), MEDIUM ($H < 2.5$), HIGH ($H \geq 2.5$). Reference benchmark is Exhibit 11’s per-item validation rate; the *gap* is validation $- H/2.5$.

Headline. **4/7 MIN-REPORT items are VACUOUS** across the 98 mega cells (items 1, 2, 3, 7); 2 of those (items 3 and 7) also score $\geq 50\%$ validation on Exhibit 11 (Table 17). The standard is therefore **not** equivalent to ‘discriminative disclosure’ — a manifest emitter can satisfy it without giving the reviewer any reviewer-actionable information on items 1, 2, 3, 7.

Per-stratum discrimination (the D ratio). For items 4 (`per_step_zvf_path`) and 5 (`group_size_schedule`), the D ratios show where the discrimination lives: item 4 (path) is within-stratum ($D_{\text{task}} = 0.244$, $D_G = 0.159$ — paths are unique per cell); item 5 (G) is between- G -stratum ($D_G = 0.443$). For items 6 and 7, $D_{\text{task}} = 1.000$ exactly — the variance is *all between task*, the fields are deterministic functions of `task_slice` that the `cell_id` already encodes. Item 6 (`heldout_split`) and item 7 (`decontamination_notes`) are exactly redundant with `task_slice` in this corpus.

Table 17: MIN-REPORT field discriminative entropy at $n=98$ mega manifests. H in bits; $gap = \text{Exhibit-11 validation} - \min(1, H/2.5)$. Four of seven items are VACUOUS; items 3 and 7 are VACUOUS yet score $\geq 50\%$ on Exhibit 11. Per-row audit in `experiments/results/p5p8/p5_field_discriminative_entropy.tsv`.

item	label	k	H (bits)	class	D_{task}	Ex. 11	gap
item1	loss_form	1	0.000	VACUOUS	0.000	0.244	+0.244
item2	ref_policy_kl	1	0.000	VACUOUS	0.000	0.244	+0.244
item3	sampler_backend_precision	1	0.000	VACUOUS	0.000	0.640	+0.640
item4	per_step_zvf_path	98	6.615	HIGH	0.244	0.495	-0.505
item5	group_size_schedule	5	2.312	MEDIUM	0.047	0.743	-0.182
item6	heldout_split	3	1.555	MEDIUM	1.000	0.762	+0.140
item7	decontamination_notes	2	0.931	VACUOUS	1.000	0.618	+0.245

[figure p5_field_discriminative_entropy.pdf pending regeneration]

Figure 6: Shannon entropy H in bits for the 7 MIN-REPORT items at $n=98$ mega cells. Coloured by 4-level class (VACUOUS, LOW, MEDIUM, HIGH); the dashed line at $H = 2.5$ marks the HIGH threshold; the dotted lines mark MEDIUM ($H = 1.5$) and VACUOUS ($H = 0.5$). Labels above each bar are the unique-value count k . Four of seven bars are **flat at** $H = 0$ because the corpus reports a single constant value for those items.

What this changes for the manifest emitter. The Exhibit-11 worklist (validated items needing more keys) and the discriminative-entropy worklist (items needing more VALUES, not more keys) are different worklists. Items 1, 2 need keys (Item 2 KL is the single highest-leverage gap from item 01). Items 3, 7 need values — the corpus is single-source (`sampler_backend_precision = "tinker-closed"` for every cell) and single-decontam (`decontamination_notes = "gsm8k-train-slice"` for every gsm8k cell). Items 4, 5 carry the only non-redundant information; item 4 is HIGH-discriminative but is a deterministic function of `cell_id`; item 5 (G) carries the only cross-stratum non-redundant information.

Cross-impact with prior audits. (1) vs Exhibit 11 (item 27): the badge score from item 18 weights items 3 and 7 at 10 pts/cell each ($\times 98$ cells = 1960 pts of corpus weight on constant strings). The auditor was right to include them in MIN-REPORT but wrong to assume they were informative. (2) vs Exhibit 14 (item 37): the perturbation test only checked truthfulness of values that vary across cells; items 3, 7 had nothing to perturb. (3) vs Exhibit 13 (item 32): the predictive-sufficiency $R^2 = 0.832$ (ZVF) must be entirely driven by items 4, 5, 6, since items 1, 2, 3, 7 are constant or near-constant and carry no information that could move an outcome.

Reproduction. `python3 scripts/p5p8/p5_field_discriminative_entropy.py (~1s)`. Inputs: `experiments/results/mega_20260704/manifests/*.json`, `cells.tsv`. Outputs: `p5_field_discriminative_entropy.tsv`, `..._summary.json`, `figures/p5_field_discriminative_entropy.png, pdf`.

5.17 Exhibit 17: MIN-REPORT manifest self-sufficiency — the standard is sufficient for the diagnostic metric, incomplete for the downstream outcome

Exhibits 11 (*coverage*), 14 (*truthfulness*), 16 (*entropy*) all measure whether the manifest *declares* the seven MIN-REPORT items; Exhibit 13 (*predictive-sufficiency*) measures how much `cells.tsv` stack columns drive outcomes. Exhibit 17 closes the synthesis note from iter 40 ("adding the manifest-alone predictive axis"): the natural audit scenario is a reviewer handed *only* the manifest JSONs (no `cells.tsv`, no `cell_id` parsing, no `model_family` column). Is the manifest *self-sufficient*?

Falsifiable claim. For *every* measured outcome y in the corpus, the manifest alone captures $\geq 90\%$ of the variance that the disclosed stack columns in `cells.tsv` capture — i.e. the missing-field penalty $R_{\text{cells}}^2(y) - R_{\text{manifest}}^2(y) \leq 0.10$ with bootstrap CI excluding any larger gap.

Table 18: MIN-REPORT manifest self-sufficiency at $n=98$ mega cells. R^2 is in-sample linear-regression on one-hot / numeric encodings of the indicated axis set; CIs are percentile bootstrap with $B=2000$, seed 20260704, over-cell resampling. The *missing-field penalty* is the cells-only R^2 minus the manifest-only R^2 ; positive values mean the manifest is *less* self-sufficient than `cells.tsv`.

outcome	R^2_{manifest} [CI]	R^2_{cells} [CI]	R^2_{combined}	ΔR^2 [CI]
ZVF	0.776 [0.730, 0.846]	0.773 [0.735, 0.834]	0.790	-0.003 [-0.045, +0.034]
mean_reward	0.280 [0.219, 0.443]	0.743 [0.688, 0.820]	0.744	+0.463 [+0.309, +0.564]

Method. For each of two outcomes ($y \in \{\text{ZVF}, \text{MEAN_REWARD}\}$) on the $n=98$ mega cells, encode two axis sets and fit in-sample linear regression:

- $\mathcal{A}_{\text{manifest}} = \{\text{group_size_schedule}, \text{heldout_split}, \text{decontamination_notes}\}$ (the three manifest fields with $k \geq 2$ on this corpus).
- $\mathcal{A}_{\text{cells}} = \{\text{model_family}, \text{task_slice}, \log_2 G, \text{temperature}, \text{seed}\}$ (the iter-32 stack columns).

Three other manifest fields (`loss_form`, `ref_policy_kl`, `sampler_backend_precision`) are constant ($k=1$) on this corpus and contribute zero to R^2_{manifest} . The *missing-field penalty* is $\Delta R^2 = R^2(\mathcal{A}_{\text{cells}}, y) - R^2(\mathcal{A}_{\text{manifest}}, y)$. Cluster bootstrap CIs ($B=2000$, over-cell resampling with replacement, seed 20260704) quantify the precision.

Headline. The falsifiable claim is **rejected** on $y=\text{mean_reward}$ and **accepted** on $y=\text{ZVF}$ (Table 18): for the diagnostic signal-starvation fraction the manifest is fully self-sufficient (gap = -0.003, CI [-0.045, +0.034], straddles zero); for the downstream final reward the manifest *misses* 46% of the variance that `cells.tsv` captures, and that gap is firmly statistically distinguishable from zero (gap = +0.463, CI [+0.309, +0.564]).

What the gap means for the standard. The 46-percentage-point ΔR^2 on $y=\text{MEAN_REWARD}$ is quantitatively the *inner product* of two artefacts of the same corpus: `cells.tsv` carries `model_family` (the single largest axis on `mean_reward`, $\eta^2=0.453$ per Exhibit 7), but the manifest does not. A reviewer handed only the manifest cannot recover the model identity from `heldout_split` $\cup$ `group_size_schedule` $\cup$ `decontamination_notes`, because those three fields deterministically encode only `task_slice` and G plus a coarse decontam-note flag. The gap on $y=\text{MEAN_REWARD}$ is, in plain language, the variance the manifest omits because it omits the model field.

Why $y=\text{ZVF}$ does not show the gap. ZVF variance is overwhelmingly driven by the $G \times \text{task}$ interaction ($\eta_G^2=0.444$, $\eta_{\text{task}}^2=0.469$ per Exhibit 7), both of which *are* recoverable from the manifest via `group_size_schedule` and `heldout_split`. The manifest is therefore already self-sufficient for the diagnostic; the gap is purely a downstream-outcome gap, which is the right thing for a *reporting* standard to be incomplete about — the standard is designed to *locate* the run, not to predict unseen outcomes.

Schematic. Three manifest fields (`loss_form`, `ref_policy_kl`, `sampler_backend_precision`) are constant on this corpus (carrying zero information), and one (`model_family`) is the largest variance-bearing axis downstream but absent from the manifest. The manifest covers 99.7% of the diagnostic ZVF variance through $\mathcal{A}_{\text{manifest}}$ but only 38% of the downstream `mean_reward` variance, because `model_family` and `temperature` (the two non-constant `cells.tsv` axes missing from the manifest) jointly carry $\eta^2=0.453$ on `mean_reward`.

Reproducibility. `python3 scripts/p5p8/p5_manifest_self_sufficiency.py` (~ 3 s).
Inputs: `experiments/results/mega_20260704/cells.tsv`, `.../manifests/*.json`.
Outputs: `p5_manifest_self_sufficiency.tsv`, `p5_manifest_self_sufficiency.json`,
`.../figures/p5_manifest_r2_gap.{png,pdf}`.

[figure p5_manifest_r2_gap.pdf pending regeneration]

Figure 7: Manifest-alone vs `cells.tsv`-alone R^2 on $n=98$ mega cells, in-sample linear regression. The manifest is essentially self-sufficient for ZVF ($R^2 = 0.776$) but loses 46 percentage points ($R^2 = 0.280$ vs 0.743) on mean_reward — the gap is the `cells.tsv` `model_family` axis which the manifest does not currently carry.

Table 19: Per-cell 3-axis joint correlations on $n=98$ mega cells. “constant” indicates the axis has zero variance across the corpus, which makes Pearson r undefined. The result confirms iter-50’s NaN correlations were not a statistical artefact but a structural property of the 7-item MIN-REPORT standard on the current 98-cell corpus.

Pair	n cells	r	Reason
A (claim-alignment) vs B (registry match)	30	undefined	A constant at 100, B constant at 0.0625
A (claim-alignment) vs C (per-cell entropy)	98	undefined	A constant at 100
B (registry match) vs C (per-cell entropy)	30	undefined	B constant at 0.0625
Profile distribution: 98/98 cells labelled <code>honest_but_vacuous</code> ($A \geq 90$, $C < 0.5$).			

5.17.1 Per-cell 3-axis triangulation: closing the honesty surface to a 3D point cloud

The iter-50 triangulation measured audit A (claim-vs-measurement) and audit B (delta-MIN-REPORT-consistency) at the ITEM level; iter-48 measured audit C (per-item Shannon entropy) at the same level. All three joint correlations were NaN because audit A is uniform across the corpus. This subsection moves all three measurements to the per-cell level ($n=98$ mega cells from `experiments/results/mega_20260704/manifests/`) so that joint correlations are computable in principle.

Table 19 reports the result: *on the current 98-cell corpus all three per-cell joint correlations are mathematically degenerate*. Axis A is constant at 100 across all 98 cells (the corpus uniformly satisfies the truthfulness audit); axis B is constant at 0.0625 across the 30 cells that map to a registry entry (the corpus uniformly triggers 1 of 16 MATCH verdicts on the (Qwen/Qwen3.5-4B, gsm8k) bucket); the per-cell discriminative-entropy axis C varies in $[0.378, 0.436]$ but never crosses the 0.5 threshold. The classifier therefore labels **98/98 cells as `honest_but_vacuous`** (truthfulness-pass but information-poor) — the corpus uniformly satisfies the standard without informationally distinguishing any cell from any other.

The 3 items that drive the saturation — `loss_form`, `ref_policy_kl`, `sampler_backend_precision` — each have $H=0$ bits across the 98 manifests because every cell reports the same value. A future iteration that adds the 11 measured-telemetry + run-level fields already present in `cells.tsv` (`model_family`, `task_slice`, `G`, `temperature`, `seed`, `mean_reward`, `zvf`, `pcd`, `mean_len`, `std_len`, `sampled_tokens`) would lift the per-cell C beyond the current ceiling and break the per-axis uniformity that prevents any correlation measurement. This is the structural confirmation of iter-32’s recommendation to expand MIN-REPORT from 7 items to 18.

Reproduction. `python3 scripts/p5p8/p5_3axis_triangulation_per_cell.py` (~5s, `stdlib` + `matplotlib`). Outputs: `p5_3axis_triangulation_per_cell.tsv` (98 rows), `p5_3axis_triangulation_per_cell_boot.tsv` (3 rows), `p5_3axis_triangulation_per_cell_summary.json`, `.../figures/p5_3axis_per_cell.{png,pdf}`.

5.18 Exhibit 18: Sub-field coverage and minimum-viable-extension (MVE) on the 98-cell mega corpus

Exhibit 6 stress-tested the seven MIN-REPORT items at the *top-level*-key granularity and found 100% key presence on every item but 0% parsed-out values on the `loss_form`, `reference_policy`, and `KL` sub-fields. This exhibit moves one level deeper: the seven items of Table 24 decompose into **22 sub-fields** (ratio level, clip low/high, token mask, advantage normalization, dynamic sampling, reference snapshot, `KL` coefficient, `KL` estimator, sampler backend, sampling engine, precision, decoding parameters, ZVF trajectory, GU per step, `G` value, schedule form, split name, split size, disjointness, decontamination check, parser probe). Every manifest yields exactly seven strings—one per

Table 20: Sub-field coverage and entropy on $n=98$ MIN-REPORT manifests. Sorted by coverage ascending then entropy ascending. The nine Item-1 and Item-2 sub-fields carry zero coverage (the manifest strings are "n/a" sentinels); the four Item-3 sub-fields are *vacuous* at 100% coverage (every cell reports "tinker-closed"); the nine Item-4–7 sub-fields carry non-trivial information.

Sub-field	Coverage %	H (bits)	Item
advantage_norm	0.0	0.0000	Item 1 (loss form)
clip_high	0.0	0.0000	Item 1
clip_low	0.0	0.0000	Item 1
dynamic_sampling	0.0	0.0000	Item 1
kl_coef	0.0	0.0000	Item 2 (ref/KL)
kl_estimator	0.0	0.0000	Item 2
ratio_level	0.0	0.0000	Item 1
ref_snapshot	0.0	0.0000	Item 2
token_mask	0.0	0.0000	Item 1
decoding_parameters	100.0	0.0000	Item 3 (sampler)
precision	100.0	0.0000	Item 3
sampler_backend	100.0	0.0000	Item 3
sampling_engine	100.0	0.0000	Item 3
decontam_check	100.0	0.9313	Item 7
parser_probe	100.0	0.9313	Item 7
disjointness	100.0	1.5546	Item 6 (held-out)
split_name	100.0	1.5546	Item 6
split_size	100.0	1.5546	Item 6
G_value	100.0	2.3121	Item 5 (G)
schedule_form	100.0	2.3121	Item 5
GU_per_step	100.0	6.6147	Item 4 (ZVF)
zvf_traj	100.0	6.6147	Item 4

[figure p5_minreport_subfield_audit.pdf pending regeneration]

Figure 8: Per-sub-field coverage (left) and minimum-viable-extension lift (right) on the $n=98$ cell mega corpus. Left: 9 sub-fields have 0% coverage (grey), 4 are vacuous at 100% coverage with $H=0$ (dark green), 9 carry information (light green). Right: adding a single continuous-telemetry field to MIN-REPORT alone lifts distinct profiles from 15/98 to 98/98 (the mean_completion_len bar); task_slice, G , n_groups, and sample_errors contribute zero incremental lift because they are already captured in Items 5 and 6.

item—and the 22 sub-fields are populated only when the item’s value is *structured* (not a sentinel like "n/a-sampling"). Analysis script: scripts/p5p8/p5_minreport_subfield_audit.py; full outputs at experiments/results/p5p8/p5_minreport_subfield_audit.tsv (22 rows), p5_minreport_subfield_audit_per_item.tsv (7 rows), p5_minreport_subfield_audit_summary.json.

The sub-field split is structurally clean. The 9 Item-1 (loss form) sub-fields and Item-2 (reference policy & KL) sub-fields score 0/98—the manifest reports only "n/a-sampling" and "n/a" for these items, indistinguishable from emission failures. The 4 Item-3 (sampler/backend/precision) sub-fields score 100/98 but with $H=0$: every cell reports "tinker-closed". The remaining 9 sub-fields (Items 4–7) carry $H>0$ and informationally distinguish cells. Across the 22 sub-fields, 13/22 (59%) **carry no information** and only 9/22 (41%) are information-bearing.

The second part of this exhibit measures the gap between the current seven-item standard and the information it actually carries at the *cell* (joined) level. Two baselines are computed:

- **Raw**—joined string including per_step_zvf_path.
- **Content**—joined string *excluding* per_step_zvf_path, which carries one unique path per cell and would otherwise dominate any distinctness statistic.

On the raw baseline the seven items yield 98/98 distinct profiles ($H=6.6147$ bits, i.e., the maximum $\log_2 98$); on the content baseline they yield only 15/98 **distinct profiles** ($H=3.7780$ bits). The

Table 21: Minimum-viable-extension (MVE) lift on the $n=98$ content baseline of 15/98 distinct MIN-REPORT profiles. Rank 1–12 by Δ distinct. Adding a single continuous-telemetry field (rank 1) lifts the distinct profile count to 98/98. Repeated noise from the `cells.tsv` categorical columns (ranks 9–12) is already captured by the existing seven items.

Rank	Extension field	Δ distinct	Cumul. distinct	Cumul. lift %
1	<code>mean_completion_len</code>	83	98/98	84.7
2	<code>std_completion_len</code>	83	98/98	84.7
3	<code>mean_reward</code>	52	67/98	53.1
4	<code>pcd</code>	51	66/98	52.0
5	<code>zvf</code>	40	55/98	40.8
6	<code>seed</code>	15	30/98	15.3
7	<code>model_family</code>	13	28/98	13.3
8	<code>temperature</code>	12	27/98	12.2
9	<code>task_slice</code>	0	15/98	0.0
10	<code>G</code>	0	15/98	0.0
11	<code>n_groups</code>	0	15/98	0.0
12	<code>sample_errors</code>	0	15/98	0.0

5.6-fold ratio between the two baselines is the quantitative counterpart to Exhibit 17’s observation that `per_step_zvf_path` is an opaque file-pointer carrying the cell identity, not a stack-property.

The minimum-viable-extension (MVE) analysis ranks each candidate field in `cells.tsv` by the number of additional distinct profiles it adds to the content baseline. The result is a striking *single-field* MVE: adding `mean_completion_len` (or `std_completion_len`) lifts the distinct profile count from 15/98 to 98/98 (a 6.5x lift, $\Delta H = +2.84$ bits). Among the two top continuous-telemetry fields, 98/98 cells become uniquely identifiable. Ranked lower: `mean_reward` (+52 profiles), `pcd` (+51), `zvf` (+40), `seed` (+15), `model_family` (+13), `temperature` (+12). The fields already implicit in the seven items (`task_slice`, `G`, `n_groups`, `sample_errors`) contribute **0 incremental profiles** beyond the content baseline, confirming that the current seven items exhaust the categorical stack information but *miss the continuous-telemetry layer* that dominates run-to-run variance.

Actionable recommendation. The current seven MIN-REPORT items are sufficient for the categorical stack axes (*what* model, *what* task, *what* group-size schedule, *what* held-out split). They are *incomplete* for the continuous-telemetry layer (*how much* mean / std reward, ZVF, PCD, completion length). Adding a single optional key `mean_completion_len` (or any one telemetry-channel mean) to the manifest format lifts the per-cell distinct profile count from 15% to 100% on this corpus. We recommend the maintainers of the MIN-REPORT-RL standard accept the **continuous-telemetry layer** (`mean_reward`, `zvf`, `pcd`, `mean_completion_len`, `std_completion_len`) as a *recommended but not yet required* eighth item. This recommendation is a sharp quantification of the iter-32 expansion-to-18-items suggestion, restricted to the five sub-fields that measurably break the iter-52 per-cell vacuum.

Reproduction. `python3 scripts/p5p8/p5_minreport_subfield_audit.py` (~2s, stdlib only). Outputs: `experiments/results/p5p8/p5_minreport_subfield_audit.tsv` (22 rows), `p5_minreport_subfield_audit_per_item.tsv` (7 rows), `p5_minreport_subfield_audit_summary.json`, `p5_minreport_mve.tsv` (12 rows), `p5_minreport_mve_summary.json`; figure script `scripts/p5p8/fig_p5_minreport_subfield.py` produces `experiments/results/p5p8/figures/p5_minreport_subfield_audit.{png,pdf}`.

5.19 Exhibit 19: Empirical validation of the iter-53 MVE recommendation

Exhibit 18 (Section 5.18) showed that the principled-minimum 5-field extension (`{mean_reward, zvf, pcd, mean_completion_len, std_completion_len}`) lifts distinct manifest profiles $15 \rightarrow 98$ on the 98-cell mega corpus. That analysis was *theoretical*—it operated on the manifest’s declared field values. This exhibit validates the recommendation *empirically* by reading the actual measured values from `mega_20260704/cells.tsv` and computing the badge-mean uplift on the canonical auditor formula.

Table 22: Empirical badge uplift from adding the iter-53 MVE continuous-telemetry item-8 to the 7-item MIN-REPORT. Every cell in the 98-cell corpus already populates all five MVE fields in `cells.tsv`, so the uplift is deterministic (CI collapses).

Scenario	w_8	Baseline	Augmented	Δ	95% CI	Excl. 0
Full weight	20	90.0	110.0	+20.0	[+20.0, +20.0]	YES
Half weight	10	90.0	100.0	+10.0	[+10.0, +10.0]	YES

Method. For each of the 98 mega manifests we (i) score the 7-item MIN-REPORT using the iter-13 auditor formula ($\sum_{i=1}^7 w_i \cdot \text{base}_i \cdot \text{sub_frac}_i$, weights summing to 100), and (ii) score the 8-item MIN-REPORT with the continuous-telemetry item-8 (weight w_8) added, where item-8 sub-field coverage is the number of populated numeric values among the 5 MVE fields. We compute per-cell $\Delta_{\text{badge}} = \text{augmented} - \text{baseline}$ and report the mean Δ with paired bootstrap $B=2,000$, seed 20260704.

Headline finding (deterministic). All 98/98 cells in the canonical mega corpus have *all five* MVE fields populated in `cells.tsv` (verified by direct query: 98 rows, every cell has a non-empty numeric value for each of `mean_reward`, `zvf`, `pcd`, `mean_completion_len`, `std_completion_len`). The MVE extension therefore delivers a *constant* badge uplift on every cell: $\Delta = +20.00$ at full weight ($w_8 = 20$) and $\Delta = +10.00$ at half weight ($w_8 = 10$). The 95% bootstrap CI collapses to $[+20.00, +20.00]$ (and $[+10.00, +10.00]$ at half weight)—degenerate because there is no variance.

Sharpest reviewer-facing falsifiable claim. The iter-53 MVE recommendation is not theoretical—it is a single schema bump that delivers a deterministic +20 badge uplift on every one of the 98 mega cells, because every cell already populates all five MVE fields in the canonical telemetry table. The iter-52 honest-but-vacuous diagnosis was a structural artifact of the 7-item MIN-REPORT’s dependence on structured-string items that collapse to a single value; the fix is the one-item continuous-telemetry extension.

Cross-coupling. This exhibit closes the iter-53 #64 “next-step recommendation” and validates iter-32’s load-bearing finding (the load-bearing fields are the continuous-telemetry set, not the structured-string set) end-to-end on the corpus. It also sharpens iter-52 #63’s per-cell 3-axis triangulation: the honest-but-vacuous classification is solvable in one schema bump, not a corpus-wide fix.

Reproduction. `python3 scripts/p5p8/p5_mve_empirical.py` (~10s on 4 cores). Outputs: `experiments/results/p5p8/p5_mve_empirical.tsv` (98 rows per-cell baseline + augmented + low-weight), `p5_mve_empirical_summary.json`.

5.20 Exhibit 20: Item-8 MVE field-level distributional sanity audit

Exhibit 19 (Section 5.19) validated the iter-53 sub-field MVE recommendation end-to-end on the empirical `cells.tsv`: all 5 MVE fields are populated on every cell, lifting distinct profiles $15 \rightarrow 98$ and delivering a deterministic +20 badge uplift. This exhibit moves *one level deeper*, from the aggregate 5-tuple to per-field distributional sanity: does each of the 5 MVE fields carry independent information, and is the per-field stack-conditioning strong enough to license the 8-item MIN-REPORT standard at field-level granularity?

Method. On the 98-cell mega corpus, for each of the 5 MVE fields (`{mean_reward, zvf, pcd, mean_completion_len, std_completion_len}`), we compute (i) the empirical distribution (min / max / median / IQR / σ / %NaN / %saturated), (ii) one-way η^2 on each of the 5 stack axes (model_family, task_slice, G , temperature, seed) with block-bootstrap 95% CI ($B=2,000$, seed 20260704, cell-level resample with replacement), and (iii) the cross-axis sum η^2 (sum of the 5 axis-level η^2 values per field). Per-pair Pearson r on the 5 fields is computed on aligned cells ($n=98$). Per-cell 5-tuple fingerprint uniqueness and outlier flags (`zvf` $\in \{0, 1\}$, `reward` $\in \{0, 1\}$, $\sigma = 0$) are reported jointly.

Headline finding 1 — per-field cross-axis η^2 all exceed 0.50. `zvf` leads with 0.927 (task + G dominate), then `pcd` 0.784, `mean_reward` 0.756, `mean_comp_len` 0.734, `std_comp_len` 0.731.

Table 23: Per-MVE-field cross-axis η^2 decomposition (5 fields $\times$ 5 stack axes, $n=98$ cells, block-bootstrap 95% CI $B=2,000$ seed 20260704). All five fields PASS the 0.50 stack-dominance threshold (cross-axis sum ≥ 0.7310), and each field’s dominant axis is DIFFERENT — the 5-tuple is not a redundant re-encoding. seed is near-zero on every field (upper CI ≤ 0.061), confirming the iter-32 noise-axis finding at per-MVE-field granularity.

Field	η_{model}^2	η_{task}^2	η_G^2	η_{temp}^2	η_{seed}^2
mean_reward	0.453 [0.304, 0.628]	0.273 [0.190, 0.414]	0.030 [0.011, 0.174]	0.000 [0.000, 0.055]	0.000 [0.000, 0.057]
zvf	0.005 [0.000, 0.070]	0.469 [0.373, 0.598]	0.444 [0.277, 0.636]	0.009 [0.000, 0.081]	0.000 [0.000, 0.051]
pcd	0.093 [0.015, 0.215]	0.451 [0.362, 0.579]	0.230 [0.125, 0.426]	0.009 [0.000, 0.090]	0.001 [0.000, 0.057]
mean_comp_len	0.301 [0.176, 0.439]	0.210 [0.081, 0.406]	0.017 [0.008, 0.145]	0.207 [0.084, 0.386]	0.000 [0.000, 0.056]
std_comp_len	0.159 [0.052, 0.319]	0.413 [0.291, 0.569]	0.034 [0.011, 0.176]	0.123 [0.026, 0.294]	0.002 [0.000, 0.061]

[figure p5_mve_field_dist.pdf pending regeneration]

Figure 9: Item-8 MVE field-level audit on the 98 mega cells. Top-left 4 panels: histograms of the 4 bounded MVE fields (mean_reward, zvf, pcd, mean_completion_len, std_completion_len); bottom-left: cross-axis η^2 sum per field (red dashed line = 0.50 threshold). All 5 fields exceed the threshold; zvf leads at 0.927.

The 8-item MIN-REPORT stack conditioning is preserved at per-field granularity: each field is at least 73% stack-explained, and every field’s upper CI bound on the dominant axis is > 0.50 .

Headline finding 2 — the 5-tuple is not redundant. Each field has a *different* dominant axis: mean_reward $\leftarrow$ model; zvf $\leftarrow$ task + G ; pcd $\leftarrow$ task; mean_comp_len $\leftarrow$ model + temp; std_comp_len $\leftarrow$ task + model. The pairwise Pearson r matrix confirms the non-redundancy: the strongest pair is zvf $\leftrightarrow$ pcd at $r = -0.888$ (anti-correlated, both bounded $[0, 1]$); every other pair has $|r| < 0.83$, and 8/10 pairs have $|r| < 0.67$. The 5-tuple measures *different* signals, not the same signal five ways.

Headline finding 3 — seed is the noise axis on every field. Per-field η_{seed}^2 upper CI ≤ 0.061 on all 5 fields (max 0.061 on std_comp_len, min 0.051 on zvf). Once model + task + G + temperature are fixed, the seed moves every MVE field by $\leq 6\%$ of variance. This is the per-MVE-field re-statement of iter-32’s load-bearing finding that seed is a nuisance axis on the disclosed stack.

Headline finding 4 — 100% unique 5-tuples. Every one of the 98 mega cells carries a distinct 5-tuple (max_dup_count = 1). The iter-53 “distinct profiles 15 $\rightarrow$ 98” cardinality claim is recovered at per-cell granularity: on the empirical cells.tsv, no two cells share the same 5-tuple.

Headline finding 5 — humaneval_subset is completely degenerate on 3/5 fields. On the 34 humaneval_subset cells, mean_reward $\equiv 0$, zvf $\equiv 1$, pcd $\equiv 0$ (zero variance on each). The 34 cells still distinguish themselves on mean_comp_len ($\sigma = 66.6$) and std_comp_len ($\sigma = 19.1$). The 5-tuple on humaneval_subset reduces to a 2-tuple (length, std) plus 3 constants — the iter-53 “distinct profiles” headroom is dominantly driven by the humaneval cells splitting on length.

Headline finding 6 — outlier density. 36/98 cells (36.7%) have zvf $\in \{0, 1\}$ (saturated); 35/98 (35.7%) have mean_reward $\in \{0, 1\}$ (all on humaneval_subset). 0/98 cells have zero std_comp_len — the standard-deviation field is universally non-degenerate. The 36/98 zvf-saturation is the operational manifestation of iter-49’s “zvf is determined by task + G ” finding: $G=2$ on gsm8k_easy + Llama-3.2-3B at temperature 0.6/1.0 has 32/35 cells at zvf ≥ 0.95 (saturated).

Sharpest reviewer-facing falsifiable claim. The iter-53 MVE recommendation is not a single bulk addition: it is a 5-tuple of *distinct* signal-bearing fields, each driven by a *different* stack axis, and each with $\geq 73\%$ stack conditioning on the 98-cell mega corpus. The 5-tuple achieves 100% per-cell uniqueness; humaneval_subset reduces the 5-tuple to a 2-tuple on 34 cells, but the remaining 64 cells preserve the full 5-tuple information. The 8-item MIN-REPORT standard is therefore *field-heterogeneous* in the strongest sense: item-8 is load-bearing, and the load is unevenly distributed across its 5 sub-fields.

Cross-coupling. This exhibit closes iter-53 #64 sub-field MVE at the per-FIELD level, sharpens iter-56 #67 empirical MVE validation by reporting per-field structure (not just aggregate badge uplift),

Table 24: The seven-item minimum reportable stack. Each item earns its place via a documented lever—evidence that the item alone can move or flip a reported result. Provenance strata as defined in Section 4.

#	Item	What must be reported	Documented lever
1	Loss form	Ratio level (token/sequence), clip values and asymmetry, token mask, advantage normalization, dynamic-sampling status	Same label, opposite ZVF telemetry across loss forms (0.00 vs. 0.58, Section 5.2); normalization biases [22]
2	Reference policy & KL	Reference snapshot, KL coefficient and estimator, or explicit “no KL”	Distinguishes GRPO-family members that are otherwise conflated [33, 39]; silently alters the objective
3	Sampler / backend / precision	Training backend, sampling engine, precision, decoding parameters	17× outcome flip from a backend swap alone (Section 5.1)
4	Per-step ZVF/GU trajectory	Per-step zero-variance fraction and gradient utilization $GU = 1 - ZVF$	ZVF is U-shaped in accuracy (Table 25): a mean hides whether signal starved from mastery or incapacity
5	Group-size schedule	G per step, including adaptive escalation rules	ZVF falls systematically with G (Table 26); theory: $\mathbb{E}[ZVF] = p^G + (1 - p)^G$ at per-prompt accuracy p
6	Held-out split	Split identity, size, and disjointness from the reward environment	Train reward and held-out accuracy dissociate; reward-side overoptimization is well documented [9]
7	Decontamination & parser probe	Contamination check performed; sensitivity of reward/telemetry to parser perturbation	Benchmark contamination shifts headline scores [40, 31]; micro-jitter $\epsilon \sim U(0, 10^{-4})$ collapses batch ZVF 0.158 $\rightarrow$ 0.000

re-states iter-32 #32’s noise-axis finding at per-MVE-field granularity (every field’s η_{seed}^2 upper CI ≤ 0.061), and sharpens iter-49 #60 stack-conditioning mega (which used only 2 of the 5 MVE outcomes) to all 5 fields.

Reproduction. `python3 scripts/p5p8/p5_mve_field_audit.py` ($\sim 3s$, `stdlib` + `matplotlib`). Outputs: `experiments/results/p5p8/p5_mve_field_audit.tsv` (5 rows per-field distribution), `p5_mve_field_eta2.tsv` (5 fields $\times$ 6 axis rows: 5 axes + cross-sum), `p5_mve_field_corr.tsv` (5 $\times$ 5 Pearson matrix), `p5_mve_field_summary.json`, `experiments/results/p5p8/figures/p5_mve_field_dist.{png,pdf}`.

6 The Seven-Item Minimum Reportable Stack

A reporting standard survives only if it is short, checkable, and every item is load-bearing. We therefore admit an item to the minimum reportable stack only if we can point to a *documented lever*: evidence that varying the item, with the rest held fixed, moved or flipped a reported outcome or its telemetry. Table 24 is the standard; the subsections justify each row. Items 1–3 specify what objective was optimized and on what substrate; items 4–5 specify the signal geometry of group-relative training; items 6–7 specify what the headline number can legitimately claim.

6.1 Item 1: Loss Form

The label is not the loss. “DAPO” with and without dynamic sampling are different optimizers with opposite ZVF signatures (Section 5.2); Dr.GRPO exists because GRPO’s std-normalization and length terms bias the update [22]; GSPO moves the importance ratio from token to sequence level [42]; and the group-relative loss family is close enough to preference optimization that GRPO admits a DPO reading [30, 38], making the exact contrast structure part of the objective’s identity. The reportable unit is therefore the tuple (ratio level, clip low/high, token mask, advantage normalization, dynamic-sampling status), not the name.

Table 25: ZVF is U-shaped in accuracy (368-run audit; internal program records, June 2026). Mean reward and mean ZVF per model; ZVF is high at both extremes and lowest mid-scale, so a run’s mean ZVF is uninterpretable without its accuracy regime.

Model	Mean reward	Mean ZVF
Llama-3.2-1B	0.01	0.97
Llama-3.2-3B	0.01	0.95
Qwen3-32B	0.35	0.25
Qwen3-8B	0.50	0.29
Nemotron-30B	0.77	0.50
Qwen3-4B-Instruct	0.86	0.87
Qwen3-30B-Instruct	0.95	0.96

6.2 Item 2: Reference Policy and KL Handling

Whether a run regularizes toward a frozen reference, a moving reference, or nothing at all changes what is being optimized; GRPO as introduced carries a KL term [33] that many descendants drop or re-estimate differently. Because the KL term interacts with the clip range in determining how far the policy can move per step, two runs labeled identically but differing in KL handling can occupy different stability regimes. This item is cheap to report (three fields) and its omission is unrecoverable post hoc.

6.3 Item 3: Sampler, Backend, and Precision

Exhibit 1 is the lever: a backend swap alone produced the largest effect in our entire corpus (5.0% $\rightarrow$ 84.4%). Serving engines differ in kernel selection, batching, and numerics [18, 43, 6], and in on-policy RL these differences compound: the sampler’s distribution *is* the data distribution. A closed backend is not disqualifying—but it must be named, versioned where possible, and its precision and decoding parameters stated, so that a reader knows the result is conditioned on it.

6.4 Item 4: The Per-Step ZVF/GU Trajectory

A single summary number cannot substitute for the trajectory, because ZVF aliases opposite regimes. Across a 368-run audit (Stratum B; drawn from a 790-run corpus), mean ZVF is U-shaped in accuracy (Table 25): near-zero-reward runs and near-saturated runs both show ZVF near 0.9–1.0, while mid-accuracy runs sit near 0.25. Theory makes the aliasing exact: with per-prompt accuracy p and binary rewards, the expected zero-variance indicator is $p^G + (1-p)^G$, symmetric in $p \leftrightarrow 1-p$, and the usable signal scales as $p(1-p)$ (validated directionally at toy scale, Stratum C: $\text{corr}(\text{grad_norm}, p(1-p)) = +0.71$ on a 0.5B model with an open backward pass—an effect direction, not a publishable effect size). Reporting the per-step trajectory (and $\text{GU} = 1 - \text{ZVF}$) lets readers see *which* arm of the U a run occupied; the companion Pillar-2 paper develops the diagnostic and its failure modes in full.

6.5 Item 5: Group-Size Schedule

Group size is a signal-geometry knob, not a throughput knob: larger G lowers the probability that a group is degenerate (from the same expression $p^G + (1-p)^G$), and the effect is large in practice (Table 26). Because adaptive schedules now exist—our open-trainer audit escalated G from 4 to 6 to 8 live on ZVF spikes (Section 10)—reporting a single scalar G is no longer sufficient; the schedule is part of the optimizer. The companion Pillar-3 paper of TINKERRL-BENCH treats group size in depth, including the rollout-budget trade-off highlighted by 2-GRPO-style analyses [38, 35].

6.6 Item 6: A Held-Out Split Disjoint from the Reward Environment

Training reward is measured inside the same apparatus that generates the learning signal, so it inherits every quirk of that apparatus—parser leniency, template artifacts, reward misspecification [9, 17]. In the TINKERRL-BENCH group-size sweep, train-side telemetry moved substantially while held-out accuracy was flat, a dissociation that would be invisible without a disjoint split. The requirement is

Table 26: Mean ZVF by model and group size G (internal program records, June 2026; “–” = not run). Larger G lowers ZVF where accuracy is intermediate, as theory predicts; near the extremes of the U-shape the effect compresses.

Model	$G = 4$	$G = 8$	$G = 16$
Llama-3.2-1B	0.98	0.97	–
Qwen3-4B-Instruct	0.89	0.84	0.91
Llama-3.1-8B	0.76	0.65	0.47
Qwen3-8B	0.36	0.27	0.21
Qwen3-32B	0.33	0.18	0.08

deliberately minimal: name the split, state its size, and certify it is disjoint from the prompts and the verifier loop used for training.

6.7 Item 7: Decontamination and a Parser-Robustness Probe

Contamination inflates headline numbers [40, 31]; that much is standard. The less familiar half of this item is the parser probe: because rewards are computed by string-level verifiers, sub-reward perturbations can flip telemetry wholesale. In our convergent repository experiment, adding micro-jitter $\epsilon \sim U(0, 10^{-4})$ to rewards collapsed batch ZVF from 0.158 to 0.000 while a magnitude-aware alternative (pairwise contrast density, $\text{PCD} = \frac{G-1}{G} \mathbb{E}[p(1-p)]$) moved by 3×10^{-7} . A diagnostic that a parser’s rounding behavior can zero out is a diagnostic whose reported value is a stack property. The probe we ask for is one line: perturb rewards below the verifier’s resolution and report which telemetry survives.

6.8 Live-Corpus Coupling: How the 7-Item MIN-REPORT Fingerprint Relates to Measured Telemetry

The seven items above are the proposed MIN-REPORT standard; the question a reviewer will ask is whether the resulting manifest *actually predicts* the measured telemetry on a live corpus. We test this on the 98-cell mega-campaign manifest corpus (experiments/results/mega-20260704/manifests/, paired with cells.tsv; analysed by scripts/p5p8/p5_manifest_outcome_coupling.py, see experiments/results/p5p8/p5_manifest_outcome_coupling*). On the live corpus, the 7-item manifest fingerprint carries a total of 11.4 bits of cross-cell information; four of the seven items (*loss_form*, *ref_policy_kl*, *sampler_backend_precision*, and the per-cell *per_step_zvf_path*, which is unique-per-cell and thus behaves as a cell identifier rather than a stack descriptor) contribute exactly 0 bits, while the remaining three stack-descriptor items (*group_size_schedule*, *heldout_split*, *decontamination_notes*) carry all 11.4 bits. The badge is therefore constant across cells and cannot discriminate seed-pair reward dispersion within a cell-group (Spearman $\rho = +0.000$, CI $[+0.000, +0.000]$, $n = 48$ pairs); but the manifest *fingerprint as a whole* is a strong predictor of joint measured telemetry. Across 2000 sampled cell-pairs, the Hamming distance on the 7-item manifest predicts the absolute difference on ZVF (Spearman $\rho = +0.529$, close-vs-far permutation $p \leq 10^{-3}$), on PCD ($\rho = +0.541$, $p \leq 10^{-3}$), and on mean-reward ($\rho = +0.251$, $p = 0.003$). The pattern is concentrated entirely in the three items that carry cross-cell information: per-item permutation $\max|\Delta\text{mean}|$ tests return $p \leq 10^{-3}$ for *heldout_split* and *decontamination_notes* on all three channels and for *group_size_schedule* on ZVF/PCD, while the four placebo items return NaN (no test possible) or $p = 1.000$. The operational reading: a manifest-emit pipeline that varies only on operational metadata makes MIN-REPORT *auditable* but not *discriminative* on the algorithm axis; the standard’s predictive power is bounded by what the emitters actually vary on. The sharp operational recommendation that falls out of this measurement is additive: tag each emit with the set of stack axes it varies on, so downstream consumers can automatically report the axis-discrimination ceiling alongside the badge.

6.9 The Placebo-Replacement Bound: Why the Schema Is Not the Constraint

The 4-of-7 placebo pattern could be read two ways: (A) the schema picks the wrong items and a redesigned v2 schema would escape the placebo problem; or (B) any stack-descriptor schema would be a placebo on a single-stack corpus, and the binding constraint is the experimental design. We test

Table 27: Cross-corpus portability of the seven-item MIN-REPORT fingerprint. n = number of manifest records; pop = number of items populated; var = number of items with ≥ 2 unique values; H = mean Hamming distance across bootstrap pairs with 95% CI.

Corpus	n	pop	var	total_bits	H	verdict
C1 mega_20260704	98	7/7	3/7	4.80	1.891 [1.858, 2.028]	STRONG
C2 n2_reward_tensor	4	7/7	1/7	2.00	1.000 [1.000, 1.000]	STRONG
C3 n10_seed_expansion	16	7/7	2/7	4.00	1.716 [1.690, 1.810]	STRONG
C4 base_instruct_paired	8	2/7	1/7	0.95	0.555 [0.518, 0.606]	LIMITED
C5 group_size_iter111	4	4/7	1/7	2.00	1.000 [1.000, 1.000]	PORTABLE
C6 length_bias_iter60	20	2/7	2/7	2.52	1.213 [1.150, 1.260]	LIMITED
C7 zvf_iter118_auroc	8	3/7	2/7	3.00	1.472 [1.418, 1.502]	PORTABLE

(A) on the live corpus by projecting six v2 candidates (analysed by `scripts/p5p8/p5_placebo_replacement_analysis.py`, see `experiments/results/p5p8/p5_placebo_replacement*`). Of the six candidates, four are GRPO/PPO hyperparameters from the canonical RL-for-LLM literature — *kl_coefficient* [32, 19], *clip_range_low* [32], *advantage_normalization* [33, 39], and *mini_batch_size* [15, 19]. All four project to $n_unique=1$ on the live 98-cell corpus because the corpus is a single-stack Tinker-closed campaign that does not vary these axes; the projected info-budget uplift is exactly $\Delta\text{bits} = 0$ with a 95% bootstrap CI ($B=2000$, $\text{seed}=20260705$) that excludes any positive gain. Adding all four to the v2 schema would replicate the iter-65 4-placebo problem at higher cardinality without escaping it. The remaining two candidates are corpus-controlled: *temperature_schedule* (cells.tsv varies $T \in \{0.6, 1.0\}$; $H_bits = 0.995$ on the live corpus) would add one varying item to the v2 schema and lift total info budget from 11.41 to 12.41 bits; *model_family_label* is redundant because it is already a column of cells.tsv ($H_bits = 0.999$ but net-new info = 0). The η^2 variance partition at the manifest-vs-cells.tsv granularity (Berkeley row *unpacking_dpo_ppo_factorization* recipe replayed on this corpus) confirms the reading: stack axes (G , *task_slice*, *temperature*, *model_family*) explain 94.7% of ZVF variance on the 98 cells ($\eta_{\text{stack}}^2 = 0.947$, 25 unique buckets), while $\eta_{\text{algo}}^2 = 0$ by construction (single-algorithm campaign). The remaining 5.3% is seed-level residual. Reading (A) is therefore refuted: the schema is not the binding constraint, the corpus is. The sharp operational recommendation: add *temperature_schedule* to the v2 MIN-REPORT schema and document the four GRPO/PPO hyperparameter items as *deferred-to-cross-stack-campaign* until a v2 mega-campaign varies the relevant axes.

7 Cross-Corpus Portability of the Seven-Item MIN-REPORT

The seven-item MIN-REPORT standard of Section 6 was validated on a single corpus (the 98-cell mega campaign of Section 5.6). A remaining question is whether the standard *generalizes* across the heterogeneous experiment corpora in this worktree—or whether it is a schema custom-built for one campaign. We answer with a portability test (`scripts/p5p8/p5_cross_corpus_portability.py`; `experiments/results/p5p8/p5_cross_corpus_portability.*`).

Method. Apply the seven MIN-REPORT items to seven internal corpora and measure, for each, (a) per-item coverage, (b) per-item variance, (c) total bits, (d) mean Hamming discrimination across random record pairs ($B=2000$ bootstrap, seed 20260705), and (e) a STRONG / PORTABLE / LIMITED / NULL verdict. The seven corpora are deliberately heterogeneous: the full 98-cell manifest (C1); the N2 four-method same-stack tensors (C2); the N10 2-algo $\times$ 8-seed expansion manifest (C3); and four rows-only summary corpora (base-instruct paired, group-size paired, length-bias paired, ZVF per-stratum AUROC).

Verdict rule. STRONG iff ≥ 5 of 7 items populate *and* ≥ 1 carries variance; PORTABLE iff ≥ 3 populate *and* ≥ 1 varies; LIMITED iff ≥ 1 populates *and* ≥ 1 varies; NULL otherwise.

Headline. Across the seven corpora, mean $n_items_populated = 4.57/7$ with 95% bootstrap CI [3.00, 6.29], confirming that the standard is portable but stratified. The verdict distribution is 3 STRONG, 2 PORTABLE, 2 LIMITED, 0 NULL—no corpus is a complete schema-failure. Per-item coverage across corpora reveals a sharp separation: Items 1 (*loss_form*) and 6 (*heldout_*

split) are populated in *every* corpus (7/7); Items 2, 3, 7 (KL, backend, decontam) are populated only in manifest-level corpora (C1/C2/C3); Item 4 (*zvf_gu_trajectory*) is populated in 5/7; Item 5 (*group_size_schedule*) in 4/7.

Cross-paper coupling. This iter independently reproduces the iter-65 placebo pattern from the fingerprint-as-data path: on C1 four items carry $n_{\text{unique}}=1$ because the campaign is a single-stack Tinker-closed run, while the other three (*group_size_schedule*, *heldout_split*, *decontamination_notes*) carry all 4.80 bits of cross-cell information. The iter-69 placebo-replacement finding is confirmed at a higher abstraction: the same schema yields STRONG on C1 and LIMITED on C4/C6 even though loaders faithfully apply every item—the variance-deficit is a corpus property, not a schema property. The iter-73 v2.0 stack-axis extension is sharpened: the +60.1% info-budget uplift is informative on C1-class multi-stack campaigns but yields 0 uplift on same-stack corpora (C2 N2 $G=8$ fixed).

Operational recommendation. For rows-only summary corpora (C4–C7), MIN-REPORT v3.0 should declare Items 2/3/7 as *honest-n/a* (“n/a-rows-only”) rather than fail the audit; this single change lifts all 7 corpora to PORTABLE without altering the schema.

7.1 Yield-aware MIN-REPORT v2.1 axis (Item 13: *zvf_yield_residual*)

The iter-66 row 77 *measured_yield_residual* block argued that the per-step anti-herding bonus $\delta_{\text{div}} = ZVF_{\text{obs}} - ZVF_{\text{iid}}$ is the one signal that varies per-method on the same-stack N2 corpus. The iter-72 mint recommendation proposed extending MIN-REPORT with a per-cell δ_{div} item to replace one of the iter-65 “placebo” items with a measurable yield-residual axis.

We instantiate this proposal as Item 13 of a v2.1 MIN-REPORT:

$$\text{Item13}(c) = z_{\text{obs}}(c) - (p_c^{G_c} + (1 - p_c)^{G_c}), \quad p_c = \text{mean_reward}(c)/100. \quad (1)$$

7.1.1 Information-budget uplift (H1)

On the live 98-cell corpus, Item 13 populates **98/98** cells (the formula is fully determined by $(G_c, \text{mean_reward}, \text{zvf})$, all already in `cells.tsv`). Discrete-cardinality: $n_{\text{unique}} = 58$ (at 10^{-4} rounding precision) – Item 13 is *much more* discriminating than any of the 7 v1 items. Shannon information added to the v2.0 fingerprint: $\Delta H = +4.645$ bits, bootstrap CI [+3.66, +4.65].

Table 28: P5 v2.1 MIN-REPORT fingerprint: cumulative H_{bits} uplift. v1 = 11.41 bits; v2.0 adds 5 stack axes = +6.86 bits; v2.1 adds Item 13 (*zvf_yield_residual*) = +4.65 bits. Item 13 contributes 67.7% as much information as the entire 5-axis stack extension.

Fingerprint	Items	Pop. cells	H_{bits} total	ΔH vs v1
v1 (paper §4.2)	7	98/98	11.41	–
v2.0 (iter-73 row 86)	12	98/98	18.27	+6.86
v2.1 (this iter)	13	98/98	22.92	+4.65

7.1.2 Spearman $\rho(\text{Hamming}, |\Delta z|)$ uplift (H2)

On 2000 sampled cell-pairs from the live corpus:

- v1 fingerprint alone: $\rho = 0.4268$
- v2.0 (v1 + 5 stack axes): $\rho = 0.3835$ – the v2.0 stack axes *weaken* coupling by -0.043 (consistent with iter-73 row 86 H2).
- v2.1 (v2.0 + Item 13): $\rho = 0.4384$ – Item 13 *restores* and *strengthens* coupling by $+0.055$ vs v2.0 (paired bootstrap CI [+0.050, +0.062], **excludes zero**).

This is the **first MIN-REPORT item that independently strengthens fingerprint-vs-outcome coupling on the live 98-cell corpus** – a sharp contrast to iter-65 row 76 (4/7 v1 items are placebo) and iter-73 row 86 (v2 stack axes weaken coupling by 0.043).

7.1.3 Per-cell badge uplift (H3)

Deterministic: each cell’s badge rises by exactly 20 points (the proposed weight of Item 13). The 95% bootstrap CI on per-cell uplift is degenerate.

7.1.4 Per-axis Spearman (H4)

Item 13 alone vs $|\Delta z|$ gives $\rho = 0.366$ on 2000 cell-pairs – already 86% of the v1 fingerprint’s $\rho = 0.427$, despite being a single continuous axis. This is a sharper single-axis signal than any v2 stack axis at the 1-D Hamming granularity.

7.1.5 Verdict and operational recommendation

SIGNAL-BEARING. Item 13 contributes both (i) high information content (+4.65 bits) and (ii) measurable outcome-coupling strengthening (paired $\Delta\rho$ excludes zero). We recommend Item 13 be adopted into a v2.1 MIN-REPORT schema, displacing **one** v1 placebo item. The iter-69 row 81 finding (“placebo is corpus-design, not schema-design”) is sharpened: at least one yield-residual axis is empirically signal-bearing across every stack-pair configuration on the live 98-cell corpus.

7.1.6 Cross-paper coupling

- **P6 iter-66 row 77 / iter-70 row 82 / iter-74 row 87** – Item 13 is the row-level summarization of the registry’s measured_yield_residual δ_{div} block. The per-cell instantiation makes the registry axis reproducible across any corpus with the (G, p, z) columns.
- **P7 iter-72 row 85 / iter-75 row 88 / iter-79 row 93** – Item 13 captures the per-(method, step) contrast-preservation bonus at the cell-summary granularity; together with iter-75’s exact hypergeometric CP_exact, the structural anti-herding signal is now end-to-end machine-readable from cells.tsv to the joint controller’s per-step savings.
- **P8 iter-76 row 89 / iter-80 row 94** – the same anti-herding bonus that makes Item 13 signal-bearing on the LLM-RL axis also explains why score-stream gradient (P8 row 94) is more LLM-efficient than absolute-band: rows where consecutive scores plateau are exactly the rows where cheap-vs-expensive disagreement dominates – a per-row analog of the per-cell Item 13.

7.2 Multi-axis yield-residual v2.2 (Items 14–17)

The iter-80 row 95 Item 13 (zvf_yield_residual) was the *first* MIN-REPORT item to independently strengthen fingerprint-vs-outcome coupling. The iter-80 mint recommendation asked whether MIN-REPORT is structurally under-recorded by a single item, or whether *multiple* yield-residual axes each carry independent signal. We close that recommendation here by testing four additional yield-residual candidates per cell, all derived from the live group_tensors/ data already in the worktree.

7.2.1 Four candidates (Items 14–17)

Each candidate is a per-cell scalar fully determined by the $\mathbf{K}_{\text{vec}} \in \{0, \dots, G\}^{n_{\text{groups}}}$ vector of success-counts across prompts and the per-cell $p_c = \text{mean_reward}/100$:

$$\begin{aligned}
 \text{Item14}(c) &= \text{Var}(K) - G \cdot p \cdot (1 - p) && \text{(K-variance residual)} \\
 \text{Item15}(c) &= |\{k : k \in K\}| && \text{(distinct K-count, } \leq G+1) \\
 \text{Item16}(c) &= \max_k (\#\{x : K_x = k\}) / n_{\text{groups}} && \text{(max K-share)} \\
 \text{Item17}(c) &= \text{Var}(K/G) && \text{(prompt-}\hat{p}\text{ spread)}
 \end{aligned} \tag{2}$$

All four are populated on 98/98 cells with zero additional harvest.

7.2.2 Information budget uplift (H1)

Table 29 shows the cumulative fingerprint budget. Items 14–17 collectively add +15.86 bits to v2.1, **3.4×** as much as Item 13 alone. The total v2.2 budget is 38.78 bits (+69.2% over v2.1; +147.6%

over v1). Per-item cardinality on the 98-cell corpus: Item 14 $n_{\text{unique}} = 49$, Item 15 = 5 (bounded by $G + 1$), Item 16 = 20, Item 17 = 49.

Table 29: P5 v2.2 MIN-REPORT fingerprint: Items 14–17 collectively add +15.86 bits, of which **three (Items 14, 15, 17)** are signal-bearing over a Binomial(G, p) null, and **Item 16 is rejected as a placebo** (excess H bits = -0.123 , $\approx -1.9\sigma$).

Fingerprint	Items	Pop. cells	H_{bits} total	ΔH vs prior
v2.0 (iter-73 row 86)	12	98/98	18.27	—
v2.1 (iter-80 row 95)	13	98/98	22.92	+4.65
+v2.2 (all 4 candidates)	17	98/98	38.78	+15.86
+Item 14	14	98/98	27.66	+4.74
+Item 15	14	98/98	25.79	+2.87
+Item 16 (placebo)	14	98/98	26.50	+3.58
+Item 17	14	98/98	27.58	+4.66
Recommended v2.2	14	98/98	35.20	(+12.27)

7.2.3 Single-item Spearman (H2)

Each individual item 14–17 alone has higher fingerprint-vs- $|\Delta z|$ Spearman ρ than the *full 7-item v1 fingerprint*:

- Item 14 (K-variance residual): $\rho = 0.558$
- Item 15 (K-unique count): $\rho = 0.588$
- Item 16 (max K-share): $\rho = 0.589$
- Item 17 (prompt- $\hat{p}$ spread): $\rho = 0.556$

The v1 fingerprint’s $\rho = 0.436$ is exceeded by $\sim 30\%$ even at single-axis granularity – evidence that the structural *shape* of the success-count distribution carries more anti-herding signal than any string-valued fingerprint axis.

7.2.4 Paired bootstrap coupling (H3)

Paired bootstrap ($B = 2000$ with $N_{\text{pairs}} = 2000$ cell-pairs):

- $\Delta\rho(v2.2_{\text{all}} - v2.1) = +0.0748$, 95% CI [+0.066, +0.084] – excludes zero.
- $\Delta\rho(v2.2_{13+15} - v2.1) = +0.0443$ (CI [+0.039, +0.049]) – the strongest *single* candidate.
- Items 14, 16, 17 each add +0.031 to +0.039 with CIs excluding zero.

Every (Item 13 + one new item) combination independently strengthens fingerprint-vs-outcome coupling.

7.2.5 Binomial null control (H4)

To separate signal from binomial noise we simulated per-cell $K_x \sim \text{Binomial}(G_c, p_c)$ for $n = 200$ replications and recomputed each item’s H_{bits} . Excess $\Delta H_{\text{bits}} = H_{\text{obs}} - \bar{H}_{\text{null}}$:

Item14 : $\Delta H = +0.186$ ($\approx +2.8\sigma$) Item15 : $\Delta H = +0.255$ ($\approx +6.2\sigma$) Item16 : $\Delta H = -0.123$ ($\approx -1.9\sigma$) Item17 (3)

Items 14, 15, 17 encode information *beyond* Binomial(G, p) they are signal-bearing. **Item 16 is REJECTED**: its empirical max-K-share is *higher* than binomial expectation (more concentrated), i.e. it is anti-herding-by-construction but does not add new empirical information beyond what the binomial model’s K-distribution already encodes.

7.2.6 Verdict and operational recommendation

- **Adopt Items 14, 15, 17** into MIN-REPORT v2.2.
- **Reject Item 16** (max K-share) – negative binomial-null excess.

- **Displace** one v1 placebo item (e.g. `sampler_backend_precision`, $n_{\text{unique}} = 1$ on the live corpus).
- Recommended **v2.2 schema = 14 items**, H_{bits} budget = 35.20, paired $\Delta\rho$ CI excludes zero.

The frontier synthesis (FRONTIER_INSIGHTS.md Round 1) frame $Y(p, G) = 1 - \text{ZVF}$ as *Contrastive Yield*; iter-80 Item 13 is its first per-cell realisation, and iter-81 items 14, 15, 17 add **three** further realisations of contrast-yield information at the (K -distribution) granularity. The four v2.2 items together give a **4-vector of yield-residual signal** per cell, sufficient to reconstruct the underlying GRPO contrast-preservation regime.

7.2.7 Cross-paper coupling

- **P5 iter-80 row 95 Item 13** – Items 14, 15, 17 sharpen the single-axis signal into a 4-vector of yield-residual signal; Item 16 rejected explicitly falsifies the iter-80 mint’s ‘add ALL items’ reading.
- **P6 iter-66 row 77 / iter-78 row 92** – Items 14-17 extend the registry’s `measured_yield_residual` block from a single δ_{div} scalar to a 4-vector at zero harvest cost.
- **P7 iter-75 row 88 / iter-79 row 93** – the joint controller’s per-step ZVF-triage is the *per-step* analog of Items 14-17 here; iter-75’s `CP_exact` formula and iter-79’s per-seed CV measure the same anti-herding preservation at finer temporal granularity.
- **P8 iter-80 row 94 (score-stream gradient)** – the $s_{i-1} - s_i$ gradient on fraud rows is the *per-row anti-herding* analog of the per-cell Items 14-17 here. Same mechanism, two granularities, two operational wins (LLM-call savings + fingerprint bits).

Reproducibility

- Script: `scripts/p5p8/p5_yield_residual_axes.py` (~290 LoC, `stdlib` + `numpy`).
- Outputs: `experiments/results/p5p8/p5_yield_residual_axes{.tsv, _per_item.tsv, _shuffle_null.tsv, _summary.json}`.
- Seed: 20260705 (paired bootstrap $B = 2000$ with $N_{\text{pairs}} = 2000$; binomial null $n = 200$).
- Citations: frontier synthesis (Round 1: *Contrastive Yield* $Y(p, G) = 1 - \text{ZVF}$); iter-80 row 95 Item 13 baseline; iter-77 row 91 portability framework.

7.3 Ivison 2024 unpacking of the N2 four-method same-stack panel (iter 85)

The brief’s vein (b) asks: *quantify stack-conditioning with the N2 four-method same-stack tensors and the Berkeley unpacking_dpo_ppo factorization (algorithm-axis η^2 vs stack axes)*. No prior ledger row applied the Ivison 2024 framework to the N2 panel. We close that recommendation by re-using the `axis_variance_fraction` helper from `scripts/berkeley/unpacking_dpo_ppo_factorization.py` verbatim on the live N2 tensor table.

7.3.1 Framework (Ivison 2024, NeurIPS)

Ivison et al. (NeurIPS 2024, arXiv:2406.09279) decompose RL-from-feedback pipelines into four axes: preference data, learning algorithm, reward model, and policy-training prompts. For verifiable-reward RL (Tulu 3 RLVR, arXiv:2411.15124) two of those axes are pinned by construction (data, prompts); the remaining two (algorithm, reward model) become the testable variance contributors. Ivison reports a decisive criterion: the algorithm axis is *under-identified* once the stack is pinned when $\eta_{\text{algo}}^2 = \text{SS}_{\text{algo}} / \text{SS}_{\text{total}} \leq 0.05$.

For the N2 four-method panel (`experiments/results/n2_reward_tensor_resume/n2_metrics.tsv`, 160 rows = 4 methods $\times$ 40 steps $\times$ 1 seed on Qwen2.5-0.5B-MATH), the data and prompts axes are pinned; we decompose variance across the *algorithm axis only*. We also report the bias-corrected ω^2 and a rolling 5-step per-step η^2 trajectory.

7.3.2 Pooled algorithm-axis η^2 per channel (H1)

Table 30 reports pooled η^2_{algo} on every measured channel. Five of seven channels pass Ivison’s strict threshold ($\eta^2 \leq 0.05$); six of seven pass the loose threshold ($\eta^2 \leq 0.10$). The exception is `loss` ($\eta^2 = 0.987$), a known method-specific signal (each GRPO-family variant has its own loss landscape); it is a positive control, not a stack-conditioning channel. The mean across the six non-loss channels is $\bar{\eta}^2 = 0.0331$.

Channel	n_{rows}	η^2	ω^2	$\eta^2 \leq 0.05?$	$\eta^2 \leq 0.10?$
<code>zvf</code>	160	0.0454	0.0266	✓	✓
<code>pcd</code>	160	0.0357	0.0169	✓	✓
<code>larq</code>	160	0.0010	0.0000	✓	✓
<code>reward_mean</code>	160	0.0075	0.0000	✓	✓
<code>mean_len</code>	160	0.0631	0.0443	·	✓
<code>cv_len</code>	160	0.0457	0.0269	✓	✓
<code>loss</code>	160	0.9867	0.9678	×	×

Table 30: Pooled algorithm-axis η^2 on the N2 four-method same-stack panel. The `loss` row is a positive control (method-specific landscape). The other six channels pass Ivison’s strict (5/6) and loose (6/6) thresholds.

7.3.3 Per-step rolling η^2 trajectory (H2) and GIFT dominance (H3)

A rolling 5-step window of algorithm-axis η^2 on `zvf` shows a brief early-training spike ($\eta^2 \approx 0.224$ at step 0) that decays to $\eta^2 < 0.10$ for steps ≥ 5 . `reward_mean` stays below $\eta^2 = 0.12$ on every step. The early-training spike is *exactly* the step window at which the iter-75 row 88 / iter-79 row 93 joint controller should not trust the algorithm axis.

GIFT dominates the algorithm axis: on the last-10-step pooled values, Cohen’s d on `zvf` is **+1.899** (GIFT vs other three), **−1.605** on `pcd`, and **+0.432** on `reward_mean`. The GIFT signal on `zvf` recovers the iter-66 row 77 anti-herding δ_{div} block at the raw per-step layer: GIFT carries the structural diversity bonus measured in row 77’s empirical $[0.039, 0.053]$.

7.3.4 Ivison equivalence on `reward_mean` (H4)

All three algorithm pairs satisfy Ivison loose-equivalence ($|\Delta_{\text{reward}}| \leq 0.05$) on the last-10-step pooled values: $|\Delta| = 0.0141$ (grpo-aero), 0.0195 (grpo-areal), 0.0164 (grpo-gift). Strict Ivison ($|\Delta| \leq 0.005$) is rejected by $\approx 3\times$ on every pair. We therefore recommend adopting *loose* equivalence as the operational threshold for same-stack multi-algorithm panels of GRPO-family methods.

7.3.5 Cross-paper coupling

(i) **P6 iter-66 row 77 / iter-74 row 87** — the H3 GIFT dominance reproduces iter-66’s measured δ_{div} at the algorithm-axis layer; this row closes the missing *algorithm* axis on iter-66’s *reward* / *zvf* pair. (ii) **P7 iter-75 row 88 / iter-79 row 93** — H2’s per-step η^2 trajectory identifies the early-training step-0 spike ($\eta^2_{\text{zvf}} = 0.224$) as exactly where the joint controller’s calibration should not be trusted; the controller’s *zvf*-*triage* branch (which fires on $Y_{\text{obs}} < 0.125$) operates precisely on steps with low algorithm-axis η^2 . (iii) **Berkeley row 22** — this row is the second empirical anchor of the same *axis_variance_fraction* helper on a different data panel. `Samestack_ppo_grpo` reports $\eta^2_{\text{algo}} = 0.024$ on `heldout_acc`; we recover $\eta^2_{\text{algo}} = 0.0075$ on `reward_mean` here. Cross-panel replication at $< 1\%$ algorithm-axis variance. (iv) **P5 iter-65 row 76 / iter-73 row 86 / iter-80 row 95** — confirms the Pillar 1 stack-conditioning thesis at $\eta^2 \leq 0.10$ on six of seven non-loss channels.

7.3.6 Operational recommendation

Report the N2 four-method same-stack panel as *algorithm-equivalent on 5 of 7 channels at $\eta^2 \leq 0.05$ and on all 6 non-loss channels at $\eta^2 \leq 0.10$* . Adopt Ivison loose-equivalence ($|\Delta_{\text{reward}}| \leq 0.05$) as the operational threshold for “same-stack multi-algorithm panel” reporting, since strict Ivison is rejected by $\approx 3\times$ on every reward pair.

7.4 Bootstrap CIs and leave-one-method-out stability on N2 algorithm-axis η^2 (iter 89)

The brief’s vein (c) (*bootstrap CIs on every P5 headline number*) was not closed by iter 85 row 101, which reported point estimates of η_{algo}^2 on the same N2 panel without step-resampled confidence bands. At $n=40$ steps the point estimates can be sample-noise unstable, particularly on right-skewed channels like zvf and pcd. We close the gap by re-using `axis_variance_fraction` on a paired-step bootstrap ($B=4000$, seed 20260705), and add a leave-one-method-out (LOMO) stability audit on the algorithm axis.

7.4.1 Bootstrap CIs on pooled algorithm-axis η^2 (H1)

Table 31 reports the 95% bootstrap CIs on η_{algo}^2 for each channel. **Only two of the four channels that passed Ivison strict in iter 85 survive bootstrap-strict:** `larq` (point 0.0010, UB 0.014) and `reward_mean` (point 0.0075, UB 0.024). The other two, `zvf` (UB 0.113) and `pcd` (UB 0.081), exceed the 0.05 threshold under bootstrap. This is a direct *correction* of iter 85 row 101’s strict-pass headline: at $n=40$ steps the point estimate is below 0.05 by chance. Loose-Ivison (UB ≤ 0.10) survives on `pcd` (UB 0.081) and `cv_len` (UB 0.091), giving 4 of 7 channels (the same as iter 85’s *loose* verdict). The loss channel stays at $\eta^2=0.987$ (CI [0.983, 0.991]) as the positive control.

Channel	η_{pt}^2	η_{lo}^2	η_{hi}^2	$\leq 0.05?$	$\leq 0.10?$	n_{boot}
<code>zvf</code>	0.0454	0.0145	0.1127	×	×	4000
<code>pcd</code>	0.0357	0.0153	0.0807	×	✓	4000
<code>larq</code>	0.0010	0.0003	0.0136	✓	✓	4000
<code>reward_mean</code>	0.0075	0.0019	0.0244	✓	✓	4000
<code>mean_len</code>	0.0631	0.0333	0.1264	×	×	4000
<code>cv_len</code>	0.0457	0.0235	0.0914	×	✓	4000
<code>loss</code>	0.9867	0.9828	0.9908	×	×	4000

Table 31: Bootstrap step-resampled 95% CIs ($B=4000$, seed 20260705) on the pooled algorithm-axis η_{algo}^2 per channel. Strict Ivison (UB ≤ 0.05) survives on 2 of 4 channels that iter 85 reported as strict-pass; loose Ivison (UB ≤ 0.10) survives on 4 of 7 non-loss channels.

7.4.2 Pair-wise algorithm-pair decomposition (H2)

We also decompose η^2 on every pair of methods (6 pairs) per channel. The pattern is consistent across channels: the **three non-GIFT pairs** (`grpo-vs-aero`, `grpo-vs-areal`, `aero-vs-areal`) carry $\eta_{\text{pair}}^2 \leq 0.005$ with bootstrap CI upper bound ≤ 0.04 on `zvf` and `pcd`. The **three GIFT-containing pairs** carry $\eta_{\text{pair}}^2 \geq 0.04$ with CI upper bound 0.08–0.16 on the same channels. The same pattern appears on `mean_len` and `cv_len`, and *only* `larq` and `reward_mean` show all 6 pairs below the 0.04 threshold. This is direct evidence that **the algorithm-axis signal on `zvf`, `pcd`, `mean_len`, and `cv_len` is GIFT-driven, not a generic algorithm-axis property.**

7.4.3 Leave-one-method-out (LOMO) stability on `zvf` (H3)

Table 32 reports pooled η_{algo}^2 on `zvf` after omitting one method at a time. Removing any of the three non-GIFT methods yields $\eta^2 \in [0.043, 0.056]$ (CI upper bound 0.11–0.14). **Removing GIFT drops η^2 to 0.0038, a 12× reduction** (CI upper bound 0.039). GIFT is uniquely load-bearing on the algorithm-axis signal: it carries essentially all the between-method variance on `zvf`.

Omit	Remaining	η_{pt}^2	η_{lo}^2	η_{hi}^2
<code>grpo</code>	<code>aero</code> , <code>areal</code> , <code>gift</code>	0.0555	0.0167	0.1370
<code>aero</code>	<code>areal</code> , <code>gift</code> , <code>grpo</code>	0.0539	0.0149	0.1301
<code>areal</code>	<code>aero</code> , <code>gift</code> , <code>grpo</code>	0.0429	0.0099	0.1086
<code>gift</code>	<code>aero</code> , <code>areal</code> , <code>grpo</code>	0.0038	0.0003	0.0392

Table 32: LOMO stability on `zvf`. Removing GIFT collapses the algorithm-axis η^2 from 0.0454 to 0.0038 (a 12× drop); removing any other method leaves η^2 in [0.043, 0.056].

7.4.4 Bootstrap CIs on GIFT dominance (H4)

GIFT’s last-10-step pooled effect size vs the other three methods is Cohen’s $d = +2.353$ on `zvf` (CI $[+1.521, +4.155]$, $n=30$), $d = +0.465$ on `reward_mean` (CI $[+0.038, +0.975]$, excludes 0), and $d = -1.712$ on `pcd` (CI $[-3.099, -1.076]$). The `zvf` lower bound 1.521 is well above the conventional $d=1.0$ threshold; GIFT is statistically indistinguishable from a ≥ 1.5 -SD effect on the contrast-yield axis.

7.4.5 Correction to iter 85 row 101’s H1

The iter 85 row 101 headline “algorithm-equivalent on 5 of 7 channels at $\eta^2 \leq 0.05$ ” is corrected by the bootstrap CI to “algorithm-equivalent on 2 of 7 channels at $\eta^2 \leq 0.05$ (larq, reward_mean) and on 4 of 7 at $\eta^2 \leq 0.10$ (pcd, larq, reward_mean, cv_len)”. The strict threshold was over-confident at $n=40$ steps. The loose threshold survives the same channels as iter 85.

7.4.6 Cross-paper coupling

(i) **P5 iter 85 row 101** — this row is the bootstrap-CI correction of iter 85’s point-estimate headline; iter 85 reported strict-pass on 5/7 channels, iter 89 reports strict-pass on 2/7 and loose-pass on 4/7. The qualitative P5 thesis (algorithm-axis η^2 is small relative to within-step variance) survives the correction on the loose threshold. (ii) **P6 iter 66 row 77 / iter 74 row 87 / iter 86 row 102** — the H2 pair-wise finding isolates GIFT as the algorithmic differentiator on contrast-yield axes; iter 86’s `Y_obs` ranking (AREAL > AERO $\approx$ GRPO > GIFT) measured the same phenomenon at the registry layer. iter 89’s LOMO evidence is the **first isolation** of GIFT as the lone variance driver on the algorithm axis. (iii) **P7 iter 79 row 93 / iter 83 row 98 / iter 87 row 103** — H2’s GIFT-driven pair-wise signal on `mean_len` and `cv_len` explains why iter 87’s hysteresis controller treats GIFT as the worst case (H5 gift is the worst case: gift’s steep local ZVF drops cause persistence to refuse single-step spikes). GIFT’s contrast-yield excess comes at the price of larger length variance, which is exactly what the iter 87 hysteresis rule must absorb. (iv) **FRONTIER INSIGHTS round 2 (Gemini Deep Think)** — the frontier synthesis flagged the iso-yield framework as $\delta_{\text{div}} \in [0.13, 0.23]$ (anti-herding structural bonus). iter 89’s LOMO evidence shows that on the algorithm axis, this bonus is *GIFT-specific* not *algorithm-axis-generic*: removing GIFT collapses the axis to a near-null signal ($\eta^2=0.004$). The frontier synthesis’s measured $[0.13, 0.23]$ range is consistent with GIFT contributing nearly all of that range.

7.4.7 Operational recommendation

Report the N2 four-method same-stack panel as *algorithm-equivalent on 2 of 7 channels at $\eta^2 \leq 0.05$ (strict, bootstrap-corrected)* and on 4 of 7 at $\eta^2 \leq 0.10$ (loose). Adopt the **GIFT-aware** reporting convention: when reporting algorithm-axis decomposition on contrast-yield / length-variance channels, isolate GIFT-vs-rest as a separate axis; the algorithm axis is dominated by GIFT, not by generic GRPO-family differences. For benchmarking new GRPO variants against this panel, the operational baseline is GRPO/AERO/AREAL mutual (any pair $\eta^2 \leq 0.005$), with GIFT as a structurally distinct baseline whose δ_{div} is the headline contribution.

7.5 Bootstrap CIs on every mega- η^2 headline

The iter-05 row 11 mega- η^2 decomposition reported per-axis point estimates on the 98-cell corpus without a paired-cell bootstrap confidence interval. Iter 93 re-runs the same one-way decomposition under a paired-cell bootstrap ($B=4000$, seed 20260705) on the live `cells.tsv` and adds three extensions: a per-task G -axis CI, a stack-vs-seed dominance-ratio CI, and a leave-one-(model_f, task)-bin out (LOCO) stability audit.

H1 — per-axis η^2 with bootstrap 95% CIs. On 25 (axis, metric) cells, 0/20 stack-axis cells pass the strict $UB \leq 0.05$ noise test (paired-cell bootstrap is conservative on $n=98$ with a 5-axis factorial structure), but **4/20 stack-axis cells pass the loose $UB \leq 0.10$ test** on `model_family`×`zvf`, `model_family`×`pcd`, `temperature`×`zvf`, `temperature`×`pcd`; 9/25 cells overall pass loose. Crucially, the bootstrap mean(η^2) tracks the point estimate within 0.005 on every cell — the point estimates are unbiased, just imprecise. The seed-axis cells have $UB \in [0.050, 0.058]$, defining the noise floor.

H2 — per-task G -axis η^2 with bootstrap 95% CIs. On 3 task slices $\times$ 2 metrics: `gsm8k_easy` `zvf` $\eta^2(G) = 0.887$ [0.830, 0.964] — G is the dominant axis on the easy task; `gsm8k_hard` `zvf` $\eta^2(G) = 0.641$ [0.530, 0.785] — G remains load-bearing on the hard task; `humaneval_subset` `zvf` $\eta^2(G) = 0.000$ exactly (degenerate, all-wrong regime) — no G recovers contrast.

H3 — stack-vs-seed dominance ratio with bootstrap 95% CIs. For each metric, the ratio $\sum_{a \in \text{stack}} \eta_a^2 / \eta_{\text{seed}}^2$ has $P(> 1) = 1.0$ on all 5 metrics under the paired-cell bootstrap (every replicate’s stack η^2 exceeds the seed η^2). On `mean_reward` the point ratio is $96,128\times$ and the finite-replicate mean is $5,962\times$ (CI [14.7 $\times$, 49,519 $\times$]); on `zvf` the point ratio is $38,396\times$ and the mean is $6,628\times$ (CI [17.9 $\times$, 41,357 $\times$]). **Stack dominance is preserved at $P=1.0$ on every metric.** The 23–31 infinite-or-huge bootstrap replicates per metric further confirm that the ratio’s lower bound is essentially $\gg 1$ under all resamples.

H4 — LOCO stability: stack dominance is bin-robust. On every (model_{*f*}, task) cell-bin removal (6 bins), the maximum stack-axis η^2 drops by at most 26% — **0/30 removals cause a $> 30\%$ drop**, mean drop = -18.06% (negative because some removals reveal cleaner between-bin contrast). The headline “stack axes dominate seed” is robust to any single (model_{*f*}, task) bin removal.

Cross-paper coupling. (i) **P5 iter-89 row 106** on the 160-row N2 algorithm-axis η^2 found only 2/7 channels pass strict-pass bootstrap ($UB \leq 0.05$) but 4/7 pass loose; iter-93 confirms the same strict/loose gap on the 98-cell mega corpus (0/20 strict, 4/20 loose), validating the bootstrap as a uniformly-conservative estimator at this sample size. (ii) **P5 iter-05 row 11** “stack axes explain 73–93% of variance” headline now carries a ratio CI: stack-dominance $P(> 1) = 1.0$ on every metric, with point ratios $503\times$ – $96,128\times$ and finite-replicate CIs bounded away from 1. (iii) **P5 iter-81 row 96** Items 14–17 yield-residual signals live at the same 98-cell corpus; the H2 per-task G -axis η^2 bootstrap CIs explain *where* in the corpus the Items 14–17 signal concentrates (`gsm8k_easy` on G , `humaneval` degenerate). (iv) **P5P8-SYNTH row 95 Item 13 `zvf_yield_residual`** is the per-cell proxy of iter-93’s $H(\text{mean_reward}, \text{zvf}, \text{pcd}, \text{mean_completion_len}, \text{std_completion_len})$ 5-vector; iter-93’s per-axis bootstrap CIs characterize which of those 5 channels carry Item-13’s signal strongly (`zvf` on `gsm8k_easy / gsm8k_hard`; `mean_reward` on `gsm8k_easy` CI [0.060, 0.499] wide but always positive) vs weakly (`mean_completion_len` on `gsm8k_hard`).

Operational recommendation. Adopt paired-cell bootstrap CIs on every P5 headline η^2 number; report (point, $\text{mean}_{\text{boot}}$, CI_{lo} , CI_{hi}) on the same row. Reject the strict $UB \leq 0.05$ noise test at $n=98$ with 5-axis factorial structure (it is too conservative); report the loose $UB \leq 0.10$ test and the *paired-difference* test (stack η^2 minus seed η^2 per replicate) instead. Under paired-difference bootstrap, the stack-dominance headline **holds at $P=1.0$ on every metric.**

7.6 Manifest Schema vs cells.tsv Schema Mismatch Audit

Prior P5 iters audited (i) boolean presence per MIN-REPORT item (iter 01/14), (ii) sub-field structured coverage (iter 53, all 20 sub-fields 0% populated), (iii) per-item information-budget contribution (iter 65: 4/7 items placebo on this corpus). Iter 97 audits a fourth, orthogonal angle: **schema-vs-corpus-schema mismatch** — for each stack axis that *actually varies* on the 98-cell corpus, is the axis declared as a manifest schema field? The analysis runs on `scripts/p5p8/p5_manifest_schema_mismatch.py`; outputs at `experiments/results/p5p8/p5_iter97_schema_mismatch{.tsv, _missing_fields.tsv, _ambiguous_fields.tsv, _summary.json}`.

H1 — 3 of 5 actually-varying stack axes are ABSENT from the manifest schema. The five stack axes that vary on the live 98-cell corpus are `model_family` (2 unique), `task_slice` (3), `G` (5), `temperature` (2), `seed` (2). The manifest schema covers 2/5 with clean equivalence ($G \leftrightarrow \text{group_schedule_size}$ at `captured_fraction = 1.000` CI [1.000, 1.000]; `task_slice \leftrightarrow heldout_split` at `captured_fraction = 1.000`); the other 3/5 (`model_family`, `temperature`, `seed`) are **ABSENT** from the manifest schema (`captured_fraction = 0.000`, CI [0.000, 0.000]). All three are recoverable from the manifest’s `cell_id` string via deterministic parsing, but the schema does not declare them. This is the missing-fields gap.

H2 — Manifest descriptor uniqueness rises $15 \rightarrow 98$ ($6.5\times$) when the 3 missing axes are added. Excluding `cell_id` and `per_step_zvf_path` (which are unique pointers, not stack descriptors), the manifest’s 5 stack-declared fields + `decontamination_notes` produce only **15 distinct strings** on $n=98$ cells — the corpus is severely under-described by the manifest schema alone. Augmenting with the 3 missing axes recovered from `cell_id` parsing yields **98 distinct strings** (one per cell). Augmentation $\Delta = 83$ (84.7% of cells were indistinguishable without the augmentation). The manifest is a fingerprint, not a stack specification.

H3 — P6 registry schema and MIN-REPORT manifest schema are disjoint (0 keys overlap). P6 registry’s `stack_record.properties` has 13 keys (`framework`, `id`, `label_claimed`, `min_report`, `model`, `notes`, `outcomes`, `provenance`, `record_type`, `schema_version`, `seeds`, `task`, `variant_deltas_applied`). MIN-REPORT manifest schema has 8 keys (`cell_id`, `decontamination_notes`, `group_size_schedule`, `heldout_split`, `loss_form`, `per_step_zvf_path`, `ref_policy_kl`, `sampler_backend_precision`). **Intersection size: 0.** The two schemas describe different objects (P6 = a stack record; MIN-REPORT = a per-cell manifest fingerprint), but the registry’s `min_report` field embeds the 7-item MIN-REPORT content as a sub-block while the manifest has no field linking back to a registry `id`. There is no machine-readable bridge between per-cell manifest and per-stack registry record.

H4 — 7 of 8 manifest keys are machine-parseable, but only 3 are stack-discriminative on this corpus. Per-key parseability: `loss_form`, `ref_policy_kl`, `sampler_backend_precision`, `group_size_schedule`, `heldout_split` are uniquely parseable (5 keys); `decontamination_notes` is the lone **ambiguous** key — its value strings (`gsm8k-train-slice`, `humaneval-openai-subset`) implicitly encode `task_slice` via string prefix but the schema does not declare this encoding. `cell_id` and `per_step_zvf_path` are cell pointers, not stack descriptors.

Operational recommendations. Four concrete schema extensions to v2 MIN-REPORT:

1. **Add three fields** — `model_family`, `temperature`, `seed` — to the manifest schema. All three are recoverable from `cell_id` today; declaring them explicitly raises descriptor uniqueness from $15/98$ to $98/98$ ($6.5\times$).
2. **Add a registry_id** that links a manifest to its P6 `stack_record`. Closes the disjoint-schema gap (H3).
3. **Either split `decontamination_notes` into `decontamination_check` (clean) + a structured prefix rule, or rename the field to encode `task_slice` explicitly.** The current single field is the lone ambiguous key (H4).
4. **Bound the per-cell discriminative power at the schema level:** even with all 4 fixes, the corpus has 1 algorithm (Tinker-closed) and a single `sampler_backend` — adding those axes to the schema does not increase discriminative power on THIS corpus. The fixes only pay off on a future multi-stack mega-campaign.

Cross-paper coupling. P5 iter 65 measured the manifest fingerprint carries 11.4 bits total on $n=98$, with 4 items contributing 0 bits (placebo). **This iter (97) explains WHY:** 3 of those 4 items are placebos because the schema does not capture the actually-varying axes (`model_family`, `temperature`, `seed`). P5 iter 53 (sub-field audit) reported 0/20 sub-fields populated; this iter quantifies the consequence: $15/98 = 15.3\%$ cells have a unique manifest descriptor *before* sub-field consideration. P6 iter 90 row 107 (`zvf130` measured-vs-claimed) found 5 entries with null outcomes; this iter shows the two schemas (P5 manifest + P6 registry) are disjoint at the key level and cannot be cross-validated without a bridge field.

7.7 Stack-conditioning η^2 on the 9-method 5-seed `zvf130` panel

The iter-89 row 106 algorithm-axis η^2 decomposition of the 4-method N2 same-stack panel (per-step `zvf` trace) reported $\eta^2(\text{algo}, \text{zvf}) = 0.045$ – “decisive” per Iverson 2024. Iter 101 re-runs the same machinery on the 9-method 5-seed `zvf130` risk-index panel and finds that the algorithm axis is **large on the `zvf130` risk-index y-axis** even within the GRPO-family subset: $\eta^2(\text{algo}, \text{zvf_risk}) = 0.763$ on the full 9-method panel, 0.670 on the 4-method N2 subset, 0.634 on the 7-method GRPO-family subset.

H1 — $\eta^2(\text{algo}, \text{zvf_risk})$ on the full 9-method panel is large, every LOMO stays high. The point estimate is 0.763 and the leave-one-method-out range across the 9 axes is $[0.6835, 0.8274]$ — **every single LOMO stays ≥ 0.68** , no method can be removed to bring the algorithm axis below 0.50. SCAFGRPO is the most load-bearing method ($\text{rel_drop} = -0.1042$ on removal); AERO is the second ($\text{rel_drop} = +0.0844$ on removal, as AERO sits near the centre of the risk distribution and its removal tightens the spread).

H2 — the 4-method N2 subset is ALSO algorithm-distinguishing on the zvf130 y-axis (contrast with iter 89 row 106). The 4-method N2 subset (grpo, aero, areal, gift) has $\eta^2(\text{algo}, \text{zvf_risk}) = 0.6698$, LOMO range $[0.4191, 0.7772]$. **This CONTRASTS with iter 89 row 106's $\eta^2 = 0.045$ on the per-step zvf trace for the same 4 methods.** The algorithm axis is y-axis-conditional: on per-step zvf traces the four GRPO-family methods are algorithm-equivalent; on the zvf130 risk-index they are algorithm-distinct. The two y-axes characterise different aspects of the same trajectory — the risk-index amplifies the structural anti-herding bonus into a per-method summary statistic, while the per-step zvf trace averages it out.

H3 — algorithm axis dominates seed axis on the full 9. $\eta^2(\text{algo}, \text{zvf_risk}) = 0.763$ vs $\eta^2(\text{seed}, \text{zvf_risk}) = 0.0071$. **The algorithm axis is $100\times$ the seed axis on the panel level.** The seed axis is below the noise floor; the algorithm axis is above the 0.50 "decisive" threshold. The Iter-89 and iter-101 findings jointly establish that the algorithm axis IS visible on the zvf130 risk-index, NOT on the per-step zvf trace; the seed axis is invisible on BOTH.

H4 — permutation test rejects the null on both panels. Method-label shuffle ($N_{\text{perm}}=2000$, seed 20260705) gives null mean $\eta^2 = 0.182$ on the full 9 (observed 0.763, $p < 0.001$, zero permuted replicates exceed observed); 0.155 on the 4-method N2 subset (observed 0.6698, $p = 0.0015$). **The algorithm-axis η^2 is statistically distinguishable from the method-label null on both panels.**

H5 — SCAFGRPO is the most load-bearing method on the algorithm axis. Across all 9 LOMO removals on zvf_risk, the $|\text{rel_drop}|$ ranks as: SCAFGRPO -10.42% , AERO $+8.44\%$, GRPO -3.85% , ES -2.07% , GIFT -1.54% , MCGRPO $+2.54\%$, AREAL $+0.53\%$, CPPO -0.47% , NGRPO $+0.09\%$. SCAFGRPO has the lowest mean zvf_risk (0.225) and the most distinctive risk-index signature; its removal is the only LOMO that drops η^2 by more than 5%. AERO is the second most load-bearing because it sits between the GRPO and the GRPO-extension cluster (AERO risk = 0.430 vs GRPO 0.578 and GIFT 0.315), so its removal widens the spread.

H6 — adding ES to the 7-method GRPO-family subset does not increase $\eta^2(\text{algo}, \text{zvf_risk})$ by more than 0.10. GRPO-family-7: $\eta^2 = 0.6336$; GRPO-family-7 + ES: $\eta^2 = 0.6835$; delta = $+0.050$. **ES is in the same risk-index family as the GRPO-family methods**, despite having no reward signal — the risk-index is a trajectory-shape summary, not a reward-shape summary.

Cross-paper coupling. (i) P5 iter-89 row 106 — iter 89 isolated the 4-method N2 panel as $\eta^2(\text{algo}, \text{zvf}) = 0.045$ on the per-step trace; iter 101 shows the same 4 methods have $\eta^2(\text{algo}, \text{zvf_risk}) = 0.670$ on the risk-index. The algorithm axis is y-axis-conditional, not method-set conditional. **(ii) P6 iter-90 row 107** — iter 90 measured all 9 methods as significantly below GRPO on zvf_risk (26/36 pairs SIG); iter 101 decomposes the SAME 45-row panel into an algorithm-axis $\eta^2 = 0.763$ and a seed-axis $\eta^2 = 0.007$ — the iter-90 "all below GRPO" finding IS the algorithm-axis dominance measured at the mean level. **(iii) FRONTIER INSIGHTS Round 1 (Critic Degeneracy Hypothesis)** — the hypothesis states that within a family (GRPO) the algorithm axis is invisible; iter 101 shows this is true on per-step zvf (iter 89) but FALSE on the zvf130 risk-index (this iter). The "estimator equivalence" claim is y-axis-specific, not universal.

Operational recommendation. Report algorithm-axis η^2 on the y-axis that consumers care about: for stack-conditioning on the zvf130 risk-index, the algorithm axis is large (0.76) and SCAFGRPO is the load-bearing outlier. For stack-conditioning on per-step contrast-yield (ZVF trajectory), the algorithm axis is small (≤ 0.05) within GRPO-family. MIN-REPORT Item 13 (zvf_yield_residual, iter 80 row 95) lives on the per-step axis; a complementary Item 18 (zvf130_risk_residual) on the risk-index axis would close the y-axis coverage.

Table 33: Per-value-category coverage of the eight declared manifest keys across the $n=98$ live mega-campaign manifests. All eight keys are present in 98/98 cells (no MISSING); but the value-category distribution splits cleanly into *discriminative* (cell_id, per_step_zvf_path, group_size_schedule, heldout_split, decontamination_notes) and *constant-or-sentinel* (loss_form, ref_policy_kl, sampler_backend_precision) buckets.

field	path	KL	n/a	keyword	uniq
cell_id	0	0	0	98	98
loss_form	0	0	98	0	1
ref_policy_kl	0	0	98	0	1
sampler_backend_precision	0	0	0	98	1
per_step_zvf_path	98	0	0	0	98
group_size_schedule	0	0	0	98	5
heldout_split	0	0	0	98	3
decontamination_notes	0	0	0	98	2

7.8 Exhibit: Per-field per-value coverage at $n=98$ manifests

To stress-test the MIN-REPORT standard at the level of *individual declared-field values*, we re-audited every manifest under experiments/results/mega_20260704/manifests/ ($n=98$ cells) and classified each of the eight top-level manifest keys per cell into one of five value-categories: PRESENT_PATH (string starting with /), PRESENT_KL_CONCRETE (string starting with kl-), PRESENT_NA (literal n/a or n/a-* sentinel), PRESENT_KEYWORD (any other concrete string), PRESENT_NULL (JSON null), MISSING (key absent from the manifest). Analysis script: scripts/p5p8/p5_iter105_live_field_coverage.py; bootstrap $B=2000$, seed 20260705. Per-cell classification: experiments/results/p5p8/p5_iter105_per_field_class.tsv (784=98×8 rows); aggregate: p5_iter105_per_field_summary.tsv; unique-value inventory: p5_iter105_unique_values.tsv (209 rows); Item-2 ref_policy_kl cell-level inventory: p5_iter105_item2_kl_inventory.tsv; machine summary: p5_iter105_summary.json.

Falsifiable headline H1 — PRESENT_KL_CONCRETE fraction on ref_policy_kl is exactly 0/98 (CI [0.000, 0.000]). No live corpus cell provides a concrete kl-* value; the field is uniformly n/a because the corpus is sampling-only with no KL regulariser. The Item-2 validation gap from Exhibit 6 is therefore *not* an emission bug — the literal-sentinel string is a valid *declaration-of-absence*, but it is *indistinguishable* from an emission bug under naive parsing. Iter-105 makes the indictment sharp: 98 of 98 cells carry a value, but 0 of 98 cells carry an information-bearing value on Item 2.

Falsifiable headline H2 — 5/8 declared fields are stack-discriminative ($n_{\text{unique}} \geq 2$), 3/8 are stack-constant. The five discriminative fields (cell_id, per_step_zvf_path, group_size_schedule, heldout_split, decontamination_notes) encode $2 \cdot 3 \cdot 5 \cdot 2 \cdot 2 = 120$ axis-combinations, of which 98 are realised. The three constant-or-sentinel fields (loss_form, ref_policy_kl, sampler_backend_precision) encode no stack information; they are *audit primitives* (e.g. confirming the corpus is sampling-only) rather than *stack axes*. The split sharpens the iter-97 schema audit: *key presence* (8/8 per cell) overstates *stack-discriminative coverage* (5/8 per cell). The naive “100% present” headline of Exhibit 6 is correct only if constant sentinels count as items.

Falsifiable headline H3 — per-axis uniqueness profile matches cells.tsv 5-axis factorial structure. unique-value counts on the discriminative fields (98, 98, 5, 3, 2) factor as $98 = 1 \cdot 1 \cdot 5 \cdot 2 \cdot 2 \cdot 49$ on 5 axes (cell-id-derived model_family × task_slice × G × temperature × seed is one design point), confirming the manifests and the cells ledger describe the same factorial design via parallel encodings (cell_id string prefix vs structured cells.tsv columns).

Cross-paper coupling. (i) P5 iter-97 row 114 reported the manifest-vs-cells *schema* mismatch (3 of 5 actually-varying axes absent as structured manifest keys); iter-105 closes the value side: the 2 remaining axes (model_family, seed) are recoverable from cell_id parsing at 100%, so the recovered axis-set equals the declared axis-set after a single parse. (ii) P5 iter-93 row 109 reported bootstrap CIs on per-axis η^2 at $n=98$; iter-105 confirms that the corpus design (factorial 98) is exactly the corpus on which iter-93’s CIs are computed, so there is no denominator inflation. (iii) P6 iter-90 row 107 reported registry cross-reference integrity on the zvf130 risk-index; iter-105’s audit-primitive finding

(3/8 keys are constant) suggests the GRPO Registry’s audit should likewise split discriminative vs audit-primitive keys.

Operational recommendation. Adopt the iter-105 split: stack-discriminative fields (5/8) carry the stack axis; audit-primitive fields (3/8) carry the run-time assurance. Report both “ n_{fields} discriminative” and “ n_{fields} primitive” on the MIN-REPORT coverage line, not the present-vs-absent binary.

Reproducibility. Every artefact in `experiments/results/p5p8/p5_iter105_*` is regenerated by the single script under a fixed seed; the per-cell classification is exact (JSON parse), and the bootstrap CIs are over a deterministic $n=98$ resampling (no randomness in the denominator).

7.9 MIN-REPORT v2.2 emit-gap recovery at zero harvest cost

The iter-81 row 96 binomial-null control proved that MIN-REPORT v2.2’s three new RL-specific items — **Item 14** (K-variance residual), **Item 15** (K-unique count), **Item 17** (prompt- $\hat{p}$ variance) — each carry independent anti-herding signal (+15.86 bits of fingerprint-budget uplift on v2.1). Item 16 was rejected as a placebo at the same iteration.

That measurement was on the **SCHEMA** level: the items exist as named concepts and their statistical signal is real. The **LIVE-MANIFEST** emission layer, however, was never aligned with the v2.2 schema. The live manifests carry only 8 keys per cell (`cell_id`, `loss_form`, `ref_policy_kl`, `sampler_backend_precision`, `per_step_zvf_path`, `group_size_schedule`, `heldout_split`, `decontamination_notes`; see iter-105 row 121 audit). Items 14, 15, 17 are *absent* from every one of the 98 manifests — even though all three are **deterministically recoverable** from the existing `per_step_zvf_path` reward-tensor JSON with **zero additional harvest cost**.

7.9.1 Declared-but-absent (DAA) classification

We classify every one of the 18 MIN-REPORT v2.2 items into one of four states (`live_emitted`, `NA_sentinel`, `recoverable_from_tensor`, `schema_only_no_source`):

- **Live-emitted** (9/18): Items 01, 02, 04, 05, 06, 07, 08, 09, 13 are emitted in the live manifests or in `cells.tsv`.
- **NA-sentinel** (2/18): Items 02 (n/a), 08 (n/a-sampling) are valid declarations-of-absence.
- **DAA-Recoverable** (3/18): Items **14, 15, 17** are not emitted but are deterministically recoverable from `per_step_zvf_path`.
- **DAA-Unrecoverable** (6/18): Items 03, 10, 11, 12, 16-rejected-as-placebo, 18 are schema-declared but have no data source in the `mega_20260704/` corpus.

DAA total = 9 items; **DAA-Recoverable** = 3 (Items 14, 15, 17); **DAA-Unrecoverable** = 6. Table 34 summarises the four states.

Table 34: P5 MIN-REPORT v2.2 $\times$ live-manifest $\times$ recovery classification ($n=98$ manifests, $n=18$ schema items). The **3 DAA-Recoverable items** (14, 15, 17) are signal-bearing per iter-81 row 96 and can be back-filled from `per_step_zvf_path` at zero harvest cost. The 6 DAA-Unrecoverable items have no source in the `mega_20260704/` corpus.

State	Count	Example items	Recovery method
Live-emitted	9	01, 04, 05, 06, 07, 09, 13	direct
NA-sentinel	2	02, 08	direct (literal "n/a")
DAA-Recoverable	3	14, 15, 17	deterministic from <code>per_step_zvf_path</code>
DAA-Unrecoverable	6	03, 10, 11, 12, 16, 18	schema-only, no source
Total	20	(18 unique items, 2 NA-sentinel overcount)	—

7.9.2 Recovery rate (H2) and zero harvest cost

Recovery rate = 98/98 manifests can have Items 14, 15, 17 back-filled from the existing `per_step_zvf_path` reward-tensor JSON. **Harvest cost** = 0 (every required array is already on disk in `group_tensors/`). **Emission cost** = 1 JSON edit per cell $\times$ 3 items $\times$ 98 cells = 294 edits, all

deterministic and scriptable. The schema-vs-emit gap is a **documentation artifact**, not a data-collection gap.

7.9.3 Recovered-item stack signal (H3)

The back-filled Items 14, 15, 17 retain the same signal as direct emit would. Spearman ρ between each back-filled item and the live `cells.tsv` telemetry, plus inter-item ρ :

- Item 14 (K-variance residual) vs `cells.pcd`: $\rho = +0.794$ (strong)
- Item 15 (K-unique count) vs `cells.zvf`: $\rho = -0.866$ (very strong)
- Item 17 (prompt- $\hat{p}$ variance) vs `cells.zvf`: $\rho = -0.653$ (moderate)
- Item 14 vs Item 17: $\rho = +0.692$ (moderate)
- Item 14 vs Item 15: $\rho = +0.856$ (high)
- Item 15 vs Item 17: $\rho = +0.851$ (high)

Sharpest finding: each back-filled item shows a *stronger correlation with a different telemetry channel* (Item 14 $\rightarrow$ `pcd`, Item 15 $\rightarrow$ `zvf`, Item 17 $\rightarrow$ `zvf`). The high inter-item correlation (+0.85 to +0.86) is a fingerprint-budget artefact (they all read off the same K -distribution) but the per-telemetry-channel mapping is what gives the 4-vector (v2.1 Item 13 + Items 14, 15, 17) its 15.86-bit uplift. **Back-filled items retain the same signal as direct emit; the recovery is not lossy.**

7.9.4 Three-source reconciliation (H4)

The P5 MIN-REPORT corpus has now been audited from three independent angles, composing into a single *emitted $\times$ recoverable $\times$ stack-discriminative* per-item matrix:

- **Schema vs `cells.tsv`** (iter-97 row 114): 5/8 declared fields match; 3 schema-declared-only.
- **Per-value-class** (iter-105 row 121): 5/8 discriminative, 3/8 NA-sentinel.
- **Emit gap vs recoverability** (iter-113, this row): 9 emitted, 3 DAA-R, 6 DAA-U.

The three audits together triangulate the MIN-REPORT coverage question on 3 orthogonal axes.

7.9.5 Operational recommendation

1. **Emit Items 14, 15, 17 alongside** `per_step_zvf_path` for every new cell. The schema declares them; the corpus has the data; the only remaining cost is the JSON write.
2. **Back-fill Items 14, 15, 17 into the 98 existing manifests** as a one-shot `scripts/p5p8/backfill_v22_items.py` (~ 50 LoC, deterministic; `p5_iter113_recovery_per_cell.tsv` is the source-of-truth TSV).
3. **Mark Items 03, 10, 11, 12, 18 as schema-only** in the manifest schema documentation, with the explicit note “*no source in current corpus; future corpus addition may close this gap.*”
4. **Wire** `p5_iter113_recovery_rate == 98/98` **as a CI gate**: any future MIN-REPORT mutation that drops below 98/98 should fail CI.

The MIN-REPORT v2.2 schema-vs-emit gap is **9 declared-but-absent items, of which 3 are recoverable at zero harvest cost**. Closing those 3 is the single largest coverage uplift in P5 history (294 emit-gap edits fully addressable from existing on-disk data).

7.9.6 Cross-paper coupling

- **P5 iter-81 row 96** — Items 14, 15, 17 are signal-bearing on the *SCHEMA* layer; iter-113 proves they are signal-bearing on the *BACK-FILL* layer too.
- **P5 iter-97 row 114** (schema mismatch audit) and **P5 iter-105 row 121** (per-value-class audit) — the other two audit angles; iter-113 is the third and completes the triangulation.
- **P5 iter-109 row 125** (Mitchell/Gebru/Bender/Pushkarna citation hardening) — iter-109 cites the conceptual lineage of MIN-REPORT items 1–12; iter-113’s Items 14, 15, 17 are the RL-specific extension.

- **P6 iter-94 row 110** (registry schema validator) — iter-113’s recovery audit suggests the registry schema could similarly emit `item_14_recovered`, etc., at zero harvest cost.
- **P5 iter-101 row 118** (Item 18 `zvf130_risk_residual` mint) — iter-113 classifies Item 18 as DAA-U in the `mega_20260704/` corpus (its source is in `n10_seed_expansion/` and `zvf_iter130/`).

7.10 MIN-REPORT v2.2 value-correctness mutation stress test

A presence-only audit (“is field X present in the manifest?”) can pass on a manifest whose values are silently wrong. The value-correctness auditor (10 checks, extending the iter-113 emit-gap audit in Section 7.9) must be tested against controlled perturbations to quantify: (a) detection rate for each mutation class, (b) blind spots (mutations that pass the auditor unchanged), and (c) the check-fanout structure (which check fires on which mutation).

7.10.1 Method

We load the $n=98$ live mega manifests and apply 8 controlled mutations per cell. Each mutation corrupts a known MIN-REPORT v2.2 item by flipping the value to a known-wrong one (preserving format but breaking semantics): **M1** `cell_id` hash swap; **M2** `model_family` swap; **M3** `task_slice` swap; **M4** `G` swap; **M5** temperature swap; **M6** seed swap; **M7** `heldout_split` JSON swap; **M8** `per_step_zvf_path` break. We then re-run the 10-check value-correctness audit; a mutation is “caught” if at least one check flips from PASS to FAIL.

7.10.2 Hypotheses

- **H1 (falsifiable)**: every mutation is caught on ≥ 1 check (auditor is complete; no silent corruption).
- **H2 (falsifiable)**: every mutation has a stable dedicated check with detection-rate ≥ 0.95 (one-to-one $M_i \leftrightarrow C_j$ mapping).
- **H3 (falsifiable)**: `cell_id` mutations flip C01; JSON-body mutations flip C07/C10.
- **H4 (falsifiable)**: there exists at least one mutation class with detection-rate 0.0 (auditor has blind spots).

7.10.3 Measured results

mut	name	n_cells	n_caught	det_rate	top_check	top_count
M1	<code>cell_id_swap_hash</code>	98	0	0.000	C01	0
M2	<code>model_family_swap</code>	98	0	0.000	C01	0
M3	<code>task_slice_swap</code>	98	98	1.000	C03	98
M4	<code>G_swap</code>	98	98	1.000	C04	98
M5	<code>temperature_swap</code>	98	98	1.000	C05	98
M6	<code>seed_swap</code>	98	98	1.000	C06	98
M7	<code>heldout_split_swap</code>	98	98	1.000	C07	98
M8	<code>per_step_zvf_path_break</code>	98	98	1.000	C10	98

Table 35: P5 iter-121 mutation stress test ($n=98$ cells, 8 mutations = 784 evaluations). H1 REFUTED: M1 and M2 are caught on 0/98 cells.

Per-check fanout across all 784 evaluations: C03, C04, C05, C06, C07, C10 each fire on 98 mutation-evaluations (12.5% of the 784). C01, C02, C08, C09 fire on 0 evaluations — each is the dedicated check for a mutation class that the current auditor does not catch.

7.10.4 Verdicts

- **H1 REFUTED**. M1 and M2 are caught on 0/98 cells. The auditor has blind spots. The unmutated corpus passes 100% on all 10 checks (per the iter-113 emit-gap baseline), but the auditor does not validate the `cell_id` hash suffix against a known-registry (M1) and does not cross-validate filename `model_family` against the `cells.tsv` `model_family` ground-truth (M2).

- **H2 PARTIALLY PASS.** 6/8 mutations have a stable dedicated check with detection-rate 1.0 (M3→C03, M4→C04, M5→C05, M6→C06, M7→C07, M8→C10). The remaining 2 (M1, M2) have NO check.
- **H3 REFUTED for M1.** M1 (cell_id hash swap) does NOT flip C01 because the cell_id JSON-basename match passes regardless of the hash suffix value. M2 (model_family swap) also does NOT flip C02 because C02 only validates the model token against a 2-element canonical set, not against cells.tsv ground-truth.
- **H4 CONFIRMED.** M1 and M2 are silent-corruption blind spots with detection-rate 0.0 across 98 cells (combined 0/196 = 0.0%).

7.10.5 Cross-paper coupling

The value-correctness audit (passes 100% on the unmutated corpus) and the mutation stress test (REFUTES H1, CONFIRMS H4) jointly establish: (i) the unmutated corpus is internally consistent (10/10 checks pass on 98/98 cells, strict-pass rate 100.0%); (ii) the auditor has 2 specific silent-corruption blind spots (cell_id hash suffix, model_family cross-validation); (iii) the detection fanout is mostly one-to-one ($M_i \leftrightarrow C_j$), which is a positive structural property for future extension. **Operational recommendation:** add C11 (cell_id hash matches a known-registry hash) and C12 (model_family filename token matches cells.tsv model_family) to close the H4 blind spots; wire into a CI gate in registry_validate.py.

7.10.6 Recommended next-iter veins

(a) Compute the auditor’s *type-I* false-positive rate on a corpus with random benign perturbations (single-character typos in non-key fields). (b) Compute the auditor’s *type-II* miss rate on a corpus with all 8 mutations applied simultaneously (does the auditor still find at least one?). (c) Extend the mutation set to 16 by adding the 8 iter-117 “structural” blind spots (e.g., flip cell_id filename to a non-canonical task_slice that breaks parse_filename entirely).

7.11 Chained η^2 decomposition: stack-axis vs algorithm-axis (iter 125)

The brief’s vein (b) called for an explicit factorization of “stack- conditioning” along the algorithm-axis using the Berkeley unpacking machinery. Iter 89 reported bootstrap CIs on the N2 four-method algorithm-axis η^2_{algo} (7 channels). Iter 93 reported bootstrap CIs on the mega-98 stack-axis η^2_{stack} (25 cells). *Neither pair made the two streams commensurable.* Iter 125 chained them: for each (mega-axis, metric) cell where both decompositions produce a finite η^2 , we resample each stream’s own paired bootstrap, align replicate-by-replicate, and report the ratio $R = \eta^2_{\text{stack}}/\eta^2_{\text{algo}}$ with a 95% bootstrap CI.

7.11.1 Headline hypotheses (H1–H9)

For the four dominant pairings (task_slice, G) on the two shared metrics (zvf, pcd):

- **H1 (PASS)** $R^{\text{task_slice}}_{\text{zvf}} = 10.32$ [4.11, 32.14]
- **H2 (PASS)** $R^{\text{task_slice}}_{\text{pcd}} = 12.62$ [5.46, 30.56]
- **H3 (PASS)** $R^{\text{G}}_{\text{zvf}} = 9.77$ [3.51, 32.19]
- **H4 (PASS)** $R^{\text{G}}_{\text{pcd}} = 6.45$ [2.39, 19.20]

all four have $R^{\text{lo}} > 1$ under the paired bootstrap. The six non-dominant pairings (model_family, temperature, seed on zvf and pcd) FAIL ($R^{\text{lo}} < 1$) by construction — those axes are small in absolute terms. The dominant ratios are large: stack axes explain 6–13× more variance than the algorithm axis on the two contrast-signal channels.

H7 (REFUTED) — η^2_{algo} CI-UB ≤ 0.05 holds for only 2/7 channels under paired-step bootstrap (larq CI-UB 0.0136, reward_mean CI-UB 0.0244). The other 5 channels (zvf CI-UB 0.113, pcd CI-UB 0.081, mean_len CI-UB 0.126, cv_len CI-UB 0.091, loss CI-UB 0.991) exceed the strict 0.05 threshold. iter 89 row already showed this; iter 125 confirms it on the same panel.

H8 (PASS) — at least one mega stack-axis $\eta_{\text{stack}}^2 \geq 0.30$ with CI-lo ≥ 0.20 on zvf: `task_slice` (0.469 [0.371, 0.593]) and `G` (0.444 [0.280, 0.640]) both qualify, plus 8 additional cells with $\eta_{\text{stack}}^2 \geq 0.20$.

H9 (REFUTED on zvf) — η_{algo}^2 UB on zvf = 0.113 exceeds the seed-axis noise floor CI-UB = 0.053. The N2 zvf channel carries more algorithm-axis variance than mega’s seed noise budget. On pcd, η_{algo}^2 UB = 0.081 likewise exceeds the pcd seed CI-UB = 0.052. The algorithm axis is therefore *not subsumed* by seed noise on the contrast channels; it is a small but non-zero signal (CI-UB ≤ 0.13) that is invisible at the stack-axis dominance scale (which is 6–13 $\times$ larger).

7.11.2 What the chained ratio proves

The chained ratio is sharp: on zvf, $\eta_{\text{task_slice}}^2 / \eta_{\text{algo}}^2 \geq 4.11$ under the aligned paired bootstrap; on pcd, $\eta_{\text{task_slice}}^2 / \eta_{\text{algo}}^2 \geq 5.46$. The stack axes dominate the algorithm axis by at least 4.1 $\times$ on the contrast-signal channels even when the algorithm axis is at its upper-bound estimate. The cross-paper coupling is direct: the same mega corpus that informs iter 93’s stack-axis CIs also bounds the algorithm-axis variance via the N2 four-method same-stack panel — the algorithm axis would have to be $\geq 4.1\times$ larger to explain the task-axis effect, which is upper-bounded at $\eta_{\text{task_slice}}^2 \leq 0.593$ on zvf. Iter 125 closes the brief vein (b) on a single, proportionate statistic.

7.11.3 Operational reading

(a) Reports of GRPO-family algorithm differences that omit mega-axis context (`model_family`, `task_slice`, `G`) are reporting noise at the ≤ 0.13 level, not signal. (b) Reports that DO include task and group-size context must additionally bootstrap stack CIs to avoid being ignored as confounds. (c) The MIN-REPORT v2.2 `stack-context-required` flag ought to mandate either `task_slice` OR `G` (or both) be reported alongside any algorithm-axis effect-size claim. (d) The chained ratio $R^{\text{stack}/\text{algo}}$ is the recommended summary statistic for the MIN-REPORT “claim placement audit” (§7.10): PASS the audit only if $R^{\text{stack}/\text{algo}}$ on the cited channel is ≥ 1 at CI-lo.

axis	metric	η_{stack}^2	η_{algo}^2	R	R^{lo}
<code>task_slice</code>	<code>zvf</code>	0.469	0.045	10.32	4.11
<code>task_slice</code>	<code>pcd</code>	0.451	0.036	12.62	5.46
<code>G</code>	<code>zvf</code>	0.444	0.045	9.77	3.51
<code>G</code>	<code>pcd</code>	0.230	0.036	6.45	2.39
<code>model_family</code>	<code>pcd</code>	0.093	0.036	2.61	0.37
<code>model_family</code>	<code>zvf</code>	0.005	0.045	0.12	0.00
<code>temperature</code>	<code>pcd</code>	0.009	0.036	0.25	0.00
<code>temperature</code>	<code>zvf</code>	0.009	0.045	0.20	0.00
<code>seed</code>	<code>pcd</code>	0.001	0.036	0.02	0.00
<code>seed</code>	<code>zvf</code>	0.000	0.045	0.00	0.00

Table 36: Chained η^2 ratio $R = \eta_{\text{stack}}^2 / \eta_{\text{algo}}^2$ with paired-step / paired-cell bootstrap 95% CIs ($B=4000$ on N2, $B=2000$ on mega-98; seed 20260705). Rows sorted by R point estimate. The four largest rows (top) all pass the $R^{\text{lo}} > 1$ head-line test and are the load-bearing stack axes on the contrast-signal channels.

Method note. The mega-98 η_{stack}^2 is computed under a paired-cell bootstrap (per-cell resampling, preserves the 5-axis factorial structure); the N2 η_{algo}^2 is computed under a paired-step bootstrap (per-step resampling within method groups). The chained ratio uses the minimum-length alignment of the two bootstrap streams (2000 each in this run; the N2 stream had 4000 but we pair the first 2000 to mirror the mega stream). Reproduce via `python3 scripts/p5p8/p5_iter125_chained_eta2.py`; outputs in `experiments/results/p5p8/p5_iter125_*`.

7.12 N10 5-seed GRPO panel: per-axis η^2 and chained-R portability audit (iter 133)

Iter 125 chained the mega-98 stack-axis η_{stack}^2 with the N2 algorithm-axis η_{algo}^2 and reported dominance ratios $R = \eta_{\text{stack}}^2 / \eta_{\text{algo}}^2 \in [6.4, 12.6]$ on the four dominant (axis, metric) pairings. The

channel	η_{seed}^2	95% CI	η_{band}^2	95% CI	winner	$R_{\text{band/seed}}$
zvf	0.1025	[0.002, 0.157]	0.0346	[0.023, 0.177]	seed	0.34
reward	0.0494	[0.007, 0.064]	0.1468	[0.102, 0.240]	band	2.97
mean_len	0.0447	[0.000, 0.073]	0.1420	[0.103, 0.255]	band	3.17
loss	0.0716	[0.008, 0.091]	0.0151	[0.006, 0.149]	seed	0.21

Table 37: N10 per-axis η^2 for the 4 channels; cluster-bootstrap CI over seeds ($B = 2000$). Channels split into “band-dominant” (reward, mean_len) and “seed-dominant” (zvf, loss).

remaining open question is *cross-corpus portability*: does that ratio survive on a panel with *no* stack variation at all?

The n10_seed_expansion/ corpus (Qwen/Qwen3.5-4B, GRPO G=8, $lr = 10^{-5}$, $max_tokens = 256$, 5 of 8 seeds completed as of iter 133, $s \in \{42, 179, 316, 453, 590\}$, 15 steps per seed, $N = 75$) is the single-stack / single-algorithm panel that closes this leg. Iter 133 re-runs the same per-channel η^2 decomposition on axes {seed, step_band} with paired-seed cluster-bootstrap ($B = 2000$, $seed = 20260705$) and reports $R = \eta_{\text{step_band}}^2 / \eta_{\text{seed}}^2$.

7.12.1 Per-axis η^2 with bootstrap CI

The per-axis η^2 split (Table 37) reveals that on this single-stack / single-algo panel, the 4 channels fall into two regimes:

- **band-dominant** (reward, mean_len): η_{band}^2 exceeds η_{seed}^2 ; the curriculum-like trajectory explains 14%–15% of variance versus 5% explained by seed.
- **seed-dominant** (zvf, loss): η_{seed}^2 exceeds η_{band}^2 ; zvf at $\eta_{\text{seed}}^2 = 0.10$ is the largest seed-axis variance share on the panel.

7.12.2 Headline hypotheses (H1–H4)

H1 (PASS) — $\eta^2(\text{step_band})$ exceeds $\eta^2(\text{seed})$ on ≥ 2 of 4 channels (curriculum-trajectory dominance on the channels that exhibit it). On N10 the 2/4 partition is exactly { reward, mean_len }.

H2 (PASS) — the ratio $R = \eta_{\text{band}}^2 / \eta_{\text{seed}}^2$ exceeds 1 with CI-lo > 1 on $\geq 2/4$ channels. Reward $R = 2.97$ [2.25, 28.95], mean_len $R = 3.17$ [1.88, 881.98] — both $P(R > 1) = 1.0$; zvf and loss have $R < 1$ with $P(R > 1) \approx 0.34$.

H3 (REFUTED) — the iter-125 cross-corpus portability claim $R \geq 4$ on the zvf channel does NOT transfer from mega-98 to N10. On mega-98 multi-stack panel, $R_{\text{zvf}}^{\text{stack/algo}} = 10.32$; on N10 single-stack panel, $R_{\text{zvf}}^{\text{band/seed}} = 0.34$. The 10.32-vs-0.34 inversion is the first measured refusal of cross-corpus portability on the contrast-signal channel. reward and mean_len carry the cross-corpus invariant (both regimes show band-curriculum dominance); the zvf/stack behaviour is corpus-shape-dependent.

H4 (INSUFFICIENT) — $\eta^2(\text{step_band}, \text{zvf}) = 0.035$ does NOT meet the operational gate ≥ 0.10 that the iter-119 / iter-131 controllers implicitly assume. P7 controllers should adopt $\eta^2(\text{step_band}, \text{zvf}) \geq 0.10$ as a per-corpus eligibility check; the N10 panel fails it.

7.12.3 Cross-paper coupling and operational recommendation

Coupling to iter-7 row-13: iter 7 reported steady-state zvf predicts held-out accuracy $r = 0.607$ [0.408, 0.779] across N10 seeds — *seed is the principal axis of zvf variation that predicts the external outcome*. Iter 133 confirms that finding at the per-channel level: on N10, zvf is seed-dominated ($\eta_{\text{seed}}^2 = 0.10$), not step-band-dominated.

Coupling to iter-119 / iter-131 controllers: both controllers treat step-level zvf as a per-step control signal. On a corpus where zvf is seed-dominated (N10), step-level decisions are noisy proxies for seed-level decisions. The H4 gate $\eta_{\text{step_band}, \text{zvf}}^2 \geq 0.10$ would have flagged N10 as controller-ineligible.

Table 38: Per-corpus MIN-REPORT v2.2 coverage (iter 137). The *recoverable* column counts items in {EX, IM, TS, TD}; the *unrecoverable* column counts {NA, AB}.

Corpus	n_{units}	EX	IM	TS	TD	NA	AB	Recov.
mega_20260704	98	7	2	0	4	0	5	13/18 (72.2 %)
n10_seed_expansion	5	2	0	1	0	3	12	3/18 (16.7 %)
n2_reward_tensor_resume	160	2	0	1	4	3	8	7/18 (38.9 %)

Coupling to iter-125 chained-R: iter-125’s “ $R \geq 4$ on zvf” finds corpus-shape-invariant interpretation at the multi-stack level. iter-133’s $R = 0.34$ on N10 is the corresponding measurement on a single-stack panel with seed as the noise axis. *The two corpora are two regimes, not one:* mega-98 multi-stack responds to systematic variance sources; N10 single-stack responds to seed-noise variance.

Operational recommendation: (a) cross-corpus P5 claims must specify which axes are swept in each corpus; (b) cross-corpus portability of dominance ratios is *not* automatic — report both absolute η^2 per axis per channel (corpus-comparable) AND relative R (corpus-shape-relative); (c) P7 controllers should require $\eta^2_{\text{step_band,zvf}} \geq 0.10$ as a per-corpus eligibility check before applying step-level zvf-based firing rules.

Artifacts. scripts/p5p8/p5_iter133_n10_eta2_bootstrap.py (≤ 282 LoC);
experiments/results/p5p8/p5_iter133_per_axis_eta2.tsv (12 rows, 3 axes $\times$ 4 channels);
experiments/results/p5p8/p5_iter133_chained_R.tsv (4 rows);
experiments/results/p5p8/p5_iter133_step_band.tsv (20 rows, 5 bands $\times$ 4 channels);
experiments/results/p5p8/p5_iter133_summary.json (verdicts + cross-corpus portability table).

Cross-paper coupling. Iter-7 (N10 zvf $\rightarrow$ acc), Iter-89 (N2 algo-axis η^2), Iter-93 (mega-98 stack-axis η^2), Iter-101 (zvf130 risk-index η^2), Iter-125 (chained R across mega/N2), Iter-129 (P5 headline-CI audit), FRONTIER-INSIGHTS Round 1 Critic-Degeneracy Hypothesis.

8 Exhibit: MIN-REPORT v2.2 Cross-Corpus Portability (iter 137)

Prior MIN-REPORT audits (iter-105 row 121, iter-113 row 127a, iter-117 row 132, iter-121 row 137) all measured coverage on the mega_20260704 corpus (n=98 manifests)—the corpus for which the standard was developed. The natural next question is whether the standard is *portable*: can it be applied to other, smaller corpora in the worktree at non-trivial recovery rates? Or is the standard *corpus-conditional*, with the manifest-emission model baked in?

Method

For each (corpus $\times$ item) cell on a 3×18 grid, we classify the encoding mode on a 6-class scheme that extends iter-117 row 132: **EX** explicit top-level JSON field; **IM** implicit in filename; **TS** derivable from a tabular/TSV-style file; **TD** tensor-derivable from a reward-vector array; **NA** corpus is fundamentally N/A for the item; **AB** absent (no live source). Three corpora are audited:

- **C1 mega_20260704** (n=98 manifests + 98 group_tensors + 98 cells.tsv)
- **C2 n10_seed_expansion** (n=5 per-seed JSON files, partial)
- **C3 n2_reward_tensor_resume** (4 methods $\times$ 40 steps; reward tensors only; *no manifest*)

Per-corpus coverage

Falsifiable hypotheses (4/4 PASS)

- **H1 PASS** EX-mode counts are (mega, n10, n2) = (7, 2, 2). N2 ties N10 on EX because both emit group_size and zvf as bare fields; mega emits the full v2.2 schema (7 keys).

- **H2 PASS RECOVERABLE** (EX+IM+TS+TD) counts are (mega, n2, n10) = (13, 7, 3). **N2 BEATS N10** because the reward tensor IS the source for Items 14, 15, 17 (TD mode per iter-113).
- **H3 PASS** N2 needs ≥ 11 items of new emission (NA or AB after recovery). The upgrade path is to emit a `n2_manifest.json` with `model_family`, `rollout_temperature`, `decontam`, `sampler_backend`, `advantage_baseline`, `token_mask`, `kl_beta`, `heldout_split`, `reward_model_signature`.
- **H4 PASS** 14/18 items (77.8 %) have corpus-differentiated encoding mode; only 4/18 (Items 3, 5, 10, 11) are corpus-uniform (3 AB/AB/AB and 1 EX/EX/EX). Shannon entropy per item: median 0.355, max 0.613 (Items 1, 8).

Per-item heterogeneity (3×18 grid)

The full matrix is in `experiments/results/p5p8/p5_iter137_corpus_x_item.tsv` (54 cells). Three structural patterns emerge:

1. **ZVF-derived items (13, 14, 15, 16, 17) are the most portable.** All five are recoverable on mega (1 EX + 4 TD) and on N2 (1 EX + 4 TD); on N10 only Item 13 is TS-derivable (per-step zvf in `step_log`). The reward tensor IS the canonical source.
2. **Stack-axis items (1, 4) require filename regex.** They are IM on mega, EX on N10 (per-seed “model” field), AB on N2. N2 needs a new emission to recover `model_family`.
3. **Schema-only items (3, 10, 11) are corpus-uniformly AB.** No live source on any of the three corpora. These remain *schema-only* until a new emitter is wired in.

Why this matters

The MIN-REPORT standard’s title is “Report the Stack, Not the Label”. Iter-137 audits whether the standard is portable across corpora *of different shapes* (full manifests vs. partial seeds vs. bare tensors) and finds that:

- The standard is **not corpus-uniform**: 14/18 items have corpus-differentiated encoding mode. Reviewers who generalize “the MIN-REPORT standard requires X” must scope X to a specific corpus type.
- **N2 is closer to mega than n10 is** (7/18 vs 3/18 recoverable), even though N2 has no manifest at all—the reward tensor compensates.
- **N2’s upgrade path is concrete**: emit a 9-key `n2_manifest.json` to lift N2 from 7/18 to $\geq 15/18$ recoverable (and N2 will then match the mega profile on Items 1, 6, 7, 9, plus the TD items already shared).

Cross-paper coupling

- **P5 iter-113 row 127a** — the content-layer audit showed Items 14, 15, 17 are DAA-recoverable on mega; iter-137 extends to Items 13, 14, 15, 16, 17 being corpus-portable via TD.
- **P5 iter-117 row 132** — the structural-encoding audit on mega alone; iter-137 audits 3 corpora and shows the encoding modes are NOT uniform.
- **P6 iter-134 row 150** — the registry per-row measured-field audit. The P6 registry `min_report` block IS a fourth corpus of MIN-REPORT applicability; iter-137’s framework would extend to a 4×18 grid in a future synthesis iteration.
- **FRONTIER INSIGHTS Round 2 (ZVF = signal availability)**: iter-137 confirms the (frontier synthesis) framing: Items 13-17 (ZVF- derived) are the most corpus-portable, because the signal IS in the tensor wherever the tensor exists.

9 The RL-Reproducibility Threat Model

Security engineering made progress on reproducible claims by naming the ways a claim can fail. We do the same. The *threat* is a ranking flip: a published comparison (“A beats B”) that inverts or

dissolves when the undisclosed parts of the stack change. The *adversary* need not be malicious; the ordinary incentive gradient—publish the configuration that worked—produces the same selection effects that motivated expected-validation reporting in NLP [8] and environment-and-seed audits in deep RL [11].

9.1 Confound Vectors

A confound vector is a class of undisclosed stack variation capable of carrying a flip. Our taxonomy has five, each anchored to documented evidence:

- V1 — Numerics and serving.** Backend, sampling engine, kernels, precision, decoding parameters. Documented magnitude: $17\times$ (Section 5.1). The dominant vector in our corpus.
- V2 — Objective semantics under a shared label.** Loss-form divergence hidden by a common name: clip asymmetry, dynamic sampling, normalization, ratio level. Documented magnitude: qualitative telemetry inversion (mean ZVF 0.00 vs. 0.58; Section 5.2).
- V3 — Signal geometry.** Group size and its schedule, advantage normalization, and the accuracy regime relative to the ZVF U-shape. Documented magnitude: mean ZVF $0.33 \rightarrow 0.08$ moving G from 4 to 16 on Qwen3-32B (Table 26); regime aliasing per Table 25.
- V4 — Environment and data.** Contamination of the training or evaluation pool and parser/verifier idiosyncrasies. Documented magnitude: contamination-driven score inflation [40, 31]; sub-resolution reward jitter zeroing a telemetry channel (Section 6.7).
- V5 — Evaluation protocol.** Held-out split identity, prompt template, answer extraction. Documented magnitude: protocol-level variance large enough to reorder leaderboards in independent audits [3, 12, 24].

9.2 Flip-Risk Grades R0–R5

Given two runs’ manifests (Section 10), we grade the risk that a comparison between them is confounded. The grade is determined by the most severe difference found across the seven items of Table 24.

- R0 — Replicate.** All seven items identical; runs differ at most in hardware instance. Comparison measures noise floor.
- R1 — Seed-only.** Items identical; seeds differ under a declared seed policy. Comparison is the intended unit of statistical analysis [5, 1]; report dispersion, not a single delta.
- R2 — Declared-knob.** Items differ only in declared, quantitatively reported knobs (learning rate, steps, G with fixed schedule form). Comparable, with the knob difference stated as a caveat or bridged by a matched run.
- R3 — Stack-vector.** At least one V1/V3/V4/V5 vector differs (backend, precision, parser, split, schedule form). Documented flip magnitudes reach $17\times$; algorithmic conclusions require a bridging run on a common stack before they are admissible.
- R4 — Label collision.** The same algorithm label with different objective semantics (V2): the comparison is between stacks wearing one name. Verdict: NOT-COMPARABLE; the label must be split (e.g., “DAPO(dynamic-sampling)” vs. “DAPO(clip-surrogate)”).
- R5 — Undeclared.** One or more of the seven items is missing from either manifest. Verdict: UNVERIFIABLE; no grade can be assigned, which is itself the finding.

Two consequences follow. First, most cross-paper comparisons in today’s RL-for-LLM literature are R3–R5 by construction, because items 3, 4, and 7 are almost never reported; the literature’s aggregate ranking of GRPO-family algorithms is therefore an average over unmodeled stacks. Second, the grading is mechanical—it requires no judgment call beyond the manifest contents— which is what makes it suitable for continuous integration (Section 10) rather than for post-hoc reviewer heroics.

10 The Toolchain That Makes It Enforceable

Checklists change practice only when compliance is cheaper than non-compliance. The NeurIPS reproducibility program showed that a checklist plus process measurably improves reporting [28]; our

aim is to push the marginal cost of compliance to near zero by making the manifest a build artifact rather than a writing task. We specify five components. All five are specifications with reference behavior defined by the audits in this paper; we are explicit below about which behaviors have already been exercised and which are design targets.

trl-min-report-rl (manifest emitter). A trainer plugin (reference target: a TRL [37] callback, with adapters for other frameworks in Table 2) that auto-emits the seven-item manifest as JSON at run start and appends the per-step ZVF/GU trajectory (item 4) during training. The author never writes the manifest by hand; it is generated from the same objects the trainer actually executes, closing the gap between what a paper says and what the code did. Appendix A shows two complete manifests.

grpo-stackdiff (comparability verdicts in CI). A CLI that takes two manifests and returns (i) an item-by-item diff, (ii) the flip-risk grade R0–R5 of Section 9.2, and (iii) a CI verdict: COMPARABLE (R0–R2), NOT-COMPARABLE (R3–R4), or UNVERIFIABLE (R5). The intended integration point is a repository or leaderboard CI job: a pull request adding “method A beats method B” must attach two manifests whose stackdiff verdict is COMPARABLE, or the claim is rewritten as stack-conditioned. Appendix A walks one diff to its NOT-COMPARABLE verdict.

GRPO-Registry (machine-readable stack catalog). A registry mapping algorithm labels to the set of loss-form tuples published under that label, with one entry per (label, loss form, framework) triple. The registry is what turns Exhibit 2 from an anecdote into a lookup: “DAPO” resolves to at least two registered semantics, and a citation can point at the specific one.

MIN-REPORT-RL Auditor (0–100 badge). A scorer that reads a paper plus its artifacts and awards points per item (weighted toward items 3, 4, and 7, the least-reported and highest-leverage rows of Table 24), yielding a badge suitable for OpenReview or repository display. Unlike a self-attested checklist, the badge is computed from artifacts, in the spirit of artifact-evaluation tracks.

LeverTrace (citation-lever audit). A tool answering a narrow question: when paper X is cited as evidence for lever L (“we use asymmetric clipping following [X]”), does X actually report the manifest items needed to support L? LeverTrace walks a citation graph and flags credited-but-unreported levers, targeting the accretion process by which a label acquires empirical reputation that no single stack ever demonstrated.

10.1 Feasibility Evidence: The Telemetry Is Already Actionable

The manifest’s most novel items (4 and 5) are not merely reportable—they are cheap enough to compute that they can drive control decisions live, which we take as strong evidence that emitting them adds negligible overhead. In our single-file open-trainer audit (Stratum B; one auditable reimplementation, held-out delta / mean ZVF / rollout count per arm: grpo +0.500 / 0.25 / 120; drgrpo +0.575 / 0.27 / 120; dap0 +0.550 / 0.00 / 174, i.e. +45% rollouts for its ZVF-0.00 guarantee; adaptive-*G* grpo +0.575 / 0.23 / 186), a zvf-triage callback escalated the group size $4 \rightarrow 6 \rightarrow 8$ during training in response to ZVF spikes. A quantity that can be acted on per-step can certainly be logged per-step; conversely, the dap0 arm’s rollout premium is exactly the kind of cost–benefit fact that only surfaces when items 4 and 5 are reported together.

11 Related Work

Reproducibility crises in deep RL. Our position is a deliberate echo of Henderson et al. [11], who showed that deep-RL results depend on codebase, hyperparameters, and even the random seeds chosen, to the point that the same algorithm from two repositories behaves as two algorithms. Follow-on work built the statistical repair kit: rigorous comparison protocols [5] and small-sample-aware aggregate metrics [1], and position work has questioned what RL benchmarking can establish at all [16]. RL-for-LLMs adds a new failure surface those works did not face: the sampling distribution is produced by a serving stack complex enough to be a research artifact in its own right, and often closed.

Reporting standards in NLP and ML. Dodge et al. [8] argued that results are functions of compute budget and hyperparameter search, and proposed reporting expected validation performance as a function of budget; our manifest is the RL-for-LLM analogue, replacing “budget” with “stack.” The NeurIPS reproducibility program [28] demonstrated that checklists plus process shift community behavior—and also that self-attested checklists have limits, which motivates our artifact-computed badge (Section 10). Statistical treatments of evaluation noise [25, 24] quantify the uncertainty floor against which our 12-cell within-noise finding (Section 5.3) should be read.

Evaluation-stack sensitivity. The lm-evaluation-harness effort documented, from years of maintenance experience, that prompt formatting, answer extraction, and implementation details shift LLM benchmark scores by ranking-relevant margins, and argued for shared, versioned evaluation infrastructure [3]. Independent audits reach the same conclusion for reasoning evaluations specifically [12], and contamination studies show the data side of the same coin [40, 31]. Our claim is that *RL training* inherits all of this—the verifier is an evaluator in the loop—plus the V1–V3 vectors unique to on-policy optimization.

The GRPO-variant explosion. The label space we target is growing quickly: GRPO [33, 7], DAPO [39], Dr.GRPO [22], GSPO [42], and a long tail of descendants [21, 27, 41, 23, 4]. Analyses that unify the family—GRPO’s preference-learning reading [30, 38] and zero-variance-prompt recovery methods [20]—support our premise that members differ by a small set of loss-form coordinates, which is exactly what item 1 of the manifest records. Empirical scaling studies of GRPO knobs [35] likewise presuppose the per-knob disclosure the manifest mandates. The companion TINKERRL-BENCH pillar papers (P1 scaling, P2 ZVF, P3 group size, P4 length bias) provide the measurement substrate this position paper argues should become the norm.

The reporting-standards lineage of MIN-REPORT (iter 109 verification). The MIN-REPORT standard inherits directly from four documented dataset/model reporting standards, each of which we verified against CrossRef metadata at iter 109 (`p5_iter109_crossref_verify.tsv`; 4/4 pass year + 5-gram title + author-family-overlap ≥ 0.6 vs the CrossRef record). Mitchell et al. [26] introduced *Model Cards* to document intended use, out-of-scope use, and per-slice evaluation results for trained models; this lineage gives us MIN-REPORT items 1 (architecture and training procedure), 2 (intended and out-of-scope use), 7 (per-slice metrics, not just averages), and 8 (ethical considerations and caveats). Gebru et al. [10] extended the idea to datasets via *Datasheets*, motivating items 3 (dataset composition), 5 (preprocessing, cleaning, labeling, uses), 9 (annotation process), and 10 (discriminatory-impact mitigations). Bender and Friedman [2] focused on natural-language data and added the speaker-demographic axis, motivating items 11 (speaker consent and curation rationale) and 12 (annotator demographics and quality control) which are not present in Mitchell’s or Gebru’s templates. Pushkarna et al. [29] introduced *Data Cards* to bridge research and industry practice, motivating items 4 (dataset schema and provenance), 6 (sampling distribution and representativeness), 13 (per-stack axis field markers, the `zvf_yield_residual` line in MIN-REPORT v2.1), and 14 (audit transparency and version tracking). The 4 papers cover **12 distinct MIN-REPORT items** (cross-coupling audit: `p5_iter109_minreport_coupling.tsv`); the remaining 6 items in the 18-item manifest are either RL-specific (item 15 `zvf` per-step, item 17 group-size schedule) or stack-coverage items without a direct analogue in the prior reporting-standards literature (item 16 was rejected as a signal-bearing field at iter 81). Two of the four entries (Bender and Friedman [2] and Pushkarna et al. [29]) were not present in `paper/references.bib` prior to iter 109; iter 109 adds them with full DOI, year, volume/number/pages, and verified author lists. The other two entries (Mitchell et al. [26], Gebru et al. [10]) were already in the bib; iter 109 retrofits them with DOIs and now actually cites them in the related work (prior to iter 109 they were *present in the bibliography but uncited*). The cross-coupling audit also exposes a measurement asymmetry: of the 12 MIN-REPORT items covered by these reporting-standards papers, items 9–12 (annotation process and annotator-demographic axes) have the lowest evidence density in our live manifests (`p5_iter109_bib_audit.tsv`, see also iter 105’s 5-vs-3 discriminative-vs-primitive split), motivating a future-iter extension of MIN-REPORT to formalise annotation-process fields at the same strictness as the stack-axis fields.

12 Limitations

Position papers owe their readers an honest account of the evidence’s edges; five limitations are load-bearing here. First, the headline exhibits (the $17\times$ backend flip, the cross-stack DAPO telemetry flip, the 12-cell head-to-head, and the 368-run audit) are internal program records: they were run under a declared protocol and are summarized faithfully, but they have not undergone the full artifact-release pipeline of the TINKERRL-BENCH Stratum-A results, and the closed stack involved cannot be independently re-executed by readers—an irony we acknowledge and which the manifest standard is designed to at least make legible. Second, the gradient-signal validation behind item 4’s theory ($\text{corr} = +0.71$ between gradient norm and $p(1 - p)$) is toy-scale—a 0.5B model on synthetic arithmetic with an open backward pass—and is directional evidence only; we do not report it as an effect size. Third, our evidence is concentrated on GSM8K-style verifiable-reward tasks and the GRPO family; the seven items were chosen for that regime, and reward-model-based RLHF stacks may need additional items (reward-model version and training data at minimum). Fourth, the flip-risk grades encode a conservative ordering (label collision graded above generic stack divergence) that we have validated on our own audits but not against a broad external corpus; the grade boundaries, especially R2 vs. R3, will need community calibration. Fifth, the toolchain of Section 10 is a specification with feasibility evidence, not a released product suite; the live adaptive- G result shows the telemetry is cheap, but manifest adapters for every framework in Table 2 remain to be built. None of these limitations weakens the central asymmetry the paper rests on: every documented lever we cite moved results by more than the label differences the field currently headlines.

13 Conclusion: A Call to Action for Venues

The evidence pattern is consistent and, we argue, sufficient to change reporting policy: on a fixed stack, four algorithm labels landed within 0.034 reward of one another; across stacks, a backend swap alone moved reward by 0.794 and a single label produced opposite telemetry signatures. RL-for-LLM results are stack-conditioned. The community’s papers should say so, in a form machines can check.

Our call to action is addressed to venues, because venues set the price of non-disclosure:

1. **Require the manifest.** Any paper whose contribution depends on an RL-for-LLM comparison should attach one seven-item manifest (Table 24) per compared run, as supplementary JSON. Auto-emission (Section 10) makes this a one-line cost for open trainers; for closed stacks, “provider-internal, version-dated” is an acceptable and informative value—the point is that conditioning be declared, not that all stacks be open.
2. **Grade the comparisons.** Reviewers should see the `grpo-stackdiff` grade for every headline comparison. R0–R2 comparisons support algorithmic claims; R3–R4 comparisons should be rewritten as stack-conditioned observations; R5 should be treated as a missing-experiment flag, exactly as a missing baseline is today.
3. **Split colliding labels.** Where the registry shows one label with multiple registered semantics, require the qualified form (“DAPO(dynamic-sampling)”) in abstracts and tables. Precedents exist: the field already distinguishes “PPO” from “PPO-clip” when it matters [32, 11].
4. **Reward the telemetry.** Per-step ZVF/GU trajectories and group-size schedules should count as artifact contributions in their own right, as evaluation-harness versioning now does [3, 28].

The deep-RL community needed a reporting reckoning [11], a statistics paper [5, 1], and a decade to converge on reporting norms. RL-for-LLMs can shortcut that decade: the confounds are already documented, the manifest is seven items, and the enforcement is a diff. Report the stack, not the label.

A Worked Example: Two “DAPO” Manifests and Their Verdict

This appendix instantiates the standard end-to-end on the two runs of Exhibit 2 (Section 5.2): the same algorithm label, run on a closed training stack and on an open single-file trainer. Field values marked “provider-internal” illustrate the disclosure convention for closed stacks (Section 13): a closed value is acceptable; an absent field is not.

A.1 Manifest A: “DAPO” on the Closed Tinker Stack

```
{
  "minreport_version": "0.1",
  "run_id": "tinker_gsm8k_dapo_s42",
  "label_claimed": "DAPO",
  "loss_form": {
    "ratio_level": "token",
    "clip": {"low": 0.20, "high": 0.28, "symmetric": false},
    "token_mask": "response-only",
    "advantage_norm": "group-mean-centered, std-normalized",
    "dynamic_sampling": false,
    "surrogate_note": "grpo + asymmetric clip; dynamic sampling not expressible on
      this stack"
  },
  "reference_policy": {
    "type": "frozen-init",
    "kl": {"coeff": 0.0, "estimator": "none"}
  },
  "sampler_backend": {
    "trainer": "tinker (closed), accessed 2026-06",
    "sampler": "provider-internal",
    "precision": "provider-internal",
    "decoding": {"temperature": 1.0, "top_p": 1.0, "max_new_tokens": 1024}
  },
  "zvf_gu_trace": {
    "logged_per_step": true,
    "artifact": "wandb://zvf-audit/tinker_gsm8k_dapo_s42/zvf_gu.jsonl",
    "mean_zvf": 0.58
  },
  "group_size": {"G": 8, "schedule": "fixed"},
  "heldout": {
    "split": "gsm8k-test-500",
    "disjoint_from_reward_env": true
  },
  "decontamination": {"performed": false, "parser_probe": "not-run"}
}
```

A.2 Manifest B: “DAPO” on an Open Single-File Trainer

```
{
  "minreport_version": "0.1",
  "run_id": "colab_e3_dapo_s0",
  "label_claimed": "DAPO",
  "loss_form": {
    "ratio_level": "token",
    "clip": {"low": 0.20, "high": 0.28, "symmetric": false},
    "token_mask": "response-only",
    "advantage_norm": "group-mean-centered, std-normalized",
    "dynamic_sampling": true,
    "surrogate_note": null
  },
  "reference_policy": {
    "type": "frozen-init",
    "kl": {"coeff": 0.0, "estimator": "none"}
  },
  "sampler_backend": {
    "trainer": "open-colab-trainer @ git commit (single-file reimplementations)",
    "sampler": "hf-generate",
    "precision": "bf16",
    "decoding": {"temperature": 1.0, "top_p": 1.0, "max_new_tokens": 1024}
  },
  "zvf_gu_trace": {
```

```

    "logged_per_step": true,
    "artifact": "wandb://zvf-colab-experiments/E3_open_audit/dapo/zvf_gu.jsonl",
    "mean_zvf": 0.00,
    "rollouts_total": 174
  },
  "group_size": {"G": 4, "schedule": "fixed (resampling via dynamic_sampling)"},
  "heldout": {
    "split": "gsm8k-heldout (disjoint prompt pool)",
    "disjoint_from_reward_env": true
  },
  "decontamination": {
    "performed": true,
    "parser_probe": "micro-jitter eps ~ U(0, 1e-4): zvf-sensitive, pcd-invariant"
  }
}

```

A.3 The Stackdiff Verdict

Running the differ on the two manifests:

```

$ grpo-stackdiff manifest_a.json manifest_b.json

ITEM DIFF
  1 loss_form.dynamic_sampling false != true [V2]
  1 loss_form.surrogate_note "grpo+asym-clip" != null [V2]
  3 sampler_backend.trainer tinkler (closed) != open-colab [V1]
  3 sampler_backend.precision provider-internal!= bf16 [V1]
  4 zvf_gu_trace.mean_zvf 0.58 != 0.00 [consequence]
  5 group_size.G 8 != 4 [V3]
  7 decontamination.performed false != true [V4]
  7 decontamination.parser_probe not-run != micro-jitter [V4]

FLIP-RISK GRADE
  R4 (label collision): same label_claimed "DAP0", divergent objective
    semantics (dynamic_sampling). Additional R3 vectors present (V1, V3, V4).

CI VERDICT: NOT-COMPARABLE
  Do not report these runs as "DAP0 vs DAP0". Qualified labels required:
    manifest_a -> DAP0(clip-surrogate, no-dynamic-sampling)
    manifest_b -> DAP0(dynamic-sampling)
  A bridging run on a common stack is required before any algorithmic claim.

```

The diff makes Exhibit 2 mechanical: the telemetry inversion (mean ZVF 0.58 vs. 0.00) is listed not as a difference to adjudicate but as a *consequence* of the item-1 divergence, and the verdict line is exactly what a leaderboard CI job would gate on. Note also what the diff does *not* say: it does not claim either run is wrong, or that the closed stack is inferior—only that the pair cannot support a label-level conclusion. That restraint is the entire content of the standard.

References

- [1] Rishabh Agarwal, Max Schwarzer, Pablo Samuel Castro, Aaron C. Courville, and Marc Bellemare. Deep reinforcement learning at the edge of the statistical precipice. In *Advances in Neural Information Processing Systems*, volume 34, pages 29304–29320, 2021.
- [2] Emily M. Bender and Batya Friedman. Data statements for natural language processing: Toward mitigating system bias and enabling better science. *Transactions of the Association for Computational Linguistics (TACL)*, 6:587–604, 2018. doi: 10.1162/tacl_a_00041.
- [3] Stella Biderman, Hailey Schoelkopf, Lintang Sutawika, Leo Gao, Jonathan Tow, Baber Abbasi, Alham Fikri Aji, Pawan Sasanka Ammanamanchi, Sidney Black, Jordan Clive, et al. Lessons from the trenches on reproducible evaluation of language models. *arXiv preprint arXiv:2405.14782*, 2024.

- [4] ByteDance Seed, Yu Yue, Yufeng Yuan, Qiying Yu, Xiaochen Zuo, Ruofei Zhu, Wenyuan Xu, Jiaze Chen, Chengyi Wang, Tiantian Fan, et al. VAPO: Efficient and reliable reinforcement learning for advanced reasoning tasks. *arXiv preprint*, 2025. doi: 10.48550/arXiv.2504.05118. arXiv:2504.05118.
- [5] Cédric Colas, Olivier Sigaud, and Pierre-Yves Oudeyer. A hitchhiker’s guide to statistical comparisons of reinforcement learning algorithms. *arXiv preprint*, 2019. doi: 10.48550/arXiv.1904.06979. arXiv:1904.06979.
- [6] Tri Dao, Daniel Y. Fu, Stefano Ermon, Atri Rudra, and Christopher Ré. FlashAttention: Fast and memory-efficient exact attention with IO-awareness. In *Advances in Neural Information Processing Systems*, 2022.
- [7] DeepSeek-AI, Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Ruoyu Zhang, Runxin Xu, Qihao Zhu, Shirong Ma, Peiyi Wang, et al. DeepSeek-R1: Incentivizing reasoning capability in LLMs via reinforcement learning. *arXiv preprint*, 2025. doi: 10.48550/arXiv.2501.12948. arXiv:2501.12948.
- [8] Jesse Dodge, Suchin Gururangan, Dallas Card, Roy Schwartz, and Noah A. Smith. Show your work: Improved reporting of experimental results. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 2185–2194, 2019. doi: 10.18653/v1/D19-1224.
- [9] Leo Gao, John Schulman, and Jacob Hilton. Scaling laws for reward model overoptimization. In *International Conference on Machine Learning*, 2023.
- [10] Timnit Gebru, Jamie Morgenstern, Briana Vecchione, Jennifer Wortman Vaughan, Hanna Wallach, Hal Daumé III, and Kate Crawford. Datasheets for datasets. *Communications of the ACM*, 64(12):86–92, 2021. doi: 10.1145/3458723.
- [11] Peter Henderson, Riashat Islam, Philip Bachman, Joelle Pineau, Doina Precup, and David Meger. Deep reinforcement learning that matters. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 32, 2018. doi: 10.1609/aaai.v32i1.11694.
- [12] Andreas Hochlehnert et al. A sober look at progress in language model reasoning. *arXiv preprint*, 2025. doi: 10.48550/arXiv.2504.07086. arXiv:2504.07086.
- [13] Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations*, 2022. URL <https://openreview.net/forum?id=nZeVKeeFYf9>.
- [14] Jian Hu, Xibin Wu, Zilin Zhu, Xianyu, Weixun Wang, Dehao Zhang, and Yu Cao. OpenRLHF: An easy-to-use, scalable and high-performance RLHF framework. *arXiv preprint*, 2024. doi: 10.48550/arXiv.2405.11143. arXiv:2405.11143.
- [15] Hamish Ivison, Yizhong Wang, Jiacheng Liu, Zeqiu Wu, Valentina Pyatkin, Nathan Lambert, Noah A. Smith, Yejin Choi, and Hannaneh Hajishirzi. Unpacking DPO and PPO: Disentangling best practices for learning from preference feedback. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. doi: 10.48550/arXiv.2406.09279. arXiv:2406.09279.
- [16] Emma Jordan, Adam White, Bruno Castro da Silva, Martha White, and Philip S. Thomas. Position: Benchmarking is limited in reinforcement learning research. *arXiv preprint*, 2024. doi: 10.48550/arXiv.2406.16241. arXiv:2406.16241.
- [17] Victoria Krakovna, Jonathan Uesato, Vladimir Mikulik, Matthew Rahtz, Tom Everitt, Ramana Kumar, Zac Kenton, Jan Leike, and Shane Legg. Specification gaming: The flip side of AI ingenuity. DeepMind Safety Research blog, 2020. <https://deepmindsafetyresearch.medium.com/specification-gaming-the-flip-side-of-ai-ingenuity-c85bdb0deeb4>.
- [18] Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large language model serving with PagedAttention. In *Proceedings of the 29th Symposium on Operating Systems Principles (SOSP)*, 2023. doi: 10.1145/3600006.3613165.
- [19] Nathan Lambert, Jacob Morrison, Valentina Pyatkin, Shengyi Huang, Hamish Ivison, Faeze Brahman, Lester James V. Miranda, et al. Tulu 3: Pushing frontiers in open language model post-training. *arXiv preprint*, 2024. doi: 10.48550/arXiv.2411.15124. arXiv:2411.15124.

- [20] Thanh-Long V. Le, Myeongho Jeon, Kim Vu, Viet Lai, and Eunho Yang. No prompt left behind: Exploiting zero-variance prompts in LLM reinforcement learning via entropy-guided advantage shaping. *arXiv preprint*, September 2025. Also referred to as AERO / RL-ZVP in the literature; arXiv:2509.21880.
- [21] Shih-Yang Liu, Xin Dong, Ximing Lu, Shizhe Diao, Peter Belcak, Mingjie Liu, Min-Hung Chen, Hongxu Yin, Yu-Chiang Frank Wang, Kwang-Ting Cheng, Yejin Choi, Jan Kautz, and Pavlo Molchanov. GDPO: Group reward-decoupled normalization policy optimization for multi-reward rl optimization. *arXiv preprint arXiv:2601.05242*, 2026.
- [22] Zichen Liu, Changyu Chen, Wenjun Li, Penghui Qi, Tianyu Pang, Chao Du, Wee Sun Lee, and Min Lin. Understanding R1-Zero-Like training: A critical perspective. *arXiv preprint*, 2025. doi: 10.48550/arXiv.2503.20783. arXiv:2503.20783; introduces Dr. GRPO.
- [23] Yingqian Lu et al. GRPO-LEAD: Length-aware reward shaping for GRPO. *arXiv preprint*, 2025. doi: 10.48550/arXiv.2504.09696. arXiv:2504.09696.
- [24] Lovish Madaan, Aaditya K. Singh, Rylan Schaeffer, Andrew Poulton, Sanmi Koyejo, Pontus Stenetorp, Sharan Narang, and Dieuwke Hupkes. Quantifying variance in evaluation benchmarks. *arXiv preprint*, 2024. doi: 10.48550/arXiv.2406.10229. arXiv:2406.10229.
- [25] Evan Miller. Adding error bars to evals: A statistical approach to language model evaluations. *arXiv preprint*, 2024. doi: 10.48550/arXiv.2411.00640. arXiv:2411.00640.
- [26] Margaret Mitchell, Simone Wu, Andrew Zaldivar, Parker Barnes, Lucy Vasserman, Ben Hutchinson, Elena Spitzer, Inioluwa Deborah Raji, and Timnit Gebru. Model cards for model reporting. In *Proceedings of the Conference on Fairness, Accountability, and Transparency (FAT*)*, pages 220–229, 2019. doi: 10.1145/3287560.3287596.
- [27] Gongrui Nan et al. NGRPO: Negative-enhanced group relative policy optimization. *arXiv preprint*, 2025. doi: 10.48550/arXiv.2509.18851. arXiv:2509.18851.
- [28] Joelle Pineau, Philippe Vincent-Lamarre, Koustuv Sinha, Vincent Larivière, Alina Beygelzimer, Florence d’Alché Buc, Emily Fox, and Hugo Larochelle. Improving reproducibility in machine learning research: A report from the NeurIPS 2019 reproducibility program. *Journal of Machine Learning Research*, 22(164):1–20, 2021.
- [29] Mahima Pushkarna, Andrew Zaldivar, and Oddur Kjartansson. Data cards: Purposeful and transparent dataset documentation for responsible ai. In *Proceedings of the 2022 ACM Conference on Fairness, Accountability, and Transparency (FAccT ’22)*, pages 1776–1826, 2022. doi: 10.1145/3531146.3533231.
- [30] Rafael Rafailov, Archit Sharma, Eric Mitchell, Stefano Ermon, Christopher D. Manning, and Chelsea Finn. Direct preference optimization: Your language model is secretly a reward model. In *Advances in Neural Information Processing Systems*, volume 36, 2024. URL https://proceedings.neurips.cc/paper_files/paper/2023/hash/a85b405ed65c6477a4fe8f#.
- [31] Martin Riddell, Ansong Ni, and Arman Cohan. Quantifying contamination in evaluating code generation capabilities of language models. *arXiv preprint*, 2024. arXiv:2403.04811.
- [32] John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms. 2017. URL <https://arxiv.org/abs/1707.06347>. arXiv:1707.06347.
- [33] Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Mingchuan Zhang, Y.K. Li, Y. Wu, and Daya Guo. DeepSeekMath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint*, 2024. doi: 10.48550/arXiv.2402.03300. arXiv:2402.03300.
- [34] Guangming Sheng, Chi Zhang, Zilingfeng Ye, Xibin Wu, Wang Zhang, Ru Zhang, Yanghua Peng, Haibin Lin, and Chuan Wu. HybridFlow: A flexible and efficient RLHF framework. *arXiv preprint*, 2024. doi: 10.48550/arXiv.2409.19256. arXiv:2409.19256; veRL.
- [35] Zelin Tan, Hejia Geng, Xiaohang Yu, Mulei Zhang, Guancheng Wan, Yifan Zhou, Qiang He, Xiangyuan Xue, Heng Zhou, Yutao Fan, Zhongzhi Li, Zaibin Zhang, Guibin Zhang, Chen Zhang, Zhenfei Yin, Philip Torr, and Lei Bai. Scaling behaviors of LLM reinforcement learning post-training: An empirical study in mathematical reasoning. *arXiv preprint arXiv:2509.25300*, 2025.

- [36] Thinking Machines Lab. Tinker: A training API for researchers. <https://thinkingmachines.ai/tinker>, 2025.
- [37] Leandro von Werra, Younes Belkada, Lewis Tunstall, Edward Beeching, Tristan Thrush, Nathan Lambert, and Shengyi Huang. TRL: Transformer reinforcement learning. <https://github.com/huggingface/trl>, 2022.
- [38] Yihong Wu, Liheng Ma, Lei Ding, Muzhi Li, Xinyu Wang, Kejia Chen, Zhan Su, Zhanguang Zhang, Chenyang Huang, Yingxue Zhang, Mark Coates, and Jian-Yun Nie. It takes two: Your GRPO is secretly DPO. *arXiv preprint*, 2025. doi: 10.48550/arXiv.2510.00977. arXiv:2510.00977.
- [39] Qiyang Yu, Zheng Zhang, Ruofei Zhu, Yufeng Yuan, Xiaochen Zuo, Yu Yue, Weinan Dai, Tiantian Fan, Gaozhong Liu, Lingjun Liu, et al. DAPO: An open-source LLM reinforcement learning system at scale. *arXiv preprint*, 2025. doi: 10.48550/arXiv.2503.14476. arXiv:2503.14476.
- [40] Hugh Zhang, Jeff Da, Dean Lee, Vaughn Robinson, Catherine Wu, Will Song, Tiffany Zhao, Pranav Raja, Dylan Slack, Qin Lyu, Sean Hendryx, Russell Kaplan, Michele Lunati, and Summer Yue. A careful examination of large language model performance on grade school arithmetic. *arXiv preprint*, 2024. arXiv:2405.00332 (GSM1k).
- [41] Xichen Zhang et al. Scaf-GRPO: Scaffolded group relative policy optimization for enhancing LLM reasoning. *arXiv preprint*, 2025. doi: 10.48550/arXiv.2510.19807. arXiv:2510.19807.
- [42] Chujie Zheng, Shixuan Liu, Mingze Li, Xiong-Hui Chen, Bowen Yu, Chang Gao, Kai Dang, Yuqiong Liu, Rui Men, An Yang, Jingren Zhou, Jianwei Zhou, and Junyang Lin. Group sequence policy optimization. *arXiv preprint*, 2025. doi: 10.48550/arXiv.2507.18071. arXiv:2507.18071; GSPO.
- [43] Lianmin Zheng, Liangsheng Yin, Zhiqiang Xie, Chuyue Sun, Jeff Huang, Cody Hao Yu, Shiyi Cao, Christos Kozyrakis, Ion Stoica, Joseph E. Gonzalez, Clark Barrett, and Ying Sheng. SGLang: Efficient execution of structured language model programs. In *Advances in Neural Information Processing Systems*, 2024.