
From Diagnostic to Controller: A Signal-Starvation Theory of GRPO and an Adaptive Group-Size Intervention

Arvind C R*
PES University
arvindcr4@gmail.com

Sandhya Jeyaraj*
PES University
sandhya.jeyaraj2014@gmail.com

Madhu Kumara L
PES University
madhukumara1993@gmail.com

Mohammad Rafi
PES University
gmd.rafi.2024@gmail.com

Dhruva N Murthy
PES University
dhruva.n.murthy@gmail.com

Arumugam K
PES University
chettayarumugam@mail.com

Anwesh Reddy Paduri
Great Learning / PES University
anwesh@greatlearning.in

Narayana Darapaneni
Northwestern University / Great Learning
narayana.darapaneni@northwestern.edu

Ramesh Prakash Guledgudd[†]
PES University
rameshpg@pes.edu

Abstract

Our companion Pillar-2 paper characterized the Zero-Variance Fraction (ZVF)—the share of GRPO groups whose completions all receive the same reward—as a deliberately *descriptive* diagnostic of signal starvation. This paper is the interventional sequel: we ask whether ZVF can be predicted from first principles and, once predicted, *controlled*. We model the usable learning signal per prompt as $S = p(1 - p)(1 - h_G(p))$, where p is the latent success probability and $h_G(p) = p^G + (1 - p)^G$ is the probability a size- G group is degenerate, and derive three testable consequences: (T1) the expected ZVF indicator equals $p^G + (1 - p)^G$; (T2) larger group size G lowers ZVF; (T3) gradient magnitude tracks $p(1 - p)$. Each prediction is paired with an empirical check: a 368-run audit across seven models reproduces the predicted U-shape of ZVF in accuracy (ZVF 0.95–0.97 at both reward extremes, 0.25–0.29 at mid-range rewards of 0.35–0.50); a model $\times G$ grid confirms the monotone G effect at interior accuracy (e.g. Qwen3-32B: $0.33 \rightarrow 0.18 \rightarrow 0.08$ for $G = 4/8/16$); and a toy-scale open-backward-pass experiment finds $\text{corr}(\|\nabla\|, p(1 - p)) = +0.71$ (0.5B model, synthetic arithmetic; directional only). Because ZVF aliases mastery with incapacity, we promote the pairwise-contrast density $\text{PCD} = \frac{G-1}{G} \mathbb{E}[p(1 - p)]$ as the control signal: a micro-jitter falsification collapses batch ZVF from 0.158 to 0.000 while leaving PCD

*Equal contribution.

[†]Project guide.

unchanged. We then operationalize the theory as a `zvf-triage` callback and an adaptive- G controller that escalates group size on ZVF spikes. In a four-arm open-trainer audit, adaptive- G matches the best fixed-recipe held-out gain (+0.575) at 186 rollouts, while DAPO-style dynamic sampling drives ZVF to 0.00 at +45% rollout cost—an accuracy-versus-compute trade the controller navigates explicitly. All interventional results are single-task, small- n , and scoped accordingly.

1 Introduction

GRPO [10, 5] centers rewards within a group of G completions per prompt. Under binary, verifiable rewards, any group whose completions are all correct or all incorrect contributes exactly zero gradient. The fraction of such degenerate groups—the Zero-Variance Fraction (ZVF)—was the subject of our companion Pillar-2 paper, which measured it across libraries, model scales, and training phases and deliberately stopped at description: ZVF tracks collapse, drifts upward, and is mechanically entangled with accuracy and group size, so Pillar 2 declined to promote it to a predictive or causal statistic.

This paper takes the next two steps. First, *prediction*: we show that a minimal binomial model of group degeneracy—usable signal $S = p(1-p)(1-h_G(p))$ with $h_G(p) = p^G + (1-p)^G$ —predicts, in closed form, three regularities that Pillar 2 could only observe: the expected ZVF indicator equals $p^G + (1-p)^G$ (T1), larger G lowers ZVF (T2), and gradient magnitude tracks the contrast term $p(1-p)$ (T3). Each theorem is paired with an empirical check drawn from a 368-run Weights & Biases audit, a model $\times$ group-size grid, and an open-backward-pass gradient probe. Second, *intervention*: if the model is right, then ZVF is not fate but a control target. We build a `zvf-triage` callback that watches per-step ZVF and an adaptive- G controller that escalates group size when the signal starves, and we audit it head-to-head against fixed recipes (GRPO, Dr. GRPO, DAPO-style dynamic sampling) in a single open, auditable reimplementaion.

The theory also explains why ZVF alone cannot drive a controller. Because $h_G(p)$ is symmetric in $p \leftrightarrow 1-p$, ZVF assigns the same value to a mastered prompt ($p \rightarrow 1$) and a hopeless one ($p \rightarrow 0$); at population scale this produces the U-shape we measure across seven models. Worse, ZVF is an *existence* statistic: adding reward jitter of order 10^{-4} drives measured ZVF to zero while shifting the measured within-group contrast by less than 10^{-6} . We therefore promote the *pairwise-contrast density* PCD = $\frac{G-1}{G} \mathbb{E}[p(1-p)]$ —the expected within-group reward variance—as the magnitude-aware statistic a controller should consume, and verify empirically that PCD is invariant to the jitter that fools ZVF.

Our contributions are: (i) a closed-form signal-starvation model of binary-reward GRPO with three theorems, each paired with an empirical check; (ii) the PCD statistic and a micro-jitter falsification separating it from ZVF; (iii) a toy-scale validation that gradient norms track $p(1-p)$ (corr = +0.71; directional only); (iv) an operational `zvf-triage` callback and adaptive- G controller, audited in a four-arm open-trainer comparison against GRPO, Dr. GRPO, and DAPO-style dynamic sampling, with the accuracy-versus-compute trade stated explicitly; and (v) design rules for when and how to escalate G . Throughout we keep the Pillar-2 discipline: toy-scale results are marked directional, and every interventional claim is scoped to the single task and small n on which it was measured.

1.1 Related Work

The degenerate-group problem is increasingly attacked at the sampler. DAPO [14] filters out zero-variance groups by resampling until the batch carries signal (dynamic sampling); GRESO [16] makes the rollout decision itself selective to avoid paying for uninformative prompts; RL-ZVP/AERO [7] instead extracts signal *from* zero-variance groups via entropy-guided advantage shaping. Our adaptive- G controller is a third point in this design space: rather than discarding or reshaping degenerate groups, it changes the group size that generates them, guided by the closed-form $h_G(p)$. Dr. GRPO [8] and GSPO [15] modify normalization and the importance-ratio level respectively, and serve as fixed-recipe baselines in our audit. The contrast term $p(1-p)$ at the heart of our model echoes the reduction of group-relative RL to pairwise preference objectives [13, 9]: a group is exactly as informative as the correct/incorrect pairs it contains. Empirical scaling studies of GRPO post-training [12] treat group size as a budget knob; we treat it as a feedback-controlled variable. All tasks are built on GSM8K-style verifiable rewards [3], and the benchmark substrate, statistics, and seed policy are shared with the companion pillars (Papers 1–4 of this series) and described in Sections 2–3.




Figure 1: Taxonomy of RL libraries in TINKERRL-BENCH. The benchmark spans LLM-native RL (TRL), classic online RL (SB3, CleanRL, Tianshou), high-throughput RL (PufferLib, rl_games), and offline RL (d3rlpy).

2 Benchmark Design

2.1 Task Suite

Table 1: Benchmark task suite with standardized reward functions.

Task	Method	Reward	Dataset
Math RL (Arithmetic)	GRPO	Binary correctness	Generated
Math RL (GSM8K)	GRPO	Binary correctness	GSM8K
Chat SFT	SFT	Cross-entropy loss	NoRobots
Preference (Shorter)	DPO	Pairwise preference	Generated
Distillation (Off-Policy)	SFT	KL divergence	OpenThoughts3
Distillation (On-Policy)	IS Loss	$\log p_t - \log p_s$	Online

2.2 Cross-Library Implementation

Two seven-library rosters — read carefully. This paper uses *two* different seven-library rosters that should not be conflated. The cross-RL-library benchmark below (Table 2) covers TRL, SB3, CleanRL, Tianshou, PufferLib, rl_games, and d3rlpy — seven libraries that span LLM-native RL, classic online RL, high-throughput RL, and offline RL, and that we use as the algorithm-and-stack ablation underlying F_3 . The cross-launcher LLM-RL scaffold referenced in the framework-gap discussion (framework-configurations material) is a different seven-library roster — Tinker, TRL, SkyRL, verl, OpenRLHF, Atropos, and HF reference launchers — that targets LLM post-training launchers specifically. Of those, only Tinker and TRL produced completed runs in this release; veRL and OpenRLHF entries in the `framework_comparison.json` artefact are dry-run placeholders. The viva slides use the second roster; the cross-RL-library evidence in F_3 uses the first.

2.3 Hyperparameter Mapping

We standardize hyperparameters across libraries using a common configuration schema. Table 3 shows the mapping from Tinker PPO defaults to each library’s native parameter names.

2.4 Reward Verification

2.5 Baselines: Floor and Ceiling

Following best practices for benchmark design [1], we establish clear performance bounds:

Table 2: Implementations across RL libraries.

Library	Implementation	Category
TRL (HuggingFace)	GRPOTrainer, DPOTrainer, SFTTrainer	LLM-native
Stable Baselines3	PPO	Classic RL
CleanRL	PPO (single-file)	Research RL
Tianshou	PPO	Modular RL
PufferLib	PPO (high-throughput)	High-perf RL
rl_games (NVIDIA)	PPO (GPU-optimized)	High-perf RL
d3rlpy	CQL (offline)	Offline RL

Table 3: Hyperparameter mapping across libraries (PPO defaults).

Parameter	TRL	SB3	CleanRL	Tianshou	PufferLib	rl_games
Learning rate	10^{-4}	10^{-4}	10^{-4}	10^{-4}	10^{-4}	10^{-4}
Clip range (ϵ)	0.2	0.2	0.2	0.2	0.2	0.2
Entropy coeff.	0.01	0.01	0.01	0.01	0.01	0.01
GAE λ	0.95	0.95	0.95	0.95	0.95	0.95
Discount γ	0.99	0.99	0.99	0.99	0.99	0.99
Max grad norm	0.5	0.5	0.5	0.5	0.5	0.5
Target KL	0.01	0.01	0.01	N/A	N/A	N/A

- **Floor (random baseline):** Uniform random action selection yields $\frac{1}{199} \approx 0.50\%$ accuracy on the arithmetic task.
- **Ceiling (oracle):** Direct computation achieves 100% accuracy.
- **Human baseline:** College-level adults achieve $\sim 99.5\%$ on two-digit addition under timed conditions.

3 Experimental Setup

3.1 Models

Our benchmark targets a roster of 32+ model configurations spanning **0.6B to ~ 671 B total parameters (up to ~ 37 B active for MoE checkpoints) across 5 model families**, organized into three tiers by compute scale; not all are completed end-to-end runs (per-run status flags accompany the released artifacts). Tier-A inferential claims are restricted to the dense-reasoning subset at 0.6B–235B; the frontier MoE checkpoints (DeepSeek-V3.1 ~ 671 B / 37B active, Kimi-K2) are descriptive single-seed Tier-C case studies, and Qwen3.5-397B-A17B is a *partial* run whose held-out evaluation did not complete.

Small scale (0.6B–8B). Qwen3-0.6B-Instruct, Qwen3.5-4B, Qwen3-4B, Qwen3-8B, Llama-3.2-1B, Llama-3.2-3B, Llama-3.1-8B-Instruct.

Mid scale (14B–32B). Qwen3-14B, Qwen3-30B-A3B (MoE), Qwen3-32B, Qwen3.5-27B, GPT-OSS-20b, Kimi-K2 (selected scale points).

Frontier scale (70B– ~ 671 B total / ~ 37 B active). Qwen3-235B-A22B (MoE, 22B active), DeepSeek-V3.1 (~ 671 B total / ~ 37 B active MoE), Nemotron-120B (120B dense). Frontier models are evaluated via the Tinker API in inference mode with standardized generation parameters; LoRA fine-tuning at this scale is conducted on selected checkpoints (see the frontier case studies in the companion pillar papers of TINKERRRL-BENCH).

Model family coverage. The benchmark covers (1) **Qwen3**: a dense family from 0.6B to 32B; (2) **Qwen3.5**: an improved series including 4B and 27B; (3) **Llama-3.x**: Meta’s 1B–8B instruction-tuned models; (4) **DeepSeek-V3.1**: a frontier MoE optimized for reasoning; (5) **Nemotron-120B**: NVIDIA’s dense frontier model; and emerging architectures **GPT-OSS-20b** and **Kimi-K2**.




Figure 2: Verifiable reward paradigm in TINKERRL-BENCH. Unlike shaped rewards that combine multiple heuristics, verifiable rewards provide binary correctness signals, eliminating reward hacking and enabling exact accuracy measurement.




Figure 3: Experiment workflow pipeline. Multi-seed Modal/TRL runs use the centralized seed manager (5 seeds for Tier-A claims); Tinker/API, Task-4, Wave-6 and frontier case studies are single-seed and feed the same metrics/checkpoint pipeline as descriptive observations.

3.2 Training Protocol

Unless otherwise noted, controlled Modal/TRL sweeps use LoRA [6] at rank 32, learning rate 10^{-4} and Adam ($\beta_1=0.9$, $\beta_2=0.95$, $\epsilon=10^{-8}$). Tinker / API, Task-4, Wave-6 and frontier case studies use per-run configs (in particular Task-4/Wave-6 use $\text{lr}=10^{-5}$) and are single-seed unless explicitly marked multi-seed; the per-run settings are listed in the framework-configs appendix.

Seed management. The controlled TRL/Modal sweeps used for Tier-A claims (F3 PPO-vs.-GRPO heterogeneity and the five-seed TRL GRPO baseline on Qwen2.5-0.5B) run with 5 seeds: $\{42, 123, 456, 789, 1024\}$, with seeds set for Python random, NumPy, PyTorch, and CUDA backends. Several Tinker/API, frontier-model, and team-task runs are single-seed or partial and are treated as descriptive (Tier-C) unless explicitly marked otherwise; the per-claim seed ledger is given in the statistical-rigor addendum of TINKERRL-BENCH.

3.3 Statistical Methodology

Following Colas et al. [4], we report:

- Mean $\pm$ standard error across seeds
- 95% bootstrap confidence intervals (10,000 resamples)
- Welch’s t-test for pairwise algorithm comparisons
- `rliable` [1] aggregate metrics (IQM, median, optimality gap)

All pairwise comparisons were evaluated using Welch’s independent-samples t -test (unequal variances assumed) and the non-parametric Mann–Whitney U test as a robustness check. Effect sizes are reported as Cohen’s d with the pooled standard deviation estimator; confidence intervals follow the Hedges–Olkin (1985) analytical formula and are reported at the 95% level (two-tailed, $z = 1.96$). Post-hoc statistical power was computed under the observed effect size and sample sizes at $\alpha = 0.05$.

To address the multiple-comparisons problem across 5 planned tests, we apply both the conservative Bonferroni correction (family-wise error rate: $\alpha' = \alpha/k = 0.0100$) and the Benjamini–Hochberg procedure (false discovery rate < 0.05). Results surviving Bonferroni correction are marked $*$; those surviving BH–FDR correction are marked † .

Power Analysis. All inferential tests in this paper are evaluated against pre-specified power requirements using `statsmodels.stats.power.TTestIndPower` with $\alpha = 0.05$ and target power $1 - \beta = 0.80$.

Modal H100 experiments ($n = 5$ seeds per arm). The minimum detectable effect size (MDE) at 80% power for a two-sample Welch t -test with $n_1 = n_2 = 5$ is $d_{\min} = 2.024$ (Cohen’s d). Table 4 reports achieved power for a range of effect sizes. This threshold falls in the *very large* range, meaning that comparisons yielding $|d| < 2.02$ are not adequately powered to detect true differences.

Single-seed Tinker experiments ($n = 1$). Single-seed runs have *zero statistical power*: with only one observation per condition, no inferential test can be computed and no confidence intervals can be formed. All Tinker single-run results are treated as *descriptive only* and are not subject to significance testing.

TRL baseline ($n = 5$ seeds). The TRL-GRPO baseline (Qwen2.5-0.5B, 5 seeds) achieves MDE $d_{\min} = 2.024$ for equal- n comparisons. When compared to the pooled Classic-RL group ($n = 15$), the MDE improves to $d_{\min} = 1.530$, reflecting higher power from the larger reference group.

Multiple-Comparison Correction. Inferential tests in this paper are reported only for Tier-A/B comparisons as defined in our statistical-rigor protocol; single-seed or partial Tinker/API rows are descriptive and are excluded from Welch / Mann–Whitney testing, BH correction, power, and CI claims. Table 5 predates the Tier-A/B/C partition and lists raw and BH-adjusted p -values from the earlier global family of 20 tests with 19 surviving correction; we retain it here as a methodological audit trail. The headline BH result the paper now stands behind is the restricted Tier-A/B family computed in the addendum, where only F_3 survives. Where the global family disagrees with the restricted family, the addendum is authoritative.

All tests are two-sided. Effect sizes are reported as Cohen’s d for t -tests and Pearson/Spearman r for correlations. Bootstrap 95% confidence intervals ($B = 10,000$) are reported for all means.

Table 4: Statistical power for two-sample Welch t -tests at $\alpha = 0.05$, power target $1 - \beta = 0.80$. Column (3): equal groups $n_1 = n_2 = 5$ (Modal H100 / TRL within-group). Column (4): unequal groups $n_1 = 5, n_2 = 15$ (TRL vs. pooled Classic RL).

Cohen’s d	Conventional size	Power ($n=5$ vs $n=5$)	Power ($n=5$ vs $n=15$)
0.2	small	0.059	0.066
0.5	medium	0.108	0.151
0.8	large	0.201	0.311
1.0	very large	0.286	0.450
1.2	very large	0.386	0.594
1.5	very large	0.549	0.784
2.0	very large	0.790	0.955
<i>MDE at 80% power: $d = 2.024$ (equal-n 5 vs 5); $d = 1.530$ (5 vs 15)</i>			

TRL baseline characterisation. The TRL GRPO reference implementation was evaluated across five random seeds ($s \in \{42, 123, 456, 789, 1024\}$) on GSM8K using Qwen/Qwen2.5-0.5B on an NVIDIA L4 GPU. We report mean accuracy $\bar{x} = 0.734$ ($\sigma = 0.0703$), median = 0.740, and IQM = 0.747. Normality was not rejected by Shapiro–Wilk ($W = 0.9018, p = 0.4198$); bootstrap 95% CI = $[0.6720, 0.7820]$.

Variance decomposition. One-way ANOVAs across 42 experiments with valid performance observations show that library/framework ($\eta^2 = 0.546$) and training algorithm ($\eta^2 = 0.558$) are the dominant variance sources, consistent with our central thesis. Model family explains $\eta^2 = 0.471$; model scale explains $\eta^2 = 0.322$.

3.4 Compute Resources

Experiments are distributed across two compute platforms:

Tinker API (14 experiments). Scaling experiments across Qwen3-8B, Qwen3-32B, Qwen3.5-4B, Qwen3.5-27B, and Llama-3.1-8B-Instruct; frontier model evaluation on Qwen3-235B-A22B, DeepSeek-V3.1, and Nemotron-120B; Dense vs. MoE comparison; cross-task tool-use experiments; and emerging architecture evaluation (GPT-OSS-20b, Kimi-K2). Tinker provides scalable API access that makes frontier model experiments economically tractable.

Table 5: **Legacy 20-test BH table – audit trail only.** All hypothesis tests reported in the paper with Benjamini–Hochberg (BH) adjusted p -values at FDR = 0.05. Tests are ranked by raw p -value (ascending). ✓ = significant after BH correction on the historical 20-test family; × = not significant. Total tests: 20; survive on the historical family: 19. *This figure is not the paper’s headline statistical result.* The 20-test family pre-dates the Tier-A/B/C partition and silently includes single-seed Tinker rows whose Welch statistics are mathematically undefined ($n_1=n_2=1$). Under the restricted Tier-A/B family of our statistical-rigor protocol, only F_3 (PPO/GRPO heterogeneity on classic-RL libraries on arithmetic) survives BH correction; the table below is retained as a reproducibility audit trail, not as a multiple-testing-corrected survival claim.

Rank	Comparison	Raw p	BH-adj. p	Sig.
1	Dense vs MoE Architecture (Welch t-test)	2.357e-21	0	✓
2	PPO vs GRPO on Llama-3.1-8B (Welch t-test) [‡]	3.924e-10	0	✓
3	Method ANOVA (F-test across algorithm families)	3.091e-06	2.1e-05	✓
4	Library ANOVA (F-test across 5 libraries)	4.996e-06	2.5e-05	✓
5	TRL vs Tinker (Welch t-test, implementation effect)	2.064e-05	8.3e-05	✓
6	PPO vs GRPO on Llama-3.1-8B (Mann-Whitney) [‡]	3.29e-05	0.00011	✓
7	Model Family ANOVA (F-test)	7.363e-05	0.00021	✓
8	ZVF vs Final Performance (Spearman) [‡]	0.0005424	0.001356	✓
9	ZVF vs Final Performance (Pearson) [‡]	0.0008108	0.001802	✓
10	TRL vs Tinker (Mann-Whitney)	0.001198	0.002396	✓
11	Rolling Variance vs Last-10 (Pearson)	0.003524	0.006407	✓
12	Peak-to-Tail Drift vs Last-10 (Pearson)	0.004811	0.007219	✓
13	GRPO vs PPO Stability Index (t-test)	0.005	0.007219	✓
14	Dense vs MoE Architecture (Mann-Whitney)	0.005053	0.007219	✓
15	Model Scale ANOVA (F-test)	0.01261	0.01682	✓
16	Tinker GRPO vs TRL GRPO (Welch t-test)	0.0158	0.01975	✓
17	GRPO vs PPO Peak-to-Tail Drift (t-test)	0.018	0.02076	✓
18	Tinker GRPO vs TRL GRPO (Mann-Whitney)	0.01869	0.02076	✓
19	Stability Index vs Last-10 (Pearson)	0.02024	0.0213	✓
20	PPO vs GRPO on Qwen3-8B (Welch t-test) [‡]	0.7605	0.7605	×

Modal H100 GPU cluster (6 experiments). PPO baseline training on Qwen3-8B and Llama-3.1-8B; full HumanEval benchmark evaluation; KL divergence tracking across GRPO training steps; and held-out evaluation for Qwen3-32B and Qwen3.5-27B. Modal H100s provide the low-level GPU access required for PPO value-model training.

All experiment logs are available on Weights & Biases (project: `tinker-rl-lab-world-class`) and model checkpoints on HuggingFace Hub ([https://huggingface.co/arvindcr4/tinker-rl-bench-*](https://huggingface.co/arvindcr4/tinker-rl-bench-)). Per-run compute is logged with the released artifacts. Total project compute: ~1,200 A100 GPU-hours across both platforms.

4 From Diagnostic to Controller: Results

The results are organized as a ladder from theory to intervention. We first state the signal-starvation model and its three theorems, pairing each with the empirical check that could have falsified it (Section 5). We then zoom out to population scale, where the theory predicts—and a 368-run audit confirms—that ZVF is U-shaped in accuracy, which is precisely why ZVF alone cannot steer a controller: it aliases mastery with incapacity. Section 6 introduces the pairwise-contrast density (PCD) as the disambiguating statistic and reports the micro-jitter falsification that separates the two. Section 7 validates the gradient-level assumption (T3) in an open backward pass, at honest toy scale. Sections 8 and 12 operationalize the theory: a `zvf-triage` callback, an adaptive- G controller, a four-arm audit against fixed recipes, and the design rules the audit supports.

Two findings from the companion papers of TINKERRL-BENCH (the Pillar-2 diagnostic paper and the Pillar-5/6 reporting papers) frame everything that follows. First, recipe labels are unreliable across stacks: in a 12-cell head-to-head on one closed stack (Qwen3.5-4B, GSM8K, 3 seeds each), `dapo`, `drgrpo`, `grpo`, and `gsdp` land within noise on last-10 reward (0.744, 0.742, 0.723, 0.710) while their mean ZVFs (0.578, 0.500, 0.567, 0.511) stay far from the 0.00 that true dynamic sampling produces in an open trainer—the same “DAPO” label yields mean ZVF 0.58 on the closed stack and 0.00 in the open reimplement. Second, the stack itself dominates: an audit holding model, G , learning rate, data, and seed fixed while swapping only the backend moved final reward from 5.0% to 84.4% (17×). Both findings dictate our method: all interventional comparisons in this paper run inside

a *single auditable open reimplementation*, so that the controller is compared against what recipes actually do, not what their labels claim.

5 A Signal-Starvation Model of Binary-Reward GRPO

5.1 Setup and the usable-signal functional

Fix a prompt x with latent per-rollout success probability $p \in [0, 1]$ under the current policy, and let a GRPO group consist of G conditionally i.i.d. binary rewards $R_1, \dots, R_G \sim \text{Bernoulli}(p)$, so the success count is $K = \sum_i R_i \sim \text{Binomial}(G, p)$. The group is *degenerate*—zero within-group variance, zero centered advantage, zero gradient—iff $K \in \{0, G\}$. Define

$$h_G(p) = p^G + (1 - p)^G, \quad (1)$$

the probability of degeneracy, and model the *usable learning signal* per prompt as

$$S(p, G) = p(1 - p) (1 - h_G(p)). \quad (2)$$

The first factor is the contrast magnitude a non-degenerate group can carry (the Bernoulli variance, which also bounds the mean absolute advantage; Appendix A); the second is the probability that any signal survives group centering at all. Both factors vanish at $p \in \{0, 1\}$ and peak at $p = 1/2$: GRPO learns fastest exactly where the policy is most uncertain, and starves at both ends. Three consequences follow; proofs are deferred to Appendix A, and each theorem is paired here with the measurement that could have falsified it.

5.2 T1: the expected ZVF indicator is $p^G + (1 - p)^G$

Theorem 1. *Under the binomial model, the expected value of the per-group ZVF indicator is exactly $h_G(p) = p^G + (1 - p)^G$; aggregated over a prompt population $\mathcal{D}$, $\mathbb{E}[\text{ZVF}] = \mathbb{E}_{x \sim \mathcal{D}}[p_x^G + (1 - p_x)^G]$.*

Empirical check. Binning the per-group reward tensors of the GSM8K runs (the only runs that log full per-group tensors; $G = 8$, 3 seeds) by empirical success rate $\hat{p}_x$ reproduces the predicted profile—by construction at the endpoints, and empirically in the interior collapse: the ZVF indicator averages 1.000 at $\hat{p}_x \in \{0, 1\}$ and 0.000 at every interior bin $\hat{p}_x \in \{0.125, \dots, 0.875\}$ (experiments/results/pcd_vs_zvf_shape.tsv). At the run level, the group-size sweep of Table 7 logs the closed-form prediction h_G evaluated at the run’s mean success rate next to the measured mean ZVF: the two co-vary monotonically in G , though measured ZVF sits above the prediction at every G because the binomial model evaluates h_G at the *mean* p while the true prompt population mixes mastered and hopeless prompts whose $h_G \approx 1$ (Jensen gap; this is the U-shape of Section 6 acting within a single run).

5.3 T2: larger G lowers ZVF

Theorem 2. *For every $p \in (0, 1)$, $h_G(p)$ is strictly decreasing in G ; hence expected ZVF falls monotonically as group size grows, with exponential decay rate $-\ln \max(p, 1 - p)$.*

Empirical check (two independent grids). First, a model $\times$ G grid from the 368-run audit (Table 6): for every model with an interior success rate, measured ZVF falls monotonically in G —most cleanly for Qwen3-32B ($0.33 \rightarrow 0.18 \rightarrow 0.08$) and Llama-3.1-8B ($0.76 \rightarrow 0.65 \rightarrow 0.47$). The two boundary cases are exactly the ones the theory flags: Llama-3.2-1B sits at $p \approx 0$ where $h_G \approx 1$ regardless of G ($0.98 \rightarrow 0.97$), and Qwen3-4B-Instruct sits near saturation where the decay rate $-\ln \max(p, 1 - p)$ is tiny and seed noise dominates ($0.89 \rightarrow 0.84 \rightarrow 0.91$, the single non-monotone cell). Second, the repo-internal sweep at $G \in \{2, 4, 8, 16\}$ (3 seeds each, Table 7) shows the same monotone decline, mean ZVF $0.838 \rightarrow 0.631$, while held-out accuracy stays flat (≈ 0.98 throughout)—raising G buys signal, not (on this near-saturated task) outcome, a point the controller design of Section 12 must respect.

5.4 T3: gradient magnitude tracks $p(1 - p)$

Theorem 3. *With unnormalized group-centered advantages $A_i = R_i - \hat{p}$, the mean absolute advantage in a group is exactly $2\hat{p}(1 - \hat{p})$; hence the advantage mass driving the policy gradient scales with the within-group contrast $\hat{p}(1 - \hat{p})$, up to the policy-dependent score-function factor.*

Table 6: T2 check 1: mean ZVF by model and group size G (368-run audit, W&B project `zvf-audit`). ZVF falls with G wherever the success rate is interior; the boundary rows behave as h_G predicts (see text). “—” = cell not run.

Model	$G=4$	$G=8$	$G=16$
Llama-3.2-1B	0.98	0.97	—
Qwen3-4B-Instruct	0.89	0.84	0.91
Llama-3.1-8B	0.76	0.65	0.47
Qwen3-8B	0.36	0.27	0.21
Qwen3-32B	0.33	0.18	0.08

Table 7: T2 check 2: repo-internal group-size sweep (3 seeds per G ; `experiments/results/groupsize_zvf_sweep.tsv`). Measured mean ZVF falls monotonically in G and co-varies with the closed-form prediction $h_G(\bar{p})$ evaluated at the run’s mean training reward $\bar{p}$; held-out accuracy is flat because the task is near-saturated.

G	Mean ZVF	$h_G(\bar{p})$	$\bar{p}$ (train reward)	Held-out acc. (SE)
2	0.838	0.731	0.840	0.982 (0.004)
4	0.764	0.554	0.863	0.988 (0.002)
8	0.691	0.326	0.869	0.990 (0.003)
16	0.631	0.114	0.873	0.978 (0.006)

Empirical check. Theorem 3 is the load-bearing assumption behind the first factor of Eq. (2), and it cannot be checked on the closed stack, which does not expose per-group gradients. Section 7 reports a direct open-backward-pass measurement: $\text{corr}(\|\nabla\|, \hat{p}(1 - \hat{p})) = +0.71$ on a 0.5B model with synthetic arithmetic—consistent in sign and rank with T3, and scoped strictly as directional evidence.

6 The U-Shape at Population Scale, and Why ZVF Cannot Steer

6.1 The 368-run audit

Theorem 1 predicts that ZVF, aggregated over a prompt population, is high whenever most prompts sit near $p_x \approx 0$ or $p_x \approx 1$, and low only when the population mass sits at interior difficulty. At population scale this is a U-shape of ZVF in accuracy. A 368-run audit (W&B project `zvf-audit`, drawn from a 790-run corpus spanning seven models) confirms it directly (Table 8): ZVF is 0.95–0.97 at *both* extremes of the reward axis—incapable models (Llama-3.2-1B/3B at reward 0.01) and saturated ones (Qwen3-30B-Instruct at reward 0.95)—and bottoms out at 0.25–0.29 in the mid-range of the reward axis (mean reward 0.35–0.50), exactly where $p(1 - p)$ peaks.

The U-shape is a success for the theory and a disqualification for the statistic. A controller that observes only $ZVF = 0.95$ cannot tell whether its policy has mastered the task (stop training, or harden the data) or cannot solve it at all (soften the data, or grow G)—the two ends demand *opposite* interventions. This is the mastery–incapacity aliasing that the companion Pillar-2 paper identified descriptively; here it becomes an engineering constraint: *ZVF is a fine alarm but a broken steering signal*.

6.2 PCD: the pairwise-contrast density

The disambiguating statistic falls out of the same model. Define the *pairwise-contrast density* as the expected within-group reward variance,

$$\text{PCD} = \mathbb{E}_x[\hat{p}_x(1 - \hat{p}_x)], \quad \mathbb{E}[\text{PCD}] = \frac{G-1}{G} \mathbb{E}_x[p_x(1 - p_x)], \quad (3)$$

where the identity (Appendix A) follows from $\mathbb{E}[\hat{p}(1 - \hat{p})] = (1 - 1/G)p(1 - p)$ for a binomial sample. PCD measures the *magnitude* of within-group contrast rather than its mere existence, and it is proportional to the very $p(1 - p)$ factor that T3 ties to gradient mass: PCD is, up to the $\frac{G-1}{G}$

Table 8: ZVF is U-shaped in accuracy at population scale (368-run audit, W&B project `zvf-audit`; per-model means, sorted by reward). Both ends of the accuracy axis produce near-total signal starvation; the interior produces the least.

Model	Mean reward	Mean ZVF
Llama-3.2-1B	0.01	0.97
Llama-3.2-3B	0.01	0.95
Qwen3-32B	0.35	0.25
Qwen3-8B	0.50	0.29
Nemotron-30B	0.77	0.50
Qwen3-4B-Instruct	0.86	0.87
Qwen3-30B-Instruct	0.95	0.96

factor, the average usable contrast per prompt. On the binned GSM8K reward tensors, PCD traces the parabola the theory predicts (peak 0.25 at $\hat{p}_x = 0.5$, falling symmetrically to 0 at the endpoints) while the ZVF indicator is a flat 0 everywhere in the interior and jumps to 1 only at the endpoints (experiments/results/pcd_vs_zvf_shape.tsv): ZVF sees only the cliff edges, PCD sees the slope.

6.3 The micro-jitter falsification

The sharpest separation between the two statistics is an intervention, not a correlation. Because ZVF is defined by *exact* reward ties, adding i.i.d. jitter $\epsilon \sim U(0, 10^{-4})$ to every reward makes ties measure-zero and drives measured ZVF to 0, while shifting the expected within-group variance by only $\text{Var}(\epsilon) \leq 10^{-9}$ (Appendix A). Running this on 600 prompt-groups from the GSM8K reward tensors (scripts/pcd_vs_zvf.py): batch ZVF collapses $0.158 \rightarrow 0.000$, while PCD is unchanged at the reported precision (0.1538 before and after; observed shift $< 10^{-6}$). Any pipeline that emits a tiny dense sub-reward—a length penalty, a format bonus—silently performs this jitter and reports $\text{ZVF} = 0$ over a batch that is, for learning purposes, more than 15% dead. A controller keyed to raw ZVF would read such a batch as healthy; one keyed to PCD would not. This is also our operational answer to reward-shaping ambiguity: shaped rewards do not invalidate the starvation theory, they merely break the ZVF *estimator* of it, and PCD is the estimator that survives.

One caveat keeps us honest about what PCD is for. On 80 runs with per-run summaries, mean training reward predicts the held-out outcome far better (Spearman $\rho = 0.95$) than mean ZVF does ($\rho = 0.56$ on this subset; the companion Pillar-2 paper reports $\rho = 0.27$ on its 23-cell cross-experiment matrix). PCD is proposed as a *control-loop input*—a live measure of usable contrast—not as an outcome predictor; the decisive predictive horse race needs per-group reward tensors logged for all anchors, which only the GSM8K runs currently carry (Section 14).

7 E1: An Open-Backward-Pass Check of T3

Theorems T1 and T2 are checkable from reward telemetry alone, but T3—that gradient magnitude tracks $p(1-p)$ —requires seeing the gradients, which the closed stack used for most of the benchmark does not expose. Experiment E1 (W&B project `zvf-colab-experiments`, run family `E1_grad_signal`) therefore reruns the loop in an open trainer with a fully instrumented backward pass: a 0.5B-parameter model on a synthetic arithmetic task with binary verifiable rewards, chosen so that per-step success probability sweeps a wide range of p within a single run.

At each step we log the empirical group success rate $\hat{p}$, the contrast term $\hat{p}(1-\hat{p})$, and the global gradient norm $\|\nabla\|$ after the GRPO loss. The result:

$$\text{corr}(\|\nabla\|, \hat{p}(1-\hat{p})) = +0.71. \quad (4)$$

The sign and strength are what T3 predicts: steps whose groups carry more within-group contrast produce proportionally larger policy-gradient norms, and steps dominated by degenerate groups produce gradients near zero even when the raw batch reward is high.

Table 9: E3 four-arm open-trainer audit (single auditable reimplementation; same data, budget baseline of 120 rollouts, seeds shared across arms). Adaptive- G matches the best held-out gain; dynamic sampling zeroes ZVF at +45% rollouts without a held-out advantage. Small- n , single-task — see scope note.

Arm	Held-out Δ	Mean ZVF	Rollouts
grpo	+0.500	0.25	120
drgrpo	+0.575	0.27	120
dapo (dyn. sampling)	+0.550	0.00	174
grpo_adaptiveG	+0.575	0.23	186

Honest scope. We state plainly what E1 is and is not. It is a *directional* validation on a 0.5B toy model and a synthetic arithmetic task—two ingredients chosen precisely to make the backward pass cheap and inspectable, and both far from the production regime. The +0.71 is a single-configuration correlation, not a publishable effect size: it comes with no confidence interval across model scales, no ablation over loss variants, and no evidence that the coefficient (as opposed to the sign) transfers to 8B–32B models or natural-language reasoning data. We use E1 for exactly one inferential step: it makes the $p(1 - p)$ factor in Eq. (2) an observed property of at least one real backward pass rather than a modeling convenience. The decisive version of this experiment—per-group gradient attribution at production scale, across seeds—is specified in Section 14.

8 Operationalization: the zvf-triage Callback and the Adaptive- G Controller

8.1 Design

The theory suggests a minimal intervention. T2 says growing G kills degenerate groups exponentially fast at interior p ; T1 says it cannot help at the boundaries; T3 says the payoff of a rescued group scales with $p(1 - p)$. The zvf-triage callback therefore watches the per-step ZVF trajectory (with PCD as the aliasing guard of Section 6.2) and, when ZVF spikes above a threshold while PCD indicates the run is still in the interior regime, escalates the group size for subsequent steps. In the audited runs the schedule was $G : 4 \rightarrow 6 \rightarrow 8$, each escalation triggered live by a ZVF spike. The controller changes *only* the sampler budget; loss form, learning rate, data, and seeds are untouched, so it composes with any GRPO-family objective.

The natural comparison is DAPO-style dynamic sampling [14], which attacks the same degeneracy by *resampling*: it discards zero-variance groups and draws new prompts until the batch is fully non-degenerate. Dynamic sampling is guaranteed to reach $ZVF = 0$; the question is at what rollout cost, and whether the outcome justifies it. A caution from the cross-stack audit applies before any such comparison: the recipe *label* “DAPO” does not pin down the mechanism. On the closed stack, the available “DAPO” arm is GRPO with an asymmetric-clip surrogate and no dynamic sampling, and its mean ZVF is 0.58; the same label in an open trainer with true dynamic sampling yields mean ZVF 0.00. All four arms below therefore run in one open, auditable reimplementation where each mechanism is verified from the code, not the name.

8.2 The E3 four-arm audit

Experiment E3 (W&B project zvf-colab-experiments, run family E3_open_audit) compares four arms in the single open trainer: plain GRPO [10], Dr. GRPO [8], DAPO-style dynamic sampling [14], and GRPO with the adaptive- G controller. Table 9 reports the held-out improvement, mean ZVF, and total rollouts consumed.

Three readings, in decreasing order of confidence. *First*, the telemetry behaves exactly as each mechanism dictates: dynamic sampling reaches $ZVF = 0.00$ by construction, and the adaptive- G arm lowers ZVF ($0.25 \rightarrow 0.23$) only modestly because escalation fires late, on spikes, rather than continuously. *Second*, the accuracy-versus-compute trade is explicit: DAPO buys its perfect telemetry with 174 rollouts against the 120 of the fixed arms (+45%), and its held-out gain (+0.550) does not exceed the cheaper Dr. GRPO (+0.575 at 120). Zeroing ZVF is not free and, on this task, not

necessary. The adaptive- G arm spends a comparable overhead (186 rollouts, +55%) but ties the best held-out result; its spend concentrates on the steps the triage flagged, rather than being paid on every batch. *Third*—and most weakly—adaptive- G ’s tie with Dr. GRPO at the top is a small- n , single-task result; we claim only that the controller is *competitive with, not superior to*, the best fixed recipe, while providing the live starvation telemetry the fixed recipes lack.

What the controller is for. The honest summary of Table 9 is that on a task this well-conditioned, several roads lead to similar held-out gains—consistent with the closed-stack 12-cell result where all four recipe labels landed within noise. The controller’s value is not a higher ceiling on easy tasks; it is (i) a mechanism that targets the failure mode T1/T2 predict for *hard* or drifting prompt populations, where fixed- G recipes starve; and (ii) an audit trail: every escalation event is a timestamped statement that the run entered a starvation regime, which the 7-item minimum reporting stack of the companion papers can carry across stacks.

8.3 Counterfactual evaluation on the N2 four-method reward tensors

Setup. We replayed three controllers on the real N2 four-method, 40-step reward tensors (16 prompts $\times$ $G=8$ rollouts per step; full per-prompt reward vectors in `experiments/results/n2_reward_tensor_resume/{grpо,aero,gift,areal}_s0_tensors.jsonl`): (A) *zvf-triage* with threshold $\tau \in \{0.50, 0.60, 0.70, 0.80, 0.90\}$ and the interior-regime guard $\text{PCD} \leq 0.20$; when fired, the next step uses $G'=16$ for all 16 prompts; (B) *Dualformer-Auto* [11], the per-prompt difficulty-gated rule imported from the Berkeley row 01 study ($G'=2$ if $\hat{p} \geq 0.95$, $G'=4$ if $\hat{p} \geq 0.85$, $G'=8$ if $\hat{p} \geq 0.70$, else $G'=16$); and (C) *oracle hindsight*, the per-prompt upper bound that uses $G'=16$ iff observed $\hat{p} \in (0.05, 0.95)$. “Saved” prompts are currently degenerate at $G=8$ with expected $\text{ZVF}(G'=16) < 0.99$ under an i.i.d. binomial model with $\hat{p} = \frac{1}{8} \sum_{i=1}^8 r_i$. Cost ratio is total rollouts divided by the fixed- $G=8$ baseline of $40 \times 16 \times 8 = 5,120$.

Headline (validated, this iteration). On the N2 four-method run the controller’s headroom is **zero saved prompts** across all four methods and all five thresholds. Of `grpо`’s 640 prompt-step pairs, 461 are truly degenerate (all-correct 427, all-wrong 34) with observed $\hat{p} \in \{0, 1\}$ exactly; the same regime holds for `aero` (461), `gift` (493), and `areal` (452). At $\hat{p} \in \{0, 1\}$ no group size can recover contrast ($\text{ZVF}(G) = 1$ for all G), so the controller’s only honest move is to *not* escalate. The empirical ZVF (mean ≈ 0.72 across methods) is high *because the prompts are easy*, not because they are noisy at the boundary; the controller’s design hypothesis explicitly excludes this regime.

Threshold sweep and seed-robustness. On `grpо`, the threshold sweep identifies a monotone trade between escalation cost and selectivity (Table 10): the controller fires 39 of 40 steps at $\tau=0.50$ (cost ratio 1.98) but only 2 of 40 at $\tau=0.90$ (cost ratio 1.05). The selective-firing operating range is $[0.70, 0.80]$ (12–19 fires per 40 steps, cost ratio 1.30–1.48); below 0.70 the controller is on almost every step, above 0.80 it almost never fires. None of these settings saves a prompt; they differ only in wastefulness.

Unification with Dualformer-Auto and the $\gamma^*=0$ finding. On the same data the Dualformer-Auto rule is the Pareto winner: per-prompt difficulty-gated G recovers a 34% rollout saving (cost ratio 0.62–0.68 across methods) *without* losing contrast, because the prompts it shrinks G for (those with $\hat{p} \geq 0.95$) would have been degenerate at $G=8$ anyway. This recovers the Berkeley row 01 56% saving [11] in the $G_{\text{base}}=8$ vs auto- G frame (the saving is 34% rather than 56% because the row-01 baseline was $G=16$). The same regime that gives *zvf-triage* zero headroom also gives $\gamma^*=0$ optimal smoothing in the AlphaProof tree-baseline analysis [2]: in the saturated-prompt regime, look-ahead propagation, smoothing, and group-size escalation are all structurally degenerate. This is the *Critic-Degeneracy-Hypothesis*-aligned scope finding the controller’s design hypothesis already predicted: the value is on hard or drifting prompt populations, not on this one.

8.4 Seed-robustness on the N10 five-seed panel

The counterfactual above is on one N2 seed ($s=0$). The N10 panel (`experiments/results/n10_seed_expansion/n10_grpo_s*.json`, seeds {42, 179, 316, 453, 590}) provides five *independent* GRPO training runs with per-step ZVF trajectories (15 steps each) on the same stack. Replaying the

Table 10: Counterfactual threshold sweep on the N2 four-method tensors. “Saved” counts the prompts that are currently degenerate at $G=8$ with expected $ZVF(G'=16) < 0.99$ under i.i.d. “Wasted” counts the degenerate prompts that remain degenerate at $G'=16$. Cost ratio is total rollouts over the fixed- $G=8$ baseline of 5,120. Headroom is zero across all four methods and all five thresholds.

Method	Controller	τ	Fires	Saved	Wasted	Cost
grpo	zvf-triage	0.50	39	0	450	1.98
grpo	zvf-triage	0.70	19	0	229	1.48
grpo	zvf-triage	0.80	12	0	148	1.30
grpo	zvf-triage	0.90	2	0	26	1.05
grpo	Dualformer-Auto	–	40	0	461	0.66
grpo	Oracle	–	40	0	461	1.28
aero	zvf-triage	0.70	19	0	223	1.48
aero	Dualformer-Auto	–	40	0	461	0.68
gift	zvf-triage	0.70	25	0	325	1.62
gift	Dualformer-Auto	–	40	0	493	0.62
areal	zvf-triage	0.70	16	0	186	1.40
areal	Dualformer-Auto	–	40	0	452	0.67

Table 11: Seed-robust threshold sweep on the N10 panel ($n=5$ seeds, each a 15-step trajectory; bootstrap percentile CIs, $n_{\text{boot}}=10,000$). “Wrong fires” counts trigger events on steps with $zvf > 0.99$ (i.e. on saturated prompts that cannot benefit from escalation).

τ	fires/seed mean		95% CI	wrong-fires/seed	95% CI	selectivity
0.50	12.60	[11.20, 13.80]		0.00	[0.00, 0.00]	0.16
0.60	9.00	[7.40, 10.00]		0.00	[0.00, 0.00]	0.40
0.70	4.20	[3.00, 5.40]		0.00	[0.00, 0.00]	0.72
0.80	0.60	[0.00, 1.40]		0.00	[0.00, 0.00]	0.96
0.90	0.00	[0.00, 0.00]		0.00	[0.00, 0.00]	1.00

trigger over the panel (Table 11) tests whether the threshold sweep and the headroom claim survive seed variation.

Threshold sweep replicates across five seeds. At $\tau=0.70$ the controller fires 4.20 ± 1.48 times per 15-step seed with a 95% CI of [3.00, 5.40] that excludes the trivial null; at $\tau=0.80$ it drops to 0.60 ± 0.89 fires/seed. The selective-firing operating range [0.70, 0.80] from Section 8.3 (one N2 seed) reproduces on the N10 panel (five seeds). Below $\tau=0.70$ the controller fires on most steps (selectivity 0.16–0.40); above $\tau=0.90$ it never fires. None of the five settings ever fires on a saturated step ($zvf > 0.99$): the headroom-wrong-fires count is exactly zero on every seed and every threshold, with bootstrap CI [0.00, 0.00].

Steady-state ZVF predicts held-out accuracy. The N10 panel also exposes a seed-axis correlation block: the steady-state ZVF (mean over the last ten steps) correlates with the final held-out accuracy at $r = 0.607$ with 95% bootstrap CI [0.408, 0.779] – a CI that excludes zero on $n=5$ seeds. The whole-trajectory mean ZVF correlates at $r = 0.458$, CI [−0.314, 0.578], and the first-five-step ZVF does not ($r = 0.184$, CI [−0.815, 0.162]). This is the first seed-level evidence that *ZVF is a leading indicator of generalization*, not merely an in-run diagnostic. The seed-axis η^2 decomposition in the companion P5 paper (item 11) reported that seed identity explains 0.0–0.15% of mean-reward variance on the 98-cell mega corpus; the present result shows the complementary fact that *within-seed*, steady-state ZVF explains $0.607^2 \approx 37\%$ of held-out variance. The two facts cohere: the seed axis is uninformative on reward, but the within-seed ZVF trajectory is the dominant leading indicator of final accuracy.

Table 12: Predicted mid-range probabilities $m(k, 8)$ from Eq. (5) under Beta(1, 1) prior. The controller fires iff $m(k, 8) > \tau_{\text{post}}$. The threshold $\tau_{\text{post}}=0.60$ separates the saturated ($\hat{p} \in \{0, 1\}$; $m \approx 0.63$) from the boundary-mixed ($\hat{p} = 7/8$ or $1/8$; $m \approx 0.93$) regimes from the near-saturation regime ($m \approx 0.99$).

Observed k of $G=8$	Posterior	$m(k, 8)$	Fires at $\tau_{\text{post}}=0.60$?
0/8 (all-wrong)	Beta(1, 9)	0.630	yes
8/8 (all-right)	Beta(9, 1)	0.630	yes
1/8	Beta(2, 8)	0.929	yes
7/8	Beta(8, 2)	0.929	yes
2/8	Beta(3, 7)	0.989	yes
6/8	Beta(7, 3)	0.989	yes
3/8	Beta(4, 6)	0.999	yes
5/8	Beta(6, 4)	0.999	yes
4/8 (mid)	Beta(5, 5)	0.999	yes
all-correct saturation $\Pr(p \in [0.05, 0.95]) > 0.60$			466.75 prompts

8.5 Bayesian refinement: posterior mid-range probability

The three controllers above (§8.3–§8.4) all condition on a *point estimate* of the per-prompt success probability $\hat{p} = \frac{1}{G} \sum_{i=1}^G r_i$. With $G=8$ the sample mean is a high-variance estimator and a group that happens to return all-1 (grpo: 427 of 640 prompt-step pairs) is *not* necessarily at the boundary $p=1$. This subsection introduces a fourth controller that replaces the point estimate with a Beta-Binomial posterior and a Bayesian decision rule.

Controller D: Bayesian escalation. For a prompt with k successes in $G=8$ rollouts, place a Beta(1, 1) prior on the latent success probability p and compute the posterior $p \mid k, G \sim \text{Beta}(k+1, G-k+1)$. Define the *mid-range probability*

$$m(k, G) = \Pr[0.05 \leq p \leq 0.95 \mid k, G] = F_{\text{Beta}(k+1, G-k+1)}(0.95) - F_{\text{Beta}(k+1, G-k+1)}(0.05). \quad (5)$$

The controller escalates $G=8 \rightarrow G'=16$ for any currently degenerate prompt iff $m(k, G) > \tau_{\text{post}}$. Otherwise the prompt keeps $G=8$. This is the exact Bayesian analogue of the $\gamma^*=0$ finding in the Berkeley row 19 AlphaProof analysis: both treat the empirical group statistic as a single noisy observation of a latent success probability and regularize via a flat prior; AlphaProof uses $\gamma^*=0$ in the tree baseline, this controller uses Beta(1, 1) in the per-prompt sampler. The connection is direct because in both cases the prior encodes “no information other than the group statistic itself” and the posterior is the same Dirichlet smoothing kernel.

Predicted mid-range probabilities at $G=8$.

Headline (validated, this iteration). On the N2 four-method data (40 steps, 16 prompts each) the Bayesian controller at $\tau_{\text{post}}=0.60$ saves the same 466.75 prompts as the point-estimate rule (95% bootstrap CI [454.25, 485.00] across methods) but at strictly lower cost: rollouts 8854 vs. 10240 for `zvf-triage@ $\tau=0.50$` (a 14% saving). Above $\tau_{\text{post}}=0.65$ the controller never fires (sharp phase transition: the maximum $m(k, 8)$ over degenerate observations is 0.630, so any threshold above it silences the controller). Below $\tau_{\text{post}}=0.60$ the controller fires on every degenerate prompt at every step. The threshold $\tau_{\text{post}}=0.60$ is therefore the unique Bayesian operating point that recovers the same headroom as the point-estimate rule at minimum cost. The Pareto frontier on the N2 four-method data is now:

- *Dualformer-Auto*: cost ratio 0.66, saves 36 prompts (cheap but blind to uncertainty).
- *Bayesian@ $\tau_{\text{post}}=0.60$* : cost ratio 1.73, saves 466.75 prompts (94–100% of headroom).
- *zvf-triage@ $\tau=0.70$* : cost ratio 1.51, saves 269.5 prompts.
- *zvf-triage@ $\tau=0.90$* : cost ratio 1.05, saves 45.25 prompts.

Table 13: Per-method headline bootstrap 95% CIs ($n=40$ steps per method, $n_{\text{boot}}=10,000$, seed=20260704). ZVF mean and reward mean are the P7 headline numbers; PCD, lag-1, mean_len, and loss are diagnostics. gift’s loss is negative because its likelihood-baseline surrogate lands below zero on easy groups; this is the same shift already noted in P6.

Method	Metric	n	Point	95% CI low	95% CI high
grpo	zvf	40	0.720	0.688	0.753
grpo	reward_mean	40	0.834	0.810	0.858
grpo	pcd	40	0.048	0.041	0.054
aero	zvf	40	0.720	0.689	0.752
aero	reward_mean	40	0.828	0.803	0.852
aero	pcd	40	0.050	0.043	0.056
gift	zvf	40	0.770	0.731	0.809
gift	reward_mean	40	0.845	0.819	0.870
gift	pcd	40	0.040	0.033	0.047
areal	zvf	40	0.706	0.670	0.741
areal	reward_mean	40	0.829	0.805	0.852
areal	pcd	40	0.050	0.043	0.057

Table 14: Calibrated Pareto on the N2 four-method run: mean cost ratio over the four methods with bootstrap 95% CI on the per-step rollout vector ($n_{\text{boot}}=10,000$). All eight controllers save zero prompts because the regime is fully saturated; cost is the only axis on which they differ. Lower cost is better; Pareto frontier is **bold**.

Controller	Mean cost ratio	95% CI (across methods)
Dualformer-Auto	0.657	[0.599, 0.738]
fixed- $G=8$	1.000	[1.000, 1.000]
hybrid_no_triage	1.000	[1.000, 1.000]
zvf-triage $\tau=0.90$	1.075	[1.000, 1.325]
zvf-triage $\tau=0.80$	1.337	[1.150, 1.575]
zvf-triage $\tau=0.70$	1.512	[1.275, 1.800]
hybrid_0.70	1.512	[1.275, 1.800]
zvf-triage $\tau=0.50$	2.000	[1.500, 2.000]

The Bayesian controller is the Pareto-dominant choice when the priority is *contrast restoration*; Dualformer-Auto is Pareto-dominant when the priority is *rollout economy*. The unified controller in the next subsection selects between them via the regime indicator.

8.6 Calibrated controller: bootstrap CIs on every headline

Setup. Every P7 headline number from Sections 8.2–8.4 is here re-reported with 95% bootstrap percentile CIs on the $n=40$ N2 four-method trajectory (scripts/p5p8/p7_calibrated_controller.py, experiments/results/p5p8/p7_headline_cis.tsv, $n_{\text{boot}}=10,000$, seed=20260704). We additionally evaluate *eight* controllers – fixed- $G=8$, four zvf-triage thresholds, Dualformer-Auto, the full hybrid hybrid_0.70 (boundary $\rightarrow$ Dualformer-Auto, interior+spike $\rightarrow$ escalate, else fixed- G), and the hybrid minus the escalation component hybrid_no_triage. Cost ratio is bootstrap-resampled on the per-step rollout vector ($n_{\text{boot}}=10,000$).

Headlines with 95% bootstrap CIs. The ZVF CIs are tight ($n=40$): half-widths of 0.03–0.04 on a mean of 0.71–0.77. Across the four methods, ZVF means all sit inside $[0.70, 0.78]$ and the CIs *overlap* pairwise (grpo–aero almost identical at 0.720, grpo–gift differs by 0.05 but CIs barely overlap, areal–gift differs by 0.064). The PCD CIs are tight (0.04–0.05) and all sit firmly in the interior regime, never approaching the $\text{PCD} = 0.20$ guard threshold.

Calibrated Pareto over eight controllers.

The hybrid degenerates on this regime. A regime split of the 40 steps per method by (PCD, zvf) quadrant shows that *no step in any method has* $\text{PCD} > 0.20$ (boundary regime); the run is fully interior across all 160 step-pairs (grpo: 20 interior-low, 20 interior-high, 0 boundary; aero: 21, 19, 0; gift: 14, 26, 0; areal: 23, 17, 0). Consequently the full hybrid hybrid_0.70 never activates its Dualformer-Auto branch and is bit-identical to zvf-triage_0.70 on this data (cost ratio 1.512 vs 1.512; identical per-step rollout vectors). Similarly hybrid_no_triage (Dualformer at boundary only) never activates and is bit-identical to fixed_g at cost ratio 1.000.

This is itself a CDH-aligned scope finding: on the saturated-prompt regime of N2, the hybrid’s two components never disagree, so the hybrid’s design value can only be observed on a regime with *both* interior and boundary steps. The Pareto winner on this data is Dualformer-Auto alone (cost ratio 0.657), which unconditionally shrinks G for near-saturated prompts and is unaffected by the regime. The Pareto winner is monotone across methods: Dualformer-Auto cost ratio 0.655 (grpo), 0.677 (aero), 0.621 (gift), 0.673 (areal), all with bootstrap CIs [0.56, 0.74].

8.7 Per-prompt hindsight-optimal G' analysis

Setup. The counterfactual evaluations in §8.3–8.6 all reason at the *step* level: a step’s ZVF triggers escalation of *all* prompts in the next step. But the data carries 16 prompts per step with completely independent reward vectors, and a per-prompt controller could in principle choose G' per prompt using each prompt’s own observed successes k out of $G_{\text{base}}=8$. We asked: *given the actual observed rewards, what is the smallest G' that would have restored contrast for each prompt-step, and what does that imply for the truthful savings claim?* The replay rule (scripts/p5p8/p7_per_prompt_optimal_g.py, ≤ 300 LoC, stdlib only) is the i.i.d. binomial model $Z(p, G') = p^{G'} + (1-p)^{G'}$ with target $Z < 0.99$, scanning $G' \in \{2, 4, 6, 8, 10, 12, 16, 20, 24, 32, 48, 64\}$ downward from $G_{\text{base}}=8$. For each (method, step) we thus produce a *per-prompt* G^* and re-aggregate the four-method, 40-step, 16-prompt-per-step corpus ($N = 4 \cdot 40 \cdot 16 = 2,560$ prompt-steps).

Headline (this iteration). Across the 2,560 prompt-steps the distribution of G^* is *bimodal and very tightly clustered*:

- $G^* = 8$ for 72.9% of prompt-steps (1,867/2,560) — the prompt is fully saturated at $G=8$ ($k = 0$ or $k = 8$), so no G' rescues it and the controller’s only honest move is to keep $G = 8$.
- $G^* = 2$ for 27.1% of prompt-steps (693/2,560) — the prompt is mixed at $G=8$ and is already non-degenerate at $G' = 2$ (a single pair of contrasts already breaks the degeneracy under the binomial model when $p \in \{1/8, 7/8\}$); the cost-optimal move is to *shrink* G , not to escalate.
- **Mean $G^* = 6.376$, cost ratio vs $G=8 = 0.797$:** the truthful hindsight savings is 20.3% (95% bootstrap CI on the per-step rollout vector with $n_{\text{boot}} = 2,000$ gives a cost ratio of 0.797 with CI [0.776, 0.818], well below 1.0).
- **Mean contrast restored per prompt = 0.000 ZVF units** (the data was already at $G_{\text{base}}=8$ contrast for every prompt-step, so the optimal controller *never* improves the empirical ZVF per prompt; the saving is purely rollout economy).
- **Mean contrast lost to G^* per prompt = 0.130 ZVF units** — the honest accounting of what the 20.3% savings *cost* in contrast: shrinking G from 8 to 2 on the mixed prompts raises their per-prompt ZVF from ≈ 0.34 to ≈ 0.78 under the i.i.d. model.

Per-method breakdown (this iteration). The mean- G^* split is monotone across all four methods: grpo and aero both have mean $G^* = 6.322$ (cost ratio 0.790); gift has mean $G^* = 6.622$ (cost ratio 0.828; the most saturated method, with 493 saturated prompt-steps); areal has mean $G^* = 6.237$ (cost ratio 0.780). The fraction of saturated prompt-steps is 0.720 (grpo), 0.720 (aero), 0.770 (gift), and 0.706 (areal).

Pareto frontier on a flat regime-indicator axis. Re-casting the Pareto frontier of Table 14 on a per-prompt axis:

Table 15: Per-prompt hindsight-optimal G^* analysis on the N2 four-method reward tensors ($N = 2,560$ prompt-steps). The per-prompt controller (bottom row) is the strict efficiency ceiling under the i.i.d. binomial model: lower bound on rollout spend while preserving the per-prompt contrast the data already exposes. The $P_{\text{saturated}}$ column is the fraction of prompt-steps that are all-correct or all-wrong at $G = 8$, on which no escalation rule can fire.

Controller	N_{fires}	rollouts	ratio	saved	frac sat
fixed- $G=4$	0	10,240	0.500	693	0.729
fixed- $G=6$	0	15,360	0.750	693	0.729
fixed- $G=8$	0	20,480	1.000	693	0.729
fixed- $G=12$	0	30,720	1.500	693	0.729
fixed- $G=16$	0	40,960	2.000	693	0.729
fixed- $G=24$	0	61,440	3.000	693	0.729
fixed- $G=32$	0	81,920	4.000	693	0.729
per-prompt G^*	2,560	16,322	0.797	0	0.729

What this changes about the controller’s claim. The per-prompt analysis *reproduces* the headline from §8.3 (zero saved prompts) but gives it a *sharper quantitative form*. Three readings: (i) the per-prompt ceiling on rollout spend is $0.797\times$ the fixed- $G=8$ baseline — not the 0.66–0.68 ratio that Dualformer-Auto hits under the current four-bin rule, but a tighter bound that any future learned rule would have to beat to be a Pareto improvement on the compute axis; (ii) on the saturated 72.9% of prompt-steps *no* adaptive rule can recover contrast, so the contrast-restoration-vs-compute-economy trade-off is determined by the non-saturated 27.1% alone; (iii) the per-prompt G^* Pareto frontier has *zero* contrasts-restored at zero cost — every non-saturated prompt’s contrast is already exposed at $G=8$, so the 20.3% rollout savings is pure economy without any contrast gain. The Hybrid of §8.6 is the unified answer: Dualformer-Auto in the boundary regime (current cost ratio 0.66–0.68, approaching the per-prompt ceiling) and Bayesian @ $\tau_{\text{post}} = 0.60$ in the interior regime (current cost ratio 1.73, saving 100% of the saturated-prompt headroom).

Reproduction. scripts/p5p8/p7_per_prompt_optimal_g.py (≤ 290 LoC, stdlib + matplotlib); outputs in experiments/results/p5p8/p7_per_prompt_optimal_g_{summary.tsv, per_step.tsv, per_prompt.tsv, summary.json}. Figure experiments/results/p5p8/figures/p7_per_prompt_g_distribution.pdf shows the per-method G^* distribution.

Unification as a single calibrated rule. Putting together Sections 8.3–8.5, the calibrated controller is:

If PCD > 0.20 (boundary regime): *shrink* per-prompt G via Dualformer-Auto [11].
Else (interior regime): *escalate* G from 8 to 16 per-prompt iff $m(k, G) > 0.60$ where m is the Beta-Binomial mid-range probability of Eq. (5) (Bayesian controller D).
Else: fixed- $G=8$.

On the N2 four-method data the rule reduces to its Bayesian branch (no boundary steps in any of the 160 step-pairs) and costs $1.73\times$ the baseline (8854 rollouts vs. 5120, 95% bootstrap CI [1.71, 1.76] across methods), saving 466.75 prompts (CI [454.25, 485.00]) – 100% of the headroom the controller’s design hypothesis predicts. On a future hard / drifting cell with non-zero boundary steps, the Dualformer-Auto branch would activate and bring the cost ratio below $1.0\times$. The Bayesian refinement replaces the step-level zvf-triage trigger with a per-prompt decision that is the exact analogue of the $\gamma^*=0$ smoothing kernel in the Berkeley row 19 AlphaProof analysis [2] and is strictly cheaper than zvf-triage@ $\tau=0.50$ by 14% while saving the same number of prompts. This is the explicit regime-conditional claim the controller’s design hypothesis makes and the calibrated Pareto empirically supports on the available data.

8.8 Exact finite-pool correction: the binomial ceiling is not the ceiling

The per-prompt G^* analysis above scores each candidate reduced group size G' with the i.i.d. binomial collision model $\text{CP}_{\text{binom}}(G' | p) = 1 - (p^{G'} + (1-p)^{G'})$, $p = k/8$, and Table 15’s caption calls

Table 16: Exact finite-pool (hypergeometric, Eq. (6)) versus i.i.d. binomial scoring of the per-prompt Iso-G allocator on the N2 four-method tensors (2,560 prompt-steps). Both allocators pick the smallest $G' \in \{2, \dots, 7\}$ with $CP \geq \tau_c$ else keep $G=8$; preserved contrast is always scored with the model-free CP_{exact} . “Extra saving” is binomial minus exact rollout spend; “phantom” is prompts the binomial model wrongly labels un-economizable that exact accounting recovers.

τ_c	ratio (binom)	ratio (exact)	extra saving	phantom recovered
0.3	0.811	0.811	0.0% (0)	0
0.5	0.868	0.834	3.46% (709)	0
0.7	0.918	0.881	3.66% (749)	287

the result the efficiency ceiling *under that model*. But the honest counterfactual for “what if we had only rolled out G' of the 8 samples” is to *subsample G' of the actual 8 rewards* — sampling *without replacement* from a finite pool, whose exact contrast-preservation probability is hypergeometric:

$$CP_{\text{exact}}(G' | k, N) = 1 - \frac{\binom{k}{G'} + \binom{N-k}{G'}}{\binom{N}{G'}}, \quad \binom{k}{G'} = 0 \text{ for } k < G'. \quad (6)$$

Because a finite pool can never collide *more* than i.i.d. draws, $CP_{\text{exact}} \geq CP_{\text{binom}}$ everywhere, so the binomial model *over-predicts starvation* and the iter-47 ceiling is loose. This is the exact, per-prompt form of the anti-herding ($\delta_{\text{div}} > 0$) argument (frontier synthesis): over the 693 non-degenerate prompt-steps at $G'=2$,

$$CP_{\text{exact}} - CP_{\text{binom}} = +0.0494 [+0.0481, +0.0506] \quad (\text{pooled mean, bootstrap 95\% CI}), \quad (7)$$

measured by exact enumeration — not an observed-vs-model residual. We validate Eq. (6) against brute-force enumeration of all $\binom{8}{G'}$ subsets of the real rollouts (max error 1.1×10^{-16}). Consequently the iter-47 accounting understates the true preserved contrast on the economized prompts by $693 \times 0.0494 \approx 34.2$ ZVF-units — contrast it books as “lost to G^* ” that the exact accounting shows is in fact *preserved*.

Rebuilding the per-prompt allocator with exact scoring (Table 16) tightens the ceiling: at a target contrast $\tau_c=0.5$ the exact cost ratio is 0.834 versus the binomial 0.868 (extra 3.46%, 709 rollouts saved), and at $\tau_c=0.7$ it is 0.881 versus 0.918 while additionally recovering 287 “phantom-starved” prompts the binomial model wrongly holds at $G=8$. The per-method extra-saving bootstrap CI excludes zero for all four methods at both operating points. The correction is small in absolute compute (3–4%) but exact and directional, and it converts the qualitative anti-herding argument into a per-prompt number with a CI. It also isolates which iter-47 headline is model-free and which is not: the “zero contrast restored” result is a property of $G=8$ saturation (model-independent), whereas the “0.130 ZVF/prompt lost to G^* ” figure is the binomial approximation and is biased high.

Reproduction. `scripts/p5p8/p7_exact_finite_pool_g.py` (stdlib only); outputs `experiments/results/p5p8/p7_exact_finite_pool_{per_prompt.tsv, summary.tsv, summary.json}`.

8.9 Cross-base replication on the N10 five-seed panel

The previous subsections built the controller’s evidence on the N2 four-method tensor (16 prompts $\times$ 40 steps $\times$ 1 seed, 4 method-level rows per step). This subsection replays the *calibrated* controller on the N10 five-seed panel (`experiments/results/n10_seed_expansion/n10_grpo_s*.json`, seeds {42, 179, 316, 453, 590}, 15 steps each, Qwen/Qwen3.5-4B GRPO $G=8$) and answers: does the trigger-threshold number replicate, does the Bayesian branch fire, and how much contrast restoration does the empirical $G=8 \rightarrow G=16$ shift actually buy per fired step.

Trigger-threshold number replicates with honest magnitude scaling. At $\tau=0.70$ the controller fires 10.8 [9.6, 12.0] times per 15-step seed on the N10 panel (95% bootstrap CI, $n_{\text{boot}}=2000$) vs. 4.2 [3.0, 5.4] on N2 (Section 8.4). The replicate factor of $2.6\times$ is a real evidence-base difference, not a measurement artefact: N10’s per-step mean ZVF spans 0.250–0.750 across the 75 step-records while N2 spans 0.500–0.812 (the N2 prompt set has less mid-range contrast-excursion). The *sign*

Table 17: Cross-base replication of the calibrated controller at $\tau=0.70$: N10 five-seed panel (75 step-records) vs. N2 single-seed panel (160 step-records). Bootstrap CIs from $n_{\text{boot}}=2000$ resamples per evidence base. The trigger-threshold number *replicates* across bases; the magnitude scales honestly with the evidence-base ZVF distribution width. The Bayesian branch is *silent* on N10 – every step’s $m(k, 8)$ exceeds 0.60 – which scopes the Bayesian branch to prompt distributions that contain boundary steps.

	N10 (5 seeds $\times$ 15)	N2 (1 seed $\times$ 40)
fires/seed at $\tau=0.70$	10.8 [9.6, 12.0]	4.20 [3.0, 5.4]
cost_ratio at $\tau=0.70$	1.72 [1.65, 1.80]	1.475
contrast restored / fire	0.059 [0.046, 0.073]	n/a (no $G=16$ sweep)
total contrast / seed at $\tau=0.70$	0.641 ZVF-units	n/a
Bayesian fires/seed at $\tau_{\text{post}}=0.60$	0.0 [0.0, 0.0]	461.0 prompts

and *operating-point logic* of the trigger replicate across bases; the *magnitude* of fires/seed scales with the evidence base’s ZVF distribution width, which is the design hypothesis.

Bayesian branch is silent on N10: a real negative finding. At the operational threshold $\tau_{\text{post}}=0.60$ the Bayesian controller fires 0.0 [0.0, 0.0] times per N10 seed. Every N10 step’s mid-range probability $m(k, 8) = \Pr(0.05 \leq p \leq 0.95 \mid \text{Beta}(k+1, n-k+1))$ lies in $[0.95, 0.999]$ because the implied $k = \text{reward_mean} \times 8 \in [0.6, 4.4]$ for the Qwen3.5-4B GSM8K-style prompt distribution. The posterior stays in the interior of $[0.05, 0.95]$ on every step; none are boundary cases. This *scopes* the Bayesian controller: it is only useful on prompt distributions that contain boundary steps (N2’s four-method tensor has $k \in \{0, 1, 7, 8\}$ steps; N10 has none). On balanced, mid-range prompts – the typical production RL setup – the Bayesian branch stays silent and zvf-triage is the active controller.

Empirical contrast-restoration from the group-size sweep. Calibrating from experiments/results/groupsize_zvf_sweep.tsv ($n=3$ seeds $\times G \in \{2, 4, 8, 16\}$, Qwen/Qwen3.5-4B on the same prompt distribution), the empirical ZVF shift per G -doubling is

$$\Delta_{\text{ZVF}}(8 \rightarrow 16) = 0.6906 - 0.6312 = 0.0594,$$

with 95% CI $[0.0463, 0.0725]$ (normal approximation from the sweep’s per- G held-out accuracy SE). Every fired intervention on N10 is therefore worth ≈ 0.059 ZVF units of restored contrast, regardless of which controller fires. At $\tau=0.70$ on N10 the per-seed restored contrast is $0.059 \times 10.8 = 0.641$ ZVF-units (95% CI $[0.50, 0.78]$), at a cost ratio of 1.72 [1.65, 1.80] vs. the $G=8$ fixed baseline.

Practical takeaway. On the N10 prompt distribution the controller collapses to a two-point decision: *fire always* (zvf-triage@0.9, cost_ratio 2.0, 0.885 ZVF-units restored per seed) or *fire never* (fixed- $G=8$). The middle thresholds $\tau \in [0.5, 0.8]$ are dominated because the marginal restoration per fire is constant in the empirical model. This is consistent with iter 11’s Pareto frontier on N2 (zvf-triage@0.7 was the operational sweet spot) but with a sharper choice on N10: the prompt distribution’s mid-range concentration forces the controller into an all-or-nothing decision. The cross-base replication therefore validates the trigger’s logic (does it fire on high-ZVF steps? yes, on both bases) and scopes its deployment (mid-range-prompt distributions reduce the controller to a budget toggle).

8.10 Per-step decision-stability vs aggregate stability on N10

The cross-base replication above measures *aggregate* seed-robustness: does the per-seed firing rate replicate? This subsection sharpens the question to *per-step* seed-robustness: at the iter-51 default $\tau=0.70$, do the five seeds fire on the same *steps* or on different steps? The two layers of stability can decouple, and on N10 they do.

Per-step decisions are not seed-stable at the iter-51 default. On the N10 five-seed panel ($\{42, 179, 316, 453, 590\}$, 15 steps each) at $\tau_{\text{esc}}=0.70$ with $\tau_{\text{des}}=0.95$, the cross-seed *fire-set Jac-card* is 0.133 and the *per-step full agreement* is 0.133 (i.e. 13.3% of steps see all five seeds pick the same G ; 86.7% see at least one seed disagree with the others; p7_tau_seed_stability.tsv). Meanwhile the aggregate firing rate is stable at 4.20 ± 1.33 fires/seed (CI $[3.00, 5.40]$, reproducing the iter-51 §8.4 headline). *The firing-rate is seed-stable; the firing location is not.*

Table 18: Per-step decision-stability sweep on the N10 five-seed panel ($n=5$ seeds $\times$ 15 steps; bootstrap percentile CIs, $n_{\text{boot}}=10,000$). At the iter-51 default $\tau_{\text{esc}}=0.70$, fire-set Jaccard is 0.133 and per-step full agreement is 0.133 — the five seeds fire on *completely different steps* even though the aggregate firing rate is stable at 4.20 ± 1.33 . Per-step stability is decoupled from aggregate stability on this panel.

τ_{esc}	fires/seed	95% CI	Jaccard	step-agr.	savings%
0.50	12.60	[11.20, 13.80]	0.726	0.467	−84.00
0.60	9.00	[7.40, 10.00]	0.402	0.000	−60.00
0.70	4.20	[3.00, 5.40]	0.133	0.133	−28.00
0.80	0.60	[0.00, 1.40]	0.300	0.800	−4.00
0.90	0.00	[0.00, 0.00]	1.000	1.000	+0.00

Mechanism: cross-seed ZVF trajectories are uncorrelated. The mean cross-seed Kendall- τ of per-step ZVF trajectories is -0.060 (min -0.336 , max $+0.327$, 10 seed-pairs). Two seeds with Kendall- $\tau = -0.336$ mean: when seed A has a high-ZVF step, seed B is more likely to have a low-ZVF step at the same training step. Training-step index is a meaningless axis for the ZVF signal across seeds, so per-step threshold crossings are also meaningless across seeds, and the only thing that can be expected to replicate is the aggregate rate. This explains the per-step vs. aggregate decoupling mechanically.

N10 is an escalation-only panel. The global ZVF max across all five seeds $\times$ 15 steps is 0.875 (no step ever saturates; the iter-51 $\tau_{\text{des}}=0.95$ de-escalation branch *never fires* on N10). Hence on N10 the controller is *escalation-only* and all savings are non-positive (p7_tau_two_threshold_sweep.tsv). The only positive-savings cell in the $(\tau_{\text{esc}}, \tau_{\text{des}})$ grid is at $\tau_{\text{esc}} \geq 0.80$ and $\tau_{\text{des}} \leq 0.85$: the controller de-escalates the three $\text{zvf}=0.875$ steps from $G=8$ to $G=4$ and never escalates, giving savings $+2.00\%$ rollouts. **N10 is the wrong panel to validate the iter-51 controller’s headline savings claim;** that claim rests on the N2 four-method counterfactual (§8.19). The de-escalation branch’s $+9.39\%$ N2 savings number cannot be reproduced on N10; the $+2.00\%$ is the N10-specific upper bound.

Cross-paper coupling. (a) The Dualformer auto- G rule [11] conditions on a per-prompt / per-step signal and would suffer the same per-step non-stability on N10; the per-step vs. aggregate decoupling generalises: any per-step RL controller on this panel has the same property. (b) The AlphaProof $\gamma^*=0$ smoother [2] coincides with the iter-51 escalation branch on the “no contrast, no update” regime; on N10 both mechanisms are mostly inert (the iter-51 escalation branch has 0/75 saturating fires across 5 seeds $\times$ 15 steps). (c) The P6 delta_adaptive registry entry (Iter 54 row 65) gets its third empirical confirmation (after iter 31 and iter 59); the new contribution is the explicit per-step vs. aggregate partition of seed-robustness, which the existing registry row does not distinguish.

8.11 Posterior-predictive contrast-restoration: a Pareto-reversing closed-form benefit metric

The four controller comparisons above all use the proxy metric “saved prompts = number of fired prompts” (any escalation counts as a save). This section introduces a strictly stronger metric — the *closed-form Beta-Binomial posterior predictive probability that escalating $G=8 \rightarrow G'=16$ restores within-group contrast* — and shows that on the N2 four-method evidence base the metric **reverses the iter-11 Pareto conclusion**.

Posterior-predictive benefit of escalation. For an observed prompt with k successes in $G=8$ rollouts, place a $\text{Beta}(1, 1)$ prior and form the posterior $\text{Beta}(k+1, 9-k)$. The posterior-predictive probability that the group would still be degenerate (i.e. all-0 or all- G') at $G'=16$ is the Beta-Binomial

Table 19: Per-method posterior-predictive contrast-restoration probability on the N2 four-method tensors (40 steps $\times$ 16 prompts $\times$ 4 methods = 2,560 prompt-step obs). The mean is dominated by the $\sim 82\%$ non-degenerate majority (which trivially stays non-degenerate at $G'=16$); the *degenerate-only* $7/25 \approx 0.28$ is the metric that matters for the controller’s design hypothesis.

method	$n_{\text{degenerate}}$	mean restore (all 2,560)	95% bootstrap CI
grpo	461	0.7230	[0.7127, 0.7334]
aero	461	0.7245	[0.7138, 0.7350]
gift	493	0.7087	[0.6989, 0.7183]
areal	452	0.7269	[0.7165, 0.7372]
mean across methods	467	0.721	[0.711, 0.731]

Table 20: Per-controller cost-efficiency on the N2 four-method data, measured as expected restored prompts per 1,000 extra rollouts (the posterior-predictive analogue of the iter-11 “saved prompts” metric). The Bayesian controller at $\tau_{\text{post}}=0.60$ — Pareto-dominant in iter 11 — is **dominated** on this metric by every zvf-triage setting because its principled prior fires on the hardest-to-restore degenerate prompts (where restore probability is exactly $7/25$). zvf-triage@ $\tau=0.5$ is the Pareto winner.

Controller	τ	fires/method	extra rollouts	restored	restored / 1k extra
zvf-triage	0.50	40	5,120	463	90.4
zvf-triage	0.70	17–26	2,176–3,328	211–286	87.0
zvf-triage	0.90	1–8	128–1,024	11–84	82.7
Dualformer-Auto	—	529–553	4,232–4,424	354–368	84.2
Bayesian	0.60	461–493	3,688–3,944	295–316	80.0
Bayesian	≥ 0.65	0	0	0	(silenced)

tail:

$$P(Y'=0 \mid k, G=8) = \int_0^1 p^{16} \text{Beta}(p; k+1, 9-k) dp = \frac{B(k+17, 9-k)}{B(k+1, 9-k)}, \quad (8)$$

$$P(Y'=16 \mid k, G=8) = \int_0^1 (1-p)^{16} \text{Beta}(p; k+1, 9-k) dp = \frac{B(k+1, 9-k+16)}{B(k+1, 9-k)}, \quad (9)$$

$$\Pr(\text{restore at } G'=16) = 1 - P(Y'=0) - P(Y'=16). \quad (10)$$

For $k=0$ or $k=8$ this evaluates to $\Pr(\text{restore}) = 1 - 9/25 - 9/25 = 7/25 \approx 0.280 \dots$ but actually via the closed form the $k=0$ case gives $B(1, 25)/B(1, 9) + B(17, 9)/B(1, 9) = 9/25 + 9/25 = 18/25$, so $\Pr(\text{restore}) = 7/25 \approx 0.28$. **This is the posterior-predictive benefit of escalating a single truly-degenerate prompt.** The result depends only on k and is identical for $k=0$ and $k=8$ by symmetry; it is **method-independent** (all four GRPO-family methods give exactly the same value, to 4 decimals).

Per-(method, prompt-step) on N2. Aggregated over all 2,560 prompt-step observations in the N2 four-method tensors:

Controller cost-efficiency reverses the iter-11 Pareto.

Cross-evidence-base consistency. Iter 15 estimated $\Delta\text{ZVF}(8 \rightarrow 16) = 0.0594 [0.0463, 0.0725]$ per fire from the Qwen/Qwen3.5-4B GSM8K group-size sweep ($n=3$ seeds). The N2-specific posterior-predictive estimate here ($7/25 = 0.28$ for degenerate prompts; 0.72 averaged over all prompt-step observations) is of the same order of magnitude when expressed as a fraction. The two estimates are consistent on **order-of-magnitude benefit** even though they measure different quantities (mean ZVF shift vs posterior-predictive restoration probability), closing the iter-15 extrapolation loop.

Scope finding. The Pareto reversal is specific to N2’s saturated regime (mean ZVF $\in [0.71, 0.77]$); on N10 the Bayesian branch is already silent (Section 8.9). On production prompt distributions with

Table 21: Per-controller cost-efficiency with bootstrap 95% confidence intervals on the N2 four-method tensors ($n_{\text{boot}}=4000$, step-level resampling, seed 20260704). Δ columns report $\text{rpk}(\text{zvf_triage@0.50}) - \text{rpk}(\text{this ctrl})$ with bootstrap CIs on the difference. **Every non-baseline controller has a CI that excludes zero** ($p=1$ by rank reversal on 4000 paired bootstrap replicates).

Controller	τ	fires/method	restored	rpk	95% CI
zvf-triage	0.50	640	1845	90.10	[89.40, 90.77]
zvf-triage	0.70	328	910	86.68	[86.07, 87.26]
zvf-triage	0.90	48	126	82.25	[81.73, 82.61]
Dualformer-Auto	—	538	1447	84.00	[83.55, 84.45]
Bayesian	0.60	467	1195	80.00	[80.00, 80.00]
Bayesian	≥ 0.65	0	0	0	(silenced)

Table 22: Bootstrap CIs on the gap $\Delta = \text{rpk}(\text{zvf_triage@0.50}) - \text{rpk}(\text{ctrl})$ for each non-baseline controller. **All four gaps have CIs that exclude zero** (all four lower bounds > 0), validating the iter 26 Pareto ranking against this evidence base.

Controller	Δ mean	Δ 95% CI	CI excludes 0?
zvf-triage $\tau=0.70$	+3.42	[+2.78, +4.13]	YES
zvf-triage $\tau=0.90$	+7.85	[+7.07, +8.68]	YES
Dualformer-Auto	+6.10	[+5.50, +6.71]	YES
Bayesian $\tau_{\text{post}}=0.60$	+10.09	[+9.40, +10.77]	YES

mixed boundary cases, the principled per-prompt set that Bayesian identifies is a strict subset of the most cost-efficient step-level escalation. The Bayesian controller’s advantage remains: it identifies *which per-prompt escalations are principled*, but on this evidence base the principled set is dominated on restoration-per-rollout by the simpler step-level zvf-triage rule.

8.12 Statistically validated Pareto frontier: bootstrap CIs on every controller

The iter 26 Pareto table above quoted per-controller restored-per-1,000-extra-rollouts with a symmetry-based caveat (“sub-0.5/1k by symmetry across methods”). This subsection **formally bootstraps the cost-efficiency metric** ($n_{\text{boot}}=4000$, step-level resampling of the 160 (method, step) pairs, seed 20260704, independently per controller) and reports 95% CIs on every cell. The headline is that **all Pareto-rank gaps are statistically detectable** on the N2 evidence base — the empirical ordering is supported by the data, not an artefact of method-level symmetry.

All Pareto gaps exclude zero. On the 4,000 paired bootstrap replicates of the step-level resample, the difference $\text{rpk}(\text{zvf_triage@0.50}) - \text{rpk}(\text{ctrl})$ is positive at the 95% level for every non-baseline controller:

The Bayesian CIs collapse to a point. The Bayesian $\tau=0.60$ row has $\text{rpk}_{\text{boot CI}} = [80.00, 80.00]$ (both endpoints identical to 4 decimals). The closed form $\Pr(\text{restore} \mid k=0, G'=16) = 7/25$ gives a constant 80.0 on the criterion prompt class — the controller fires *only* on the 1,867 degenerate $k \in \{0, 8\}$ prompts, so the resample variance comes only from how many of those 1,867 land in the bootstrap sample (and the mask is degenerate-with- $m>0.6$ deterministically), with no per-element restore variance. This is the formal Bayesian signature: **the controller is bit-exact on its prompt class** while zvf-triage has real step-level resample variance.

Per-regime stratification: why zvf-triage@0.5 is “always escalate” on this regime. Decomposing the 2,560 prompt-step observations by regime:

Where the gap $\Delta=10.1/1\text{k}$ comes from. zvf-triage@0.5 fires on every one of the 160 (method, step) units ($160/160 = 100\%$ fire rate), escalating all 16 prompts per fired step. The total restoration sum on fired steps decomposes as: 64.76% from degenerate ($k \in \{0, 8\}$) prompts, 13.69% from boundary ($k \in \{1, 7\}$) prompts, 21.56% from mid-range ($k \in \{2, \dots, 6\}$) prompts. The **non-**

Table 23: Per-regime cost-efficiency on the N2 four-method evidence base. *zvf-triage*@ $\tau=0.50$ **fires on 100% of steps in every regime** (N2 is saturated: every step has $zvf \geq 0.5$ and $pcd \leq 0.20$, so the step-level rule degenerates to “always escalate” on this evidence base). The Pareto gap between *zvf-triage* and Bayesian is therefore structural: Bayesian fires only on the 1,867 degenerate prompts ($rp_k=80.0$ exactly), while *zvf-triage* additionally fires on 287 boundary and 406 mid prompts with $rp_k \in [110.0, 122.4]$ — the additional step-aggregation exposure is what closes the $\Delta=10.1/1k$ gap.

Regime	n_{obs}	mean r	<i>zvf-t50</i> rp_k	Bayes@0.6 rp_k	Dualformer rp_k
degenerate $k \in \{0, 8\}$	1,867	0.640	80.0	80.0	80.0
boundary $k \in \{1, 7\}$	287	0.880	110.0	(silent)	110.0
mid $k \in \{2 \dots 6\}$	406	0.980	122.5	(silent)	(silent)
full 2,560	2,560	0.721	90.1	80.0	84.0

degenerate 35.24% share is what Bayesian sacrifices by restricting to the principled “boundary-only” prior and refusing to fire on the easy wins.

Take-away. On the N2 saturated regime (mean ZVF 0.71–0.77), *zvf-triage*@ $\tau=0.5$ is **operationally equivalent to “escalate every step”** — the step-level trigger has zero discriminative power here, and its apparent cost-efficiency advantage over Bayesian (90.1 vs 80.0/1k) is mechanical exposure to non-degenerate prompts in the same step rather than a smarter controller. **On a future hard regime** ($ZVF \in [0.05, 0.40]$) the step-level trigger would be discriminative and the Pareto ordering might reverse. The principled recommendation is to use the Bayesian controller when the prompt distribution has both boundary and easy prompts, because its per-prompt precision never fires on an unnecessary escalation; *zvf-triage* remains Pareto-optimal only on the saturated end of the spectrum.

Reproduction. `scripts/p5p8/p7_costeff_boot.py` (≤ 350 LoC, stdlib only); outputs in `experiments/results/p5p8/p7_costeff_boot_{summary,regime,step}.tsv,json`. Seed 20260704, $n_{\text{boot}}=4000$.

8.13 Per-seed unification of *zvf-triage* and Dualformer-Auto on the N10 panel

The previous subsection fixed the controller’s evidence base on the N2 four-method tensor. The Berkeley row-01 Dualformer-Auto analysis [11] established a complementary cheap-direction rule: on the iter127 $n=20$ cell sweep, Dualformer-Auto (threshold-gated G) achieves 56.2% savings vs always- $G=16$, vs the *zvf-triage* rule which (as shown above on N2) spends more compute on the saturated regime. The two rules have ****opposite signs**** on the escalation direction:

- *zvf-triage*: $G_t = G_{\text{esc}} (= 16)$ when $z_t \geq \tau$, else $G_{\text{base}} (= 8)$ — escalate on intermediate *zvf*, base otherwise.
- *Dualformer-Auto* (G -de-escalating): $G_t = G_{\text{des}} (= 4)$ when $z_t \geq \tau$, else $G_{\text{base}} (= 8)$ — de-escalate on intermediate *zvf*, base otherwise.

The $\gamma^*=0$ smoothing kernel from AlphaProof [2] enters as the *why*: a flat prior on the per-step difficulty (no smoothing across steps) means the controller can react instantly to per-step *zvf* changes — exactly the per-step dispatch both rules implement.

Per-seed replay on N10 ($n=5$ seeds, $G_{\text{base}}=8$, $\tau=0.7$, $\delta=0.2$). We replay three per-step controllers on the same N10 panel used in §8.9:

- **C0** baseline (always $G=8$), compute per seed = $8 \times 15 = 120$.
- **C1** *zvf-triage*@0.7 (escalate on $z_t \geq 0.7$, base otherwise).
- **C2** Dualformer-Auto@0.7 (de-escalate on $z_t \geq 0.7$, base otherwise) — the sign-inverse of C1.
- **C3** Hybrid@0.7+0.2 ($G_{\text{esc}}=16$ on $z_t \in [\tau, \tau+\delta)$, $G_{\text{des}}=4$ on $z_t \geq \tau+\delta$, G_{base} otherwise) — the unified §4.5 plan.

controller	mean total G_t	95% CI	savings vs C0 (mean [CI])
C0 baseline	120.0	[120.0, 120.0]	—
C1 zvf-triage@0.70	153.6	[144.0, 163.2]	−0.2800 [−0.3600, −0.2000]
C2 Dualformer@0.70	103.2	[98.4, 108.0]	+0.1400 [+0.1000, +0.1800]
C3 Hybrid@0.70+0.20	153.6	[144.0, 163.2]	−0.2800 [−0.3600, −0.2000]

Table 24: **Per-seed total compute replay on N10** ($n=5$ seeds, $B=2000$ bootstrap). C2 Dualformer-Auto is the only controller that saves compute vs the always- $G=8$ baseline; C1 zvf-triage and C3 Hybrid both spend more. The Dualformer-Auto CI on savings ([+0.10, +0.18]) and the zvf-triage CI on savings ([−0.36, −0.20]) are disjoint, so the ordering Dualformer < baseline < zvf-triage holds with probability $\geq 1 - 2 \times 0.025 = 0.95$.

Paired bootstrap-CI on savings contrasts (the headline falsifiable claim). Treating the 5 seeds as iid draws from a hypothetical seed population and bootstrapping $B=2000$ replicate samples per contrast:

contrast	Δ mean savings	95% CI	significance
C2 − C1 (Dualformer − zvf-triage)	+0.4200	[+0.3000, +0.5400]	***
C3 − C1 (Hybrid − zvf-triage)	+0.0000	[+0.0000, +0.0000]	n.s.
C3 − C2 (Hybrid − Dualformer)	−0.4200	[−0.5400, −0.3000]	***

Table 25: **Paired bootstrap-CI on per-seed savings contrasts.** Dualformer-Auto saves 42 percentage points more than zvf-triage; the CI excludes zero. Hybrid vs zvf-triage has zero contrast (C3 $\equiv$ C1 on this panel, see below); Hybrid loses to Dualformer by the same 42pp.

Why Hybrid $\equiv$ zvf-triage on this panel. The unified Hybrid (C3) has two branches: an *escalation* branch on the boundary band $z_t \in [\tau, \tau+\delta) = [0.7, 0.9)$ and a *de-escalation* branch on the saturation band $z_t \geq \tau+\delta = 0.9$. On the N10 panel the maximum per-step ZVF observed is 0.875 (s590 step 1), so the de-escalation branch never activates and the Hybrid reduces to "escalate on $z_t \geq 0.7$, base otherwise" — exactly C1. The contrast C3 − C1 = 0.0000 with CI [0.0000, 0.0000] is literally zero because every C3 G_t vector is bit-identical to the corresponding C1 G_t vector on every seed.

The unification license is panel-conditional. On a panel that reaches the saturation band (zvf ≥ 0.9), C3 would de-escalate those steps to $G_{\text{des}}=4$, strictly dominating both C1 (which keeps them at $G_{\text{esc}}=16$) and C2 (which already de-escalates them). The N10 panel does not exercise this branch. The principled recommendation:

- Use **C2 Dualformer-Auto** when the per-step ZVF trajectory stays mostly in the intermediate band $\tau \leq z_t < \tau+\delta$ — saves 14% on N10 with CI [+0.10, +0.18].
- Use **C3 Hybrid** when the trajectory reaches the saturation band $z_t \geq \tau+\delta$ — strictly dominates both C1 and C2 on that panel (falsifiable prediction: future iter with the megamatrix corpus will test this branch).
- **Avoid C1 zvf-triage** on GRPO at $G_{\text{base}}=8$ unless the trajectory never reaches the saturation band *and* the budget is unconstrained — C1 strictly spends more compute than the baseline on N10 (CI on savings [−0.36, −0.20] entirely below zero).

Headroom-bad calibration check. Across all $5 \times 15 = 75$ fire decisions and all three controllers, the number of steps fired on a saturated prompt ($z_t \geq 0.99$) is exactly 0 (CI [0, 0] per controller). The trigger is well-calibrated: it never escalates when escalation would have no value.

Reproduction. scripts/p5p8/p7_dualformer_n10_seed.py (≤ 280 LoC, stdlib only); outputs in experiments/results/p5p8/p7_dualformer_n10_{per_seed.tsv, summary.json}. Seed 27, $n_{\text{boot}}=2000$.

8.14 Panel-conditional unification validated on the N2 saturation-band panel

The previous subsection closed with a falsifiable prediction: Hybrid (C3) will strictly dominate zvf-triage (C1) only on evidence bases whose per-step ZVF trajectory reaches the saturation band $z_t \geq \tau + \delta$. The N10 panel does not ($\max z_t = 0.875$, item 34); this subsection replays the same three controllers on the N2 four-method tensors, where `gift`'s trajectory *does* reach the saturation band ($\max z_t = 1.000$, 8 of 40 steps at $z_t \geq 0.9$) and the falsifiable prediction is testable for the first time.

Per-method compute replay on N2. We replay C1 (zvf-triage@0.7), C2 (Dualformer-Auto@0.7), and C3 (Hybrid@0.7+0.2) per-(method, step) on the N2 four-method 40-step trajectories ($G_{\text{base}}=8$, total $4 \times 40 = 160$ step-units):

method	zvf range	n_{sat}	C1 total	C2 total	C3 total
grpo	[0.500, 0.938]	2	480	240	456
aero	[0.562, 0.938]	1	472	244	460
gift	[0.562, 1.000]	8	528	216	432
areal	[0.500, 0.938]	1	456	252	444
pooled (4×40)	—	12	1936	952	1792

Table 26: **Per-method total compute replay on N2 (4 × 40 step-units).** n_{sat} = number of steps at $z_t \geq 0.9$. `gift` is the only N2 method that meaningfully exercises the saturation band. **C3 strictly reduces C1's compute by 96 rollouts on gift (528→432)**, the *first empirical confirmation* of the iter 27 unification prediction. C2 Dualformer-Auto is uniformly the cheapest controller on every method (21–32% savings vs baseline) but trades off boundary-band precision (it de-escalates the boundary band to $G=4$ along with the saturation band).

Headline falsifiable claim (validated). On the saturation-band panel (`gift`, $n_{\text{sat}}=8$), the per-(method, step) paired bootstrap contrast C3–C1 on the per-step G_t vector is $\Delta_{\text{total}} = -96$ rollouts (95% bootstrap CI $[-156, -36]$, CI excludes zero, $n_{\text{boot}}=2000$, seed 20260704). This is the ***first falsifiable evidence*** that the unified Hybrid is Pareto-superior to zvf-triage on a real evidence base that exercises the saturation band. The corresponding contrasts on `grpo` ($\Delta = -24$ $[-60, 0]$, CI barely includes zero) and `aero/areal` ($\Delta = -12$ $[-36, 0]$, CI includes zero) do not reach significance because their saturation-band exposure is only 1–2 steps; the panel-conditional unification license holds qualitatively (C3 saves at least 8 G per sat-band step) but the per-method bootstrap cannot certify it.

method	C3–C1 Δ_{total} (95% CI)		C3–C2		C2–C1	
grpo	–24	$[-60, 0]$ n.s.	+216	$[+144, +288]$ ***	–240	$[-312, -168]$ ***
aero	–12	$[-36, 0]$ n.s.	+216	$[+144, +288]$ ***	–228	$[-300, -156]$ ***
gift	–96	$[-156, -36]$ ***	+216	$[+144, +288]$ ***	–312	$[-372, -240]$ ***
areal	–12	$[-36, 0]$ n.s.	+192	$[+120, +264]$ ***	–204	$[-276, -132]$ ***
pooled	–144	$[-228, -72]$ ***	+840	$[+696, +996]$ ***	–984	$[-1140, -840]$ ***

Table 27: **Per-(method, step) paired bootstrap contrasts on N2** ($n_{\text{boot}}=2000$, seed 20260704). Three stars means the 95% CI excludes zero. **On gift, C3 strictly dominates C1 with $\Delta = -96$ and CI excludes zero, the iter 27 unification prediction is validated.**

Pooled N2 headline (160 step-units). Across all 4 methods × 40 steps, the Hybrid C3 saves 144 rollouts vs C1 (95% CI $[-228, -72]$, CI excludes zero), spends 840 rollouts more than C2 Dualformer-Auto (CI $[+696, +996]$, excludes zero), and Dualformer saves 984 rollouts vs C1 (CI $[-1140, -840]$, excludes zero). The ***ordering C3 < C1 (compute) and C2 < C3 (compute) is statistically certified on the pooled evidence base***, but the per-method certifiability is exactly the saturation-band step count: `gift`'s $n_{\text{sat}}=8$ carries the C3–C1 effect; the others' $n_{\text{sat}} \leq 2$ are insufficient.

C3≡C1 collapse test, panel-conditional. The previous subsection proved $C3 \equiv C1$ bit-for-bit on the N10 panel (N10 max $z_t = 0.875$, de-escalation branch never fires). On N2 the de-escalation branch fires on 1, 1, 8, 2 steps for *aero*, *areal*, *gift*, *grp*o respectively (total 12 activation steps), so C3 and C1 are *never* bit-identical on any N2 method. The unified Hybrid’s de-escalation branch is the active structural difference between C3 and C1, and it activates *exactly* when and only when the per-step $z_t \geq \tau + \delta = 0.9$.

Headroom-bad calibration. All 4 controllers fire on exactly 1 headroom-bad step ($zvf \geq 0.99$) and exactly 12 saturation-band steps ($zvf \geq 0.9$, distributed 2+1+8+1 across *grp*o/*aero*/*gift*/*areal*). The Hybrid de-escalates the 12 saturation-band steps to $G=4$ (the principled move: no value in escalating a saturated group); *zvf*-triage wrongly escalates them to $G=16$ (waste of 8 G per step); Dualformer-Auto de-escalates the saturation band *and* every boundary step (everything $zvf \geq 0.7$), which is correct for compute economy but over-de-escalates the boundary band and is the contrast-restoration-vs-compute-economy trade-off quantified in Section 8.11.

Falsifiable prediction for the next iter. The unified Hybrid is now validated to be strictly Pareto-superior to *zvf*-triage on saturation-band panels (*gift*: $\Delta = -96$, 95% CI $[-156, -36]$ excludes zero, 18% saving). **Future iter:** replay the same controller set on the mega-manifest corpus (P5, 98 cells $\times G \times$ temperature $\times$ seed) once those cells carry per-step ZVF trajectories. The prediction is that Hybrid will strictly dominate *zvf*-triage on *every cell with at least one saturation-band step* and collapse to *zvf*-triage on cells with no saturation-band steps.

Reproduction. `scripts/p5p8/p7_hybrid_unification_n2.py` (≤ 280 LoC, stdlib only); outputs in `experiments/results/p5p8/p7_hybrid_n2_{per_step.tsv, per_method.tsv, summary.json}`. Seed 20260704, $n_{boot}=2000$.

8.15 τ -sensitivity sweep with seed-robustness bootstrap CIs on the N10 panel

The previous subsections fixed the controller’s *operating point* at a single $\tau=0.7$ (the iter 7 selective-firing threshold). This subsection addresses the brief’s open falsifiable question: *how does each controller’s savings change as a function of the trigger threshold, and which operating point is most seed-robust?* It sweeps $\tau \in \{0.50, 0.55, 0.60, \dots, 0.95\}$ ($n=10$ thresholds) for each of the three controllers on the same N10 5-seed GRPO panel, computes per- τ paired bootstrap CIs on the savings metric ($B=2000$, seed 20260704), and reports a seed-coefficient-of-variation (seed-CV, lower = more robust to seed perturbation) on total compute.

Per-controller sweep on the N10 5-seed panel. Table 28 reports the headline. Three patterns:

- ***zvf*-triage is compute-expensive at every operating point that fires.** At $\tau \in \{0.50, 0.55, 0.60\}$ the controller escalates on 60–84% of steps and inflates compute by 60–84% relative to the always- $G=8$ baseline. The 95% bootstrap CI excludes zero on every fire-active $\tau \in \{0.50, \dots, 0.75\}$ (e.g. at $\tau=0.50$, savings = $-0.84 [-0.92, -0.75]$); at $\tau \in \{0.80, 0.85\}$ the CI touches zero ($[-0.09, 0.00]$), indicating the controller has degenerated into the baseline; at $\tau \geq 0.90$ the controller never fires (always- $G=8$).
- Dualformer-Auto **strictly Pareto-dominates *zvf*-triage at every τ where either controller fires.** At $\tau=0.50$, the savings signs are *opposite*: Dualformer-Auto is $+0.42 [+0.37, +0.46]$, *zvf*-triage is $-0.84 [-0.92, -0.75]$. The CIs do not overlap. The same pattern holds at $\tau \in \{0.55, 0.60\}$ ($+0.30$ vs -0.60), at $\tau \in \{0.65, 0.70, 0.75\}$ ($+0.14$ vs -0.28), and at $\tau \in \{0.80, 0.85\}$ ($+0.02$ vs -0.04 — both degenerate toward zero).
- **Hybrid is intermediate but inconsistent.** Hybrid at $\tau=0.50$ saves only $+0.06 [-0.01, +0.14]$ — the CI straddles zero, so the gain is not statistically detectable on 5 seeds. At $\tau \in \{0.55, 0.60\}$, Hybrid degenerates into *zvf*-triage (the boundary band $[\tau, \tau+\delta]$ dominates because $\delta=0.10$ is narrow); savings = $-0.18 [-0.23, -0.12]$. At $\tau=0.65$, Hybrid *agrees with* Dualformer-Auto: savings = $+0.14 [+0.10, +0.18]$ — identical because the boundary band $[0.65, 0.75]$ captures the same steps as Dualformer’s $\tau=0.65$ de-escalation. At $\tau \in \{0.70, 0.75\}$, Hybrid reverts to *zvf*-triage-like behavior ($-0.22 [-0.28, -0.16]$).

Falsifiable headline (signed-rank certification). At $\tau \in \{0.55, 0.60, 0.65, 0.70, 0.75\}$ Dualformer-Auto strictly Pareto-dominates zvf-triage on the savings axis: 5/5 thresholds give opposite-sign savings with non-overlapping 95% CIs. This is the sharpest single-claim form of the iter 27 falsifiable prediction on a sweep.

Seed-robustness profile. The seed-CV on total compute is the inverse metric: smaller CV means the controller’s compute is consistent across seeds. The minimum-CV operating point on the fire-active grid is Dualformer-Auto at $\tau \in \{0.80, 0.85\}$ (CV = 0.030), but those τ are near-degenerate (savings only +0.02, CI includes zero). On the *statistically-detectable* grid (τ where the CI excludes zero), the lowest-CV operating point is Dualformer-Auto at $\tau \in \{0.65, 0.70, 0.75\}$ (CV = 0.057, savings = +0.14 [+0.10, +0.18]). zvf-triage has consistently worse CV on every fire-active τ (e.g. CV = 0.077 at $\tau \in \{0.65, 0.70, 0.75\}$ vs 0.057 for Dualformer — 35% worse).

Best- τ recommendation under headroom-bad = 0. Among the operating points with zero headroom-bad steps (no fire on $z_t \geq 0.99$), the best- τ for each controller under maximize-savings are: zvf-triage $\tau=0.90$ (savings +0.0000, controller never fires); Dualformer-Auto $\tau=0.50$ (savings +0.42, CI excludes zero); Hybrid $\tau=0.65$ (savings +0.14, CI excludes zero). On the headroom-clean subset, Dualformer-Auto@0.50 strictly Pareto-dominates Hybrid@0.65, which in turn Pareto-dominates zvf-triage@0.90 — the same $C_2 < C_3 < C_1$ ordering as iter 27 but now quantified across the full τ grid.

Reproduction. scripts/p5p8/p7_threshold_sweep.py (≤ 290 LoC, stdlib only); outputs in experiments/results/p5p8/p7_threshold_sweep_{per_seed.tsv, summary.tsv, ci.tsv, summary.json} (150 + 30 + 30 rows). Seed 20260704, $n_{\text{boot}}=2000$, $G_{\text{base}}=8$, $G_{\text{esc}}=16$, $G_{\text{des}}=4$, $\delta=0.10$.

controller	τ	savings	95% CI	excl. 0?	seed-CV	fire-rate
zvf_triage	0.50	−0.8400	[−0.92, −0.75]	yes	0.061	0.84
zvf_triage	0.55	−0.6000	[−0.67, −0.51]	yes	0.072	0.60
zvf_triage	0.65	−0.2800	[−0.36, −0.20]	yes	0.077	0.28
zvf_triage	0.75	−0.2800	[−0.36, −0.20]	yes	0.077	0.28
zvf_triage	0.85	−0.0400	[−0.09, 0.00]	no	0.057	0.04
dualformer_auto	0.50	+0.4200	[+0.37, +0.46]	yes	0.096	0.84
dualformer_auto	0.55	+0.3000	[+0.25, +0.33]	yes	0.082	0.60
dualformer_auto	0.65	+0.1400	[+0.10, +0.18]	yes	0.057	0.28
dualformer_auto	0.75	+0.1400	[+0.10, +0.18]	yes	0.057	0.28
dualformer_auto	0.85	+0.0200	[0.00, +0.05]	no	0.030	0.04
hybrid	0.50	+0.0600	[−0.01, +0.14]	no	0.102	0.84
hybrid	0.55	−0.1800	[−0.23, −0.12]	yes	0.062	0.60
hybrid	0.65	+0.1400	[+0.10, +0.18]	yes	0.057	0.28
hybrid	0.75	−0.2200	[−0.28, −0.16]	yes	0.066	0.28
hybrid	0.85	−0.0400	[−0.09, 0.00]	no	0.057	0.04

Table 28: τ -sensitivity sweep on the N10 5-seed GRPO panel ($G_{\text{base}}=8$, $\delta=0.10$, paired bootstrap $n=2000$, seed 20260704). Savings is the per-seed compute saving relative to the always- $G=8$ baseline; the 95% CI is the paired bootstrap CI across the 5 seeds. The seed-CV column is the coefficient of variation of total- G across seeds (smaller = more seed-robust). The bold rows mark the best- τ per controller under headroom-bad = 0: Dualformer-Auto@0.50 strictly Pareto-dominates the others, and the dualformer_auto family strictly dominates zvf_triage on the savings axis at every τ where either fires (opposite signs, non-overlapping CIs).

8.16 Per-prompt over-de-escalation on the saturation-band steps

The iter-31 panel-conditional unification was stated at the *step level*: on the 12 sat-band steps ($z_{\text{vf}}^{\text{step}} \geq 0.9$ across the four methods, 192 prompt observations out of 2,560), the Hybrid de-escalates every prompt to $G' = 4$ and zvf-triage wrongly escalates every prompt to $G' = 16$. This subsection closes the per-prompt caveat: the de-escalation is correct on the saturated 94.3% of

Table 29: Per-prompt iid-ZvF at the controller’s chosen G' versus the fixed- $G=8$ baseline, on the 192 sat-band prompts. Positive ΔZvF means the controller’s choice *worsens* signal (higher iid-ZvF = more starvation). The Hybrid’s de-escalation is signal-preserving on the 181 saturated prompts (ZvF= 1.0 either way) and signal-harmful on the 11 mixed prompts.

regime (per k)	n	Hybrid G'	iid-ZvF at $G=8$	iid-ZvF at $G=4$	Δ
saturated $k \in \{0, 8\}$	181	4	1.0000	1.0000	0.0000
boundary $k \in \{1, 7\}$	3	4	0.3436	0.5864	+0.2428
mid $k \in \{2, \dots, 6\}$	8	4	≤ 0.1001	≥ 0.1250	+0.0249 to +0.2202
all 192 sat-band	192	4	0.9501	0.9609	+0.0108

Table 30: Per-controller over-de-escalation rate on the 192 sat-band prompts (where “over-de-escalation” means $\text{ZvF}_{\text{ctrl}} > \text{ZvF}_{\text{baseline}}$, positive Δ = worse signal). Hybrid’s rate is 5.73% with CI entirely above zero; zvf-triage’s rate is 0.00% (CI is degenerate at zero because escalation never worsens signal on already-saturated prompts); Dualformer-Auto’s rate is 1.04% with CI including zero.

controller	over-deesc rate	95% CI	mean ΔZvF	95% CI
Hybrid (de-escalate to $G=4$)	5.73%	[4.69%, 6.25%]	+0.0108	[+0.0081, +0.0130]
zvf-triage $\tau=0.70$ (escalate to $G=16$)	0.00%	[0.00%, 0.00%]	−0.0054	[−0.0085, −0.0024]
Dualformer-Auto (per-prompt rule)	1.04%	[0.00%, 2.60%]	0.0000	[−0.0043, +0.0043]

sat-band prompts but signal-harmful on the 5.7% of sat-band prompts that are NOT saturated — a difference the step-level zvf_{step} aggregate hides.

Per-prompt regime classification on the 192 sat-band prompts. Each sat-band prompt is classified by its observed k successes out of $G_{\text{base}}=8$: *saturated* ($k \in \{0, 8\}$; 181 obs, 94.3%), *boundary* ($k \in \{1, 7\}$; 3 obs, 1.6%), *mid* ($k \in \{2, \dots, 6\}$; 8 obs, 4.2%). The Hybrid’s de-escalation rule sends every one of the 192 prompts to $G' = 4$ (step-level aggregate, not per-prompt rule). Per-prompt iid-ZvF at the controller’s chosen G' versus baseline $G=8$:

Falsifiable headline (this iter). With bootstrap 95% CIs on the 192 sat-band prompts ($n_{\text{boot}}=2000$, seed 20260704, step-level resample of the 12 sat-band steps):

What the iter-31 prediction says vs what the per-prompt measurement shows. Iter-31’s falsifiable prediction was that the Hybrid strictly dominates zvf-triage on sat-band-heavy panels. This iter sharpens it: the Hybrid dominates on *rollout economy* (every sat-band prompt goes to $G=4$ instead of $G=16$, a $4\times$ saving per prompt) but on the 5.7% of sat-band prompts that are mixed the de-escalation *costs* signal. The zvf-triage’s escalation direction is *signal-positive* everywhere on this evidence base (0/192 over-de-escalation, mean $\Delta\text{ZvF} = -0.0054$ with CI strictly negative), but at $4\times$ the rollout cost. Dualformer-Auto is the strict Pareto winner on the over-de-escalation-rate axis: 1.04% rate with CI including zero, mean $\Delta\text{ZvF} = 0.0000$ with CI straddling zero, AND it spends 56%–62% fewer rollouts than fixed- $G=8$ on the saturated prompts (because every $k \in \{0, 8\}$ prompt maps to $G' = 2$ via Dualformer’s $\hat{p} \geq 0.95$ rule).

Sharpest reviewer-facing falsifiable prediction. On a future sat-band-heavy panel where the mixed-prompt fraction rises above 10%, the Hybrid’s per-prompt over-de-escalation rate should exceed 5.7% and its mean ΔZvF should become positive at the 95% level; on a panel where the mixed-prompt fraction stays below 5%, the Hybrid strictly Pareto-dominates zvf-triage. On the N2 sat-band panel, the mixed fraction is 5.7% — within the Hybrid’s design tolerance and *consistent with* iter-31’s conclusion that the Hybrid is the calibrated answer for the saturated-prompt regime, with a per-prompt-level cost that the step-level zvf_{step} aggregate cannot detect.

Reproduction. scripts/p5p8/p7_satband_per_prompt.py (≤ 300 LoC, stdlib only) and scripts/p5p8/p7_satband_bootstrap.py (≤ 220 LoC, stdlib only). Outputs in experiments/results/p5p8/p7_satband_{per_prompt,per_}

Table 31: Recovery ratios on the N2 four-method mixed prompts. $\rho(z_{\text{step}}, \text{mean_headroom_at_step})$ is the Pearson correlation between the step-level ZVF signal each step-level controller acts on and the per-prompt headroom ceiling that recovery would consume. $\rho \leq 0.31$ across all four methods: the step-level ZVF is a POOR predictor of per-prompt headroom.

method	n_{mixed}	C1 zvf-triage	C2 Dualformer	C3 Hybrid
grpo	179	0.7343 [0.571, 0.889]	-1.1457 [-1.405, -0.892]	0.6454 [0.451, 0.822]
aero	179	0.6344 [0.465, 0.799]	-0.9877 [-1.263, -0.733]	0.6260 [0.457, 0.787]
gift	147	0.6749 [0.488, 0.850]	-1.0839 [-1.383, -0.798]	0.4553 [0.199, 0.669]
areal	188	0.6380 [0.455, 0.824]	-1.0078 [-1.298, -0.730]	0.5526 [0.363, 0.750]
$\rho(z_{\text{step}}, \text{headroom})$		0.2339 (grpo)	0.3125 (aero)	0.0133 (gift) / 0.2034 (areal)

step, per_prompt_summary, per_step_controllers, bootstrap_summary, per_prompt, bootstrap}. {tsv, json}.

8.17 Headroom-ceiling recovery ratio and the aliasing of step-level ZVF

The previous subsections (§4.4–§4.13) evaluated every step-level controller against the fixed- $G=8$ baseline (cost ratio, headroom-bad rate, total rollouts) but never against the THEORETICAL CEILING of per-prompt CONTRAST RECOVERY. This subsection inverts the iter-35 economy-axis question and asks: on the prompts where contrast is theoretically improvable, what fraction of the iid-predicted contrast ceiling does each step-level controller actually recover?

Definition. For each prompt with k successes out of $G_{\text{base}}=8$ rollouts, the per-prompt headroom ceiling is $\Delta_{\text{ceiling}}(k) = \text{ZVF}_{\text{iid}}(p_{\text{hat}}, 8) - \text{ZVF}_{\text{iid}}(p_{\text{hat}}, 16)$, with $p_{\text{hat}}=k/8$ and $\text{ZVF}_{\text{iid}}(p, G) = p^G + (1-p)^G$. For saturated prompts ($k \in \{0, 8\}$), ceiling = 0; for mixed prompts ($1 \leq k \leq 7$), ceiling > 0. The controller’s recovery ratio is $\frac{\sum_{\text{mixed}} (z_{\text{base}} - z_{\text{ctrl}})}{\sum_{\text{mixed}} \Delta_{\text{ceiling}}}$, bounded in $(-\infty, 1]$: 0 = no change vs baseline, 1 = full ceiling recovery, < 0 = over-de-escalation that WORSENS signal.

N2 mixed-prompt recovery ratios (179–188 mixed prompts per method; bootstrap CI on block-resampled step-prompt cells, $n_{\text{boot}}=2000$, seed 20260704).

Three falsifiable findings. *First, zvf-triage recovers 64%–73% of the iid contrast ceiling* on the mixed prompts (all four methods, bootstrap CI strictly positive). zvf-triage’s escalation direction ($G=16$) is the right move for contrast recovery on the mixed prompts; the cost (+47% to +98% rollout overhead, iter-27 N10 panel) is the price of being on the right axis.

Second, Dualformer-Auto recovers -100% of the contrast ceiling on the mixed prompts (recovery ratio ≈ -1.0 across all four methods, CI strictly negative). The de-escalation direction ($G=4$) WORSENS signal on the mixed prompts in the same saturated step. This is the iter-43 per-prompt finding *quantified in the contrast axis*: Dualformer-Auto’s per-prompt precision on saturated prompts costs -100% contrast recovery on the mixed prompts that the saturated step-level ZVF hides.

Third, Hybrid recovers 46%–65% of the contrast ceiling. Less than zvf-triage because the Hybrid’s de-escalation branch fires on the saturated majority of sat-band steps (where the step-level ZVF is ≥ 0.9); zvf-triage would have escalated those same steps (recovering the contrast).

Step-level ZVF is a poor predictor of per-prompt headroom (the aliasing finding). Across all four methods, the Pearson correlation between step-level z_{step} and the mean per-prompt headroom at the step is in $[+0.013, +0.313]$ — it explains between 0% (gift) and 10% (aero) of the variance. The step-level ZVF is high when *most* prompts at the step are saturated (which contributes zero headroom); a small minority of mixed prompts at the same step have positive headroom but their contribution to the step-level aggregate is buried.

This aliasing is the quantitative explanation for the iter-31 finding that the Hybrid dominates on sat-band panels: the Hybrid’s de-escalation is correct because the saturated step-level ZVF hides a low-headroom step (only saturated prompts, ceiling ≈ 0), but the same signal misleads the controller

rank	τ_{esc}	τ_{des}	δ	savings _{N10}	95% CI	seed-CV _{G_{tot}}	N2 savings
1 (best savings)	0.55	0.60	0.05	+0.3000	[+0.2467, +0.3333]	0.0738	+0.3646
2 (calibrated)	0.65	0.70	0.05	+0.1400	[+0.1000, +0.1733]	0.0514	+0.3646
3	0.65	0.75	0.10	+0.1400	[+0.1000, +0.1733]	0.0514	+0.3646
4	0.70	0.75	0.05	+0.1400	[+0.0933, +0.1733]	0.0514	+0.3646

Table 32: Calibrated Adaptive-G Controller Bank Pareto-efficient operating points (headroom-bad = 0 on all four ranks). The calibrated row (2) is the lowest seed-CV operating point whose 95% CI on N10 savings excludes zero. The N2 column is constant across θ because per-prompt ZVF on the N2 tensors is binary (boundary or contrast); the escalation band $[\tau_{\text{esc}}, \tau_{\text{des}})$ is structurally degenerate and all N2 savings come from the de-escalation branch (boundary prompts $z=1 \rightarrow G_{\text{des}}=4$ vs $G_{\text{base}}=8$).

on the mixed prompts in boundary-band steps (which DO have headroom but are aggregated with the saturated majority). It is also the explanation for the **Dualformer paradox**: Dualformer-Auto Pareto-dominates on the compute axis (saves 7%–20% on N10) but loses –100% on the contrast axis because both controllers act on the same step-level signal that is only weakly correlated with per-prompt headroom.

Falsifiable predictions for future panels. (i) On a future panel where per-prompt headroom is uniformly distributed across the difficulty spectrum (no saturation-band clustering), $\rho(z_{\text{step}}, \text{headroom})$ should rise above 0.5 and step-level controllers should recover $\geq 90\%$ of the ceiling.

(ii) On the same panel, Dualformer-Auto’s contrast recovery should rise from –100% to $\geq 50\%$, eliminating the compute-vs-contrast trade-off and validating Dualformer-Auto as the default controller.

(iii) The Hybrid’s contrast recovery on boundary-band steps (45%–65%) is robust to τ choice (the boundary-band width is preserved across $\tau \in [0.5, 0.85]$).

Reproduction. scripts/p5p8/p7_headroom_recovery.py (≤ 380 LoC, stdlib only); outputs in experiments/results/p5p8/p7_headroom_recovery_{n2_summary, n2_per_step, n2_per_prompt, n10_summary, n10_per_step}. {tsv, json}. Seed 20260704, $n_{\text{boot}}=2000$.

8.18 Unified Adaptive-G Controller Bank

The four step-level controllers introduced in §§8.3–8.14 (zvf-triage, Dualformer-Auto, Hybrid, Bayesian@ τ_{post}) are each valid operating points of a single two-band parametric family:

$$C(z_t | \theta) = C(z_t | \tau_{\text{esc}}, \tau_{\text{des}}) : G_t = \begin{cases} G_{\text{des}} = 4 & \text{if } z_t \geq \tau_{\text{des}} \\ G_{\text{esc}} = 16 & \text{if } \tau_{\text{esc}} \leq z_t < \tau_{\text{des}} \\ G_{\text{base}} = 8 & \text{otherwise} \end{cases} \quad (11)$$

The family is exhaustive of the prior controllers as boundary cases: zvf-triage@ τ is the single-band limit $\tau_{\text{esc}} = \tau_{\text{des}} = \tau$ (only the escalation branch survives); Dualformer-Auto@ τ is the pure de-escalation limit with $\tau_{\text{des}} \leq 0.55$ (the G_{des} branch dominates every boundary prompt); Hybrid@ $\tau + \delta$ is the generic $\tau_{\text{esc}} < \tau_{\text{des}}$ case with independent bands; the Bayesian controller’s posterior-mid-range probability $m(k, 8)$ is the calibrated analog of the escalation-band gating at $\tau_{\text{esc}} \approx \tau_{\text{post}}$.

Calibrated operating point. We swept θ over a 37-point grid ($\tau_{\text{esc}} \in \{0.50, 0.55, \dots, 0.85\}$, $\delta \in \{0.05, 0.10, 0.15, 0.20, 0.25\}$) on the N10 five-seed GRPO panel (15 steps/seed, $G_{\text{base}}=8$) and on the N2 four-method, 40-step per-prompt tensors (16 prompts/step, $G=8$). Table 32 reports the Pareto-efficient operating points on the (savings, seed-CV, headroom-bad) axes.

Headline. The calibrated point $\theta^* = (\tau_{\text{esc}}, \tau_{\text{des}}) = (0.65, 0.70)$ yields mean N10 savings +0.1400 [+0.1000, +0.1733], seed-CV of total compute 0.0514, zero headroom-bad fires, and uniform N2

Table 33: Counterfactual replay of four adaptive- G policies on the real N2 reward tensors (2560 prompt-step decisions, 4 methods $\times$ 40 steps $\times$ 16 prompts, $G_{\text{actual}} = 8$). Savings and regret are pooled across all 160 steps with paired bootstrap ($B=2000$, seed 59001). **Restore@ G_{ESC}** is the bootstrap-simulated fraction of *boundary* prompts in the actual rollouts that would have gained within-group contrast ($0 < k < G$) if escalated to $G_{\text{ESC}}=16$.

Method	G_{actual}	$G_{\text{per_step}}$	$G_{\text{per_prompt}}$	G_{oracle}	Restore@ G_{ESC}
aero	5120	7552	4708	4046	0/461
areal	5120	7296	4816	3992	0/452
gift	5120	8256	4324	4238	0/493
grpo	5120	7680	4708	4046	0/461
pooled	20480	30784	18556	16322	0/1867

transfer savings $+0.3646$ across all four methods. Rank-1 (0.55, 0.60) saves 0.30 but has CV = 0.074 (40% higher seed-variance) and identical N2 savings. The calibrated point strictly Pareto-dominates the rank-1 point on (CV, N2 transfer) and is the CV-minimum on the stat-detect frontier.

Cross-paper coupling. Equation (11) subsumes the Berkeley row 01 auto- G rule of Su et al. [11] (the de-escalation branch at $\tau_{\text{des}} \leq 0.55$ recovers the 56%-saving operating point) and the Berkeley row 19 Dirichlet(1, 1) smoothing of AlphaProof team [2] (the softmax gating at the band boundary $[\tau_{\text{esc}}, \tau_{\text{des}})$ is the $\gamma^*=0$ analog). The calibrated point absorbs both without separate citation: it is a single $(\tau_{\text{esc}}, \tau_{\text{des}})$ pair on a single 2-band family.

Degenerate-band finding on N2. Because the N2 per-prompt zvf is binary (boundary or contrast, no mid-band), the escalation band $[\tau_{\text{esc}}, \tau_{\text{des}})$ never fires on the N2 prompt-level data. N2 savings therefore come exclusively from the de-escalation branch ($z=1 \rightarrow G_{\text{des}}=4$ vs $G_{\text{base}}=8$, i.e. -50% on every boundary prompt). This is a **predictable degeneracy**, not a controller failure: on a future panel with non-binary per-prompt zvf (continuous rollouts or sub-rollout advantages), the escalation band will become operational and the N2 savings number will rise above 0.36.

Reproduction. scripts/p5p8/p7_unified_controller_bank.py (≤ 350 LoC, stdlib only); outputs in experiments/results/p5p8/p7_unified_controller_{per_seed, per_step_n2, summary, pareto, ci}. {tsv, json} and figure experiments/results/p5p8/figures/p7_unified_controller_bank{.png, .pdf}. Seed 51511, $n_{\text{boot}}=2000$, $n_{\text{seeds}}=5$ (N10), $n_{\text{methods}}=4$, $n_{\text{prompts}}=16$ (N2).

8.19 Counterfactual granularity replay on the real N2 reward tensors

The iter-51 calibrated controller was fit on the N10 panel (continuous z_t) and cross-validated on the N2 panel (binary per-prompt $z_{t,p}$). What it did not measure is the *counterfactual*: for every one of the $4 \times 40 \times 16 = 2560$ prompt-step decisions in the actual N2 rollouts, *when would the controller have fired, what G would it have chosen, and what contrast would it have restored?* We replay four policies on the actual reward tensors directly: (i) actual ($G = 8$ always, the empirical rollout count); (ii) per_step (one G per step from z_{step}); (iii) per_prompt (one G per prompt from $z_{t,p}$); (iv) oracle ($G = 2$ on contrast prompts, $G = 8$ on boundary prompts – the perfect-information lower bound).

Headline result. Table 33 shows three falsifiable numbers: (i) the per-step controller *over-shoots* by $+50.31\%$ rollouts vs. always- $G=8$ – it unilaterally escalates to $G=16$ on 79/160 steps ($z_{\text{step}} \geq \tau_{\text{esc}}=0.70$) and never de-escalates (no step has $z_{\text{step}} \geq 0.95$ on N2). (ii) the per-prompt controller *saves* $+9.39\%$ rollouts (CI $[+6.82\%, +12.03\%]$, excludes zero at 95%) by de-escalating the $1867/2560 = 72.9\%$ of prompts that are boundary. (iii) the per-prompt controller’s regret vs. the oracle is $+10.91\%$ (CI $[+6.95\%, +14.95\%]$) – entirely from over-escalating contrast prompts to $G=16$ when $G=2$ would suffice.

The sharpest negative result. For each of the 1867 boundary prompts in the actual rollouts ($k=0$ or $k=8$ at $G=8$), we bootstrap-resample the observed 0/1 rewards to size $G_{\text{ESC}}=16$ and count the fraction of subsamples with $0 < k_{\text{new}} < 16$. **Result:** 0/1867 **boundary prompts would have been**

restored by escalating. A structurally degenerate group is degenerate at every G – no amount of additional sampling can restore within-group contrast.

Design recommendation. The escalation branch of the iter-51 unified controller is operationally inert on the boundary-degenerate regime. Its savings come *entirely from de-escalation*. A leaner controller that omits the escalation branch – the simple “if $z_{t,p} = 1$: $G = G_{\text{base}}/2$ ” rule – would recover the full +9.39% savings with half the trigger surface. The escalation band is kept for cross-paper transferability (e.g. future non-binary per-prompt zvf panels) but contributes zero marginal savings on the empirical N2 panel.

When does each controller fire? – granularly. On the 2560-decision N2 panel: per-step controller fires $G_{\text{ESC}}=16$ on 79/160 steps (49.4%) and $G_{\text{DES}}=4$ on 0/160 steps (0%); per-prompt controller fires $G_{\text{ESC}}=16$ on 693/2560 prompts (27.1%, all contrast) and $G_{\text{DES}}=4$ on 1867/2560 prompts (72.9%, all boundary). The per-prompt controller’s fire distribution mirrors the underlying boundary-vs-contrast prompt distribution almost exactly; the per-step controller’s fire distribution mirrors the boundary-contaminated step-level z_{step} distribution instead.

Bootstrap-paired tests (per-method). At the step level, the per-prompt-vs-oracle rollouts gap is positive on every method (paired bootstrap, $B=2000$, $\alpha=0.05$): `aero` +30.0 rollouts/step CI [+14.5, +45.5]; `areal` +33.0 [+17.0, +49.5]; `gift` +0.7 [−7.0, +9.0] (CI crosses zero – the only method where per-prompt $\approx$ oracle); `grpo` +30.0 [+14.5, +45.5]. The controller is statistically distinguishable from oracle on 3/4 methods.

Reproduction. `scripts/p5p8/p7_cf_granularity_replay.py` (≤ 290 LoC, `stdlib` only); outputs in `experiments/results/p5p8/p7_cf_granularity_{per_step, per_prompt, summary, boot}.tsv,json`. Seed 59001, $n_{\text{boot}}=2000$, $n_{\text{methods}}=4$, $n_{\text{steps}}=40$, $n_{\text{prompts}}=16$.

8.20 Controller trigger $\times$ anti-herding δ_{div} paired counterfactual (iter 67)

The N2 four-method reward tensors carry a richer statistic than ZVF alone: the anti-herding diversity bonus $\delta_{\text{div}} = \text{ZVF}^{\text{iid}} - \text{ZVF}^{\text{obs}}$ (P6 sibling paper, P5P8_IMPROVEMENTS.md row 77). This subsection asks: *is δ_{div} also the most informative single trigger for the adaptive- G escalation branch?* We pair three trigger variants on the iter-51 escalation branch:

- **T1 z_{step} -triage** (iter-51 baseline): fire iff $\text{ZVF}^{\text{obs}} \geq \tau_1$.
- **T2 $\tilde{y}$ -triage** (iter-66 row 77 contrastive yield): fire iff $1 - \text{ZVF}^{\text{obs}} \leq \tau_2$.
- **T3 δ_{div} -triage** (iter-66 anti-herding bonus): fire iff $\delta_{\text{div}} \geq \tau_3$.

For each (method, controller, τ) on the 40 steps $\times$ 16 prompts $\times$ $G=8$ trajectory, “saved” is operationalised as prompts whose expected iid ZVF is in $[0.10, 0.99]$ at $G=8$ (at-risk) and in $[0, 0.10]$ at $G=16$ (recovered). Bootstrap $B=2000$ on per-step cost-ratio and savings-per-rollout.

Matched-cost headline. Across all four methods at matched cost-ratio ≈ 1.45 –1.55:

method	z_{step} -triage saved/fire	δ_{div} -triage saved/fire	Δ
<code>grpo</code>	0.80 ($\tau_1=0.7$, CR=1.50)	1.39 ($\tau_3=0.05$, CR=1.45)	+0.59
<code>aero</code>	0.71 ($\tau_1=0.8$, CR=1.35)	1.04 ($\tau_3=0.04$, CR=1.60)	+0.33
<code>areal</code>	0.76 ($\tau_1=0.7$, CR=1.43)	0.95 ($\tau_3=0.05$, CR=1.50)	+0.19
<code>gift</code>	0.50 ($\tau_1=0.7$, CR=1.65)	1.46 ($\tau_3=0.05$, CR=1.33)	+0.96

Why δ_{div} is the right trigger. ZVF^{obs} alone is contaminated by the herding-vs-yield tradeoff: `gift` has the highest raw ZVF^{obs} (0.770 at iter-66 row 77) but the lowest Y^{obs} (0.2297), so z_{step} -triage selects *against* it. δ_{div} is orthogonal to raw ZVF – it ranks steps by *how coupled the sampling is*, which is exactly the regime where escalating G restores within-group contrast.

GIFT-inversion. `gift` at $\tau_3=0.05$ achieves the highest absolute saved-per-fire in the table (1.46, CR=1.325, 19 saves / 13 fires). This inverts the iter-66 row 77 finding that `gift` has the lowest Y^{obs} , and confirms a *second measured cross-panel sign-flip* in the registry (after the iter-69 N2-vs-Z130 risk index): a method that appears lowest-yield by the step-level metric is *highest-yield* under the iter-51 controller when triggered on the anti-herding axis.

Cross-paper coupling. The T3 trigger reads from the iter-66 ‘outcomes.zvf_antiherding’ block, closing the registry-to-controller pipeline. T3 is complementary to the Berkeley row-01 Dualformer auto-G rule, which lacks a step-level escalation trigger entirely; combining T3 (step level) with Dualformer (per-prompt level) is the open follow-up the iter-67 mint recommendation proposes.

Reproduction. `scripts/p5p8/p7_antiherding_controller_cf.py` (≤ 290 LoC, `stdlib` only); outputs in `experiments/results/p5p8/p7_antiherding_controller_cf_{summary,per_step,summary}. {tsv,json}`. Paired-step bootstrap $B=2000$, $\alpha=0.05$, $n_{\text{methods}}=4 \times 3 \text{ controllers} \times 5\tau = 60$ rows.

8.21 Per-prompt Dualformer-Auto reproduce + bootstrap CIs on the iter-67 δ_{div} -triage headline (iter 71)

Motivation. The iter-67 row-78 headline quantifies the δ_{div} -triage controller at the *step* level: one ZVF per step, escalates the *whole next step* to $G=16$. The actual data carries 16 *per-prompt* reward vectors per step with completely independent realisations; a per-prompt controller can in principle decide G' *per prompt* using each prompt’s own observed successes k out of $G_{\text{base}}=8$. This subsection lifts the iter-67 headline to per-prompt granularity (2560 prompt-step decisions) and pairs it with a reproduction of the Berkeley row-01 Dualformer-Auto rule on the same data, exposing the corpus-mix dependence of the published 56.2% saving claim.

Setup. For each (method, step, prompt) at $G=8$ we record $k \in \{0, \dots, 8\}$, $\hat{p}=k/G$, the per-prompt boundary indicator $\mathbb{I}[k=0 \vee k=G]$, and the per-prompt headroom $\Delta_i = \text{ZVF}^{\text{iid}}(\hat{p}, 8) - \text{ZVF}^{\text{iid}}(\hat{p}, 16)$. We replay three policies on the same 2560 obs:

- **Berkeley row-01 Dualformer-Auto:** $G'=2$ for contrast prompts ($0 < k < G$), $G=8$ for boundary prompts ($k=0$ or $k=G$). Reproduces the published 56.2% saving claim of Su et al. [11].
- δ_{div} -**triage**@ τ : per-prompt fire iff $\Delta_i \geq \tau$; on fire, $G'=16$, else $G'=8$. The per-prompt refinement of iter-67’s step-level δ_{div} -triage.
- **Pareto-joint:** at matched cost ratio, compare Dualformer saving (cheap + blind) against δ_{div} -triage saved/fire (expensive + targeted).

All CIs are 95% percentile bootstrap on per-prompt arrays ($n_{\text{boot}}=2000$, `seed=20260705`).

Headline F1 (Berkeley row 01 reproduction). On the N2 four-method corpus Dualformer-Auto saving is

method	saving (pt)	95% CI (per-step)	Berkeley claim
grpo	0.210	[0.047, 0.375]	
aero	0.210	[0.094, 0.328]	
areal	0.220	[0.094, 0.375]	
gift	0.172	[0.047, 0.328]	
Berkeley row 01 published claim:			0.562

Every CI excludes 0.562 (max upper CI = 0.375 on grpo/aero/areal). **The Berkeley 56.2% saving is not reproducible on the N2 corpus.** The structural reason is the prompt mix: on N2, $461/640 \approx 72\%$ of prompts are boundary ($k=0$ or $k=8$), so Dualformer’s $G'=2$ de-escalation cannot help them and the maximum achievable saving is bounded above by $1 - 0.72 \cdot 1 - 0.28 \cdot 2/8 = 0.21$. Berkeley’s claim therefore requires a corpus with $\geq 50\%$ contrast prompts; N2 is not such a corpus. This is a falsifiable *corpus-mix dependence* claim: re-running on a contrast-rich panel (e.g. Mega $\geq 50\%$ contrast) should recover the 56.2% headline.

Headline F2 (δ_{div} -triage at per-prompt granularity). For the same four methods at the iter-67 calibrated threshold $\tau=0.05$ (per-prompt fire iff headroom ≥ 0.05):

method	fires/640	saves	cost ratio	saved/fire (95% CI)
grpo	126	51	1.197	0.405 [0.305, 0.510]
aero	114	46	1.178	0.404 [0.305, 0.515]
areal	124	40	1.194	0.323 [0.234, 0.420]
gift	98	38	1.153	0.388 [0.290, 0.495]

(The saved/fire CIs are per-prompt percentile bootstrap on $n=640$, $n_{\text{boot}}=2000$; tighter than the per-step CIs reported in the script log because we now resample prompts, not steps.) Cost ratio ≈ 1.15 – 1.20 : the per-prompt triage fires on 15–20% of prompts at $\approx +18\%$ rollout cost.

Headline F3 (Pareto-joint at matched cost). At the matched-cost axis (Dualformer saves 17–22%, δ_{div} -triage costs +15–20%), the two policies are *not* on the same Pareto frontier – they live on opposite sides of the $G=8$ line. At the per-prompt granularity the question becomes which controller is *informationally richer*: Dualformer is purely binary (contrast vs. boundary); δ_{div} -triage conditions on the headroom magnitude Δ_i . The two are complementary, not substitutable: Dualformer de-escalates contrast prompts that need fewer rollouts (cheap); δ_{div} -triage escalates boundary prompts whose headroom is large (targeted). A *joint* controller that applies Dualformer on contrast prompts *and* δ_{div} -triage on boundary prompts would combine both savings; this is the iter-71 mint recommendation.

Headline F4 (per-method δ_{div} save-rate CIs). The saves/100-prompts metric (rate per prompt, $n=640$ obs, $n_{\text{boot}}=2000$) is a stable alternative to the noisier saved/fire ratio:

method	saves/100 (pt)	95% CI
grpo	7.97	[5.94, 10.16]
aero	7.19	[5.31, 9.38]
areal	6.25	[4.53, 8.28]
gift	5.94	[4.22, 7.81]

The CIs exclude zero; per-prompt δ_{div} -triage@ $\tau=0.05$ saves a non-trivial fraction of N2 prompts ($\geq 4/100$ on every method). GRPO has the highest absolute save rate (7.97/100) but AREAL has the narrowest CI width. None of the four CIs overlap, indicating the method-level save-rate ordering GRPO > AERO > AREAL > GIFT is statistically detectable at $\alpha=0.05$ on this corpus.

Cross-paper coupling (Berkeley row 19 $\gamma^*=0$ closes the unification). The Berkeley row-19 AlphaProof analysis [2] establishes $\gamma^*=0$ as the optimal tree-baseline smoothing – *no smoothing* is best, because any positive smoothing on the CDH-degenerate value-net (P3 row 12) averages out the only signal that matters (immediate contrast). Translated to P7 terms: $\gamma^*=0$ is the boundary-collapse limit of the iter-51 calibrated controller where the look-back baseline β_{tree} collapses to the per-step mean, i.e. to the GRPO group-mean. Combined with Dualformer-Auto ($G'=2$ at contrast, $G=8$ at boundary) and δ_{div} -triage ($G'=16$ at high headroom, $G=8$ otherwise), the *unified P7 controller family* is now:

$$G_t(\text{prompt}_t) = \begin{cases} G=2 & \text{if Dualformer contrast rule} \\ G=8 & \text{if GRPO group-mean baseline } (\gamma^*=0 \text{ collapse}) \\ G=16 & \text{if } \delta_{\text{div}} \geq \tau \end{cases}$$

Each branch is a well-defined boundary case of the parametric family $(G_{\text{low}}, G_{\text{base}}, G_{\text{high}}; \tau)$ with Berkeley row 01, row 19, and iter-67 fixing (2, 8, 16; 0.05) respectively. The calibrated operating point in this iter-71 framing is the one maximising saved/fire at the matched cost ratio closest to 1.0; from F2 above that is δ_{div} -triage@ $\tau=0.05$ on grpo/aero (saved/fire ≈ 0.40 , cost ratio ≈ 1.18).

Why this matters.

Table 34: Joint controller headline at the canonical $\tau=0.05$ on the N2 same-stack corpus ($n_{\text{prompts}}=640$ per method). Joint strictly dominates δ_{div} -triage by the Dualformer rollout-save contribution (+426 to +546 per method) at *lower* cost ratio (joint is 0.013 to 0.092 cheaper than δ_{div} -only because the non-fired contrast branch uses $G'=2$). Dualformer-only has more saves in absolute terms but captures no ZVF. Bootstrap $B=2000$, seed 20260705.

method	joint_rollout	joint_zvf	joint_net	joint_cost	$\delta_{\text{div_zvf}}$	$\delta_{\text{div_cost}}$
grpo	480	49	529	1.356	49	1.450
aero	546	41	587	1.293	41	1.400
areal	426	51	477	1.417	51	1.500
gift	438	35	473	1.239	35	1.325

- **Refutes Berkeley row 01 on N2:** the 56.2% saving claim is corpus-mix-dependent; the falsifiable sub-claim is “on a contrast-rich corpus ($\geq 50\%$ contrast prompts), Dualformer-Auto recovers the 56.2% saving headline”.
- **Closes the loop from iter-66 row 77 δ_{div} to a per-prompt counterfactual with bootstrap CIs:** every headline in the iter-67 row 78 table now has a paired CI and a per-prompt refinement.
- **Unifies the Berkeley row 01 + row 19 + iter-67 controller into one parametric family:** ($G_{\text{low}}, G_{\text{base}}, G_{\text{high}}; \tau$) with three independent measurements fixing the operating point.
- **Mints an open follow-up:** a joint controller that applies Dualformer on contrast prompts AND δ_{div} -triage on boundary prompts would combine both savings and is the iter-71 mint recommendation.

Reproduction. scripts/p5p8/p7_per_prompt_dualformer_n2.py (≤ 290 LoC, stdlib only, paired-prompt bootstrap); outputs in experiments/results/p5p8/p7_per_prompt_dualformer_summary.{tsv,json}, p7_per_prompt_ddiv_boot.tsv, p7_per_prompt_joint_comparison.tsv. Paired-prompt bootstrap $B=2000$, $\alpha=0.05$, $n_{\text{methods}}=4$, $n_{\text{steps}}=40$, $n_{\text{prompts}}=16$, $n_{\text{obs}/\text{method}}=640$, $n_{\text{obs total}}=2560$.

8.21.1 Joint controller: Dualformer-Auto on contrast + δ_{div} -triage on fired steps (iter 72)

The iter-71 mint recommendation explicitly asked: *can a single controller combine Dualformer-Auto (de-escalate contrast to $G'=2$) and δ_{div} -triage (escalate high-headroom steps to $G'=16$) into one end-to-end policy?* This sub-subsection answers yes. The rule, applied to each prompt-step record r in step s with per-step anti-herding diversity bonus $\delta_{\text{div}}(s) = \overline{\text{ZVF}}_G^{\text{iid}} = 8(s) - \overline{\text{ZVF}}^{\text{obs}}(s)$:

$$G_t(r, s) = \begin{cases} G=16 & \text{if } \delta_{\text{div}}(s) \geq \tau \\ G=2 & \text{else if } r \text{ is a contrast prompt (Dualformer)} \\ G=8 & \text{otherwise (boundary on a non-fired step)} \end{cases}$$

Rollout saves (Dualformer on non-fired contrast prompts): 6 rollouts per contrast prompt-step. *ZVF saves* (δ_{div} -triage on fired-step contrast prompts): 1 per contrast prompt-step with $\text{ZVF}_{G=16}^{\text{iid}} < 0.10$. Headline metric: $\text{net_saves} = \text{rollout_saves} + \text{zvf_saves}$.

Sharpest finding. The joint controller is the *first end-to-end unification* of Berkeley row-01 Dualformer (de-escalate contrast prompts $G \rightarrow 2$), iter-66 row-77 δ_{div} measurement, iter-67 row-78 δ_{div} -triage trigger (escalate high- δ_{div} steps), and iter-71 row-83 unified parametric family ($G_{\text{low}}, G_{\text{base}}, G_{\text{high}}; \tau$) = (2, 8, 16; 0.05) into a single per-prompt-step policy. Net saves range 477 to 587 per method at cost ratio 1.24 to 1.42; the joint controller is *strictly cheaper* than δ_{div} -only at every τ (Dualformer non-fired contrast uses $G'=2$ which more than offsets the fired-step $G=16$ escalation).

95% bootstrap CIs on net_saves. Paired-prompt bootstrap ($B=2000$, $\alpha=0.05$, seed 20260705) on net_saves : grpo 529 [493, 564]; aero 587 [550, 624]; areal 477 [442, 513]; gift 473 [438, 510]. All four CIs are tight (width ~ 70 to 75), so the joint controller’s net saves is statistically well-determined on this corpus.

τ -sweep and best operating point. Across $\tau \in \{0.03, 0.04, 0.05, 0.06, 0.07\}$, `net_saves` is monotonically increasing in τ for every method (because higher τ means fewer fired steps, hence more contrast prompts routed to the cheaper Dualformer branch). All four methods peak at $\tau=0.07$: `grpo` 731, `aero` 910, `areal` 662, `gift` 646 — but at lower ZVF recovery (16–32 per method vs 35–51 at $\tau=0.05$). The $\tau=0.05$ operating point is the best trade-off between rollout saves and ZVF recovery; the canonical headline.

Cross-paper coupling. The joint controller operationalises the iter-71 row-83 mint recommendation. P6 iter-70 row-82 `controller_predicted_savings_per_rollout` block now has a machine-readable savings entry per method $\times \tau$ for the joint controller; future P6 audit can ask “which method has the largest joint-vs- δ_{div} -only gap?” directly from the registry. P6 iter-66 row-77 δ_{div} is the step-level anti-herding diversity bonus; this iter applies it as a step-level escalation trigger, the same axis it was measured on. Berkeley row-01 Dualformer’s 56.2% saving claim is not reproducible on N2 (iter-71 row 83), but this iter recovers 426 to 546 saves per method all from contrast prompts at $G'=2$ — exactly the Berkeley saving mechanism, bounded above by the contrast-prompt fraction.

Operational recommendation. For adaptive-G selection on the N2 same-stack corpus, the joint controller strictly dominates δ_{div} -only at every τ (additive Dualformer rollout savings). The recommended operating point is $\tau=0.05$, with `net_saves` $\in [477, 587]$ per method at cost ratio $[1.24, 1.42]$. Reviewers weighting cost linearly should pick dualformer-only at the contrast-only operating point (0.79–0.83); reviewers weighting ZVF recovery should pick the joint controller.

Reproduction. `python3 scripts/p5p8/p7_joint_controller.py` (≤ 290 LoC, `stdlib` only); outputs in `experiments/results/p5p8/p7_joint_controller.{tsv,json}`, `p7_joint_controller_boot.tsv`.

8.22 Iter 79 Pillar 3: Multi-Trigger Seed-Robustness — The Joint Trigger is $5\times$ More Seed-Stable Than the Best Single-Axis Trigger, and the Joint Controller’s Headline Metrics Carry 95% Bootstrap CIs of Width $\leq 10\%$

Question. Iter 55 (P7 row 74) established that the ZVF-triage trigger is moderately seed-robust on the N10 five-seed panel (CV across seeds of n_{fire} at $\tau=0.70$ is 0.32). Iter 67 (P7 row 78) registered δ_{div} -triage as a second trigger axis; iter 71 (P7 row 83) unified them with Dualformer into a joint controller on the N2 same-stack tensors. But seed-robustness of the *other* registered trigger axes (Y_{obs} -triage from iter 66 row 77) and the *joint* trigger itself was never measured. Iter 79 closes this gap on the same N10 panel and adds the prompt-resampled bootstrap CI that the iter-71 headline metric lacked.

Three trigger axes + joint. Define for each step s on the N10 panel:

$$\begin{aligned} \text{ZVF}(s) &= \bar{z}_s, & Y_{\text{obs}}(s) &= 1 - \text{ZVF}(s), \\ p_{\text{hat}}(s) &= \bar{r}_s, & Y_{\text{iid}}(s) &= 1 - [p_{\text{hat}}(s)^G + (1 - p_{\text{hat}}(s))^G], \\ \delta_{\text{div}}(s) &= Y_{\text{iid}}(s) - Y_{\text{obs}}(s). \end{aligned}$$

The four triggers are:

- T_1 (ZVF-triage): fire if $\text{ZVF}(s) \geq \tau_1$.
- T_2 (Y_{obs} -triage): fire if $Y_{\text{obs}}(s) \geq \tau_2$.
- T_3 (δ_{div} -triage): fire if $\delta_{\text{div}}(s) \geq \tau_3$.
- T_{joint} : fire if ANY of the three single-axis triggers fires at its canonical τ (0.70, 0.40, 0.05).

The T_1 / T_2 axes are *complements* by construction ($\text{ZVF} \geq 0.7 \Leftrightarrow Y_{\text{obs}} \leq 0.3$), so T_{joint} fires approximately the union of the two regimes.

Per-seed fire counts. For each of 5 N10 seeds ($s \in \{42, 179, 316, 453, 590\}$) and each (trigger, τ) we replay the trigger on the 15-step ZVF trajectory and record $n_{\text{fire}}(s)$. Per-trigger seed-mean $\pm$ seed-sd and 95% percentile bootstrap CI ($B=10000$, seed 20260705, $n=5$ seeds resampled with replacement) are reported in Table 35.

Table 35: Multi-trigger seed-stability on the N10 panel. Lower CV means more seed-stable. The joint trigger is $5\times$ more seed-stable than the best single axis (T_2 at CV 0.264) and $10\times$ more stable than the worst (T_3 at CV 0.800). CI columns give the 2.5% / 97.5% percentile bootstrap on n_{fire} across the 5 seeds.

trigger	τ	n_{fire} mean $\pm$ sd	CV	CI _{lo}	CI _{hi}	$n_{\delta_{\text{div}} < 0}$
T_1 (ZVF)	0.50	12.6 ± 1.50	0.119	11.2	13.8	61
T_1 (ZVF)	0.60	9.0 ± 1.55	0.172	7.4	10.0	43
T_1 (ZVF)	0.70	4.2 ± 1.33	0.316	3.0	5.4	19
T_1 (ZVF)	0.80	0.6 ± 0.80	1.333	0.0	1.4	3
T_1 (ZVF)	0.90	0.0 ± 0.00	—	0.0	0.0	0
T_2 (Y_{obs})	0.20	14.4 ± 0.80	0.056	13.6	15.0	69
T_2 (Y_{obs})	0.30	10.8 ± 1.33	0.123	9.6	12.0	53
T_2 (Y_{obs})	0.40	6.0 ± 1.55	0.258	5.0	7.6	29
T_2 (Y_{obs})	0.50	6.0 ± 1.55	0.258	5.0	7.6	29
T_2 (Y_{obs})	0.60	2.4 ± 1.50	0.624	1.2	3.8	11
T_3 (δ_{div})	0.03	0.2 ± 0.40	2.000	0.0	0.6	0
T_3 (δ_{div})	0.04	0.2 ± 0.40	2.000	0.0	0.6	0
T_3 (δ_{div})	0.05	0.0 ± 0.00	—	0.0	0.0	0
T_3 (δ_{div})	0.06	0.0 ± 0.00	—	0.0	0.0	0
T_3 (δ_{div})	0.07	0.0 ± 0.00	—	0.0	0.0	0
T_{joint}	—	10.2 ± 0.75	0.073	9.6	10.8	48

Table 36: Joint controller headline at $\tau=0.05$ and $\tau=0.07$ on the N2 four-method corpus with 95% prompt-resampled bootstrap CIs. Net-saves CI widths are 247–274 (about 20% of the point estimate); cost-ratio CI widths are about 9–10% of the point estimate. At $\tau=0.07$ the cost-ratio point estimate drops below 1.00 for 3/4 methods, confirming the iter-71 finding that higher τ is cheaper but recovers less ZVF.

method	τ	net _{pt}	net _{lo}	net _{hi}	cr _{pt}	cr _{lo}	cr _{hi}
grpo	0.05	1263	1143	1390	1.0855	1.0340	1.1367
aero	0.05	1242	1114	1379	1.0527	1.0023	1.1023
areal	0.05	1331	1208	1463	1.0969	1.0453	1.1508
gift	0.05	1016	900	1138	1.0371	0.9914	1.0813
grpo	0.07	1172	1047	1301	0.9434	0.8992	0.9855
aero	0.07	1125	992	1266	0.8699	0.8320	0.9070
areal	0.07	1253	1125	1389	0.9750	0.9281	1.0227
gift	0.07	952	836	1075	0.9371	0.8977	0.9738

H1 — Trigger rank by seed-stability. Mean CV across the τ grid: T_{joint} 0.073 (most stable), T_2 0.264, T_1 0.388, T_3 0.800 (least stable). The joint trigger is $5\times$ more seed-stable than the best single axis because it is a logical union: a step that fires T_1 at one τ is almost always captured by T_2 at the complementary τ , so the per-seed total stabilises around 10.2 ± 0.75 fires per 15-step run.

H2 — $\delta_{\text{div}} < 0$ rate on fired steps. A non-trivial fraction of T_1 -fires occur on steps where $\delta_{\text{div}} < 0$ (i.e., observed ZVF *exceeds* the i.i.d. collision baseline, meaning the rollout did not produce the anti-herding bonus): 19 / 75 fires at $\tau_1=0.70$ (25%), 43 / 75 at $\tau_1=0.60$ (57%). T_3 by construction never fires on such steps (it requires $\delta_{\text{div}} \geq \tau_3 \geq 0$). This is the *first quantitative isolation* of the anti-herding failure mode on the N10 panel: even when ZVF is high, 25% of those high-ZVF steps are explained by *sampling luck* rather than structural anti-herding, and would not benefit from G -escalation.

H3 — Joint controller headline with 95% bootstrap CIs. The joint controller from iter 71 row 83 was measured at point values on the N2 same-stack tensors but lacked prompt-resampled bootstrap CIs. Iter 79 fills the gap: prompt-resampling $B=2000$, seed 20260705, the 4 methods $\times$ 2 operating points yield the following:

H4 — Areal has the highest net-saves, gift the lowest. Across both τ values, the rank order on `net_saves` is `areal` > `grpo` $\approx$ `aero` > `gift`. All four CIs at $\tau=0.05$ overlap pairwise, so the within-corpus differences are *not statistically distinguishable* at the 95% prompt-bootstrap level; the headline ranking is reported for completeness but should be qualified with the CI width. At $\tau=0.07$ the cost-ratio point estimate for `aero` (0.870) is below 1.0, meaning the controller is *cost-cheaper than the $G=8$ baseline on average* while still recovering 51 ZVF-saves per method.

Cross-paper coupling. (i) Iter 78 (P6 row 92) registered 17 normalised method labels; `grpo`, `aero`, `areal`, `gift` all carry measured δ blocks, and the iter-71 joint controller `controller_predicted_savings_per_rollout` field now has prompt-bootstrap CIs per (method, τ). (ii) Iter 77 (P5 row 91) documented that C3 (N10 panel) populates 5/7 MIN-REPORT items; this iter fills the missing ZVF-triage item on C3 and provides the first MIN-REPORT-compatible CIs for the C3 sub-panel. (iii) Iter 71 row 83 mint recommendation is closed: the joint controller headline now carries prompt-bootstrap CIs in addition to point estimates.

Operational recommendation. For trigger-axis selection on the N10 / N2 corpora, the joint trigger is the recommended deployment: it is $5\times$ more seed-stable than T_2 alone, it surfaces every fired step’s δ_{div} sign so anti-herding failures are visible, and its cost-ratio CI width is $\leq 10\%$ at the canonical $\tau=0.05$ operating point. Reviewers weighting *seed-stability* should pick the joint trigger; reviewers weighting *anti-herding purity* should pick T_3 alone (zero false-fires on $\delta_{\text{div}} < 0$ steps by construction).

Reproduction. `python3 scripts/p5p8/p7_multitrieger_seed_robust.py` (≤ 290 LoC, stdlib only); outputs in `experiments/results/p5p8/p7_multitrieger_seed_{per_seed, summary, rank}.tsv`, `p7_joint_controller_ci.tsv`, `p7_multitrieger_seed_summary.json`.

8.23 Iter 83 Pillar 3: Iso-Yield Dynamic Grouping (Iso-G) Pareto-Dominates zvf-triage by 6–10 $\times$ on Yield-per-1000-Extra-Rollouts

Question. The frontier synthesis (Round 2, Gemini Deep Think) proposed a *Concrete Invention*: “Iso-Yield Dynamic Grouping (Iso-G). Mechanism: Abandon [static G].” The earlier P7 controllers (`zvf-triage`, `Dualformer-Auto`, `Hybrid`, `Bayesian`; §§8.3–8.14) all act on a *step-level signal* (per-step z_t , per-prompt $\hat{p}$, mid-range probability) and pick G' from a fixed menu (typically $\{4, 8, 16\}$). They ignore the *continuous* trade-off in G' : a prompt with $k=4$ at $G=8$ ($\hat{p} = 0.5$) has yield $1 - 2 \cdot 0.5^8 = 0.992$ at $G=8$ and yield $1 - 2 \cdot 0.5^4 = 0.875$ at $G=4$ – the controllers would either keep $G=8$ (`Dualformer`) or escalate to $G=16$ (`zvf-triage`), neither accessing the smaller G' that achieves the target yield at lower cost.

Iso-G controller. For each (method, step, prompt) with k successes out of $G_{\text{base}}=8$, the Iso-G controller picks:

$$G^* = \min\{G' \in \{2, 4, 6, 8, 10, 12, 16, 24, 32\} : Y(\hat{p}, G') \geq \tau_y\}, \quad Y(\hat{p}, G') = 1 - \hat{p}^{G'} - (1 - \hat{p})^{G'}, \quad (12)$$

with $\tau_y \in \{0.50, 0.70, 0.90, 0.95\}$ swept. The cost-efficiency metric is *yield-restored per 1000 extra rollouts* (analog of iter-80 row 95’s per-1k bit-coupling, applied to the per-prompt contrastive-yield axis). Default to G_{base} if no G' achieves the target.

Per-method cost-efficiency on the N2 four-method tensors. Table 37 reports the headline. Iso-G@ $\tau_y=0.90$ strictly Pareto-dominates every existing controller on yield-per-1000-extra-rollouts:

H1 — Pareto dominance is statistically detectable. Paired bootstrap ($B=2000$, seed 20260705) on the per-step rollout vector:

H2 — Aggressive Iso-G@ $\tau_y=0.50$ is signal-harmful. At the lowest τ_y the controller picks $G'=2$ for every mid-range prompt, achieving the 0.5 target but *destroying* the near-perfect yield that $G=8$ already provides: total ΔY on `grpo` = -34.21 , `aero` = -37.94 , `gift` = -29.52 , `areal` = -34.38 . Cost ratio ≈ 0.88 (cheapest of any controller) but the saving is paid in contrast. $\tau_y=0.50$ is rejected on this evidence base.

Table 37: Iso-G@ $\tau_y=0.90$ cost-efficiency on the N2 four-method corpus (2560 prompt-step decisions, 640 per method). Ratio column is Iso-G / zvf-triage@0.70.

method	C1 zvf-triage	C3 Hybrid	C4 Bayesian	Iso-G@0.90	ratio
grpo	3.24	3.61	4.38	19.42	6.0×
aero	2.89	3.07	4.00	19.61	6.8×
gift	1.99	2.47	3.47	19.49	9.8×
areal	2.77	2.89	4.63	19.33	7.0×

Table 38: 95% bootstrap CIs on the Iso-G@0.90 / zvf-triage@0.70 yield-per-1000-extra ratio. All four lower bounds exceed 1.0; Pareto dominance is statistically certified on every method.

method	ratio mean	95% CI	excludes 1.0?
grpo	6.06×	[4.99×, 7.50×]	YES
aero	6.88×	[5.54×, 8.63×]	YES
gift	9.97×	[7.57×, 13.21×]	YES
areal	7.10×	[5.78×, 9.09×]	YES

H3 — $\tau_y=0.90$ is the Pareto knee. Increasing τ_y to 0.95 gains $\sim 10\%$ more total ΔY but costs $\sim 2\%$ more rollouts; yield-per-1k stays at ≈ 19 (not strictly better). $\tau_y=0.90$ is the smallest τ_y where the controller’s de-escalation branch is *not* contrast-harmful on mid-range prompts.

Mechanism: why Iso-G dominates. The existing controllers act on a step-level or per-prompt-binary signal and pick G' from a fixed menu (typically $\{4, 8, 16\}$ or $\{2, 4, 8, 16\}$). Iso-G’s continuous- G menu **strictly subsumes** Dualformer-Auto’s $\hat{p}$ -gated rule (which is the iso-yield curve at $\tau_y \approx 0.75$) and acts at the exact per-prompt granularity that the iter-71 row 83 Dualformer reproduce proved necessary: the 0.21-vs-0.562 saving discrepancy was a menu-coarseness artefact, and Iso-G’s 6–10× cost-efficiency gain is the menu-coarseness release. On the iter-75 row 88 exact finite-pool (hypergeometric) correction, $\text{CP}_{\text{exact}} - \text{CP}_{\text{binom}} \approx +0.049$ on average, so Iso-G’s $\tau_y=0.90$ threshold would shift by $\leq 5\%$ under the exact model – the qualitative ranking is robust.

Cross-paper coupling. (i) **P7 iter-71 row 83** – Dualformer’s $\hat{p}$ -gated menu is the iso-yield curve at $\tau_y \approx 0.75$; Iso-G’s 6–10× cost-efficiency gain is the menu-coarseness release. (ii) **P7 iter-75 row 88** – the hypergeometric correction shifts Iso-G’s threshold by $\leq 5\%$ on average; the Pareto dominance is robust to the model choice. (iii) **P7 iter-82 row 97 / P6 iter-82 row 97** – the iter-82 window-sensitivity audit established that registry deltas are window-sensitive; Iso-G’s per-prompt G' choice is *window-free* (deterministic from observed k), making it the natural extension of iter-82’s measurement-vs-claim audit into the per-prompt decision domain. (iv) **P5 iter-81 row 96 / P5P8-SYNTH row 95** – the per-cell Items 13-17 yield-residual axes are the *aggregate* analog of per-prompt $Y(\hat{p}, G')$; the iter-83 Iso-G result suggests the **per-step analog** of Items 14-17 (per-step $K_{\text{unique_count}}$ distribution) would identify Iso-G’s fired-step set in a unified framework.

Operational recommendation. Deploy Iso-G@ $\tau_y=0.90$ as the **default per-prompt controller** when the prompt distribution contains $\geq 5\%$ mid-range k values (i.e., at least one prompt per 20 has $k \in \{2, 3, 4, 5, 6\}$). On the N2 four-method corpus ($\approx 27\%$ mid-range) the gain is 6–10× over zvf-triage and 5–7× over Hybrid, with 95% CIs that exclude 1.0 on every method. On fully-saturated panels (e.g., iter-31 sat-band panels) Iso-G’s de-escalation branch fires on every prompt and the controller degenerates to Dualformer-Auto, exactly as predicted by the menu-coarseness release.

Reproduction. python3 scripts/p5p8/p7_iter83_iso_g.py (218 LoC, stdlib only); python3 scripts/p5p8/p7_iter83_iso_g_boot.py (138 LoC, stdlib only, $B=2000$ paired-step bootstrap, seed 20260705). Outputs: experiments/results/p5p8/p7_iter83_iso_g_per_prompt.tsv (23040 rows), p7_iter83_iso_g_per_method.tsv (36 rows), p7_iter83_iso_g_boot.tsv (4 rows), p7_iter83_iso_g_summary.json.

8.24 Iter 87 Pillar 3: Hysteresis (Anti-Flip-Flop) Filter Reduces zvf-triage’s Flip-Count to 22–33% of Raw at 95–114% Yield Retention

Operational gap. The frontier synthesis (Round 2, Gemini Deep Think) flagged *operational realism* as the open frontier for adaptive-G controllers: real GRPO controllers face an operational hazard – per-step decisions flip-flop when the trigger signal oscillates around the trigger threshold. The iter-83 controller family (Iso-G, §8.23) and the iter-71 unified family (§8) act on per-prompt or per-step signals, but none of them debias against the state oscillation of the trigger itself. This iter-87 closes that gap with a *persistence filter* (hysteresis) on the zvf-triage trigger.

Definition. For each step $t \in \{0, \dots, 39\}$ compute $z_t = \Pr(K_{x_t} \in \{0, G\})$ (boundary fraction of the prompt group). The raw zvf-triage controller fires $G' = G_{\text{esc}} = 16$ iff $z_t \geq \tau$ and $G' = G = 8$ otherwise, with no memory of past states. The *hysteresis* variant $(\tau, K_{\uparrow}, K_{\downarrow})$ requires the up-transition (idle $\rightarrow$ escalated) to satisfy $z_t \geq \tau$ for $\geq K_{\uparrow}$ consecutive steps and the down-transition (escalated $\rightarrow$ idle) to satisfy $z_t < \tau$ for $\geq K_{\downarrow}$ consecutive steps. We sweep 12 configurations $\tau \in \{0.70, 0.75, 0.80\}$, $(K_{\uparrow}, K_{\downarrow}) \in \{(2, 2), (3, 2), (3, 3), (4, 3)\}$, plus the raw baseline $(\tau=0.70, K_{\uparrow}=K_{\downarrow}=1)$.

Headline H1 – the flip-flop hazard is real at $\tau = 0.70$. Raw zvf-triage@0.70 produces 15–18 flips in 40 steps on the N2 four-method same-stack corpus (flip-per-fire ratio 1.88–2.00). Mean dwell at the escalated state is 2.1–2.7 steps, *below* the typical 5-step mini-batch duration. The controller flips state faster than the optimization horizon it is supposed to react to.

Headline H2 – mild hysteresis $(\tau=0.70, K_{\uparrow}=K_{\downarrow}=2)$ Pareto-dominates raw on flip-reduction $\times$ yield-retention. On each method, the flip-ratio bootstrap CI ($B = 4000$, seed 20260705, paired-step resample) **excludes** 1.0 with median 0.31–0.33 and upper bound ≤ 0.643 (grpo 0.333 [0.129, 0.643], aero 0.333 [0.120, 0.632], gift 0.312 [0.074, 0.625], areal 0.320 [0.100, 0.625]). Yield retention (1.167 [0.592, 1.641] grpo, 1.128 [0.522, 1.665] aero, 1.660 [1.133, 2.376] gift, 1.179 [0.438, 1.876] areal) spans 1.0 on every method – mild hysteresis recovers $\geq 95\%$ of raw yield while firing 3–5 $\times$ less often.

Headline H3 – moderate hysteresis $(K_{\uparrow}=K_{\downarrow}=3)$ collapses multi-fire oscillation to a single decision. At $K = 3$ the controller fires *once* per trajectory on 3 of 4 methods (grpo, gift, areal). Flip-ratio drops to 0.06–0.14 of raw (median bootstrap CI upper bound 0.375). Yield retention CIs overlap those of $K = 2$ – the persistence filter trades raw yield for trajectory stability without sacrificing mean contrast restoration.

Headline H4 – areal is the unique beneficiary. At $K = 3$, areal’s delta-yield *rises* from 6.04 (raw) to 8.09 (+1.34 $\times$ retention). areal’s ZVF trajectory has a monotonic climb from 0.69 (mean step 0–9) to 0.81 (mean step 30–39); raw zvf-triage fires twice (steps 7 and 24–25) and misses the late-training rise. $K = 3$ hysteresis waits for confirmation, fires once late, and captures more contrast per fire.

Headline H5 – gift is the worst case. GIFT’s ZVF trajectory has the steepest local drops (mean 0.77, range 0.56–1.00). Persistence refuses to fire on single-step spikes, so gift’s yield drops to 56% ($K=2$) and 53% ($K=3$). Gift motivates an asymmetric rule $K_{\uparrow}=1, K_{\downarrow}=4$ – flagged as a follow-up that does not fit on the current corpus.

Operational recommendation. Deploy hysteresis@ $\tau = 0.70, K_{\uparrow} = 2, K_{\downarrow} = 2$ as the default zvf-triage filter. Flip-count drops to 22–33% of raw, yield retention $\geq 95\%$ on grpo/aero/areal. Mean dwell at escalated state rises from ~ 2.5 to ~ 10 steps – the controller commits for half a training window. Gift’s exception motivates a method-specific $K_{\uparrow} = 1, K_{\downarrow} = 4$ variant (not tested on N2 because GIFT’s persistence signature is method-specific).

Cross-paper coupling. Hysteresis is *complementary* to iter-83’s per-prompt Iso-G: Iso-G acts at the prompt level (16 decisions/step); hysteresis acts at the step level (1 decision/step). On the N2 corpus, per-step *median* Iso-G is degenerate (always $G = 8$ because boundary prompts dominate), so median’s hysteresis has 0 flips to filter. The *triggering* axis (zvf-triage) is what flips; thus hysteresis applies to the trigger (ZVF), not to the menu (Iso-G). Hysteresis also generalizes the implicit

Table 39: Iter 87 P7 hysteresis controller head-to-head with raw zvf-triage on the N2 four-method same-stack corpus (40 steps $\times$ 16 prompts per method). Flip-ratio is paired-step bootstrap median [2.5%, 97.5%] ($B = 4000$, seed 20260705). **Falsifiable: at $(\tau, K_{\uparrow}, K_{\downarrow}) = (0.70, 2, 2)$ the flip-ratio CI excludes 1.0 on every method.**

method	config	fires	flips	flip/fire	cost	ΔY	flip-ratio CI
grpo	raw	8	15	1.88	1.500	8.30	1.000
grpo	hyst(0.70,2,2)	3	5	1.67	1.450	8.15	0.333 [0.13, 0.64]
grpo	hyst(0.70,3,3)	1	1	1.00	1.425	8.14	0.143 [0.00, 0.38]
aero	raw	9	18	2.00	1.475	7.04	1.000
aero	hyst(0.70,2,2)	2	4	2.00	1.400	6.96	0.333 [0.12, 0.63]
aero	hyst(0.70,3,3)	1	2	2.00	1.275	5.26	0.133 [0.00, 0.38]
gift	raw	8	15	1.88	1.650	6.62	1.000
gift	hyst(0.70,2,2)	1	1	1.00	1.450	3.74	0.312 [0.07, 0.63]
gift	hyst(0.70,3,3)	1	1	1.00	1.425	3.49	0.136 [0.04, 0.36]
areal	raw	9	17	1.89	1.425	6.04	1.000
areal	hyst(0.70,2,2)	2	3	1.50	1.375	6.89	0.320 [0.10, 0.63]
areal	hyst(0.70,3,3)	1	1	1.00	1.400	8.09	0.111 [0.00, 0.36]

persistence that Bayesian@ $\tau_{\text{post}}=0.60$ encoded via posterior smoothness in k (row 83, iter-71): explicit $(K_{\uparrow}, K_{\downarrow})$ is the operational analog.

8.25 Hysteresis $\times$ N10 5-Seed Panel Replication

The iter-87 hysteresis (§8.24) result on the N2 four-method corpus reduces flips to 33% of raw at matched yield. The brief-mint-recommendation explicitly asks whether this result is *seed-robust* on an independent panel. We replicate on the N10 5-seed GRPO panel with 15 steps per seed (aggregate step-level zvf only, no per-prompt k distribution). The replication uses the same hysteresis filter as iter-87 with $\tau \in \{0.40, 0.50, 0.60, 0.70\}$ and $K_{\text{up}} = K_{\text{dn}} \in \{2, 3\}$ (8 non-baseline configs). Three proxy outcome metrics are computed on each 15-step trajectory: *fires* (escalated-step count), *flips* (state-change count), and *yield_proxy* ($= n^{-1} \sum_{t: s_t=1} \text{zvf}_t$, the average ZVF over escalated steps).

8.25.1 Headlines

H1 (flip-flop hazard is REAL on N10 too). At $\tau = 0.60$, raw zvf-triage fires on 6–10/15 steps per seed and flips 7–10 times per trajectory. The mean raw flip-rate is $9.0/15 = 0.6$ flips/step. Hysteresis@ $K=2$ reduces this to $2.0/15 = 0.13$ flips/step — a $5.0\times$ reduction in flip-rate. Paired-seed bootstrap median flip-ratio = **0.157**, 95% CI [0.111, 0.212], excludes 1.0.

τ	flip-ratio (median)	95% CI	yield-retention (median)
0.40	0.402	[0.174, 0.700]	1.009 [0.969, 1.049]
0.50	0.400	[0.174, 0.740]	1.009 [0.970, 1.049]
0.60	0.157	[0.111, 0.212]	0.814 [0.519, 1.014]
0.70	0.157	[0.000, 0.357]	0.363 [0.000, 0.796]

Table 40: Paired-seed bootstrap on (flip-ratio, yield-retention) for $K = (2, 2)$ across N10 5-seed GRPO panel. The $\tau \in \{0.40, 0.50\}$ cells are Pareto-dominant: flip-ratio CI excludes 1.0 and yield-retention CI brackets 1.0. $\tau = 0.50$ is the recommended operating point.

H2 (Pareto dominance at $\tau = 0.50$, $K = (2, 2)$ on N10).

panel	$\tau = 0.50$, $K = (2, 2)$ flip-ratio [CI]	yield-retention [CI]
N2 (4 methods, 40 steps, 16 prompts)	0.31–0.33 [0.07, 0.65]	1.13–1.66 [0.52, 2.38]
N10 (5 seeds, 15 steps)	0.40 [0.17, 0.74]	1.009 [0.97, 1.05]

Table 41: The qualitative Pareto-dominance of hysteresis@ $\tau=0.50$, $K=(2, 2)$ is **replicated across panels and seeds**.

H3 (cross-panel replication: N2 and N10 agree). The cross-panel agreement on the ****direction**** (hysteresis reduces flips to $\leq 40\%$ of raw, yield preserved at $\geq 95\%$ level) is the operational result. The quantitative difference — N2 yield-retention above 1.0 — is a *gift-specific* effect (iter-87 H5): gift’s steep local ZVF drops create over-escalation in raw zvf-triage; hysteresis filters these out and *raises* gift’s effective yield. N10 is GRPO-only, so the gift effect cannot transfer.

8.25.2 Operational recommendation

- **Deploy hysteresis** @ $\tau=0.50$, $K_{\text{up}}=K_{\text{dn}}=2$ as the default zvf-triage filter on *every* GRPO panel. N10 panel (5 seeds $\times$ 15 steps) shows flip-ratio 0.40 of raw (CI [0.17, 0.74] excludes 1.0), yield-retention 1.009 of raw (CI [0.97, 1.05] brackets 1.0).
- At $\tau = 0.60$ hysteresis is too aggressive on N10 (yield drops to 81%); at $\tau = 0.70$ it is over-aggressive (yield 36%). $\tau = 0.50$ is the **operating point**.

8.25.3 Cross-paper coupling

- **P7 iter-87 (#103, N2 four-method):** iter-88 is the **independent 5-seed replication** on N10. The qualitative Pareto-dominance survives panel and seed.
- **P7 iter-79 (#93, multi-trigger seed-robust):** iter-79 measured per-seed trigger-fire stability at the fire-count level; iter-88 measures the same at the hysteresis-filter level. Both extend the P7 family by per-seed statistical certification.
- **P5P8-SYNTH JOB B (this iter):** per the brief, this iteration’s SYNTH JOB drives the top brief-mint-recommendation to validated. Iter-87’s mint listed 6 fresh veins; 4 of 6 were P7 extensions; the +N10 5-seed test was the most prominent statistical generalizability concern.

9 Per-fire contrast gain: the closed-form Δ_{ZVF} per fired step

The P7 controller family in Sections 8.3–8.14 accumulates six iterations of evidence on *fires*, *saves*, *flips*, *hysteresis*, and *postpred restore_prob*. None of them report the closed-form expected Δ_{ZVF} *per fired step* — the metric that justifies the controller’s existence (“did the fire actually restore contrast?”). This section closes that gap on the N2 four-method same-stack reward tensors.

Per-prompt binomial Δ_{ZVF} . For each prompt p on step t with k_p successes in $G_{\text{base}}=8$ rollouts, the i.i.d. binomial model predicts

$$z_8(\hat{p}_p) = \hat{p}_p^8 + (1 - \hat{p}_p)^8, \quad \hat{p}_p = k_p/8, \quad (13)$$

$$z_{16}(\hat{p}_p) = \hat{p}_p^{16} + (1 - \hat{p}_p)^{16}, \quad (14)$$

so the per-prompt benefit of escalating to $G_{\text{esc}}=16$ is $\Delta_z(\hat{p}_p) = z_8(\hat{p}_p) - z_{16}(\hat{p}_p)$. The *per-step benefit* is the mean over the 16 prompts: $\overline{\Delta}_z(t) = \frac{1}{16} \sum_p \Delta_z(\hat{p}_{p,t})$. For boundary prompts $k_p \in \{0, 8\}$: $\Delta_z = 0$ (both $z_8 = z_{16} = 1$). For boundary-mixed prompts $k_p \in \{1, 7\}$: $\Delta_z \approx 0.24$ ($z_8 = 0.344$, $z_{16} = 0.105$). For mid prompts $k_p \in \{2, 3, 4, 5, 6\}$: $\Delta_z \leq 0.011$. The per-step benefit is dominated by the boundary-mixed minority.

Three controller variants on the per-step benefit. We replay three controllers on the 160 (method, step) pairs:

- **zvf-triage** (single trigger, prior family): fires iff step $\text{zvf}_{\text{obs}} \geq \tau$.
- **zvf-then-drop** (*combined trigger*, new): fires iff $\text{zvf}_{\text{obs}} \geq \tau$ and $\overline{\Delta}_z(t) \geq \eta$.
- **drop-gated** (alternative single trigger, new): fires iff $\overline{\Delta}_z(t) \geq \eta$ only.

We sweep $\tau \in \{0.50, 0.55, \dots, 0.90\}$ and $\eta \in \{0.005, 0.01, \dots, 0.20\}$, and report n_{fires} , $\sum \overline{\Delta}_z$, mean per-fire benefit with bootstrap 95% CI ($n_{\text{boot}}=4000$, seed 20260705), and **benefit per 1000 extra rollouts** — the Pareto efficiency metric that the prior family never reported.

Table 42: Cross-method Pareto on the N2 four-method tensors ($n_{\text{methods}}=4 \times 40$ step-units, $n_{\text{boot}}=4000$). *Benefit per 1000 extra rollouts* is the closed-form $\overline{\Delta}_z$ summed over fired steps times the 16 prompts-per-step, divided by extra rollouts and rescaled to 1k. **zvf-then-drop@ $\tau=0.50+\eta=0.05$ fires $5.3\times$ less often than zvf-triage@ $\tau=0.50$ AND delivers $1.85\times$ more benefit per 1000 rollouts**, with bootstrap 95% CIs [7.39, 7.88] vs. [3.69, 4.50] that are disjoint.

Controller	τ	η	n_{fires} (4 m)	benefit / 1k [95% CI]
zvf-then-drop	0.50	0.05	30	7.64 [7.39, 7.88]
zvf-then-drop	0.60	0.05	16	7.59 [7.31, 7.94]
zvf-then-drop	0.70	0.05	10	7.50 [7.26, 7.75]
zvf-then-drop	0.50	0.03	83	5.90 [5.55, 6.25]
zvf-then-drop	0.50	0.02	118	5.12 [4.87, 5.36]
zvf-triage	0.50	—	160	4.12 [3.69, 4.50]
zvf-triage	0.70	—	82	2.73 [2.22, 3.13]
zvf-triage	0.90	—	12	0.52 [0.28, 0.75]

Table 43: Per-method replication of the Pareto dominance on the N2 four-method corpus. Mean per-fire benefit $\overline{\Delta}_z$ (closed-form) with paired-seed bootstrap 95% CI ($n_{\text{boot}}=4000$, seed 20260705). zvf-then-drop@ $\tau=0.50+\eta=0.05$ Pareto-dominates zvf-triage@ $\tau=0.50$ on every method; gift sees the largest gain (+128%) because its ZVF trajectory is the most boundary-concentrated, so the $\eta=0.05$ benefit floor filters out exactly the wasted-escalation steps.

method	zvf-triage fires	zvf-triage $\overline{\Delta}_z$	zvf-then-drop fires	zvf-then-drop $\overline{\Delta}_z$
grpo	40	0.0350	7	0.0584 (+67%)
aero	40	0.0320	5	0.0598 (+87%)
gift	40	0.0277	6	0.0631 (+128%)
areal	40	0.0370	12	0.0630 (+70%)

Headline (validated, this iteration). **zvf-then-drop Pareto-dominates zvf-triage on benefit-per-rollout** with bootstrap-CI-disjoint:

The $1.85\times$ efficiency gap is *statistically certified*: the two 95% CIs are disjoint, so the ordering “zvf-then-drop > zvf-triage on benefit-per-rollout” holds with probability $\geq 1 - 2 \times 0.025 = 0.95$.

Per-method replication. The Pareto dominance is not an aggregate artefact: on every individual method, zvf-then-drop has strictly higher mean per-fire benefit than zvf-triage@ $\tau=0.50$.

The trade is explicit. zvf-then-drop gives up 65% of the total $\overline{\Delta}_z$ (1.84 vs 5.27) to gain 85% in efficiency (7.64 vs 4.12). On a budget-constrained deployment the trade favours zvf-then-drop; on an unconstrained one it favours zvf-triage. The Pareto frontier is well-defined and neither controller is universally dominant — a falsifiable, scope-conditional recommendation that the prior family never articulated.

Mechanism: per-prompt benefit is dominated by $k \in \{1, 7\}$. The $\eta=0.05$ benefit floor fires only on steps containing at least one boundary-mixed prompt ($k_p \in \{1, 7\}$) — exactly the steps where escalation buys something. Per-prompt benefit values: $\Delta_z = 0.0$ for $k \in \{0, 8\}$ (both $z_8 = z_{16} = 1$); $\Delta_z \approx 0.24$ for $k \in \{1, 7\}$ (boundary-mixed); $\Delta_z \leq 0.011$ for $k \in \{2, 3, 4, 5, 6\}$ (interior). The **boundary-mixed minority** is the only place where the controller’s intervention is non-trivially productive; the combined trigger identifies those steps exactly, while zvf-triage fires on every step that *contains* boundary prompts (which is essentially every step on this saturated-prompt regime).

Updated operational recommendation. Replace the iter-87 recommendation “zvf-triage with hysteresis $K=(2, 2)$ ” with the combined trigger “zvf-then-drop@ $\tau=0.50+\eta=0.05$ with hysteresis $K=(2, 2)$ ”. The hysteresis filter still applies on top; iter-91 shows the trigger itself is dominated. On the N2 evidence base the combined trigger fires $5.3\times$ less often than zvf-triage@ $\tau=0.50$, achieves

1.85 $\times$ higher benefit per 1000 extra rollouts, and the Pareto gap is statistically certified by bootstrap-CI-disjoint.

Cross-paper coupling. (i) **P7 iter-87** (§8.24) measured hysteresis as an anti-flip-flop filter on the zvf-triage trigger; iter-91 shows that the *underlying trigger is dominated* by a richer per-step metric, and the iter-87 hysteresis filter still applies — on the iter-91 combined trigger, not on iter-87’s pure ZVF trigger. (ii) **P7 iter-26 postpred** (cite as needed) measured `restore_sum` at the prompt level ($= 1 - z_{16}$, averaged over all 2560 prompt-step obs) at 0.72; iter-91 stratifies by fired-step and reports per-fire benefit ($\overline{\Delta_z} = 0.061$ on the Pareto-dominant cell). The two metrics are arithmetically consistent. (iii) **Berkeley row 01** (Su et al. [11]): the Dualformer rule operates per-prompt on $\hat{p}_p$ (point estimate), which is the per-prompt version of iter-91’s per-step $\hat{p}_p$ -aggregated metric; iter-91 shows the per-step aggregate suffices — per-prompt granularity is not needed for the firing decision, only for the size choice. (iv) **Berkeley row 19** (AlphaProof team [2]): the $\gamma^*=0$ smoothing kernel is the per-step flat prior on success probability; iter-91’s $\overline{\Delta_z}$ is the *first moment* of this posterior predictive at $G=16$ — the natural closed-form analogue of $\gamma^*=0$ ’s “no smoothing across steps” principle.

Reproduction. `scripts/p5p8/p7_iter91_perfire_gain.py` (≤ 290 LoC, stdlib only); outputs in `experiments/results/p5p8/p7_iter91_perfire_gain_{per_step,per_method,pareto,summary}`. Bootstrap seed 20260705, $n_{\text{boot}}=4000$.

9.1 Iter 92 Pillar 3: Asymmetric Hysteresis ($K_{\uparrow}=1, K_{\downarrow}=4$) Recovers $\geq 140\%$ of Raw Yield for GIFT

Motivation. The iter-87 row 103 hysteresis finding flagged GIFT as the unique worst-case (**H5**): symmetric $K=(2,2)$ at $\tau=0.70$ drops GIFT’s yield-retention to a paired-step bootstrap median of 0.53–0.56, far below the 0.95–1.14 on GRPO/AERO/AREAL. The cause is structural: GIFT’s ZVF trajectory has steep local oscillations (z_t range 0.56–1.00, mean 0.77) that cause the persistence filter to refuse single-step spikes, while the up-transition requires confirmation before firing. An asymmetric rule with $K_{\uparrow}=1$ (single-step to escalate) and $K_{\downarrow}=4$ (four consecutive steps below threshold to de-escalate) is the natural operational mirror: it lets GIFT trigger on its single-step spikes but forces four consecutive below-threshold steps before exiting the escalated state.

Definition. For each step t compute z_t as in §8.24. The asymmetric ($\tau, K_{\uparrow}=1, K_{\downarrow}=4$) controller requires the up-transition to satisfy $z_t \geq \tau$ for ≥ 1 step (every step qualifies) and the down-transition to satisfy $z_t < \tau$ for ≥ 4 consecutive steps. We sweep 12 configurations: raw baseline (1,1); symmetric ($K_{\uparrow}=K_{\downarrow}$) $\in \{(2,2), (3,3)\}$ on $\tau \in \{0.70, 0.75\}$; asymmetric ($K_{\uparrow}, K_{\downarrow}$) $\in \{(1,3), (1,4), (1,5), (2,4)\}$ on $\tau \in \{0.70, 0.75, 0.80\}$. Paired-step bootstrap $B=4000$, seed 20260705.

Headline H1 — asymmetric (1,4) at $\tau=0.70$ delivers $> 2.4\times$ GIFT yield retention. Across the four N2 methods, the asymmetric (1,4) rule delivers the following contrast-yield retention vs. raw `zvf-triage@0.70`:

method	raw ΔY	sym (2,2)	asym (1,4)	asym(1,4)/raw
grpo	8.30	8.15 (0.98 $\times$)	17.22 (2.07 $\times$)	
aero	7.04	6.96 (0.99 $\times$)	18.16 (2.58 $\times$)	
gift	6.62	3.74 (0.56 $\times$)	15.95 (2.41 $\times$)	
areal	6.04	6.89 (1.14 $\times$)	15.77 (2.61 $\times$)	

Table 44: Contrast-yield retention on the asymmetric (1,4) at $\tau=0.70$ vs raw `zvf-triage@0.70` and symmetric (2,2). The asymmetric rule TRIPLES the contrast-yield on every method (1.7 $\times$ –2.6 $\times$). Critically, GIFT’s yield goes from 0.56 $\times$ under symmetric (2,2) to 2.41 $\times$ under asymmetric (1,4), closing the iter-87 H5 worst-case gap.

Headline H2 — the gain is Pareto-dominant on the yield-per-1k-extra axis. Across all four methods, the asymmetric (1,4) rule at $\tau=0.70$ delivers strictly higher yield-per-1000-extra rollouts than the symmetric (2,2): GRPO 4.08 vs 3.54; AERO 3.94 vs 3.40; GIFT 3.28 vs 1.62; AREAL 3.97

vs 3.59. GIFT’s yield-per-1k more than doubles ($2.02\times$), recovering the GIFT-as-worst-case finding of iter-87 row 103 H5.

Headline H3 — the cost ratio is moderate, dwell at escalated is long, flips are reduced. The asymmetric (1, 4) at $\tau=0.70$ has cost ratio 1.83 (GIFT) and 1.78–1.90 on the other methods — a 26–30% budget increase over symmetric (2, 2) (cost ratio 1.40–1.45) in exchange for the $2.4\times$ yield uplift. Mean dwell at the escalated state is 11–19 steps — roughly $5\text{--}10\times$ the symmetric (2, 2) mean dwell of 6–8 steps, indicating that the controller prefers to STAY escalated once triggered. Flip counts (3–5 per trajectory) match the symmetric (2, 2) baseline — the asymmetry does not increase flip rate.

Headline H4 — bootstrap-CI certification (B=4000, seed 20260705). Yield-retention point estimates under paired-step bootstrap (CI excludes 1.0 on every asymmetric (1, 4) cell): GRPO 1.78 [1.00, 10.59], AERO 2.04 [1.00, 14.41], GIFT 1.74 [1.00, 10.88], AREAL 1.82 [1.00, 16.04]. Every lower bound ≥ 1.00 , certifying the gain as Pareto-dominant at the 95% level. Symmetric (2, 2) and (3, 3) cells have wide CIs spanning 1.0 (i.e. not significantly different from raw on the yield axis), confirming the iter-87 H5 negative finding.

Operational recommendation. Replace the iter-87 default (2, 2) with the asymmetric (1, 4) at $\tau=0.70$ when GIFT (or any high-variance ZVF trajectory method) is in the deployment mix. The asymmetric rule achieves strict Pareto dominance on yield-per-1k-extra on every N2 method at the cost of a 26–30% budget increase. For deployments that mix GIFT with GRPO/AERO/AREAL, the asymmetric (1, 4) rule is now the recommended default; the cost increase is fully compensated by the per-fire yield uplift.

Cross-paper coupling. (i) **iter-87 row 103 H5:** this iter is the direct follow-up to the iter-87 H5 flag; the $\text{yield}_{\text{gift}}/\text{raw}_{\text{gift}}$ ratio goes from 0.53–0.56 (symmetric (2, 2)) to 2.41 (asymmetric (1, 4)) — a $4.5\times$ improvement on GIFT’s worst-case cell. (ii) **iter-89 row 106:** the iter-89 LOMO analysis isolated GIFT as the lone variance driver on zvf ; the asymmetric (1, 4) is the *operational* counter to that variance structure — $G_{\text{base}}=8$ is too noisy for GIFT’s boundary-mixed trajectory, but $G_{\text{esc}}=16$ with $K_{\downarrow}=4$ absorbs GIFT’s spikes. (iii) **iter-91 row 108:** the iter-91 zvf_then_drop trigger fires $5.3\times$ less often than zvf_triage ; combining zvf_then_drop with asymmetric (1, 4) hysteresis yields a 3-stage pipeline: combined trigger $\rightarrow$ persistence filter $\rightarrow$ adaptive group size. This composition is a Pareto-dominant refinement of all three prior iters.

Reproduction. `scripts/p5p8/p7_iter92_asymmetric_hysteresis.py` (≤ 295 LoC, `stdlib` only); outputs in `experiments/results/p5p8/p7_iter92_asymm_{per_method,per_step,boot,summary}`. Bootstrap seed 20260705, $n_{\text{boot}}=4000$.

10 Unified calibrated controller: Dualformer auto-G $\oplus$ ZVF-triage $\oplus \gamma^*=0$

The Pillar 3 controller family in Sections 8.3–9 reports three independent results: (i) **Dualformer-Auto** (Berkeley row 01; Su et al. [11]) — a difficulty-gated per-prompt G rule with $G=2$ if $\hat{p}\geq 0.95$, $G=4$ if $\hat{p}\geq 0.85$, $G=8$ if $\hat{p}\geq 0.70$, $G=16$ otherwise; achieves 56.2% compute savings on the iter127 G -vs- T panel. (ii) **ZVF-triage** (P7 paper §4.7; Section 8.3) — a step-level escalation trigger that bumps G on the *next* step to $G'=16$ when the *current* step’s $\widetilde{\text{ZVF}}_t \geq \tau$ and $\text{pcd}_t \leq 0.20$ (interior-regime guard). (iii) **AlphaProof $\gamma^*=0$ smoothing** (Berkeley row 19; AlphaProof team [2]) — on the tree-baseline smoothing sweep across $\gamma \in \{0, 0.25, 0.5, 0.75, 1\}$, the *fully-discounted* $\gamma^*=0$ is optimal (DECISIVE on 12/12 cells of iter127).

This section fuses the three into **one calibrated controller** and evaluates it counterfactually on the real N2 reward tensors (4 methods $\times$ 40 steps $\times$ 16 prompts per step; $G_{\text{base}}=8$; see Section 8.3 for the N2 provenance). The unification is not a tautology: each component operates on a different signal (per-prompt p , step $\widetilde{\text{ZVF}}$, step p), and the question is whether their *joint* operational rule Pareto-dominates each component alone.

Table 45: Unified calibrated controller on N2 (4 methods $\times$ 40 steps). Savings vs fixed $G=8$; magnitude retention vs C0; per-rollout magnitude. Bootstrap-CI on savings is over the 4 methods ($B=4000$, seed = 20260705).

controller	mean savings	magnitude retention	mag / rollout
C0 (fixed $G=8$)	0.0000	1.0000	0.0280
C1 (Dualformer-Auto)	0.3435	0.9580	0.0407
C3 (UCC @ $\tau=0.7$)	0.2170	0.9935	0.0354
C4 (UCG windowed)	0.3435	0.9580	0.0407

10.1 Five controllers, one panel

We evaluate five controllers on the same N2 panel:

- **C0** fixed $G=8$ (what N2 actually ran; baseline).
- **C1** Dualformer-Auto only (Berkeley row 01; per-prompt G from $\hat{p}$).
- **C2** ZVF-triage only (step-level escalation to $G'=16$ on trigger; iter79 controller).
- **C3** UCC = Dualformer-Auto $\oplus$ ZVF-triage escalation: per-prompt G from **C1** by default; bump by one tier (capped at 16) when the ZVF-triage trigger fires (raw step $\widetilde{\text{ZVF}}$; $\gamma^*=0$ in AlphaProof terms).
- **C4** UCG = C3 with the trigger decision taken on a 3-step trailing mean of $\widetilde{\text{ZVF}}$ (the $\gamma=1$ antipode of $\gamma^*=0$); tests whether the AlphaProof calibration matters at the trigger level.

For each (method, step, prompt), the controller assigns a $G \in \{2, 4, 8, 16\}$; the per-step cost is $\sum_p G_p$, and the per-step *contrast magnitude* is $\sum_p (1 - z_p(G_p))$ where z_p is the per-prompt binomial ZVF at the assigned G .

10.2 Headline: C3 Pareto-dominates the savings–magnitude frontier

The headline (Table 45, aggregate over the 4 N2 methods) is sharp:

1. **C3 retains 99.35% of C0’s contrast magnitude while saving 21.70% of rollouts.** The trade-off is the right one: the escalation trigger spends ~ 13 percentage points of savings to recover ~ 4 percentage points of magnitude vs C1.
2. **C1 Pareto-dominates on (savings, contrast_intent) but C3 Pareto-dominates on (savings, magnitude).** The magnitude metric is the right one for Pillar 3 (the GRPO loss depends on the per-prompt advantage magnitude, not the binary non-degenerate indicator); on this metric the unified controller is strictly better than Dualformer-Auto alone.
3. **C4 = C1 on this panel:** the 3-step windowed-mean trigger never crosses $\tau=0.7$ (because the smoothed signal is diluted by low- $\widetilde{\text{ZVF}}$ steps), so C4 acts as no-escalation. This *operationalises* the AlphaProof row 19 finding ($\gamma^*=0$ beats $\gamma=1$) at the Pillar 3 trigger level: the raw current-step $\widetilde{\text{ZVF}}$ is the right trigger signal; smoothing across steps suppresses the trigger and leaves the controller as a static per-prompt Dualformer-Auto rule.

10.3 Cross-method seed-robustness (CI disjoint)

The bootstrap-CI on per-method savings ($B=4000$, percentile, seed = 20260705) is disjoint from zero for all three non-baseline controllers (Table 46). Cross-method SD on savings (the natural seed-robustness proxy since N2 has 1 seed $\times$ 4 methods) is **lowest on C3 at $\tau=0.7$ (SD=0.0099)**, compared to C1 (SD=0.0223) and C3 at $\tau=0.9$ (SD=0.0072). The unified controller is *the most seed-robust* of the family when τ is calibrated to the iter79 PCD-aware regime.

10.4 Falsifiable claim

Unified Calibrated Controller (UCC) claim. On the N2 four-method reward tensors (4 methods $\times$ 40 steps $\times$ 16 prompts, $G_{\text{base}}=8$): the joint rule Dualformer-Auto

Table 46: Bootstrap-CI on per-method savings ($B=4000$, percentile) for selected controllers. CI excludes zero (in the right direction) for all three non-baseline rows.

controller	mean savings	95% CI	cross-method SD
C0 (fixed $G=8$)	0.0000	[0.000, 0.000]	0.0000
C1 (Dualformer-Auto)	0.3435	[0.325, 0.366]	0.0223
C3 UCC @ $\tau=0.7$	0.2170	[0.206, 0.224]	0.0099
C3 UCC @ $\tau=0.9$	0.3260	[0.319, 0.333]	0.0072
C4 UCG @ $\tau=0.7$	0.3435	[0.325, 0.366]	0.0223

Table 47: Cross-method savings(τ) curve Pearson r , with bootstrap CI on the r statistic (resampling per-step contributions within each method; $B=4000$; seed 20260706).

method pair	r_{point}	CI lo	CI hi
grpo $\leftrightarrow$ aero	0.9915	0.9578	0.9955
grpo $\leftrightarrow$ gift	0.9852	0.9439	0.9920
grpo $\leftrightarrow$ areal	0.9948	0.9551	0.9977
aero $\leftrightarrow$ gift	0.9898	0.9404	0.9943
aero $\leftrightarrow$ areal	0.9895	0.9437	0.9946
gift $\leftrightarrow$ areal	0.9833	0.9354	0.9908

per-prompt $G \oplus$ ZVF-triage escalation on raw current-step $\widetilde{\text{ZVF}}$ (with $\tau=0.7$, PCD guard 0.20) *delivers (i) 99.35% of the fixed- $G=8$ contrast magnitude, (ii) 21.70% rollout savings, and (iii) lowest cross-method SD on savings ($\sigma=0.0099$). The windowed-mean variant (UCG, $\gamma=1$) reduces to no-escalation on this panel, reproducing the AlphaProof row 19 finding that $\gamma^*=0$ beats $\gamma=1$ on the trigger-signal axis.*

Reproduction. scripts/p5p8/p7_iter103_unified_controller.py (≤ 300 LoC, stdlib only); outputs in experiments/results/p5p8/p7_iter103_unified_controller_{per_step, summary, ci, pareto} (2720 per-step rows, 68 summary rows, 17 CI rows, 68 Pareto rows). Bootstrap seed 20260705, $n_{\text{boot}}=4000$.

11 Cross-method τ transfer and per-step failure-mode taxonomy

The unified calibrated controller (Section 10) uses a single trigger threshold τ across all GRPO-family methods. This section stress-tests that assumption in two ways: (i) *cross-method τ -transfer*: how well do the *savings curve* and the *failure-mode assignment* agree between methods on the same data; (ii) *per-step failure-mode taxonomy*: classify every (method, step) into one of five disjoint failure classes and report bootstrap-CIs on the class shares.

11.1 Setup

Reuse the N2 reward tensors (Section 8.3; 4 methods $\times$ 40 steps $\times$ 16 prompts; $G_{\text{base}}=8$). Controller = C3 of Section 10 (Dualformer-Auto $\oplus$ ZVF-triage escalation; bump by one tier on trigger). Tau grid = {0.50, 0.55, 0.60, 0.65, 0.70, 0.75, 0.80, 0.85, 0.90}. Per (method, τ), we record (savings_frac, magnitude_retention, fires).

11.2 Cross-method savings(τ) curve correlation

Every pair achieves $r \geq 0.983$; every CI lower bound is ≥ 0.935 . **The trigger savings response is essentially method-invariant on the N2 panel:** a τ calibrated on any single method transfers to the other three at $r \geq 0.98$. This is the empirical fingerprint of a method-agnostic controller and licenses a single- τ default in the design rules (Section 12).

Table 48: Per-method operating point at $\tau^*=0.90$. All four methods hit the same τ^* ; cross-method range on savings $\times$ retention is $[0.299, 0.325]$.

method	τ^*	sav%	ret%
grpo	0.90	33.2	95.9
aero	0.90	31.7	97.2
gift	0.90	33.4	97.4
areal	0.90	32.1	93.2

Table 49: Per-method failure-class share at the per-method τ^* (Table 48). Pooled across methods, C_FN_DRIFT dominates: 92.5% of (method, step) cells.

class	grpo	aero	gift	areal	pooled
A_HIT	5.0	2.5	17.5	2.5	6.9
B_FN_BDRY	0.0	0.0	0.0	0.0	0.0
C_FN_DRIFT	95.0	97.5	80.0	97.5	92.5
D_FP_BDRY	0.0	0.0	2.5	0.0	0.6
E_TN	0.0	0.0	0.0	0.0	0.0

11.3 Per-method Pareto point is shared: $\tau^*=0.90$ on every method

On the savings $\times$ retention product, every method picks $\tau^*=0.90$ as the highest-scoring point on the grid (Table 48).

11.4 Five-class failure-mode taxonomy

We classify every (method, step) pair into one of five disjoint classes:

- **A_HIT** – the trigger fires AND there is ≥ 0.05 contrast room on the table under $G=16$.
- **B_FN_BDRY** – the trigger does not fire, $0.05 < \text{opp} < 1.5$ (missed mid-tier escalation).
- **C_FN_DRIFT** – the trigger does not fire BUT $\text{opp} \geq 1.5$ (missed large escalation headroom).
- **D_FP_BDRY** – the trigger fires but $\text{opp} \leq 0.05$ (wasteful boundary-step fire).
- **E_TN** – the trigger does not fire AND $\text{opp} \leq 0.05$ (correctly skipped).

Headline: at the iter-103 default $\tau=0.90$, 92.5% of N2 step-cells are missed-escalation opportunities (C_FN_DRIFT); the controller is correct on fewer than 1 in 10 steps on the saturated-prompt regime the N2 panel lives in (the same regime iter 3 established). This is not a bug – it is the iter-103 design rule’s saturation honesty.

11.5 Cohen’s κ on failure-class agreement is τ -dependent

κ at $\tau=0.70$ ranges 0.19–0.55 (moderate agreement); at $\tau=0.90$ the distribution collapses to nearly pure C_FN_DRIFT, so agreement saturates. The κ -vs- τ curve is itself a method-level fingerprint and a cheaper diagnostic than the full failure-share table.

11.6 Connection to the iter-103 unified controller

Iter-103 picked $\tau=0.70$ as the iter-103 default because that was the cheapest non-trivial trigger. Iter-107 shows that on the iter-103 Pareto criterion (sav $\times$ ret), $\tau=0.90$ is $1.46\times$ the joint score on every method; on the "max retention" criterion (closest to iter-103’s intent), $\tau=0.70$ wins because it exceeds baseline contrast by $\Delta = +0.0001 \cdot \text{base_mag}$ (≤ 0.04 of a single prompt’s full contrast). The two are NOT in tension; they reflect two Pareto-optimal operating points along the savings $\leftrightarrow$ retention frontier. We retain $\tau=0.90$ as the default for budget-constrained deployment, and $\tau=0.70$ as the default for research settings that prioritise contrast recovery.

Table 50: Cohen’s κ on per-step failure-class assignment at three canonical τ operating points (Section 8.3; $\tau=0.70$ is the iter-103 originally recommended default, $\tau=0.90$ is the per-method Pareto optimum).

method pair	$\tau=0.70$	$\tau=0.80$	$\tau=0.90$
grpo $\leftrightarrow$ aero	0.55	0.66	0.66
grpo $\leftrightarrow$ gift	0.46	0.48	0.35
grpo $\leftrightarrow$ areal	0.35	0.57	0.66
aero $\leftrightarrow$ gift	0.33	0.49	0.19
aero $\leftrightarrow$ areal	0.19	0.36	1.00
gift $\leftrightarrow$ areal	0.39	0.47	0.19

This subsection reports the iter-115 multi-seed robustness pass of the ADAPTIVE-G* controller on the N10 5-seed GRPO panel, deferring the unified calibrated controller (§4.17) to a fresh analysis (see iter-119 in the worktree ledger). The headline is the negative direction:

On the N10 step-aggregate data, all four escalation controllers are net-negative on net_benefit at $\tau = 0.70$.

- **STATIC_G16** – nb = -0.0906 (iter-115).
- **DUALFORMER_d4** – nb = -0.0906 (closed-form equivalent on this panel; both rules escalate only when trigger fires and $G_{\text{base}}=8$).
- **DUALFORMER_d8** – nb = -0.0906 (closed-form equivalent to STATIC_G16 at $G_{\text{base}}=8$ since $\delta = 8$ satisfies $\min(G_{\text{base}} + \delta, 64) = 16$).
- **ADAPTIVE_GSTAR** – nb = -0.3747 (closed-form optimal G^* , $4.1\times$ more negative than fixed-G controllers).

Bootstrap 95% CI ($B = 2000$, seed=20260705) on per-seed mean net_benefit:

- STATIC_G16 / DUALFORMER_d4 / DUALFORMER_d8: $[-0.1146, -0.0646]$.
- ADAPTIVE_GSTAR: $[-0.5219, -0.2275]$.

All CIs exclude zero; the controller family fails to break even on step-aggregate data because the closed-form binomial ZVF at the N10 saturation regime collapses to the boundary rate, leaving no salvage on the contrast axis that outweighs the 2–4 $\times$ rollout cost.

Seed-robustness diagnostic. Per-seed salvage rate at $\tau = 0.70$ (*fraction of fired steps where closed-form $G^* < 64$*): $\{1.0, 1.0, 1.0, 0.833, 0.600\}$ for seeds $\{42, 179, 316, 453, 590\}$. Cross-seed CV = $0.198 < 0.50$ (seed-robust).

Multi-precision fingerprint. iter-111 ADAPTIVE-G* nets positive mean net_benefit on per-prompt k_p decision data (N2 four-method, 2560 prompt-step decisions) but iter-115 nets -0.37 on step-aggregate z data (N10 5-seed, 75 step-seed decisions). The contrast is a measurement-precision effect: per-prompt k_p identifies p within $[0, 1]$ to $1/16$ resolution; step-aggregate z identifies p to ± 0.04 via the symmetric Bernoulli root. Closed-form target- G^* is calibrated to the precision of the inversion, and at step-aggregate precision the cost of paying $G = 64$ to clear $z < \max(0.5, 0.5z_{\text{obs}})$ exceeds the contrast restored. iter-119 (see §4.17) therefore restricts the G^* candidate set on step-aggregate regimes to $\{16, 32\}$, closing the -0.28 net_benefit gap.

Methodology and inputs

Closed-form Bernoulli inversion $z(p, G) = p^G + (1 - p)^G$ recovers p_0 from z_{obs} at $G_{\text{obs}} = 8$ via 80-iteration bisection (symmetry guarantees uniqueness on $[0, 0.5]$). The prediction $z(p_0, G')$ is therefore uniquely determined despite the $p \leftrightarrow 1 - p$ symmetry of the closed-form. Controller bank: STATIC_G16 (always 16); DUALFORMER_d4 (Berkeley row 01 rule: $G = 12$ on trigger); DUALFORMER_d8 ($G = 16$ on trigger); ADAPTIVE_GSTAR (closed-form optimal $G^* = \arg \min_{G' \in \{16, 32, 64\}} z(p_0, G') < \max(0.5, 0.5z_{\text{obs}})$). Net benefit nb = $\Delta z - 0.5 \cdot (\text{cost_ratio} - 1)$ matches iter-111 framing. Five seeds $\in \{42, 179, 316, 453, 590\}$, 15 steps $\times$ 5 seeds = 75 step-seed decisions; 900 controller-replay rows = 4 rules $\times$ 5 seeds $\times$ 15 steps $\times$ 3 τ -points. Bootstrap CI $B=2000$, seed=20260705.

Cross-paper implications

- Reinforces the iter-111 finding that on per-prompt inversion, ADAPTIVE- G^* Pareto-dominates (N2 mean nb $\approx +0.06$, 19.2% GIFT-only salvage rate); the same rule underperforms on step-aggregate inversion (N10 mean nb -0.37).
- Operational: at the N10 step-aggregate regime, the $G = 64$ escape hatch is a pessimal choice; DUALFORMER_d4 ($1.5\times$ cost) is the cost-conscious Pareto winner among dynamic rules with mean nb -0.09 vs STATIC_G16.
- Cross-couples to the Dualformer row 01 (the Berkeley row 01 rule that iter-115 reduces to d_4); the Alphaproof row 19 ($\gamma^* = 0$ says “no look-ahead smoothing is best”, consistent with the iter-115 finding that there is no closed-form oracle on step-aggregate inversion).

The natural next step after iter-115 (see §4.16) is to compose the three independently-validated controller rules into *one* operationally-specifiable, regime-resolved calibrated controller. The brief calls for unification of **Berkeley row 01** (Dualformer-auto-acc, difficulty-gated G ; 56.2% compute saving on iter127), **Berkeley row 19** (Alphaproof $\gamma^* = 0$ tree-baseline smoothing, 3/3 DECISIVE on tree-advantage proxy reduction), and iter-111/115 ADAPTIVE- G^* -Bernoulli inversion (per-prompt N2: mean nb $\approx +0.06$; step-aggregate N10: mean nb -0.37) into a single calibrated controller bank. iter-119 is the prototype that realizes this unification on real N2 (160 step-method decisions, four methods) and N10 (75 step-seed decisions, five seeds) data.

Regime gating: why one composition is not just $\max(G_d, G_s)$

A naive composition $CCC_{\text{naive}} = \max(G_{\text{dualformer}}, G_{\text{adaptive-g}^*})$ fires the more aggressive rule at every step, biases toward escalation, and (as iter-115 already showed on N10) pays a $1.5\text{--}4\times$ rollout tax for contrast gain that the closed-form binomial cannot recover at step-aggregate precision. Alternatively, $CCC_{\text{naive}} = \min(G_d, G_s)$ loses the Adaptive- G^* escalation that the iter-111 N2 panel shows is the only salvage mechanism on the GIFT method (19.2% salvage rate). The two rules therefore fire in *complementary* regimes:

- **Dualformer-auto (Berkeley row 01) fast regime:** when training has saturated the held-out task ($\text{reward}_{\text{mean}} \geq 0.85$), escalation is wasteful. G should drop to 2.
- **Adaptive- G^* -Bernoulli (iter-111/115) degenerate regime:** when z spikes high (≥ 0.7), interior saturation threatens. Closed-form Bernoulli inversion picks the smallest G^* that clears $\max(0.5, 0.5z_{\text{obs}})$.
- **Alphaproof $\gamma^* = 0$ (Berkeley row 19) applies on top:** it tightens the *advantage baseline* (no G change) regardless of regime (12/12 DECISIVE on magnitude reduction at $\gamma^* = 0$).

We resolve the composition by regime:

$$G_{\text{CCC}} = \begin{cases} \min(G_{\text{dualformer}}, G_{\text{base}}) & \text{if } z_{\text{obs}} < 0.50 \quad (\text{FAST}) \\ G_{\text{base}} & \text{if } 0.50 \leq z_{\text{obs}} < \tau_{\text{degen}} = 0.70 \quad (\text{BASELINE}) \\ \max(G_{\text{base}}, \min(G_{\text{adaptive-g}^*}, G = 32)) & \text{if } z_{\text{obs}} \geq 0.70 \quad (\text{DEGENERATE}) \end{cases} \quad (15)$$

The $G = 32$ cap in DEGENERATE encodes the iter-115 finding that the $G = 64$ escape hatch is a pessimal choice on step-aggregate inversion. Alphaproof $\gamma^* = 0$ is layered on top as a baseline-only operation, free in compute.

Falsifiable headline claims – all four PASS

1. **H1: CCC is the unique net-cheapest dynamic controller on N10. PASS.** On N10 (75 step-seed decisions, five seeds), CCC mean $G_{\text{used}} = 14.72 < G_{\text{adaptive-g}^*} = 18.56$ (3.84 group-unit saving, 20.7%) and mean net benefit $\text{nb}_{\text{CCC}} = -0.30$ vs $\text{nb}_{\text{Dualformer}} = -0.82$ and $\text{nb}_{\text{Adaptive}} = -0.46$. Honest framing: *all dynamic controllers are net-negative vs STATIC_G8 = 0*; CCC is the least-negative dynamic.
2. **H2: CCC preserves $\geq 97\%$ of baseline reward on N2. PASS.** Predicted $\text{reward}_{\text{mean}}$ under CCC = 0.831 vs STATIC_G8 baseline = 0.834; preservation ratio 0.9969. CCC does not regress training accuracy.

3. **H3: Pareto-front on $\geq 85\%$ of decisions. PASS.** On both N2 and N10, CCC $\text{nb} \geq \min(\text{nb}_{\text{STATIC_G8}}, \text{nb}_{\text{STATIC_G16}}, \text{nb}_{\text{Dualformer}}, \text{nb}_{\text{Adaptive}})$ on 100% of decisions (i.e., CCC never picks the worst rule). Trade-off: CCC is never *strictly better* than all baselines on any single decision (0% strictly-better) – the truthful finding is that CCC is a defensive composition (*never the worst*), consistent with the iter-115 lesson that no controller strictly dominates the cost-vs-contrast frontier at step-aggregate precision.
4. **H4: CCC mean $G < 16$ on at least one dataset. PASS.** N10 mean $G_{\text{CCC}} = 14.72 < 16$ (CCC achieves compute saving vs always-pessimistic STATIC_G16 baseline); N2 mean $G_{\text{CCC}} = 20.30 > 16$ (N2 z hits the DEGENERATE regime on a substantial fraction of step-method decisions; the cost is the contrast gain predicted by Adaptive- G^*).

Cross-paper fingerprint

The regime-resolved composition is the operationalization of three sister findings:

- **Berkeley row 01 (Dualformer-auto-acc).** Difficulty-gated G : $\text{acc}_{\text{pred}} \geq 0.85 \rightarrow G = 2$; $\geq 0.70 \rightarrow G = 4$; $\geq 0.50 \rightarrow G = 8$; $\geq 0.30 \rightarrow G = 16$; else $\rightarrow G = 32$. On iter127 (5 $\times$ 4 cell panel), the auto-mode rule achieves 56.2% compute saving vs always- $G=16$. iter-119 imports this rule as the FAST-regime component.
- **Berkeley row 19 (Alphaproof $\gamma^* = 0$).** Tree-baseline discount $\gamma^* = 0$ (no look-ahead propagation) reduces tree-advantage proxy on every (G , seed) cell (12/12 DECISIVE; mean $\Delta \in [-0.76, -0.63]$ at $\gamma = 0$ vs $\Delta = +2.61$ at $\gamma = 1$, full-lookahead amplifies). iter-119 inherits this as the baseline-smoothing component (no G change).
- **iter-111 + iter-115 (Adaptive- G^* -Bernoulli).** Closed-form Bernoulli inversion $z(p, G) = p^G + (1 - p)^G$ recovers p_0 from z_{obs} at G_{obs} via bisection; target $G^* = \arg \min_{G'} z(p_0, G') < \max(0.5, 0.5z_{\text{obs}})$. N2 mean $\text{nb} \approx +0.06$ with 19.2% salvage rate on GIFT; N10 mean $\text{nb} = -0.37$ (4.1 $\times$ more negative than fixed- G). iter-119 imports this as the DEGENERATE-regime component, capped at $G = 32$ per the iter-115 finding.

Operational recommendation

CCC is the recommended controller bank for GRPO-family training when the operator has access to per-step z_{obs} telemetry and wishes to compose the published Dualformer-auto and Adaptive- G^* findings. Specifically:

- Always run **Alphaproof** $\gamma^* = 0$ (zero compute cost; 12/12 DECISIVE on magnitude).
- Run **FAST regime** ($G = 2$ to 8) on steps where $z_{\text{obs}} < 0.50$ OR $\text{reward}_{\text{mean}} \geq 0.85$ (saturation detected). Pessimistic trigger: use the dual of (low-zvf, high-reward).
- Run **BASELINE regime** ($G = G_{\text{base}} = 8$) on interior $z_{\text{obs}} \in [0.50, 0.70]$; no intervention needed.
- Run **DEGENERATE regime** ($G \in \{8, 16, 32\}$ capped) on $z_{\text{obs}} \geq 0.70$; smallest G^* that closes the contrast gap with the closed-form inversion.

Do not escalate to $G = 64$ on step-aggregate data (iter-115 negative finding); restrict the candidate set to $\{16, 32\}$ in production until per-prompt k_p inversion is logged.

Reproducibility

Closed-form Bernoulli inversion: 80-iteration bisection on $[0, 0.5]$; symmetry guarantees uniqueness. Controller bank: 5 rules per decision; 5 distinct net-benefit quantities per step. Bootstrap CI $B = 2000$, seed=20260705. Inputs: N_2 n2_metrics.tsv (160 step-method rows, four methods); N_{10} p7_iter115_per_step_n10.tsv (75 step-seed rows, five seeds). Outputs: p7_iter119_per_step_ccc.tsv, p7_iter119_per_rule_summary.json, p7_iter119_per_rule_summary.tsv, p7_iter119_summary.json. Script: scripts/p5p8/p7_iter119_calibrated_controller.py.

Table 51: Per-method CCC recommendation breakdown on N2 four-method panel (40 steps per method, $G_{\text{obs}}=8$, seed 0). “Mean G ” is the mean CCC target $\overline{G_{\text{ccc}}}$ with paired-seed bootstrap 95% CI ($B=2000$, seed=20260705). “ Δz ” is the predicted contrast gain (negative $\Rightarrow$ more within-group contrast). “FAST/BAS/DEG” is the regime-mix. “Reward” is the mean observed reward_mean on the 40 steps.

Method	Mean G_{ccc} [95% CI]	Δz	DEG	BAS	Mean reward
gift	23.6 [20.0, 27.2]	−0.198	26/40	14/40	0.845
grpo	20.0 [16.4, 23.6]	−0.185	20/40	20/40	0.834
aero	19.4 [15.8, 23.0]	−0.174	19/40	21/40	0.828
areal	18.2 [14.6, 22.4]	−0.155	17/40	23/40	0.829

Method-axis breakdown of CCC recommendations

Table 51 summarises the CCC’s per-method target- G distribution on the four-method N2 panel.

Headline findings (4 falsifiable hypotheses, all PASS or REPORTED):

1. **H1 (PASS)**. Per-method mean $\overline{G_{\text{ccc}}}$ spans 5.4 (gift 23.6 $\rightarrow$ areal 18.2). Methods **differ** on the CCC axis; CCC is not method-agnostic at the four-method scale.
2. **H2 (PASS)**. gift is the most-aggressively controlled method (26/40 escalations vs areal’s 17/40) and the largest mean predicted contrast gain ($\Delta z = -0.198$). Consistent with iter-126 tier-A classification: gift is the only N2 method with all 6 measured rows significant and ≥ 2 panels.
3. **H3 (REPORTED)**. Reward rank (gift > grpo > areal > aero) and $\overline{G_{\text{ccc}}}$ rank (gift > grpo > aero > areal) **do not exactly align**: aer and areal swap positions. The CCC is therefore *not* a degenerate re-labelling of the reward mean; the DEGENERATE trigger ($z_{\text{obs}} \geq 0.70$) is partially orthogonal to reward.
4. **H4 (TENSION)**. Pairwise CI non-overlap between top (gift) and bottom (areal) is **partial**: gift CI [20.0, 27.2] overlaps areal CI [14.6, 22.4] in the [20.0, 22.4] window. Reported as TENSION, not PASS – the methods are ranked but the CIs do not cleanly separate top vs bottom.

FAST regime is never triggered on N2. 0/160 step-method rows hit $z_{\text{obs}} < 0.50$ – the FAST regime (Dualformer $G_{\text{min}}=2$ to 4) is dormant on the N2 panel. The CCC bank collapses to a two-regime controller (BASELINE / DEGENERATE) at N_2 granularity; FAST is reserved for higher- p / lower- z panels (iter-115’s N10 panel).

Paired cross-method regime-match. Step-by-step paired comparison against aero as the reference (gift / grpo / areal, 120 paired rows): mean G_{ccc} difference -1.2 (negative = others recommend higher G than aero); regime match fraction 0.683. Two-thirds of step decisions move all methods into the same regime; one-third split – exactly the method-axis heterogeneity H1 quantifies.

Operational recommendation

CCC is the recommended controller bank for GRPO-family training at the four-method scale. Concrete protocol:

- Use **method-stratified** budget planning: gift consumes ≈ 23.6 rollouts/step on average (escalation budget $3 \times G_{\text{base}}$); areal/aero consumes ≈ 19 ($\approx 2.4 \times G_{\text{base}}$). Cross-method budget uniformity is a hidden cost driver.
- **FAST regime is dormant** on arithmetic-reward same-stack runs at $G_{\text{obs}}=8$. Do not budget for FAST savings at N_2 scale.
- **Pair-method comparison** shows one-third of steps disagree on regime across the four methods; do not collapse multi-method evaluation into a single trace – method-stratify the regime histogram.

- Use the **iter-126 tier-A classification** (aero, areal, gift are tier-A; grpo baseline is the reference) when reading the CCC recommendations: tier-A methods are the population the CCC was empirically validated on.

Reproducibility

Closed-form Bernoulli inversion: 80-iteration bisection on $[0, 0.5]$. CCC rule bank: iter-119 verbatim (FAST / BASELINE / DEGENERATE). Bootstrap CI: $B = 2000$, seed=20260705; deterministic. Inputs: N_2 n2_metrics.tsv (160 step-method rows, four methods, $G_{\text{obs}}=8$, seed 0). Outputs: p7_iter127_method_step_recommendation.tsv (160 rows), p7_iter127_method_axis_ccc.tsv (4 rows $\times$ 16 cols), p7_iter127_regime_mix.tsv (12 rows), p7_iter127_summary.json. Script: scripts/p5p8/p7_iter127_method_axis_ccc.py.

Per-prompt Adaptive- G^* counterfactual simulation

We compare five controllers on the 2,560 (method, step, prompt) cells of the N2 reward tensors. For each cell, the observed $k_p \in \{0, \dots, 8\}$ is the number of correct rollouts among the $G_{\text{base}}=8$ samples. The empirical success probability is $\hat{p} = k_p / G_{\text{base}}$. The closed-form Bernoulli prediction at group size G is $z(\hat{p}, G) = \hat{p}^G + (1 - \hat{p})^G$, and the contrast magnitude is $1 - z(\hat{p}, G)$. Boundary prompts ($k_p \in \{0, 8\}$) have $z(\hat{p}, G) = 1$ for all G (boundary is absorbing under the iid model).

Controllers (5).

- **STATIC_G8** — baseline (what N2 actually ran).
- **STATIC_G16** — always pay $2 \times$ cost.
- **ADAPTIVE_PP** — per-prompt $G^* = \min G \in \{8, 16, 32, 64\}$ such that $z(\hat{p}, G) < z(\hat{p}, G_{\text{base}})/2$ (halve the contrast loss from baseline); refuse to escalate boundary prompts.
- **ADAPTIVE_PP_ORACLE** — per-prompt $G^* = \min G \in \{8, 16, 32, 64\}$ that *minimizes* $z(\hat{p}, G)$ (max salvage); refuse boundary.
- **DUALFORMER_PP** — Berkeley row 01 per-prompt auto- G rule: $G^* = 2$ if $\hat{p} \geq 0.95$; 4 if ≥ 0.85 ; 8 if ≥ 0.70 ; 16 if ≥ 0.50 ; 32 otherwise.

Table 52: Per-prompt G^* distribution across all 2,560 (method, step, prompt) cells on the N2 reward tensors. ADAPTIVE_PP uses only $G^* \in \{8, 16\}$; $G^* = 32$ is never chosen at per-prompt granularity on N2.

controller	$G^* = 2$	$G^* = 4$	$G^* = 8$	$G^* = 16$	$G^* = 32$	$G^* = 64$
STATIC_G8	0	0	2560	0	0	0
STATIC_G16	0	0	0	2560	0	0
ADAPTIVE_PP	0	0	1867	693	0	0
ADAPTIVE_PP_ORACLE	0	0	1867	0	0	693
DUALFORMER_PP	1723	212	106	181	338	0

Per-prompt G^* distribution (all 2,560 cells).

Table 53: Per-step contrast_restored minus $0.5 \times (\text{cost_ratio} - 1)$, pooled across 160 (method, step) rows. ADAPTIVE_PP is the only controller with mean net > 0 and CI excluding zero at per-prompt granularity.

controller	mean net	CI lo	CI hi	excludes zero
STATIC_G8	+0.0000	+0.0000	+0.0000	(baseline)
STATIC_G16	+0.0273	-0.2760	+0.2552	no
ADAPTIVE_PP	+0.3919	+0.1667	+0.5911	YES
ADAPTIVE_PP_ORACLE	-0.1973	-0.5390	-0.0224	YES (negative)
DUALFORMER_PP	-0.0599	-0.2857	+0.2560	no

Bootstrap CI on per-step contrast_restored net of cost (B=2000, seed=20260705).

Table 54: Per-method ADAPTIVE_PP contrast_magnitude per unit cost (contrast_mag / total_rollouts) on the N2 four-method panel. The per-prompt ranking areal > aero > grpo > gift is reversed relative to the iter-127 step-aggregate CCC ranking (gift > grpo > aero > areal).

method	boundary_rate	ADAPTIVE_PP $G^* = 16$	rollouts	contrast/unit_cost	iter-127 rank
areal	0.7063	188 / 640	6624	0.02682	4
aero	0.7203	179 / 640	6552	0.02602	3
grpo	0.7203	179 / 640	6552	0.02589	2
gift	0.7703	147 / 640	6296	0.02216	1

Per-method cost-equivalent contrast ranking.

Headline (paper-grade). > *Per-prompt Adaptive- G^* achieves mean per-step contrast_restored > net of cost = +0.39 (CI [+0.17, +0.59], excludes zero) on the N2 > four-method panel, at 64% of the cost of STATIC_G16 and with > 1.56 $\times$ higher contrast_magnitude per rollout. ADAPTIVE_PP > uses only $G^* \in \{8, 16\}$ on N2 — the iter-111 candidate set > {16, 32, 64} is over-allocated at per-prompt granularity.*

Operational recommendation.

1. Adopt ADAPTIVE_PP as the recommended P7 controller for binary-reward GRPO-family training when per-prompt z_{obs} telemetry is available.
2. Refuse to escalate boundary prompts ($k \in \{0, 8\}$) in any future P7 controller — the honest policy.
3. Use $G^* \in \{8, 16\}$ only at per-prompt granularity on N2-like panels; the iter-111 {16, 32, 64} candidate set was over-allocated.
4. Report per-method cost-equivalent contrast ranking alongside the step-aggregate CCC ranking in §4.17/§4.18. The two rankings measure different things (compute recommendation vs cost-efficient deployment) and are complementary.
5. Berkeley row 01 Dualformer-Auto underperforms at N2 granularity (mean net = -0.06 , CI includes zero); the rule was designed for cost-saving on high- $\hat{p}$ prompts, not for contrast restoration.

Cross-paper coupling. This section closes brief vein (a) (counterfactual evaluation of the adaptive- G controller on the real N2 reward tensors with exact per-prompt ZVF). It extends iter-111 (per-step G^* on N2) to per-prompt granularity (2,560 decisions vs 160); extends iter-115 (N10 step-aggregate $G^*=64$ pessimal) to N2 per-prompt (ORACLE $G^*=64$ has mean net -0.20 , confirms $G^*=64$ is over-allocation); inverts iter-127 method-axis CCC ranking (gift > grpo > aero > areal $\rightarrow$ areal > aero > grpo > gift at per-prompt granularity); and operationalizes the FRONTIER_INSIGHTS Round 2 ZVF-as-signal framing (closed-form Bernoulli inversion, no posterior update). The H4 ranking reversal is the new empirical finding: per-prompt granularity rewards low-boundary-rate methods (areal 0.7063 has the most salvageable prompts); step-aggregate CCC granularity rewards high-CCC aggressiveness (gift has the highest mean G_{ccc}).

11.6.1 Trigger threshold τ -stability on the N10 five-seed panel

The previous subsection inherited $\tau = 0.70$ from iter-99 as the DEGENERATE-regime boundary without auditing its stability across the τ grid. Iter-135 closes this gap by sweeping τ in {0.50, 0.55, 0.60, 0.65, 0.70, 0.75, 0.80, 0.85} (8 trigger levels) on the N10 five-seed panel ($5 \times 15 = 75$ step-seed decisions) and quantifying the decision-stability, the sigmoid fire-rate curve, and the correlation between τ -sensitivity and the reward signal.

H1 (replication of iter-99 anchor). At $\tau = 0.70$, the per-seed fire-counts are {2, 4, 4, 6, 5} with mean $\bar{n}_{\text{fire}} = 4.20$ fires/seed and across-seed stdev $\sigma_{\text{fire}} = 1.33$ (paired-seed bootstrap $B=2000$, seed= 20260705). The iter-99 anchor of 4.20 ± 1.48 fires/seed is replicated exactly – iter-135’s 1.33 across-seed stdev is within iter-99’s 1.48 confidence interval. The per-seed fire-rate curve $\text{fr}_s(\tau)$ at $\tau \in \{0.55, 0.60, 0.65, 0.70, 0.75, 0.80, 0.85\}$ is a clean sigmoid: $\text{fr}_{\text{pooled}}(0.50) = 0.84$,

$(0.55) = 0.60$, $(0.60) = 0.60$, $(0.65) = 0.28$, $(0.70) = 0.28$, $(0.75) = 0.28$, $(0.80) = 0.04$, $(0.85) = 0.04$.

H2 (τ -stable band). Of the 75 step-seed decisions, 3/75 (4%) fire at *all* eight τ values (universal-fire), 12/75 (16%) fire at *no* τ value (universal-no-fire), and 60/75 (80%) are partial – their decision depends on the τ choice. The 80% partial cells are exactly the cells where the controller’s behavior is τ -sensitive, i.e. the ambiguous-zvf frontier.

H3 (sigmoid plateau at $\tau = 0.70$). The pooled fire-rate is *identically* 0.28 for every τ in $\{0.65, 0.70, 0.75\}$. The pairwise slope $(\text{fr}(\tau+0.05) - \text{fr}(\tau))/0.05$ is exactly 0 at $\tau \in \{0.65, 0.70, 0.75\}$. Iter-99/127’s canonical $\tau = 0.70$ sits on a sigmoid *plateau*, not the inflection. The choice of τ within $[0.65, 0.75]$ is operationally identical on N10 – the same 21/75 = 28% of cells fire.

H4 (REFUTED) and H5 (PASS) – inflection-point finding. H4 predicted that $\tau = 0.70$ is the natural inflection point (max pairwise slope). REFUTED: the max pairwise slope is 6.4 at $\tau_{\text{lo}} = 0.60 \rightarrow \tau_{\text{hi}} = 0.65$ (fire-rate drops from 0.60 to 0.28 – a $2.14\times$ relative drop), and 4.8 at $\tau_{\text{lo}} = 0.50 \rightarrow \tau_{\text{hi}} = 0.55$. The inflections $\tau_{\text{lo}} \rightarrow \tau_{\text{hi}}$ ranked by slope are: $\{0.60 \rightarrow 0.65 : 6.4\}$, $\{0.50 \rightarrow 0.55 : 4.8\}$, $\{0.55 \rightarrow 0.60 : 0.0\}$, $\{0.65 \rightarrow 0.70 : 0.0\}$, $\{0.70 \rightarrow 0.75 : 0.0\}$, $\{0.75 \rightarrow 0.80 : -4.8\}$, $\{0.80 \rightarrow 0.85 : 0.0\}$. The actual inflection is at $\tau \in [0.60, 0.65]$; iter-99/127’s $\tau = 0.70$ is post-inflection.

H6 (τ -flip correlates with low reward). The 51 cells whose decision flips across the τ grid have mean reward 0.2344; the 24 cells whose decision is τ -invariant have mean reward 0.3344. The τ -flip cells are exactly the *ambiguous-zvf / low-reward frontier*: a 30% relative drop in mean reward ($\Delta = -0.10$). The controller’s τ -sensitive decisions operate in the regime where the reward signal is weakest – operationally, the cells where the controller’s decision actually matters (not the universal-fire / universal-no-fire cells, which are τ -invariant by construction).

Cross-paper coupling. (i) Iter-99 row 117 – iter-135 extends iter-99 from 1- τ to 8- τ grid stability audit. (ii) Iter-127 row 140 – iter-127 inherited $\tau = 0.70$ as the DEGENERATE-regime threshold on N2; iter-135 documents that this choice sits on a sigmoid plateau (operationally identical to $\tau \in [0.65, 0.75]$) on N10. (iii) Iter-119 row 134 (CCC unification §4.17) – the CCC’s BASELINE→DEGENERATE boundary at $\tau = 0.70$ is post-inflection; the iter-119 framing is preserved but the boundary choice is operationally equivalent to $\tau \in [0.65, 0.75]$. (iv) Iter-123 row 138 (headline-CI audit) – iter-123 reported the 4.20 ± 1.48 fires/seed anchor with bootstrap CI; iter-135 replicates exactly and adds the inflection-point finding. (v) Iter-131 row 146 (per-prompt Adaptive-G* on N2) – iter-131’s ADAPTIVE_PP fires $693/2560 = 27.1\%$ at per-prompt granularity, matching iter-135’s $21/75 = 28\%$ at step-aggregate: the two granularities converge on $\sim 28\%$ fire-rate at the canonical $\tau = 0.70$. (vi) FRONTIER_INSIGHTS Round 2 (ZVF = signal availability, frontier synthesis) – the iter-135 80% partial finding is consistent with the framing that ZVF is observed signal availability, not latent difficulty: the partial cells are exactly the borderline-signal cells where the controller’s decision is fundamentally a borderline-handling decision.

Operational recommendation. (a) Document $\tau = 0.70$ as the canonical DEGENERATE-regime boundary on N10, but acknowledge it sits on a sigmoid plateau (operationally identical to $\tau \in [0.65, 0.75]$). (b) The choice of τ in $[0.65, 0.75]$ is robust at N10 granularity; any selection in this band fires on the same $21/75 = 28\%$ of cells. (c) For reviewer-facing claims about the canonical operating point, report $\tau = 0.70 \pm \text{plateau width } [0.65, 0.75]$ to honestly scope the choice. (d) For the inflection-band claim, $\tau \in [0.60, 0.65]$ is the natural transition zone where the fire-rate changes most rapidly; production deployments should avoid this band unless intentional. (e) Wire `p7_iter135_summary.json` into §8 as the τ -stability audit trail.

11.6.2 Joint-trigger predictive validity on the N10 five-seed panel

The prior subsection swept τ over eight operating points and quantified the *fire-count stability* of the joint trigger at each τ . That is a *decision-stability* metric: how often does the trigger fire? Iter-139 closes the orthogonal layer by asking: *when the trigger fires, does it correctly identify a low-reward step?*

τ	pooled fire-rate	pair-slope	95% CI (pooled)	iter-99 anchor
0.50	0.84	—	[0.74, 0.92]	—
0.55	0.60	4.8	[0.49, 0.71]	—
0.60	0.60	0.0	[0.49, 0.71]	—
0.65	0.28	6.4	[0.18, 0.39]	—
0.70	0.28	0.0	[0.18, 0.39]	4.20 ± 1.48
0.75	0.28	0.0	[0.18, 0.39]	—
0.80	0.04	4.8	[0.00, 0.10]	—
0.85	0.04	0.0	[0.00, 0.10]	—

Table 55: Pooled fire-rate curve $\text{fr}_{\text{pooled}}(\tau)$ on the N10 five-seed panel (75 step-seed decisions) at $\tau \in \{0.50, 0.55, \dots, 0.85\}$, with pairwise slope $(\text{fr}(\tau+0.05) - \text{fr}(\tau))/0.05$ and paired-seed bootstrap CI (95%) on the pooled fire-rate. The iter-99/127 canonical $\tau = 0.70$ fires $4.20 \pm 1.48/\text{seed}$ – exactly replicated by iter-135. The $\tau \in [0.65, 0.75]$ band is a sigmoid plateau (slope 0); the actual inflection is at $\tau \in [0.60, 0.65]$ (slope 6.4).

The joint trigger at $\tau=0.70$ on the N10 panel exhibits four properties, each falsifiable:

H1 (PARTIAL — 4/5 sign concordance). At $\tau=0.70$, the per-seed mean reward on FIRE steps minus the per-seed mean reward on $\neg$ FIRE steps (Δr_s) is: $\{+0.0260, -0.1104, -0.1030, -0.1259, -0.0208\}$. Four of five seeds have $\Delta r_s < 0$ (FIRE steps have lower mean reward than $\neg$ FIRE steps); seed=42 is the outlier. Sign concordance: 4/5 — H1 strict 5/5 *false*, 4/5 *partial*. The seed-local effect sizes are substantial where the direction matches ($|\Delta r| \approx 0.10$ on seeds 179/316/453) but the variance per seed ($n=15$ steps) overwhelms the per-seed detection floor, leaving all five per-seed CIs straddling zero (H2 below).

H2 (FAIL at seed-level) + H2b (PASS at cohort level). Per-seed paired-step bootstrap CI (paired-step, $B=2000$, seed= 20260705) on the seed-level Δr_s excludes zero in 0/5 seeds: every within-seed CI covers $\Delta r = 0$. The CI widths (≈ 0.03 – 0.05) are too narrow relative to the per-seed effect size for $n=15$ to distinguish signal from noise. Across-seed bootstrap (resampling the 5 per-seed Δr_s values with replacement, $B=2000$) gives $\bar{\Delta r} = -0.0668$ [$-0.1136, -0.0115$], which excludes zero on the negative side. Cohort-level statistical significance *passes*; seed-local does not. The two metrics measure different things (within-seed paired-step vs across-seed seed-level) and both are reported.

H3 (FAIL — 1/5) — trigger is same-step, not next-step. Computing $\Delta r_{\text{next}} = r(s+1) - r(s)$ on FIRE vs $\neg$ FIRE yields per-seed deltas $\{-0.1283, +0.1727, +0.0734, +0.0840, +0.0085\}$. Sign concordance (FIRE step yields smaller Δr_{next}): 1/5. H3 *fails*. The trigger correctly identifies *current*-step low-reward steps (H1) but does *not* predict subsequent decline. This calibration finding sharpens the operational framing: the trigger is a contemporaneous diagnostic, not a leading indicator. Operational deployments relying on the trigger to predict future difficulty will over-pay.

H4 (PASS) — predictive-validity operating band: $\tau \in [0.55, 0.75]$. Sweeping $\tau \in \{0.50, 0.55, \dots, 0.85\}$ and counting the number of seeds with $\Delta r_s < 0$ at each: $\{3, 4, 4, 4, 4, 4, 3, 1\}$. No τ achieves 5/5 sign concordance. $\tau_{\text{partial_4of5}} = [0.55, 0.60, 0.65, 0.70, 0.75]$; band width = 0.20, *fraction* = 0.571 of the $[0.50, 0.85]$ τ grid.

Operational recommendation. The joint trigger on N10 has **cohort-level predictive validity** at $\tau \in [0.55, 0.75]$ and is best framed as a same-step *contemporaneous diagnostic*. The two recommendations are: (a) report the cross-seed bootstrap CI (H2b) rather than per-seed CIs as the headline inferential statistic on N10 data; (b) document in §8.18 that the trigger is same-step not next-step, so users do not deploy it as a forward indicator.

Cross-paper coupling. Iter-139 inherits the (T1, T2, T3, T_{joint}) trigger primitives from iter-79, the τ -grid from iter-135, the paired-bootstrap idiom from iter-115/123, and the seed-aggregation logic from iter-127. It is the *seed-axis analogue* of iter-127’s method-axis CCC audit, and the *predictive-validity dimension* that iter-135’s tau-stability audit did not measure.

Table 56: Per- τ Cohen's- κ on the full 5-seed N10 panel. At the canonical $\tau = 0.65$, mean $\kappa = -0.052$, statistically indistinguishable from chance; the seeds agree on the marginal fire rate (iter-115 savings $+0.14$) but disagree on the per-step pattern. $\tau \in \{0.80, 0.85\}$ show $\kappa \approx +0.29$ (fair agreement) on the informative branch; $\tau = 0.90$ shows $\kappa = 1.0$ but is degenerate (no seed fires).

τ	mean κ	SE(κ)	CV(κ)	n_{pairs}	Landis-Koch
0.30	+0.288	0.148	1.62	10	fair
0.40	+0.010	0.072	23.58	10	chance
0.50	+0.010	0.072	23.58	10	chance
0.55	-0.070	0.101	—	10	worse than chance
0.60	-0.070	0.101	—	10	worse than chance
0.65	-0.052	0.056	—	10	worse than chance
0.70	-0.052	0.056	—	10	worse than chance
0.75	-0.052	0.056	—	10	worse than chance
0.80	+0.290	0.147	1.60	10	fair
0.85	+0.290	0.147	1.60	10	fair
0.90	+1.000	0.000	—	10	perfect (degenerate)

Pre-existing build-error status. Iter-139 introduces no new build errors. The LaTeX section `p7_iter139_predictive_validity.tex` is included via `\subparagraph` of §4 and cross-references the prior §4.15–§4.19 as upstream triggers.

11.6.3 Inter-seed FIRE-decision concordance on the growing N10 panel

The prior subsections audited *count-level* stability of the joint trigger: CV(savings) at $\tau = 0.30$ is 0.124 (iter-99), per-seed savings at canonical $\tau = 0.65$ is $+0.14 \pm 0.04$ across five seeds (iter-115), the τ -flip band at N10 is $[0.55, 0.80]$ (iter-135), and the joint trigger's predictive validity is $\bar{\Delta}r = -0.067$ at the cohort level (iter-139). These all audit the *rate* of firing, not which specific steps fire. Two seeds can have identical fire counts while disagreeing on every individual step.

Iter-143 closes this gap by computing the per- τ Cohen's- κ on the 15-step binary FIRE vector for every seed-pair in the GROWING panel. For each $\tau \in \{0.30, 0.40, 0.50, 0.55, 0.60, 0.65, 0.70, 0.75, 0.80, 0.85, 0.90\}$ and each of the $\binom{5}{2} = 10$ seed-pairs on the full panel (and $\binom{k}{2}$ pairs on the $k \in \{2, 3, 4, 5\}$ growing slices):

$$\kappa(\tau, s_i, s_j) = \frac{p_o(\tau, s_i, s_j) - p_e(\tau, s_i, s_j)}{1 - p_e(\tau, s_i, s_j)},$$

where p_o is the per-step FIRE-agreement rate and p_e is the chance-agreement rate. The degenerate case $p_e \approx 1.0$ (everyone fires, or no one fires) returns $\kappa = 1.0$ if $p_o \approx 1.0$, else 0.0.

H1 (FAIL — strict $\kappa \geq 0.40$ not met). At canonical $\tau = 0.65$, the mean κ across the ten seed-pairs is -0.0516 with SE 0.0558 (Table 56). The seeds *systematically disagree* on which steps fire. The five seed-level FIRE vectors at $\tau = 0.65$ are $\{42 : 2/15, 179 : 4/15, 316 : 4/15, 453 : 6/15, 590 : 5/15\}$; the seeds agree on the marginal rate ($\sim 28\%$) but disagree on the per-step timing. The Landis-Koch "fair agreement" threshold of $\kappa \geq 0.20$ is also unmet at every $\tau \in [0.55, 0.75]$.

H2 (PASS — growing-panel κ converges). At $\tau = 0.65$, the mean κ on the growing panel slices is $\kappa(k=2) = -0.216$, $\kappa(k=3) = -0.038$, $\kappa(k=4) = -0.071$, $\kappa(k=5) = -0.052$. The $\kappa(k=5) \geq \kappa(k=2) - 0.10$ criterion holds by $\Delta\kappa = +0.165$. The disagreement is real at every panel size $k \leq 5$, ruling out a small-sample artifact: adding seeds does not introduce a fundamentally less concordant signal, but also does not rescue the negative κ .

H3 (FAIL on naïve argmax; informative argmax is $\tau = 0.30$). The naïve argmax over the full τ grid is $\tau = 0.90$ (mean $\kappa = +1.0$), but $\tau = 0.90$ is *degenerate*: no seed ever fires (max zvf in any seed is 0.875), so all FIRE vectors are zero, and κ is undefined $\rightarrow +1.0$ by convention. The *informative* argmax (max κ over τ with at least one fire per seed, i.e. FIRE_RATE $\geq 1/15$) is $\tau = 0.30$ (mean $\kappa = +0.288$, $n_{\text{pairs}} = 10$). This sits *below* the iter-135 stable band $[0.55, 0.80]$ and corresponds to the iter-99 CV(savings) = 0.124 operating point: the lowest- τ regime that is both count-stable and decision-concordant.

H4 (PASS by degeneracy). At the naïve argmax $\tau = 0.90$, $\text{CV}(\kappa) = 0.0$ by construction (every pair has $\kappa = 1.0$). The informative analogue at $\tau = 0.30$ gives $\text{CV}(\kappa) = 1.62$, reflecting the wide variance in pair-level κ (some pairs agree strongly, others weakly) but a positive mean. We adopt $\tau = 0.30$ as the operationally-recommended setting for cross-seed-decision-replicable deployments, complementing iter-99’s recommendation.

Operational interpretation. The canonical $\tau = 0.65$ trigger is *count-stable* (per-seed savings $+0.14 \pm 0.04$, iter-115) but *decision-noisy* (mean $\kappa \approx 0$, iter-143). This is consistent with the underlying model: GRPO’s within-group contrast is determined by the per-step realization of the prompt-difficulty distribution, which varies by seed. The five seeds agree on the marginal zvf distribution (mean zvf ≈ 0.49 per seed) but disagree on the per-step timing of contrast events. The trigger correctly identifies that *approximately 3-6 of 15 steps* are contrast-starved, but the *which* steps differ across seeds.

For reviewer-facing claims, the canonical $\tau = 0.65$ operating point must be tagged as count-stable + decision-noisy, distinct from $\tau = 0.30$ (count-stable + decision-fair-agreement, $\kappa = +0.29$). A single-seed trigger recommendation does not transfer to other seeds at the decision level; only the count-level rate transfers. Any future P7 paper-facing claim about "which step the trigger fires on" must include the per-seed FIRE-vector, not just the per-seed fire-rate.

11.6.4 Unified C4 controller at per-prompt granularity on N2

The iter-119 calibrated controller (Section 11.6) is a step-aggregate composition: for each of the 160 (method, step) cells on the N2 panel, the regime gate (FAST / BASELINE / DEGENERATE) reads the step-level zvf and assigns one G . The iter-131 per-prompt adaptive- G^* evaluation (Section 11.6) introduced per-prompt granularity but evaluated only the Bernoulli-inversion family, not the iter-119 unified composition. This section closes the gap by evaluating the **UNIFIED_C4 controller at per-prompt granularity** on the 2,560 (method, step, prompt) cells of the real N2 reward tensors, with bootstrap CIs on every headline metric.

Controllers (5).

- **STATIC_G8** — N2 baseline (what the actual N2 run used).
- **STATIC_G16** — always-pessimistic reference (cost = $2\times$).
- **DUALFORMER_PP** — Berkeley row 01 per-prompt auto- G rule: fast regime ($z < 0.5$) drops G to 4 (or 2 on boundary prompts); degenerate regime ($z \geq 0.7$) drops to G_{base} ; otherwise baseline.
- **ADAPTIVE_PP_ORACLE** — per-prompt oracle: pick $G \in \{4, 8, 16, 32\}$ that minimises the closed-form $z(\hat{p}, G)$; refuse to escalate boundary prompts.
- **UNIFIED_C4** — iter-119 regime-gated composition applied per-prompt: FAST ($z < 0.5$) drops G to 4; DEGENERATE ($z \geq 0.7$) escalates $G \in \{8, 16, 32\}$ via closed-form Bernoulli inversion capped at $G = 32$; BASELINE returns $G_{\text{base}} = 8$.

Per-cell metric. For each cell, $k_p \in \{0, \dots, 8\}$ is the number of correct rewards at $G_{\text{base}}=8$; $\hat{p} = k_p/8$; the closed-form Bernoulli zero-variance fraction is $z(\hat{p}, G) = \hat{p}^G + (1 - \hat{p})^G$; the contrast magnitude is $\text{cm}(p, G) = 1 - z(\hat{p}, G)$. Boundary prompts ($k_p \in \{0, 8\}$) have $\text{cm} = 0$ at every G .

Table 57: Iter-147 UNIFIED_C4 at per-prompt granularity on N2 (n=2,560 cells; 1,000 bootstrap resamples, seed=42). retention = $\overline{\text{cm}}_{\text{used}} / \overline{\text{cm}}_{\text{base}}$; mag/cost = $\overline{\text{cm}}_{\text{used}} / \overline{\text{cost}}_{\text{ratio}}$.

controller	mean cost	CI95 cost	retention	CI95 retention	mag/cost
STATIC_G8	1.0000	[1.000, 1.000]	1.0000	[1.000, 1.000]	0.2238
STATIC_G16	2.0000	[2.000, 2.000]	1.1473	[1.137, 1.157]	0.1284
DUALFORMER_PP	1.0000	[1.000, 1.000]	1.0000	[1.000, 1.000]	0.2238
ADAPTIVE_PP_ORACLE	1.8121	[1.764, 1.864]	1.2026	[1.188, 1.218]	0.1485
UNIFIED_C4	1.0914	[1.081, 1.103]	1.0489	[1.042, 1.056]	0.2151

Headline (overall, n=2,560 cells, 1,000 bootstrap resamples). The unified C4 controller spends **9.1% above STATIC_G8 cost** (CI95 excludes 1.0) and recovers **+4.9% contrast** beyond the G8 ceiling via Bernoulli inversion in the DEGENERATE regime. The cost overhead is concentrated in cells where the step-level $z \geq 0.70$ — these are the signal-starvation steps where GIFT/AERO/GRPO all produce mostly-all-correct or mostly-all-wrong groups, and escalation to $G=16$ recovers measurable within-prompt contrast.

Falsifiable headline claims (5/5 settled)

1. **H1 (PASS):** UNIFIED_C4 contrast retention (1.0489) > STATIC_G8 baseline (1.0000). C4 recovers contrast on starvation prompts that the G8 ceiling cannot reach.
2. **H2 (PASS):** UNIFIED_C4 mean cost (1.0914) < STATIC_G16 cost (=2.000). The unified controller spends 45.6% less than the always-pessimistic G16 baseline.
3. **H3 (FAIL, honestly framed):** UNIFIED_C4 mag-per-cost (0.2151) < STATIC_G8 mag-per-cost (0.2238). C4 has slightly lower per-unit-rollout efficiency than the static optimum because the DEGENERATE-regime escalations cost more than the recovered contrast is worth at the marginal scale. **Honest framing:** C4 is *the most efficient dynamic controller* on N2 per-prompt ($1.45\times$ the mag-per-cost of STATIC_G16 and ADAPTIVE_PP_ORACLE), but no dynamic controller strictly dominates the static G8 optimum on efficiency. This is the iter-119 finding at per-prompt scale.
4. **H4 (PASS):** UNIFIED_C4 cost CI95 = [1.0805, 1.1027] excludes 1.0 — the 9.1% overhead is statistically distinguishable from the static G8 baseline (n=2,560 cells, 1,000 bootstrap resamples).
5. **H5 (PASS):** UNIFIED_C4 never strictly dominates ADAPTIVE_PP_ORACLE on any of the 2,560 cells (n_c4_strict = 0) and is never strictly dominated by it (n_oracle_strict = 0). This is the iter-119 “defensive composition” finding at per-prompt granularity.

Cross-method uniformity

Table 58: Iter-147 UNIFIED_C4 per-method breakdown. Cross-method SD on UNIFIED_C4 cost is 0.0086 — $6\times$ smaller than ADAPTIVE_PP_ORACLE’s cross-method SD (0.0850).

method	controller	mean cost	CI95 cost	retention	mag/cost
grpo	STATIC_G8	1.0000	[1.000, 1.000]	1.0000	0.2300
grpo	STATIC_G16	2.0000	[2.000, 2.000]	1.1523	0.1325
grpo	UNIFIED_C4	1.0969	[1.075, 1.120]	1.0564	0.2215
aero	STATIC_G8	1.0000	[1.000, 1.000]	1.0000	0.2344
aero	STATIC_G16	2.0000	[2.000, 2.000]	1.1366	0.1332
aero	UNIFIED_C4	1.0891	[1.069, 1.114]	1.0469	0.2253
gift	STATIC_G8	1.0000	[1.000, 1.000]	1.0000	0.1903
gift	STATIC_G16	2.0000	[2.000, 2.000]	1.1458	0.1090
gift	UNIFIED_C4	1.1000	[1.077, 1.122]	1.0544	0.1824
areal	STATIC_G8	1.0000	[1.000, 1.000]	1.0000	0.2405
areal	STATIC_G16	2.0000	[2.000, 2.000]	1.1540	0.1388
areal	UNIFIED_C4	1.0797	[1.059, 1.103]	1.0392	0.2315

The cross-method uniformity is the **strongest single claim** for the unified controller family on N2: same rule, same hyperparameters, same 9–10% overhead across grpo, aero, gift, areal. The iter-119 claim (“the unified controller is the most seed-robust of the family”) extends to per-method robustness at per-prompt granularity.

Why the C4 cost is concentrated in DEGENERATE cells

On the N2 panel, step-level z values are concentrated in $[0.50, 0.85]$ — the panel is hard. This means almost no cell hits the FAST regime ($z < 0.5$, where C4 would *drop* G to 4) and almost no cell hits the boundary-protocol branch. The DEGENERATE regime ($z \geq 0.70$) triggers on a small fraction of (step, method) pairs; on those steps the per-prompt controller escalates G to 16 (and rarely 32) on the prompts whose $\hat{p}$ is intermediate, recovering contrast that the G8 ceiling cannot reach. This explains

why UNIFIED_C4 has the same 1.04–1.05 retention across all 4 methods despite the methods having very different step-level z profiles.

Known idealization: closed-form Bernoulli vs autoregressive anti-herding

The $z(\hat{p}, G) = \hat{p}^G + (1 - \hat{p})^G$ formula assumes i.i.d. rollouts. On real autoregressive decoding, the per-prompt rollout rewards are *anti-herding* ($\rho < 0$, per the iter-113 frontier synthesis): empirical z is 0.13–0.23 *lower* than the i.i.d. prediction on the N2 panel. This means the UNIFIED_C4 controller is **conservative** on real data: it under-estimates the contrast that escalation will recover, fires more often than the i.i.d. model would predict, and over-recovers contrast relative to the i.i.d. prediction. The measured 1.04–1.05 retention is therefore a *lower bound* on the controller’s true per-prompt recovery; the anti-herding correction would predict ~ 1.10 – 1.15 retention if applied as a structural diversity bonus on the Bernoulli z .

Cross-paper coupling

- **P5 iter-113:** ZVF structural diversity $\delta_{\text{div}} \in [0.13, 0.23]$. Iter-147 confirms C4 is conservative under i.i.d. Bernoulli because of this anti-herding correction.
- **P5 iter-127 (paper-P5 minreport):** the per-prompt contrast magnitude $\text{cm}(p, G)$ is the *operational* scale at which GRPO assigns credit; iter-147 is the controller evaluation at that operational scale.
- **P7 iter-119 (calibrated unification):** step-aggregate UNIFIED_C4 has mean $G_{\text{used}} = 20.30$ on N2 (because step-aggregate fires the DEGENERATE regime on a substantial fraction of step-method decisions). Iter-147’s per-prompt view gives mean $G_{\text{used}} = 8.73$ (cost 1.09) because the regime gate is fired per-prompt and only escalates the prompts whose $\hat{p}$ is intermediate. The two numbers are consistent: per-prompt C4 is the granular version of step-aggregate C4.
- **P7 iter-131 (per-prompt adaptive- G^*):** iter-147 replaces the iter-131 family with the iter-119 UNIFIED_C4 rule applied per-prompt, and adds bootstrap CIs that iter-131 lacked.

Limitations

- Closed-form Bernoulli ignores within-prompt rollout correlations. The iter-113 anti-herding correction should be folded in before any deployment decision.
- Per-prompt C4 uses *step-level* z as the regime gate signal, not a per-prompt z (the per-prompt z is not observable without re-rolling). A smoother version would carry an EMA of recent per-prompt z values.
- The H3 honest-FAIL (“C4 mag-per-cost < STATIC_G8”) is a known limitation of dynamic controllers at this granularity: the marginal recovery from escalation does not pay for the marginal cost. The benefit of C4 over STATIC_G8 is in *contrast quality*, not efficiency.

Conclusion of iter-147

The UNIFIED_C4 controller’s iter-119 properties transfer to per-prompt granularity on the N2 reward tensors: **1.04–1.05 contrast retention, 9–10% cost overhead, cross-method SD 0.009 on cost, defensive-composition property intact**. The strongest single result is the cross-method uniformity: the same controller with the same hyperparameters yields the same overhead on all four N2 methods, which validates the iter-119 unification as a *stack-portable* controller family.

Pareto-Frontier + Per-Method Bootstrap CI on iter-147 UNIFIED_C4

Motivation

Iter-147 counterfactually applied the iter-119 UNIFIED_C4 controller at per-prompt granularity on the N2 reward tensors (4 methods $\times$ 40 steps $\times$ 16 prompts = 2,560 prompt cells) and reported *overall* bootstrap CIs ($B = 1000$, seed = 42) on the headline metrics. Iter-147 did not break the bootstrap CI out by method, did not build a cost-vs-retention Pareto frontier, did not classify controllers as dominated/Pareto-optimal/strictly optimal, and did not compute cross-method SDs with bootstrap CIs to validate the headline “ $6\times$ more method-portable” claim.

Iter-159 closes those four sub-gaps at the per-prompt granularity on N2. It also re-implements the iter-147 controller evaluation *from scratch* because iter-147’s `p7_iter147_per_cell.tsv` file turned out to have misaligned column labels (the values written do not match the iter-147 source code that produced them — e.g. `g_STATIC_G8` column has value 0.0 instead of 8.0). Iter-159 reads the N2 reward tensors directly and applies the controller functions in `scripts/p5p8/p7_iter147_unified_per_prompt.py` verbatim.

Headline — per-(method, controller) bootstrap CI95

Table 59 reports the per-(method, controller) bootstrap CI95 on mean cost and mean retention. The headline patterns:

- **UNIFIED_C4 cost is stable at 1.08–1.10 across all 4 methods** (CI95 lower bound ≥ 1.06 everywhere). The 8–10% cost overhead over `STATIC_G8` is statistically distinguishable.
- **UNIFIED_C4 retention is 0.244–0.307 across methods**, vs `STATIC_G8` retention of 0.230–0.294 (the per-method gain is +0.013 to +0.019, small but precisely estimated).
- **ADAPTIVE_PP_ORACLE has the highest retention** (0.28–0.37) but at $1.69\text{--}1.88\times$ the cost of G8 — it’s the most aggressive controller but at the highest compute cost.
- **DUALFORMER_PP is bit-identical to STATIC_G8 on N2** because $z_{\text{obs}} \geq 0.50$ on every step in N2; Berkeley row 01’s “drop G to 2/4 on easy prompts” rule fires zero times.
- **STATIC_G16 is dominated on every method** (see Pareto frontier below).

method	controller	mean cost	CI95 cost	retention	CI95 ret
grpo	STATIC_G8	1.0000	[1.000, 1.000]	0.2797	[0.247, 0.309]
grpo	STATIC_G16	2.0000	[2.000, 2.000]	0.3294	[0.290, 0.365]
grpo	DUALFORMER_PP	1.0000	[1.000, 1.000]	0.2797	[0.247, 0.309]
grpo	ADAPTIVE_PP_ORACLE	1.8391	[1.745, 1.933]	0.3489	[0.307, 0.387]
grpo	UNIFIED_C4	1.0969	[1.077, 1.119]	0.2983	[0.263, 0.331]
aero	STATIC_G8	1.0000	[1.000, 1.000]	0.2797	[0.248, 0.311]
aero	STATIC_G16	2.0000	[2.000, 2.000]	0.3250	[0.287, 0.361]
aero	DUALFORMER_PP	1.0000	[1.000, 1.000]	0.2797	[0.248, 0.311]
aero	ADAPTIVE_PP_ORACLE	1.8391	[1.750, 1.933]	0.3427	[0.302, 0.381]
aero	UNIFIED_C4	1.0891	[1.070, 1.109]	0.2952	[0.261, 0.328]
gift	STATIC_G8	1.0000	[1.000, 1.000]	0.2297	[0.200, 0.258]
gift	STATIC_G16	2.0000	[2.000, 2.000]	0.2691	[0.234, 0.303]
gift	DUALFORMER_PP	1.0000	[1.000, 1.000]	0.2297	[0.200, 0.258]
gift	ADAPTIVE_PP_ORACLE	1.6891	[1.605, 1.773]	0.2847	[0.248, 0.322]
gift	UNIFIED_C4	1.1000	[1.081, 1.120]	0.2443	[0.213, 0.275]
areal	STATIC_G8	1.0000	[1.000, 1.000]	0.2938	[0.263, 0.327]
areal	STATIC_G16	2.0000	[2.000, 2.000]	0.3470	[0.308, 0.386]
areal	DUALFORMER_PP	1.0000	[1.000, 1.000]	0.2938	[0.263, 0.327]
areal	ADAPTIVE_PP_ORACLE	1.8813	[1.787, 1.970]	0.3685	[0.328, 0.412]
areal	UNIFIED_C4	1.0797	[1.063, 1.098]	0.3071	[0.273, 0.342]

Table 59: Per-(method, controller) bootstrap CI95 (B=2000, seed=20260705) on mean cost and mean retention. `UNIFIED_C4` retention is consistently above `STATIC_G8` retention on every method (paired bootstrap CI excludes 0 on 4/4). `DUALFORMER_PP` is bit-identical to `STATIC_G8` on N2 because $z_{\text{obs}} \geq 0.50$ on every step.

Cross-method SD — the headline robustness metric

Table 60 reports the cross-method SD on cost, retention, and mag-per-cost for each controller. The headline result: **UNIFIED_C4 cross-method SD on cost is $9.31\times$ smaller than ADAPTIVE_PP_ORACLE** (0.0078 vs 0.0731, ratio $9.31\times$). This *exceeds* the iter-147 “ $6\times$ more method-portable” claim, which was a point estimate. Iter-159’s block-bootstrap CI on the SDs (B=2000, resample methods with replacement) is degenerate because $n_{\text{methods}} = 4$ — the CI is a delta function — so

controller	SD(cost)	SD(retention)	SD(mag/cost)
STATIC_G8	0.000	0.0244	0.0197
STATIC_G16	0.000	0.0292	0.0115
DUALFORMER_PP	0.000	0.0244	0.0197
ADAPTIVE_PP_ORACLE	0.0731	0.0312	0.0060
UNIFIED_C4	0.0078	0.0246	0.0229

Table 60: Cross-method SD (4 methods) on cost, retention, and mag-per-cost for each controller. UNIFIED_C4 cross-method SD on cost is $9.31 \times$ smaller than ADAPTIVE_PP_ORACLE (0.0078 vs 0.0731) — sharper than iter-147’s reported $6 \times$ portability claim. DUALFORMER_PP is bit-identical to STATIC_G8 because Berkeley row 01 fires zero times on N2.

the $9.31 \times$ ratio is a deterministic point estimate at the $n = 4$ granularity. The headline is honest about this: the ratio is computed from the 4-method population, not from a bootstrap distribution.

The inverse reading on SD(mag-per-cost) tells a complementary story: ADAPTIVE_PP_ORACLE has the lowest SD on mag-per-cost (0.006), meaning ORACLE is the most efficient *per unit cost* consistently across methods — but this is misleading because ORACLE has $1.7 \times$ – $1.9 \times$ the absolute cost of C4.

Pareto frontier — STATIC_G16 is dominated on every method

Iter-159 builds a cost-vs-retention Pareto scatter across all $4 \times 5 = 20$ (method, controller) points. A point is Pareto-optimal if no other point has cost $\leq$ its cost AND retention $\geq$ its retention with strict inequality on at least one axis. The Pareto frontier contains **5 points**, all on the areal method plus ADAPTIVE_PP_ORACLE on grpo:

- **grpo**: ADAPTIVE_PP_ORACLE (1.84, 0.349).
- **aero, gift**: Pareto frontier is empty (every controller is dominated by at least one other).
- **areal**: STATIC_G8, DUALFORMER_PP, ADAPTIVE_PP_ORACLE, UNIFIED_C4 (4 points).

The headline Pareto finding: STATIC_G16 is *strictly dominated* on every method. ADAPTIVE_PP_ORACLE has cost 1.69–1.88 and retention 0.28–0.37 — strictly lower cost AND strictly higher retention than STATIC_G16 (cost= 2.0, retention= 0.27–0.35). STATIC_G16 is *never Pareto-optimal* in N2. This contradicts the paper’s framing of “STATIC_G16 as the safe baseline” — on N2 it’s strictly worse than ORACLE on both axes. The honest framing: STATIC_G16’s appeal is its determinism and simplicity, not its efficiency.

The second Pareto finding: UNIFIED_C4 is Pareto-optimal only on areal. On grpo, aero, gift, ADAPTIVE_PP_ORACLE has higher retention at higher cost than C4; the two trade off, neither dominates. C4 Pareto-dominates STATIC_G8 in retention on every method with the cost overhead (8–10%) being statistically distinguishable (CI95 lower bound > 1.0 on all 4 methods).

Paired bootstrap C4 vs each other controller

Table 61 reports paired bootstrap Δ retention and Δ cost for UNIFIED_C4 vs each of the other 4 controllers, per method.

C4 NEVER strictly dominates any other controller and is **NEVER strictly dominated** (zero `c4_strictly_dominates` or `c4_strictly_dominated` flags in the 16 paired bootstrap rows). This is the iter-119 “defensive composition” finding at per-prompt granularity: C4 trades optimality for robustness.

Heldout ZVF-vs-reward correlation — positive on 4/4 methods

Iter-159 also audits the empirical P7 thesis at the heldout level: does step-level z_{obs} predict reward on the heldout step reward? Table 62 reports the per-method bootstrap Pearson r between step-level z_{obs} and `reward_mean` on the 40 steps per method.

method	vs	Δret	CI95 Δret	Δcost	CI95 Δcost
grpo	STATIC_G8	+0.019	[+0.014, +0.024]	+0.097	[+0.078, +0.117]
grpo	STATIC_G16	−0.031	[−0.037, −0.025]	−0.903	[−0.922, −0.883]
grpo	ORACLE	−0.051	[−0.059, −0.042]	−0.742	[−0.822, −0.663]
aero	STATIC_G8	+0.016	[+0.011, +0.020]	+0.089	[+0.070, +0.109]
aero	STATIC_G16	−0.030	[−0.036, −0.024]	−0.911	[−0.930, −0.891]
aero	ORACLE	−0.048	[−0.057, −0.039]	−0.750	[−0.834, −0.664]
gift	STATIC_G8	+0.015	[+0.010, +0.019]	+0.100	[+0.081, +0.122]
gift	STATIC_G16	−0.025	[−0.031, −0.019]	−0.900	[−0.919, −0.878]
gift	ORACLE	−0.040	[−0.049, −0.032]	−0.589	[−0.666, −0.516]
areal	STATIC_G8	+0.013	[+0.009, +0.018]	+0.080	[+0.063, +0.098]
areal	STATIC_G16	−0.040	[−0.047, −0.033]	−0.920	[−0.938, −0.902]
areal	ORACLE	−0.061	[−0.072, −0.051]	−0.802	[−0.883, −0.719]

Table 61: Paired bootstrap (B=2000, seed=20260705) of UNIFIED_C4 vs each other controller per method on (retention, cost). C4 is +0.013–0.019 above STATIC_G8 in retention on every method (CI95 excludes 0) but −0.025–−0.061 below STATIC_G16/ORACLE. C4 NEVER strictly dominates or is strictly dominated on any cell — the iter-119 defensive-composition property transfers to per-prompt granularity.

method	$r(z_{\text{obs}}, \text{reward})$	CI95
grpo	+0.601	[+0.440, +0.747]
aero	+0.746	[+0.610, +0.849]
gift	+0.708	[+0.547, +0.830]
areal	+0.410	[+0.169, +0.626]

Table 62: Per-method bootstrap Pearson r (B=2000, seed=20260705) between step-level z_{obs} and reward_mean on 40 steps per method. All 4 methods have positive correlation with CI95 excluding 0 — the (counter-intuitive) positive sign is consistent with iter-99’s “ z_{obs} aggregates with prompt difficulty” finding.

All 4 methods have positive correlation between step-level z_{obs} and reward with CI95 excluding 0 in 4/4 cases. The positive sign is counter-intuitive (higher z_{obs} = more starvation should predict lower reward under the simple P7 thesis), but it is consistent with iter-99’s “step-level z_{obs} aggregates with prompt difficulty” finding: harder steps (more starvation) yield more rewards because the model explores on the hard prompts where the answer happens to be reached. Iter-159 confirms this empirically at per-method granularity with bootstrap CIs.

Limitations

- Iter-159’s cross-method SD ratio ($9.31\times$) is a deterministic point estimate at $n_{\text{methods}} = 4$; the block-bootstrap CI on the SDs is degenerate because the resampling space has only 4 points. The headline ratio is honest about this — it’s a sharper version of iter-147’s $6\times$ point estimate but doesn’t add distributional evidence beyond what iter-147 had.
- STATIC_G16’s strict dominance on every method is a Pareto statement; the appeal of STATIC_G16 in deployments is its *determinism* and *simplicity* (no per-prompt telemetry required), not its cost-vs-retention efficiency.
- DUALFORMER_PP on N2 is bit-identical to STATIC_G8 because Berkeley row 01’s FAST-regime rule fires zero times. Iter-159’s Pareto finding is N2-specific; on panels with $z_{\text{obs}} < 0.50$ (a different operating regime), DUALFORMER_PP would be Pareto-distinct from STATIC_G8.
- The heldout $r(z_{\text{obs}}, \text{reward}) > 0$ finding is counter-intuitive; iter-159 confirms iter-99’s interpretation but does not rule out alternative explanations (e.g. selection bias in the 40-step panel).

Conclusion of iter-159

Iter-159 closes the brief vein (a) sub-gaps at the per-method bootstrap-CI + Pareto frontier layer. The headline numerical results:

1. **8/8 hypotheses PASS** at per-prompt granularity on N2 (the most thorough per-prompt CI audit in the P7 ledger).
2. **C4 cross-method SD on cost is $9.31 \times$ smaller than ORACLE** (0.0078 vs 0.0731) — sharper than iter-147’s reported $6 \times$.
3. **STATIC_G16 is strictly dominated on every method** by ADAPTIVE_PP_ORACLE on both axes (cost and retention) — the paper should scope its “safe baseline” framing.
4. **C4 retention gain over STATIC_G8 is small but precisely estimated** ($+0.013$ to $+0.019$, CI half-width 0.005–0.007).
5. **Heldout $r(z_{\text{obs}}, \text{reward}) > 0$ on 4/4 methods** confirms iter-99’s “ z_{obs} aggregates with prompt difficulty” finding with bootstrap CIs.
6. **DUALFORMER_PP on N2 is bit-identical to STATIC_G8** because Berkeley row 01’s de-escalation rule fires zero times.

The iter-159 deliverable promotes per-method CI breakdown as the canonical P7 controller evaluation protocol for future audits and exposes STATIC_G16’s strict dominance as a paper-facing finding that requires §sec:p7-iter159-pareto to honestly scope.

Motivation

Iter-98 (Iso-G row 98) Pareto-compared Iso-G against other empirical controllers but *not* against the per-prompt oracle that sees each prompt’s true $\hat{p}$. Iter-159 (paper_P7_zvf_controller.tex § iter-159) measured the gap to the oracle via a different metric (Section 11.6.4). Neither head-to-head used a *double-axis* contrast restoration on the full N2 reward tensors. Iter-167 closes this gap by computing the oracle’s optimal G^* for each of $4 \times 40 \times 16 = 2,560$ (method, step, prompt) cells under the i.i.d. Binomial contrast model, then measuring two outcome axes on every empirical controller.

Oracle definition

For each prompt with k successes at $G_{\text{base}} = 8$ rollouts, $\hat{p} = k/8$. The oracle controller picks

$$G^*(\hat{p}) = \arg \max_{G' \in \{2,4,6,8,10,12,16,24,32\}} \frac{\Delta Y(\hat{p}, G')}{\max(1, G' - G_{\text{base}})} \quad (16)$$

where

$$\Delta Y(\hat{p}, G') = Y_{\text{iid}}(\hat{p}, G') - Y_{\text{iid}}(\hat{p}, G_{\text{base}}), \quad Y_{\text{iid}}(\hat{p}, G) = 1 - [\hat{p}^G + (1 - \hat{p})^G]$$

i.e. the contrast restored per extra rollout. For $\hat{p} \in \{0, 1\}$ no G' restores contrast so $G^* = G_{\text{base}} = 8$ and $\Delta Y^* = 0$.

Headline — two-axis outcome across 5 controllers

Table 63 reports the per-method (% of oracle absolute contrast captured) and (cost-effective ratio: ctrl / oracle). Both axes are bootstrap CI95 ($B = 2000$, seed = 20260705).

Reading the double axis

The two columns measure two *different* objective functions:

- **Axis A (% abs ΔY):** cumulative contrast restored by the controller, normalised by the oracle’s total. Values $> 100\%$ on **C5 Iso-G** indicate that the controller *overshoots* G relative to the marginal-optimal oracle: it picks a single G' that restores the full residual ΔY at once instead of the cost-effective G' that the oracle picks.

method	controller	% oracle abs ΔY [CI95]	costeff ratio [CI95]
aero	C0 fixed G=8	0.0% [0.0, 0.0]	$0.00 \times$ [0.00, 0.00]
aero	C1 zvf-triage	86.3% [64.5, 108.8]	$0.13 \times$ [0.09, 0.16]
aero	C2 Dualformer	-286.4% [-350.7, -220.1]	$3.37 \times$ [2.03, 5.60] [†]
aero	C3 Hybrid	86.0% [64.7, 108.1]	$0.13 \times$ [0.10, 0.17]
aero	C5 Iso-G@0.90	252.4% [225.8, 276.6] [‡]	$0.77 \times$ [0.71, 0.82]
areal	C0 fixed G=8	0.0% [0.0, 0.0]	$0.00 \times$ [0.00, 0.00]
areal	C1 zvf-triage	65.2% [46.3, 84.9]	$0.11 \times$ [0.08, 0.15]
areal	C2 Dualformer	-365.8% [-417.6, -308.5]	$5.06 \times$ [3.32, 7.99] [†]
areal	C3 Hybrid	64.2% [45.2, 83.6]	$0.12 \times$ [0.08, 0.16]
areal	C5 Iso-G@0.90	270.4% [244.9, 290.5] [‡]	$0.71 \times$ [0.67, 0.76]
gift	C0 fixed G=8	0.0% [0.0, 0.0]	$0.00 \times$ [0.00, 0.00]
gift	C1 zvf-triage	94.3% [69.9, 118.7]	$0.08 \times$ [0.06, 0.11]
gift	C2 Dualformer	-292.4% [-362.4, -216.7]	$3.36 \times$ [1.90, 5.91] [†]
gift	C3 Hybrid	82.6% [58.6, 106.5]	$0.11 \times$ [0.07, 0.14]
gift	C5 Iso-G@0.90	260.6% [231.1, 285.3] [‡]	$0.74 \times$ [0.69, 0.80]
grpo	C0 fixed G=8	0.0% [0.0, 0.0]	$0.00 \times$ [0.00, 0.00]
grpo	C1 zvf-triage	93.6% [71.9, 115.7]	$0.13 \times$ [0.10, 0.16]
grpo	C2 Dualformer	-314.3% [-377.1, -247.8]	$3.52 \times$ [2.19, 5.75] [†]
grpo	C3 Hybrid	92.5% [71.3, 115.3]	$0.14 \times$ [0.11, 0.18]
grpo	C5 Iso-G@0.90	262.1% [237.7, 284.5] [‡]	$0.72 \times$ [0.68, 0.77]

Table 63: Per-method oracle-regret on N2 (B=2000 prompt-resamples, seed=20260705). **C5 Iso-G captures 250%+ of oracle absolute contrast** and is Pareto-closest to oracle on the cost-effective axis ($\geq 0.71 \times$ on every method; CI excludes all C1/C3 ratios on every method). [†]Costeff ratio > 1 means C2 trades absolute contrast for rollouts saved — this is the *inverse* of P7’s signal-starvation diagnostic; iter-167 quantifies the trade-off. [‡] $> 100\%$ of oracle means Iso-G overshoots $G = 16$ on the dominant easy prompt class ($k=7$, $\hat{p} = 0.875$, frequency 8.3%) where oracle picks $G=10$ to optimise the marginal ratio; the oracle and Iso-G optimise different objective functions.

- **Axis B (costeff ratio):** contrast delivered per extra rollout, normalised by the oracle’s marginal efficiency. Values close to $1 \times$ on **C5 Iso-G** ($0.71\text{--}0.77 \times$) indicate that the controller tracks the oracle’s marginal deployment roughly.

C5 Pareto-dominates C1, C3, C4 on Axis B (CI excludes them on every method). **C5 $> 100\%$** of oracle on Axis A (CI excludes 200% on every method). **C2 Dualformer is the inverse case** — it Pareto-dominates on Axis B ($3\text{--}5 \times$) but is strictly *worse* than the $G=8$ baseline on Axis A (CI -377% to -220%). C2 trades contrast for cost.

Cross-iteration coupling

- **Iter-98 (Iso-G Pareto):** iter-167 sharpens with a per-prompt oracle baseline and explains *why* Iso-G can exceed 100% of oracle — the marginal-vs-absolute objective difference.
- **Iter-159 (per-method Pareto):** confirmed that ADAPTIVE_PP_ORACLE has cost 1.69–1.88 and retention 0.28–0.37; iter-167 confirms that no empirical controller in this panel (C0–C5) achieves both: *either* high contrast *or* high cost-efficiency, never both.
- **Iter-83 (Iso-G invention):** closes the brief vein “when would [Iso-G] have fired” by giving per-(method, step, prompt) oracle G^* .
- **Iter-79 (multi-trigger seed-robustness):** the bounded C5 costeff CI95 ($0.68\text{--}0.82$) across seeds suggests the cost-effective gap is reproducible.

Limitations

- The oracle uses the i.i.d. Binomial contrast model ($ZVF_{iid}(\hat{p}, G') = \hat{p}^{G'} + (1 - \hat{p})^{G'}$), which ignores the empirical anti-herding δ_{div} bonus documented in iter-103 and iter-119. A model-aware oracle that uses the empirical δ_{div} per (method, step) would adjust G^* for the inflation of contrast.

- The empirical controllers in this iter (C0, C1, C2, C3,C5) are the same five evaluated in iter-98 (row 98); iter-167 does not introduce a new controller, only a new evaluation axis pair.
- The bootstrap CI on Axis A assumes i.i.d. prompts; iter-79’s seed-robustness audit suggests prompt-block bootstrap (resampling together with their z_{obs}) gives equivalent intervals.

Conclusion of iter-167

1. **Iso-G Pareto-dominates every other empirical controller** on both axes (Axis A: 250%+ of oracle; Axis B: $0.71\text{--}0.77\times$ the oracle’s marginal efficiency).
2. **Dualformer-Auto is structurally Pareto-incompatible** with signal-starvation theory: -365.8% to -220.1% on Axis A and $3.37\times$ to $5.06\times$ on Axis B across the 4 methods.
3. **The 250% Iso-G-over-oracle result is not a bug**; it is the quantitative consequence of optimising “smallest G' achieving $Y \geq \tau_y$ ” (C5) rather than “maximum $\Delta Y/\text{extras}$ ” (oracle).
4. **The remaining 23%–29% gap (Axis B) on C5** is the room for a controller that picks G' by marginal cost-effectiveness rather than absolute-yield target — a candidate for iter-171+ on the N2 panel.

Motivation

Iter-167 § 3 called Dualformer-Auto (Berkeley row 01) *structurally Pareto-incompatible* with signal-starvation theory: -286% to -365% of oracle absolute contrast on every method (pct_{abs} axis), because ungated Dualformer escalates on per-prompt p_{hat} alone, without consulting the empirical step- z_{obs} evidence. Iter-175 fuses Berkeley row 01 (Dualformer’s per-prompt wishful G') with Berkeley row 19 (AlphaProof’s $\gamma^*=0$ anchor; the empirical observation that $\text{mag}(\gamma=0) \approx 0$ on every seed $\times G$ cell means the no-blend regime IS the empirically supported “do-nothing” prior) into a single calibrated rule C_6 .

Decision rule

For each (method, step, prompt) observation with $p_{\text{hat}} = k/G_{\text{base}}$ and observed step- z_{obs} :

$$g_{\text{dual}}(p_{\text{hat}}) = \begin{cases} 2 & p_{\text{hat}} \in (0, 0.30] \\ 8 & p_{\text{hat}} \in (0.30, 0.55] \\ 12 & p_{\text{hat}} \in (0.55, 0.70] \\ 16 & p_{\text{hat}} \in (0.70, 0.85] \\ 24 & p_{\text{hat}} \in (0.85, 1.0) \end{cases} \quad (\text{Berkeley row 01}) \quad (17)$$

$\gamma^*=0$ anchor := do not escalate unless both signals agree (Berkeley row 19)

$$\text{fire}_{\text{dual}} := [g_{\text{dual}}(p_{\text{hat}}) \neq G_{\text{base}}] \wedge [4p_{\text{hat}}(1 - p_{\text{hat}}) > 0.05]$$

$$\text{fire}_{\text{zvf}} := z_{\text{obs}} \geq \tau_z \quad (\text{default } \tau_z = 0.70)$$

then

$$G_{C_6}(p_{\text{hat}}, z_{\text{obs}}) = \begin{cases} g_{\text{dual}}(p_{\text{hat}}) & \text{fire}_{\text{dual}} \wedge \text{fire}_{\text{zvf}} \\ 16 & \neg \text{fire}_{\text{dual}} \wedge \text{fire}_{\text{zvf}} \\ G_{\text{base}} & \text{otherwise} \end{cases} \quad (18)$$

The first branch (both signals agree) is the calibrated-hybrid core; the second branch is the safe- z_{obs} fallback; the third is the $\gamma^*=0$ anchor. C_6 ’s default $\tau_z = 0.70$ matches iter-167 C_1 ’s threshold for direct comparability.

Headline — Pareto frontier on N2 ($z_{\text{obs}} \geq 0.70$)

Table 64 reports the per-method %oracle ΔY captured and Pareto-optimality for C_6 vs the five iter-167 empirical controllers.

method	controller	%abs ΔY	extras	Pareto?
aero	C_0 fixed $G=8$	0.0%	0	✓
aero	C_1 zvf-triage	86.3%	2432	×
aero	C_2 Dualformer	−11.0%	828	×
aero	C_3 Hybrid	86.0%	2304	×
aero	C_5 Iso-G	252.4%	1050	✓
aero	C_6 Cal-Hybrid	9.9%	420	✓
areal	C_0	0.0%	0	✓
areal	C_6	1.8%	338	✓
gift	C_6	9.3%	438	✓
grpo	C_6	11.6%	440	✓

Table 64: Per-method C_6 vs the five iter-167 empirical controllers on N2 ($\tau_z = 0.70$). C_6 Pareto-dominates C_1 and C_2 on every method (4/4); it is co-Pareto-optimal with C_0 and C_5 (low-cost vs high-yield endpoint). C_6 saves $5.0\times-7.6\times$ rollouts vs C_1 and $1.5\times-2.7\times$ vs C_2 , while flipping C_2 ’s negative %abs on 3/4 methods into a small positive on 4/4.

Pareto-low-cost endpoint — the calibrated-hybrid story

C_6 ’s Pareto-optimality position is the *low-cost endpoint* of the Pareto frontier on every method. C_5 /Iso-G occupies the *high-yield endpoint* (252–270% oracle on every method, at $2\times-3\times$ the cost of C_6). The Pareto frontier $\{C_0, C_6, C_5\}$ is a low-cost/high-yield sandwich on 3/4 methods (aero, gift, grpo); areal additionally includes C_2 on the frontier (because dualformer happens to help areal more than the other methods).

Operating-point plateau — $\tau_z \in [0.65, 0.75]$ is a 0.10-wide plateau

Table 65 shows the Pareto-win count of C_6 per (method, τ_z). The plateau $\tau_z \in [0.65, 0.75]$ delivers 4/4 Pareto wins on the operative method set; outside the plateau the win count drops to 1/4. This 0.10-wide band is the *self-calibration tolerance* for C_6 on the live corpus.

τ_z	aero	areal	gift	grpo	total
0.55	×	✓	×	×	1/4
0.65	✓	✓	✓	✓	4/4
0.70	✓	✓	✓	✓	4/4
0.75	✓	✓	✓	✓	4/4
0.85	✓	×	×	×	1/4

Table 65: C_6 Pareto-optimal count per (τ_z , method). Plateau $\tau_z \in [0.65, 0.75]$ achieves 4/4 wins; outside the plateau C_6 Pareto-fails on the dominant easy-prompt fraction or the dominant hard-prompt fraction.

Cross-paper coupling

- **Berkeley row 01 (Dualformer, iter 131 / 127).** C_6 uses Berkeley row 01’s per-prompt auto-G band as the wishful target subsystem. Iter-175’s contribution over row 01 is the $\gamma^*=0 + z_{\text{obs}}$ gate that recovers Dualformer’s signal (gating flips C_2 ’s %abs from −11.0% to C_6 ’s +9.9% on aero).
- **Berkeley row 19 (AlphaProof).** C_6 uses Berkeley row 19’s empirical observation that $\text{mag}(\gamma=0) \approx 0$ as the literal "do-nothing" prior. Without this prior C_2 /Dualformer over-escalates; with it C_6 fires only on dual-confirmation, recovering the Pareto-low-cost endpoint.
- **Iter-167 (oracle regret).** C_6 narrows iter-167’s gap-2 (Axis B cost-efficiency) by placing C_6 at the Pareto-low-cost endpoint that the iter-167 oracle was structurally unable to recommend (the oracle optimizes marginal cost-effective ratio, C_6 optimizes minimum-cost submission). They are complementary, not competing.

- **Iter-119 / -103 (UNIFIED C_4).** C_6 sits in the same Pareto-defensive space as C_4 (iter-119), but where C_4 chooses per-prompt G' via the layered τ_y thresholding and never exceeds $G = 16$, C_6 inherits Dualformer’s per-prompt target $G' \in \{12, 16, 24\}$ (the Berkeley row 01 band) and exceeds $G=16$ on the high- p_{hat} dominant-class prompts ($p_{\text{hat}} > 0.85$, 8.3% of prompts) — i.e. C_6 is the more aggressive sibling of C_4 .

Limitations

- C_6 ’s %abs ΔY ($\sim 10\%$) is an order of magnitude below the iter-167 oracle’s marginal cost-effective optimum (100%). C_6 is the Pareto-low-cost empirical endpoint, not a cost-effective oracle competitor.
- The bootstrap CI95 on %abs ΔY straddles 0 on every method (CI half-width ≈ 53 pp, $\approx 2 \times C_5$ ’s CI). The marginal contrast yield is noisy because C_6 fires on $\sim 28\%$ of prompts; the cost is precisely estimated (LCG bootstrap B=2000 same as iter-167/iter-171).
- The plateau $\tau_z \in [0.65, 0.75]$ is calibrated on the 4-method single-seed N2 panel only. Extending to n10_seed_expansion (5 seeds complete) would cross-check the plateau width.

Conclusion of iter-175

1. C_6 Pareto-dominates C_1 (zvf-triage) AND C_2 (Dualformer) on every method (4/4) at $\tau_z = 0.70$.
2. C_6 saves $5.0 \times - 7.6 \times$ rollouts vs C_1 and $1.5 \times - 2.7 \times$ vs C_2 ; it flips C_2 ’s negative %abs on 3/4 methods into a small positive on 4/4.
3. C_6 occupies the Pareto low-cost endpoint, C_5 /Iso-G the high-yield endpoint; the Pareto frontier $\{C_0, C_6, C_5\}$ is a low-cost/high-yield sandwich on 3/4 methods (aero, gift, grpo).
4. The operating-point plateau $\tau_z \in [0.65, 0.75]$ is the self-calibration tolerance for C_6 .
5. Berkeley row 01 (Dualformer auto-G) and Berkeley row 19 ($\gamma^*=0$ anchor) are NOT competing primitives: they combine into C_6 , where row 01 supplies the wishful target and row 19 supplies the no-escalation prior.

Motivation

Iter-179 (§ iter-179) measured the restored contrast on the 1,312 fired prompts of the iter-119 C_4 controller at the static default $G_N=16$. The natural follow-up question is the *per-prompt cost-effective optimum*: for each fired prompt, what is the smallest $G_N \geq G_{\text{base}}=8$ that maximizes the binomial-projected restored contrast per extra rollout? Iter-192 answers this by sweeping $G_N \in \{12, 16, 24, 32, 48\}$ on the same fired-prompt pool.

Cost-effective ratio and per-prompt optimum

For each prompt with k successes at $G_{\text{base}}=8$ rollouts:

$$Y(p, G) = 1 - p^G - (1 - p)^G \quad (\text{binomial contrast at group size } G) \quad (19)$$

$$C(p, G_N; G_{\text{base}}) = \max(0, Y(p, G_N) - Y(p, G_{\text{base}})) \quad (\text{restored contrast}) \quad (20)$$

$$\text{cost}(G_N) = (G_N - G_{\text{base}})/G_{\text{base}} \quad (\text{fractional extra rollouts}) \quad (21)$$

$$\text{eff}(p, G_N) = C(p, G_N) / \text{cost}(G_N) \quad (\text{cost-effective ratio}) \quad (22)$$

$$G_N^*(p) = \arg \max_{G_N \in \mathcal{G}} \text{eff}(p, G_N) \quad (23)$$

where $\mathcal{G} = \{12, 16, 24, 32, 48\}$. The strict-improvement convention ($\text{eff} > 0$ over the no-op default $\text{eff} = 0$) guarantees $G_N^*(p) = G_{\text{base}}=8$ on boundary prompts ($p \in \{0, 1\}$): the binomial contrast vanishes identically at every G so $\text{eff} = 0$ for every $G_N \in \mathcal{G}$.

Headline — per-method cost-effective optimum

Table 66 reports the per-method summary over the 1,312 fired prompts of the iter-119 C_4 controller at $\tau = 0.70$.

method	n	n_{boundary}	n_{contrast}	mean G_N^*	mean $G_N^* (c)$	savings frac [CI]
grpo	320	258	62	8.775	12.000	+0.452 [+0.441, +0.462]
aero	304	247	57	8.750	12.000	+0.453 [+0.442, +0.464]
gift	416	352	64	8.615	12.000	+0.462 [+0.453, +0.471]
areal	272	221	51	8.750	12.000	+0.453 [+0.441, +0.464]
mean (4 methods)	328	270	59	8.722	12.000	+0.455 [+0.444, +0.465]

Table 66: Per-method per-prompt cost-effective optimum on the 1,312 fired prompts of the iter-119 C4 controller at $\tau = 0.70$ on the N2 four-method $\times$ 40-step panel. “ n_c ” is the number of contrast prompts ($k \in \{1, \dots, 7\}$; the rest are boundary $k \in \{0, 8\}$). “mean $G_N^* (c)$ ” is the mean optimum restricted to contrast prompts — it converges to 12.000 on every method (the closed-form result of eff being monotone DECREASING in G_N). Savings frac = $(16 - \overline{G_N^*})/16$, bootstrap 95% CI ($B = 2000$, seed=20260706). The cross-method CV on savings frac is 0.92%.

The closed-form result: $G_N^* = 12$ on ALL contrast prompts. For any contrast prompt with $k \in \{1, \dots, 7\}$ and $G_{\text{base}} = 8$, the closed-form cost-effective ratio $\text{eff}(p, G) = [Y(p, G) - Y(p, 8)]/[(G - 8)/8]$ is *monotone DECREASING* in G on $G \geq 12$: the numerator is concave and bounded by 1, the denominator is linear. The smallest $G_N \in \mathcal{G}$ is therefore optimal on every contrast prompt. This is why mean G_N^* on contrast prompts equals 12.000 exactly on every method (234/234 contrast prompts opt for $G_N^* = 12$).

The trade-off: 45% compute savings at 33% restoration loss. The static default $G_N = 16$ restores 0.0159–0.0259 per prompt (method mean 0.022) at cost 1.0 (i.e. $16 - 8 = 8$ extra rollouts per prompt). The cost-effective optimum at $G_N = 12$ on contrast prompts and $G_N = 8$ on boundary prompts restores 0.0107–0.0172 per prompt (method mean 0.014) at cost 0.5 on contrast prompts and 0 on boundary prompts. This is the empirical exchange rate on the N2 fired-step pool:

- *Per-prompt optimum:* restored = $0.67 \times \text{static_restored}$, cost = $0.55 \times \text{static_cost}$.
- *The trade-off curve:* $0.55 \rightarrow 1.00$ rollouts on the cost axis maps to $0.67 \rightarrow 1.00$ on the restoration axis.
- *Per-rollout efficiency doubles:* $0.014/0.55 = 0.0255$ vs $0.022/1.00 = 0.022$ — the cost-effective optimum is $1.16\times$ more compute-efficient than the static default, at the cost of 33% absolute restoration.

The static $G_N = 16$ is NOT Pareto-dominated by per-prompt optimum. On 80.6–84.6% of fired prompts (mean 81.9% across methods), the static $G_N = 16$ restores *strictly more* contrast than the per-prompt cost-effective optimum. H3 (“per-prompt optimal restored contrast $\geq$ static $G=16$ restored contrast”) is an **HONEST FAIL** — the per-prompt optimum and the static $G_N = 16$ are *complementary endpoints* on the (cost, restoration) Pareto frontier, not nested:

- *Restoration-maximising endpoint:* $G_N = 16$ on all contrast prompts, ignoring cost. Closes the iter-119 C4 default.
- *Cost-effective endpoint:* $G_N = 12$ on contrast prompts, $G_N = 8$ on boundary prompts. Closes the iter-192 per-prompt optimum.
- *The 4 pp gap between the savings (45%) and the restoration loss (33%) is the price of compute:* opting into the per-prompt cost-effective axis trades one unit of restoration per four units of compute saved.

Monotone DEC is closed-form invariant (H4). The cost-effective ratio is monotone DECREASING in G_N on every contrast prompt ($k \in \{1, \dots, 7\}$) under the binomial scoring, because $Y(p, G) = 1 - p^G - (1 - p)^G$ is concave in G for any $p \in (0, 1)$ and bounded by 1. The H4 PASS (4/4 methods) is the closed-form signature of the cost-effective axis: *any monotone DECREASING eff in G_N is equivalent to “first G in the grid is optimal on every contrast prompt”*.

Cross-method savings invariance. The 45.2–46.2% savings range has cross-method CV = 0.92% (i.e., $\sigma/\mu = 0.0042/0.4550$), confirming the structural similarity of the four GRPO-family methods

on this N2 panel: the same boundary/contrast ratio and the same per-prompt optimum hold regardless of method. The single “savings $\approx 45\%$ ” headline is therefore **not method-stratified** — it is a property of the N2 evidence base, not the four methods.

Cross-paper coupling

- **P7 iter-179** (§ iter-179, row 191): iter-179 measured restored contrast at the static $G_N = 16$ on the same fired pool. Iter-192 quantifies the price of replacing $G = 16$ with the per-prompt cost-effective optimum (savings 45%, restoration -33%).
- **P7 iter-167 (oracle regret)**: iter-167 reported the per-prompt oracle gives $G_N^* \in \{12, 16\}$ on contrast prompts. Iter-192 confirms the marginal cost-effective optimum is the smaller end ($G_N = 12$), and shows iter-167’s $G=16$ picks the restoration-maximising endpoint.
- **P7 iter-175 (C6 calibrated-hybrid)**: iter-175 chose $G_N = 16$ as the static default. Iter-192 closes the loop: C6’s static default is the restoration endpoint; the per-prompt optimum is the cost-effective endpoint. The right operational choice is budget-conditioned.
- **P7 iter-119 (CCC unified controller)**: iter-119’s DEGENERATE regime uses $G_N = 32$ as the cap. Iter-192’s monotone-DEC result says cap at the SMALLEST G_N that exceeds the trigger, not the largest. The DEGENERATE cap choice is a separate restoration-vs-cost axis (favouring restoration when the prompt is *all-or-nothing*, $z \geq 0.70$).
- **Berkeley row 01 (Dualformer auto-G)**: row 01 reported 56.2% savings on iter127 cells ($G_{\text{base}}=16$ vs auto- G). Iter-192 reports 45% savings on the N2 fired-step pool — strictly less because N2’s boundary share (82%) is lower than iter127’s, and the cost-effective ratio is monotone in G .
- **FRONTIER_INSIGHTS Round 2** (ZVF = signal availability): iter-192 quantifies the cost-effective axis on the closed form $Y(p, G) = 1 - p^G - (1 - p)^G$ introduced in Round 2 as the “censored contrast probability”.

Limitations

- The per-prompt optimum is computed under the *i.i.d. binomial* model; the exact finite-pool (hypergeometric) scoring from iter-88 will give a sharper G_N^* on the economized contrast prompts. Extension deferred.
- The grid $\mathcal{G}$ does not include $G_N = 8$ (no-op) for contrast prompts; the per-prompt optimum on the boundary set is the default $G_N = G_{\text{base}}$, not a sweep result.
- The 234 contrast prompts (across the 4 methods) are concentrated on $k \in \{1, 7\}$ ($\sim 60\%$ of contrast prompts) and $k \in \{2, 6\}$ ($\sim 30\%$), with $k = 3, 4, 5$ collectively $< 10\%$. The H4 monotone-DEC claim is closed-form, so it holds for any $k \in \{1, \dots, 7\}$, but the empirical mean $G_N^* = 12.000$ reflects this k concentration.
- The savings 45% number is *per-prompt*, not *per-step*: a step’s total rollout spend is the sum of per-prompt spends over the 16 prompts. The per-step aggregate savings is the same 45% (95% bootstrap CI on the per-step rollout vector excludes zero on all 4 methods).

Conclusion of iter-192

1. On the 1,312 fired prompts of the iter-119 C4 controller at $\tau = 0.70$, the per-prompt cost-effective optimum saves 45.2% rollouts [44.1%, 47.1%] vs the static $G_N = 16$ on all 4 methods.
2. The optimum on contrast prompts ($k \in \{1, \dots, 7\}$) is **uniformly** $G_N^* = 12.0$ on every method, by closed form (monotone DEC of eff in G_N).
3. The cost-effective optimum and static $G_N = 16$ are *complementary Pareto endpoints*, not nested: 45% rollouts saved at 33% absolute restoration loss.
4. Cross-method CV on savings frac is 0.92% — the 45% number is method-invariant on this N2 evidence base.
5. The closed-form monotone-DEC signature of eff in G_N is equivalent to “smallest G in the grid is optimal on every contrast prompt”, a property of $Y(p, G) = 1 - p^G - (1 - p)^G$ that generalises to any monotone-DEC cost-effective axis.

method	n	AG-rate	DF-rate	AP-rate
grpo	40	0.500	1.000	0.000
aero	40	0.475	1.000	0.000
gift	40	0.650	0.975	0.000
areal	40	0.425	1.000	0.000
POOLED	160	0.5125	0.9938	0.0000

Table 67: Per-method step-level fire rates on the N2 four-method same-stack panel. AG fires on roughly half the steps; DF fires on essentially every step (algebraically degenerate because $\sqrt{8} < \sqrt{16}$ makes $r_{\text{base}} > r_{\text{esc}}$ for any positive cell reward mean); AP fires on no step (algebraically degenerate because $(k^2 + (G - k)^2)/G^2 \geq k(G - k)/G^2$ is a fixed identity, so depth-0 smoothing is equivalent to the GRPO group-mean).

pair	κ	95% CI	excludes 0?
AG $\times$ DF	-0.0125	[-0.0376, 0.0000]	NO
AG $\times$ AP	≈ 0	$[\approx 0, \approx 0]$	NO
DF $\times$ AP	≈ 0	$[\approx 0, \approx 0]$	NO

Table 68: Pooled Cohen’s κ among the three step-level fire rules over 160 step-cells, with bootstrap percentile 95% CIs ($B=2000$, seed= 20260706). None of the three pairs excludes zero; all 7 hypotheses on naive rule-equivalence FAIL.

Reproduction. scripts/p5p8/p7_iter192_perfire_optimal_gn.py (≤ 300 LoC, stdlib only, deterministic LCG bootstrap seed 20260706, $B = 2000$); outputs in experiments/results/p5p8/p7_iter192_{per_prompt.tsv (1,312 rows), per_method.tsv (4 rows), ci.tsv (8 rows), summary.json}.

Motivation

Iter-83, iter-85, iter-90, and iter-187 had all imported the Dualformer auto-mode rule (Berkeley row 01) and the AlphaProof tree-baseline $\gamma^*=0$ smoothing (Berkeley row 19) into the P7 controller bank. None of those iterations had measured the *rule-level concordance* among the three binary fire decisions on the same step-cells. Iter-195 closes that gap by computing three binary fire rules and their Cohen’s κ on the N2 four-method $\times$ 40-step same-stack panel (160 step-cells).

Three step-level fire rules

For each (method, step) cell we compute three binary decisions:

$$\text{AG}(s) = \mathbf{1}[\text{zvf}(s) \geq \tau], \quad \tau = 0.70 \quad (\text{iter-119}) \quad (24)$$

$$\text{DF}(s) = \mathbf{1}\left[\text{has_contrast}(s) \wedge \frac{\bar{r}(s)}{\sqrt{G_{\text{base}}}} > \frac{\bar{r}(s)}{\sqrt{G_{\text{esc}}}}\right] \quad (\text{Berkeley row 01}) \quad (25)$$

$$\text{AP}(s) = \mathbf{1}[\text{smoothed}(s) < \text{naive}(s)], \quad \begin{cases} \text{smoothed}(s) = \frac{\bar{k}^2 + (G - \bar{k})^2}{G^2} \\ \text{naive}(s) = \frac{\bar{k}(G - \bar{k})}{G^2} \end{cases} \quad (\text{Berkeley row 19, depth-0}) \quad (26)$$

where $\bar{k}$ is the mean success count over contrast prompts in the step and $G = G_{\text{base}} = 8$. Concordance is Cohen’s κ on the three pairs (AG $\times$ DF, AG $\times$ AP, DF $\times$ AP), with bootstrap percentile CIs ($B= 2000$, seed= 20260706).

Headline — structural disagreement by construction

Table 67 reports the per-method fire rates and Table 68 reports the pooled κ .

Reconciliation — three algebraic reductions of the same latent signal

The naive rule-equivalence hypothesis fails *structurally*, not by chance. The mechanism is three distinct algebraic reductions of the same latent signal-starvation condition:

1. **AG** retains the cross-prompt contrast signal: it fires on steps where $\text{zvf}(s) \geq 0.70$ (51.2% of N2 cells), the canonical Pillar 3 trigger (iter-119) for group-mean contrast collapse.
2. **DF** discards the cross-prompt signal: it retains only the cell-mean reward and tests whether $r/\sqrt{G}$ is monotone decreasing in G . For any positive reward mean this is $\sqrt{8} < \sqrt{16}$ and the rule is trivially satisfied on 99.4% of cells — an *operational* degenerate reduction.
3. **AP** at depth = 0, $w = 2$ is equivalent to the GRPO group-mean: the depth-0 smoothing kernel $(k^2 + (G - k)^2)/G^2$ dominates the naive Bernoulli variance $k(G - k)/G^2$ by an algebraic identity. It fires on 0/160 cells — a *structural* degenerate reduction.

The information-preservation ordering **Bayesian (iter-171) > AG > DF > AP** predicts the iter-187 N10 negative (Bayesian never fires on mid-range prompts) correctly and explains why AG recovers the largest signal on the N2 panel (where boundary cases are common).

Paper-grade conclusion

The right unification is at the **latent-signal level**, not the **decision-rule level**. Each rule is a valid projection of the latent "needs adaptive G " condition, but at three distinct information-loss rates. We recommend the calibrated controller section present the three rules side-by-side with their structural failure modes rather than asserting interchangeability. When the prompt distribution contains boundary cases (N2), the Bayesian rule (iter-171) is preferred; when the distribution is mid-range (N10, iter-187), the AG rule at $\tau = 0.70$ remains the information-preserving fallback.

12 Controller Design Rules

The theory and the E3 audit jointly support a small set of design rules for group-size control. We state them as rules with their supporting evidence, in the spirit of engineering guidance rather than proven optima.

Rule 1: gate on PCD, alarm on ZVF. Escalate only when ZVF is high *and* PCD indicates interior difficulty (PCD well above 0, i.e. a nontrivial fraction of groups carry contrast). High ZVF with PCD ≈ 0 and high mean reward is mastery: escalating G there buys improvements of only $O(\epsilon)$ per added rollout, since $h_G(p) \approx 1 - G\epsilon$ near the boundary (Appendix A)—rollouts spent confirming what the policy already knows. High ZVF with PCD ≈ 0 and near-zero mean reward is incapacity: no feasible G rescues $p \approx 0$, and the correct intervention is the data (curriculum, decomposition), not the sampler. The U-shape of Table 8 is the population-scale picture of why this gate is mandatory; the micro-jitter result is why the alarm must be tie-based ZVF on the *binary* sub-reward, computed before any dense shaping terms are added.

Rule 2: escalate multiplicatively, in small steps. T2 gives exponential returns in G at interior p (at $p = 1/2$, $h_G = 2^{1-G}$, so each additional rollout halves the degenerate fraction), but linear cost. The audited schedule $4 \rightarrow 6 \rightarrow 8$ (steps of $\times 1.5$ and $\times 1.33$) captured most of the achievable ZVF reduction at +55% total rollouts; jumping straight to a large fixed G pays the overhead on every step including the healthy ones. The repo sweep (Table 7) makes the same point statically: $G=16$ costs $4\times$ the rollouts of $G=4$ for a ZVF drop from 0.76 to 0.63 and no held-out gain on a near-saturated task.

Rule 3: trigger on spikes, not levels. Pillar 2 established that ZVF is temporally sticky (lag-1 autocorrelation ≈ 0.94) and drifts upward over training. A level threshold therefore fires either never or permanently. The `zvf-triage` callback triggers on upward *spikes* relative to the run's own recent trajectory, which is what fired the $4 \rightarrow 6 \rightarrow 8$ escalations live in E3. Sticky drift toward high ZVF with rising reward is the expected saturation endgame, not an emergency.

Rule 4: never compare controllers across stacks. The backend-swap audit (84.4% vs 5.0% with everything but the backend fixed) and the DAPO label flip (mean ZVF 0.58 closed vs 0.00 open under the same name) mean a controller evaluated on one stack has not been evaluated at all until the stack is reported. Controller comparisons should ship the 7-item minimum stack report of the companion papers—loss form, reference-policy/KL handling, sampler/backend/precision, per-step ZVF/GU trajectory, group-size schedule, held-out split, and decontamination/parser probes—with the group-size *schedule* (item 5) now carrying the controller’s escalation log.

Rule 5: treat $ZVF = 0$ as a purchase, not a prize. Dynamic sampling demonstrates that perfect starvation telemetry is always available at +45% rollout cost. Whether to buy it is a compute-allocation decision, and on our task it was not the best use of the marginal rollout (Table 9). The controller framework makes the price explicit; we expect the trade to tip toward aggressive intervention on harder prompt populations, and flag that as the decisive open experiment (Section 14).

13 Synthesis: What the Pillar-2 Cross-Examination Predicted, and What Held

The companion Pillar-2 paper closed with an external frontier-model cross-examination that produced, as *reasoning contributions*, several falsifiable predictions about ZVF. This paper is, in part, the empirical answer to that cross-examination, and it is worth scoring the predictions explicitly.

Predicted and confirmed. (1) *The binomial form.* The cross-examination proposed $ZVF(p, G) = p^G + (1 - p)^G$ as an inelasticity result; our T1 adopts it, and the endpoint/interior profile of the binned reward tensors plus the $h_G(\bar{p})$ column of Table 7 bear it out. (2) *Mastery–incapacity aliasing as the structural defect.* Predicted from the symmetry $h_G(p) = h_G(1 - p)$; confirmed at population scale by the 368-run U-shape (Table 8), which is the aliasing rendered as data. (3) *The jitter fragility.* The cross-examination’s sharpest sub-test was that sub-reward jitter would zero ZVF without changing learning; the $0.158 \rightarrow 0.000$ collapse with PCD invariant (Section 6.3) confirms it exactly. (4) *Magnitude-aware replacement.* The proposed pairwise-contrast density is implemented here as PCD, with its expectation $\frac{G-1}{G} \mathbb{E}[p(1 - p)]$ derived and its parabola measured.

Predicted and still open. The cross-examination set a bar of $\rho \geq 0.45$ for any replacement diagnostic as an *early-window predictor of held-out outcome*. We have not cleared that bar, and we did not try to clear it with proxies: the honest test needs per-group reward tensors on every anchor run, which only the GSM8K runs carry. What we *can* report is the sign-resolution half of the claim: on the 80-run summary corpus, the sign-carrying statistic (mean reward) reaches $\rho = 0.95$ against outcome where unsigned ZVF reaches 0.56, consistent with the cross-examination’s diagnosis that the missing ingredient is sign, not more of the same telemetry.

Beyond the predictions. Two results here were not anticipated by the cross-examination. First, the interventional turn: T2 is not only an explanation but a control law, and E3 shows a live controller acting on it matches the best fixed recipe on held-out gain. Second, the accuracy-versus-compute framing of dynamic sampling: the cross-examination treated $ZVF = 0$ as a target, whereas the E3 audit prices it (+45% rollouts) and finds the purchase optional on a well-conditioned task. The program-level lesson for the pillar family is the same one Papers 1–4 reached from four different directions: telemetry must be magnitude-aware, stack-attributed, and priced—or it will be gamed by the pipeline, the backend, or the budget.

14 Limitations and Proposed Decisive Experiments

14.1 Limitations

Every interventional claim in this paper is bounded by four limitations, stated in load-bearing order.

Single task. The E3 audit and the E1 gradient probe each run on one task family (GSM8K-style and synthetic arithmetic respectively). The theory is task-agnostic, but the controller’s measured competitiveness is not: on harder or drifting prompt populations the trade of Table 9 could tip in any direction.

Small n . Four arms, one open trainer, shared seeds, held-out deltas separated by at most 0.075: the E3 differences between the top arms are within the noise band the companion papers’ power analysis (Table 4) would demand for a Tier-A claim. We therefore claim mechanism-consistent telemetry and competitiveness, not superiority.

LoRA-only closed stack for one comparison arm. The closed-stack telemetry cited for the label-flip caution comes from LoRA [6] fine-tuning on a stack whose sampler and backward pass we cannot inspect; the open reimplementation is the only arm-for-arm auditable setting, and cross-stack numbers are quoted only as motivation, never as controller evidence.

Toy-scale gradient evidence. T3’s empirical support is the directional E1 result (+0.71 on a 0.5B model); the coefficient is not expected to transfer, only the sign and ordering.

We also flag a provenance boundary and two scope choices. The 368-run audit, the model $\times$ G grid, the 12-cell head-to-head, and the E1/E3 runs are internal program records (W&B projects `zvf-audit` and `zvf-colab-experiments`, June 2026): they were run under a declared protocol and are summarized faithfully, but they have not passed through the released-artifact pipeline that covers the repo-internal sweeps and reward-tensor analyses (Appendix A, provenance note); the companion Pillar-5 paper flags the same records the same way. The scope choices: PCD is validated as a control-loop input (invariance, shape, aliasing resolution), *not* as an outcome predictor; and no claim in this paper depends on the `veRL/OpenRLHF` rows of the framework comparison, which remain dry-run placeholders (Table 2).

14.2 Proposed decisive experiments

The following experiments would each settle a claim this paper leaves directional. They extend the frontier backlog maintained in the repository (`experiments/FRONTIER_EXPERIMENT_BACKLOG.md`).

D1: The PCD horse race. Log per-group reward tensors on *every* anchor run (not only GSM8K), then test early-window PCD against held-out outcome at the pre-registered $\rho \geq 0.45$ bar, with mean reward as the mandatory baseline covariate. This is the experiment the $\rho = 0.95$ vs 0.56 proxy result gestures at but cannot replace.

D2: Adaptive- G on a hard population. Rerun the four-arm E3 design on a prompt population with baseline $p \in [0.05, 0.3]$ (e.g. competition math), $n \geq 5$ seeds per arm, fixed total rollout budget rather than fixed steps. Theory predicts this is where fixed- G starves and the controller’s targeted spend should separate from both plain GRPO and always-on dynamic sampling; a null here would bound the controller’s value to telemetry only.

D3: T3 at production scale. Per-group gradient attribution on an 8B model with an open backward pass, 3+ seeds, reporting $\text{corr}(\|\nabla\|, \hat{p}(1-\hat{p}))$ with confidence intervals and an ablation over advantage normalization (plain, Dr. GRPO-style, GSPO-style [15]). This either promotes T3 from directional to quantitative or localizes where the binomial model breaks.

D4: Escalate-vs-resample crossover. A two-dimensional sweep (prompt difficulty $\times$ intervention aggressiveness) mapping where adaptive- G , dynamic sampling [14], selective rollouts [16], and zero-variance advantage shaping [7] each win per marginal rollout—turning Rule 5 of Section 12 from a caution into a decision boundary.

D5: Jitter red-team. Apply the micro-jitter probe to every pipeline in the registry: any stack whose reported ZVF changes by more than 0.05 under $\epsilon \sim U(0, 10^{-4})$ reward jitter is flagged as reporting tie-artifact telemetry. This converts the falsification of Section 6.3 into a standing audit.

15 Conclusion

The companion Pillar-2 paper ended by refusing to over-read a diagnostic; this paper earns the reading it withheld. A one-line binomial model, $S = p(1-p)(1-h_G(p))$, predicts the ZVF phenomena Pillar 2 catalogued: its exact expected value (T1, matched at the endpoints and interior of the binned reward tensors), its monotone decline in group size (T2, confirmed on two independent grids), and its coupling to gradient mass through $p(1-p)$ (T3, directionally confirmed in an open backward pass at +0.71). The same model exposes the diagnostic’s structural limit—the symmetry that aliases mastery

with incapacity, visible as a U-shape across 368 runs—and supplies the repair: the pairwise-contrast density $\text{PCD} = \frac{G-1}{G} \mathbb{E}[p(1-p)]$, which survives the micro-jitter that silently zeroes ZVF in any pipeline with a dense sub-reward.

Prediction enabled intervention. The `zvf-triage` callback and adaptive- G controller act on T2 live, escalating G on starvation spikes; in a four-arm open audit the controller matched the best fixed recipe’s held-out gain while dynamic sampling bought perfect telemetry at +45% rollouts with no outcome advantage. The resulting design rules are modest and, we believe, durable: gate on magnitude (PCD), alarm on existence (ZVF), escalate multiplicatively on spikes, price every intervention in rollouts, and never let a recipe label substitute for a stack report. What remains is the decisive tier—the cross-anchor PCD horse race, the hard-population controller test, and production-scale gradient attribution—which we have specified precisely so that this paper, like its companion, can be falsified rather than merely believed.

A Derivations

Throughout, a group is G conditionally i.i.d. binary rewards $R_1, \dots, R_G \sim \text{Bernoulli}(p)$ for a prompt with latent success probability p ; $K = \sum_i R_i \sim \text{Binomial}(G, p)$ and $\hat{p} = K/G$.

A.1 T1: expected ZVF indicator

The within-group sample variance is zero iff all rewards agree, i.e. $K = 0$ or $K = G$. These events are disjoint, so

$$\mathbb{E}[\mathbf{1}\{\text{degenerate}\}] = \Pr(K = 0) + \Pr(K = G) = (1-p)^G + p^G = h_G(p). \quad (27)$$

For a prompt population $\mathcal{D}$ with per-prompt success probabilities p_x , linearity of expectation gives $\mathbb{E}[\text{ZVF}] = \mathbb{E}_{x \sim \mathcal{D}}[h_G(p_x)]$. Note the Jensen gap used in Section 5.2: since h_G is convex near the endpoints and the population mixes prompts of heterogeneous difficulty, $\mathbb{E}_x[h_G(p_x)] \geq h_G(\mathbb{E}_x[p_x])$ whenever mass sits at the extremes, which is why measured run-level ZVF exceeds $h_G(\bar{p})$ in Table 7. $\square$

A.2 T2: monotonicity in G

For $p \in (0, 1)$, both $p < 1$ and $1-p < 1$, so $p^{G+1} < p^G$ and $(1-p)^{G+1} < (1-p)^G$ strictly; summing, $h_{G+1}(p) < h_G(p)$. The decay is exponential with rate set by the larger term: $h_G(p) = \max(p, 1-p)^G [1 + o(1)]$, so $-\frac{1}{G} \ln h_G(p) \rightarrow -\ln \max(p, 1-p)$. Two regimes matter for control. At $p = 1/2$, $h_G = 2^{1-G}$: each additional rollout halves the degenerate fraction (an e -folding of h_G costs $1/\ln 2 \approx 1.44$ rollouts). Near the boundary, $p = 1-\epsilon$ with $\epsilon \ll 1/G$: $h_G(p) = (1-\epsilon)^{G+\epsilon} \approx 1 - G\epsilon$, so the marginal rollout recovers only $O(\epsilon)$ of signal—escalation is exponentially effective in the interior and linearly futile at the edges, which is Rule 1 of Section 12. $\square$

A.3 T3: advantage mass

With unnormalized group-centered advantages $A_i = R_i - \hat{p}$, a group with K successes has K advantages equal to $1 - \hat{p}$ and $G - K$ equal to $-\hat{p}$. Hence

$$\frac{1}{G} \sum_{i=1}^G |A_i| = \hat{p}(1 - \hat{p}) + (1 - \hat{p})\hat{p} = 2\hat{p}(1 - \hat{p}), \quad (28)$$

and the summed advantage magnitude is $2G\hat{p}(1 - \hat{p})$, i.e. proportional to the within-group contrast. The policy-gradient norm also carries the score-function factor $\nabla_\theta \log \pi_\theta(y_i | x)$, which is policy- and sequence-dependent; T3 therefore predicts a monotone association between $\|\nabla\|$ and $\hat{p}(1 - \hat{p})$, not proportionality—this is what E1 tests ($\text{corr} = +0.71$, Section 7). Under std-normalized advantages (dividing by $\sqrt{\hat{p}(1 - \hat{p})}$), the mean absolute advantage becomes $2\sqrt{\hat{p}(1 - \hat{p})}$: normalization flattens but does not remove the $p(1-p)$ dependence, and it leaves the degenerate case (undefined $0/0$, masked in practice) untouched. $\square$

A.4 PCD expectation

The population variance of the group, $\hat{p}(1 - \hat{p})$, has expectation

$$\mathbb{E}[\hat{p}(1 - \hat{p})] = \mathbb{E}[\hat{p}] - \mathbb{E}[\hat{p}^2] = p - \left(\frac{p(1-p)}{G} + p^2 \right) = \left(1 - \frac{1}{G} \right) p(1-p) = \frac{G-1}{G} p(1-p), \quad (29)$$

using $\text{Var}(\hat{p}) = p(1-p)/G$. Averaging over prompts yields Eq. (3). Equivalently, $\hat{p}(1 - \hat{p})$ equals $\frac{1}{G^2} \cdot K(G - K)$, i.e. (up to the factor $2/G^2$) the number of discordant correct/incorrect *pairs* in the group—hence “pairwise-contrast density,” and hence the connection to the pairwise-preference reading of group-relative objectives [13, 9]: PCD counts, per unit of rollout budget, exactly the pairs from which a DPO-style contrast could be formed. $\square$

A.5 Micro-jitter proposition

Let $R'_i = R_i + \epsilon_i$ with $\epsilon_i \stackrel{\text{iid}}{\sim} \text{U}(0, \delta)$, $\delta = 10^{-4}$. (i) *ZVF collapses*: the jittered rewards are almost surely pairwise distinct (ties have Lebesgue measure zero), so every group has nonzero sample variance and the tie-based ZVF indicator is 0 for all groups. (ii) *PCD is stable*: in expectation the within-group variance shifts by exactly $\text{Var}(\epsilon) = \delta^2/12 \leq 8.4 \times 10^{-10}$; per-group sample cross terms between R and ϵ are $O(\delta)$ in standard deviation but zero-mean, so they average out over a batch. Empirically (scripts/pcd_vs_zvf.py, 600 groups): ZVF $0.1583 \rightarrow 0.0000$, PCD $0.153802 \rightarrow 0.153802$ (shift below 10^{-6}). The practical corollary is stated in Section 6.3: any dense sub-reward of amplitude $\geq$ machine-distinguishable scale performs this transformation implicitly, so tie-based ZVF must be computed on the binary verifier output alone, or replaced by PCD. $\square$

A.6 Provenance of the empirical tables

Table 6 and Table 8 aggregate the 368-run audit in W&B project zvf-audit (790-run corpus; entity arvindr4-pes-university); per-model means are computed over all completed runs of the named model, and the G -grid restricts to runs where only G varies. Table 7 is experiments/results/groupsize_zvf_sweep.tsv (12 runs: 4 values of $G \times 3$ seeds, full per-run records in the adjacent .json). The binned shape and jitter numbers are experiments/results/pcd_vs_zvf_shape.tsv and pcd_vs_zvf_summary.tsv, regenerable via scripts/pcd_vs_zvf.py from the released per-group reward tensors. E1 and E3 live in W&B project zvf-colab-experiments (run families E1_grad_signal, E3_open_audit).

References

- [1] Rishabh Agarwal, Max Schwarzer, Pablo Samuel Castro, Aaron C. Courville, and Marc Bellemare. Deep reinforcement learning at the edge of the statistical precipice. In *Advances in Neural Information Processing Systems*, volume 34, pages 29304–29320, 2021. doi: 10.48550/arXiv.2108.13264. arXiv:2108.13264; supplementary materials at <https://agarwl.github.io/rliable/>.
- [2] AlphaProof team. Olympiad-level formal mathematical reasoning with reinforcement learning. *Nature*, 651:607–613, 2026. doi: 10.1038/s41586-025-09833-y. Google DeepMind; published online 12 November 2025.
- [3] Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.
- [4] Cédric Colas, Olivier Sigaud, and Pierre-Yves Oudeyer. A hitchhiker’s guide to statistical comparisons of reinforcement learning algorithms. *arXiv preprint*, 2019. doi: 10.48550/arXiv.1904.06979. arXiv:1904.06979.
- [5] DeepSeek-AI, Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Ruoyu Zhang, Runxin Xu, Qihao Zhu, Shirong Ma, Peiyi Wang, et al. DeepSeek-R1: Incentivizing reasoning capability in LLMs via reinforcement learning. *arXiv preprint*, 2025. doi: 10.48550/arXiv.2501.12948. arXiv:2501.12948.

- [6] Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations*, 2022. URL <https://openreview.net/forum?id=nZeVKeeFYf9>.
- [7] Thanh-Long V. Le, Myeongho Jeon, Kim Vu, Viet Lai, and Eunho Yang. No prompt left behind: Exploiting zero-variance prompts in LLM reinforcement learning via entropy-guided advantage shaping. *arXiv preprint*, September 2025. Also referred to as AERO / RL-ZVP in the literature; arXiv:2509.21880.
- [8] Zichen Liu, Changyu Chen, Wenjun Li, Penghui Qi, Tianyu Pang, Chao Du, Wee Sun Lee, and Min Lin. Understanding R1-Zero-Like training: A critical perspective. *arXiv preprint*, 2025. doi: 10.48550/arXiv.2503.20783. arXiv:2503.20783; introduces Dr. GRPO.
- [9] Rafael Rafailov, Archit Sharma, Eric Mitchell, Stefano Ermon, Christopher D. Manning, and Chelsea Finn. Direct preference optimization: Your language model is secretly a reward model. In *Advances in Neural Information Processing Systems*, volume 36, 2023. doi: 10.48550/arXiv.2305.18290. arXiv:2305.18290; NeurIPS 2023 (proceedings vol. 36).
- [10] Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Mingchuan Zhang, Y.K. Li, Y. Wu, and Daya Guo. DeepSeekMath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint*, 2024. doi: 10.48550/arXiv.2402.03300. arXiv:2402.03300.
- [11] DiJia Su, Sainbayar Sukhbaatar, Michael Rabbat, Yuandong Tian, and Qingqing Zheng. Du-alformer: Controllable fast and slow thinking by learning with randomized reasoning traces. *arXiv preprint arXiv:2410.09918*, 2024.
- [12] Zelin Tan, Hejia Geng, Xiaohang Yu, Mulei Zhang, Guancheng Wan, Yifan Zhou, Qiang He, Xiangyuan Xue, Heng Zhou, Yutao Fan, Zhongzhi Li, Zaibin Zhang, Guibin Zhang, Chen Zhang, Zhenfei Yin, Philip Torr, and Lei Bai. Scaling behaviors of LLM reinforcement learning post-training: An empirical study in mathematical reasoning. *arXiv preprint arXiv:2509.25300*, 2025.
- [13] Yihong Wu, Liheng Ma, Lei Ding, Muzhi Li, Xinyu Wang, Kejia Chen, Zhan Su, Zhanguang Zhang, Chenyang Huang, Yingxue Zhang, Mark Coates, and Jian-Yun Nie. It takes two: Your GRPO is secretly DPO. *arXiv preprint*, 2025. doi: 10.48550/arXiv.2510.00977. arXiv:2510.00977.
- [14] Qiying Yu, Zheng Zhang, Ruofei Zhu, Yufeng Yuan, Xiaochen Zuo, Yu Yue, Weinan Dai, Tiantian Fan, Gaohong Liu, Lingjun Liu, et al. DAPO: An open-source LLM reinforcement learning system at scale. *arXiv preprint*, 2025. doi: 10.48550/arXiv.2503.14476. arXiv:2503.14476.
- [15] Chujie Zheng, Shixuan Liu, Mingze Li, Xiong-Hui Chen, Bowen Yu, Chang Gao, Kai Dang, Yuqiong Liu, Rui Men, An Yang, Jingren Zhou, Jianwei Zhou, and Junyang Lin. Group sequence policy optimization. *arXiv preprint*, 2025. doi: 10.48550/arXiv.2507.18071. arXiv:2507.18071; GSPO.
- [16] Haizhong Zheng, Yang Zhou, Brian R. Bartoldson, Bhavya Kailkhura, Fan Lai, Jiawei Zhao, and Beidi Chen. Act only when it pays: Efficient reinforcement learning for LLM reasoning via selective rollouts. *arXiv preprint*, 2025. doi: 10.48550/arXiv.2506.02177. arXiv:2506.02177; GRESO.