

FROZEN ACTOR — one policy, shared by all fourteen lanes
Qwen/Qwen3.6-35B-A3B @ 995ad96e + LoRA campaign adapter @ withheld revision
served in bfloat16 as campaign-actor-bf16

1 → 14 lanes

Lane	Original contract (orig)	Replacement scope (repl)	Terminal state
E1	<i>SWE-bench Pro</i> : 2/731 = 0.27%	<i>SWE-bench Multilingual</i> : 1/190 = 0.53% [‡]	★ SCORED 09-27
E2	<i>FrontierSWE</i> : 1/17; replay 0.8628	<i>CORE-Bench</i> (hard): 27/45 = 0.600	★ SCORED 09-27
E3	<i>SDAB</i> : closed external, no figure	—	× EXTERNAL
E4	<i>BankerToolBench</i> : 1/100; recovery 0.3115	—	○ PARTIAL
E5	<i>APEX-Agents</i> : 7/480; mean 0.0505	<i>Tau3</i> : 8/97 = 0.0825 pass [^] 1	★ SCORED 09-27
E6	<i>WebBench</i> : no figure	<i>WebArena</i> : 90/812 = 0.111 (lower bd.)	★ SCORED 09-27
E7	<i>BinaryAudit</i> : 1/46 attempted; no grade	—	× EXTERNAL
E8	<i>LifeSciBench</i> : closed external	<i>LAB-Bench</i> : 450/1967 = 0.2288	● COMPLETE
E9	<i>MLE-bench</i> : 40/75 graded; score null	<i>MLDevBench</i> : 10/34 = 0.294	★ SCORED 09-27
E10	<i>AgentHarm</i> : closed external	<i>AgentDojo</i> : 88/97 = 0.9072	● COMPLETE
E11	<i>VerilogEval</i> : 129/312 = 41.35%	—	● COMPLETE
E12	<i>AppBench</i> : closed external, no figure	—	× EXTERNAL
E13	<i>OpenReward Games</i> : closed external	<i>BALROG</i> : 26.1% progression, 255 ep.	★ SCORED 09-27
E14	<i>FrontierMath</i> : closed external	<i>Omni-MATH</i> : 2271/4428 = 51.29% [‡]	● COMPLETE [‡]

Bold = the lane’s figure from the original campaign ledger (at most one per lane, tagged with its scope). Pass@1 for E1 and E11; E8 is the LAB-Bench public split, E10 AgentDojo benign utility. [‡]E14 closes COMPLETE with a terminal note: its 2 unjudged rows (truncated judge output) count as failures here; the official scorer reports 51.31% on the 4426 judged rows. ● complete ★ replacement scope scored on 2026-09-27 (plain figure: its own receipt, never pooled with the orig figure) ○ partial × blocked (closed external). [‡]This run’s 190 tasks; the earlier 110 (4 resolved) are a separate contract, not pooled. WebArena: 108 unjudged tasks count as failures.

NATIVE EVALUATOR

per lane: the suite’s own grader at a revision pinned in the lane’s receipt; no shared grader, no cross-lane normalisation

RECEIPT

one immutable, hash-pinned record per lane: coverage, terminal state, claim boundary, run-manifest hash; it may close with no figure

REPORTED FIGURE

at most one per receipt, tagged with its scope; a lane without a grade reports none rather than an adjacent suite’s number

14 lanes → 8 ledger figures + 6 replacement figures (2026-09-27)

3 lanes with no figure · no receipt carries two

RULE 1 — one actor, no per-lane tuning

One frozen checkpoint serves all fourteen lanes: no lane-specific fine-tune, no prompting beyond the suite’s own harness, no tuning on a lane’s evaluation set. Suite differences reflect one fixed policy meeting fourteen interfaces.

RULE 2 — orig and repl figures are never pooled

A replacement scope is a different suite, interface and often grader. Its figure is tagged with its scope and never averaged, summed or compared with an original-contract figure.