

One configuration — the canonical spec

Model Qwen/Qwen3-8B · **task** GSM8K, verified reward · GRPO · 30 steps

Group $G = 8$ · batch 2×8 · lr $1e-6$ · LoRA $r = 16$, $\alpha = 32$, all-linear

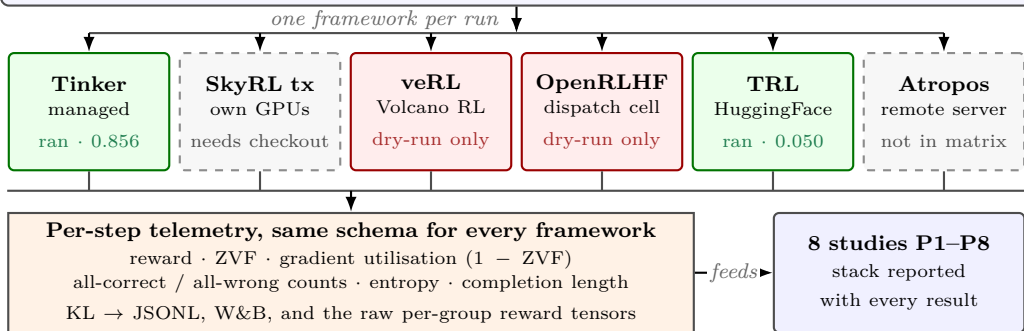
GRPO $\beta = 0.0$ · $\varepsilon = 0.2$ · max 1024 tokens · seed 211

Unified Tinker-style API and adapter layer

UnifiedLauncher — `--framework` picks the trainer, `--backend` the site.

Adapter layer — `Backend.plan()` → `LaunchPlan`, returned unexecuted.

Dispatch matrix — 5 frameworks × 6 compute backends, validated with no GPU or key.



Six RL frameworks, five of them in the dispatch matrix; **two executed** (Tinker, TRL).

Box outline: green solid = completed run; red = dry-run placeholder; grey dashed = not executed.

Status and reward figures: `framework_comparison.json`, last-10 mean training reward, Qwen/Qwen3-8B.

The two dry-run rows are launch-plan artefacts, not measurements, and are excluded from pooling.