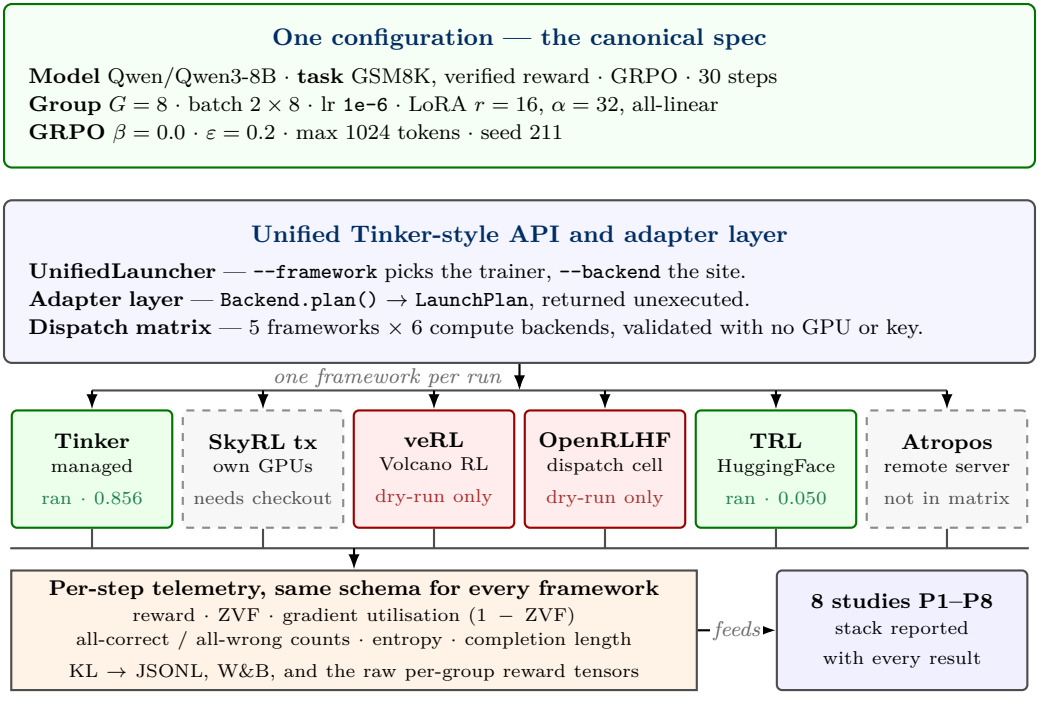

Six RL frameworks, five of them in the dispatch matrix; **two executed** (Tinker, TRL).
Box outline: green solid = completed run; red = dry-run placeholder; grey dashed = not executed.
Status and reward figures: `framework_comparison.json`, last-10 reward; Tinker: Qwen3-8B-Base, TRL: Qwen3-8B.
The two dry-run rows are launch-plan artefacts, not measurements, and are excluded from pooling.