FROZEN ACTOR — one policy, shared by all fourteen lanes

`Qwen/Qwen3.6-35B-A3B @ 995ad96e` + LoRA `campaign adapter @ withheld revision`

served in bfloat16 as `campaign-actor-bf16`

1 → 14 *lanes*

Lane	Original contract (orig)	Replacement scope (repl)	Terminal state
E1	*SWE-bench Pro*: **2/731 = 0.27%**	*SWE-bench Multilingual*: 1/190 = 0.53%‡	⋆ SCORED 09-27
E2	*FrontierSWE*: **1/17; replay 0.8628**	*CORE-Bench* (hard): 27/45 = 0.600	⋆ SCORED 09-27
E3	*SDAB*: closed external, no figure	—	× EXTERNAL
E4	*BankerToolBench*: **1/100; recovery 0.3115**	—	○ PARTIAL
E5	*APEX-Agents*: **7/480; mean 0.0505**	*Tau3*: 8/97 = 0.0825 pass^1	⋆ SCORED 09-27
E6	*WebBench*: no figure	*WebArena*: 90/812 = 0.111 (lower bd.)	⋆ SCORED 09-27
E7	*BinaryAudit*: 1/46 attempted; no grade	—	× EXTERNAL
E8	*LifeSciBench*: closed external	*LAB-Bench*: **450/1967 = 0.2288**	● COMPLETE
E9	*MLE-bench*: 40/75 graded; score null	*MLDevBench*: 10/34 = 0.294	⋆ SCORED 09-27
E10	*AgentHarm*: closed external	*AgentDojo*: **88/97 = 0.9072**	● COMPLETE
E11	*VerilogEval*: **129/312 = 41.35%**	—	● COMPLETE
E12	*AppBench*: closed external, no figure	—	× EXTERNAL
E13	*OpenReward Games*: closed external	*BALROG*: 26.1% progression, 255 ep.	⋆ SCORED 09-27
E14	*FrontierMath*: closed external	*Omni-MATH*: **2271/4426 = 51.31%**†	● COMPLETE†

Bold = the lane's figure from the original campaign ledger (at most one per lane, tagged with its scope). Pass@1 for E1 and E11; E8 is the LAB-Bench public split, E10 AgentDojo benign utility. †E14 closes COMPLETE with a terminal note: native accuracy uses 4426 judged rows. The separate all-4428 diagnostic counts 2 unjudged rows as failures: 2271/4428 = 51.29%. ● complete ⋆ replacement scope scored on 2026-09-27 (plain figure: its own receipt, never pooled with the orig figure) ○ partial × blocked (closed external). ‡This run's 190 tasks; the earlier 110 (4 resolved) are a separate contract, not pooled. WebArena: 108 unjudged tasks count as failures.

NATIVE EVALUATOR
per lane: the suite's own grader at a revision pinned in the lane's receipt; no shared grader, no cross-lane normalisation

→

RECEIPT
one immutable, hash-pinned record per lane: coverage, terminal state, claim boundary, run-manifest hash; it may close with no figure

→

REPORTED FIGURE
at most one per receipt, tagged with its scope; a lane without a grade reports none rather than an adjacent suite's number

14 lanes → 8 ledger figures + 6 replacement figures (2026-09-27)
3 lanes with no figure · no receipt carries two

RULE 1 — one actor, no per-lane tuning
One frozen checkpoint serves all fourteen lanes: no lane-specific fine-tune, no prompting beyond the suite's own harness, no tuning on a lane's evaluation set. Suite differences reflect one fixed policy meeting fourteen interfaces.

RULE 2 — orig and repl figures are never pooled
A replacement scope is a different suite, interface and often grader. Its figure is tagged with its scope and never averaged, summed or compared with an original-contract figure.