Lane terminal state, E1–E14 (ledger of 2026-09-19, updated with the runs of 2026-09-27)

E1 SWE-bench Multilingual ⋆ SCORED 1/190†	**E2** CORE-Bench (hard) ⋆ SCORED 27/45	**E3** SDAB × CLOSED EXTERNAL	**E4** Banker-ToolBench △ CLOSED PARTIAL	**E5** Tau3 ⋆ SCORED 8/97	**E6** WebArena ⋆ SCORED 90/812‡	**E7** BinaryAudit × CLOSED EXTERNAL
E8 LAB-Bench (public split) • COMPLETE	**E9** MLDevBench ⋆ SCORED 10/34	**E10** AgentDojo (benign) • COMPLETE	**E11** VerilogEval (native) • COMPLETE	**E12** AppBench × CLOSED EXTERNAL	**E13** BALROG ⋆ SCORED 26.1% (255 ep.)	**E14** Omni-MATH* (replacement) • COMPLETE

• **COMPLETE** (4) — full coverage, no open gate: E8, E10, E11; E14* terminal-complete

⋆ **SCORED** (6) — double border: replacement scope scored on 2026-09-27: E1, E2, E5, E6, E9, E13

△ **PARTIAL** (1) — closed with partial evidence: E4 (1/100 tasks)

× **BLOCKED** (3) — hatched: whole lane externally closed, external payload unavailable: E3, E7, E12 (original contracts of E8, E10, E13, E14 are also blocked; those lanes report replacement scopes)

Source: `finish/Pending_Experiments.md` terminal-state table; `E1_E14_FINAL_RESULTS_2026-09-19.md`; `finish_pending_2026-09-27/<lane>/result.json`.

† This run's 190 tasks; the earlier 110 (4 resolved) are a separate contract and are not pooled.

‡ Lower bound: 108 tasks could not be judged and count as failures.